

# Introduction to group equivariant deep learning

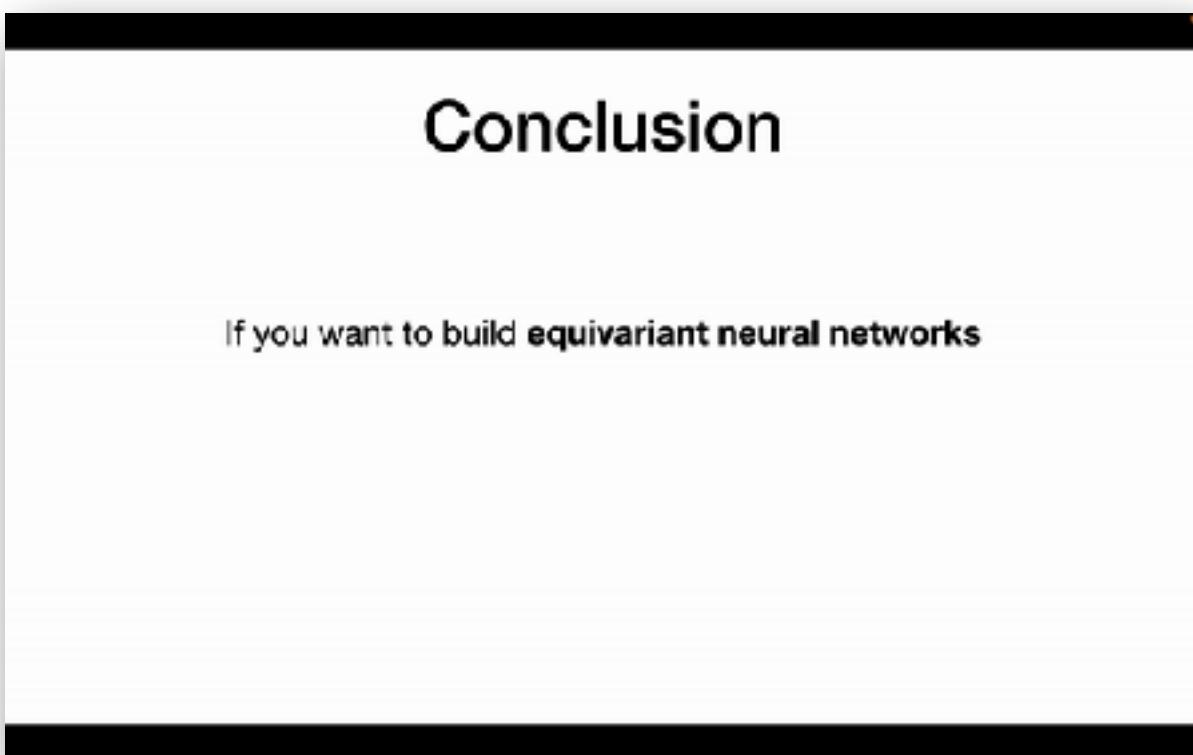
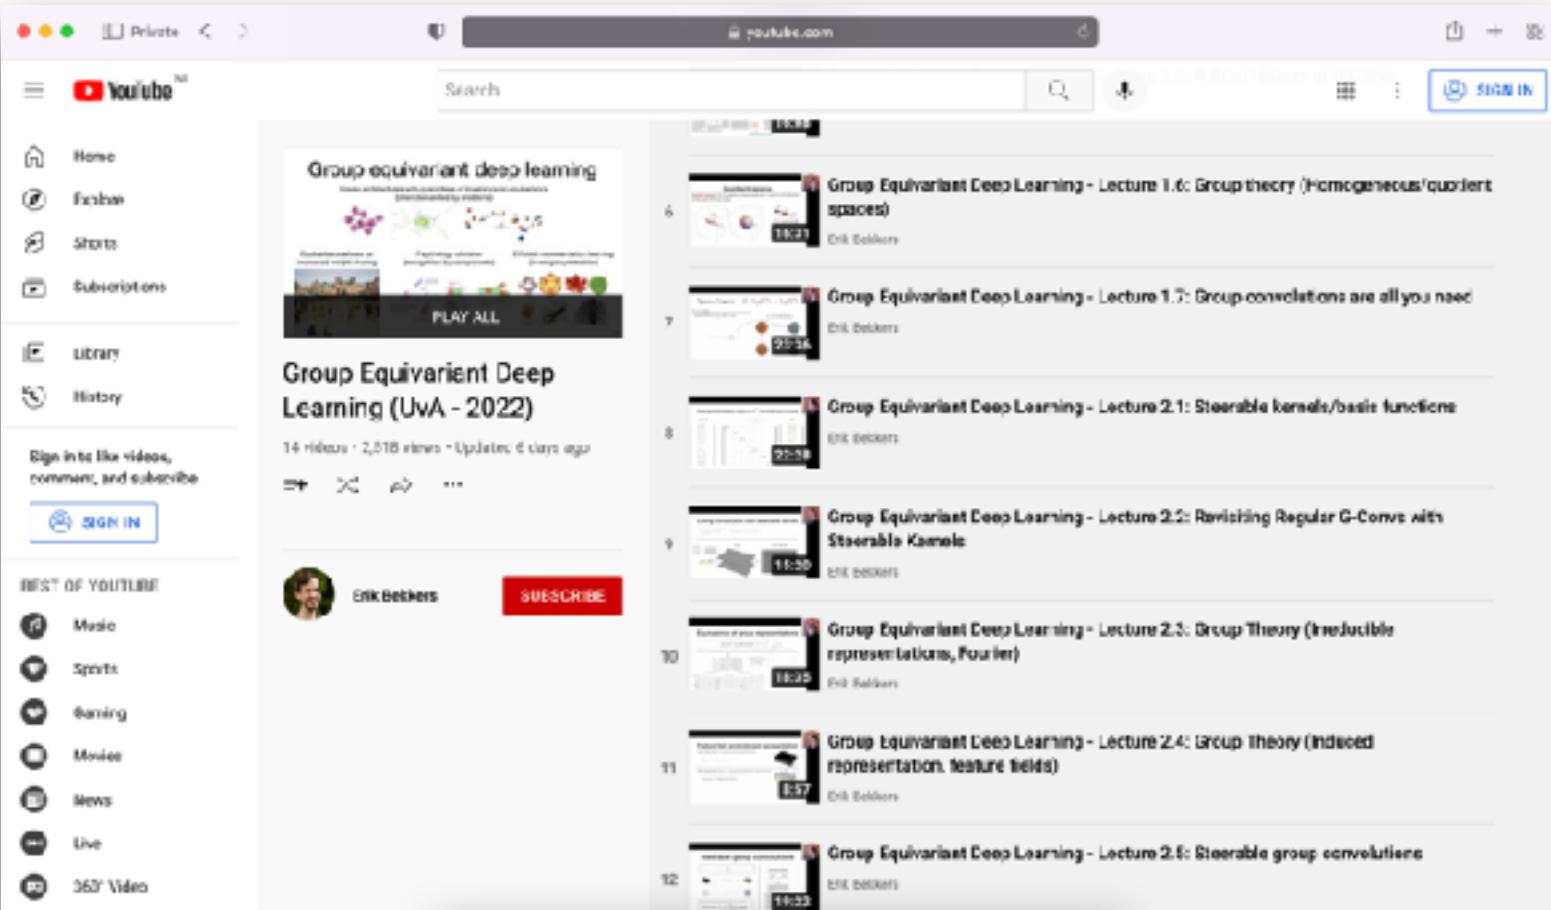
Erik Bekkers

[e.j.bekkers@uva.nl](mailto:e.j.bekkers@uva.nl)

Amsterdam Machine Learning Lab, University of Amsterdam

# UvA course on group equivariant deep learning (<https://uvagedl.github.io>)

## Youtube playlist



## Tutorial notebooks

[# UvA Notebooks](#)

[Search docs](#)

**GUIDE**

- Guide 1: Working with the UvA cluster
- Guide 2: Research projects with PyTorch
- Guide 3: Debugging in PyTorch

**DEEP LEARNING 1**

- Tutorial 1: Introduction to PyTorch
- Tutorial 2: Activation Functions
- Tutorial 3: Optimizers and Initialization
- Tutorial 4: Attention, Multi-Head Attention
- Tutorial 5: Transformers and Multi-Head Attention
- Tutorial 6: Deep Generative Models
- Tutorial 7: Deep Autoencoders
- Tutorial 8: Adversarial Attacks
- Tutorial 9: Generative Image Modeling

**GDL - Regular Group Convolutions**

Filled notebook: [View On GitHub](#) [Open in Colab](#)

Pre-trained models: [View On GitHub](#) [Open in Colab](#)

Authors: David Krueger, Galo Valdés Cesa

### 0. Introduction

In this notebook, we will be implementing regular group convolutional networks from scratch, only making use of [PyTorch](#), primarily. The goal is to get familiar with the practical considerations to take into account when actually implementing these convolutional networks. Questions and feedback may be forwarded to David Krueger: [d.krueger@uva.nl](mailto:d.krueger@uva.nl).

If you'd like a refresher of the lecture, here we give a brief overview of the operations we are going to work with / implement. These will be treated more extensively below. For simplicity of notation, here we assume each CNN layer consists of only a single channel.

#### 0.1 Brief recap on CNNs

Conventions: CNNs make use of the convolution operator, here defined over  $\mathbb{R}^2$  for a signal  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$  and a kernel  $k: \mathbb{R}^2 \rightarrow \mathbb{R}$ :  $x \in \mathbb{R}^2$ ,

$$(f * k)(x) = \int_{\mathbb{R}^2} f(s)k(x-s)ds,$$

As we can see, the convolution operation comes down to an inner product of the function  $f$  and a shifted kernel  $k$ .

Sidenote: In reality CNNs implement a discretized version of this operation:

$$(f * k)(x) = \sum_{s \in \mathbb{Z}^2} f(s)k(x-s)\Delta s,$$

[# UvA Notebooks](#)

[Search docs](#)

**GUIDE**

- Guide 1: Working with the UvA cluster
- Guide 2: Research projects with PyTorch
- Guide 3: Debugging in PyTorch

**DEEP LEARNING 1**

- Tutorial 1: Introduction to PyTorch
- Tutorial 2: Activation Functions
- Tutorial 3: Optimizers and Initialization
- Tutorial 4: Attention, Multi-Head Attention
- Tutorial 5: Transformers and Multi-Head Attention
- Tutorial 6: Deep Generative Models
- Tutorial 7: Graph Neural Networks
- Tutorial 8: Deep Energy-Based Generative Models
- Tutorial 9: Deep Autoencoders
- Tutorial 10: Adversarial Attacks
- Tutorial 11: Normalizing Flows for Image modeling
- Tutorial 12: Autoregressive Image Modeling

**GDL - Steerable CNNs**

Filled notebook: [View On GitHub](#) [Open in Colab](#)

Empty notebook: [View On GitHub](#) [Open in Colab](#)

Authors: Galo Valdés Cesa

### GDL - Steerable CNNs

During the lectures, you have learnt that the symmetries of a machine learning task can be modelled with groups. In the previous tutorial, you have also studied the framework of Group Convolutional Neural Networks (GCNNs), which describes a neural architecture design equivariant to general groups.

The feature maps of a GCNN are functions over the elements of the group. A naive implementation of group convolution requires computing and storing a response for each group element. For this reason, the GCNN framework is not particularly convenient to implement networks equivariant to groups with infinite elements.

Steerable CNNs are a more general framework which solves this issue. The key idea is that, instead of storing the value of a feature map on each group element, the model stores the Fourier transform of this feature map, up to a finite number of frequencies.

In this tutorial, we will first introduce some representation theory and Fourier theory (from quantum mechanics and linear algebra) and then we will explore how this idea is used in practice to implement Steerable CNNs.

#### Prerequisite Knowledge

Throughout this tutorial, we will assume you are already familiar with some concepts of group theory, such as groups, group actions (in particular on functions), semi-direct product and order of a group, as well as basic linear algebra.

We start by importing the necessary packages. You can run the following command to install all the requirements:

[# UvA Notebooks](#)

[Search docs](#)

**Introduction**

The goal of this tutorial is to get you familiar with the ways in which positional information is used in graph neural networks. We will firstly look at superpixel CNNs, which group pixels by similarity in order to create a graph. Then, we will apply steerable methods to MD17, a 3D dataset of molecular relaxation trajectories.

This tutorial was created by Petri van der Linden and Ruud Hesselink. If you would like to know more about these methods, feel free to contact us at [p.vanderlinden@uva.nl](mailto:p.vanderlinden@uva.nl) or [r.hesselink@uva.nl](mailto:r.hesselink@uva.nl).

Below you can find several papers that we recommend, if you're interested in some more reading.

[Neural Message Passing for Quantum Chemistry](#)  
[Directional Message Passing for Molecular Graphs](#)  
[Equivariant Graph Neural Networks for Data-Efficient and Accurate Interatomic Potentials](#)

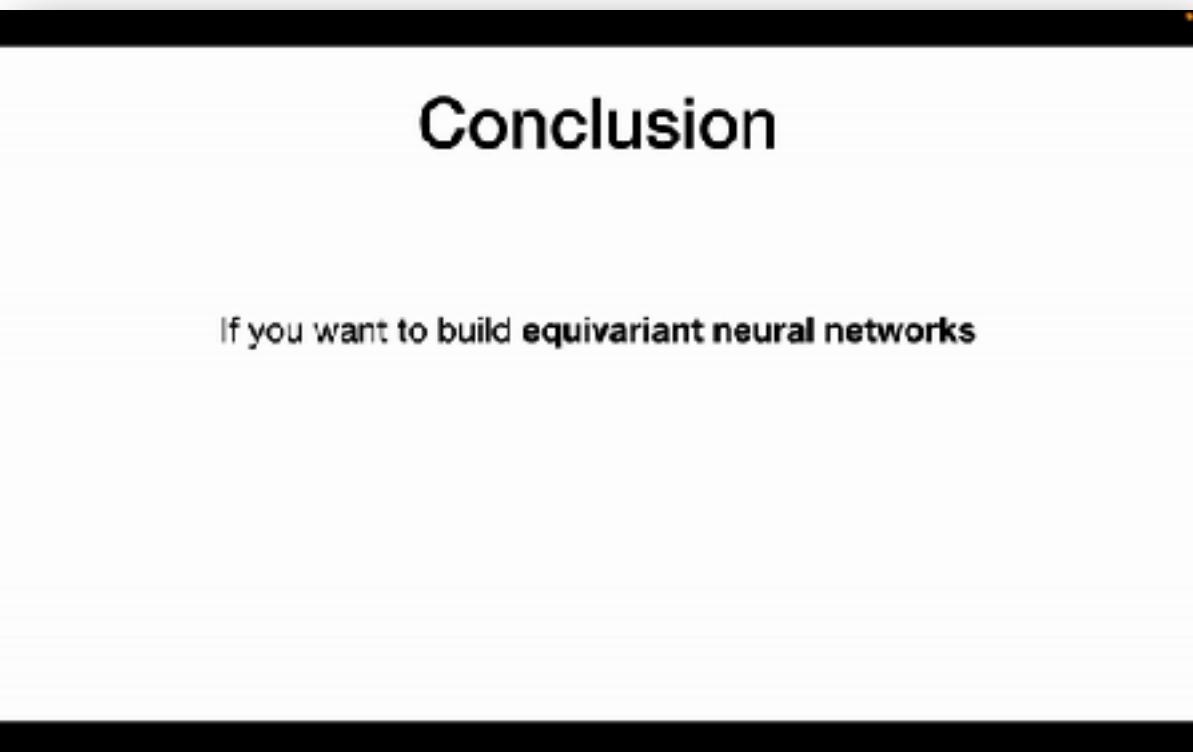
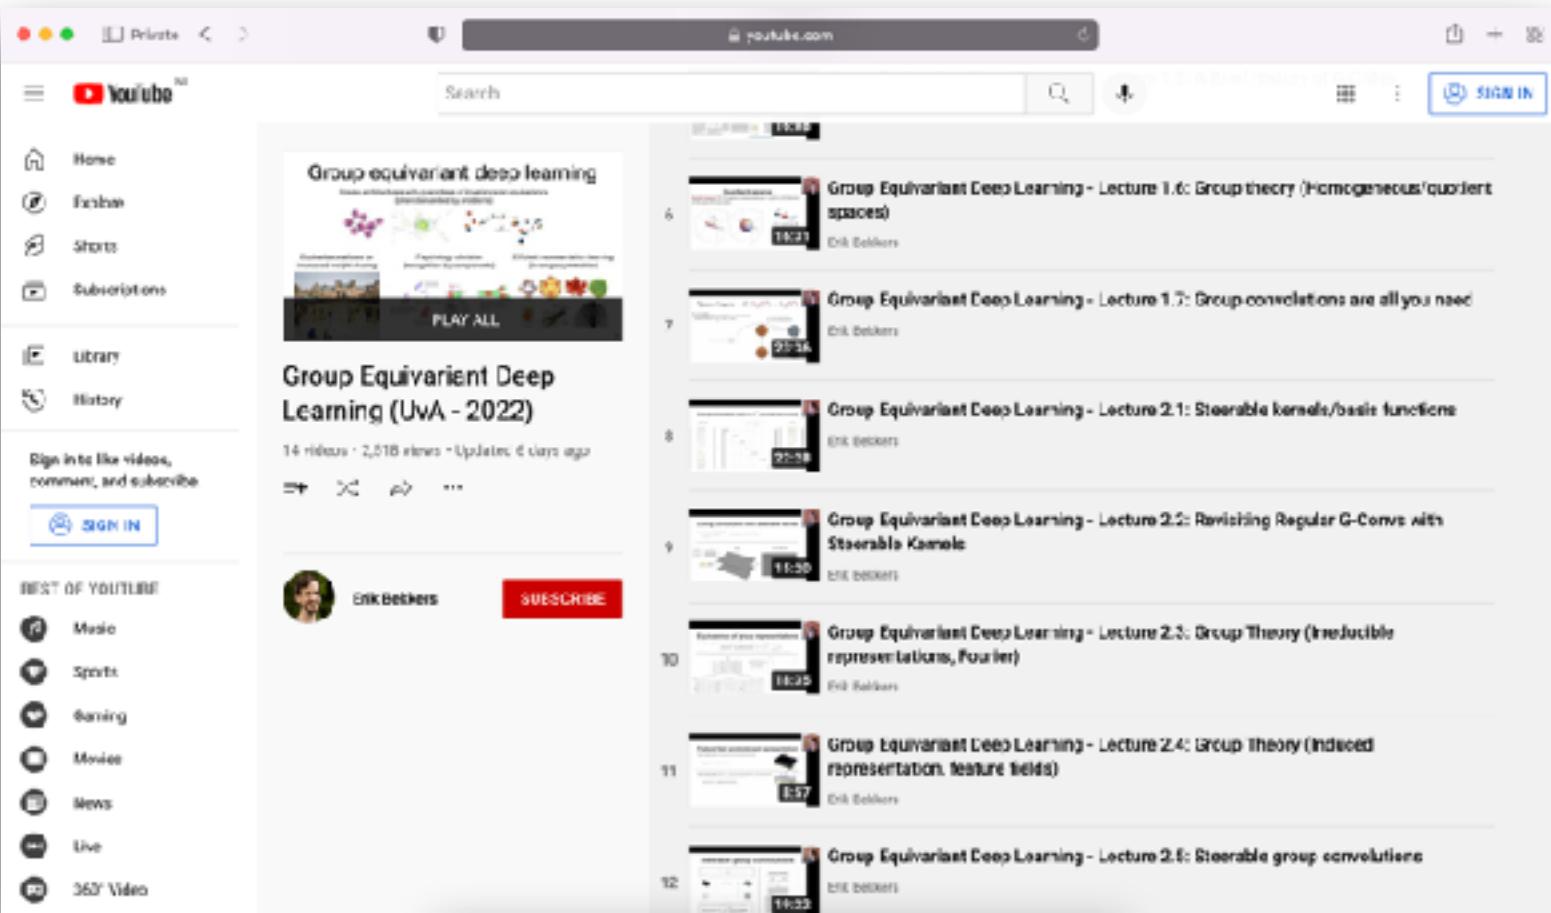
### 0. Graphs as objects embedded in Euclidean Space

In this tutorial we are going to look at graph convolution methods that can act on graphs embedded in some Euclidean space, meaning that the graph represents a some multidimensional structure, e.g. we will refer to it as a **Manifold Graph**.

Let us first consider a generic graph consisting of a node set  $V$  containing nodes  $v_i$  and a connectivity given by an edge set  $E$  consisting of

# UvA course on group equivariant deep learning (<https://uvagedl.github.io>)

## Youtube playlist



## Tutorial notebooks

[# UvA Notebooks](#)

[Search docs](#)

**GUIDE**

- Guide 1: Working with the UvA cluster
- Guide 2: Research projects with PyTorch
- Guide 3: Debugging in PyTorch

**DEEP LEARNING 1**

- Tutorial 1: Introduction to PyTorch
- Tutorial 2: Activation Functions
- Tutorial 3: Optimizers and Initialization
- Tutorial 4: Attention, Multi-Head Attention
- Tutorial 5: Transformers and Multi-Head Attention
- Tutorial 6: Deep Neural Networks
- Tutorial 7: Deep Energy-Based Generative Models
- Tutorial 8: Deep Autoencoders
- Tutorial 9: Adversarial attacks
- Tutorial 10: Autoregressive Image Modeling

**GDL - Regular Group Convolutions**

Filled notebook: [View On GitHub](#) [Open in Colab](#)

Pre-trained models: [View On GitHub](#) [Open in Colab](#)

Authors: David Krueger, Galo Valdés Cesa

### 0. Introduction

In this notebook, we will be implementing regular group convolutional networks from scratch, only making use of [PyTorch](#) primitives. The goal is to get familiar with the practical considerations to take into account when actually implementing these convolutional networks. Questions and feedback may be forwarded to David Krueger: [d.krueger@uva.nl](mailto:d.krueger@uva.nl).

If you'd like a refresher of the lecture, here we give a brief overview of the operations we are going to work with / implement. These will be treated more extensively below. For simplicity of notation, here we assume each CNN layer consists of only a single channel.

#### 0.1 Brief recap on CNNs

Conventions: CNNs make use of the convolution operator, here defined over  $\mathbb{R}^2$  for a signal  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$  and a kernel  $k: \mathbb{R}^2 \rightarrow \mathbb{R}$ :  $x \in \mathbb{R}^2$ ,

$$(f * k)(x) = \int_{\mathbb{R}^2} f(s)k(x-s)ds,$$

As we can see, the convolution operation comes down to an inner product of the function  $f$  and a shifted kernel  $k$ .

Sidenote: In reality CNNs implement a discretized version of this operation:

$$(f * k)(x) = \sum_{s \in \mathbb{Z}^2} f(s)k(x-s)\Delta s,$$

[# UvA Notebooks](#)

[Search docs](#)

**GUIDE**

- Guide 1: Working with the UvA cluster
- Guide 2: Research projects with PyTorch
- Guide 3: Debugging in PyTorch

**DEEP LEARNING 1**

- Tutorial 1: Introduction to PyTorch
- Tutorial 2: Activation Functions
- Tutorial 3: Optimizers and Initialization
- Tutorial 4: Attention, Multi-Head Attention
- Tutorial 5: Transformers and Multi-Head Attention
- Tutorial 6: Deep Neural Networks
- Tutorial 7: Deep Energy-Based Generative Models
- Tutorial 8: Deep Autoencoders
- Tutorial 9: Adversarial attacks
- Tutorial 10: Autoregressive Image Modeling

**GDL - Steerable CNNs**

Filled notebook: [View On GitHub](#) [Open in Colab](#)

Empty notebook: [View On GitHub](#) [Open in Colab](#)

Authors: Galo Valdés Cesa

### GDL - Steerable CNNs

During the lectures, you have learnt that the symmetries of a machine learning task can be modelled with groups. In the previous tutorial, you have also studied the framework of Group Convolutional Neural Networks (GCNNs), which describes a neural architecture design equivariant to general groups.

The feature maps of a GCNN are functions over the elements of the group. A naive implementation of group convolution requires computing and storing a response for each group element. For this reason, the GCNN framework is not particularly convenient to implement networks equivariant to groups with infinite elements.

Steerable CNNs are a more general framework which solves this issue. The key idea is that, instead of storing the value of a feature map on each group element, the model stores the Fourier transform of this feature map, up to a finite number of frequencies.

In this tutorial, we will first introduce some representation theory and Fourier theory (from quantum mechanics and linear algebra) and then we will explore how this idea is used in practice to implement Steerable CNNs.

#### Prerequisite Knowledge

Throughout this tutorial, we will assume you are already familiar with some concepts of group theory, such as groups, group actions (in particular on functions), semi-direct product and order of a group, as well as basic linear algebra.

We start by importing the necessary packages. You can run the following command to install all the requirements:

[# UvA Notebooks](#)

[Search](#)

**Introduction**

The goal of this tutorial is to get you familiar with the ways in which positional information is used in graph neural networks. We will firstly look at superpixel CNNs, which group pixels by similarity in order to create a graph. Then, we will apply steerable methods to MD17, a 3D dataset of molecular relaxation trajectories.

This tutorial was created by Petri van der Linden and Ruud Hesselsink. If you would like to know more about these methods, feel free to contact us at [p.vanderlinden@uva.nl](mailto:p.vanderlinden@uva.nl) or [r.hesselsink@uva.nl](mailto:r.hesselsink@uva.nl).

Below you can find several papers that we recommend, if you're interested in some more reading.

- [Neural Message Passing for Quantum Chemistry](#)
- [Directional Message Passing for Molecular Graphs](#)
- [Equivariant Graph Neural Networks for Data-Efficient and Accurate Interatomic Potentials](#)

### 0. Graphs as objects embedded in Euclidean Space

In this tutorial we are going to look at graph convolution methods that can act on graphs embedded in some Euclidean space, meaning that the graph represents a some multidimensional structure, e.g. we will refer to it as a **Manifold Graph**.

Let us first consider a generic graph consisting of a node set  $V$  containing nodes  $v_i$  and a connectivity given by an edge set  $E$  consisting of

**1. Motivation**

**2. Pattern matching using group theory**

**3. Group convolutions**

**4. Example**

**5. G-convs are all you need!**

**6. Steerable group convolutions**

**7. Feature fields and escnn library**

**8. Equivariant tensor product layers**

**9. Equivariant graph NNs**

**1. Motivation**

**2. Pattern matching using group theory**

**3. Group convolutions**

**4. Example**

**5. G-convs are all you need!**

**6. Steerable group convolutions**

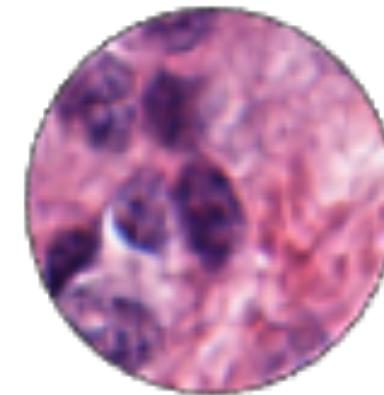
**7. Feature fields and escnn library**

**8. Equivariant tensor product layers**

**9. Equivariant graph NNs**

# Geometric guarantees (invariance)

*Example: Detection of pathological cells*



# Geometric guarantees (invariance)

*Example: Detection of pathological cells*

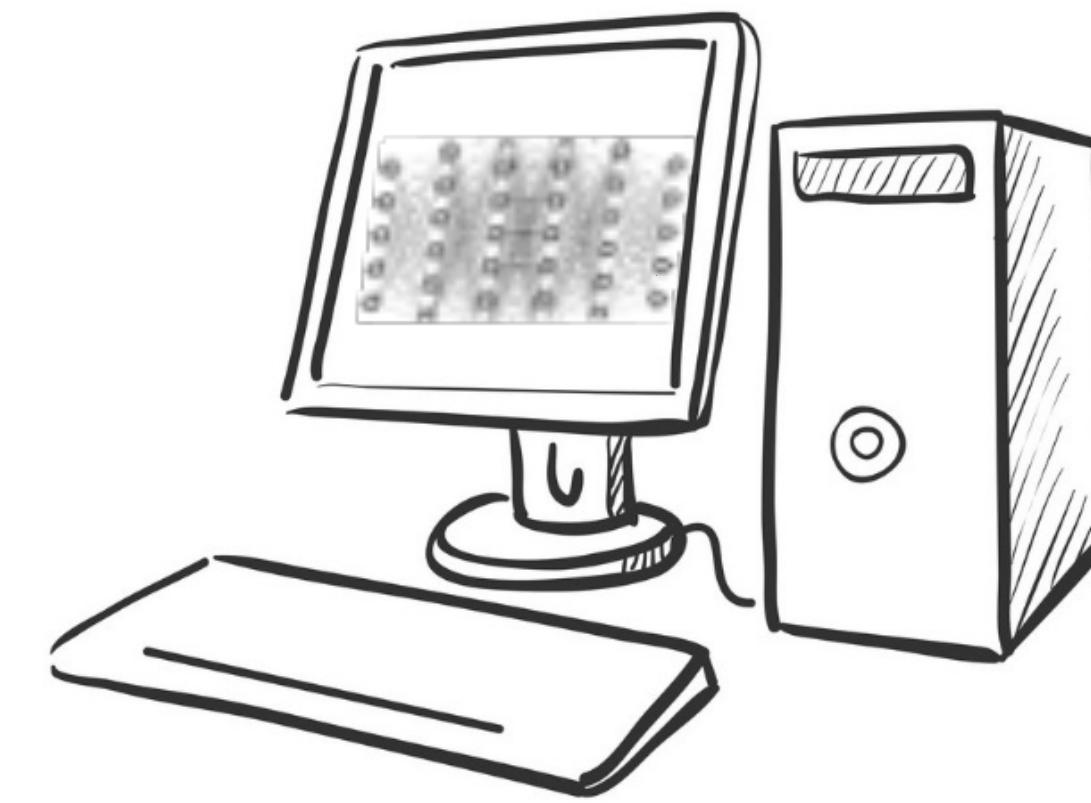


**Healthy**

?

# Geometric guarantees (invariance)

*Example: Detection of pathological cells*



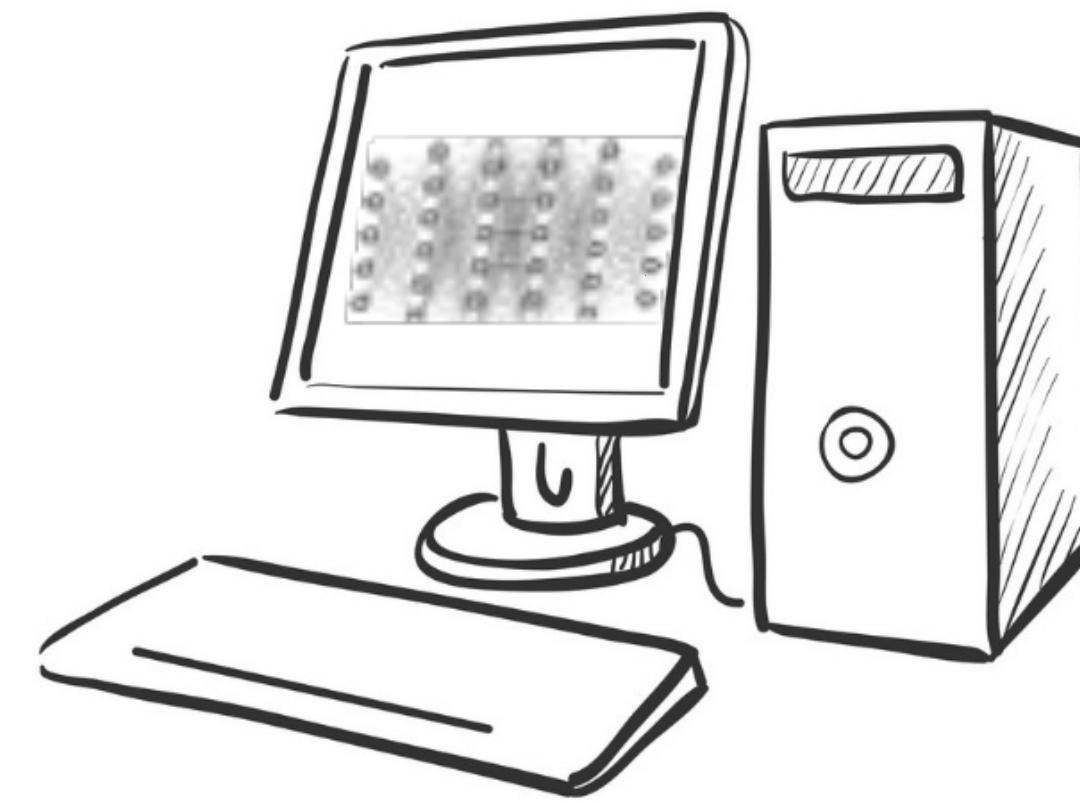
**Healthy**

?

**Pathological**

# Geometric guarantees (invariance)

*Example: Detection of pathological cells*

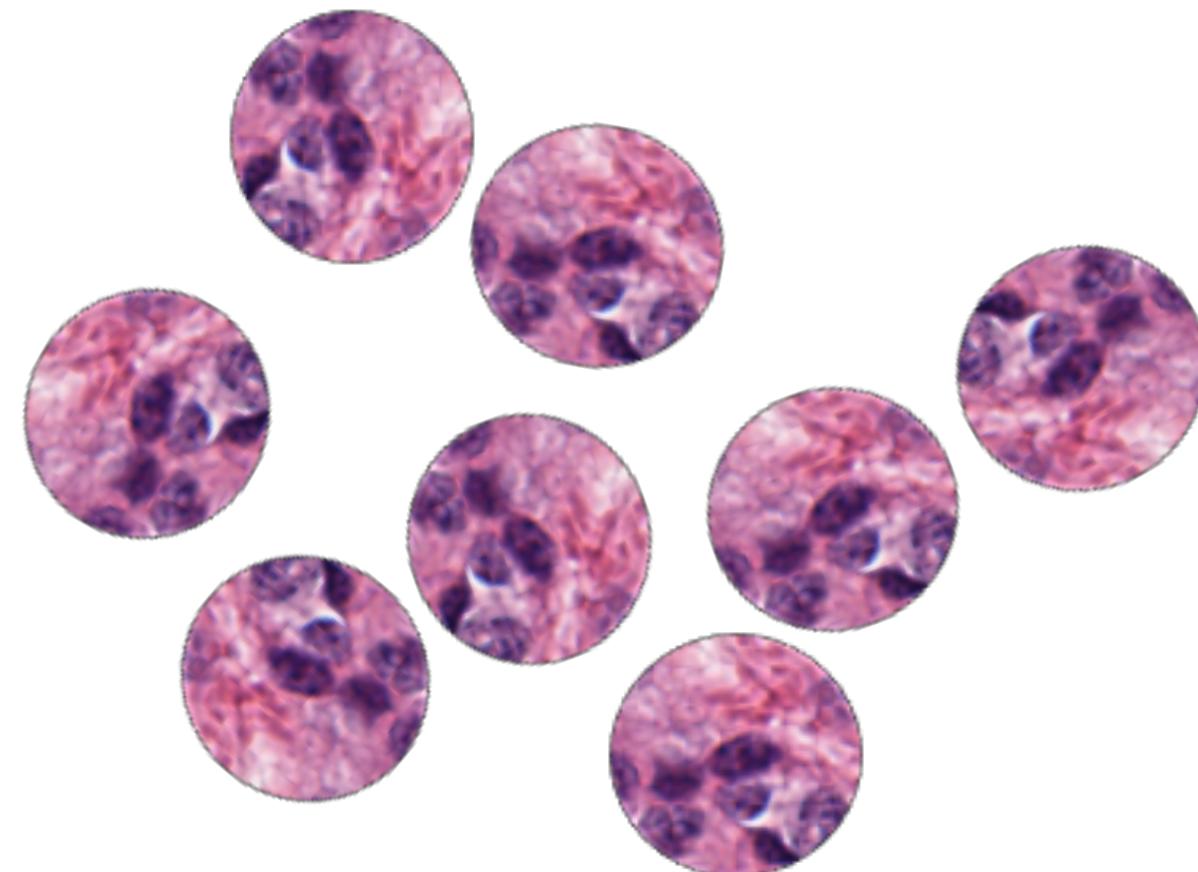


**Healthy**

?

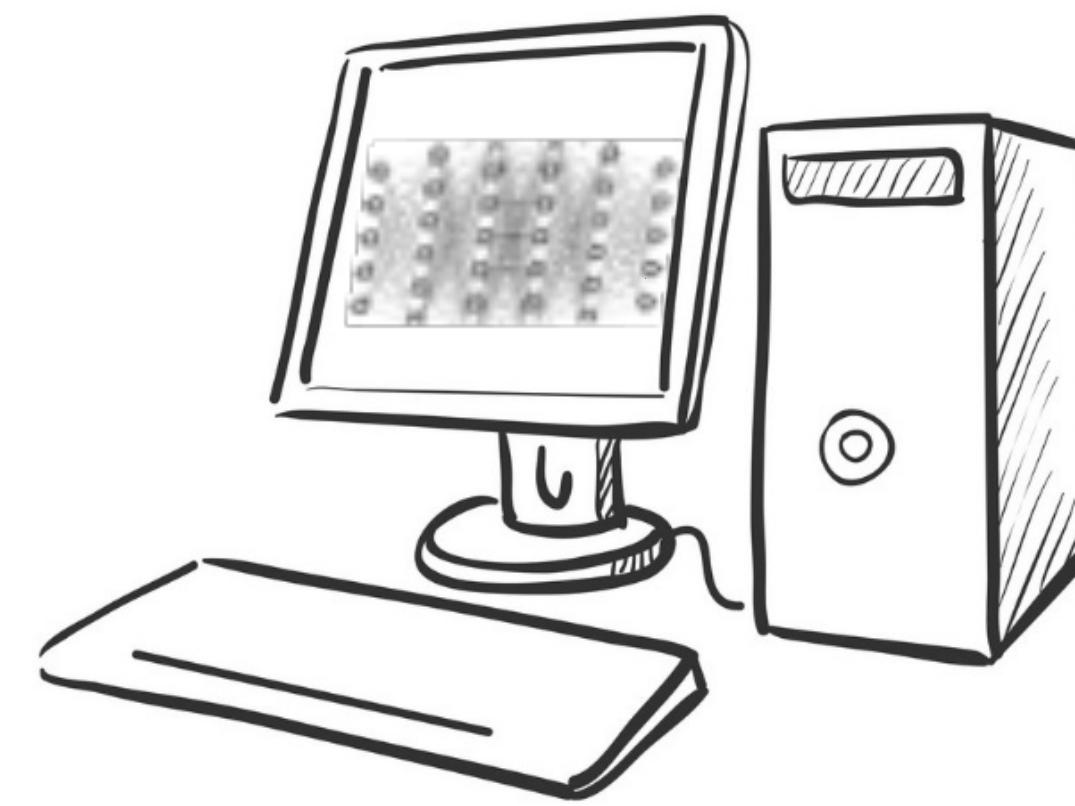
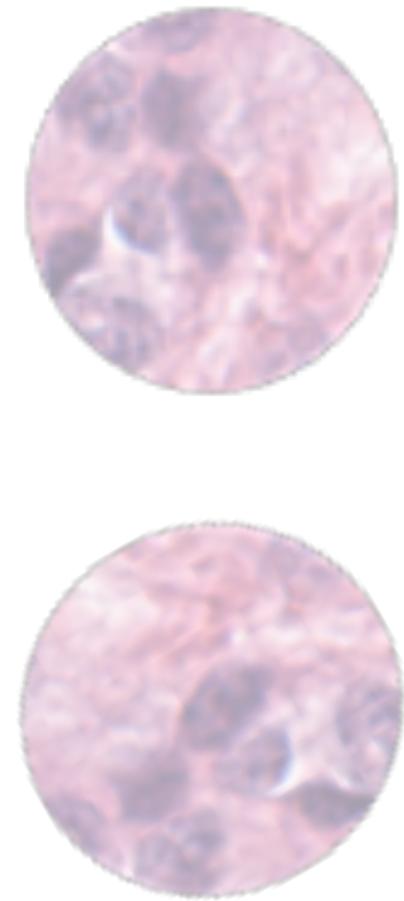
**Pathological**

Common approach: data-augmentation



# Geometric guarantees (invariance)

*Example: Detection of pathological cells*

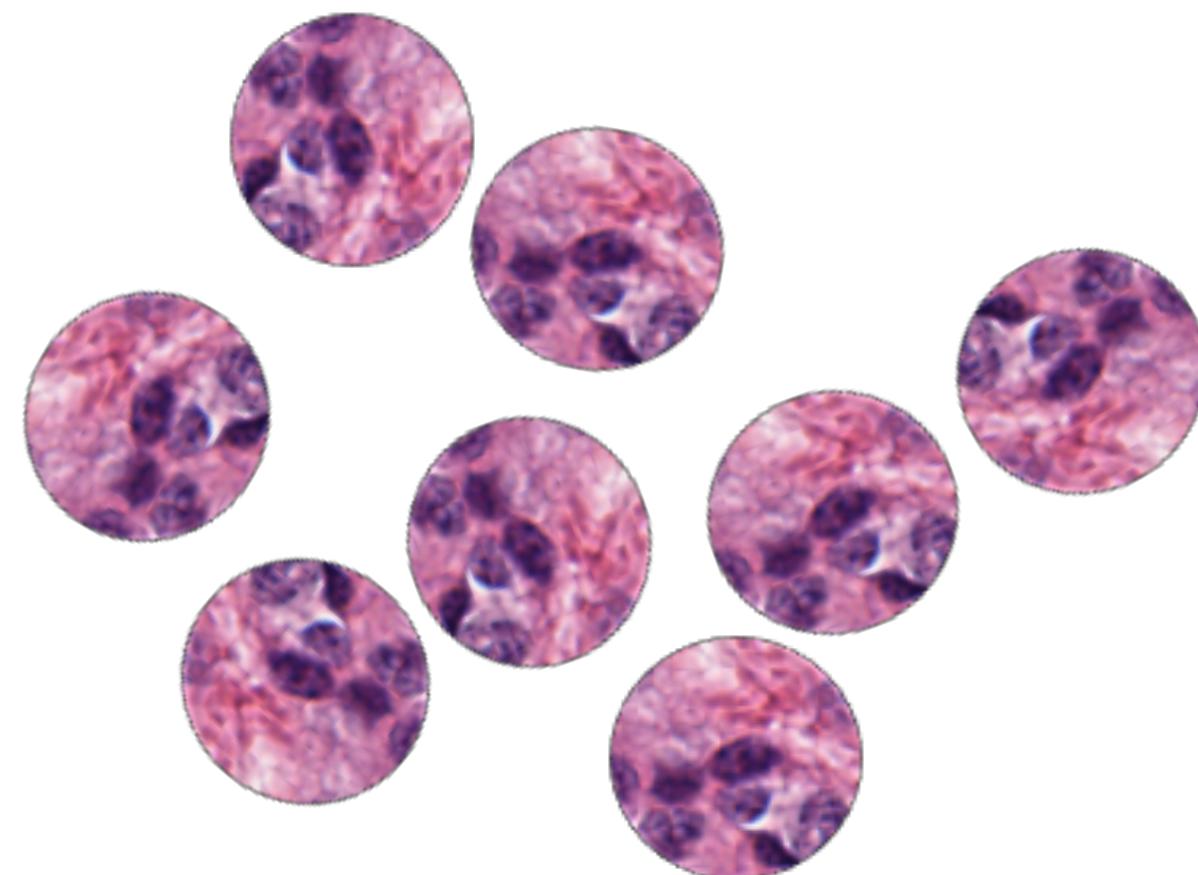


**Healthy**

?

**Pathological**

Common approach: data-augmentation



**Issues:**

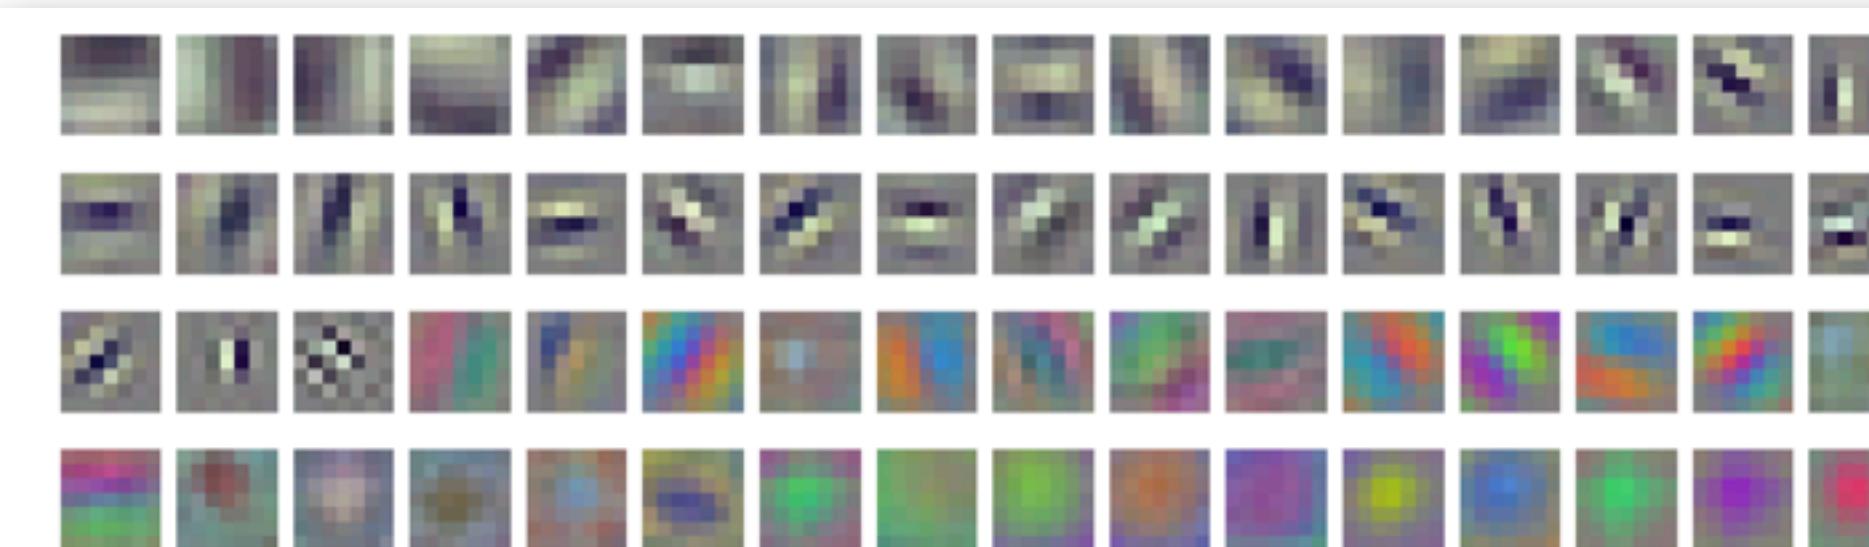
- Still no guarantee of invariance
- Valuable net capacity is spent on learning invariance
- Redundancy in feature repr.



<https://distill.pub/2020/circuits-equivariance/>

## Naturally Occurring Equivariance in Neural Networks

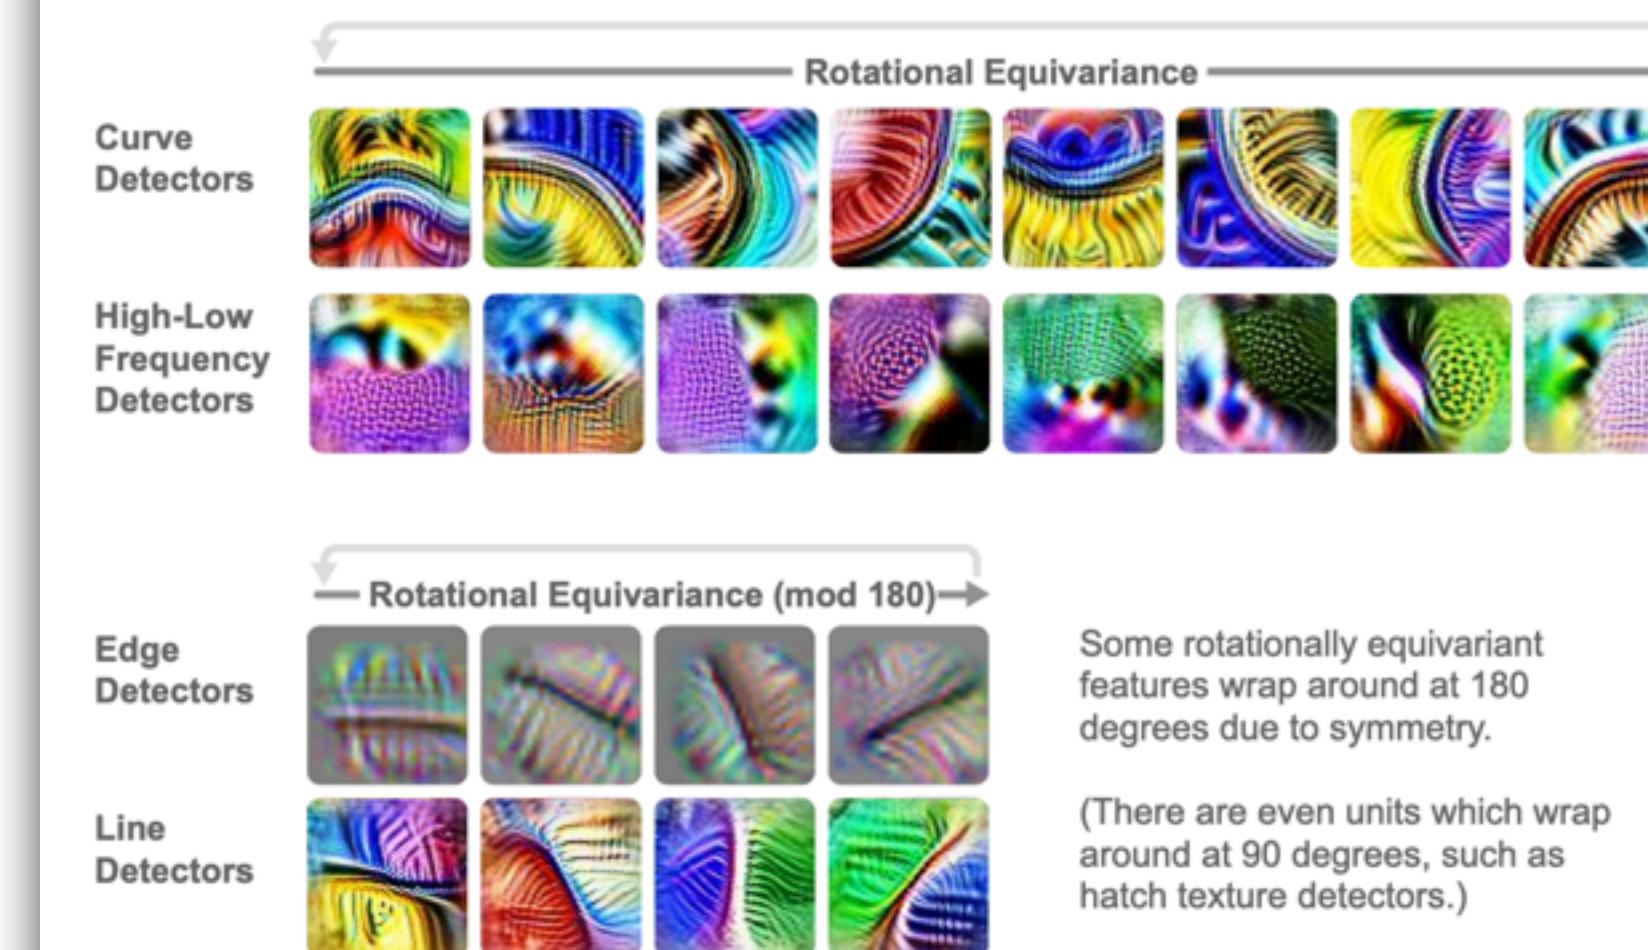
AUTHORS	AFFILIATIONS	PUBLISHED	DOI
Chris Olah	OpenAI	Dec. 8, 2020	10.23915/distill.00024.004
Nick Cammarata	OpenAI		
Chelsea Voss	OpenAI		
Ludwig Schubert			
Gabriel Goh	OpenAI		



The weights for the units in the first layer of the TF-Slim [11] version of InceptionV1 [8].<sup>5</sup> Units are sorted by the first principal component of the adjacency matrix between the first and second layers. Note how many features are similar except for rotation, scale, and hue.

## Equivariant Features

**Rotational Equivariance:** One example of equivariance is rotated versions of the same feature. These are especially common in early vision, for example curve detectors, high-low frequency detectors, and line detectors.



# Geometric guarantees (equivariance)

CNNs are translation equivariant



Via convolutions



# Geometric guarantees (equivariance)

CNNs are translation equivariant

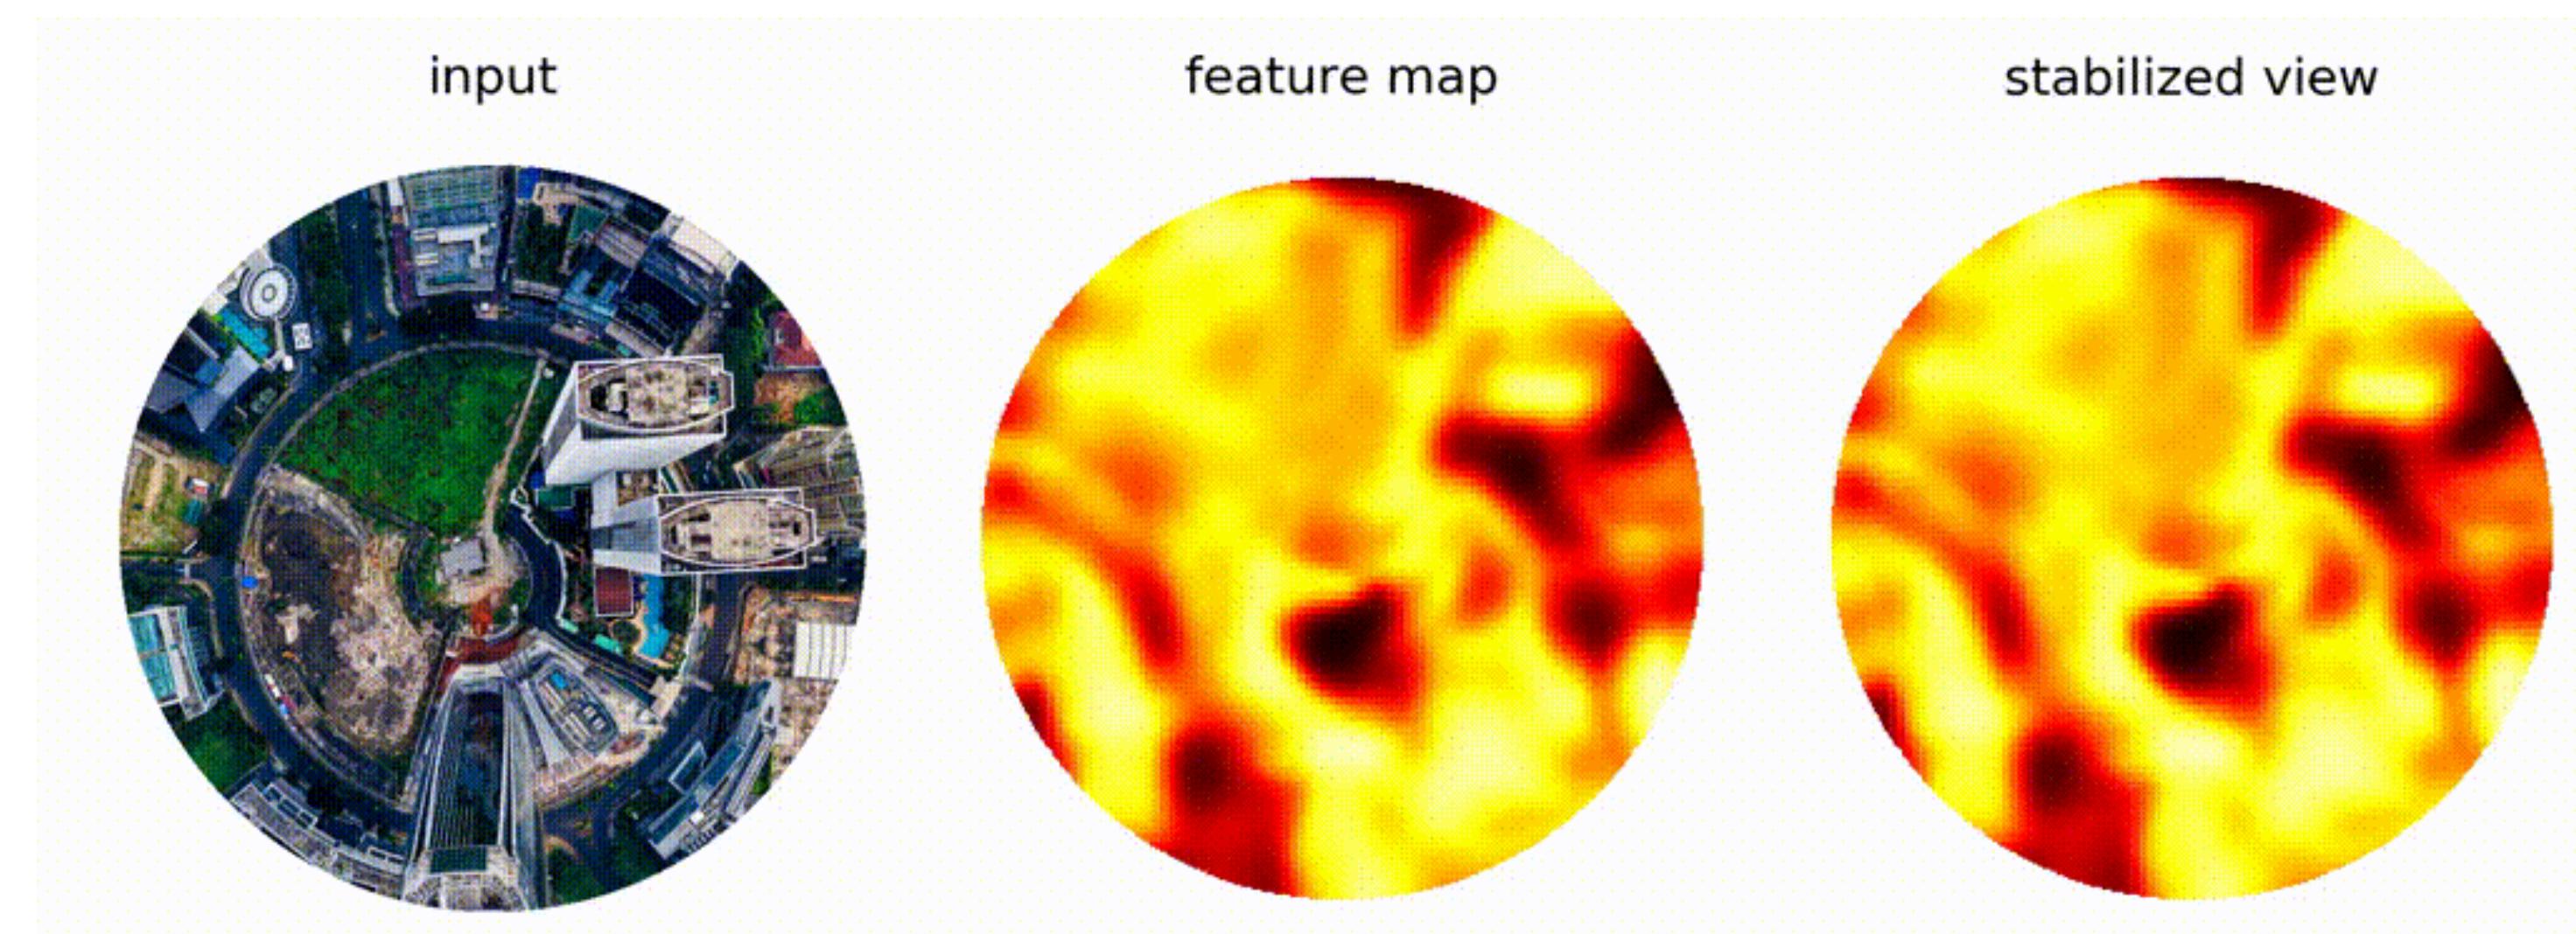


Via convolutions



# Geometric guarantees (equivariance)

Normal CNN

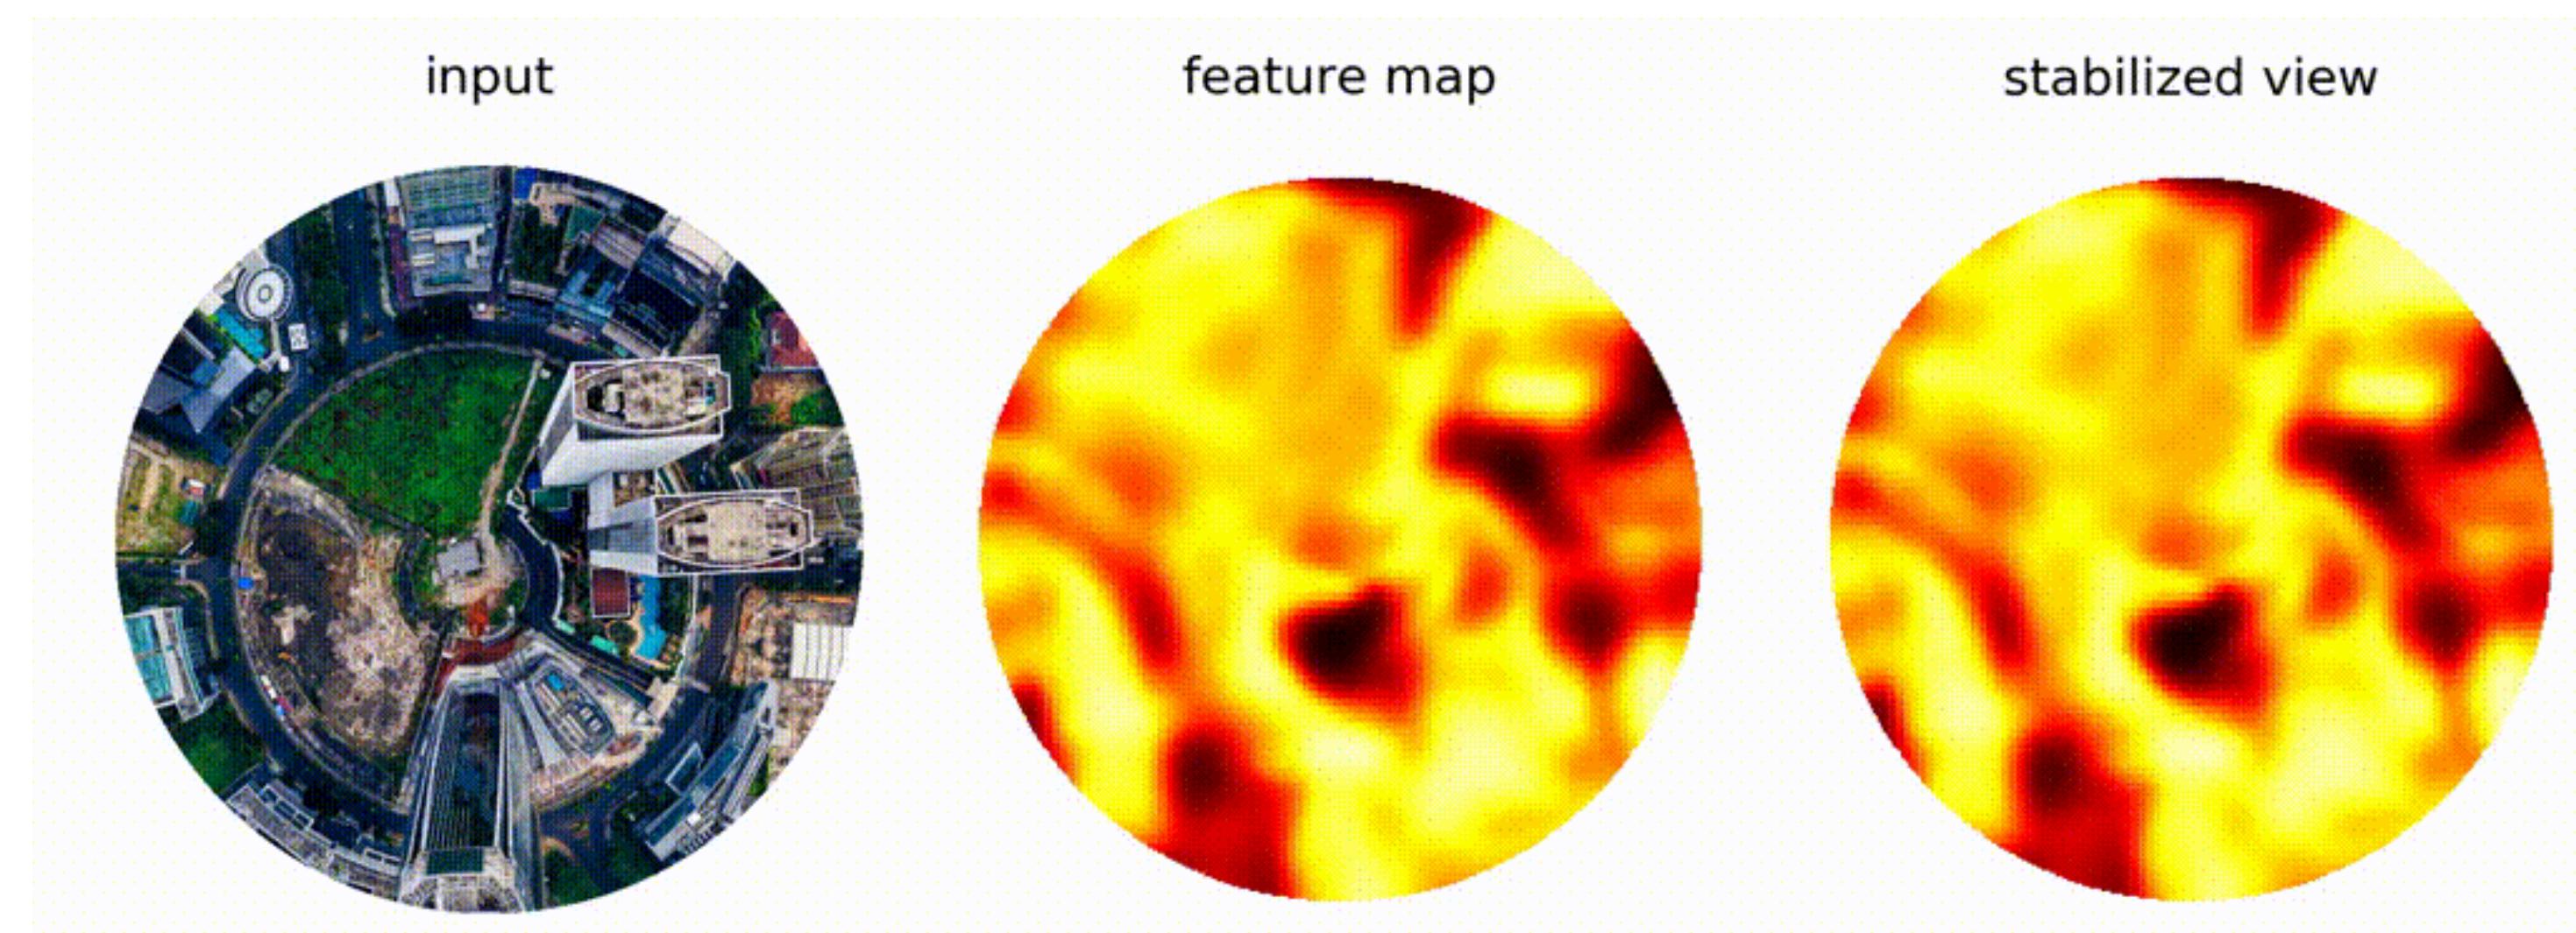


Figures source:

<https://github.com/QUVA-Lab/e2cnn>

# Geometric guarantees (equivariance)

Normal CNN

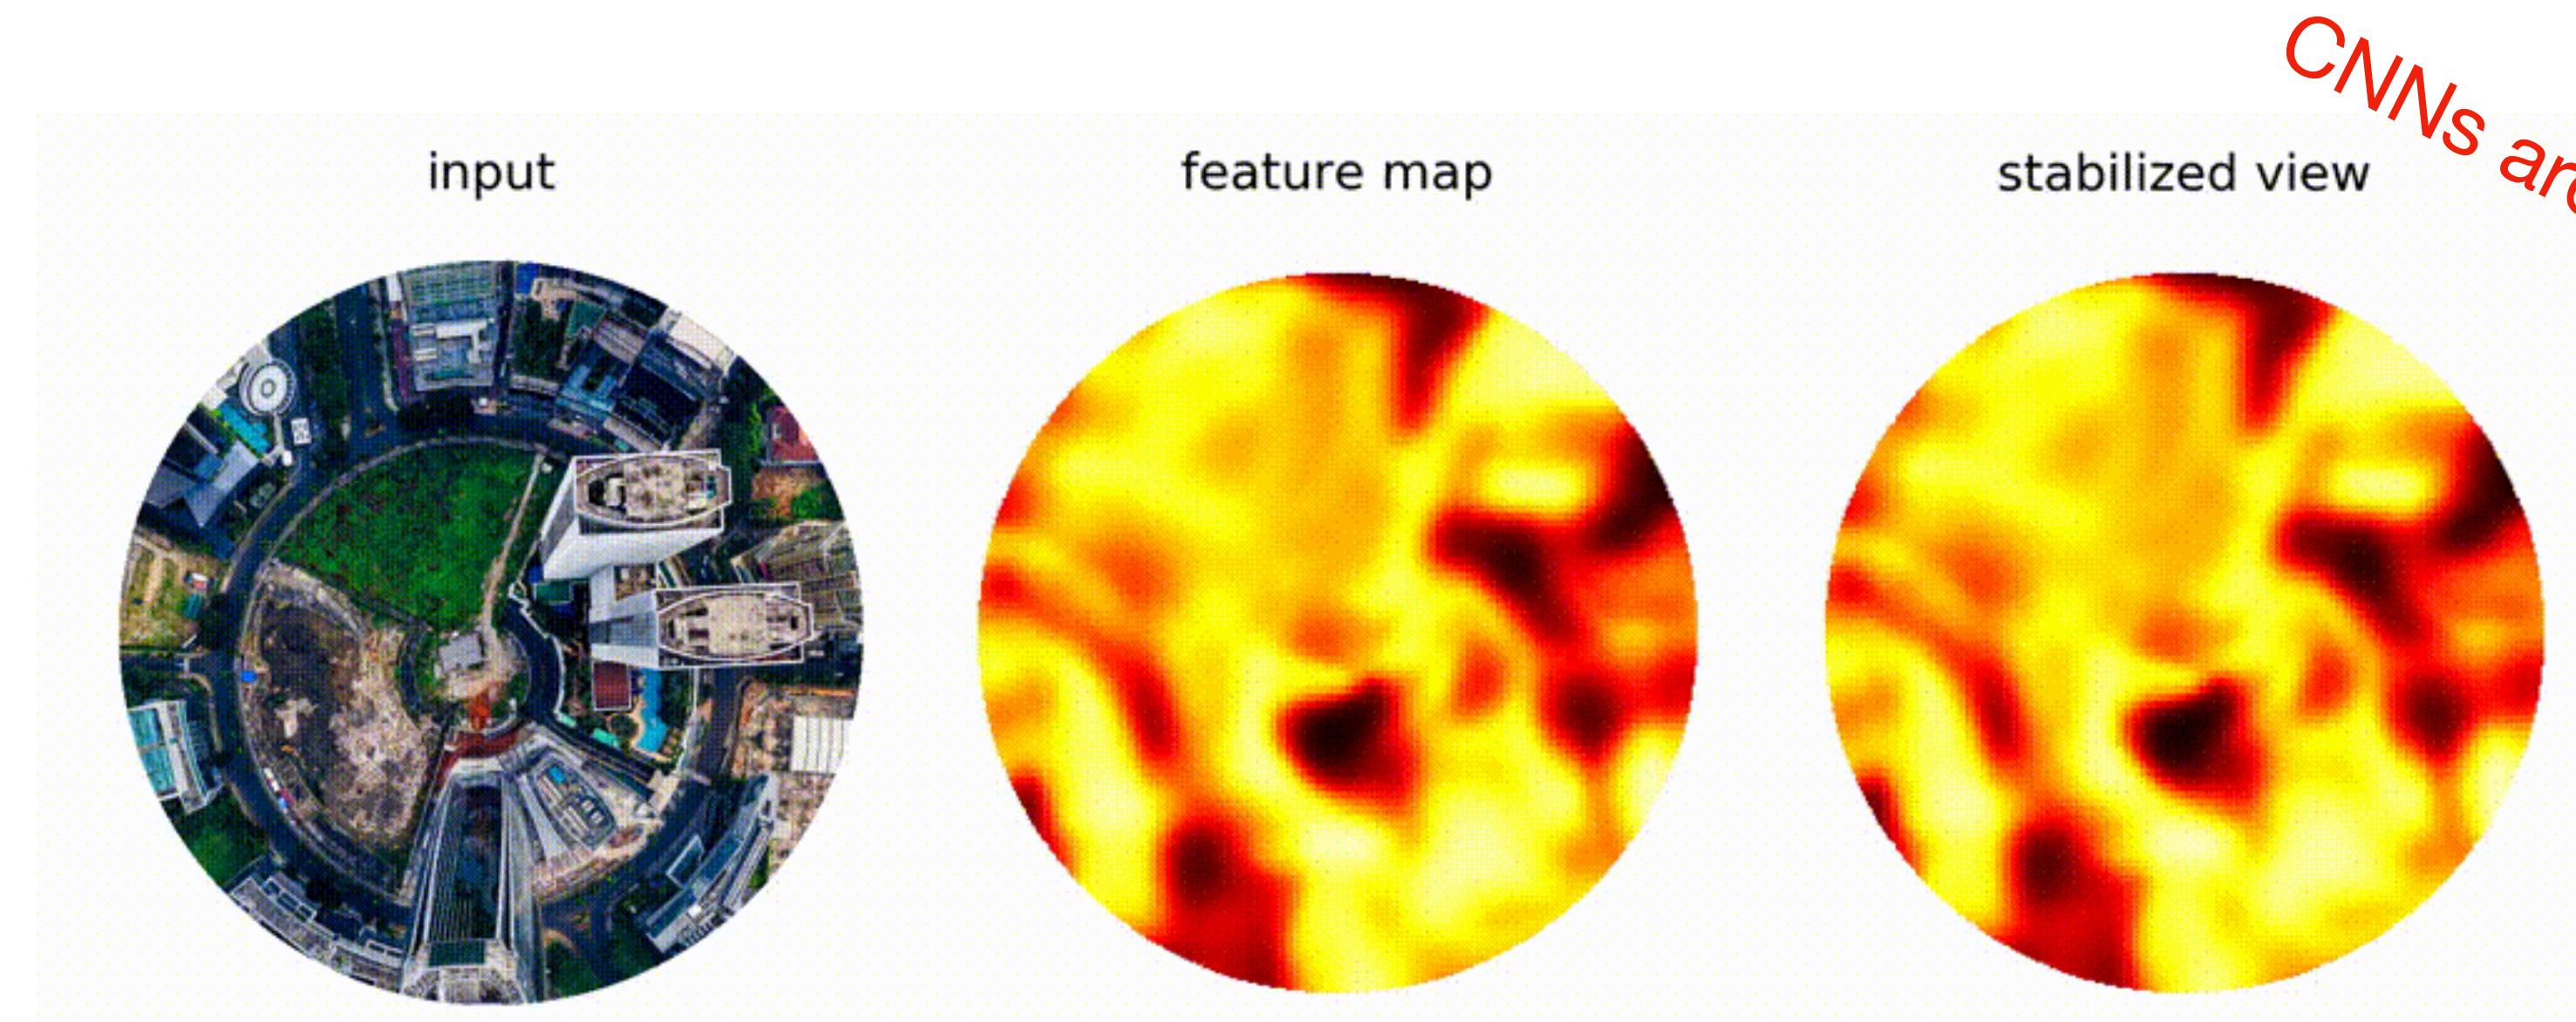


Figures source:

<https://github.com/QUVA-Lab/e2cnn>

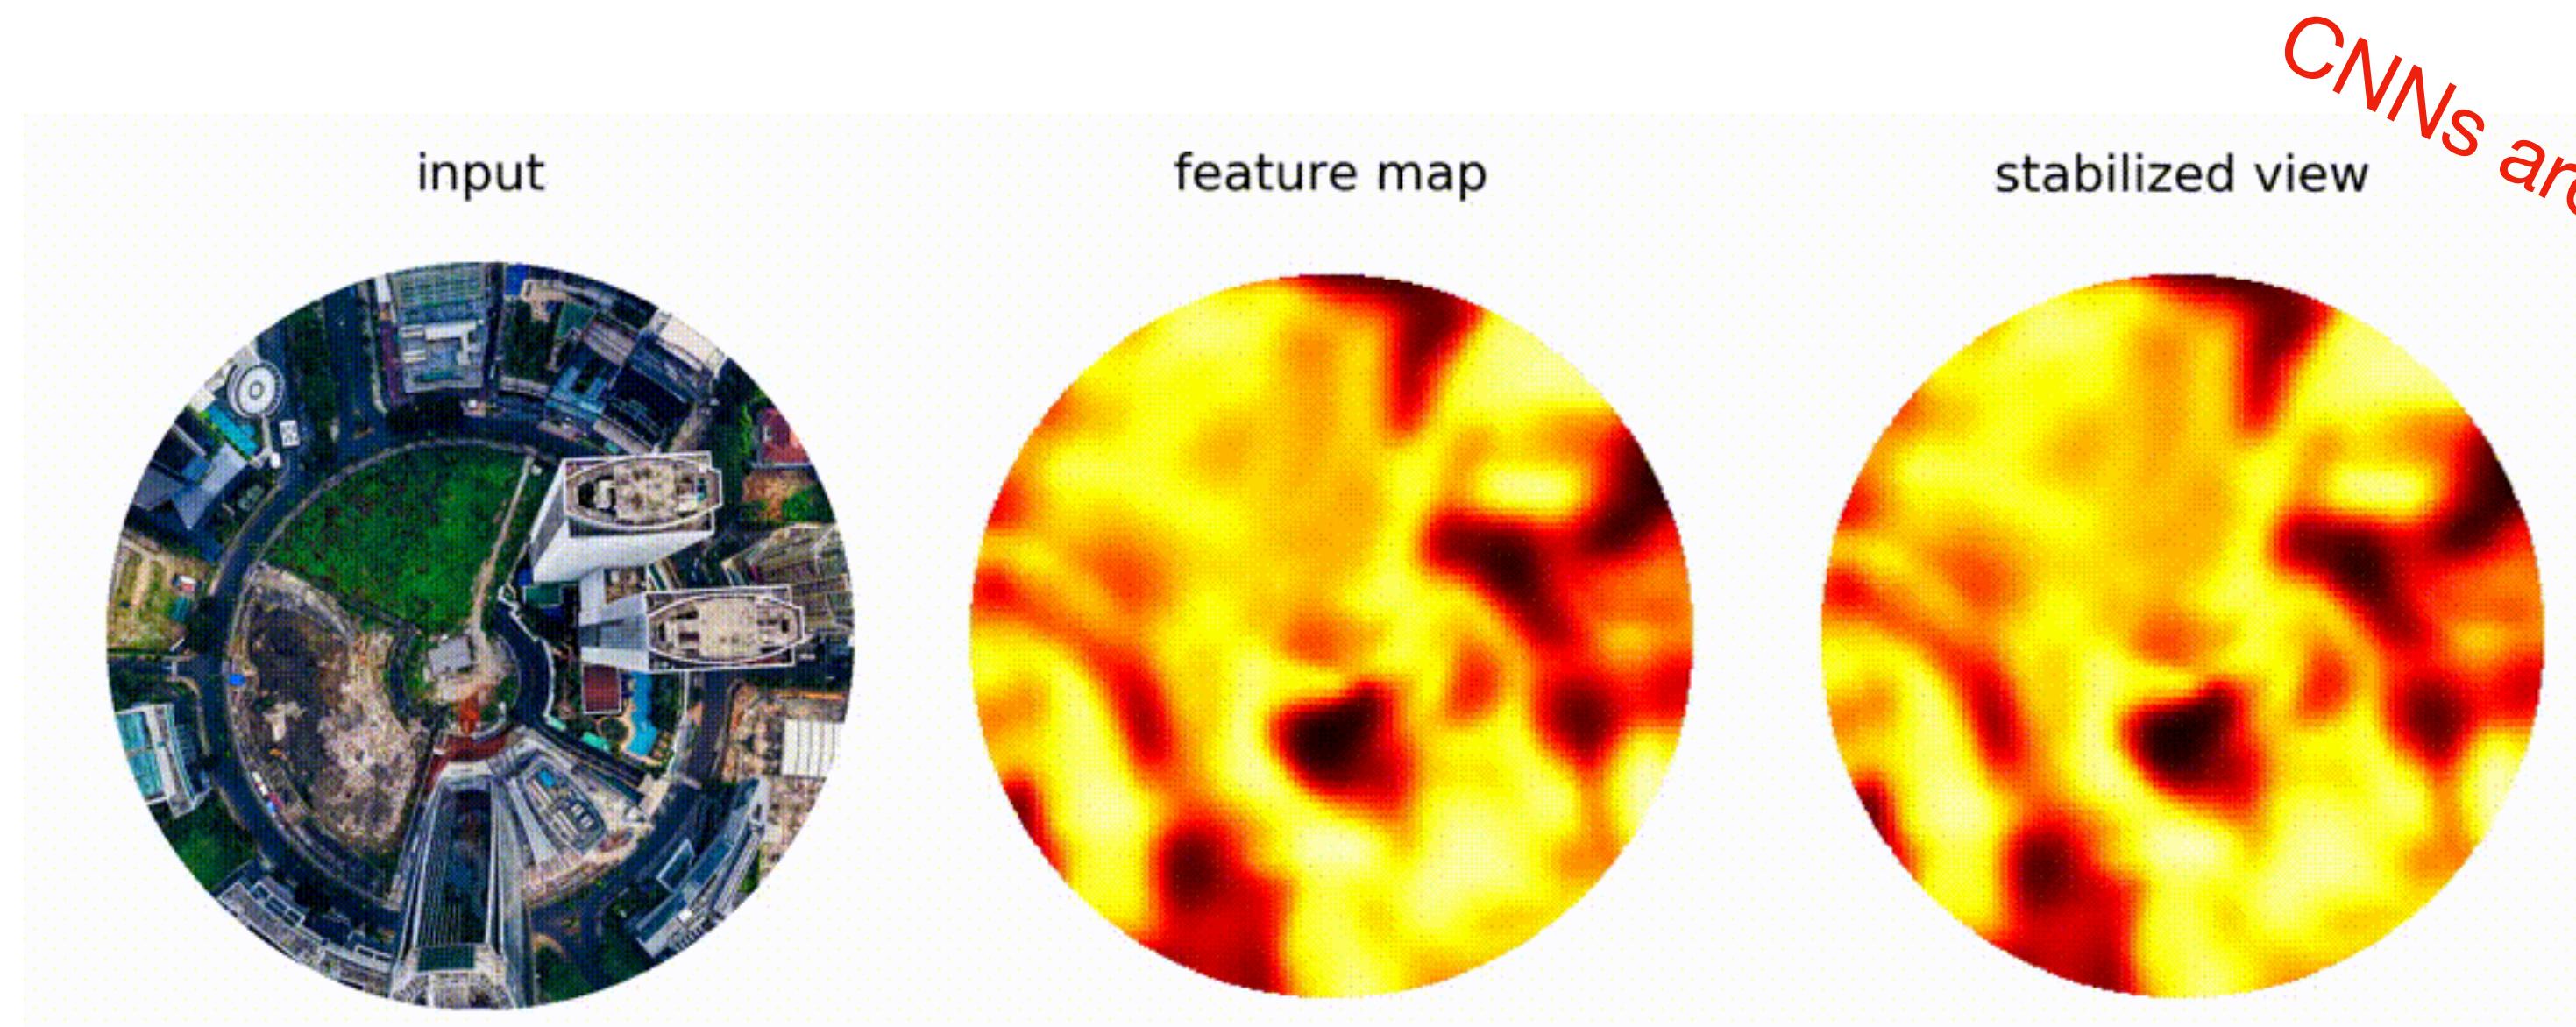
# Geometric guarantees (equivariance)

Normal CNN

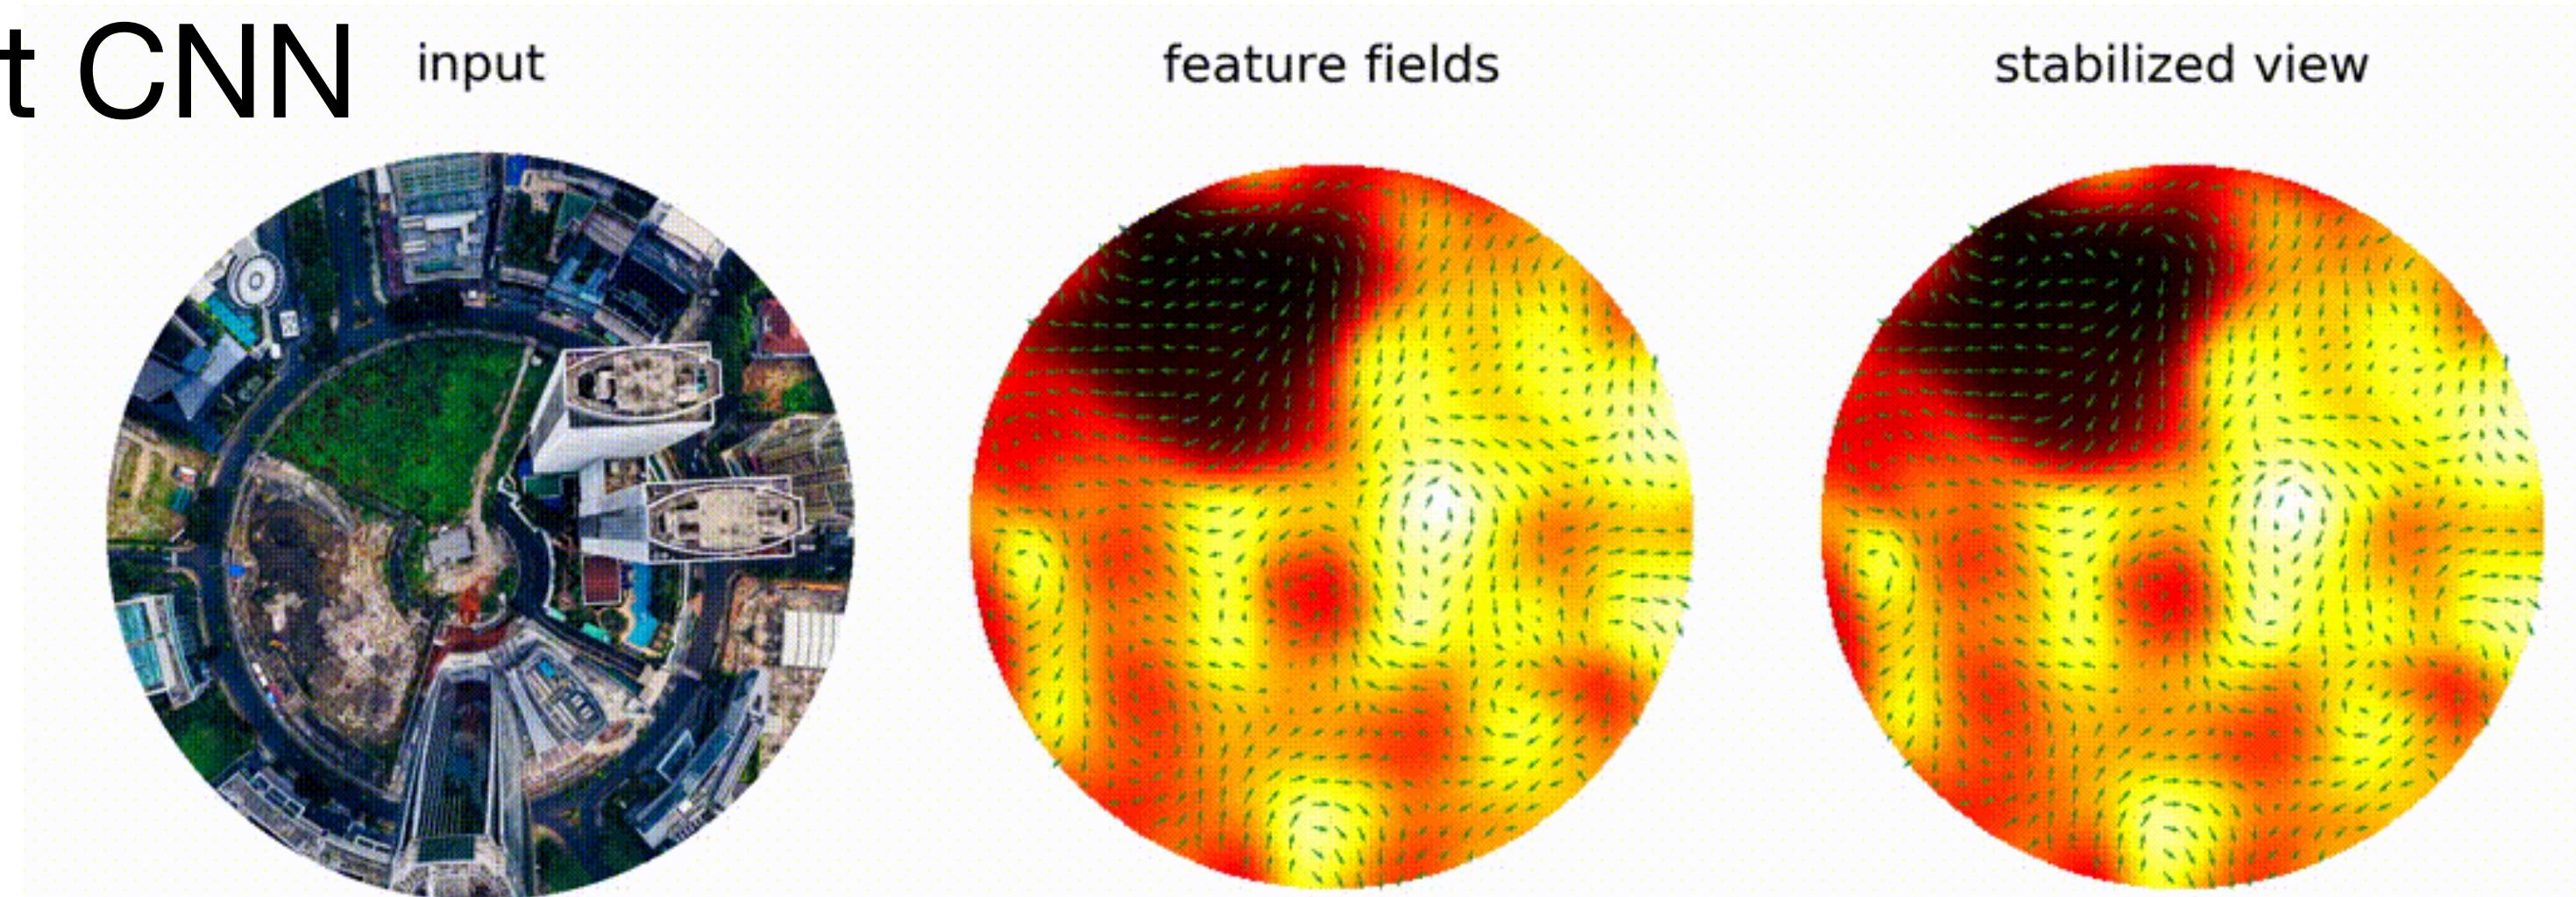


# Geometric guarantees (equivariance)

Normal CNN



Group equivariant CNN

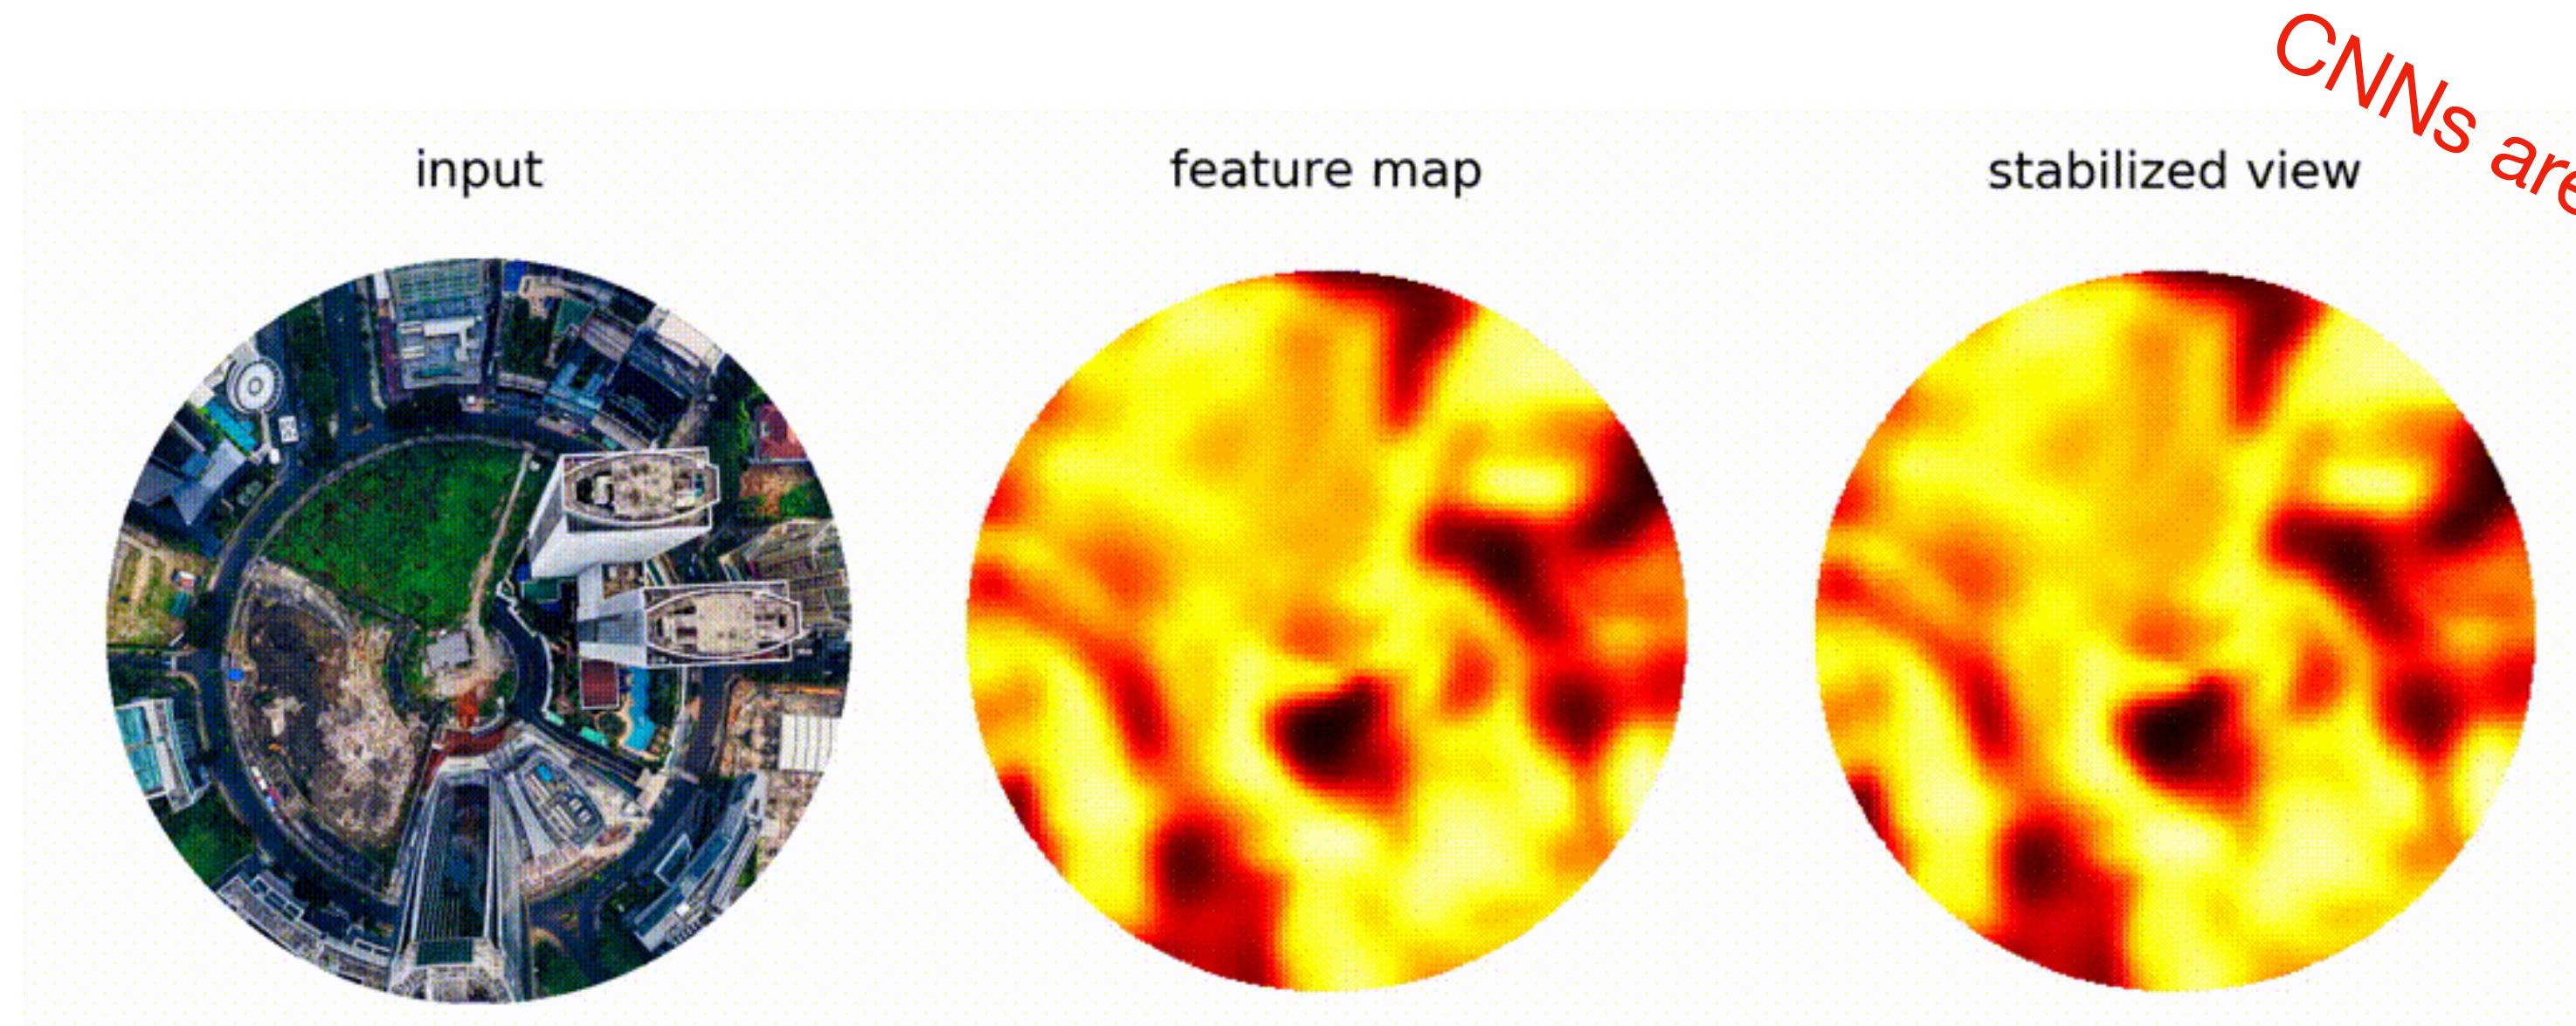


Figures source:

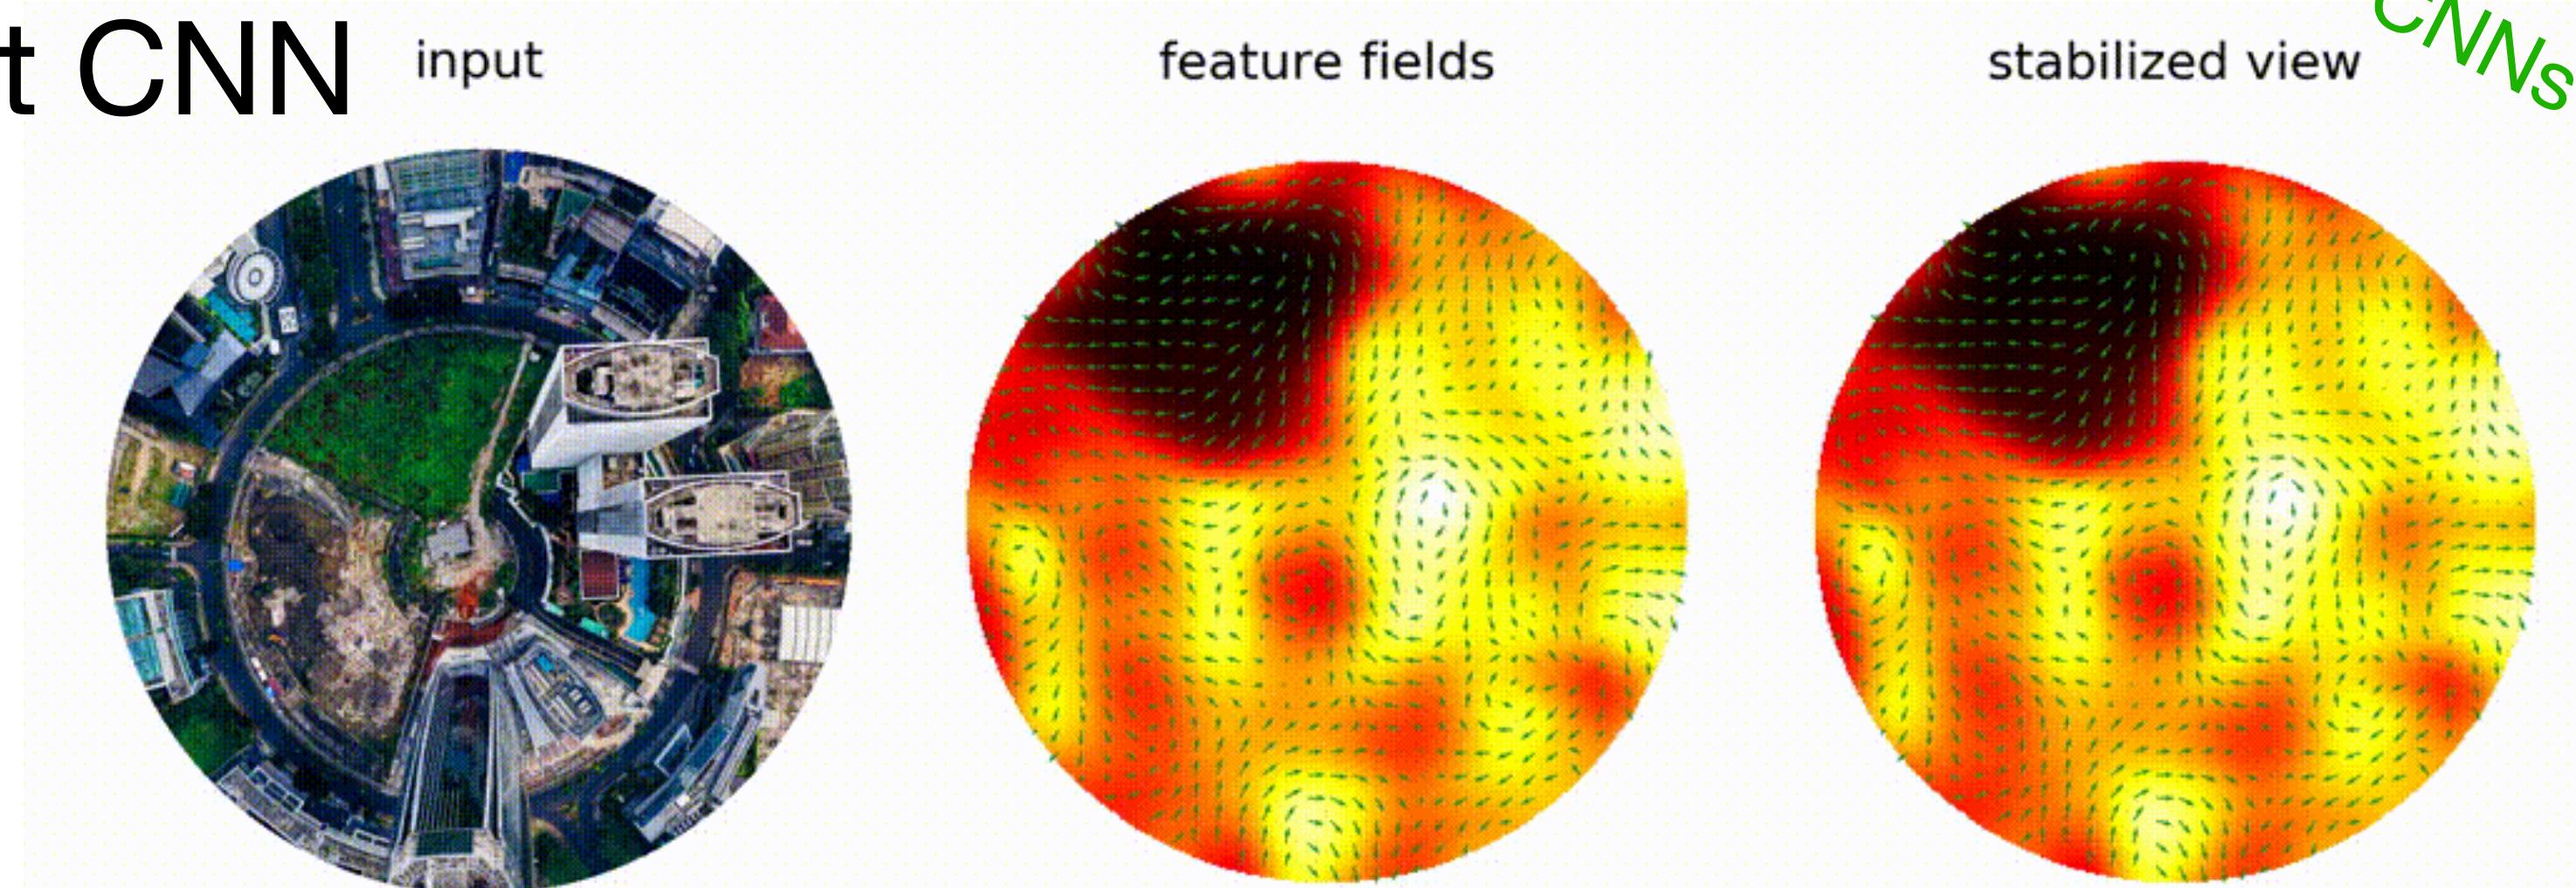
<https://github.com/QUVA-Lab/e2cnn>

# Geometric guarantees (equivariance)

Normal CNN



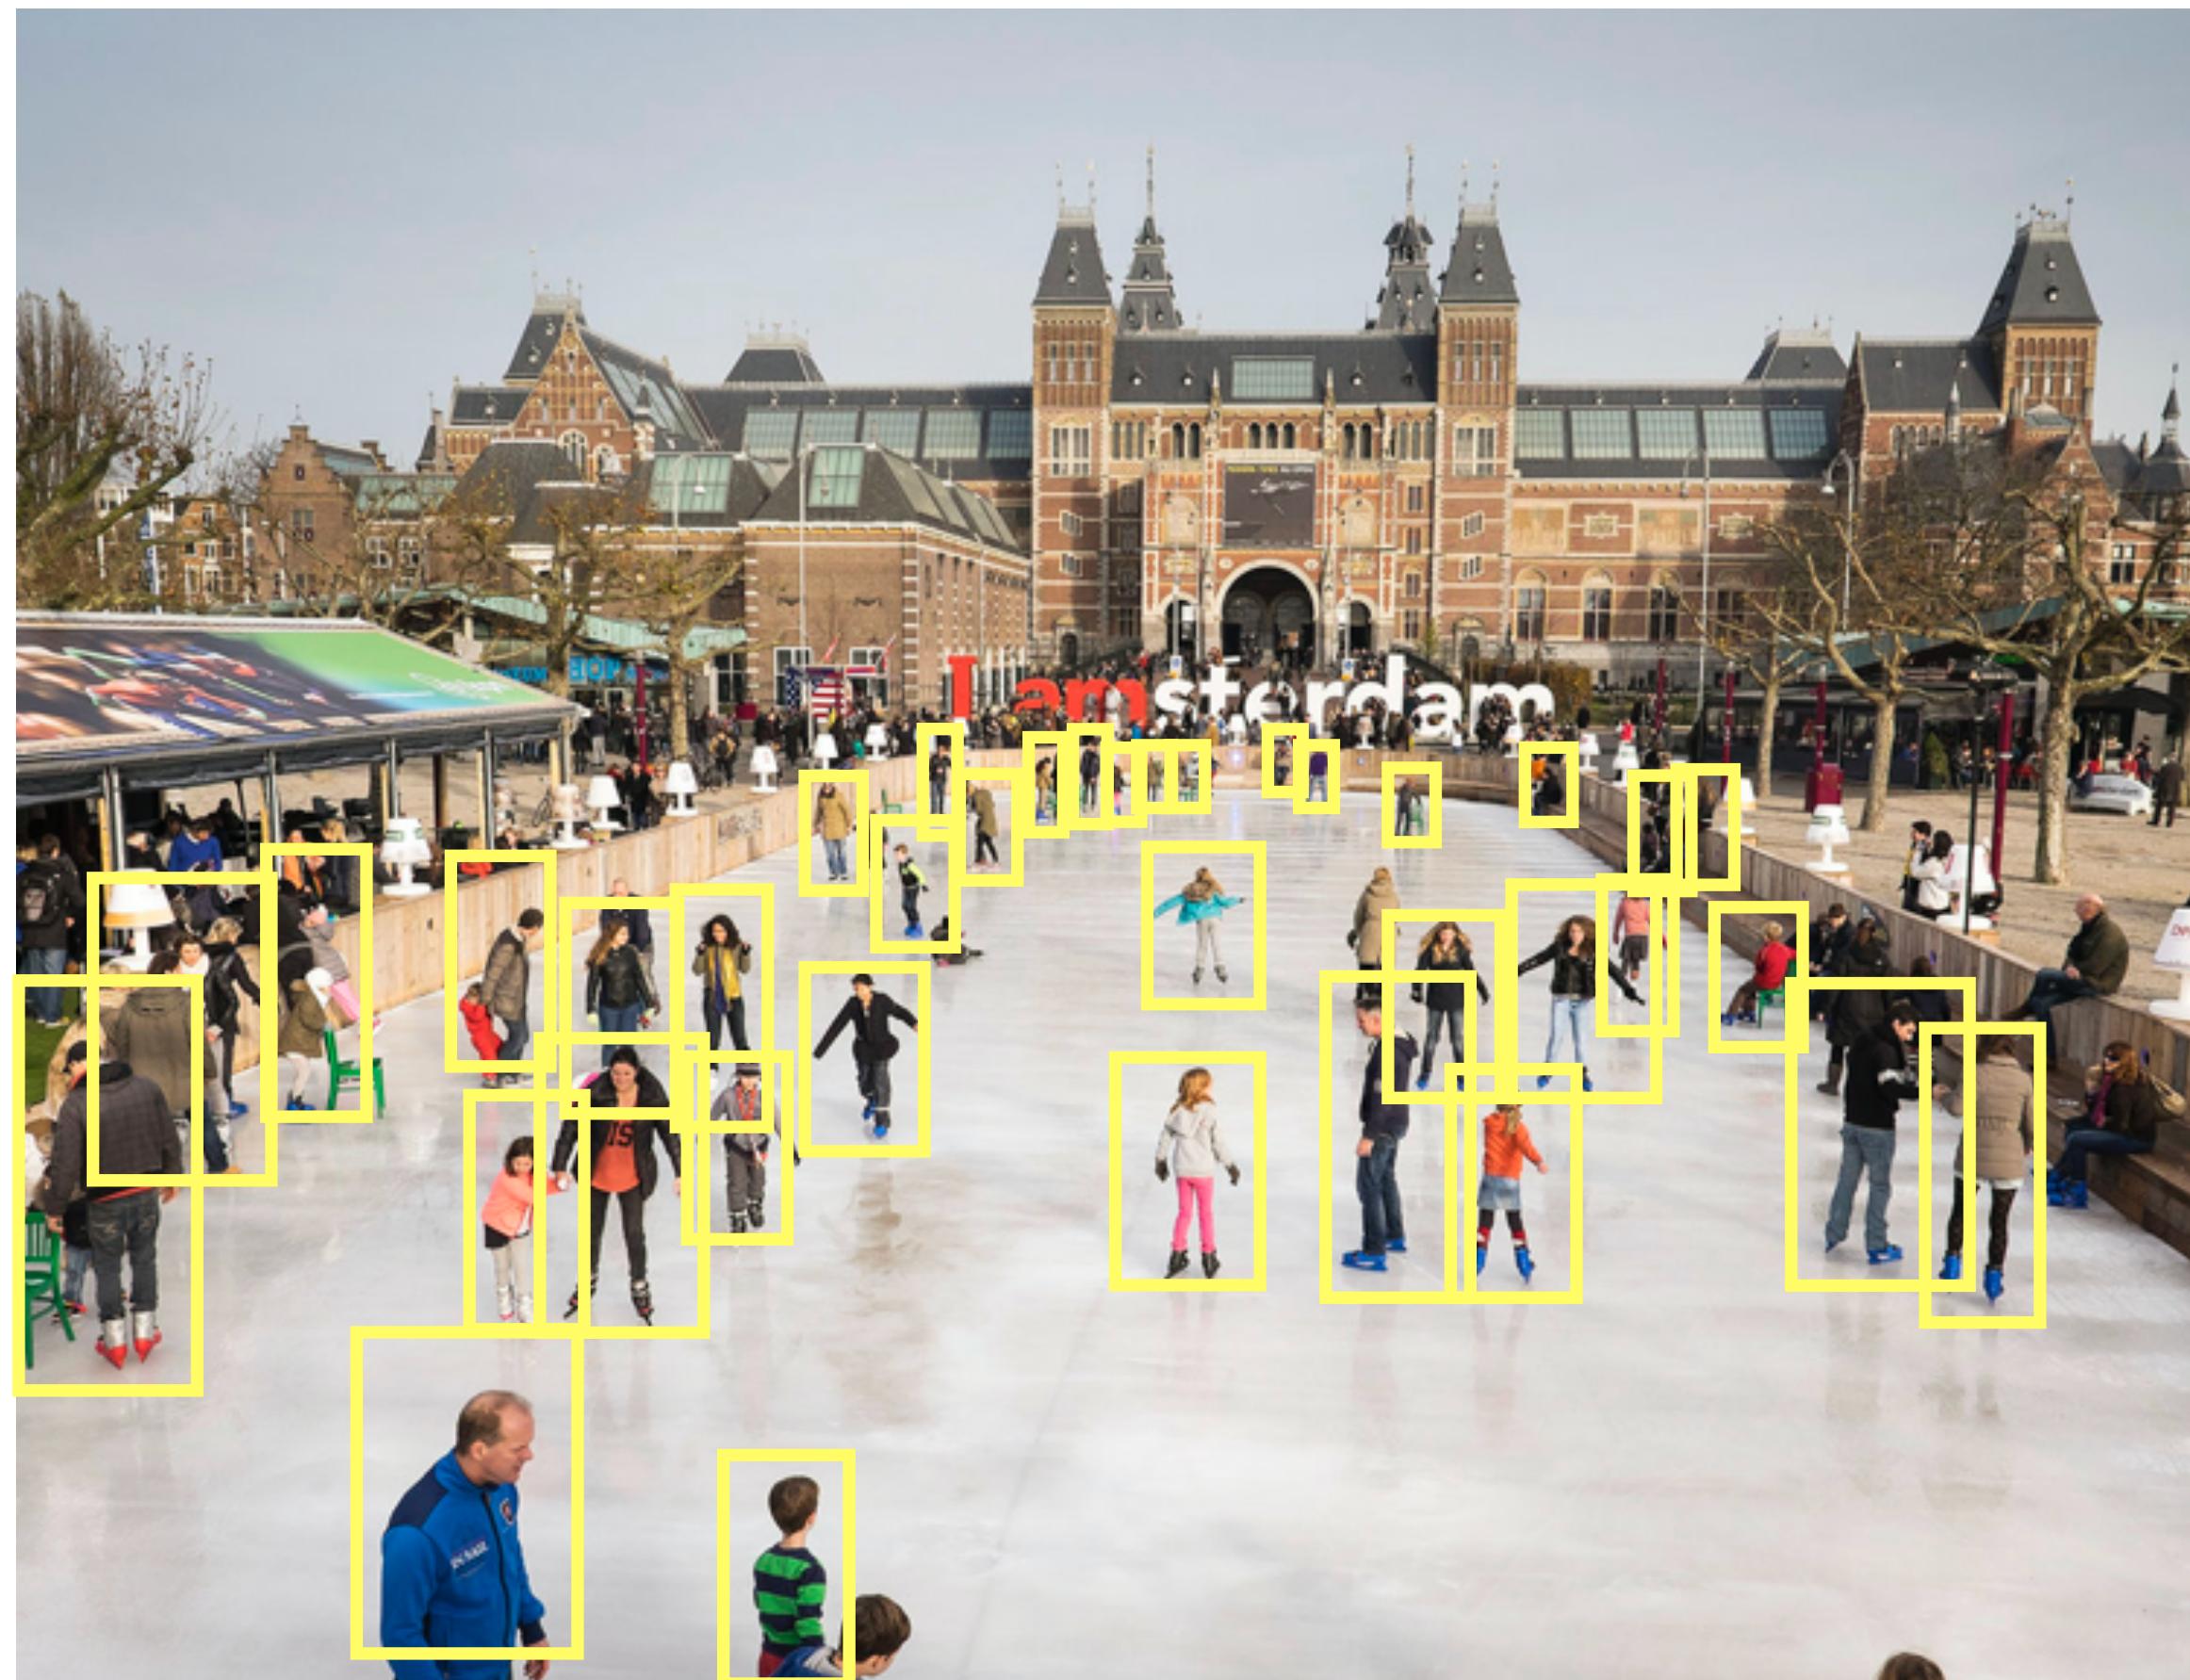
Group equivariant CNN



Figures source:

<https://github.com/QUVA-Lab/e2cnn>

# Geometric guarantees (equivariance)



Importance of equivariance:

- No information is lost when the input is transformed
- Guaranteed stability to (local + global) transformations

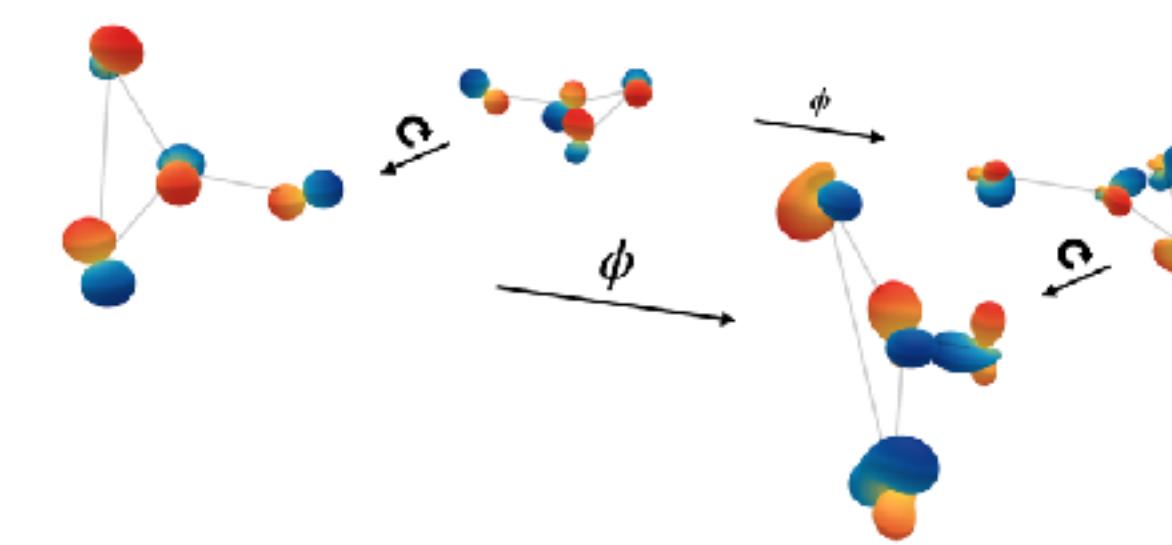
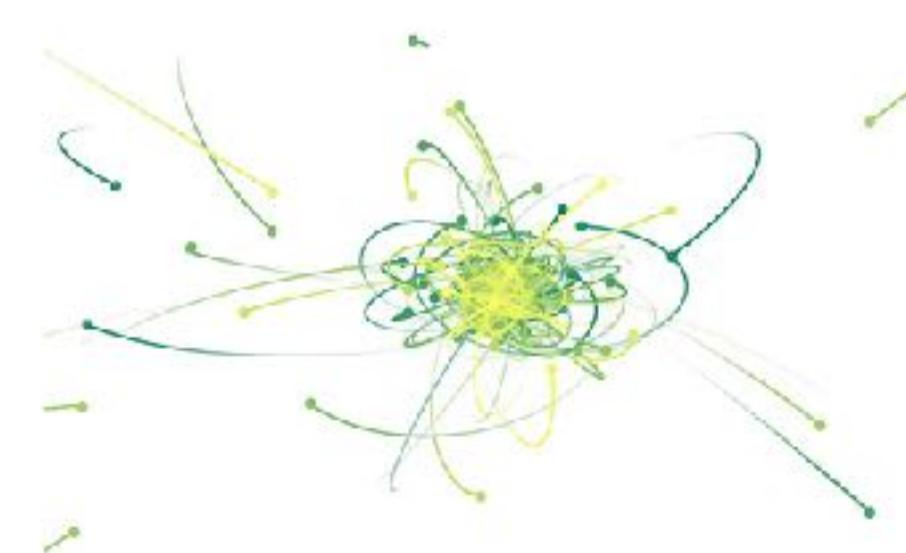
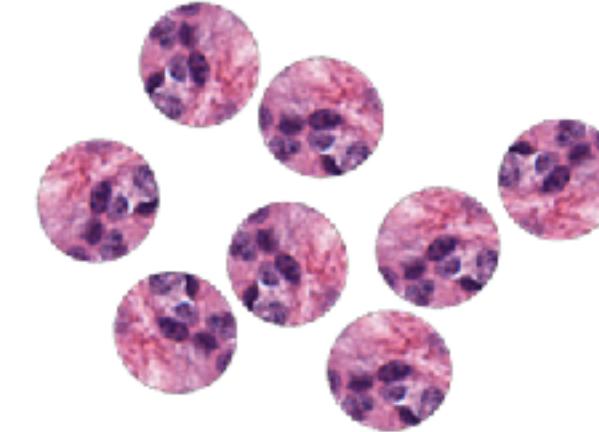
Group convolutions:

- Equivariance beyond translations
- Geometric guarantees
- Increased weight sharing

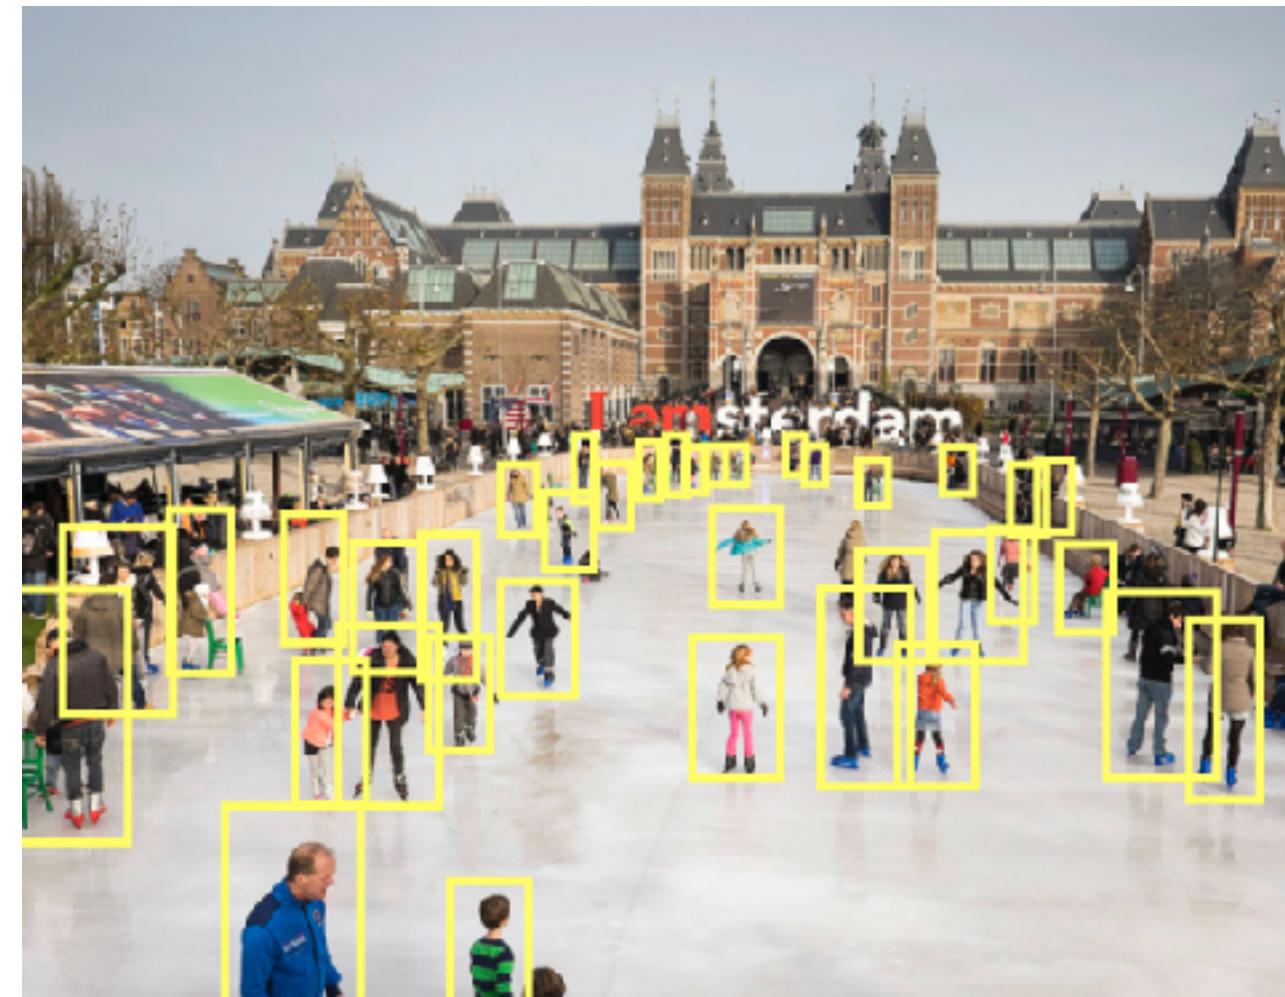
**G-CNNs are not only relevant for invariant problems but for any type of structured data!**

# Group equivariant deep learning

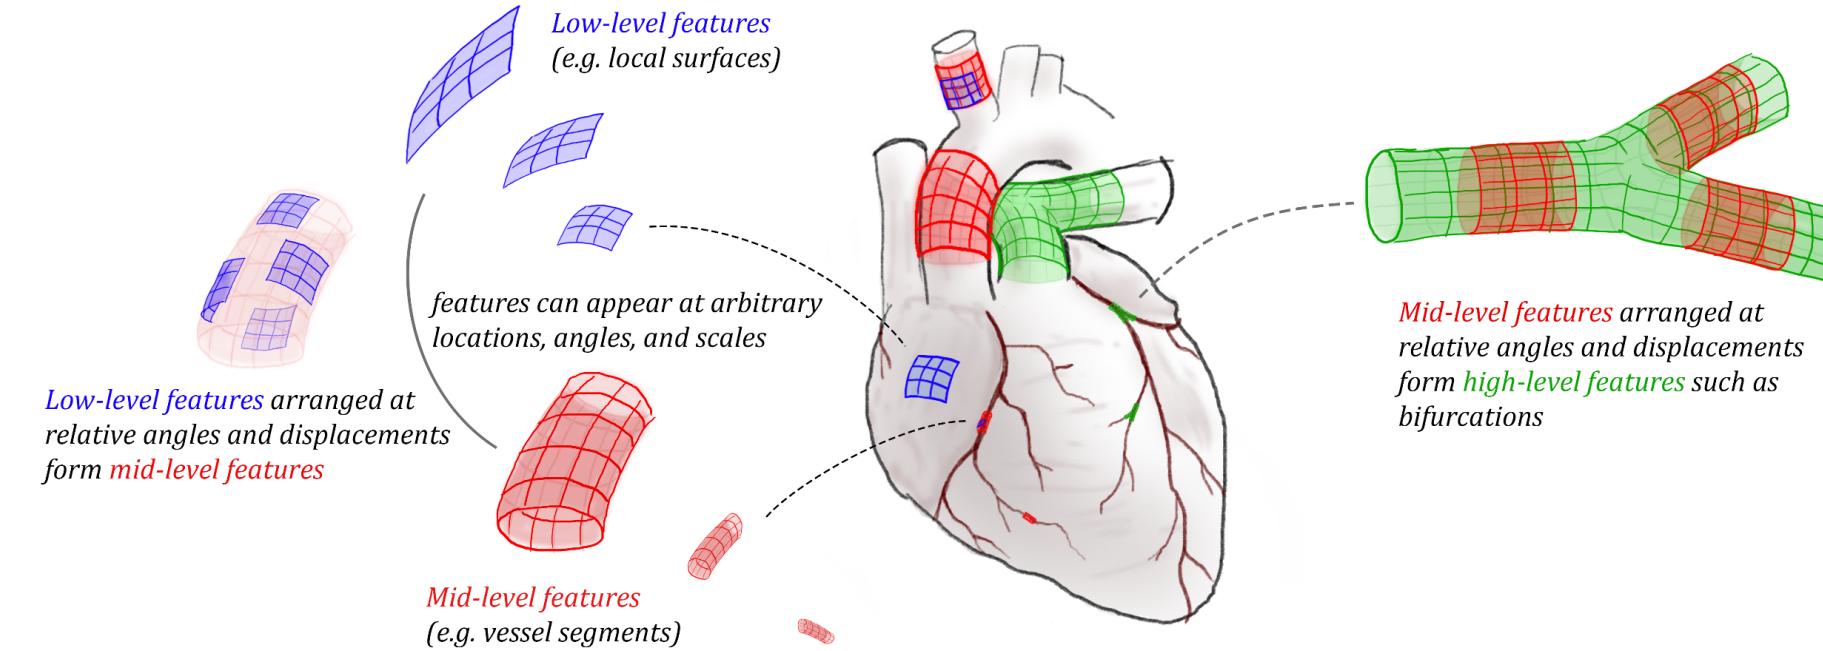
Create architectures with guarantees of invariance or equivariance  
(often demanded by problems)



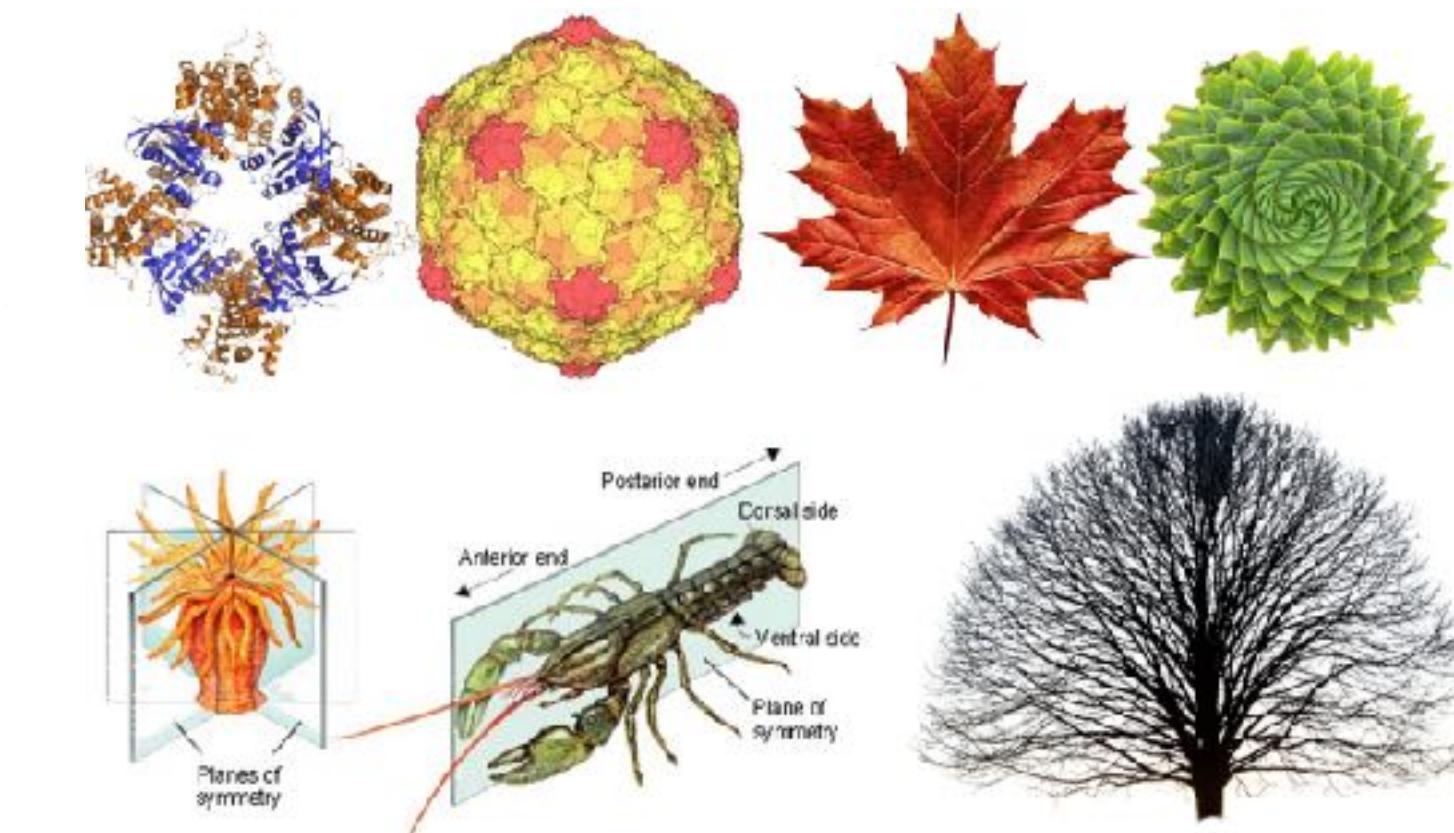
Equivariance allows for increased weight sharing



Psychology of vision  
(recognition by components)



Efficient representation learning  
(leverage symmetries)



## 1. Motivation

---

Equivariance → weight-sharing and generalization

## 2. Pattern matching using group theory

## 3. Group convolutions

## 4. Example

## 5. G-convs are all you need!

## 6. Steerable group convolutions

## 7. Feature fields and escnn library

## 8. Equivariant tensor product layers

## 9. Equivariant graph NNs

1. Motivation

---

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

3. Group convolutions

4. Example

5. G-convs are all you need!

6. Steerable group convolutions

7. Feature fields and escnn library

8. Equivariant tensor product layers

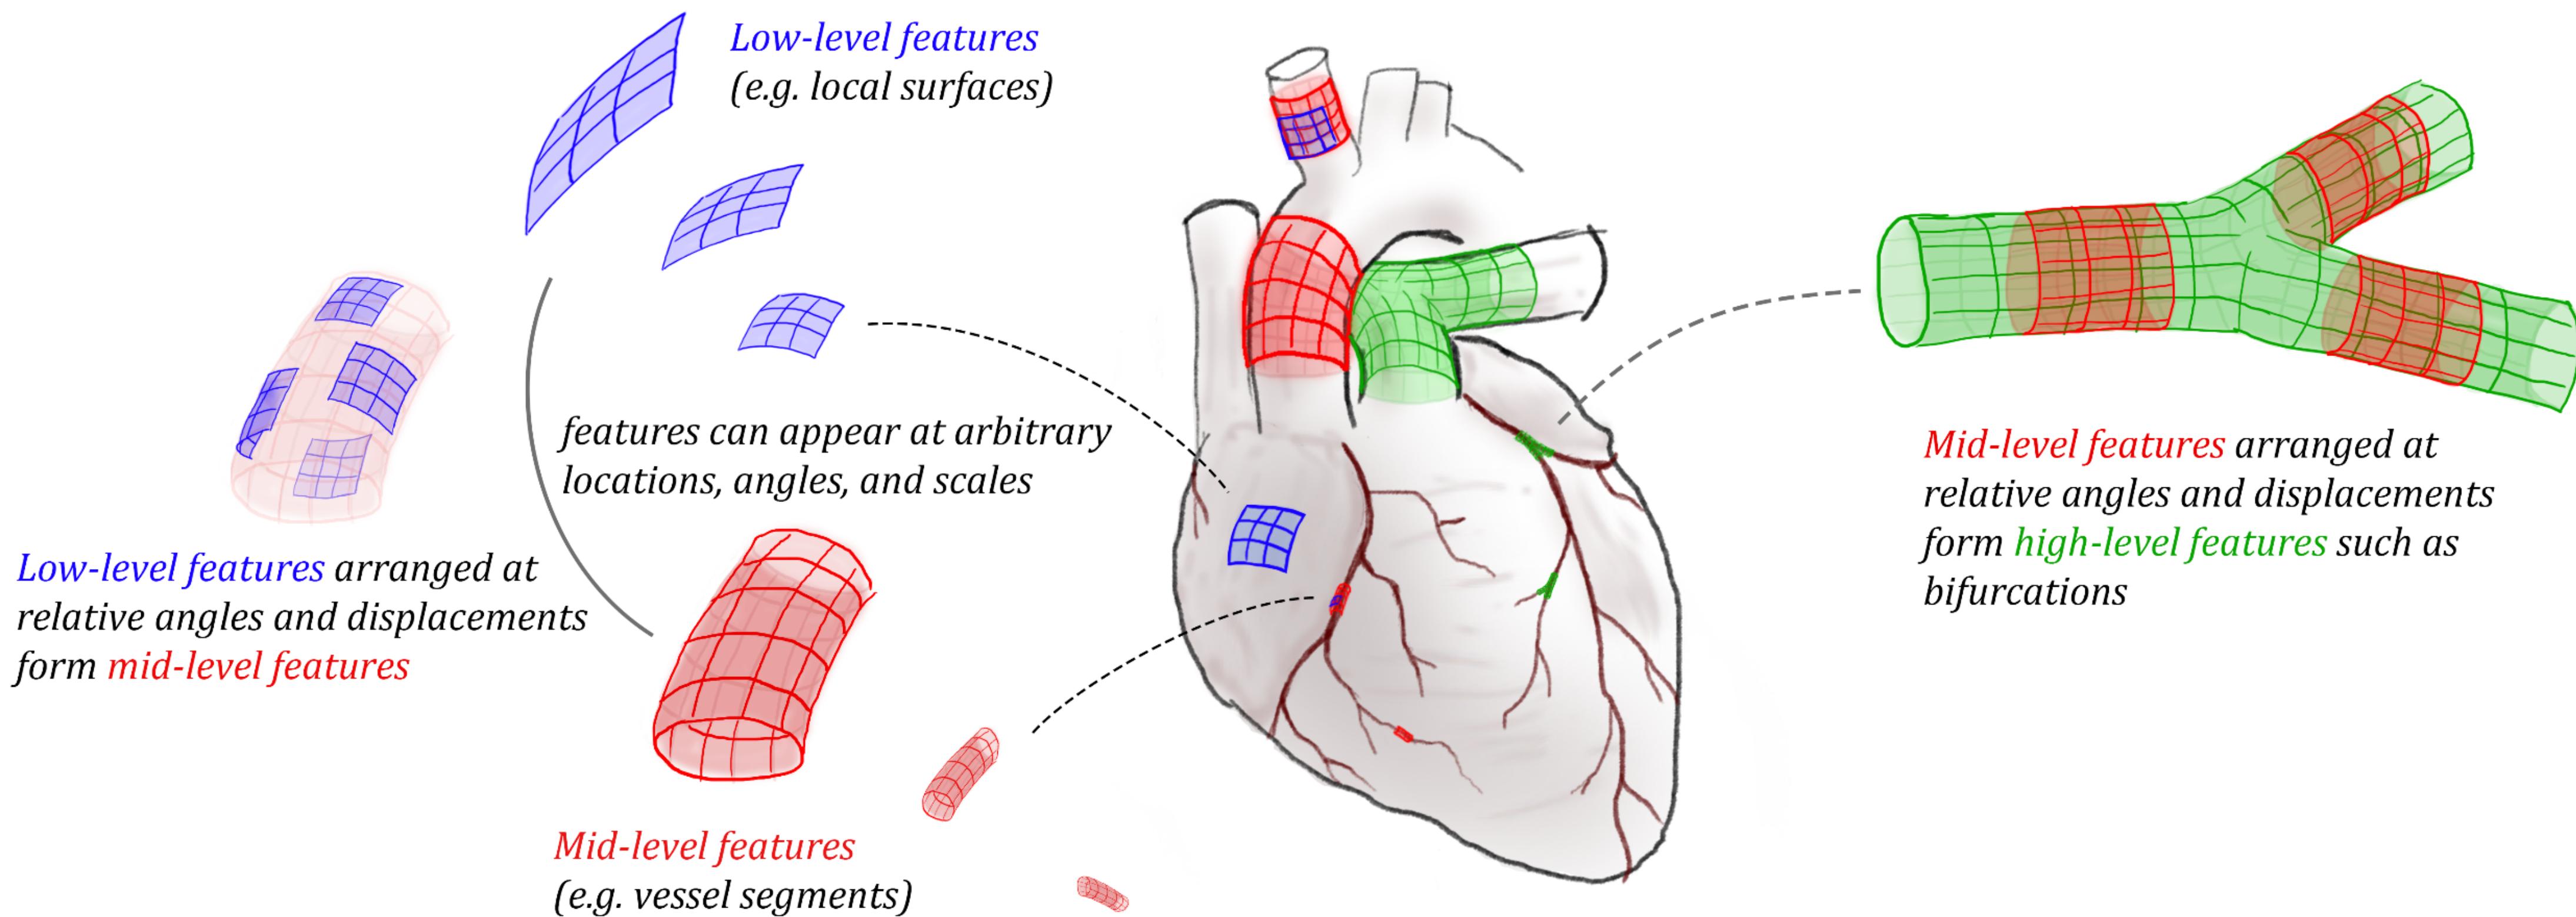
9. Equivariant graph NNs

# What is a group?

A group  $(G, \cdot)$  is a **set of elements**  $G$  equipped with a **group product**  $\cdot$ , a binary operator, that satisfies the following four axioms:

- **Closure**: Given two elements  $g$  and  $h$  of  $G$ , the product  $g \cdot h$  is also in  $G$ .
- **Associativity**: For  $g, h, i \in G$  the product  $\cdot$  is associative, i.e.,  $g \cdot (h \cdot i) = (g \cdot h) \cdot i$ .
- **Identity element**: There exists an identity element  $e \in G$  such that  $e \cdot g = g \cdot e = g$  for any  $g \in G$ .
- **Inverse element**: For each  $g \in G$  there exists an inverse element  $g^{-1} \in G$  s.t.  
$$g^{-1} \cdot g = g \cdot g^{-1} = e.$$

# Psychology of vision: recognition by components

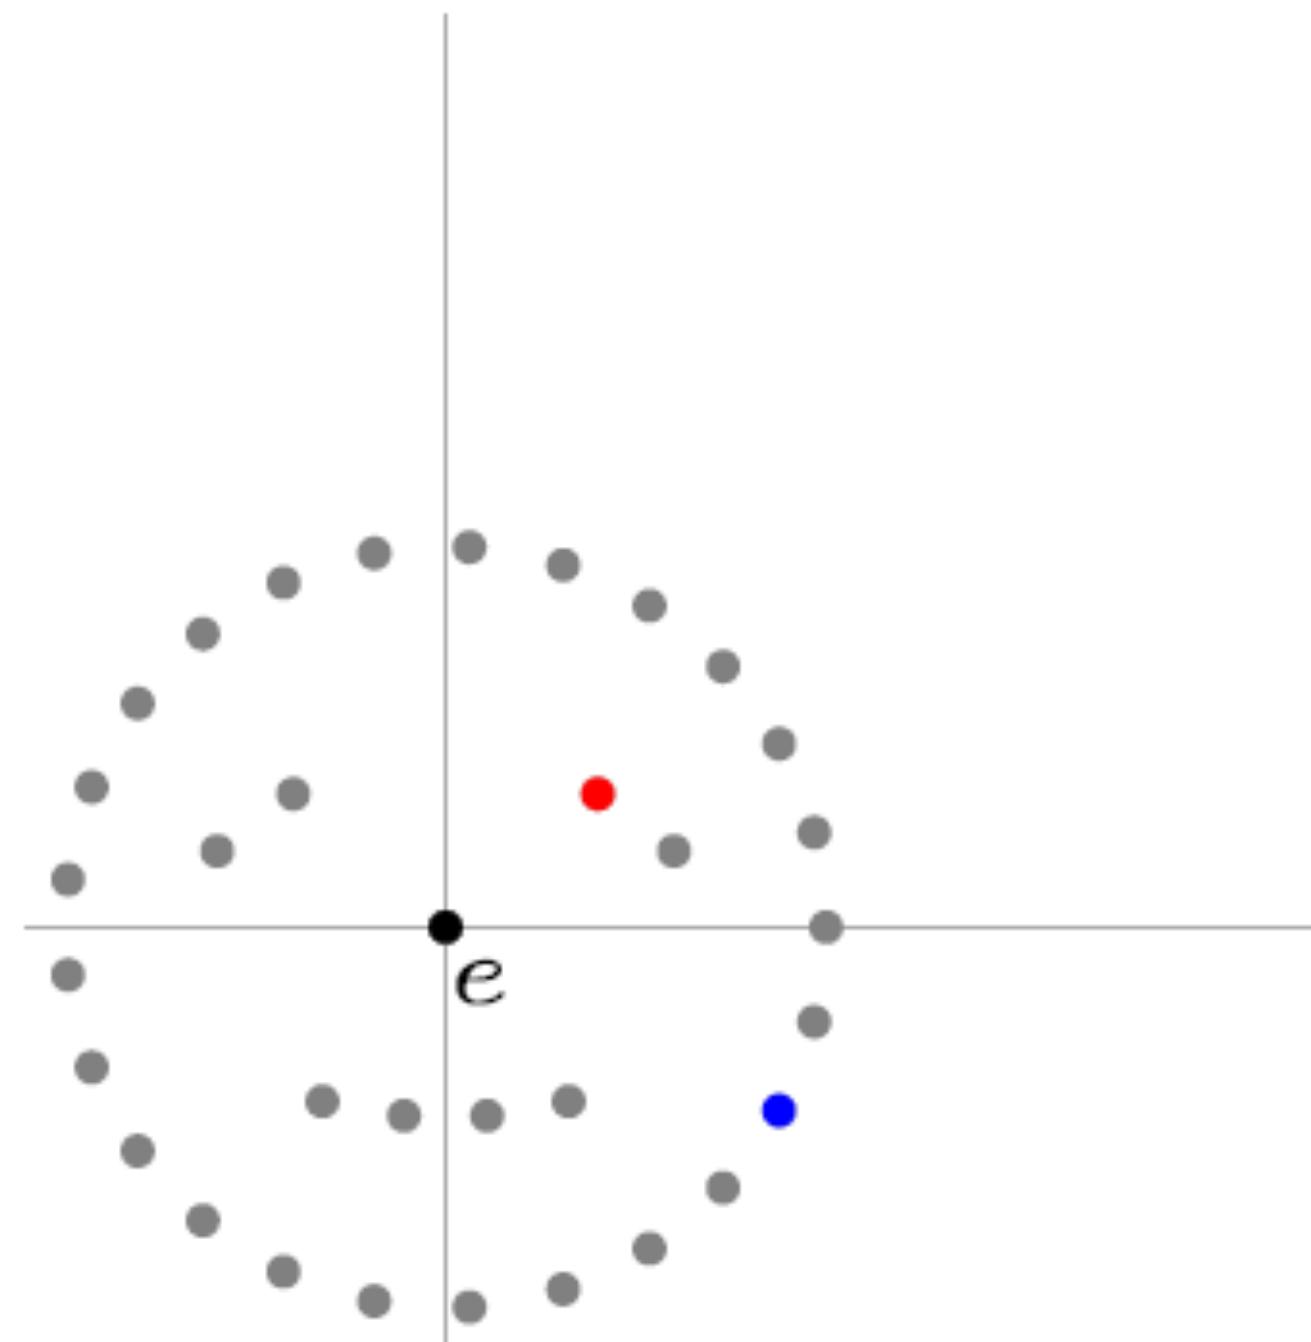


# Translation group $(\mathbb{R}^2, +)$

The translation group consists of all possible translations in  $\mathbb{R}^2$  and is equipped with the **group product** and **group inverse**:

$$\begin{aligned}g \cdot g' &= (\mathbf{x} + \mathbf{x}') \\g^{-1} &= (-\mathbf{x})\end{aligned}$$

with  $g = (\mathbf{x})$ ,  $g' = (\mathbf{x}')$  and  $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^2$ .

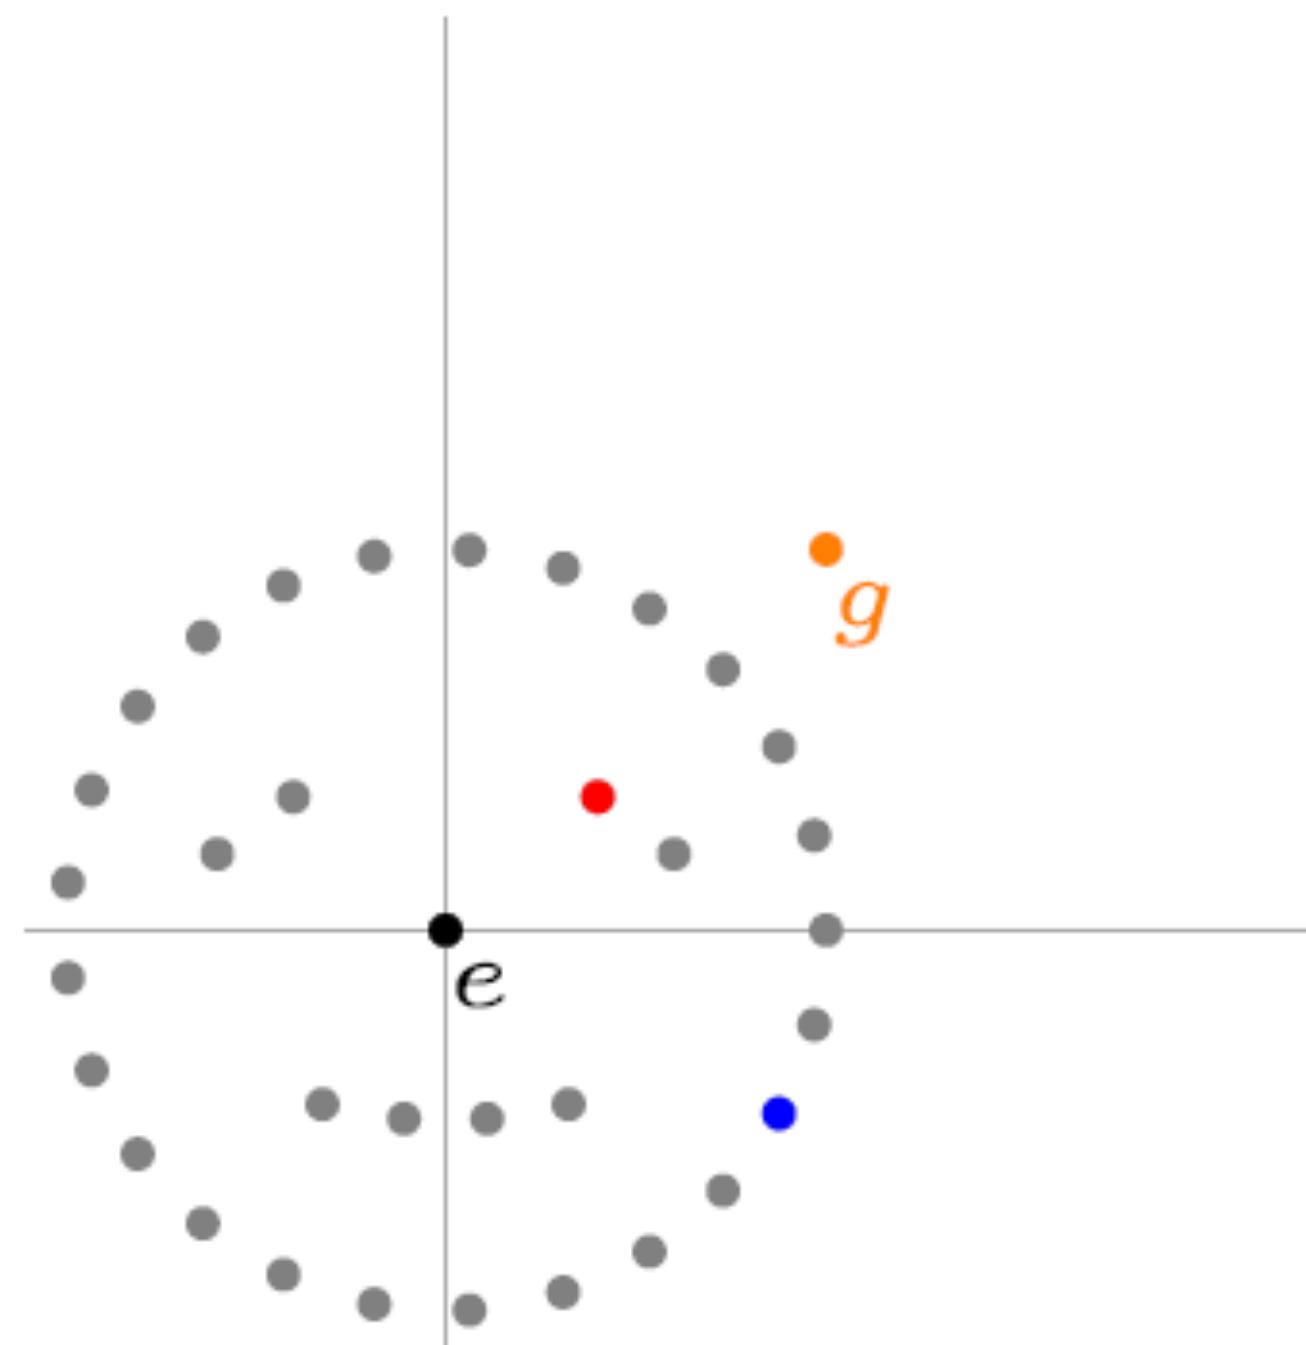


# Translation group $(\mathbb{R}^2, +)$

The translation group consists of all possible translations in  $\mathbb{R}^2$  and is equipped with the **group product** and **group inverse**:

$$\begin{aligned}g \cdot g' &= (\mathbf{x} + \mathbf{x}') \\g^{-1} &= (-\mathbf{x})\end{aligned}$$

with  $g = (\mathbf{x})$ ,  $g' = (\mathbf{x}')$  and  $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^2$ .

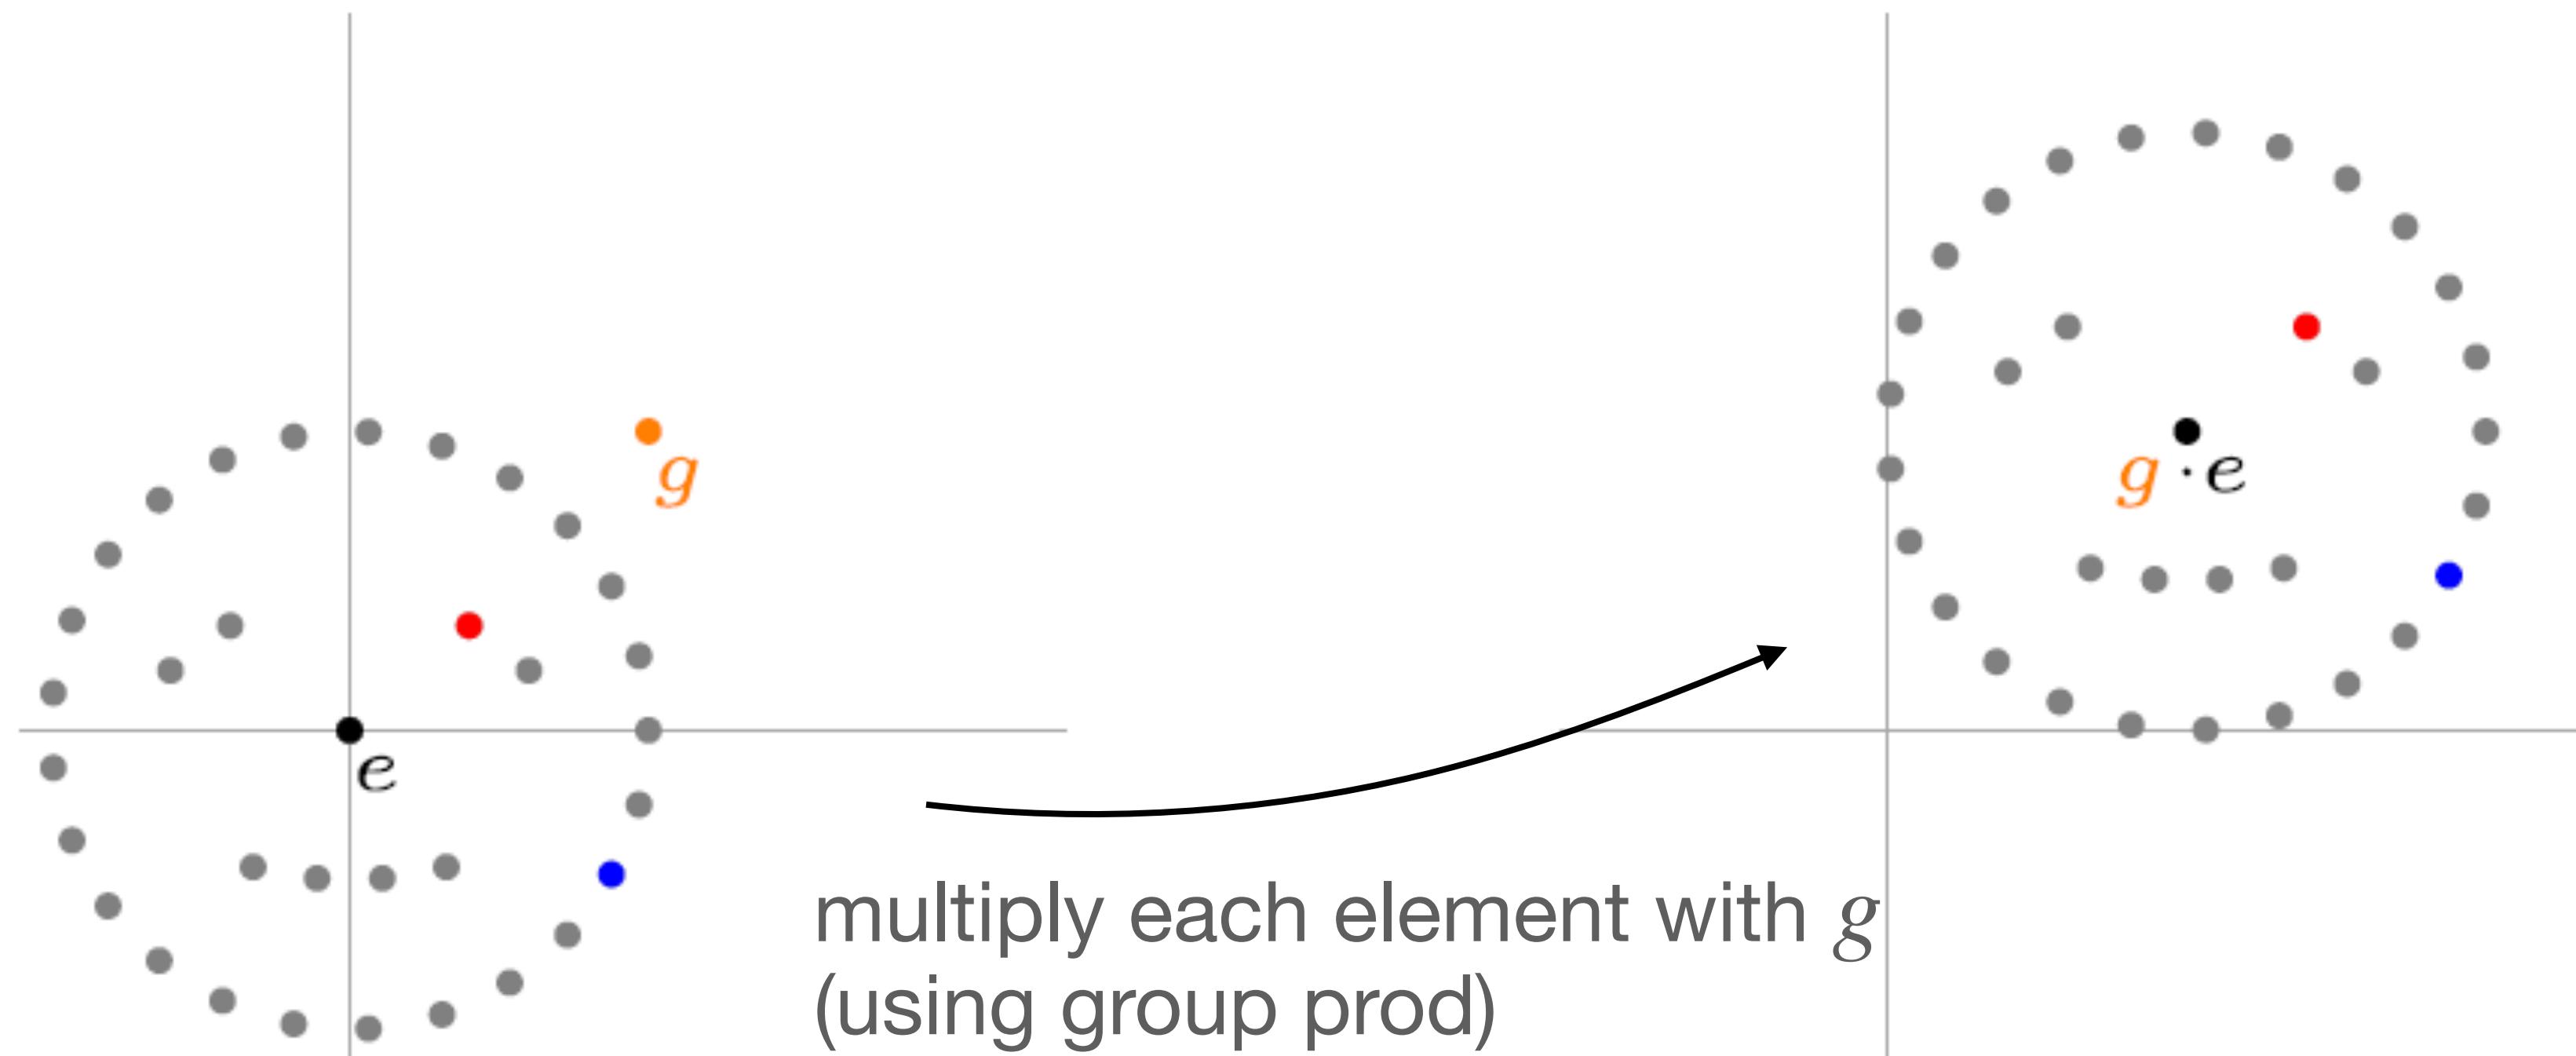


# Translation group $(\mathbb{R}^2, +)$

The translation group consists of all possible translations in  $\mathbb{R}^2$  and is equipped with the **group product** and **group inverse**:

$$\begin{aligned}g \cdot g' &= (\mathbf{x} + \mathbf{x}') \\g^{-1} &= (-\mathbf{x})\end{aligned}$$

with  $g = (\mathbf{x})$ ,  $g' = (\mathbf{x}')$  and  $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^2$ .



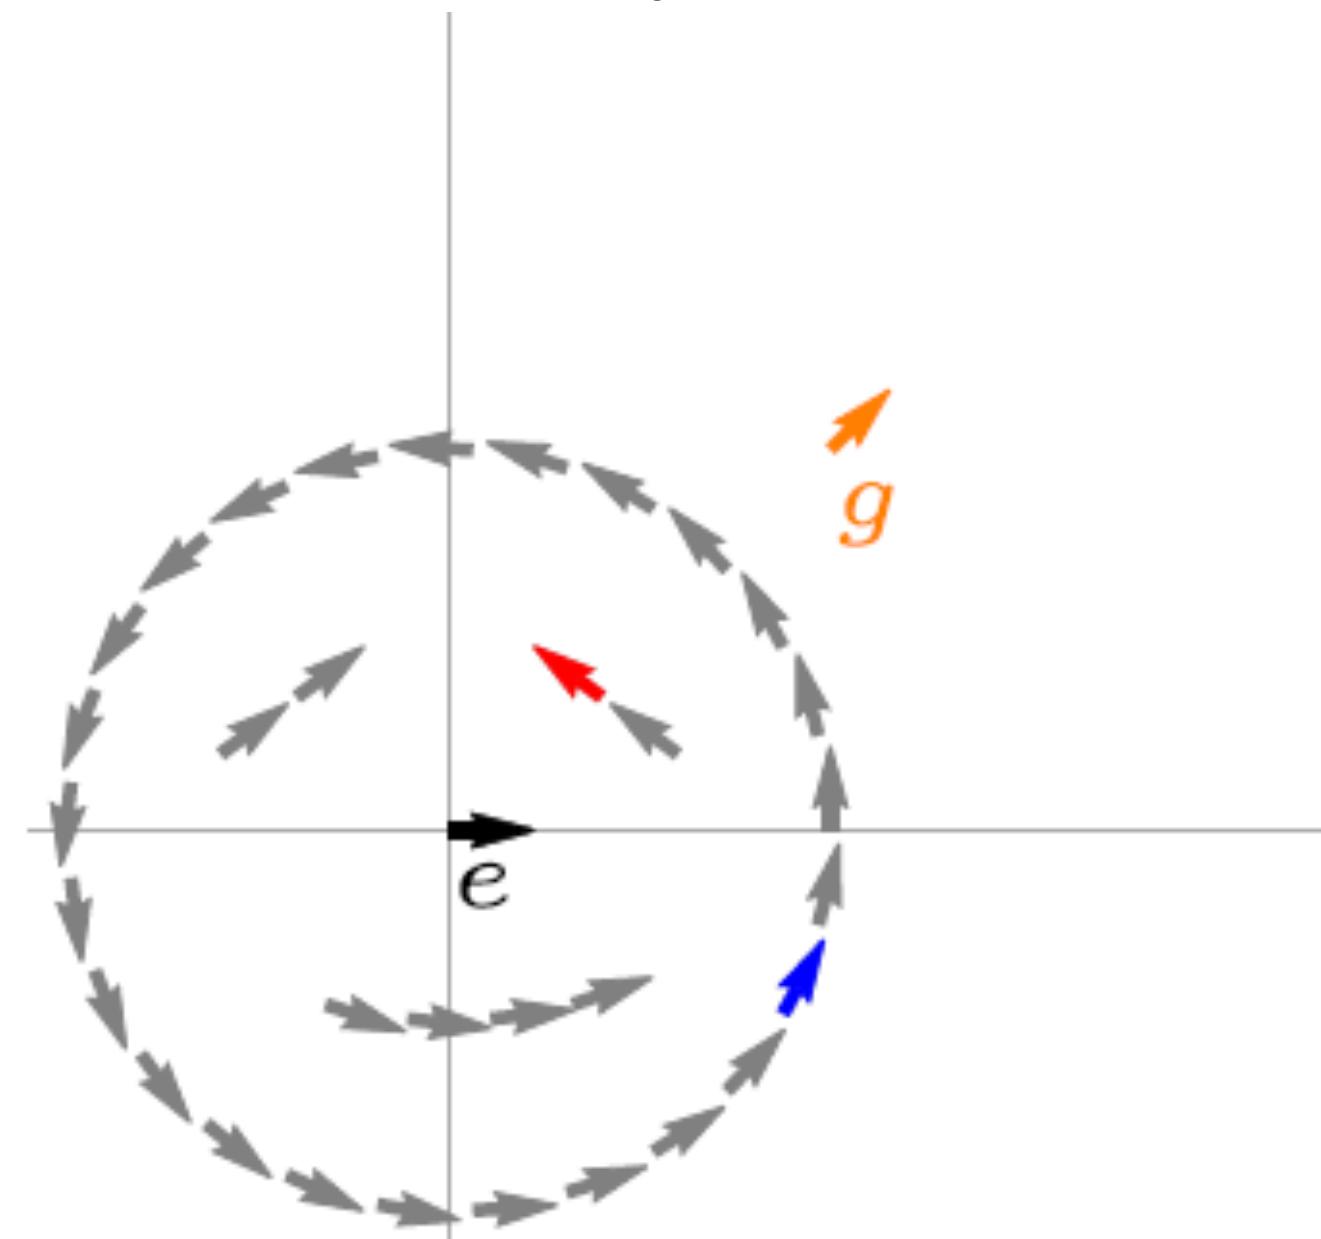
# Roto-translation group $SE(2)$

2D Special Euclidean motion group

The group  $SE(2) = \mathbb{R}^2 \rtimes SO(2)$  consists of the **coupled** space  $\mathbb{R}^2 \times S^1$  of translations vectors in  $\mathbb{R}^2$ , and rotations in  $SO(2)$  (or equivalently orientations in  $S^1$ ), and is equipped with the group product and group inverse:

$$\begin{aligned}g \cdot g' &= (\mathbf{x}, \mathbf{R}_\theta) \cdot (\mathbf{x}', \mathbf{R}_{\theta'}) = (\mathbf{R}_\theta \mathbf{x}' + \mathbf{x}, \mathbf{R}_{\theta+\theta'}) \\g^{-1} &= (-\mathbf{R}_\theta^{-1} \mathbf{x}, \mathbf{R}_\theta^{-1})\end{aligned}$$

with  $g = (\mathbf{x}, \mathbf{R}_\theta)$ ,  $g' = (\mathbf{x}', \mathbf{R}_{\theta'})$ .



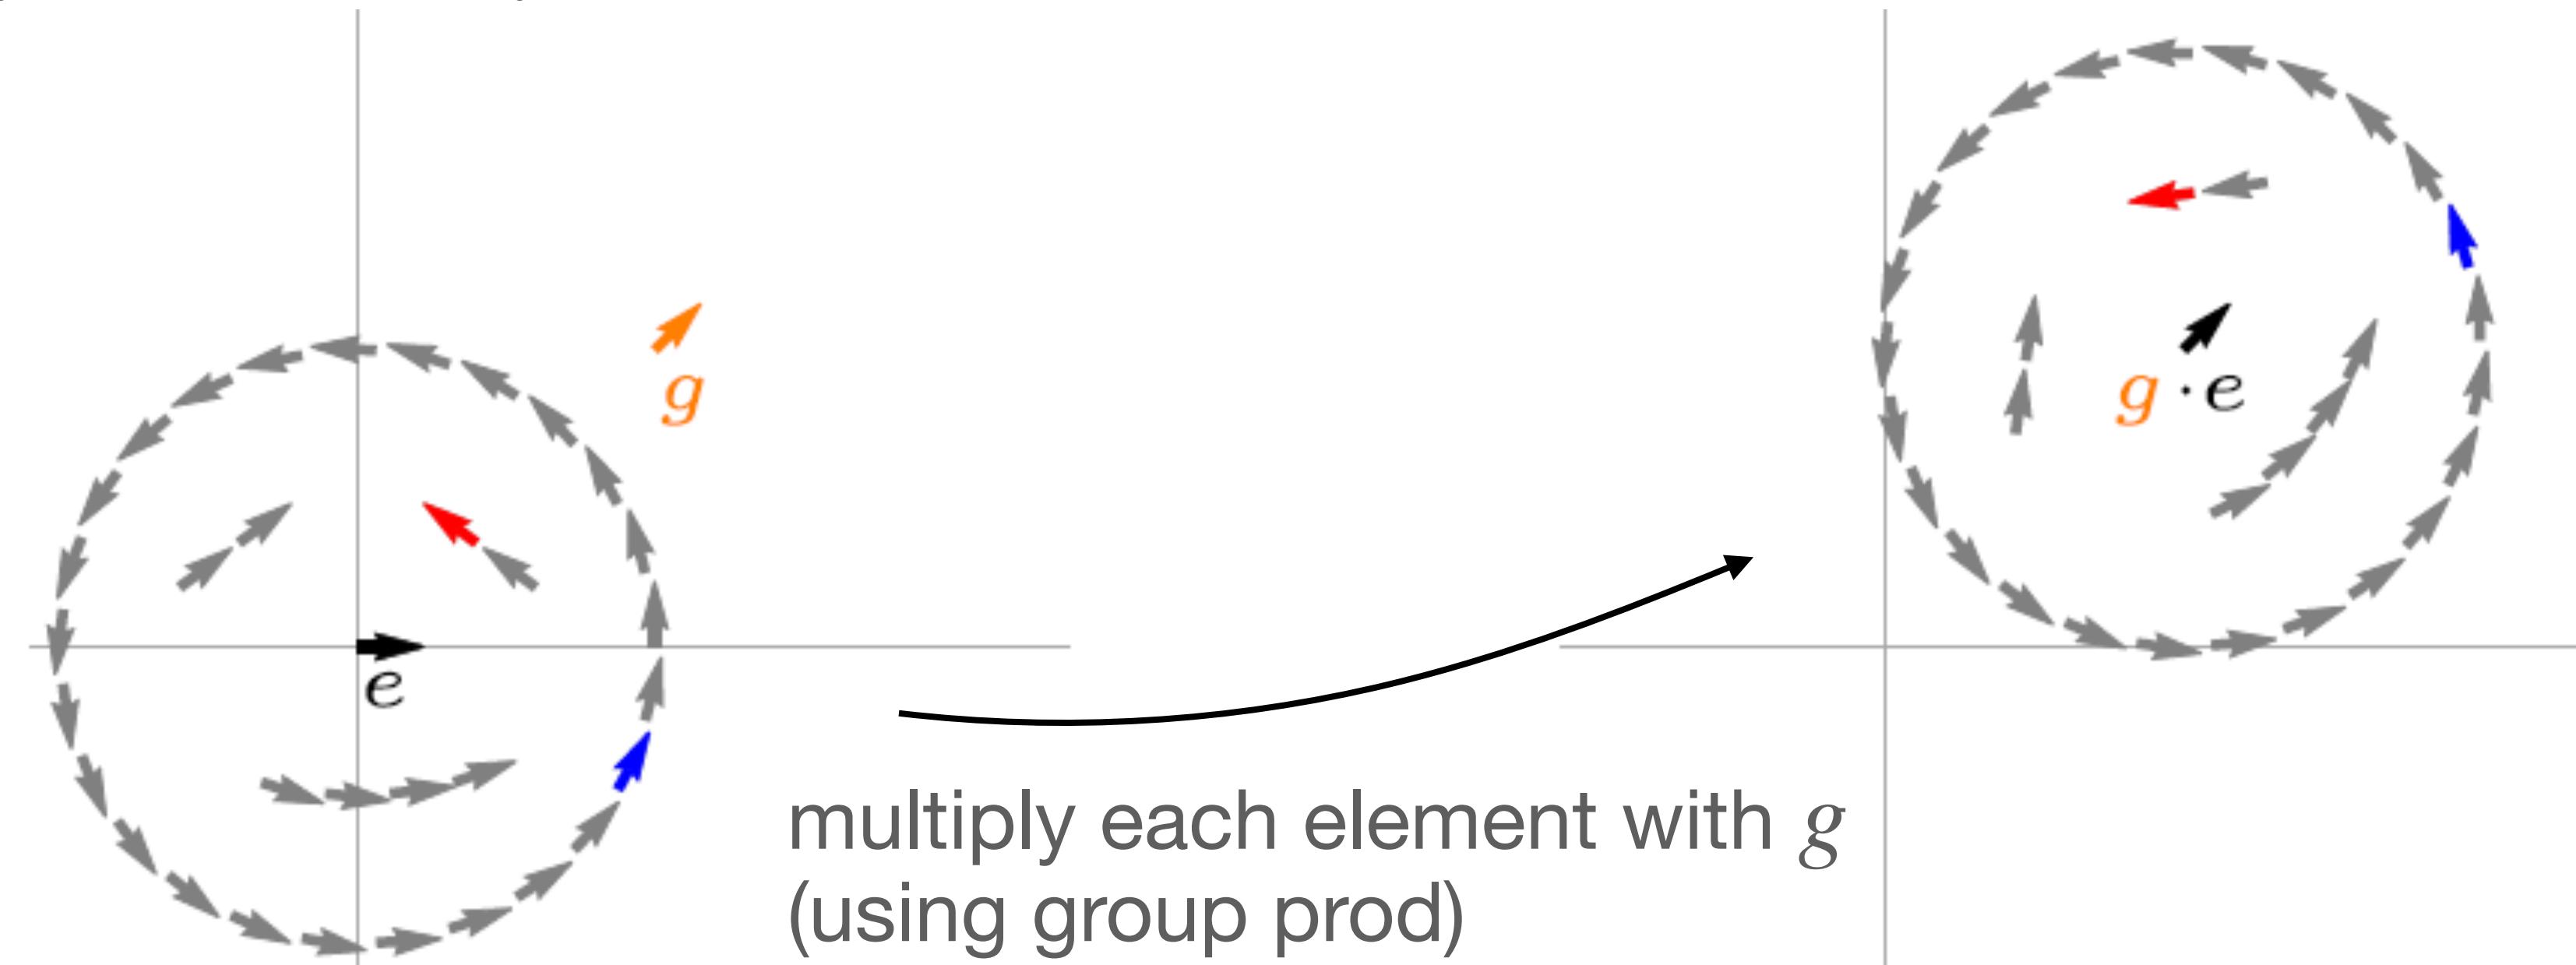
# Roto-translation group $SE(2)$

2D Special Euclidean motion group

The group  $SE(2) = \mathbb{R}^2 \times SO(2)$  consists of the **coupled** space  $\mathbb{R}^2 \times S^1$  of translations vectors in  $\mathbb{R}^2$ , and rotations in  $SO(2)$  (or equivalently orientations in  $S^1$ ), and is equipped with the group product and group inverse:

$$\begin{aligned}g \cdot g' &= (\mathbf{x}, \mathbf{R}_\theta) \cdot (\mathbf{x}', \mathbf{R}_{\theta'}) = (\mathbf{R}_\theta \mathbf{x}' + \mathbf{x}, \mathbf{R}_{\theta+\theta'}) \\g^{-1} &= (-\mathbf{R}_\theta^{-1} \mathbf{x}, \mathbf{R}_\theta^{-1})\end{aligned}$$

with  $g = (\mathbf{x}, \mathbf{R}_\theta)$ ,  $g' = (\mathbf{x}', \mathbf{R}_{\theta'})$ .



# Roto-translation group $SE(2)$

2D Special Euclidean motion group

**Matrix representation:** The group can also be represented by matrices

$$g = (\mathbf{x}, \mathbf{R}_\theta) \quad \leftrightarrow \quad \mathbf{G} = \begin{pmatrix} \cos \theta & -\sin \theta & x \\ \sin \theta & \cos \theta & y \\ 0 & 0 & 1 \end{pmatrix} = \begin{pmatrix} \mathbf{R}_\theta & \mathbf{x} \\ \mathbf{0}^T & 1 \end{pmatrix}$$

with the group product and inverse simply given by the matrix product and matrix inverse.

In parametric form:

$$(\mathbf{x}, \theta) \cdot (\mathbf{x}', \theta') = (\mathbf{R}_\theta \mathbf{x}' + \mathbf{x}, \theta + \theta' \bmod 2\pi)$$

$\leftrightarrow$

In matrix form:

$$\begin{pmatrix} \mathbf{R}_\theta & \mathbf{x} \\ \mathbf{0}^T & 1 \end{pmatrix} \begin{pmatrix} \mathbf{R}'_{\theta'} & \mathbf{x}' \\ \mathbf{0}^T & 1 \end{pmatrix} = \begin{pmatrix} \mathbf{R}_{\theta+\theta'} & \mathbf{R}_\theta \mathbf{x}' + \mathbf{x} \\ \mathbf{0}^T & 1 \end{pmatrix}$$

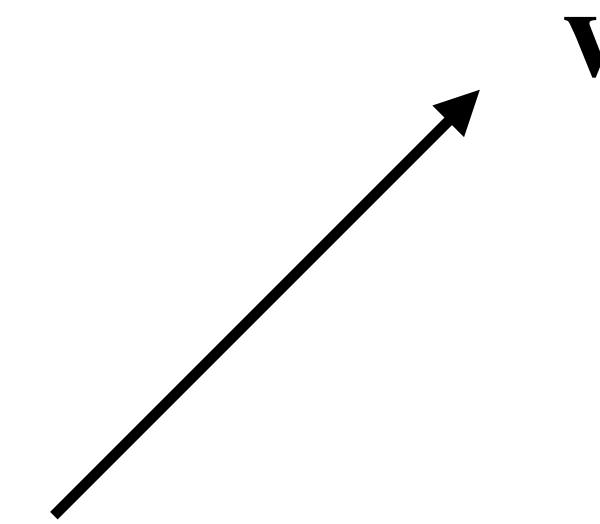
# Representations

A **representation**  $\rho : G \rightarrow GL(V)$  is a group homomorphism from  $G$  to the general linear group  $GL(V)$ .

That is  $\rho(g)$  is a linear transformation that is **parameterized by group elements**  $g \in G$  that transforms some vector  $\mathbf{v} \in V$  (e.g. an image) such that

$$\rho(g') \circ \rho(g)[\mathbf{v}] = \rho(g' \cdot g)[\mathbf{v}]$$

# Representations

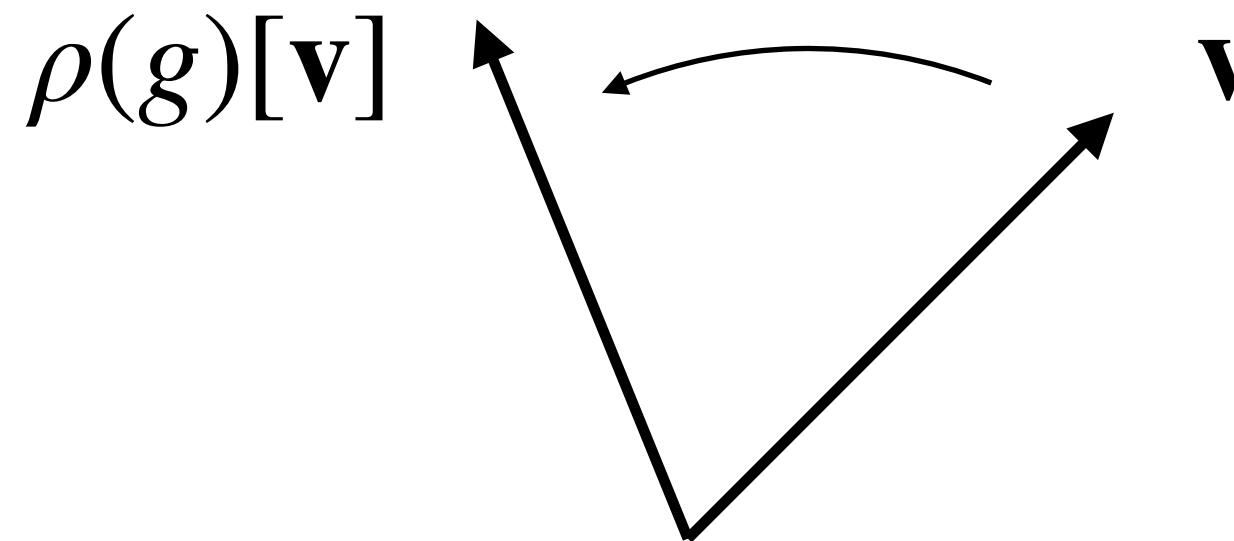


A **representation**  $\rho : G \rightarrow GL(V)$  is a group homomorphism from  $G$  to the general linear group  $GL(V)$ .

That is  $\rho(g)$  is a linear transformation that is **parameterized by group elements**  $g \in G$  that transforms some vector  $\mathbf{v} \in V$  (e.g. an image) such that

$$\rho(g') \circ \rho(g)[\mathbf{v}] = \rho(g' \cdot g)[\mathbf{v}]$$

# Representations

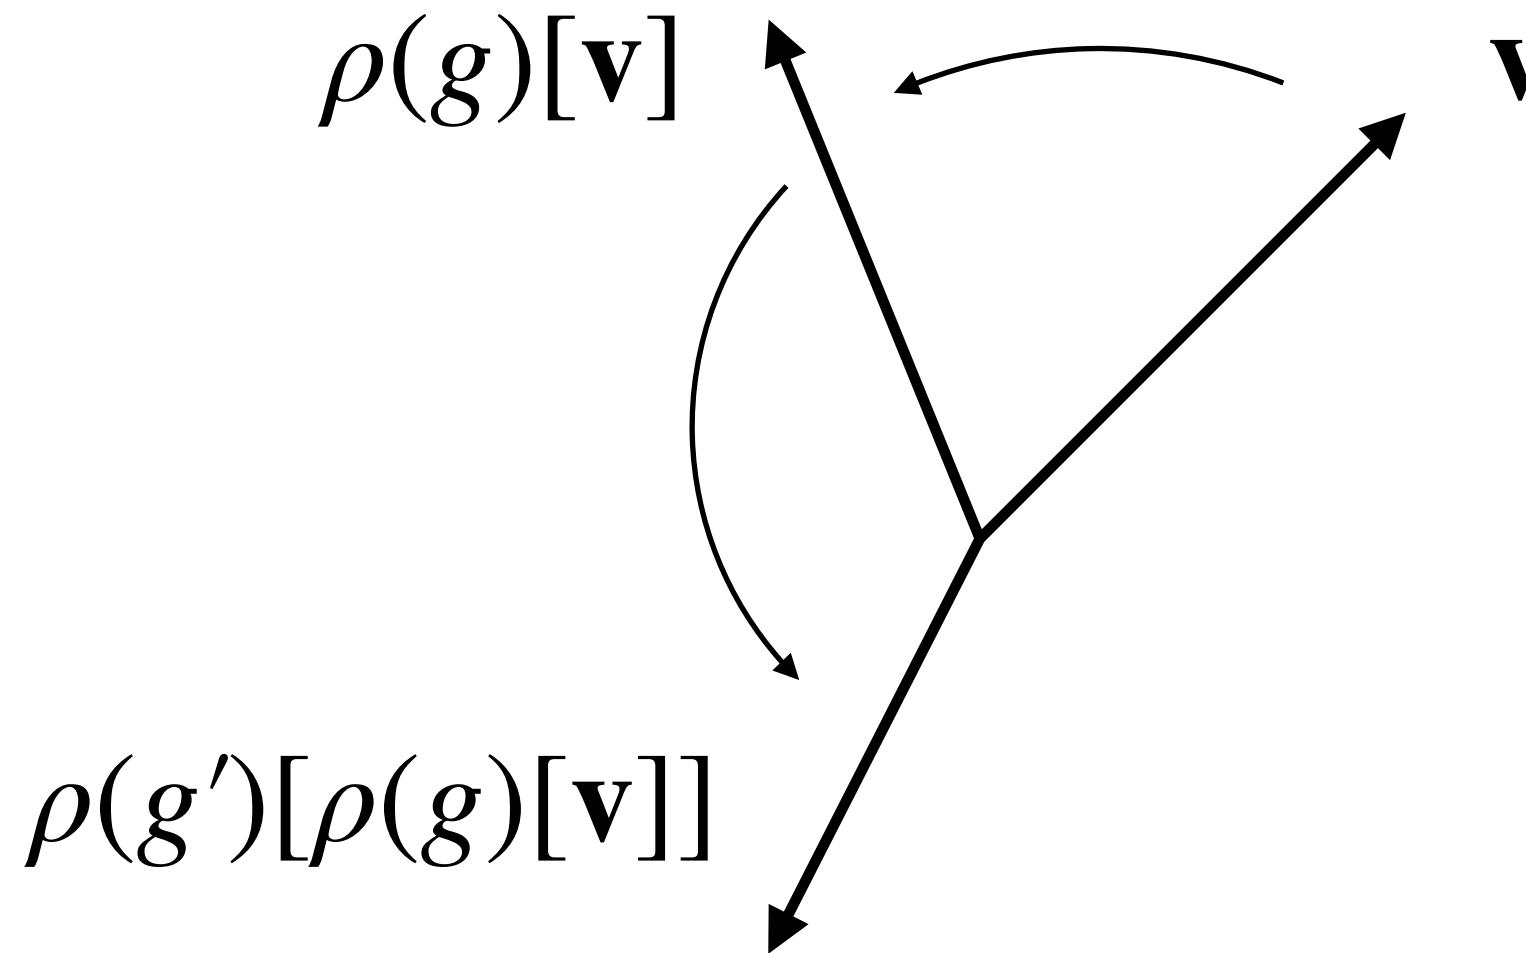


A **representation**  $\rho : G \rightarrow GL(V)$  is a group homomorphism from  $G$  to the general linear group  $GL(V)$ .

That is  $\rho(g)$  is a linear transformation that is **parameterized by group elements**  $g \in G$  that transforms some vector  $v \in V$  (e.g. an image) such that

$$\rho(g') \circ \rho(g)[v] = \rho(g' \cdot g)[v]$$

# Representations

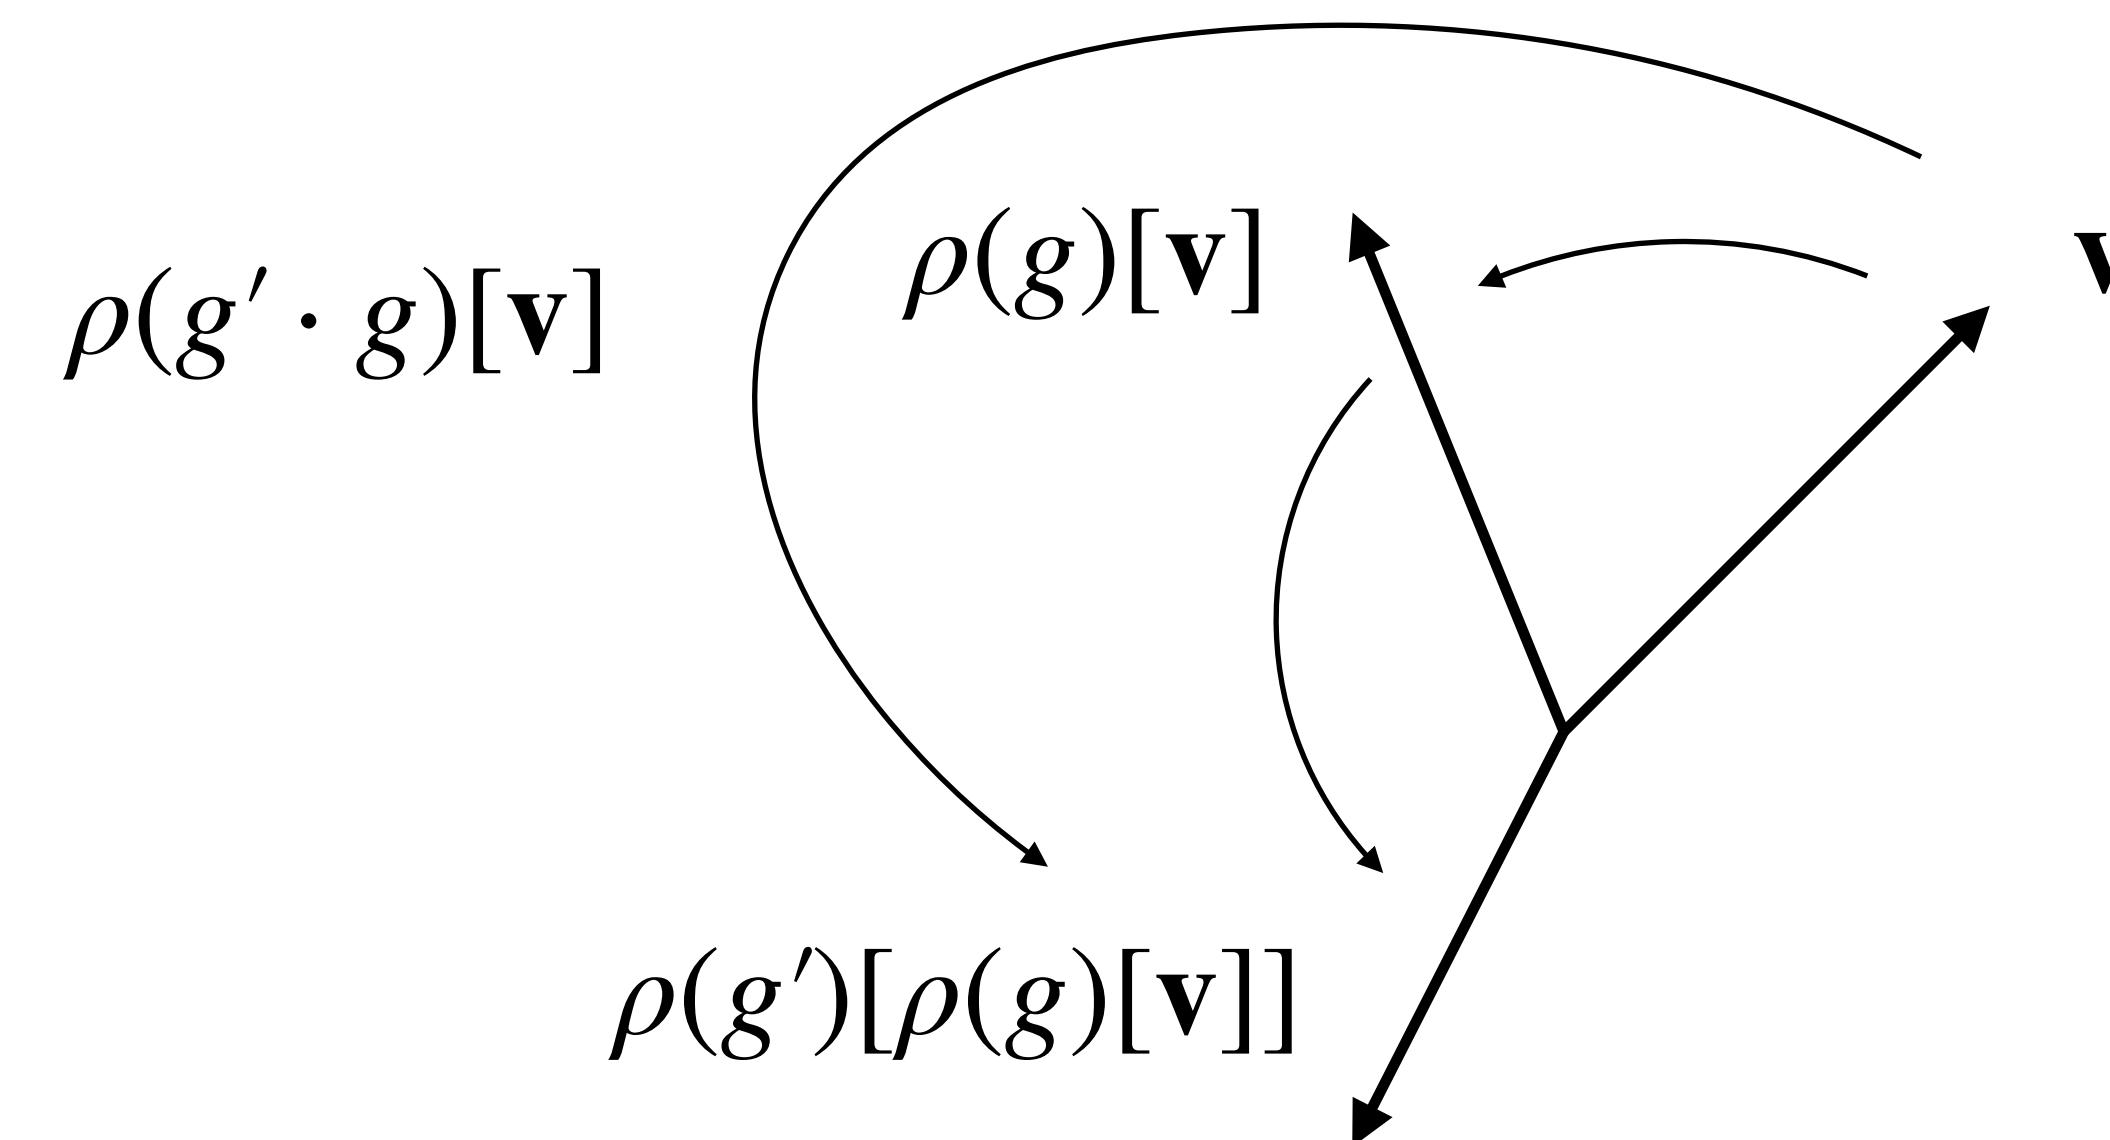


A **representation**  $\rho : G \rightarrow GL(V)$  is a group homomorphism from  $G$  to the general linear group  $GL(V)$ .

That is  $\rho(g)$  is a linear transformation that is **parameterized by group elements**  $g \in G$  that transforms some vector  $v \in V$  (e.g. an image) such that

$$\rho(g') \circ \rho(g)[v] = \rho(g' \cdot g)[v]$$

# Representations



A **representation**  $\rho : G \rightarrow GL(V)$  is a group homomorphism from  $G$  to the general linear group  $GL(V)$ .

That is  $\rho(g)$  is a linear transformation that is **parameterized by group elements**  $g \in G$  that transforms some vector  $\mathbf{v} \in V$  (e.g. an image) such that

$$\rho(g') \circ \rho(g)[\mathbf{v}] = \rho(g' \cdot g)[\mathbf{v}]$$

# Left-regular Representations

A **left-regular representation**  $\mathcal{L}_g$  is a representation that transforms functions  $f$  by transforming their domains via the inverse group action

$$\mathcal{L}_g[f](x) := f(g^{-1} \cdot x)$$

“group action” equals  
group product when  
domain is  $G$

# Left-regular Representations

Example:

$$f \in \mathbb{L}_2(\mathbb{R}^2)$$

- a 2D image

$$G = SE(2)$$

- the roto-translation group

$$\mathcal{L}_g(f)(\mathbf{y}) = f(\mathbf{R}_\theta^{-1}(\mathbf{y} - \mathbf{x}))$$

- a roto-translation of the image



A **left-regular representation**  $\mathcal{L}_g$  is a representation that transforms functions  $f$  by transforming their domains via the inverse group action

$$\mathcal{L}_g[f](x) := f(g^{-1} \cdot x)$$

“group action” equals  
group product when  
domain is  $G$

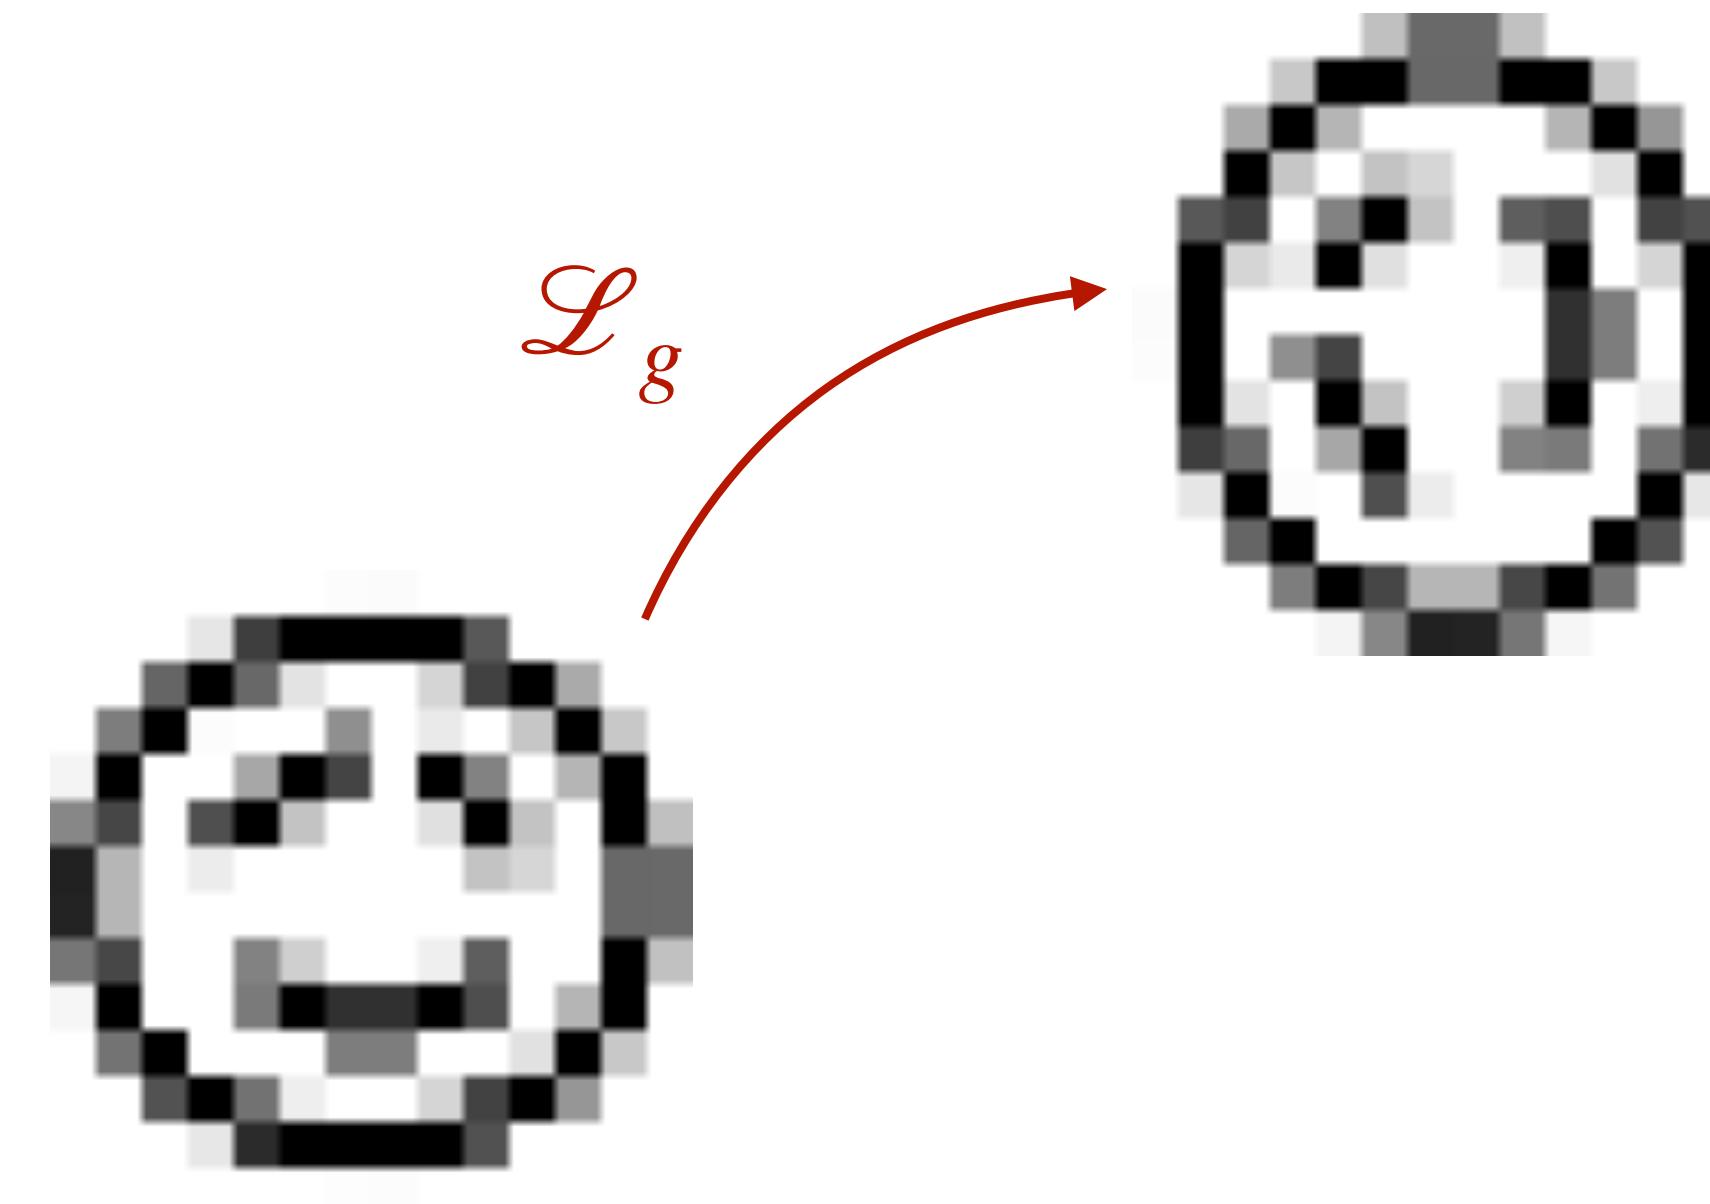
# Left-regular Representations

Example:

$f \in \mathbb{L}_2(\mathbb{R}^2)$   
- a 2D image

$G = SE(2)$   
- the roto-translation group

$\mathcal{L}_g(f)(\mathbf{y}) = f(\mathbf{R}_\theta^{-1}(\mathbf{y} - \mathbf{x}))$   
- a roto-translation of the image



A **left-regular representation**  $\mathcal{L}_g$  is a representation that transforms functions  $f$  by transforming their domains via the inverse group action

$$\mathcal{L}_g[f](x) := f(g^{-1} \cdot x)$$

“group action” equals  
group product when  
domain is  $G$

# Left-regular Representations

Example:

$$f \in \mathbb{L}_2(\mathbb{R}^2)$$

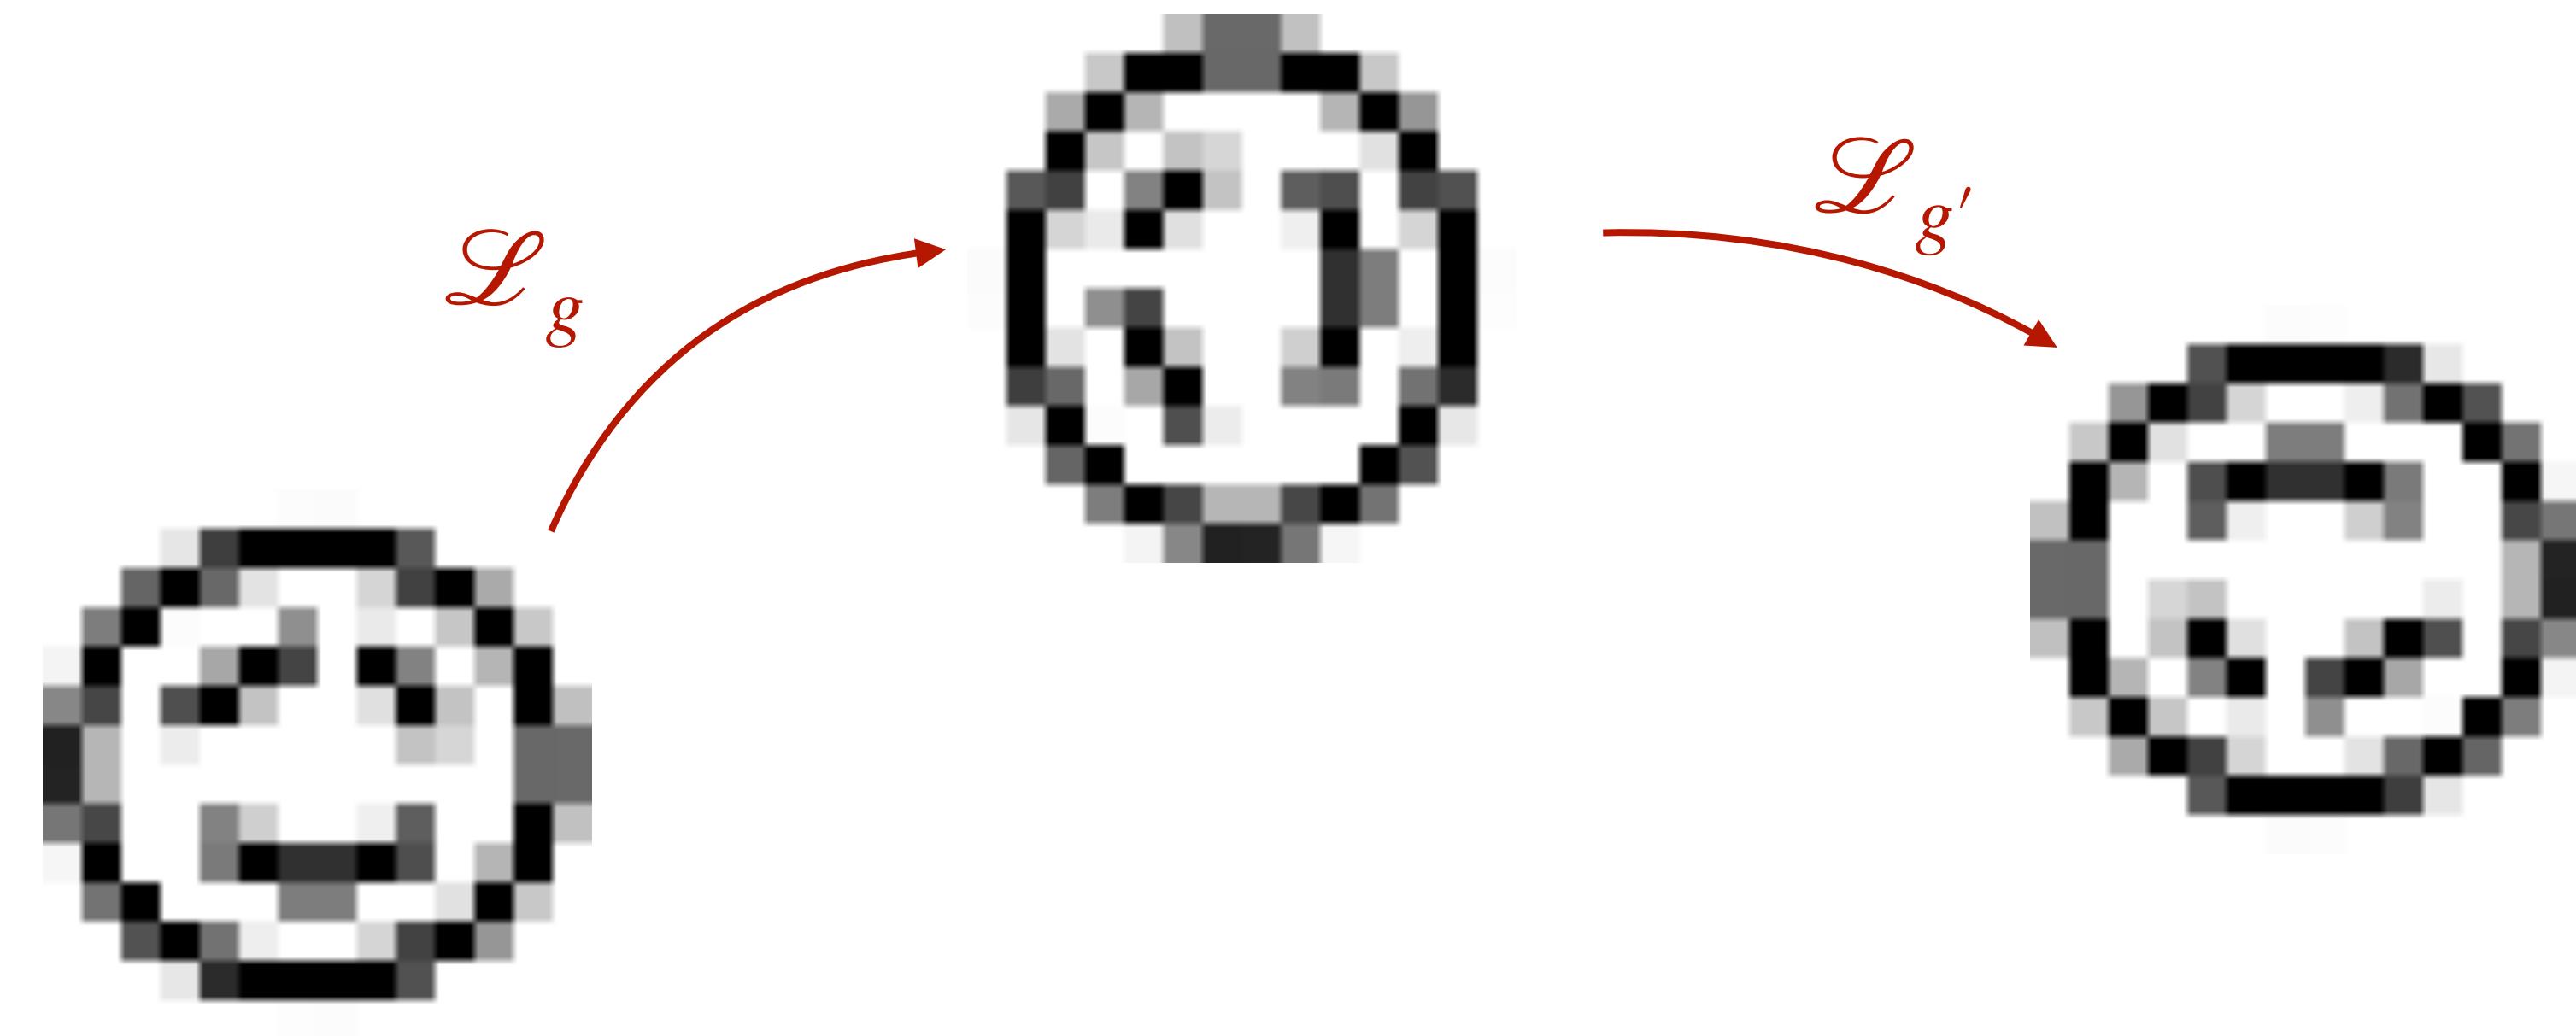
- a 2D image

$$G = SE(2)$$

- the roto-translation group

$$\mathcal{L}_g(f)(\mathbf{y}) = f(\mathbf{R}_\theta^{-1}(\mathbf{y} - \mathbf{x}))$$

- a roto-translation of the image



A **left-regular representation**  $\mathcal{L}_g$  is a representation that transforms functions  $f$  by transforming their domains via the inverse group action

$$\mathcal{L}_g[f](x) := f(g^{-1} \cdot x)$$

"group action" equals  
group product when  
domain is  $G$

# Left-regular Representations

Example:

$$f \in \mathbb{L}_2(\mathbb{R}^2)$$

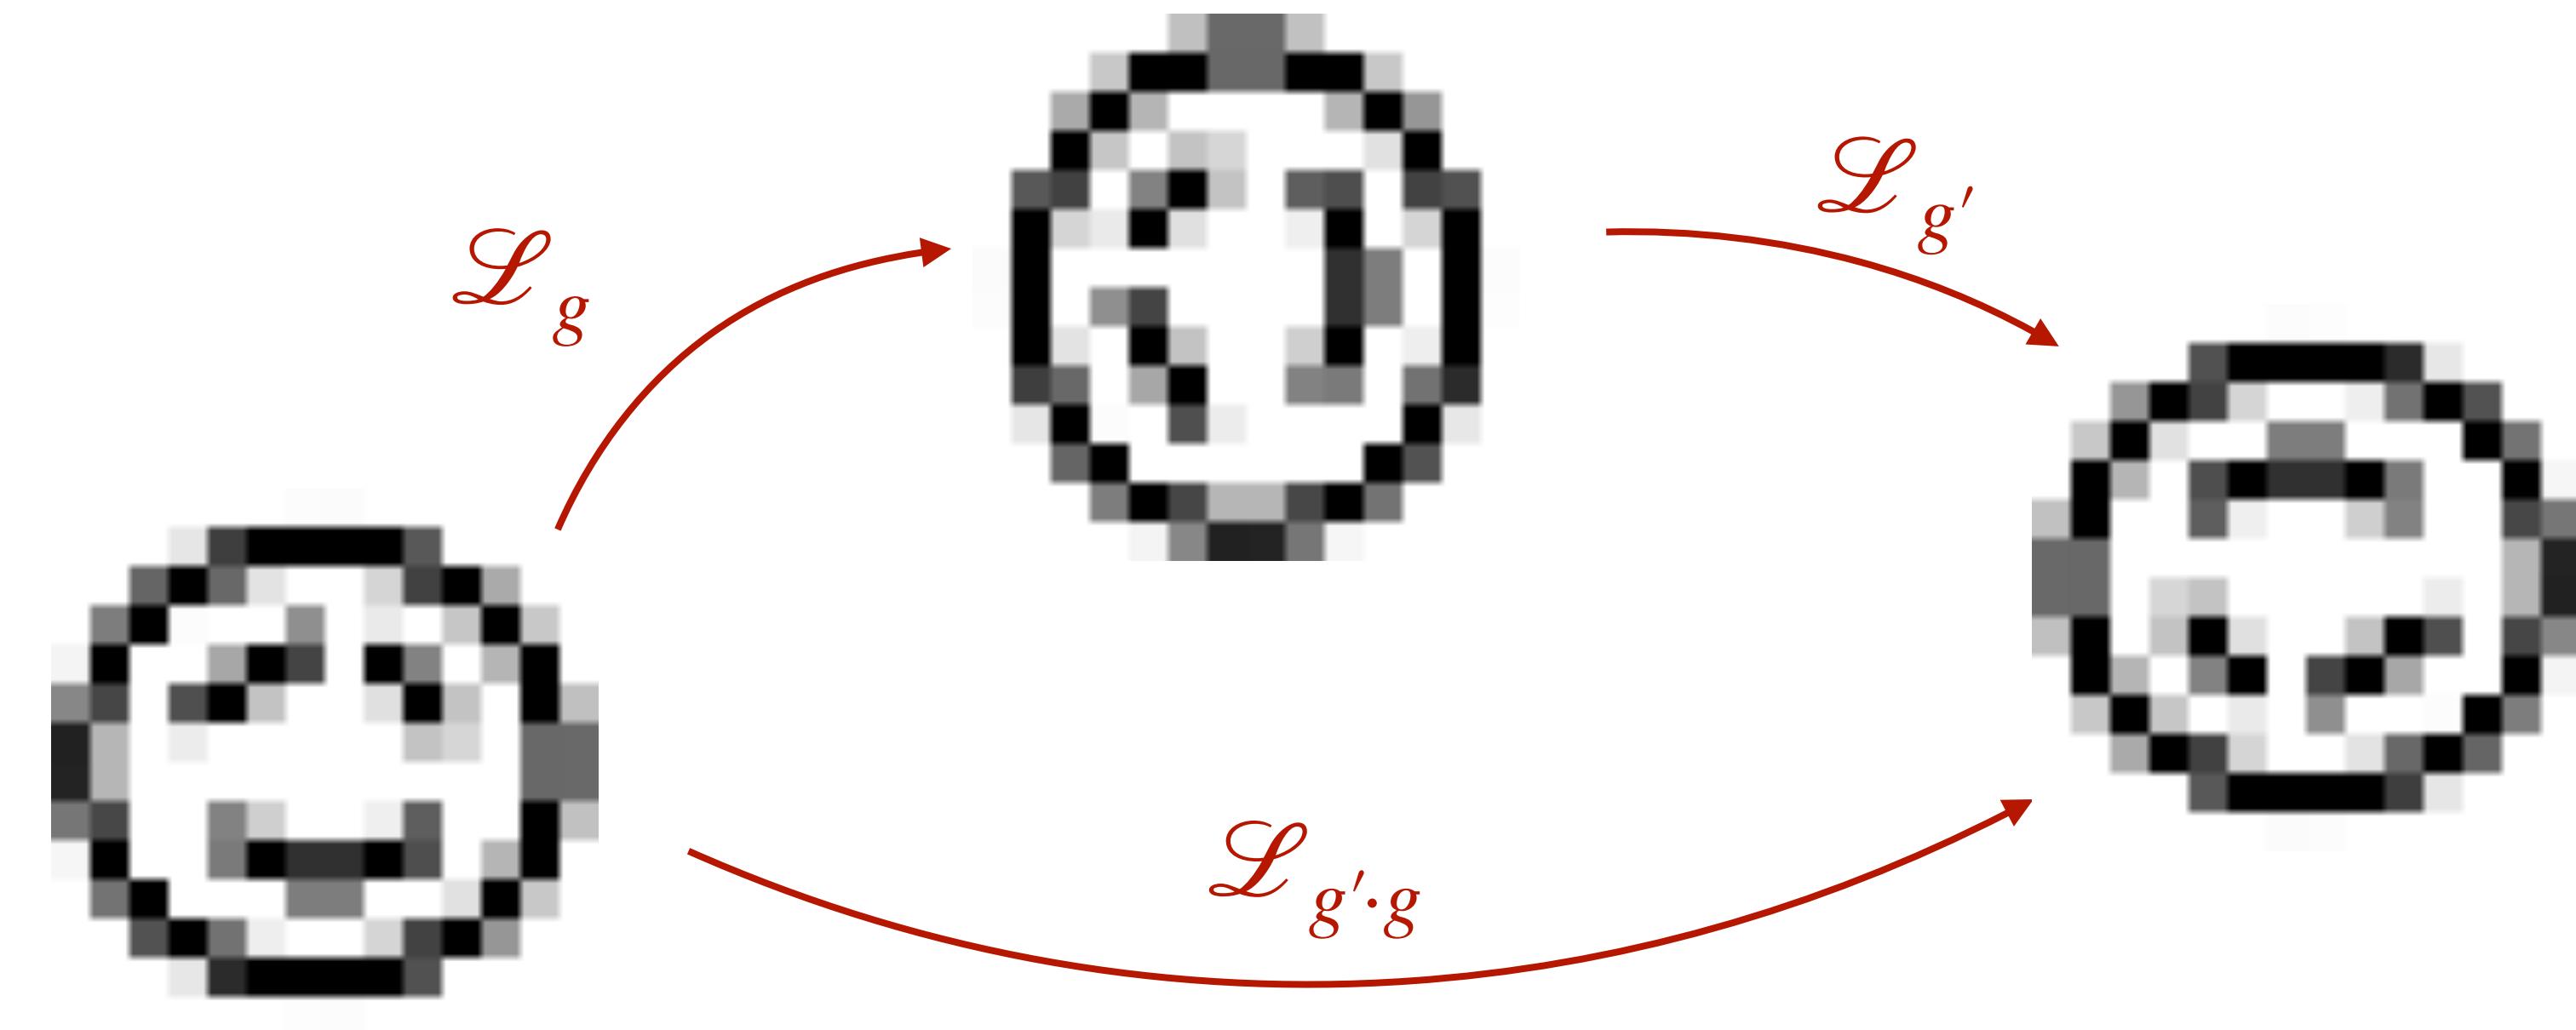
- a 2D image

$$G = SE(2)$$

- the roto-translation group

$$\mathcal{L}_g(f)(\mathbf{y}) = f(\mathbf{R}_\theta^{-1}(\mathbf{y} - \mathbf{x}))$$

- a roto-translation of the image



A **left-regular representation**  $\mathcal{L}_g$  is a representation that transforms functions  $f$  by transforming their domains via the inverse group action

$$\mathcal{L}_g[f](x) := f(g^{-1} \cdot x)$$

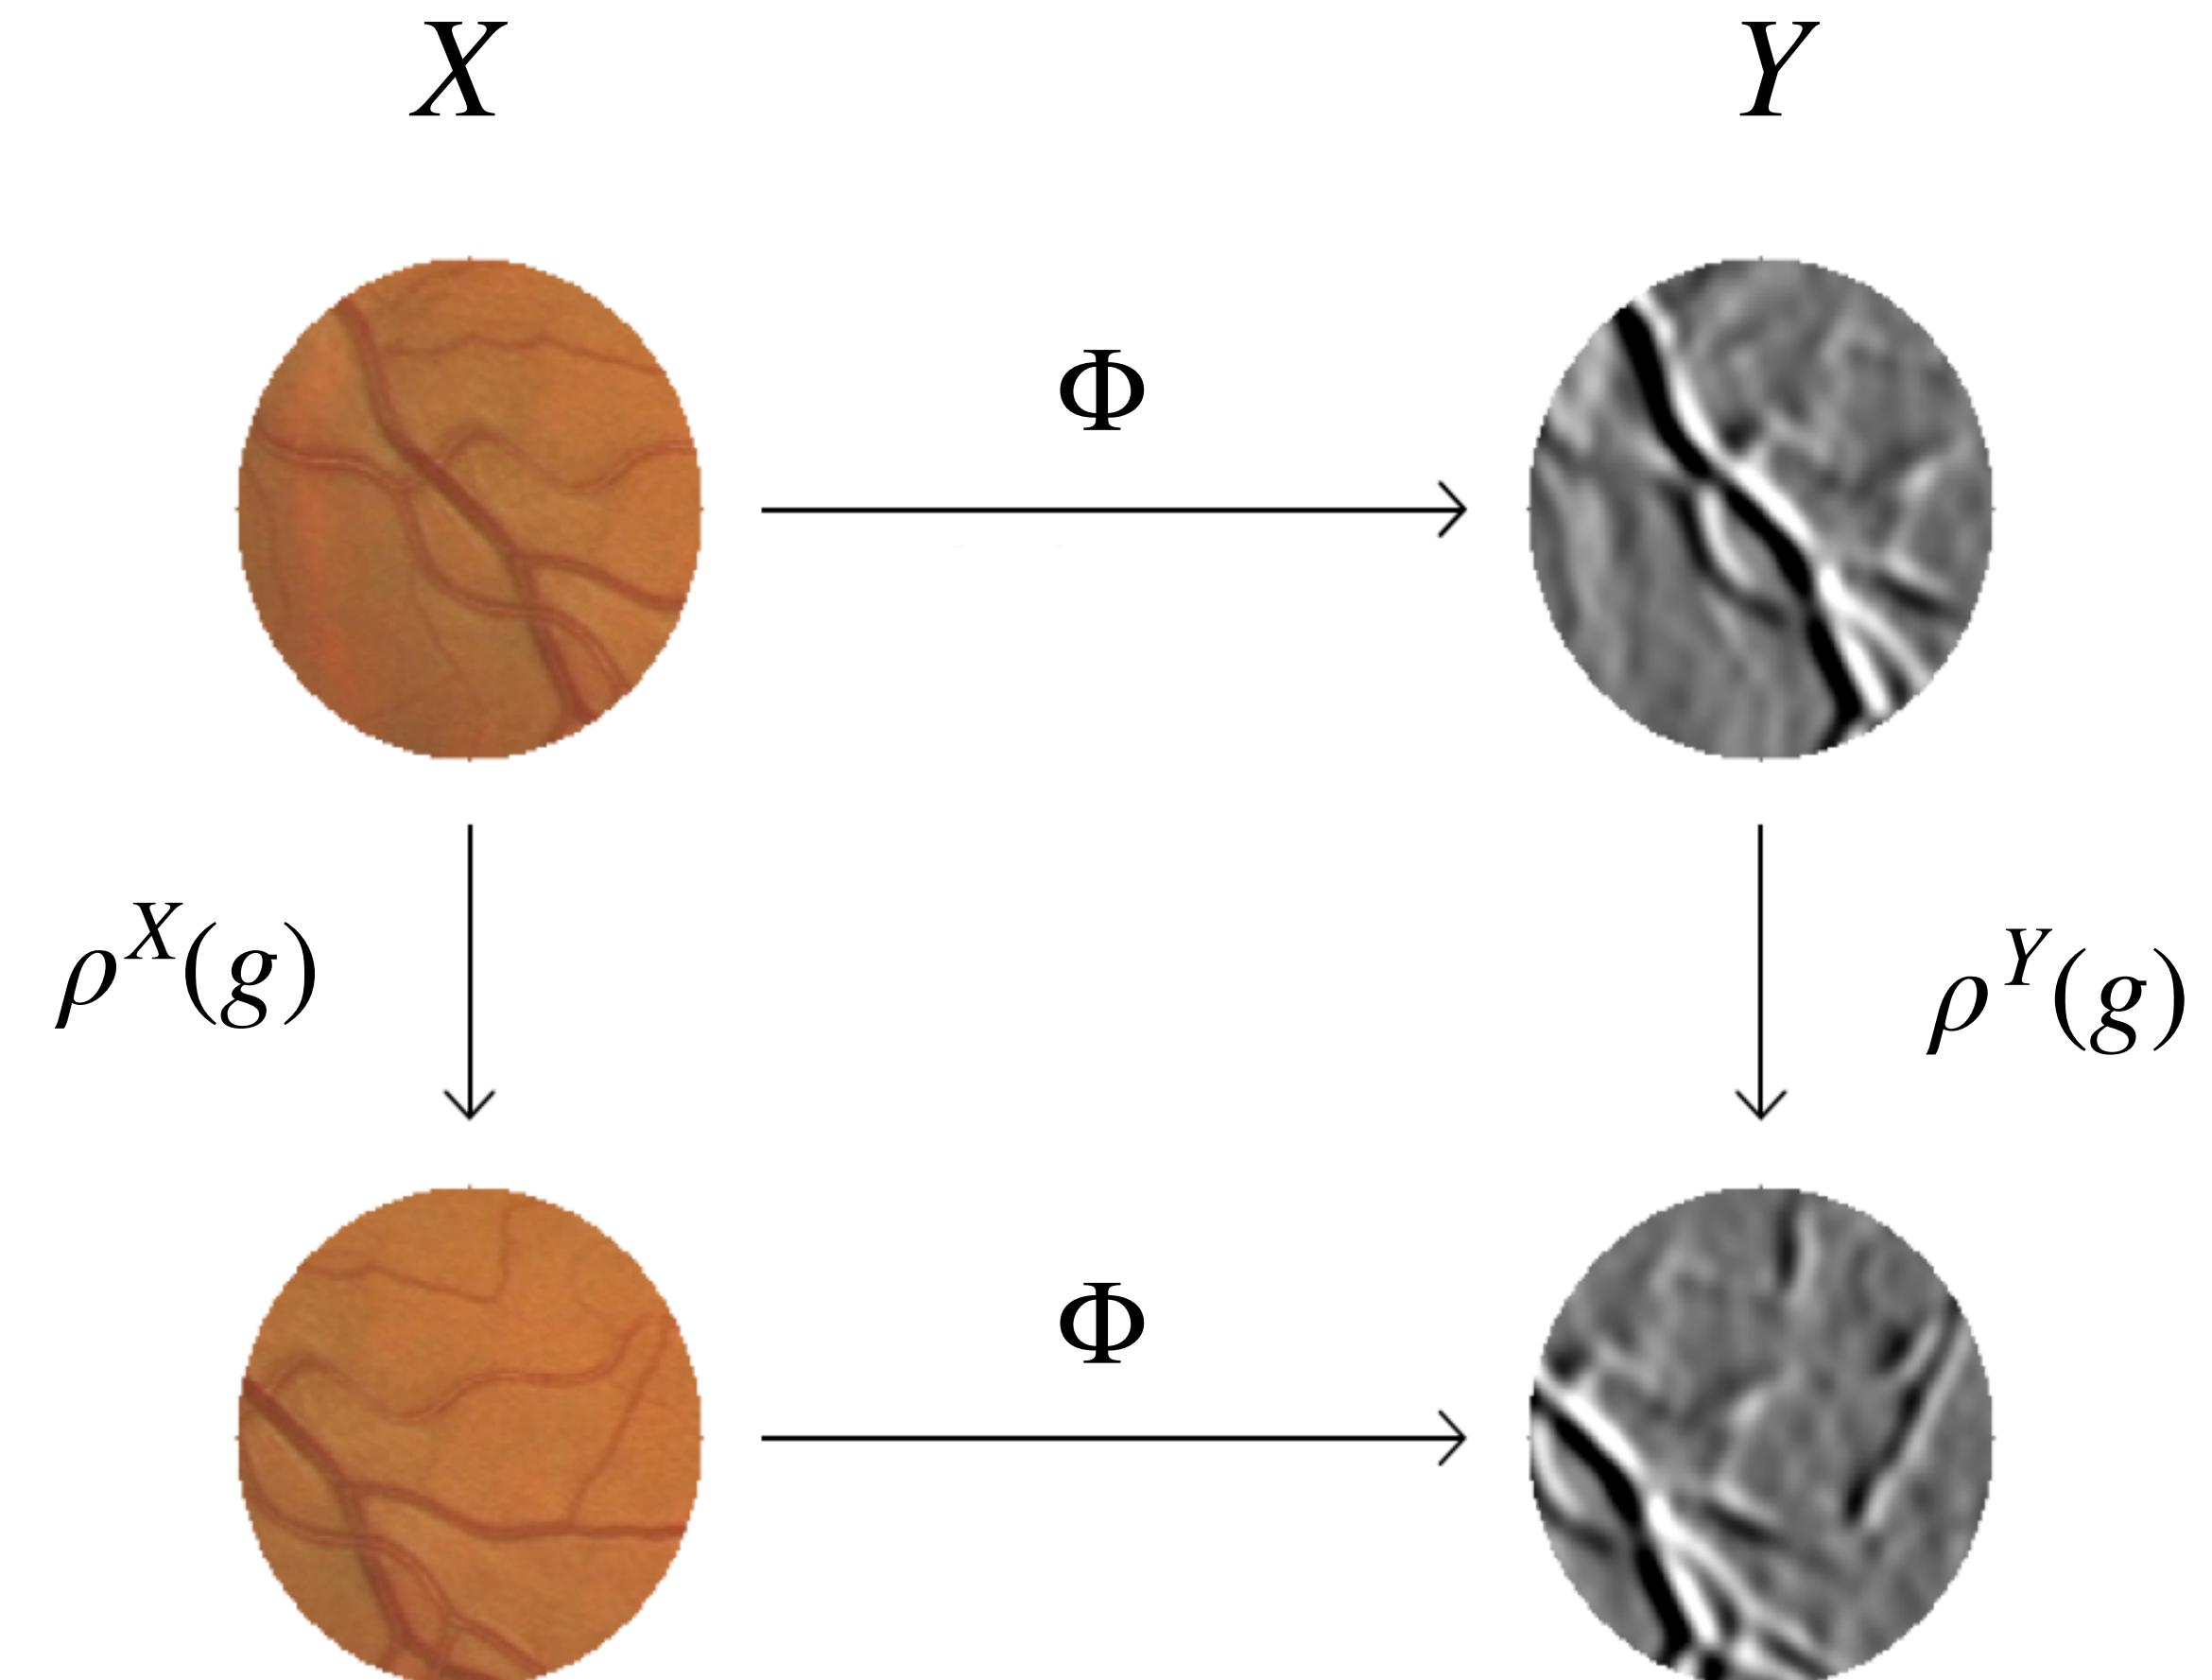
"group action" equals  
group product when  
domain is  $G$

# Equivariance

Equivariance is a property of an operator  $\Phi : X \rightarrow Y$  (such as a neural network layer) by which it commutes with the group action:

$$\Phi \circ \rho^X(g) = \rho^Y(g) \circ \Phi$$

group representation action on  $X$



1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

**Group theory: symmetries & recognition by components**  
(features have “poses”)

3. Group convolutions

4. Example

5. G-convs are all you need!

6. Steerable group convolutions

7. Feature fields and escnn library

8. Equivariant tensor product layers

9. Equivariant graph NNs

**1. Motivation**

Equivariance → weight-sharing and generalization

**2. Pattern matching using group theory**

Group theory: symmetries & recognition by components  
(features have “poses”)

**3. Group convolutions**

**4. Example**

**5. G-convs are all you need!**

**6. Steerable group convolutions**

**7. Feature fields and escnn library**

**8. Equivariant tensor product layers**

**9. Equivariant graph NNs**

Are convolutions with reflected conv kernels (and vice versa)

# Cross-correlations

$$(k \star_{\mathbb{R}^2} f)(\mathbf{x}) = \int_{\mathbb{R}^2} k(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$

Are convolutions with reflected conv kernels (and vice versa)

# Cross-correlations

Representation of the translation group!

$$(k \star_{\mathbb{R}^2} f)(\mathbf{x}) = \int_{\mathbb{R}^2} k(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}' = (\mathcal{L}_g k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$$

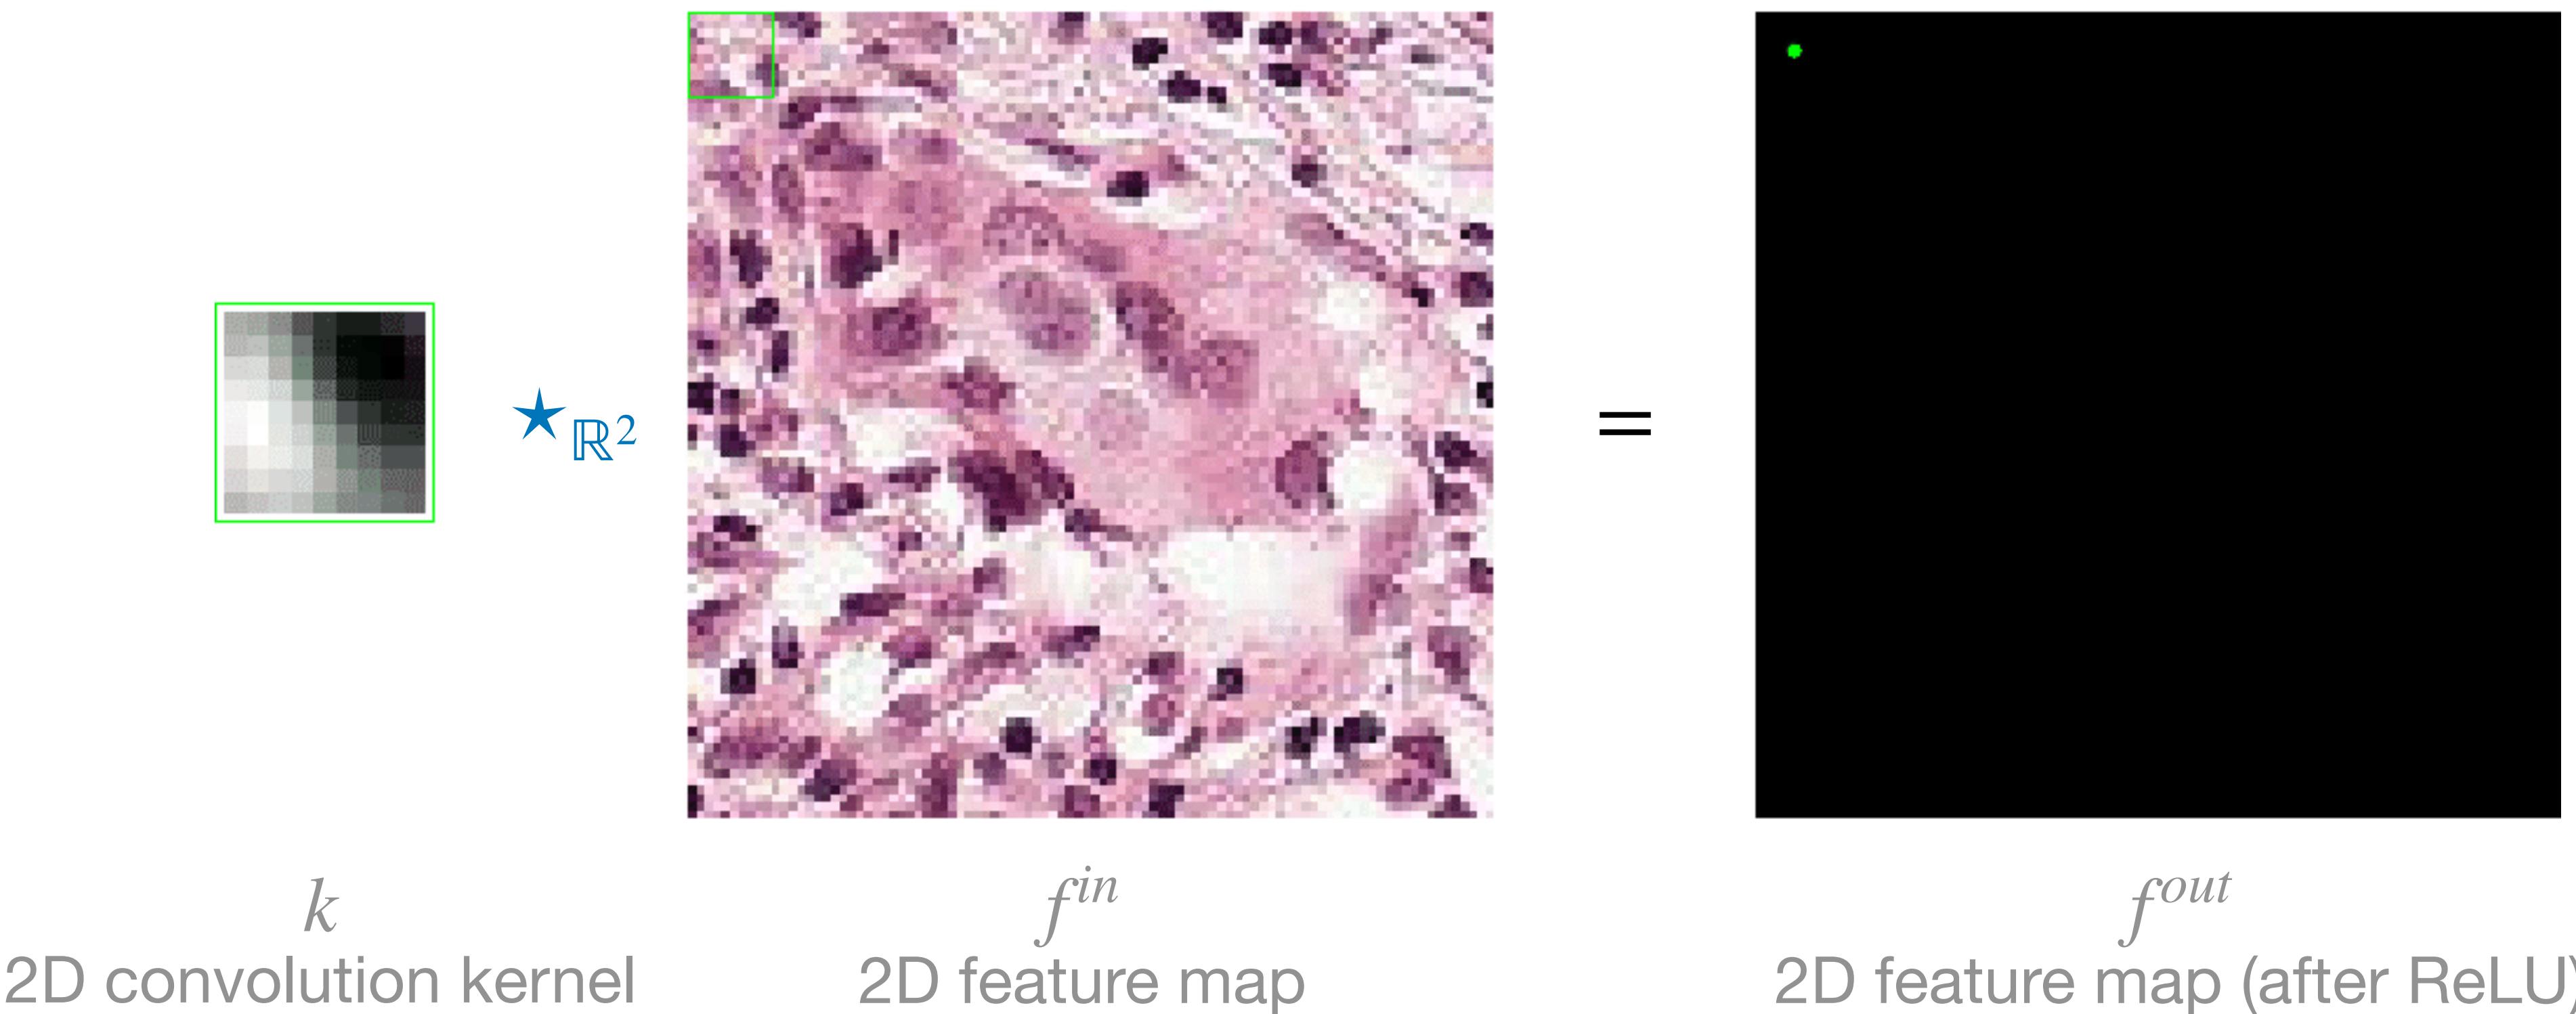


Are convolutions with reflected conv kernels (and vice versa)

# Cross-correlations

Representation of the translation group!

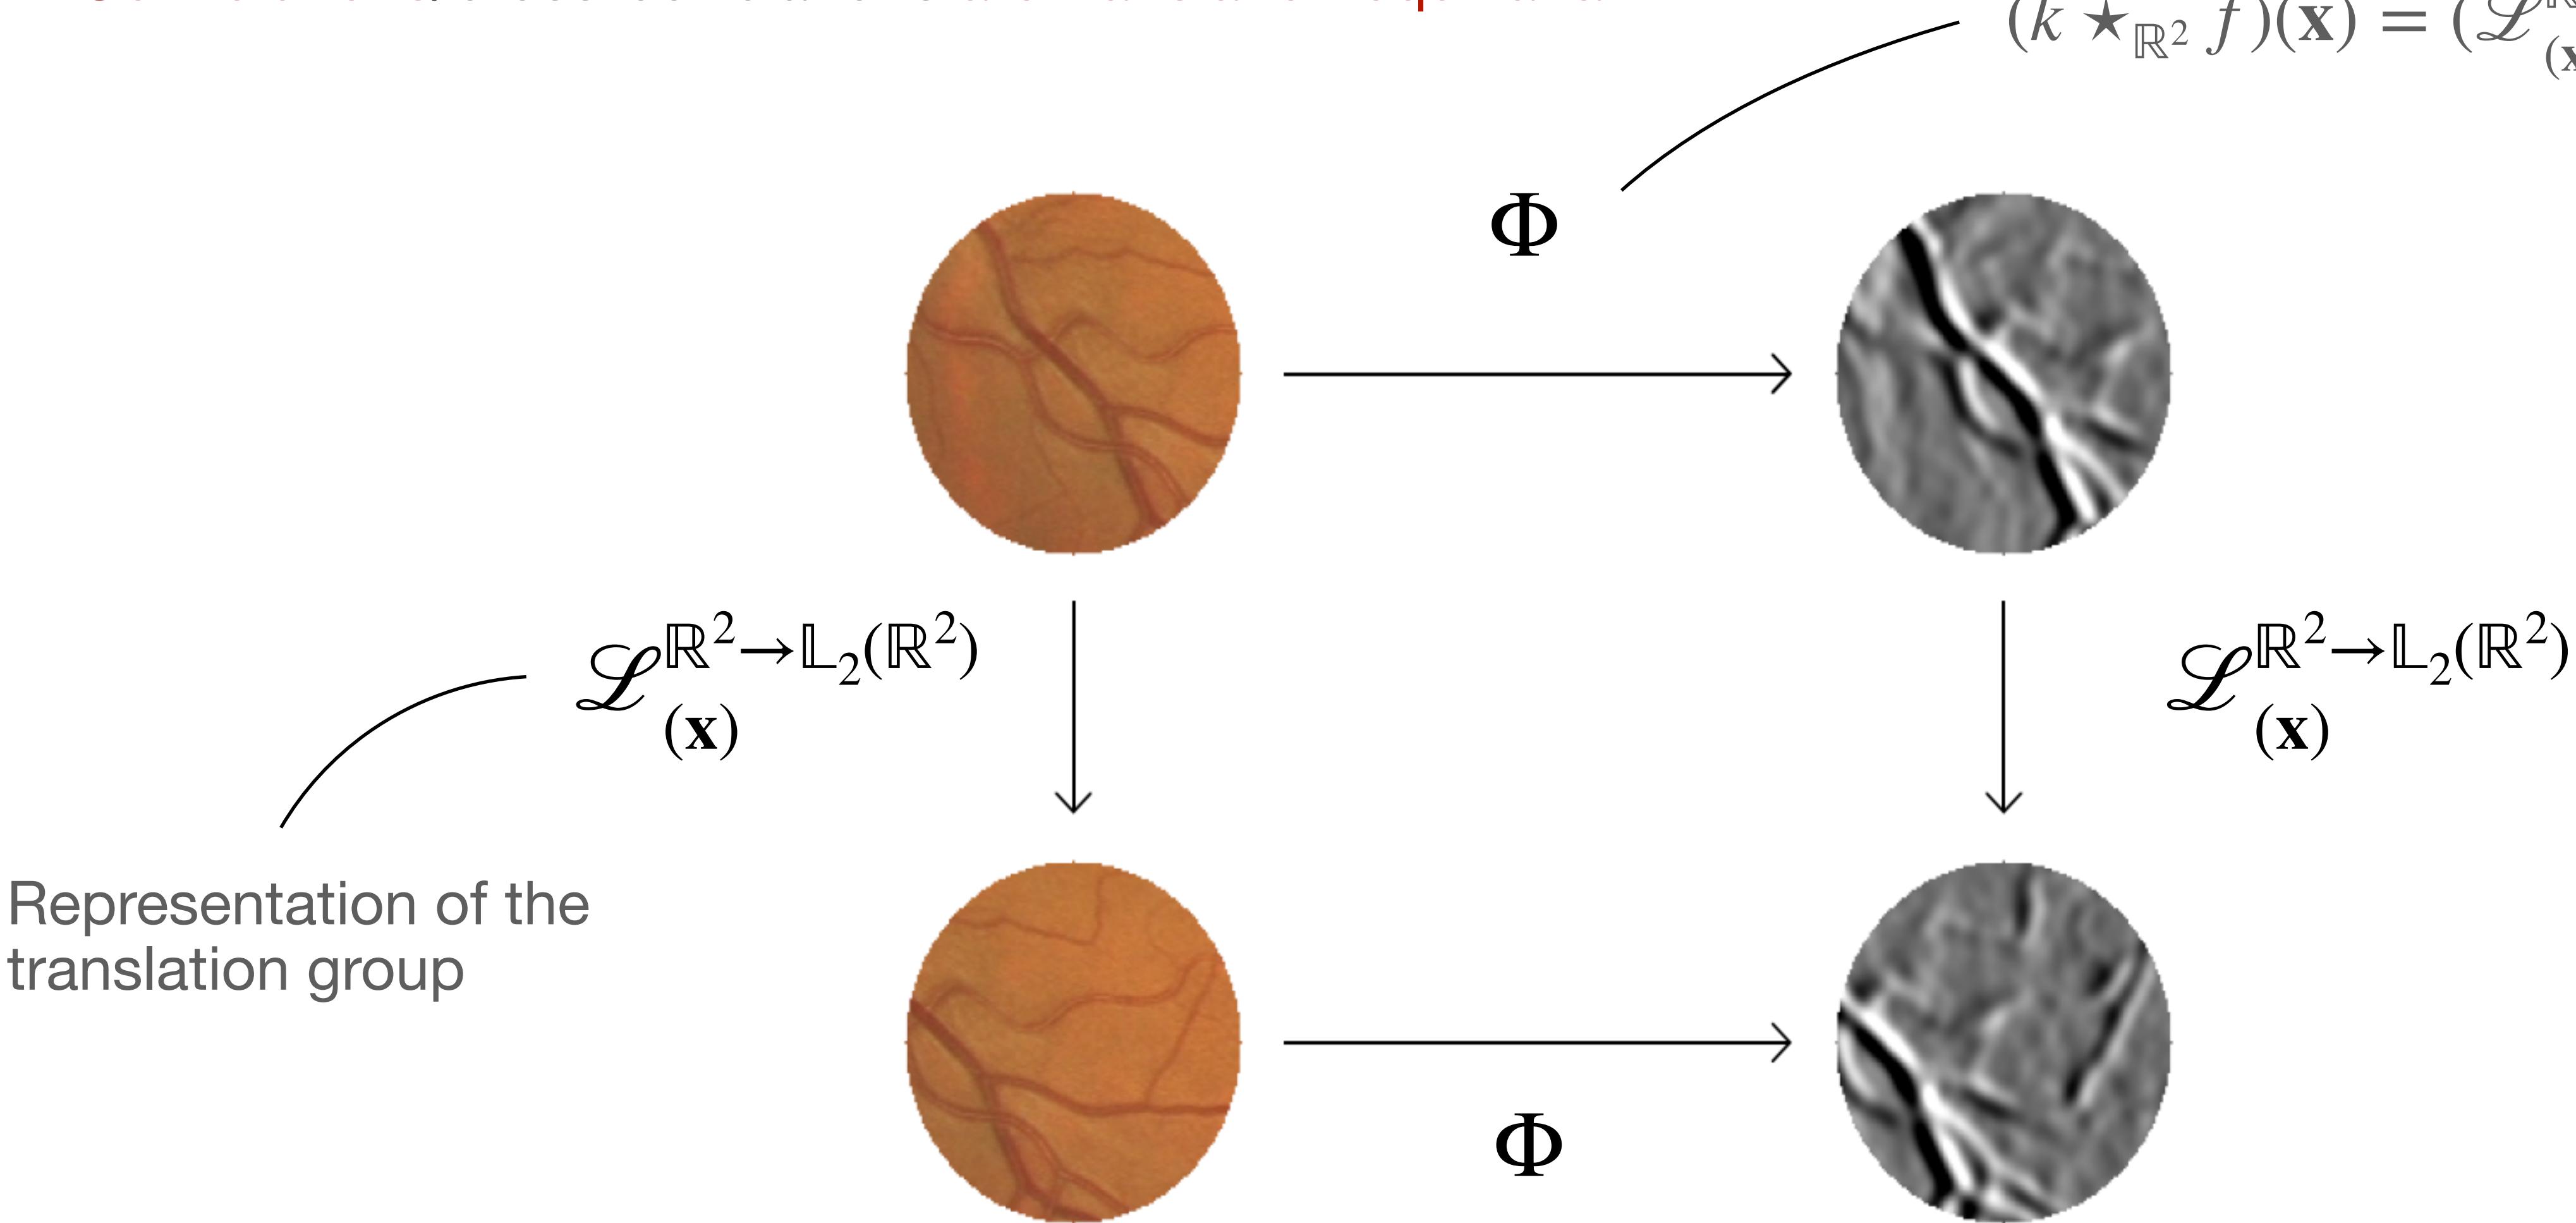
$$(k \star_{\mathbb{R}^2} f)(\mathbf{x}) = \int_{\mathbb{R}^2} k(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}' = (\mathcal{L}_g k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$$



# Equivariance

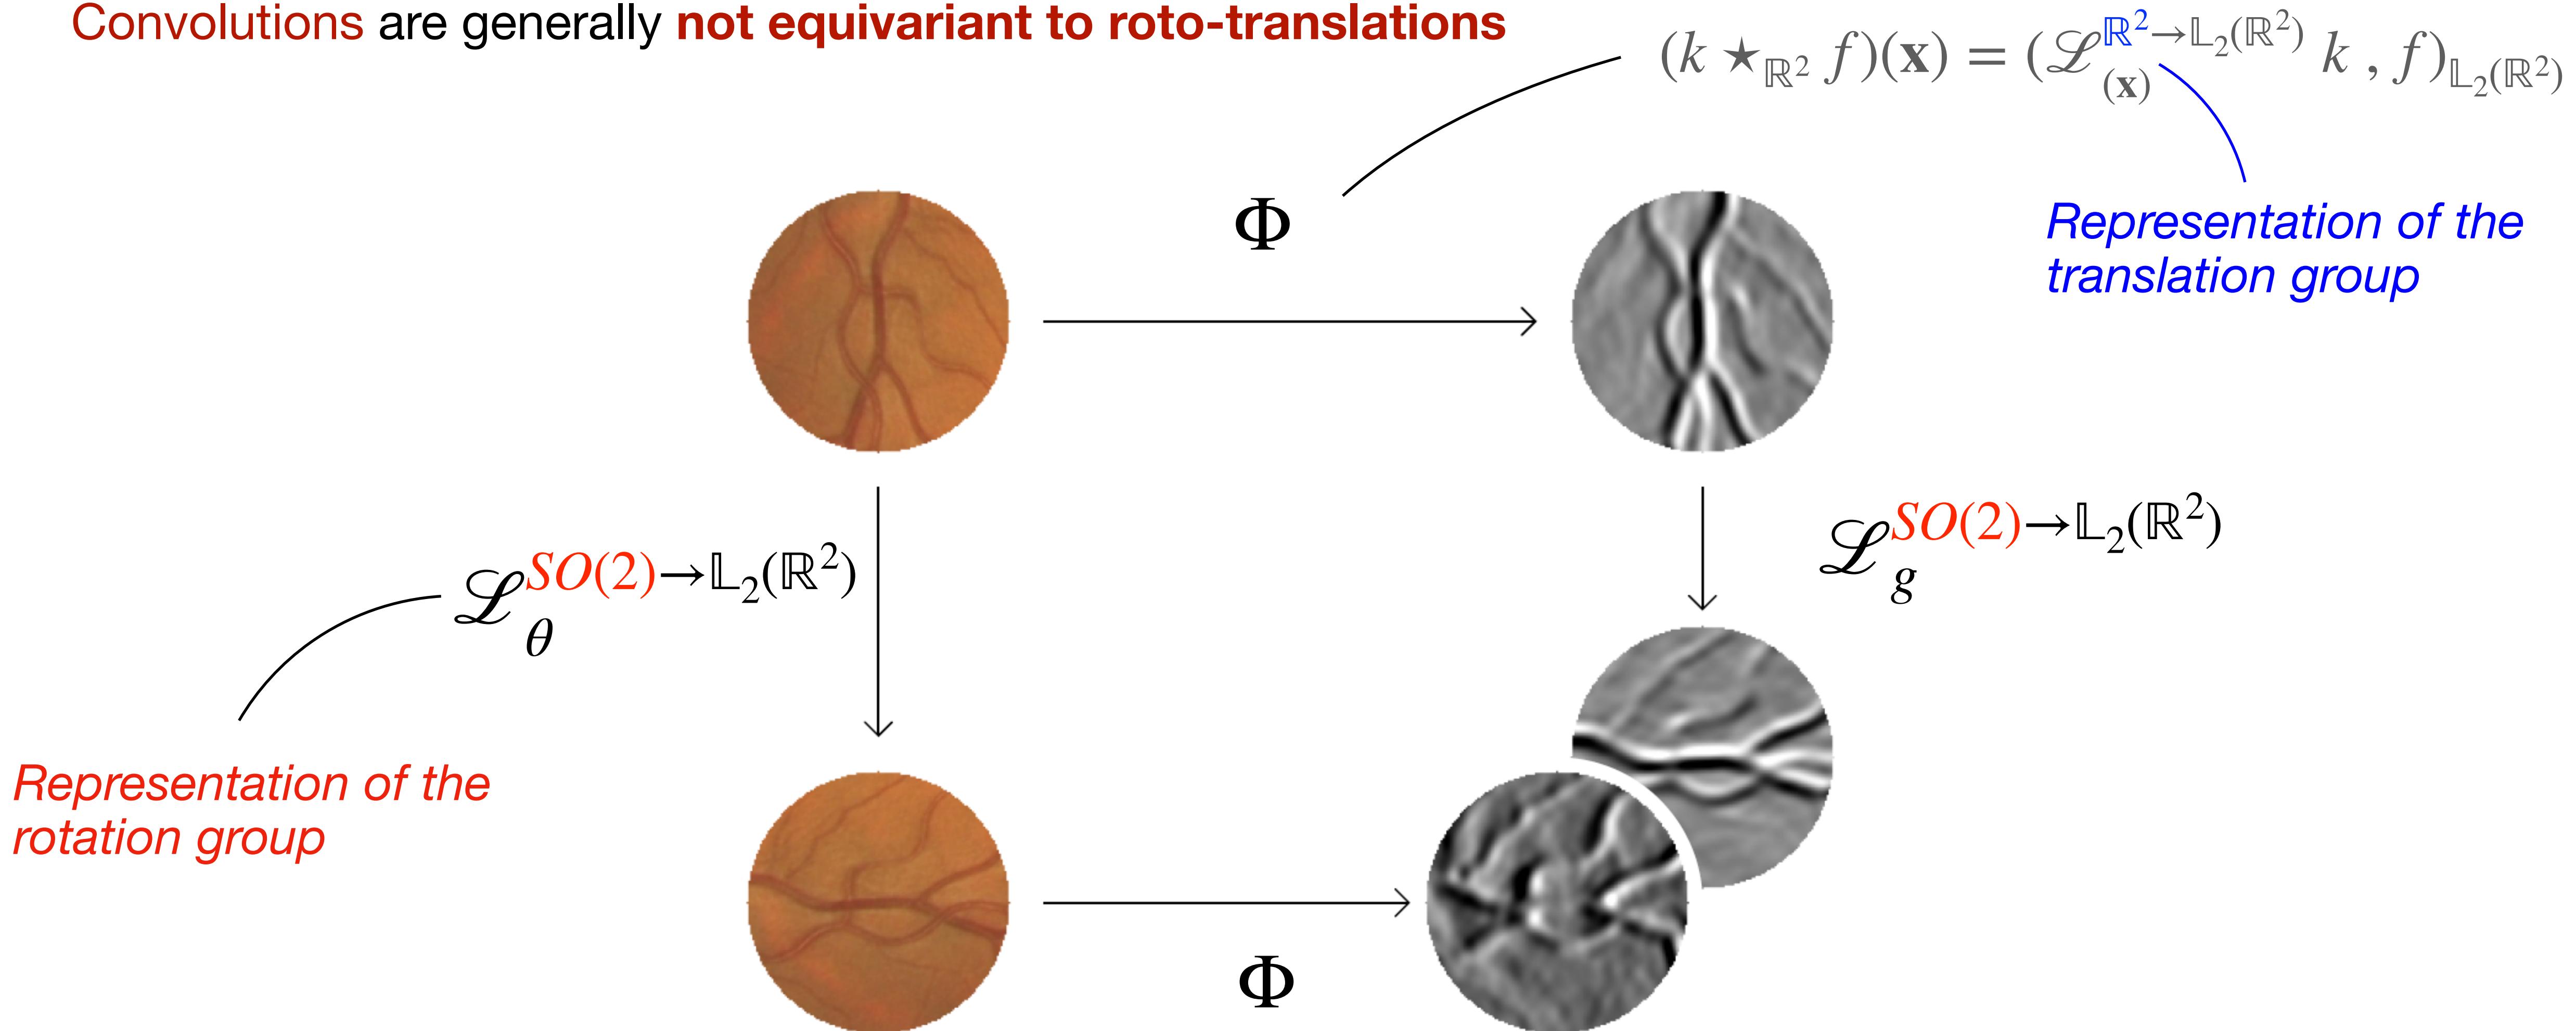
Convolutions/cross-correlations are translation equivariant

$$(k \star_{\mathbb{R}^2} f)(\mathbf{x}) = (\mathcal{L}_{(\mathbf{x})}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$$



# Equivariance

Convolutions are generally **not equivariant to roto-translations**



# SE(2) equivariant cross-correlations

*Representation of the roto-translation group!*

**Lifting correlations:**  $(k \tilde{\star} f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$

# SE(2) equivariant cross-correlations

*Representation of the roto-translation group!*

**Lifting correlations:**  $(k \tilde{\star} f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)} = (\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$

*translation      rotation*

# SE(2) equivariant cross-correlations

*Representation of the roto-translation group!*

Lifting correlations:  $(k \tilde{\star} f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$

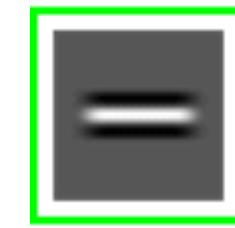
$\overbrace{\mathcal{L}_{\mathbf{x}}^{\mathbf{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)}}^{\begin{matrix} \text{translation} & \text{rotation} \end{matrix}} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$

$k(\mathbf{R}_{\theta}^{-1}(\mathbf{x}' - \mathbf{x}))$

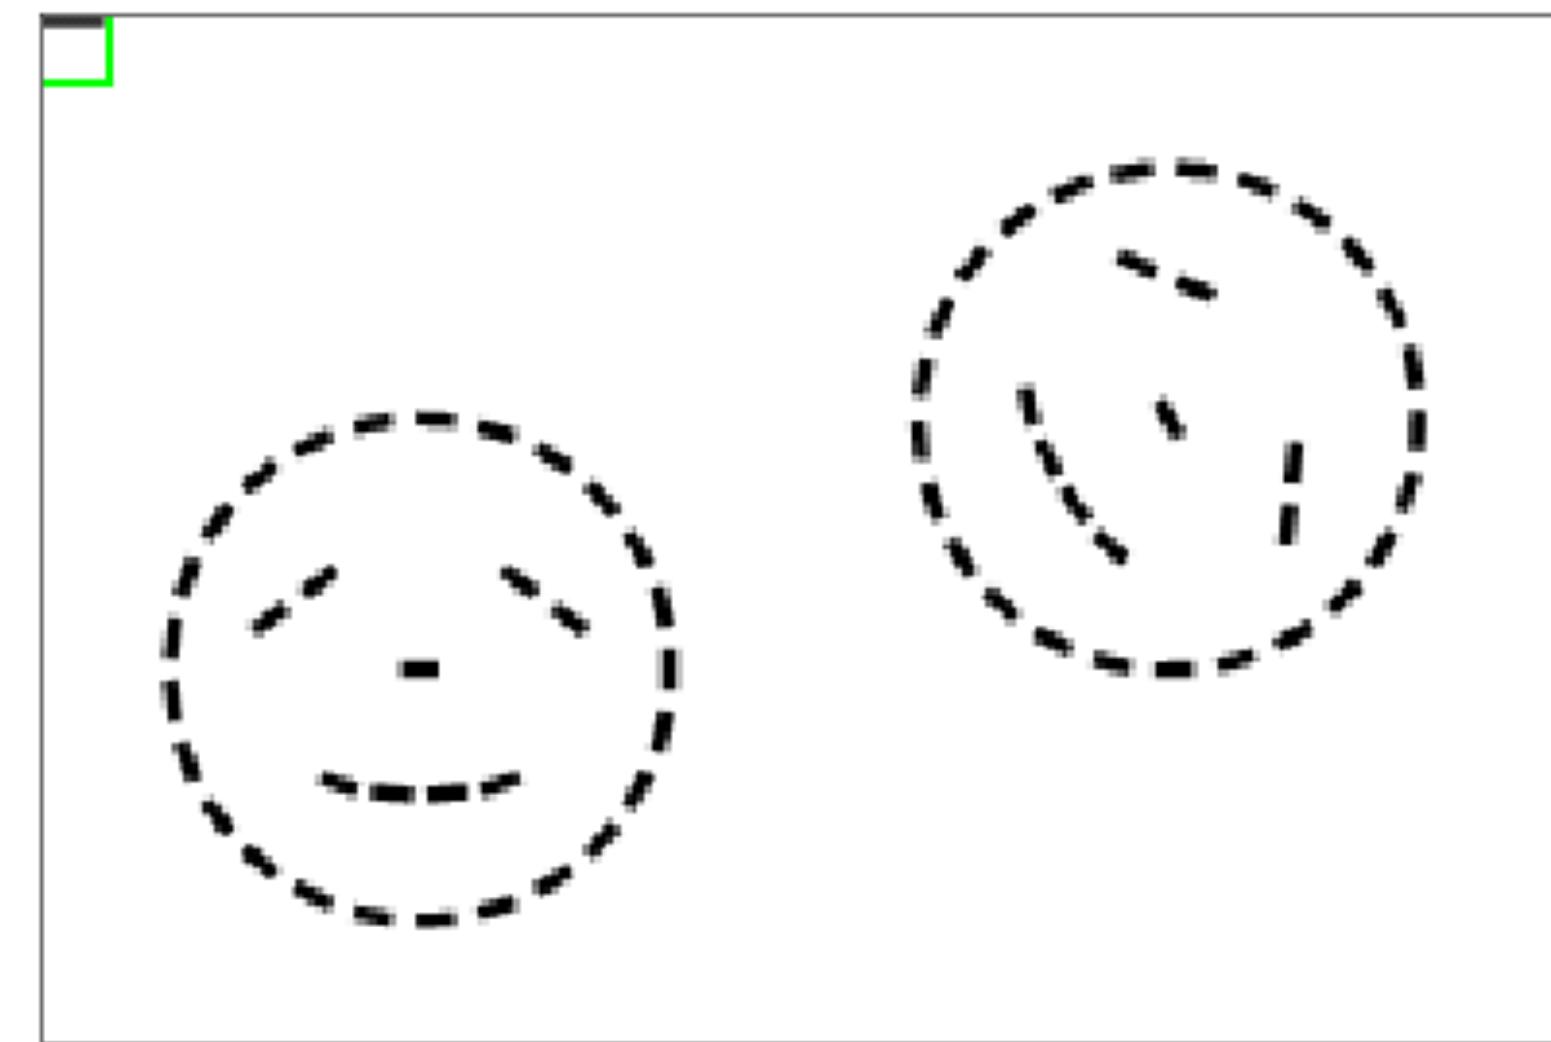
# SE(2) equivariant cross-correlations

*Representation of the roto-translation group!*

**Lifting correlations:**  $(k \tilde{\star} f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)} = (\overbrace{\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)}}^{\text{translation rotation}} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$

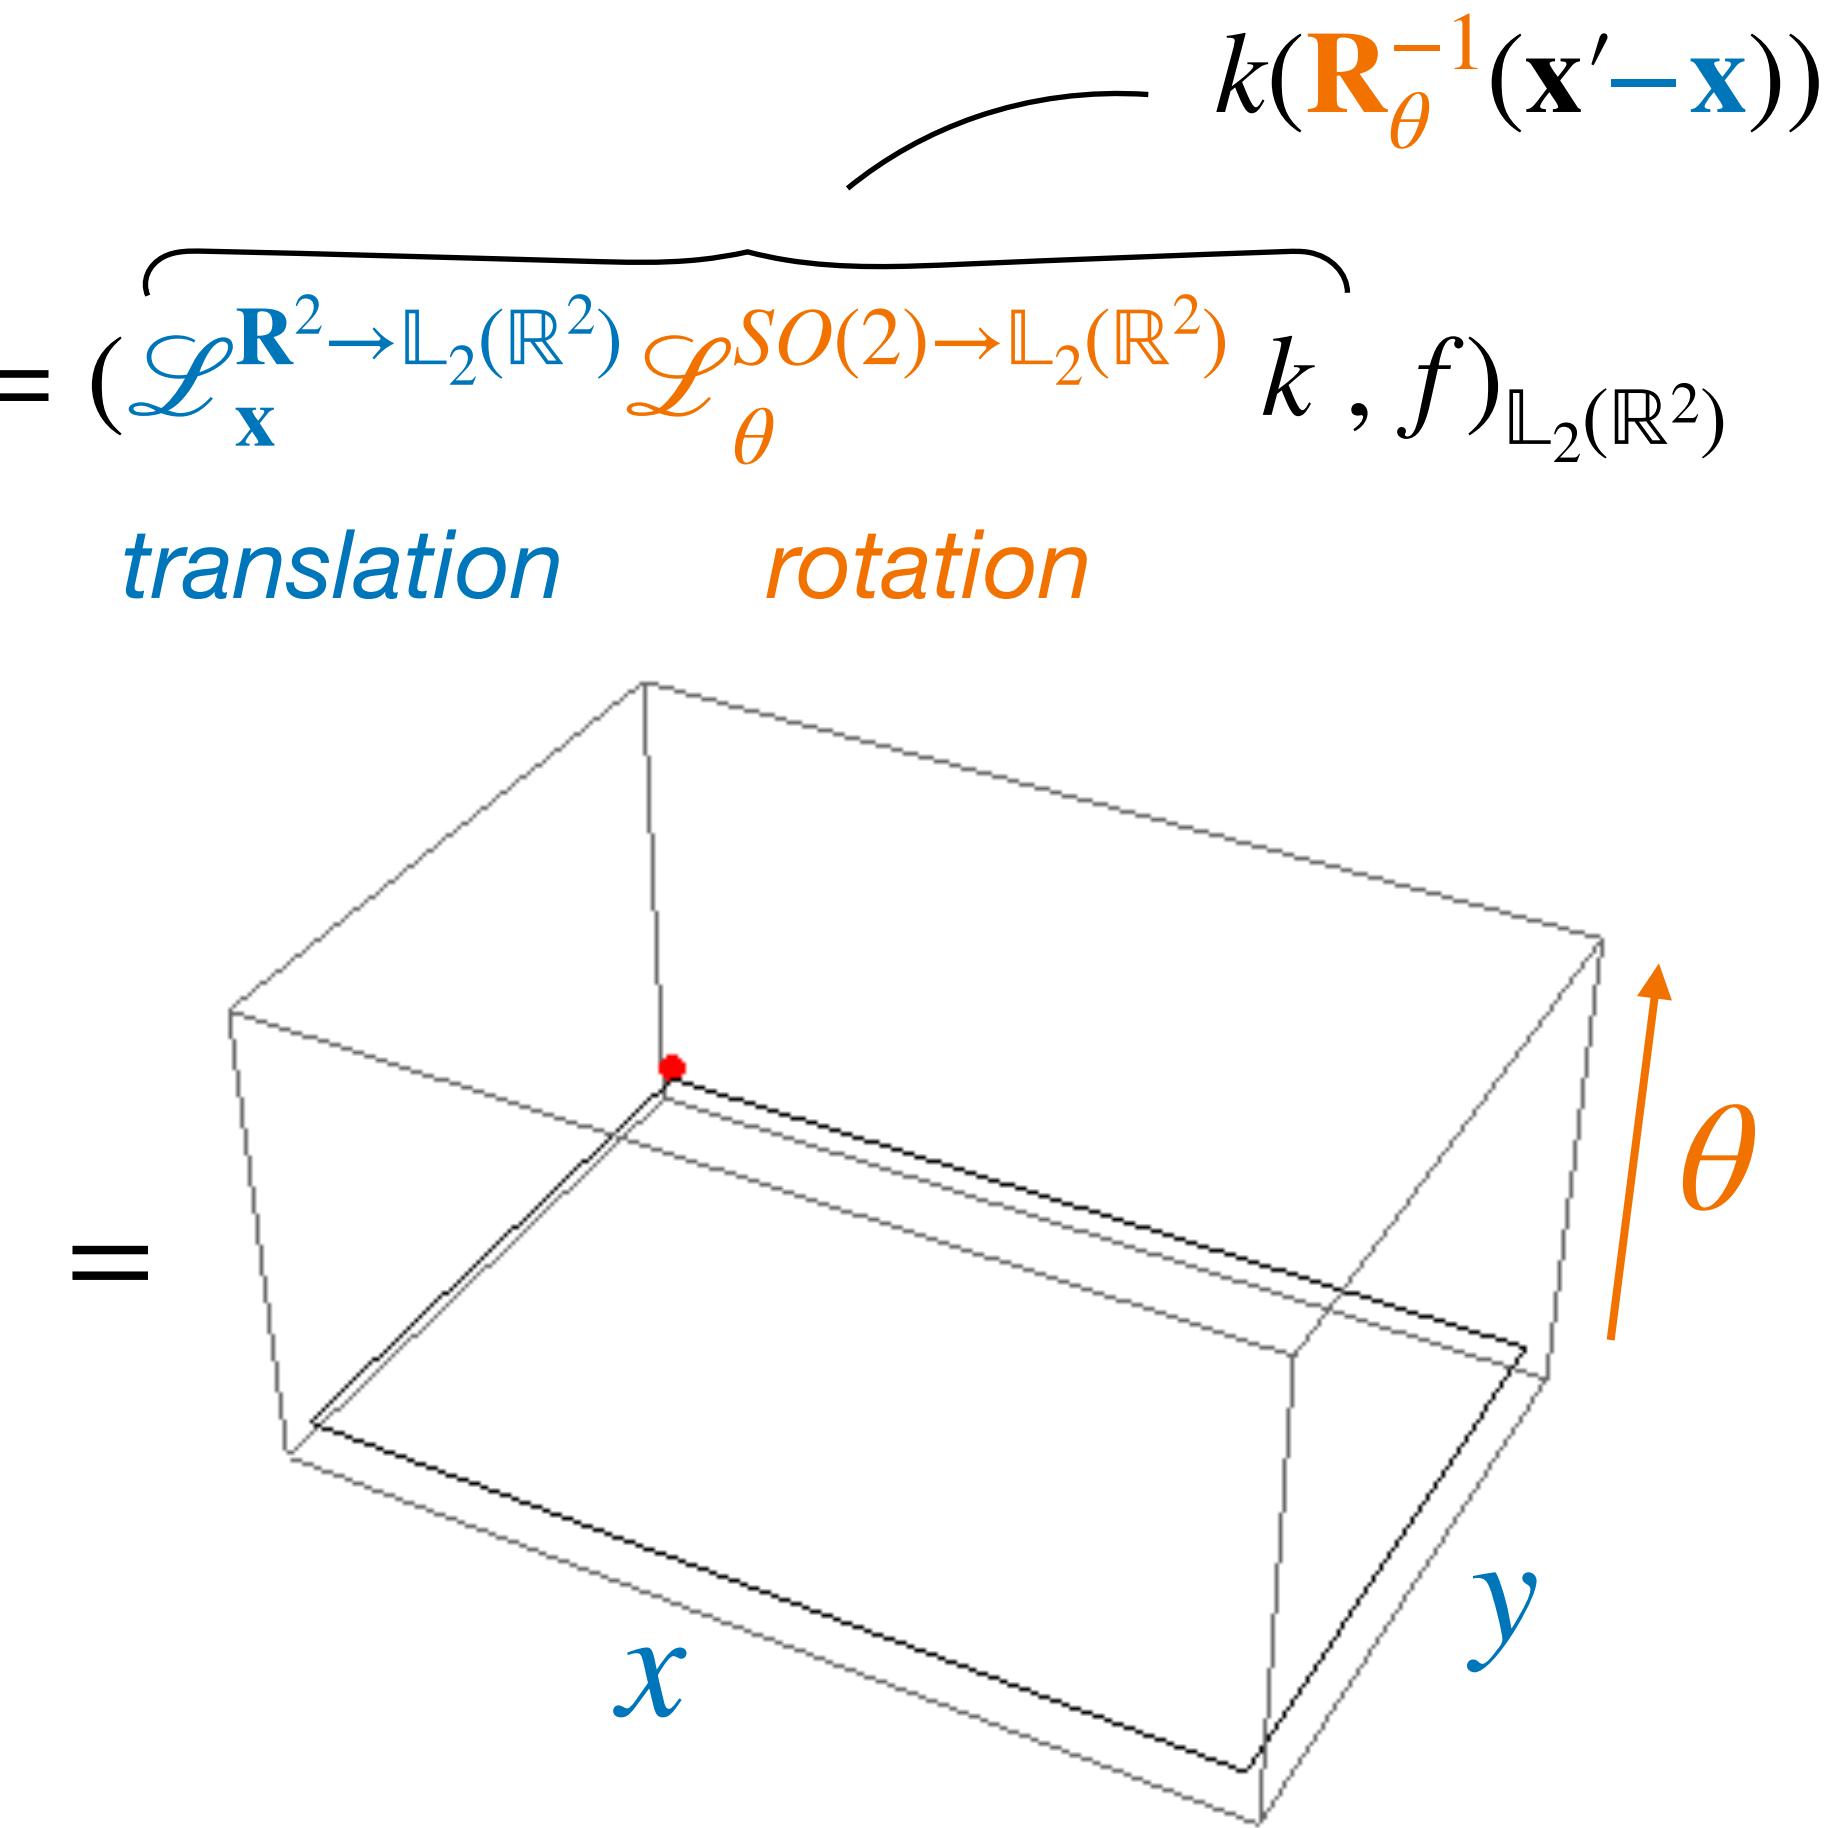


$\star_{\mathbb{R}^2}$



$\mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k$   
Rotated 2D convolution kernel

$f^{in}$   
2D feature map



$f^{out}$   
3D (SE(2)) feature map (after ReLU)

# SE(2) equivariant cross-correlations

*Representation of the roto-translation group!*

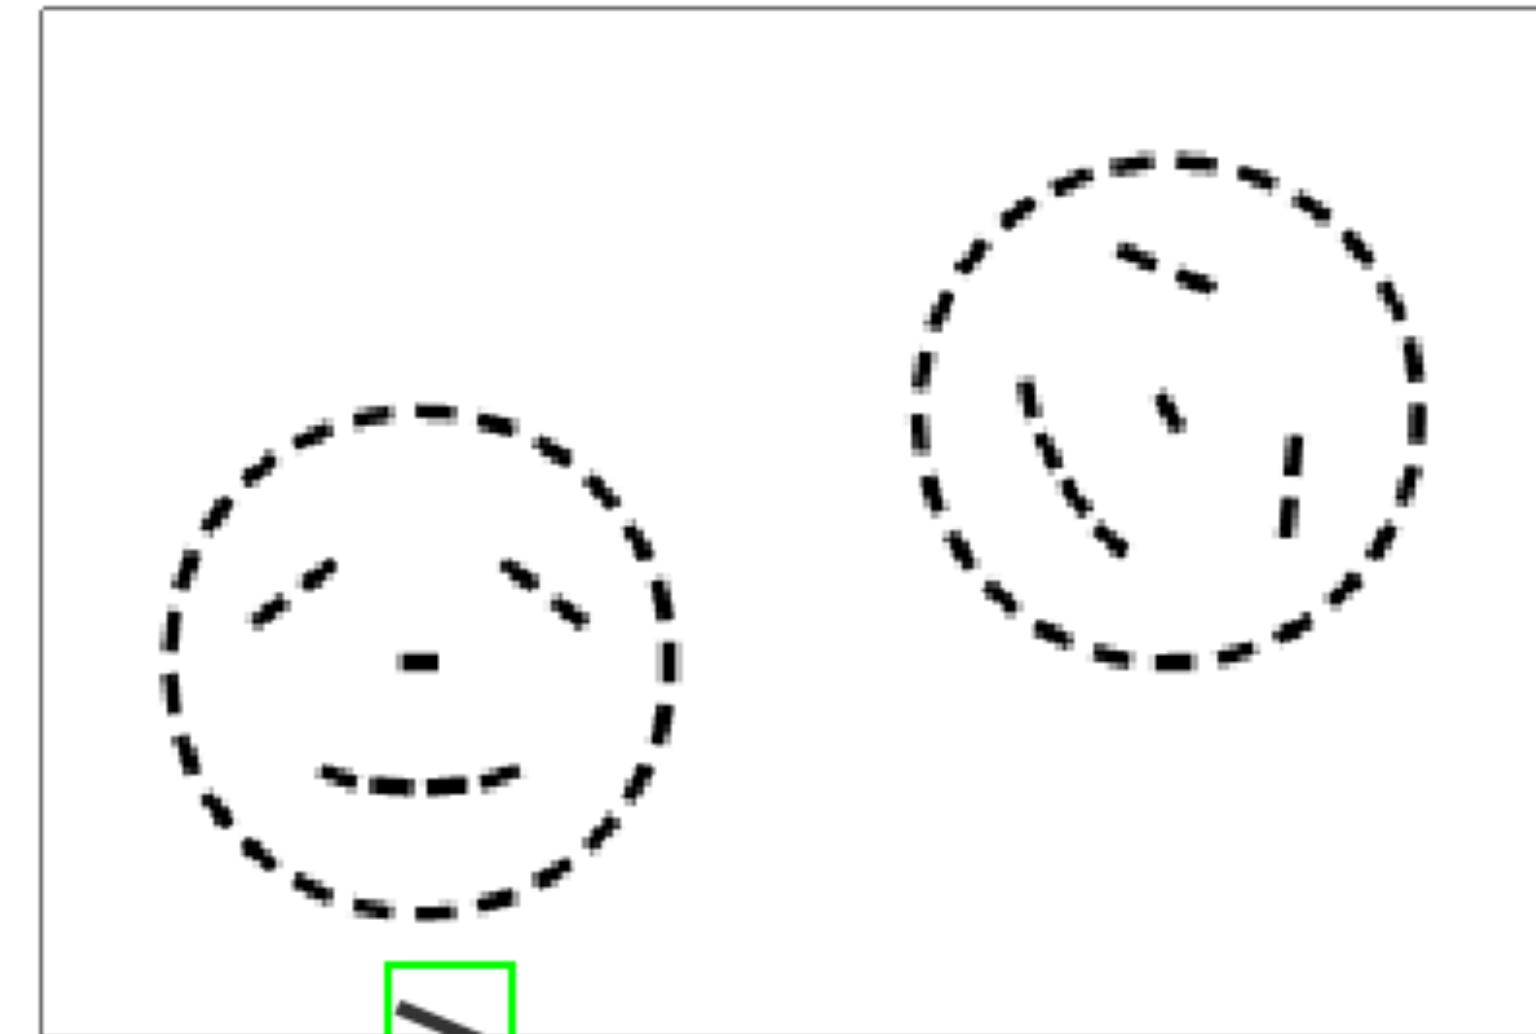
**Lifting correlations:**  $(k \tilde{\star} f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)} = (\overbrace{\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)}}^{\text{translation}} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$

$$k(\mathbf{R}_{\theta}^{-1}(\mathbf{x}' - \mathbf{x}))$$

*translation      rotation*

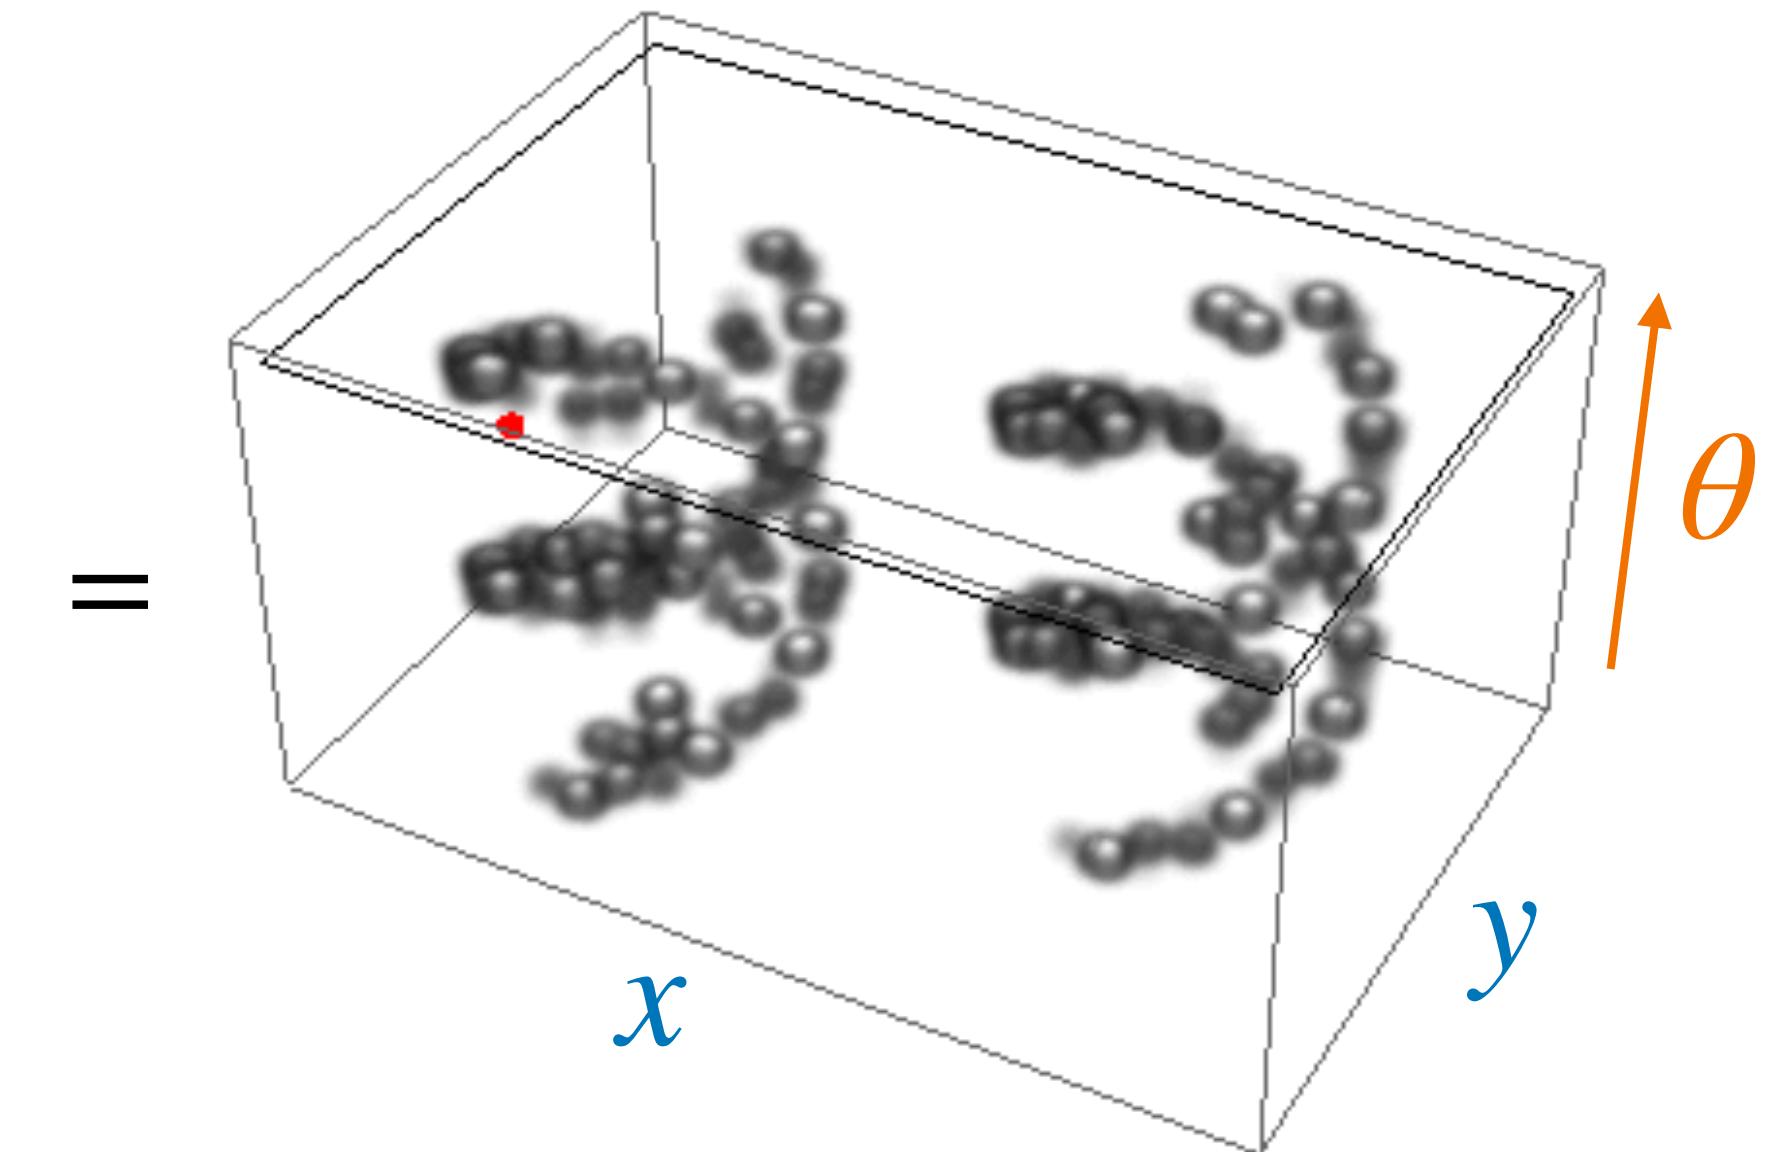


$\star_{\mathbb{R}^2}$



$\mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k$   
Rotated 2D convolution kernel

$f^{in}$   
2D feature map

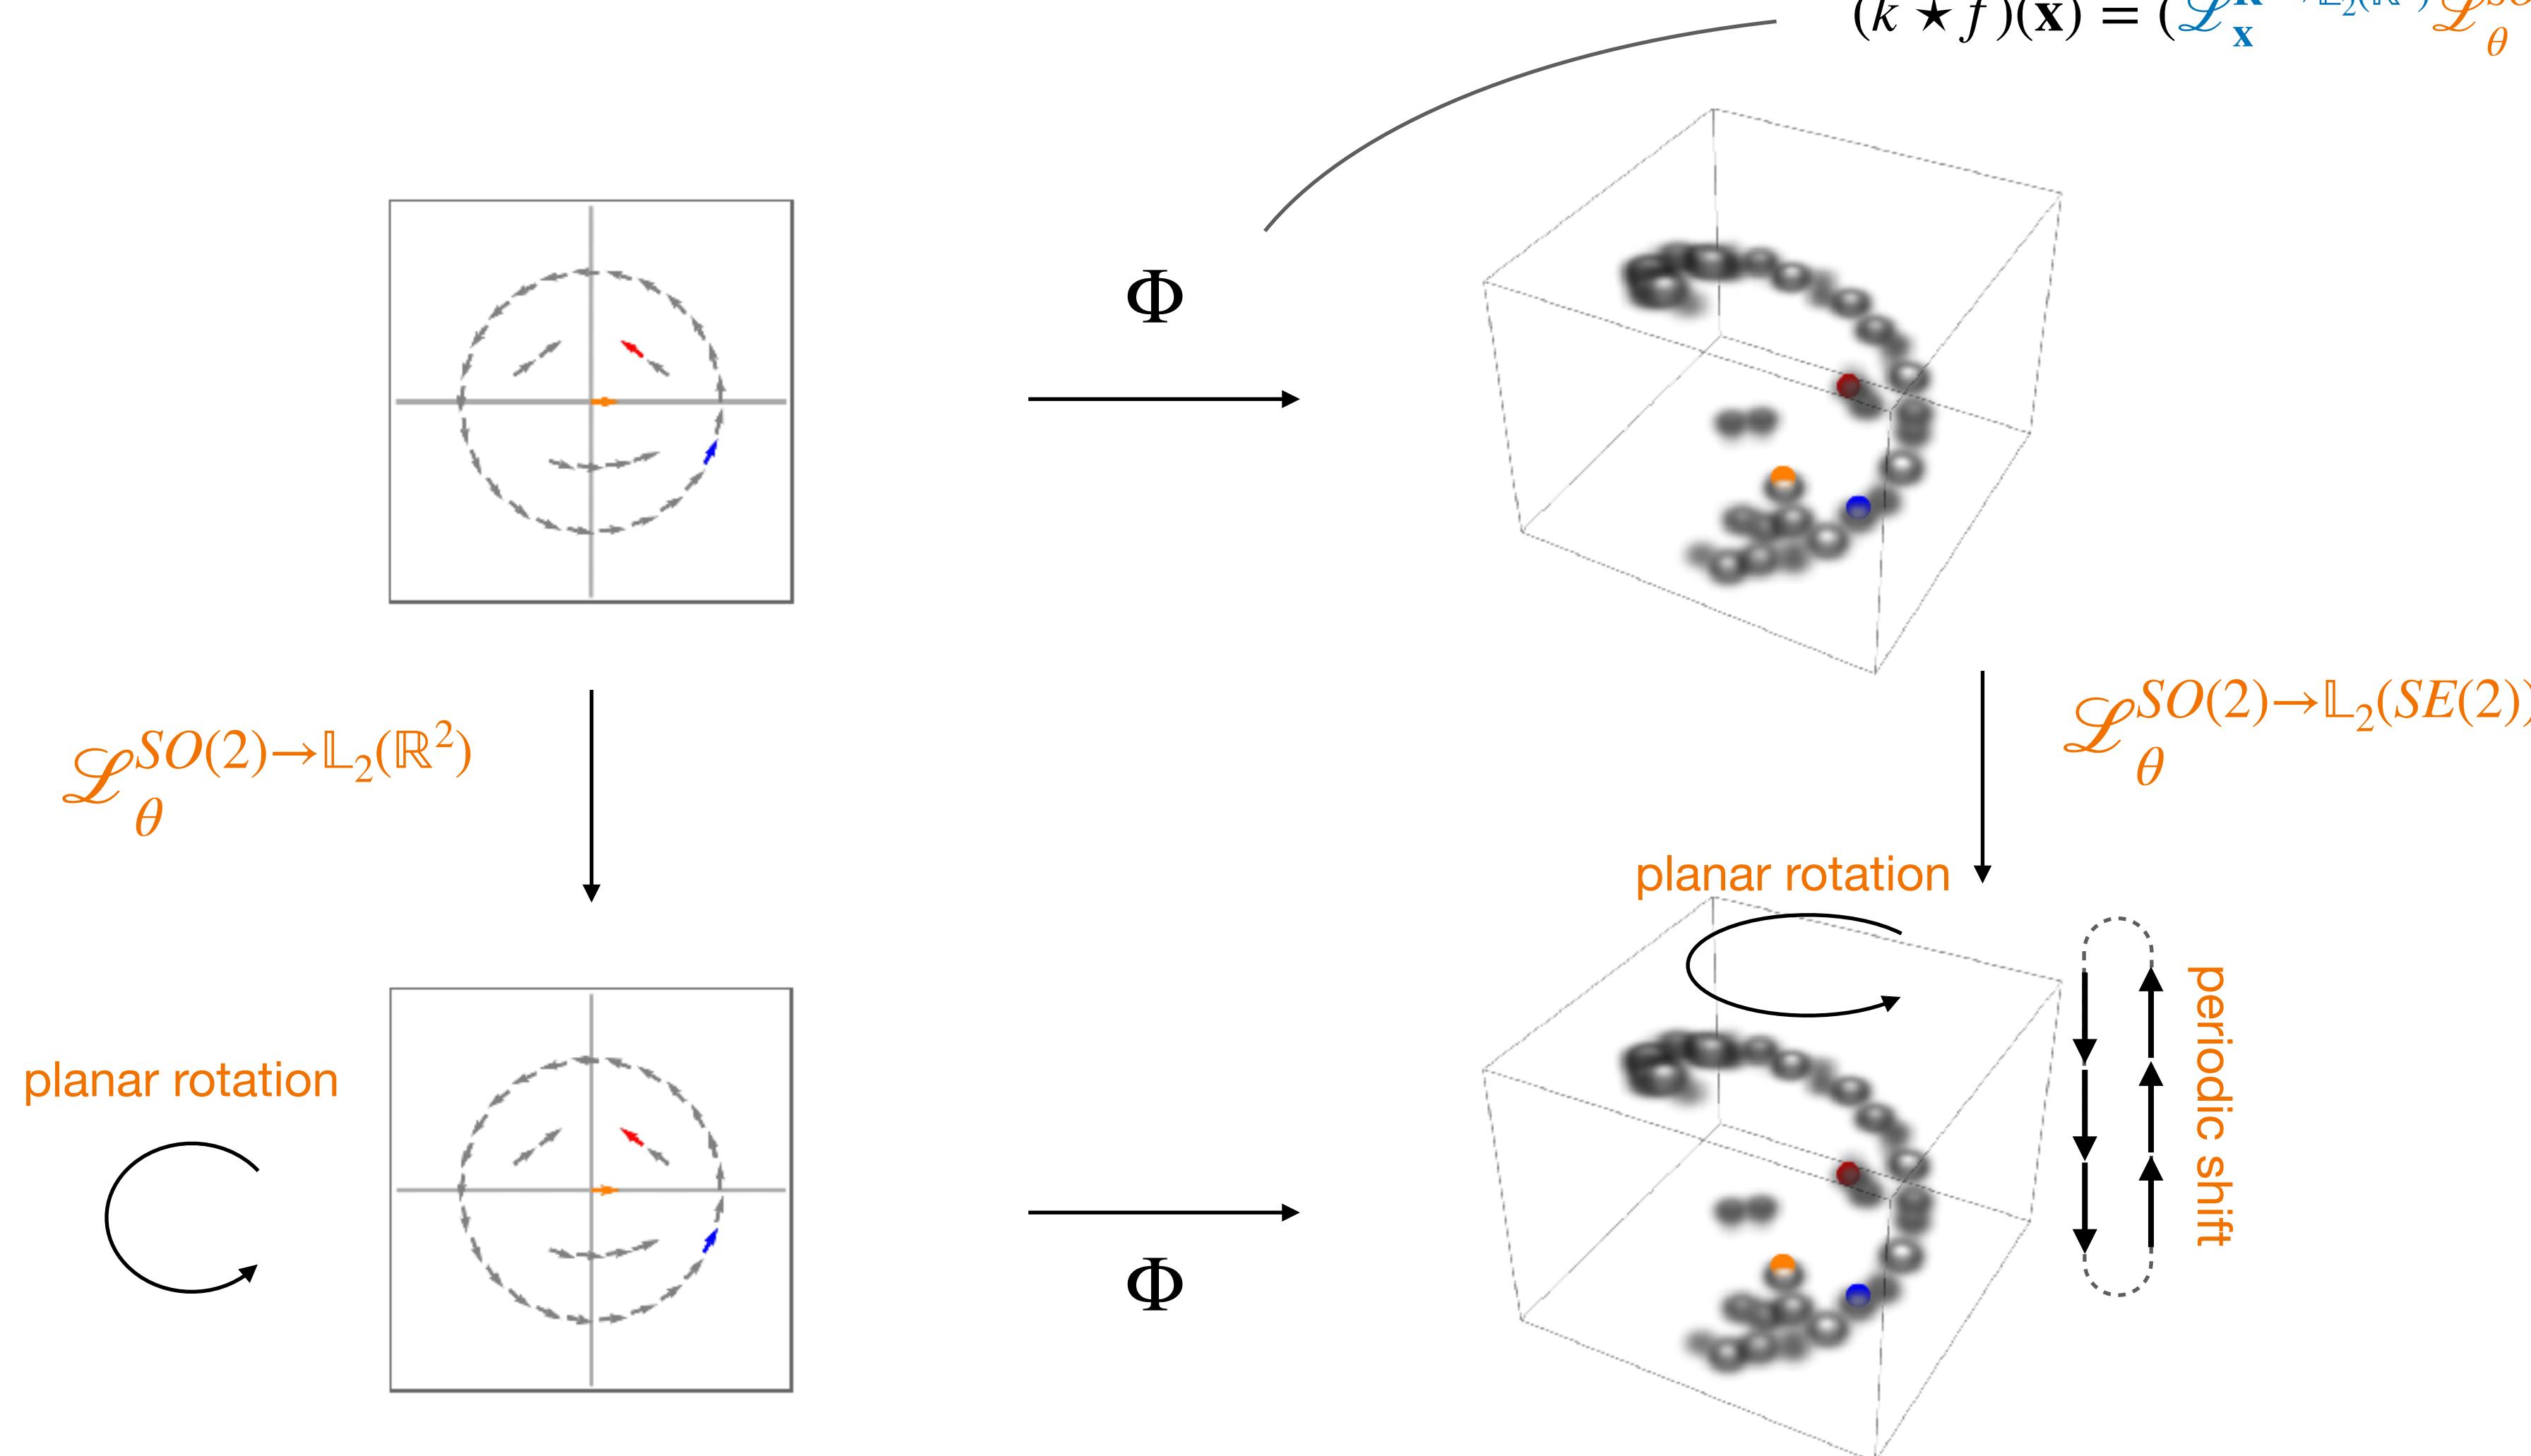


$f^{out}$   
3D (SE(2)) feature map (after ReLU)

# Equivariance

SE(2) group **lifting convolutions** are roto-translation equivariant

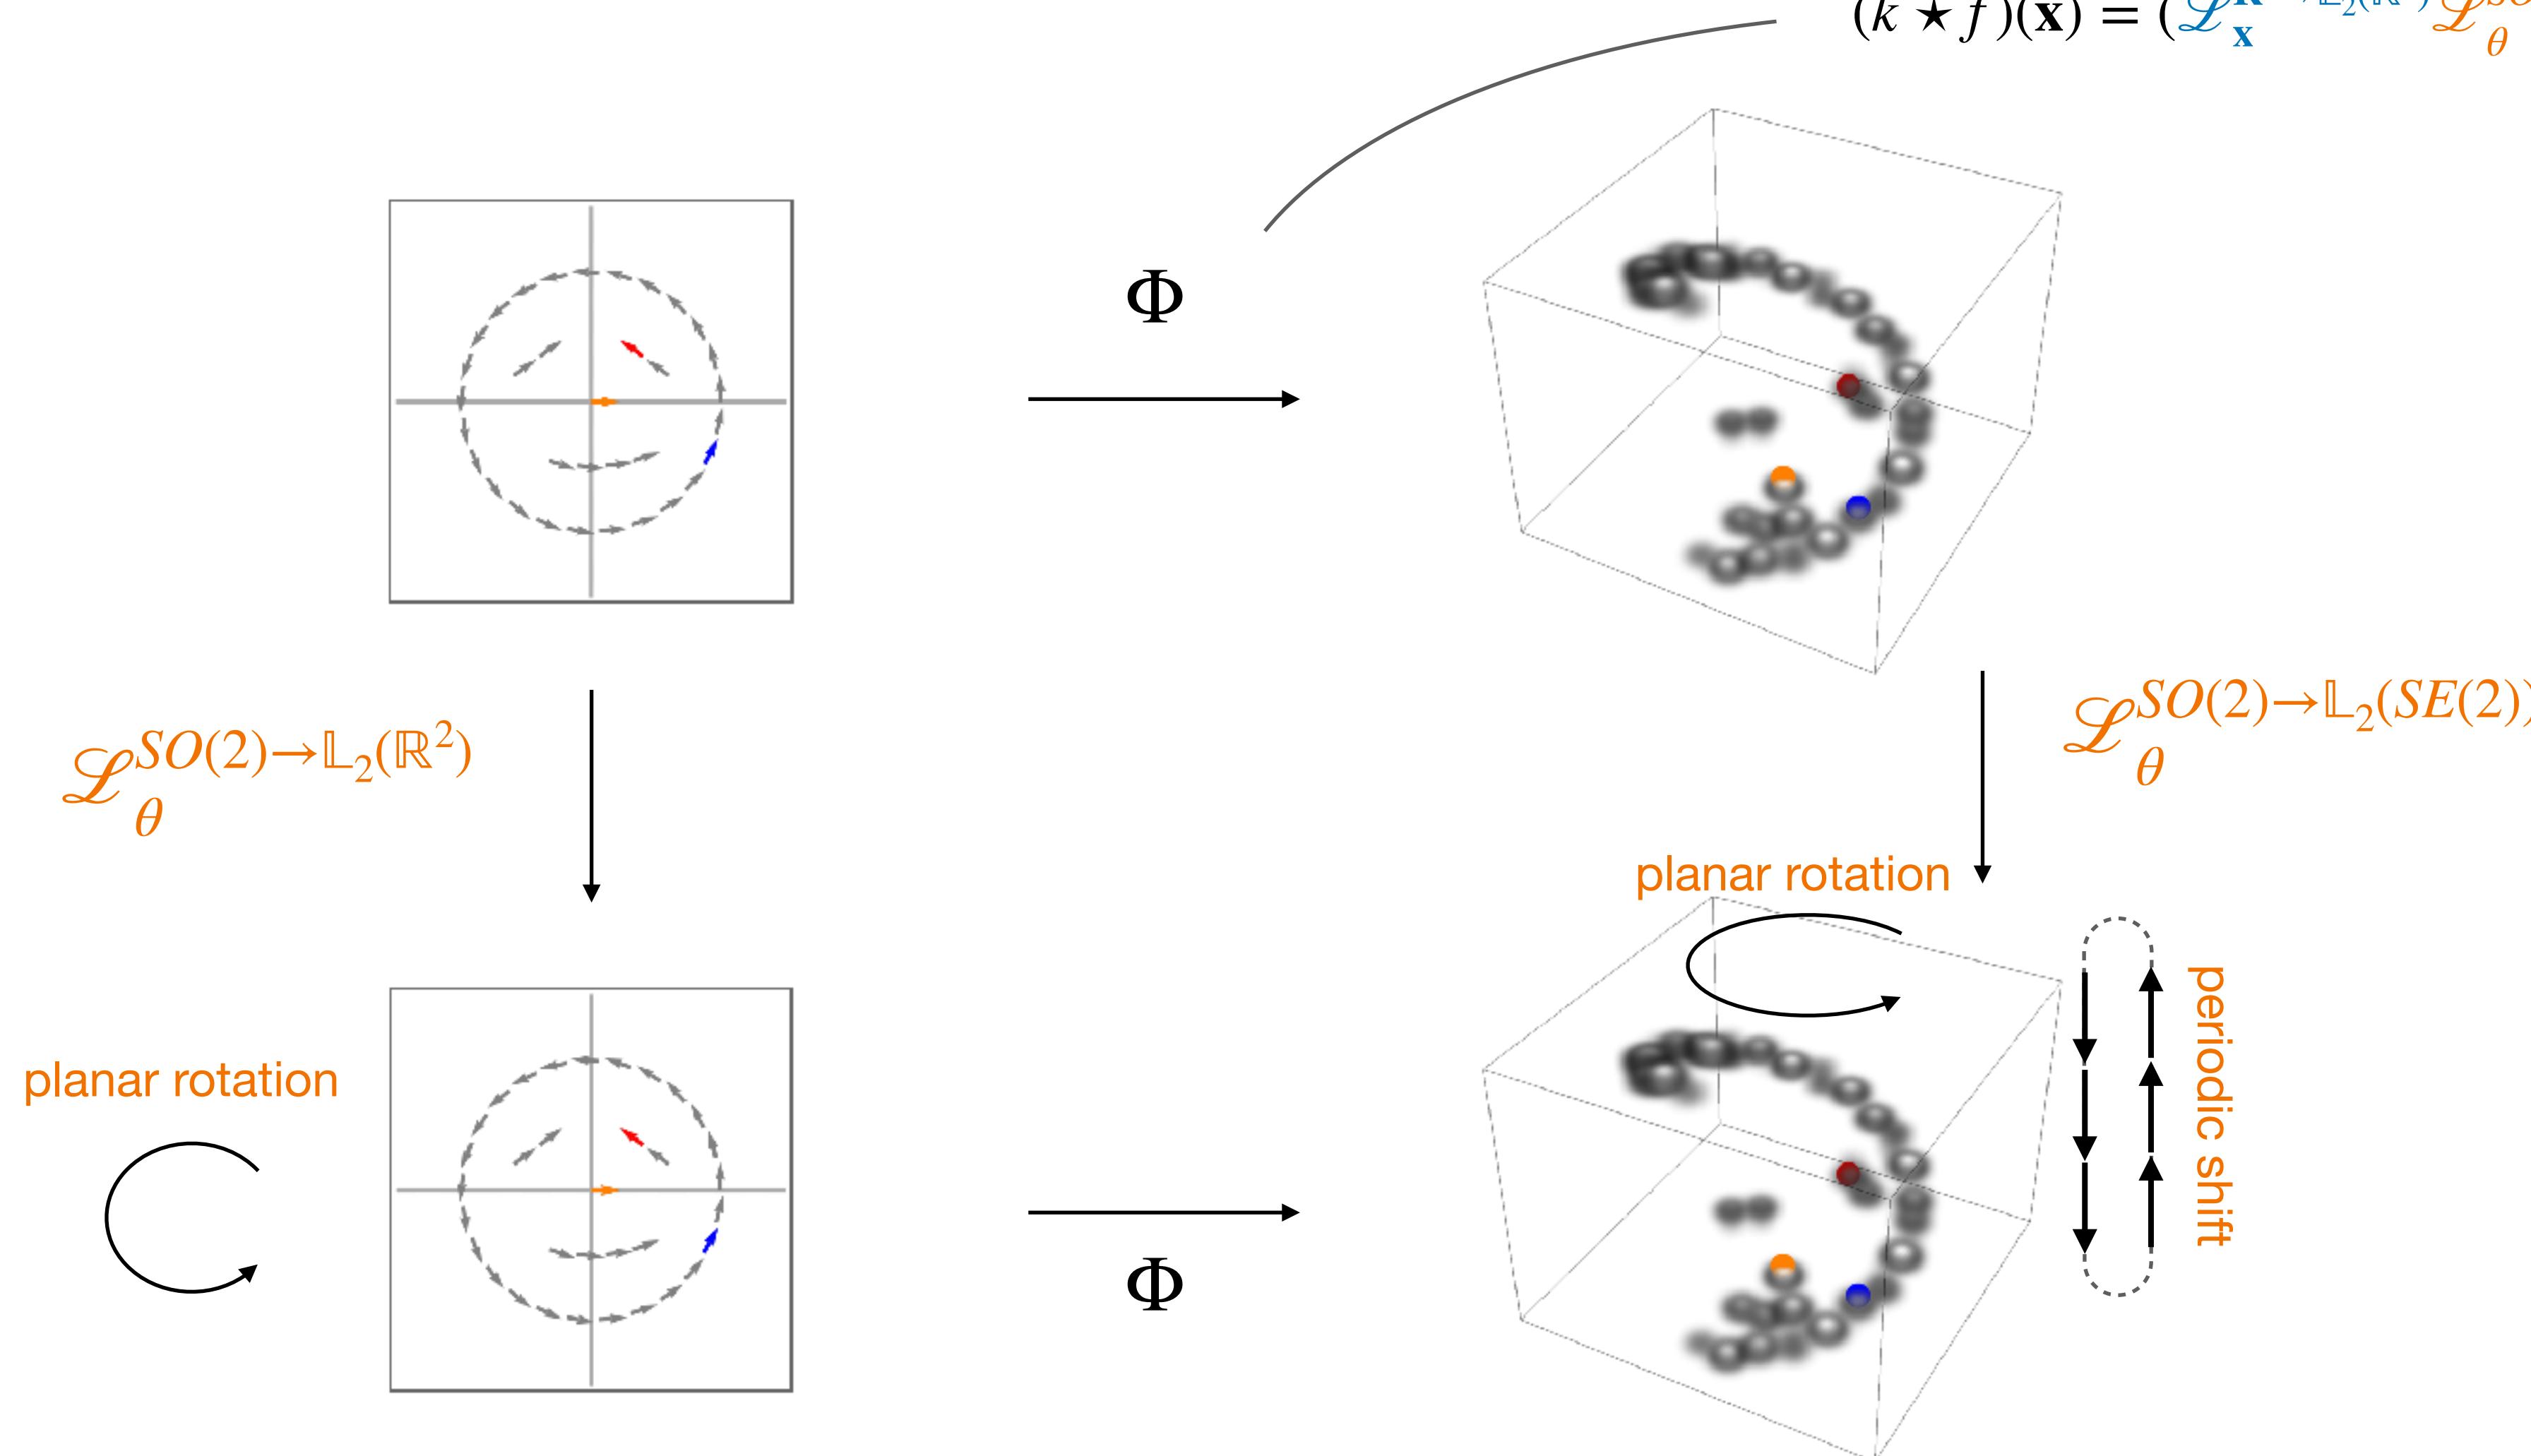
$$(k \tilde{\star} f)(\mathbf{x}) = (\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$$



# Equivariance

SE(2) group **lifting convolutions** are roto-translation equivariant

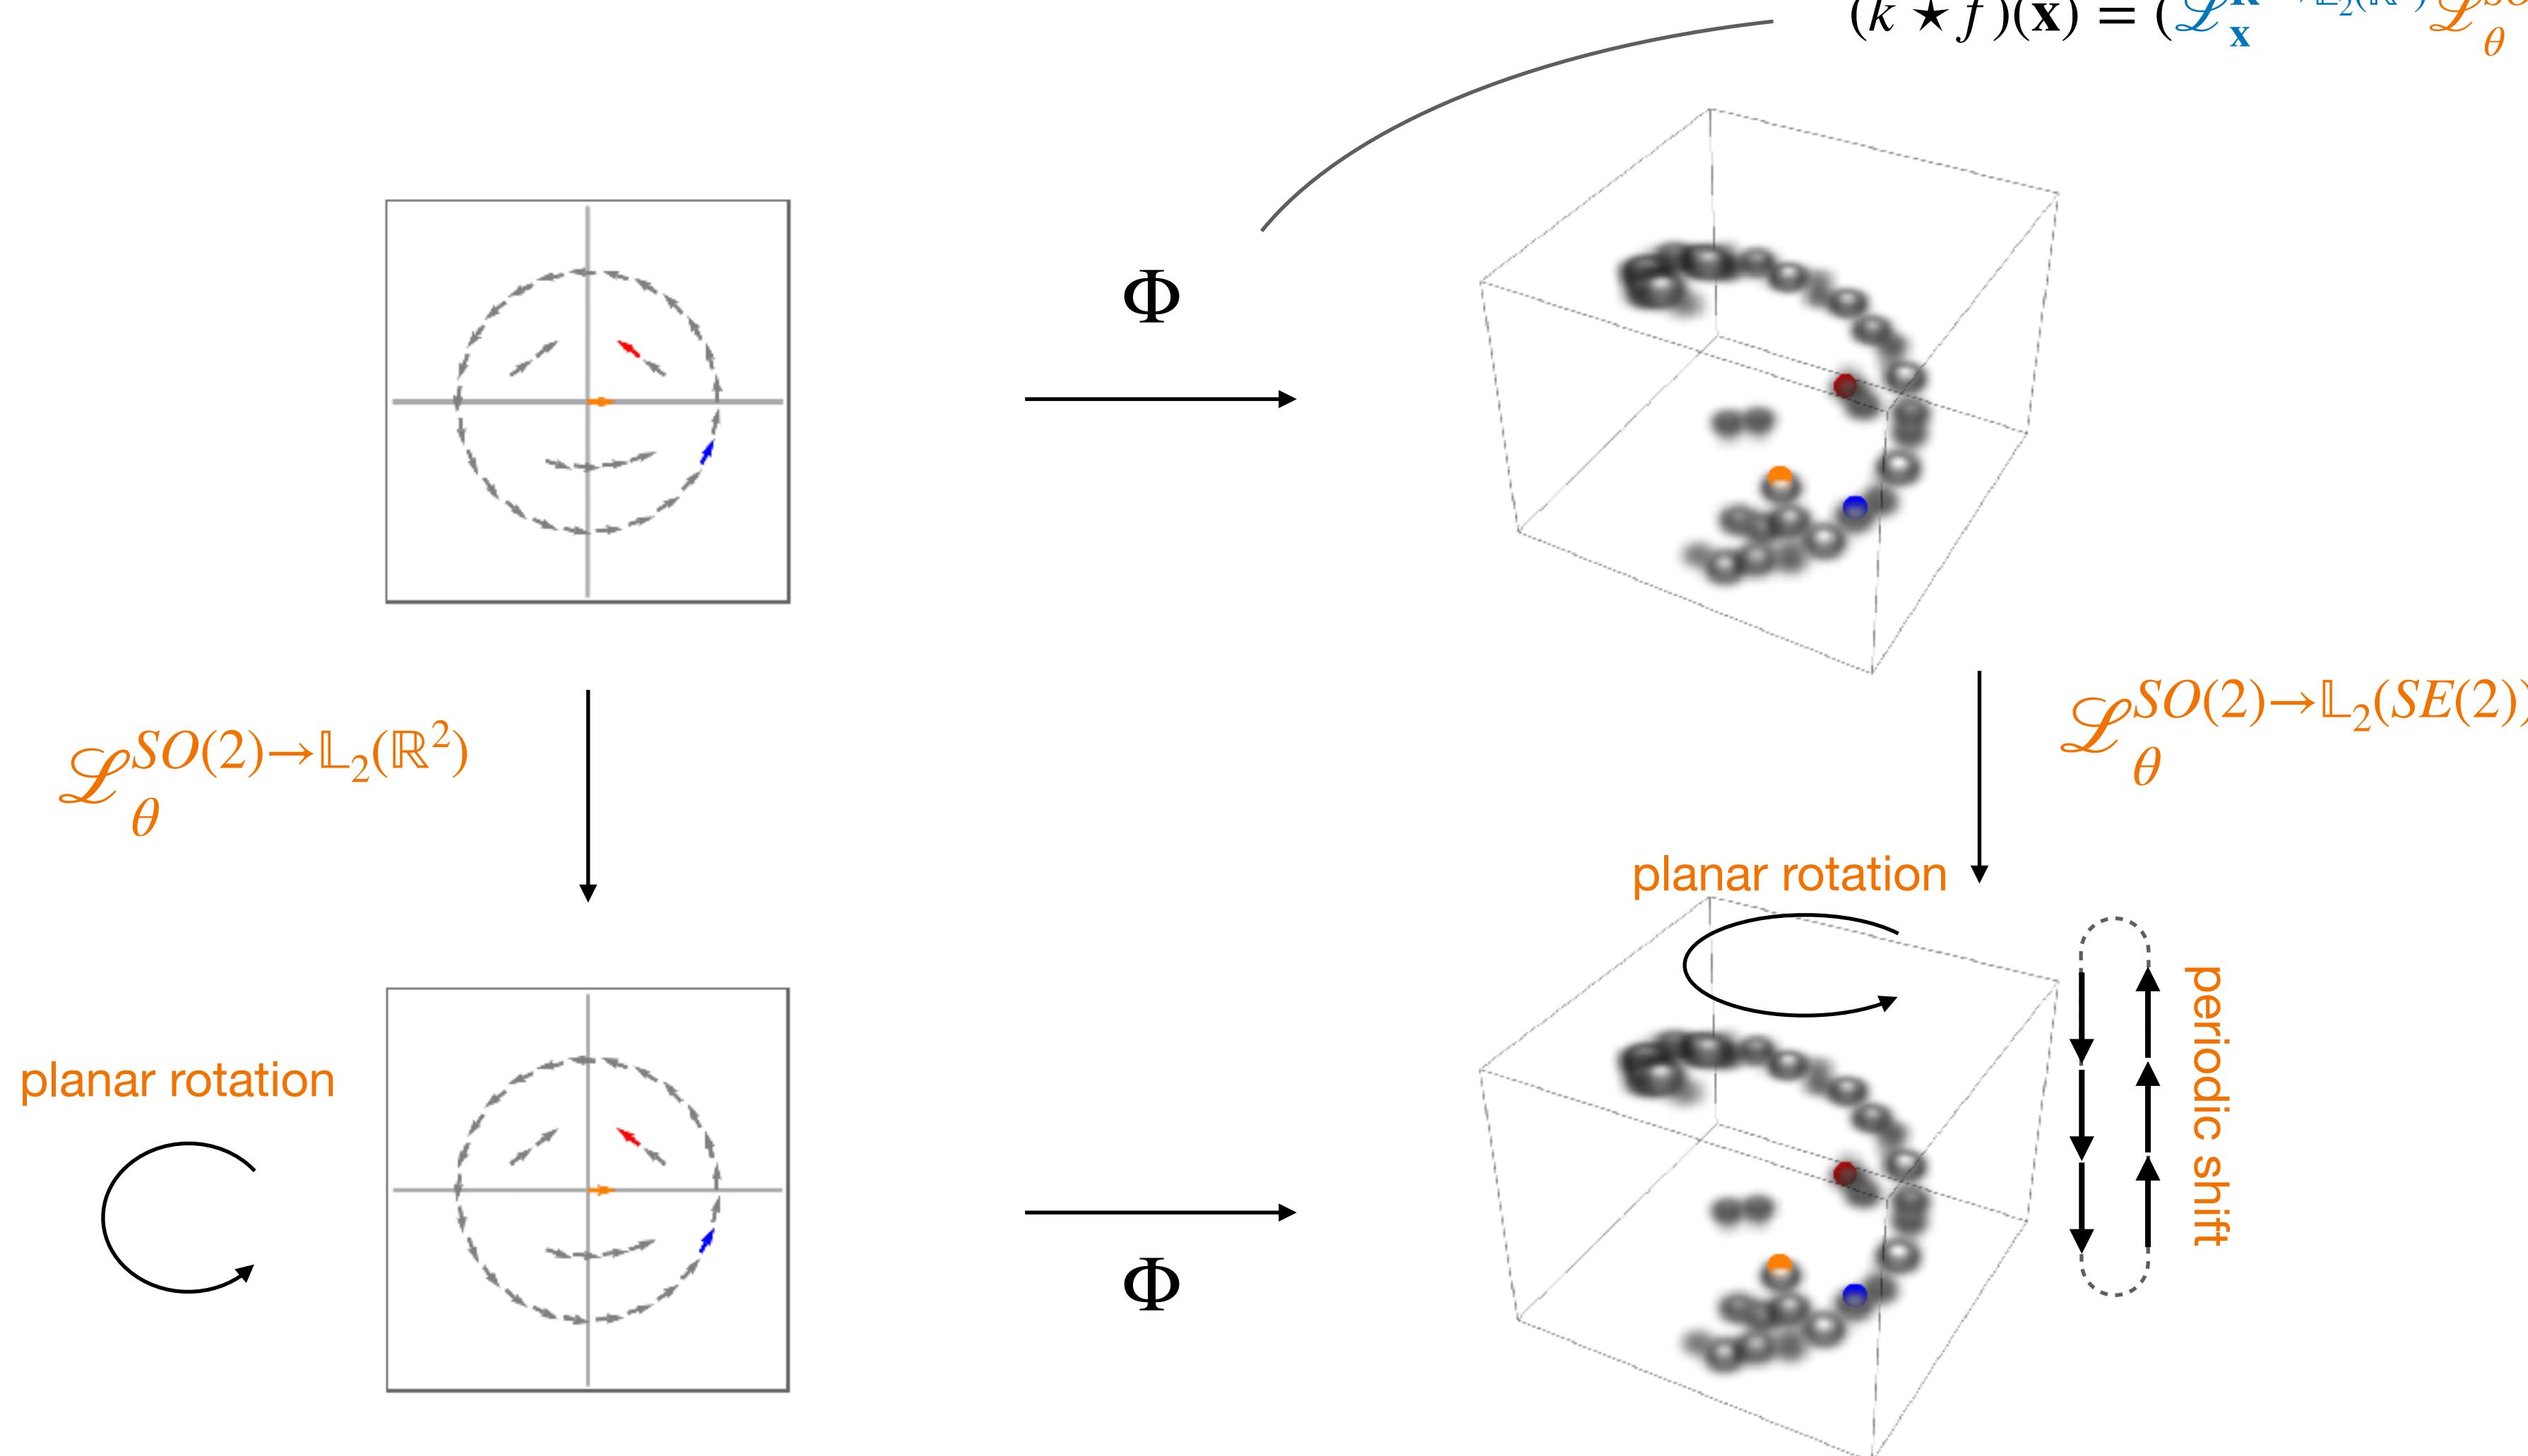
$$(k \tilde{\star} f)(\mathbf{x}) = (\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$$



# Equivariance

SE(2) group **lifting convolutions** are roto-translation equivariant

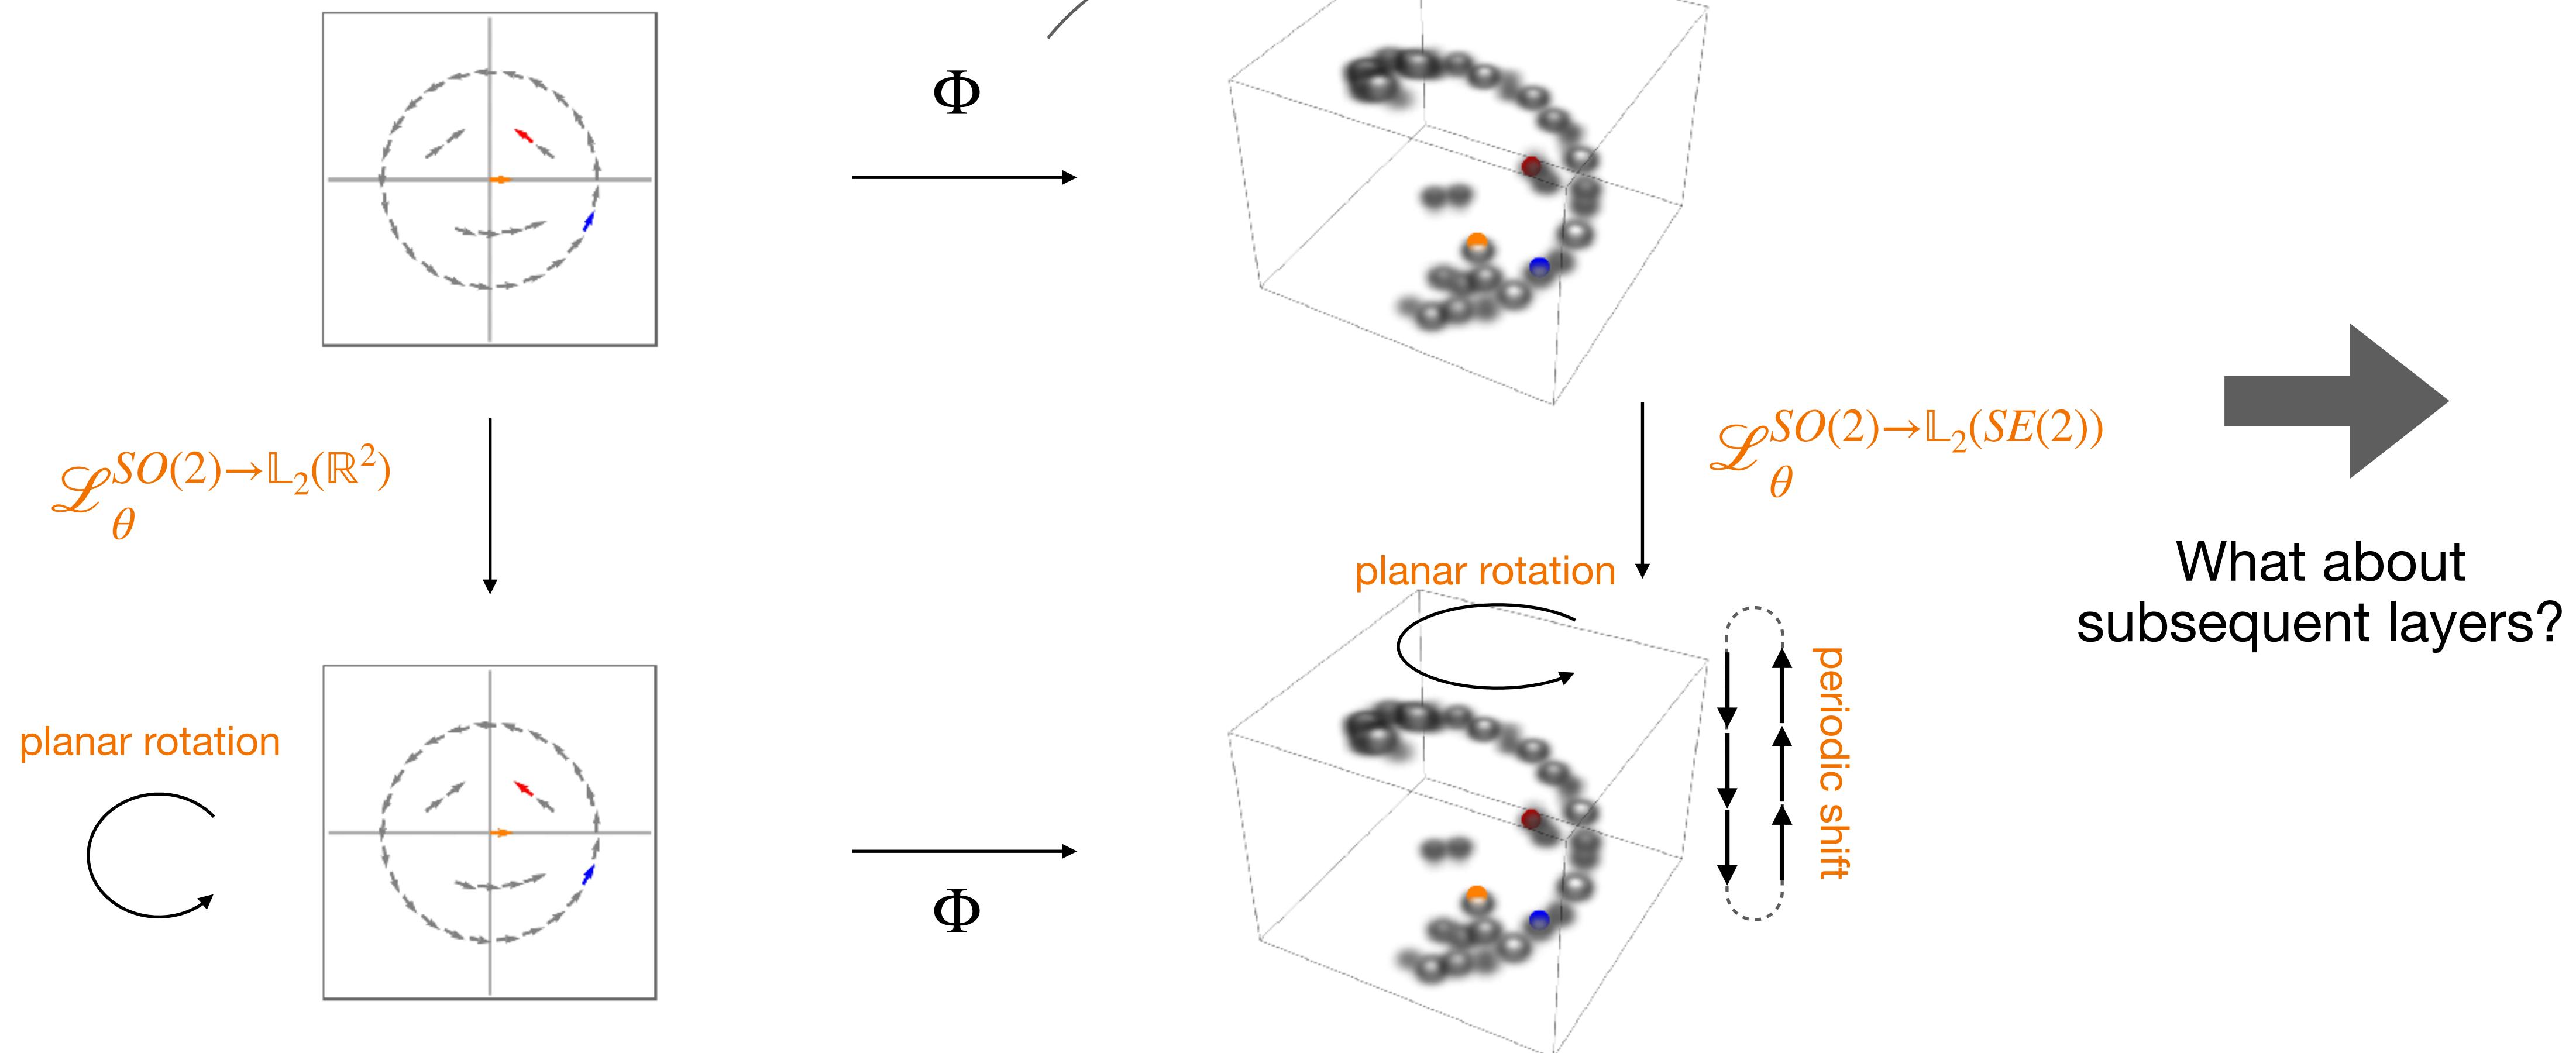
$$(k \tilde{\star} f)(\mathbf{x}) = (\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$$



# Equivariance

SE(2) group **lifting convolutions** are roto-translation equivariant

$$(k \tilde{\star} f)(\mathbf{x}) = (\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(\mathbb{R}^2)} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(\mathbb{R}^2)} k, f)_{\mathbb{L}_2(\mathbb{R}^2)}$$



# SE(2) equivariant cross-correlations

**Group correlations:**

$$(k \star f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(SE(2))} k, f)_{\mathbb{L}_2(SE(2))}$$

# SE(2) equivariant cross-correlations

$$k(\mathbf{R}_\theta^{-1}(\mathbf{x}' - \mathbf{x}), \mathbf{R}_{\theta' - \theta})$$

**Group correlations:**

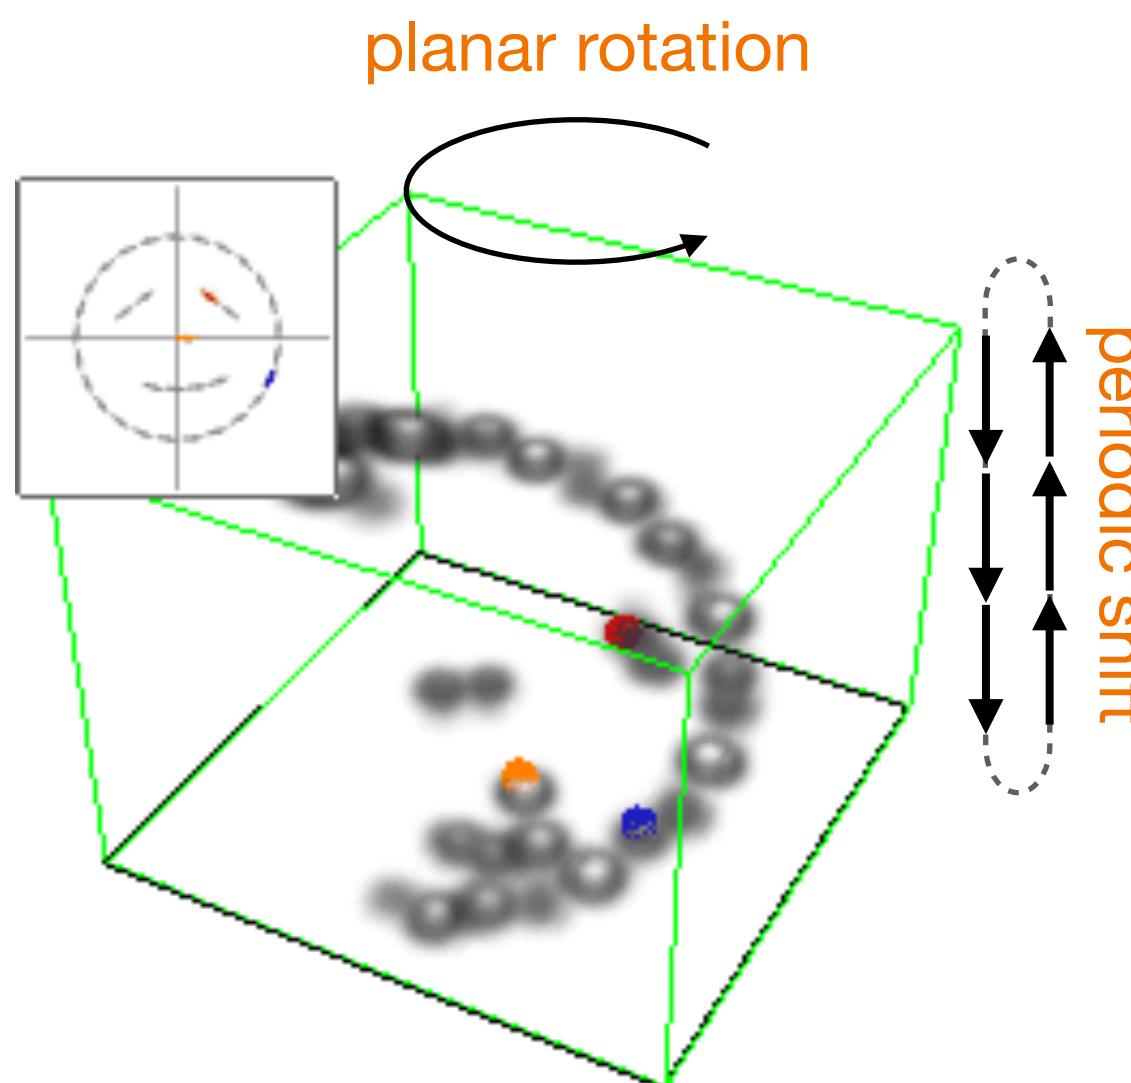
$$(k \star f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(SE(2))} k, f)_{\mathbb{L}_2(SE(2))} = (\underbrace{\mathcal{L}_{\mathbf{x}}^{\mathbf{R}^2 \rightarrow \mathbb{L}_2(SE(2))} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(SE(2))} k, f}_{\text{translation } \text{rotation}})_{\mathbb{L}_2(SE(2))}$$

# SE(2) equivariant cross-correlations

$$k(\mathbf{R}_\theta^{-1}(\mathbf{x}' - \mathbf{x}), \mathbf{R}_{\theta' - \theta})$$

**Group correlations:**

$$(k \star f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(SE(2))} k, f)_{\mathbb{L}_2(SE(2))} = (\underbrace{\mathcal{L}_{\mathbf{x}}^{\mathbf{R}^2 \rightarrow \mathbb{L}_2(SE(2))} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(SE(2))} k, f}_{\text{translation } \text{rotation}})_{\mathbb{L}_2(SE(2))}$$



$$\mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(SE(2))} k$$

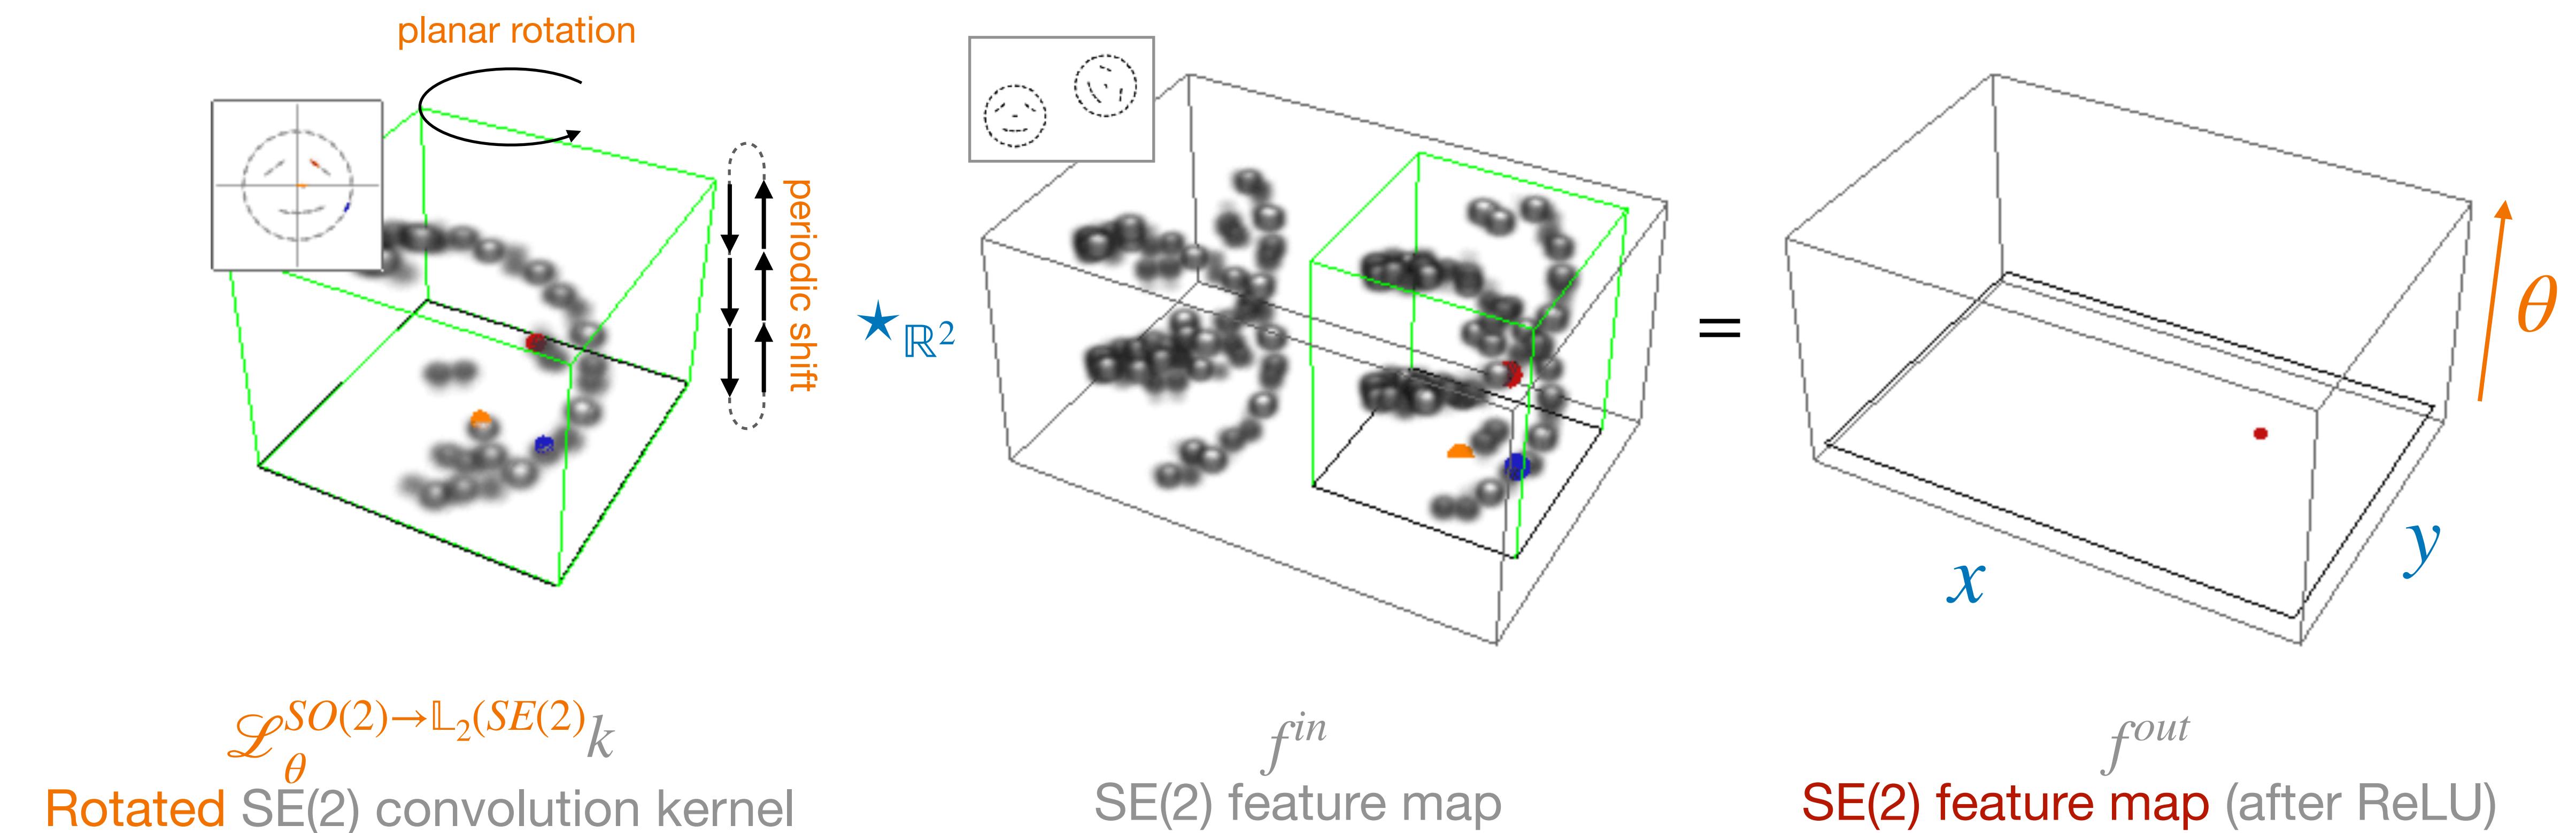
Rotated SE(2) convolution kernel

# SE(2) equivariant cross-correlations

$$(k \star f)(\mathbf{x}, \theta) = \underbrace{(k, f)_{\mathbb{L}_2(SE(2))}}_{k(\mathbf{R}_{\theta}^{-1}(\mathbf{x}' - \mathbf{x}), \mathbf{R}_{\theta' - \theta})}$$

**Group correlations:**

$$(k \star f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(SE(2))} k, f)_{\mathbb{L}_2(SE(2))} = (\underbrace{\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(SE(2))} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(SE(2))} k, f}_{\text{translation } \text{rotation}})_{\mathbb{L}_2(SE(2))}$$



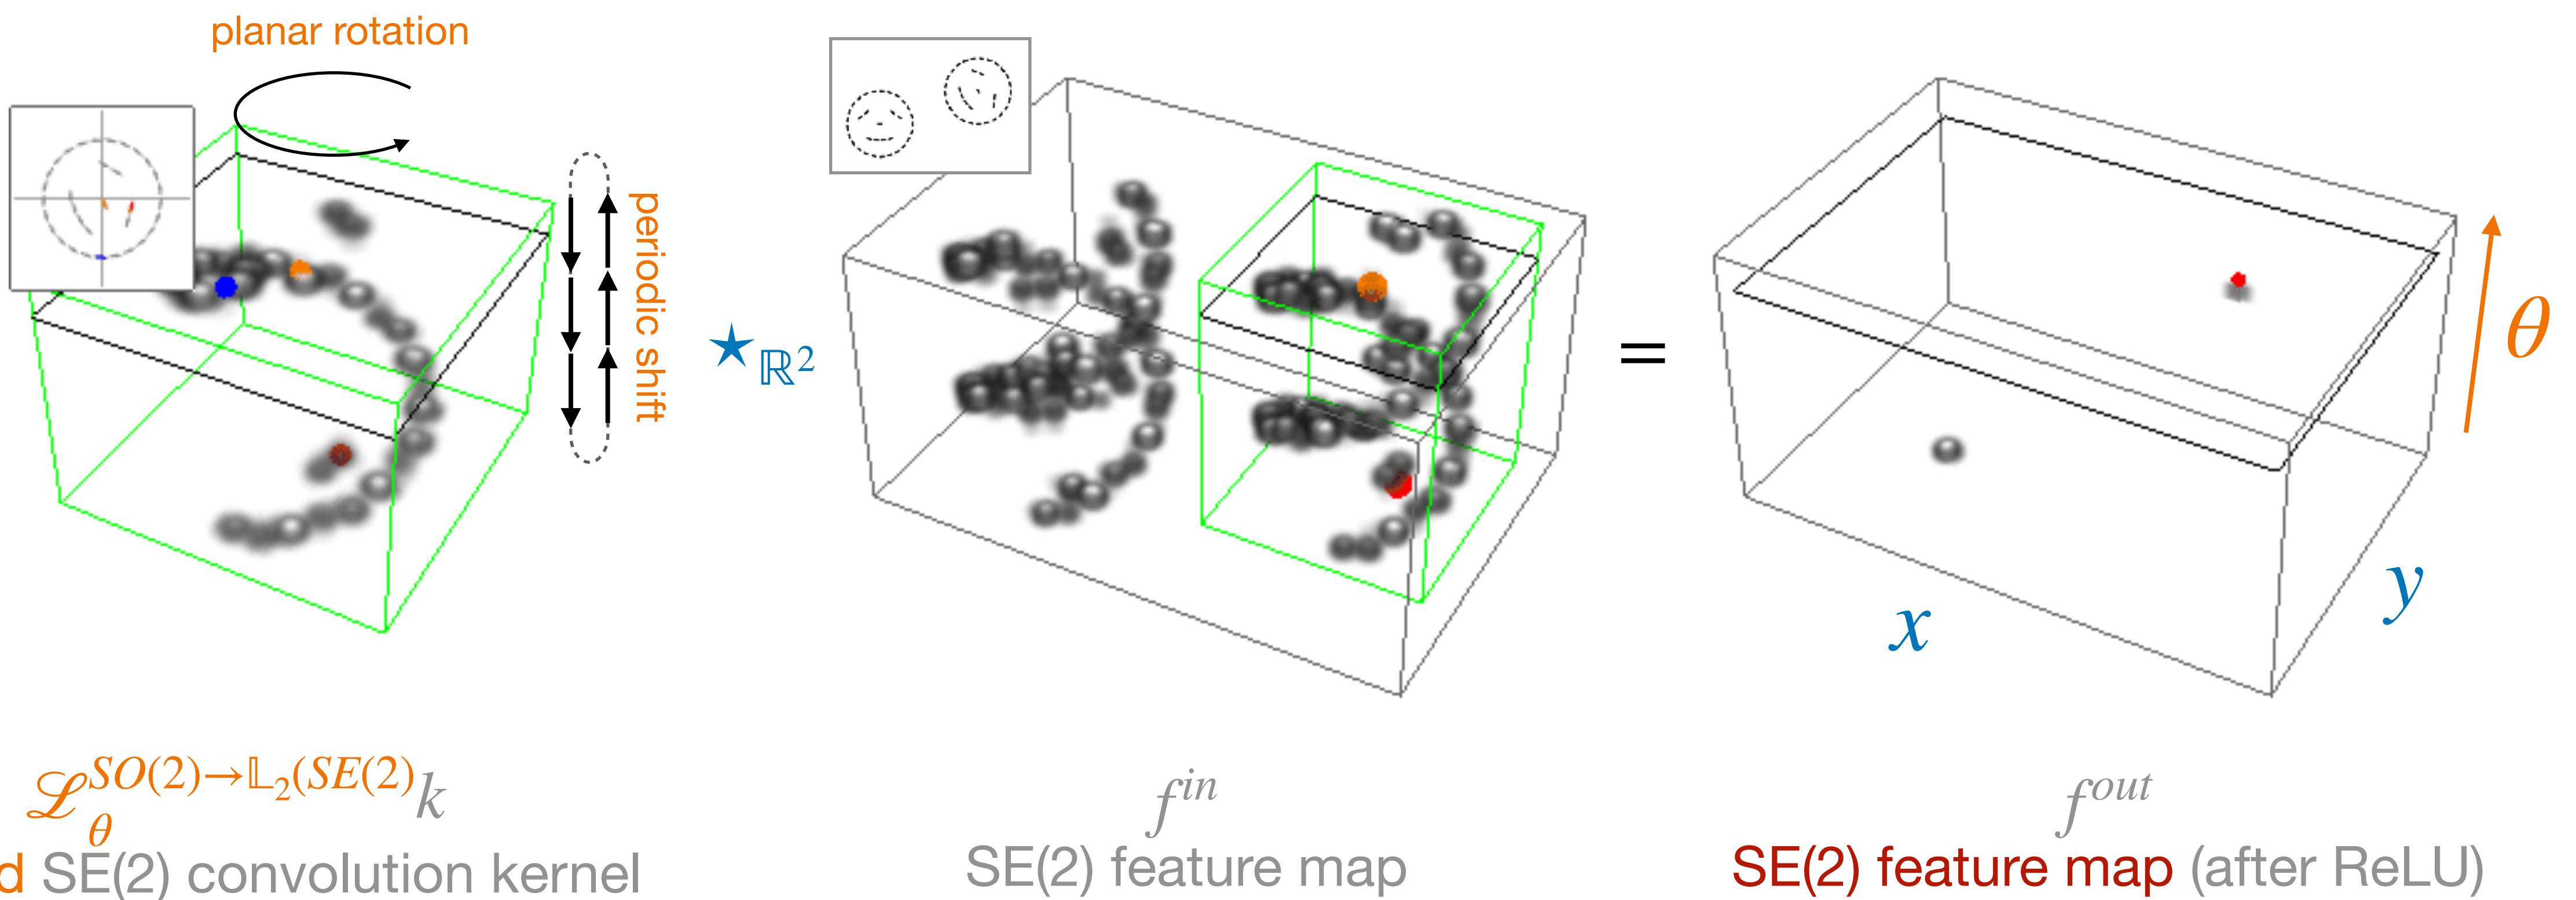
# SE(2) equivariant cross-correlations

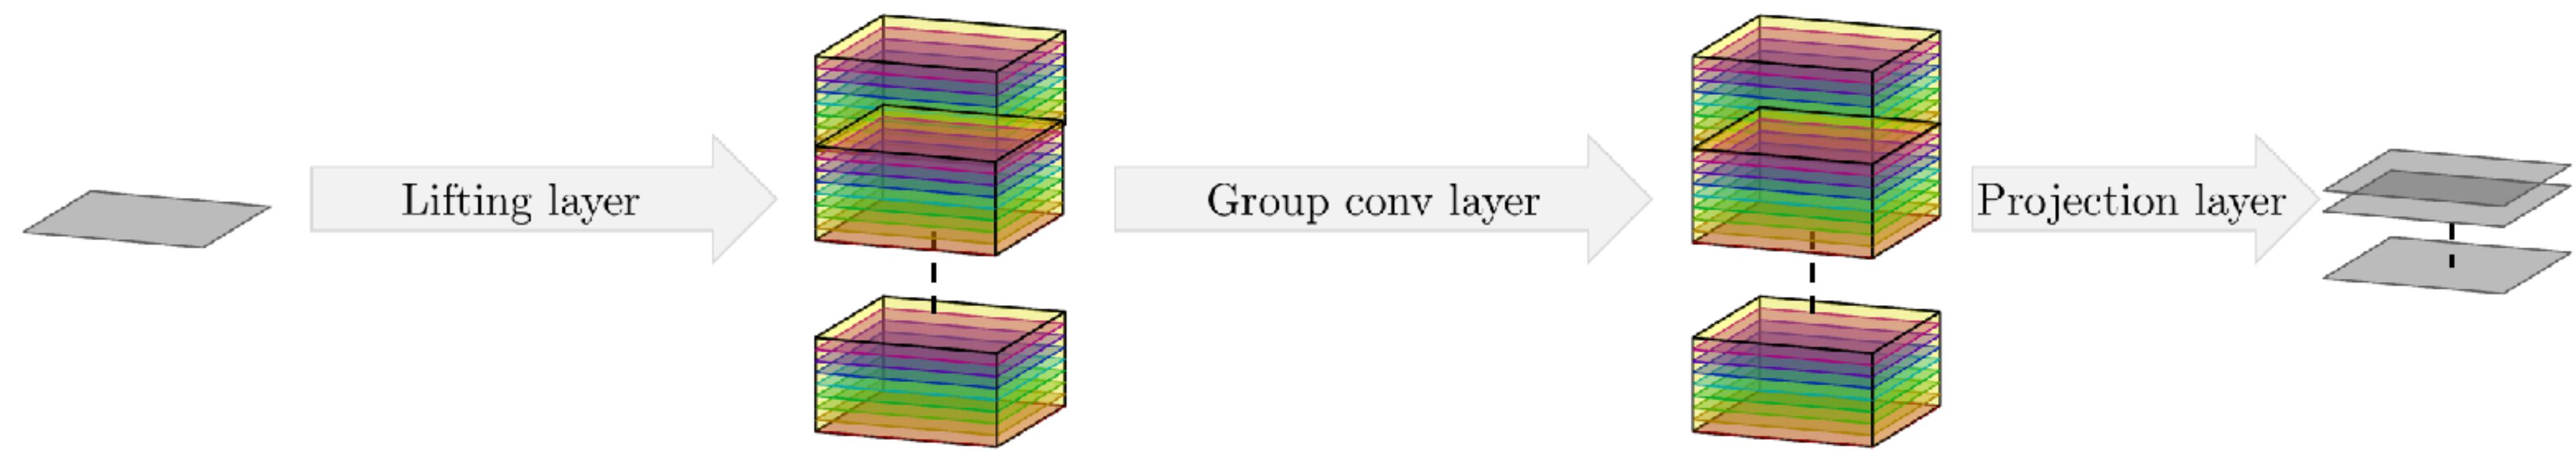
$$(k \star f)(\mathbf{x}, \theta) = (\mathcal{L}_g^{SE(2) \rightarrow \mathbb{L}_2(SE(2))} k, f)_{\mathbb{L}_2(SE(2))} = (\underbrace{\mathcal{L}_{\mathbf{x}}^{\mathbb{R}^2 \rightarrow \mathbb{L}_2(SE(2))} \mathcal{L}_{\theta}^{SO(2) \rightarrow \mathbb{L}_2(SE(2))} k, f)_{\mathbb{L}_2(SE(2))}$$

*translation*      *rotation*

$$k(\mathbf{R}_{\theta}^{-1}(\mathbf{x}' - \mathbf{x}), \mathbf{R}_{\theta' - \theta})$$

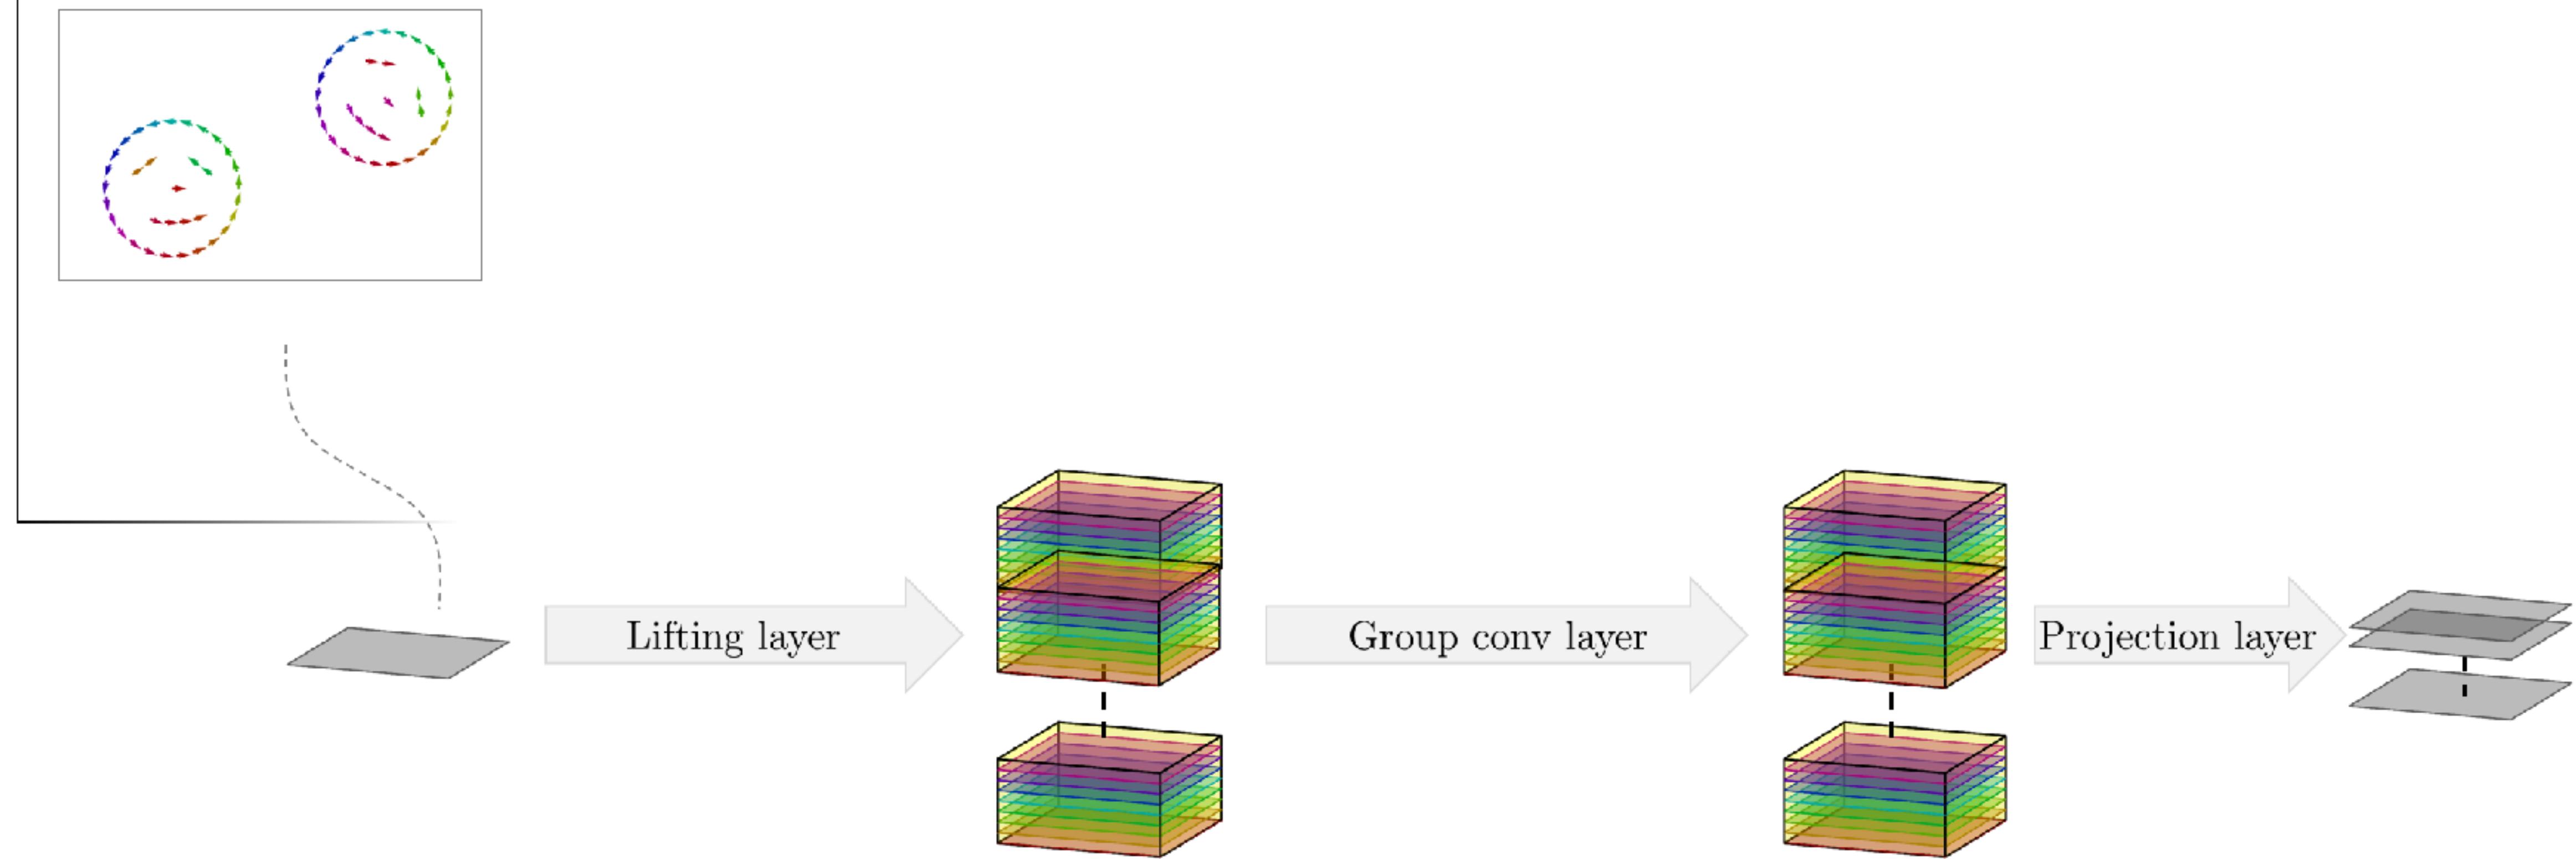
**Group correlations:**





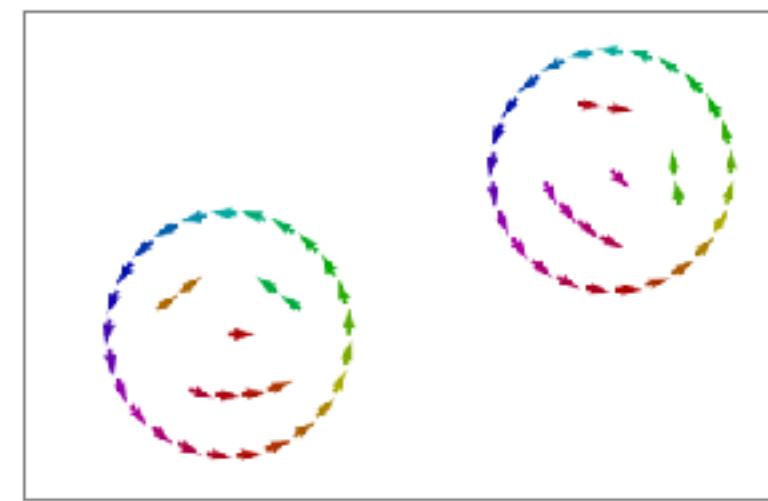
Roto-translation group  $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



## Roto-translation group $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



Using a set of transformed  
2D conv kernels

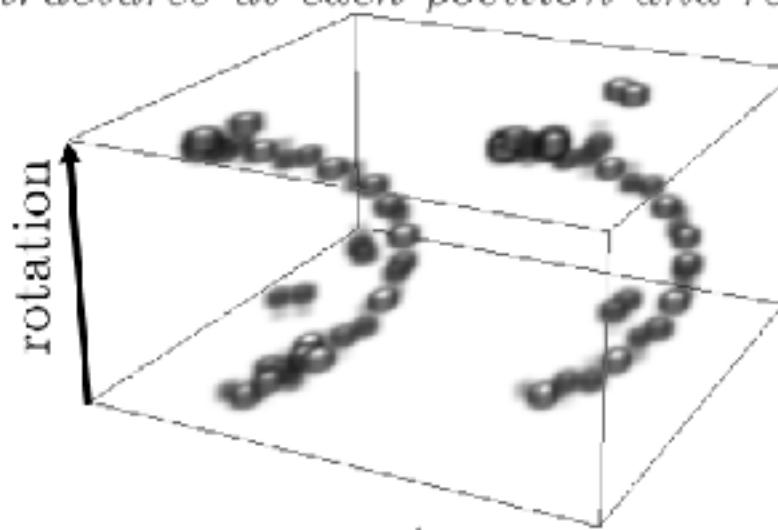
$$\theta = \frac{\pi}{2}$$

$$\theta = \frac{\pi}{4}$$

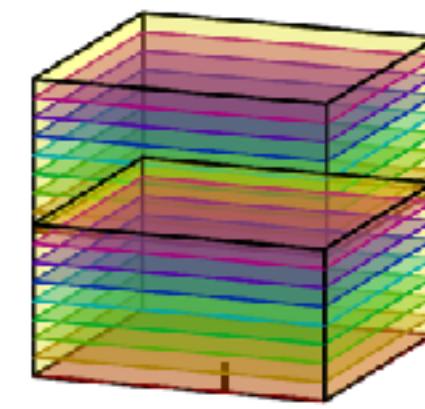
$$\theta = 0$$



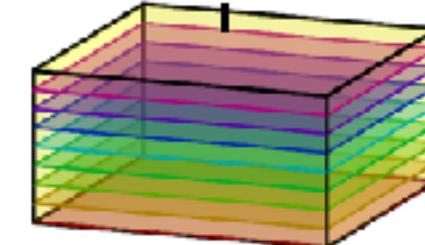
*G* feature map (activation for oriented  
structures at each position and rotation)



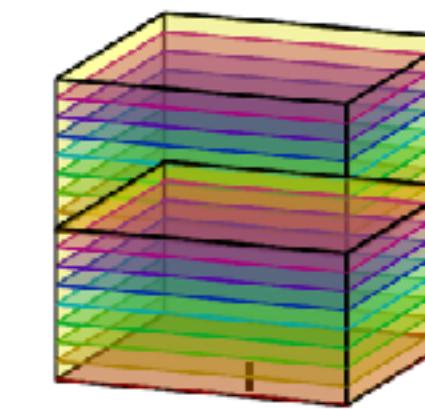
*G*-feature maps are equivariant  
w.r.t. translation and rotation  
of the input



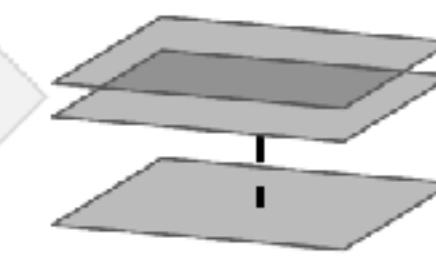
Lifting layer



Group conv layer

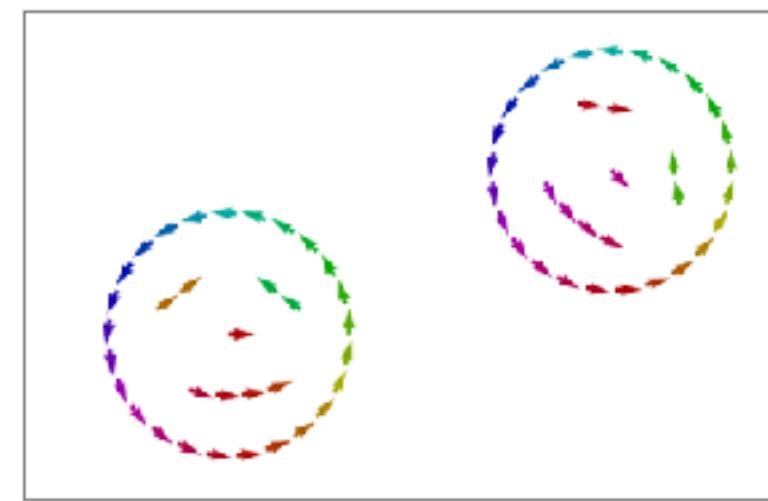


Projection layer



## Roto-translation group $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



Using a set of transformed  
2D conv kernels

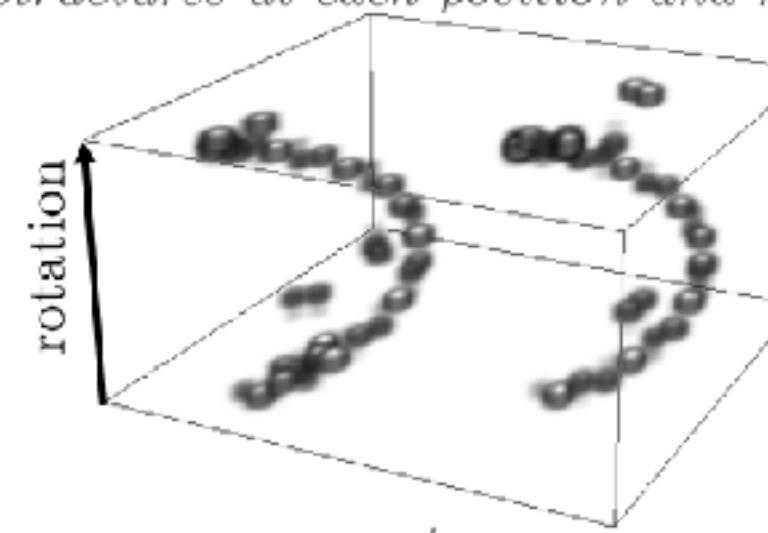
$$\theta = \frac{\pi}{2}$$

$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$



*G* feature map (activation for oriented  
structures at each position and rotation)

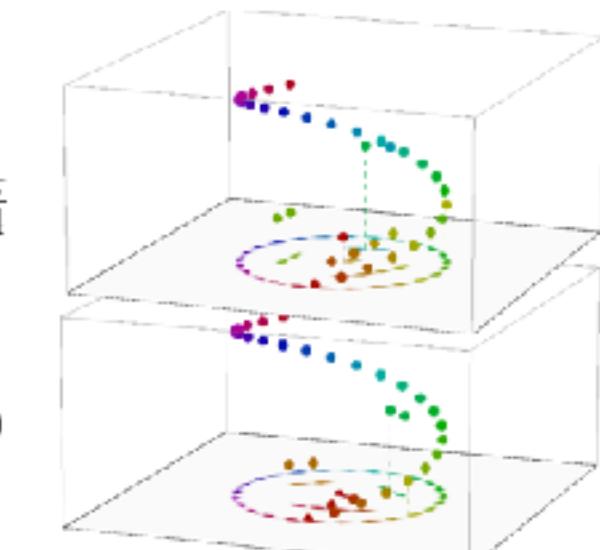


*G*-feature maps are equivariant  
w.r.t. translation and rotation  
of the input

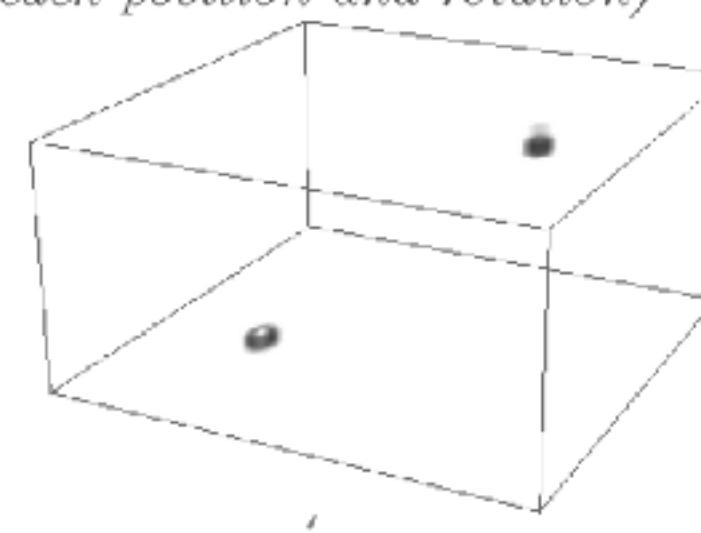
Using a set of transformed  
*G*-conv kernels

$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$

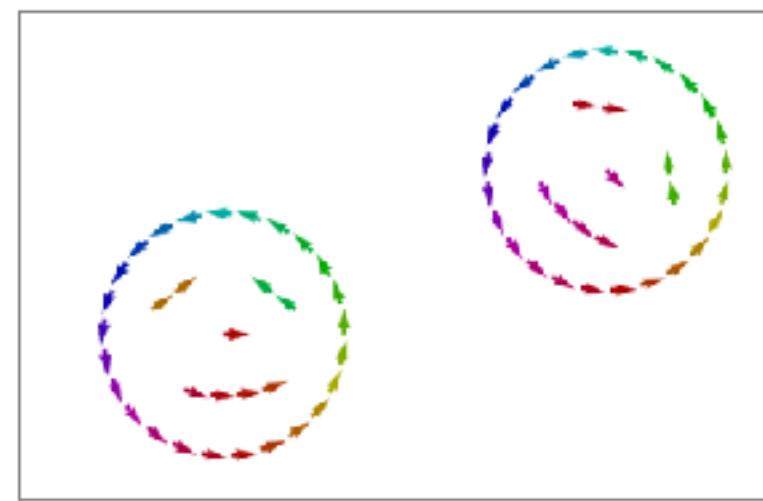


*G* feature map (activation for faces  
at each position and rotation)



# Roto-translation group $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



*Using a set of transformed  
2D conv kernels*

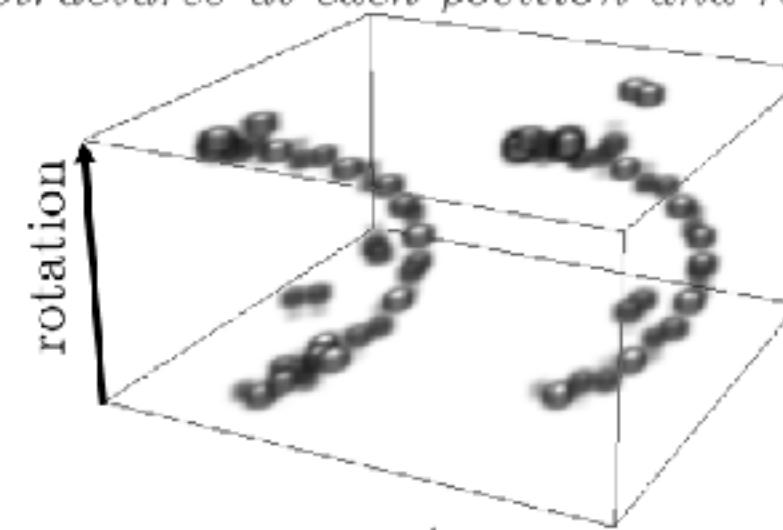
$$\theta = \frac{\pi}{2}$$

$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$



*G feature map (activation for oriented  
structures at each position and rotation)*

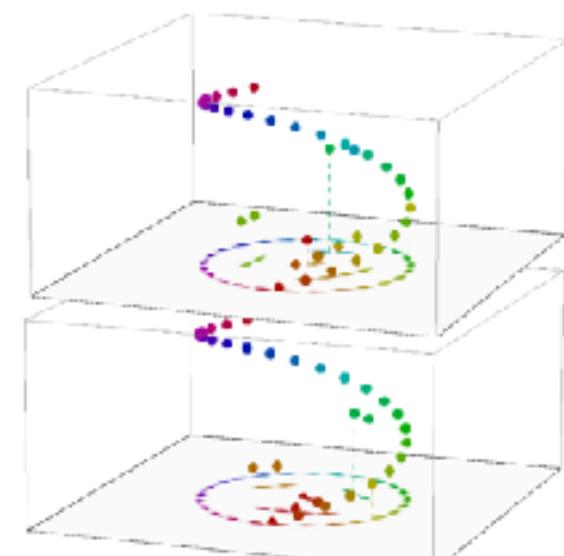


*G-feature maps are equivariant  
w.r.t. translation and rotation  
of the input*

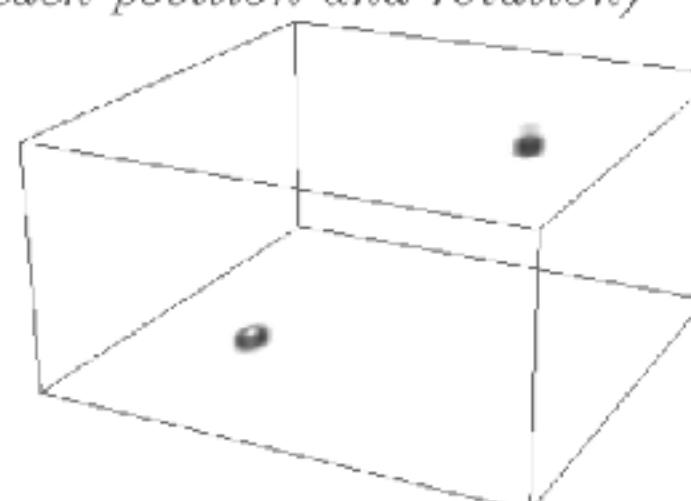
*Using a set of transformed  
G-conv kernels*

$$\theta = \frac{\pi}{4}$$

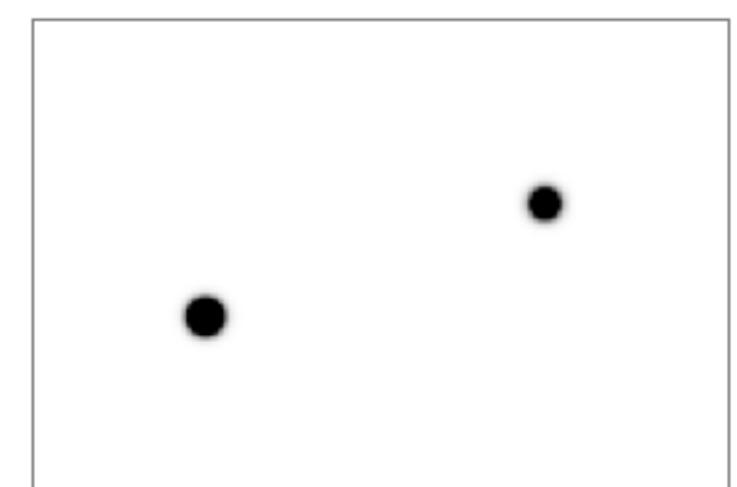
$$\theta = 0$$



*G feature map (activation for faces  
at each position and rotation)*

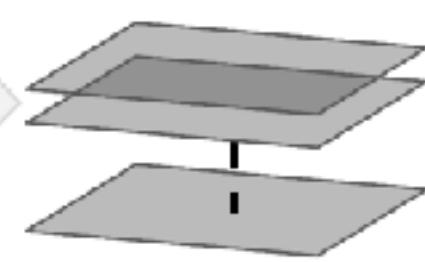


2D feature map

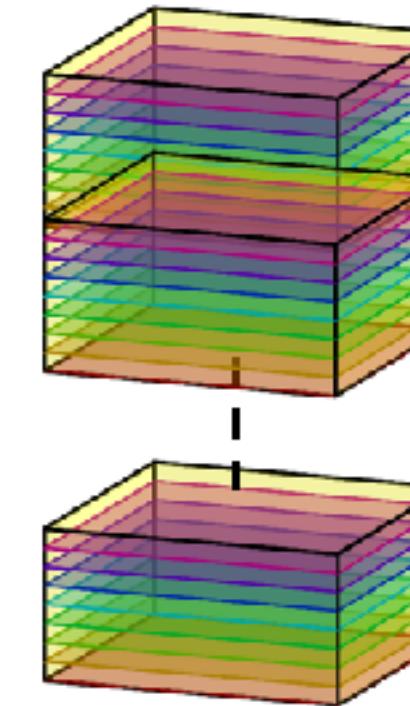


*Projection over sub-group  $H$   
guarantees local invariance*

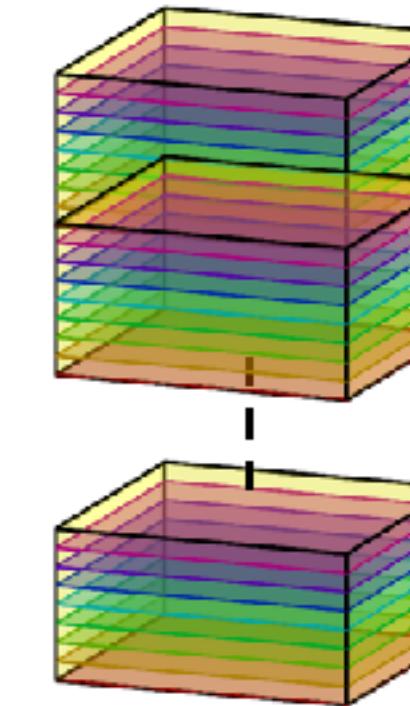
*Projection layer*



*Lifting layer*



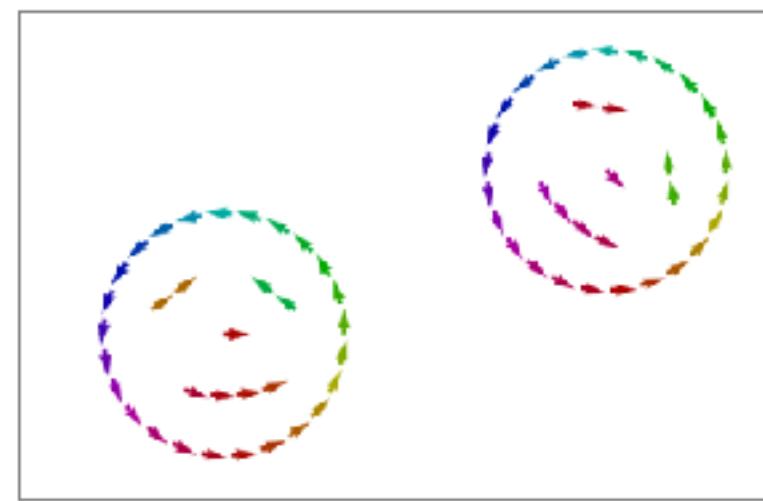
*Group conv layer*



*Projection layer*

## Roto-translation group $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



*Using a set of transformed  
2D conv kernels*

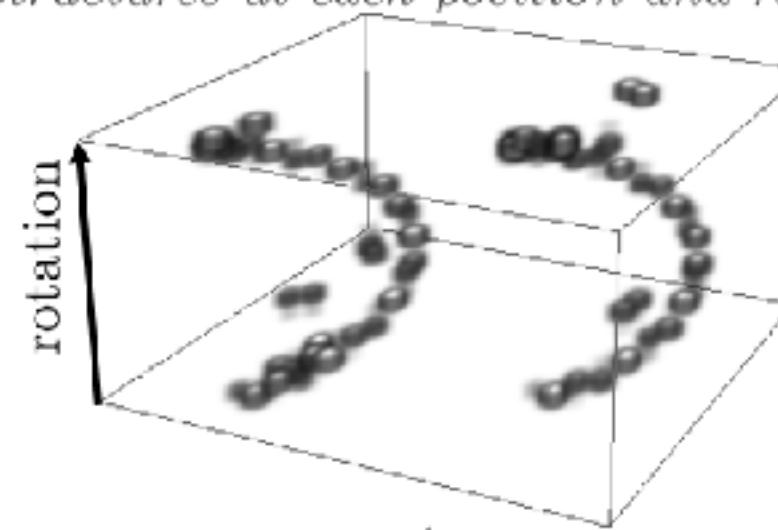
$$\theta = \frac{\pi}{2}$$

$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$



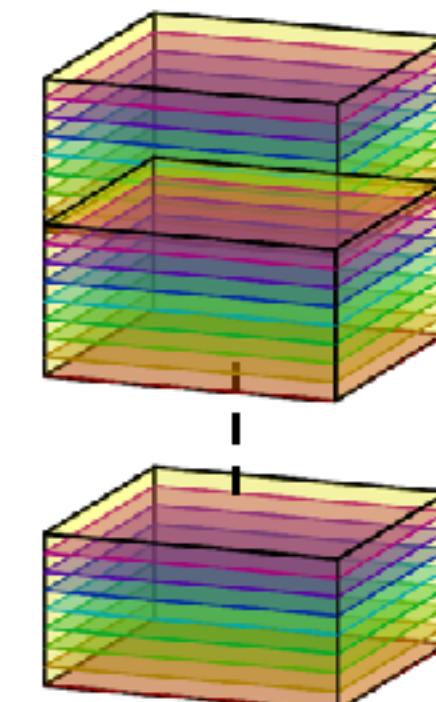
*G feature map (activation for oriented  
structures at each position and rotation)*



*Using a set of transformed  
G-conv kernels*

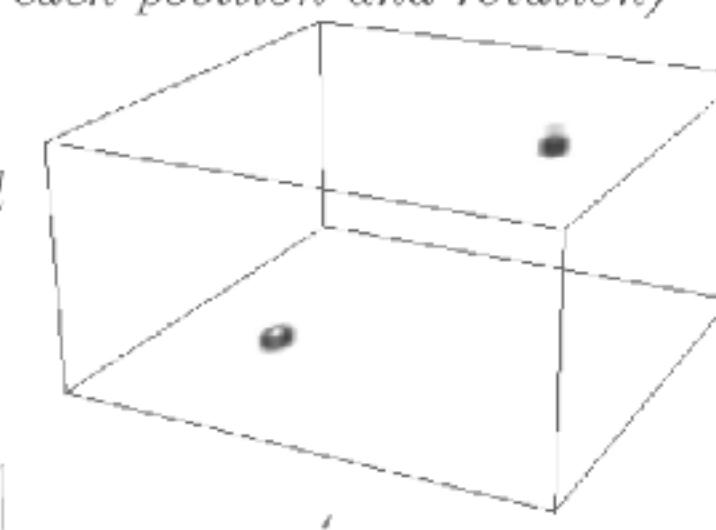
$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$

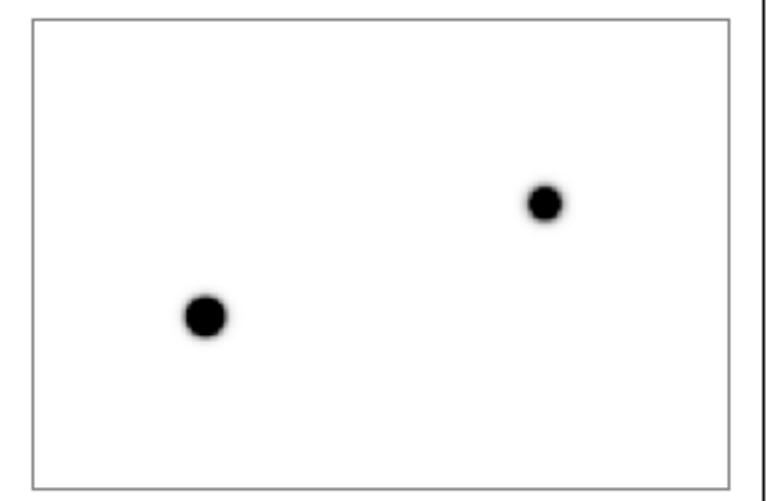


*G-feature maps are equivariant  
w.r.t. translation and rotation  
of the input*

*G feature map (activation for faces  
at each position and rotation)*

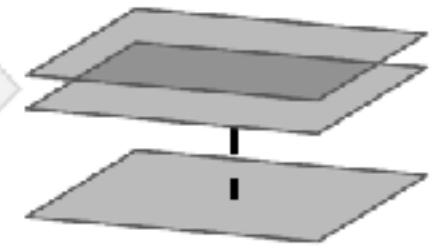


2D feature map

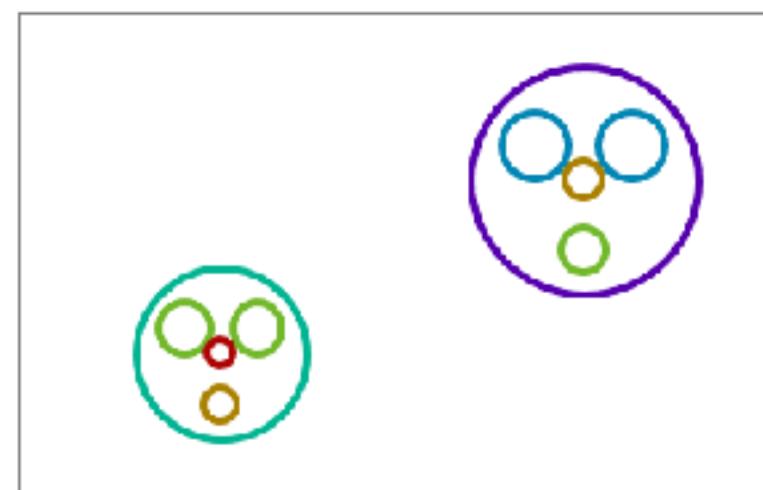


*Projection over sub-group  $H$   
guarantees local invariance*

Projection layer

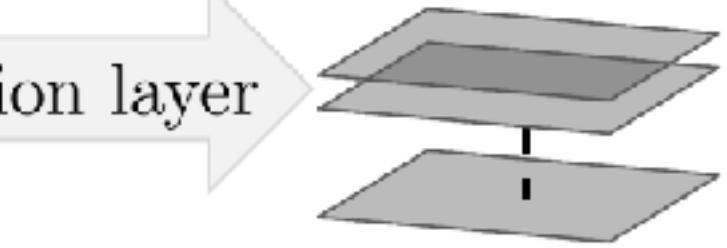


Scale-translation  
group  $\mathbb{R}^2 \rtimes \mathbb{R}^+$



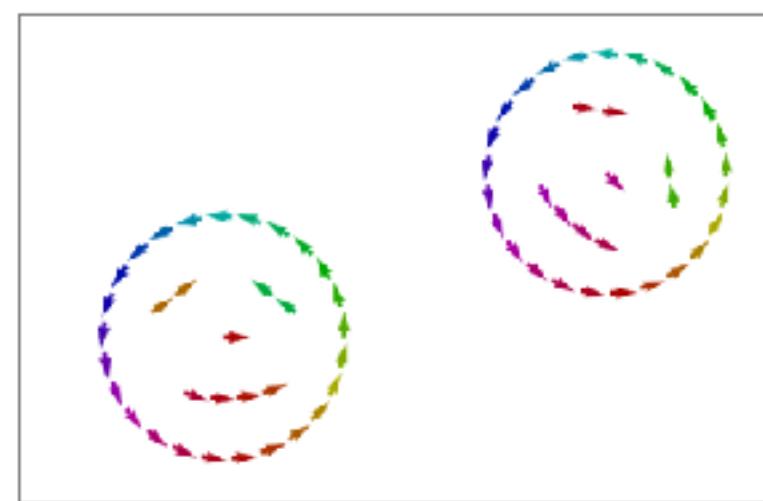
Lifting layer

Group conv layer



# Roto-translation group $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



*Using a set of transformed 2D conv kernels*

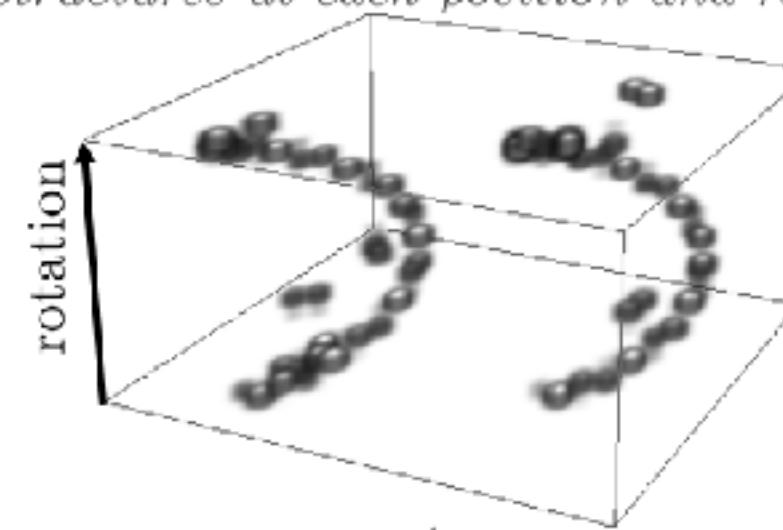
$$\theta = \frac{\pi}{2}$$

$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$



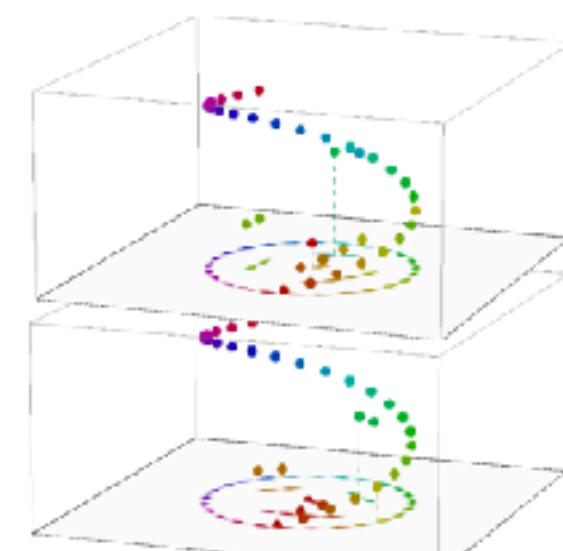
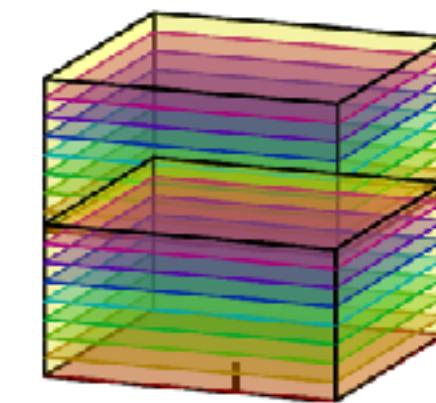
*G feature map (activation for oriented structures at each position and rotation)*



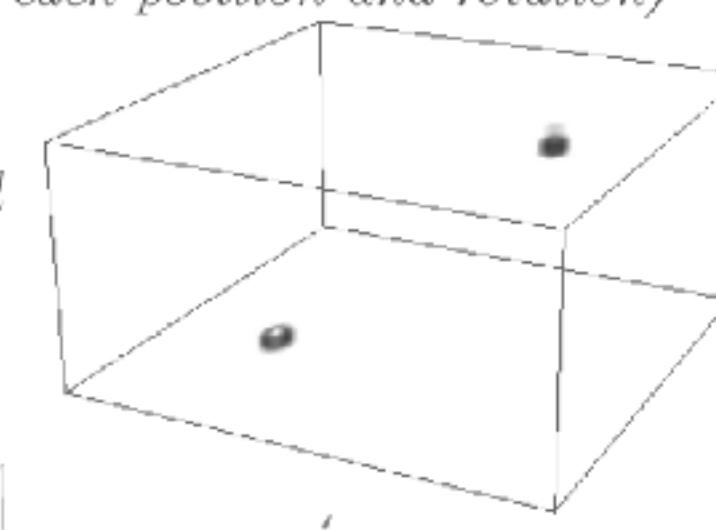
*Using a set of transformed G-conv kernels*

$$\theta = \frac{\pi}{4}$$

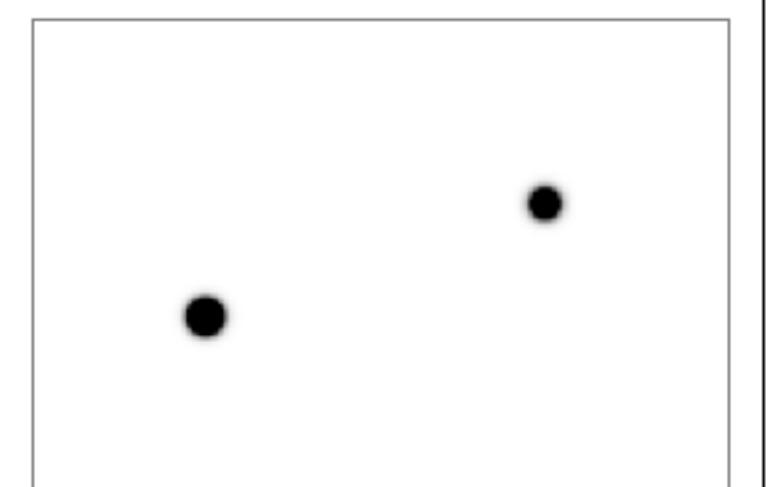
$$\theta = 0$$



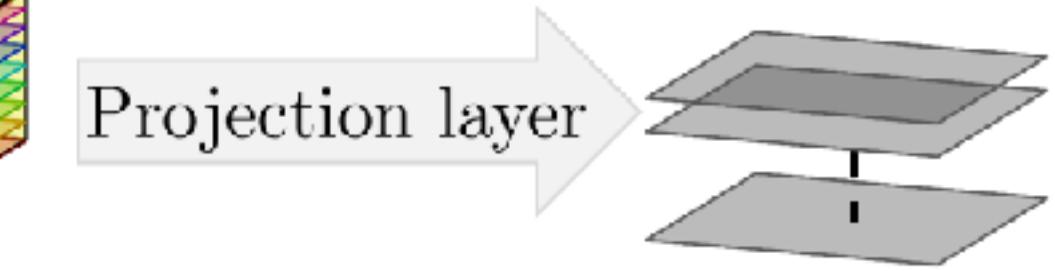
*G feature map (activation for faces at each position and rotation)*



2D feature map



*Projection over sub-group  $H$  guarantees local invariance*

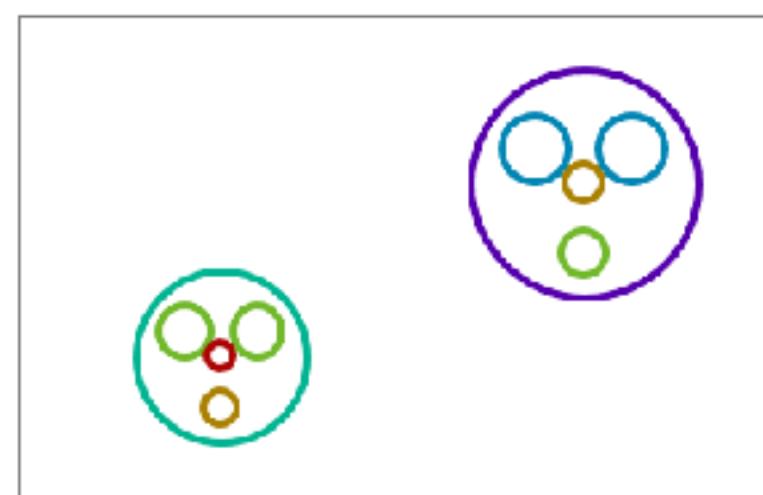


Lifting layer

Group conv layer

Projection layer

Scale-translation group  $\mathbb{R}^2 \rtimes \mathbb{R}^+$



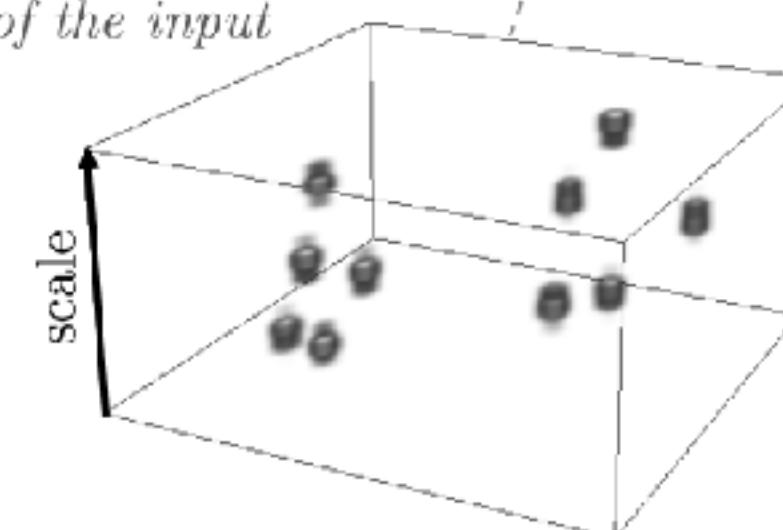
$$s = 2$$

$$s = 1.4$$

$$s = 1$$



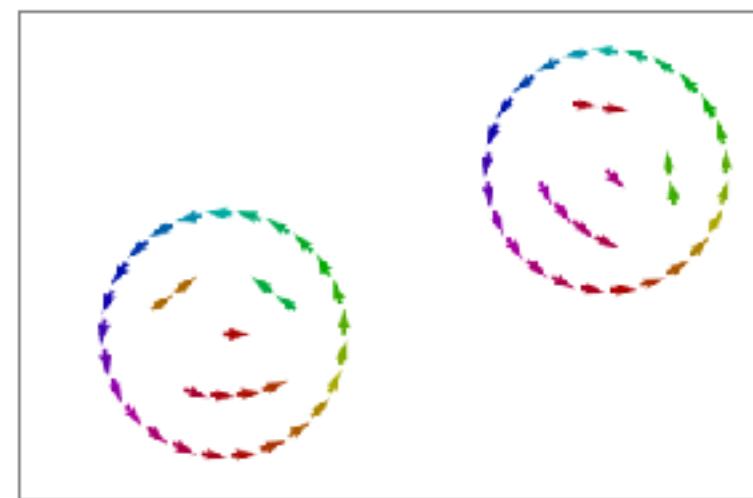
*G-feature maps are equivariant w.r.t. translation and scaling of the input*



*Activation for circles at each position and scale*

# Roto-translation group $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



*Using a set of transformed 2D conv kernels*

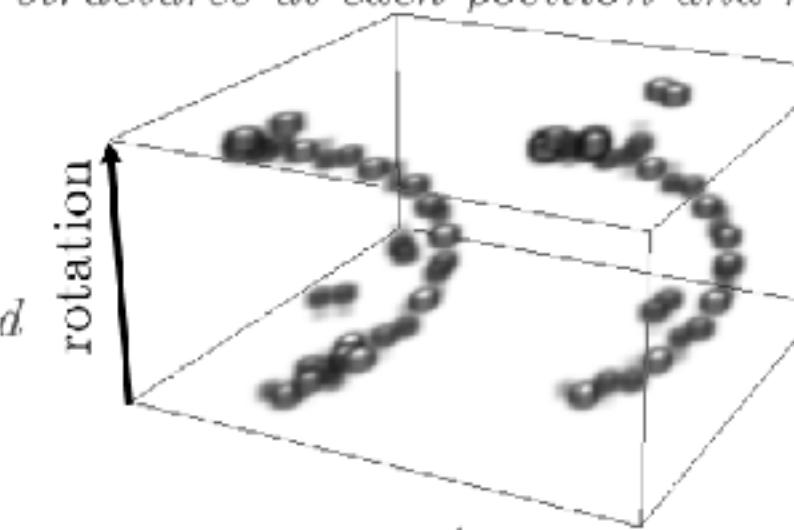
$$\theta = \frac{\pi}{2}$$

$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$



*G-feature maps are equivariant w.r.t. translation and rotation of the input*

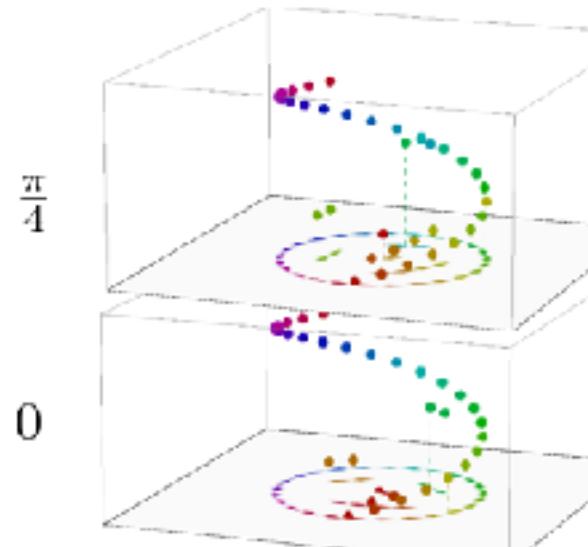


*Using a set of transformed G-conv kernels*

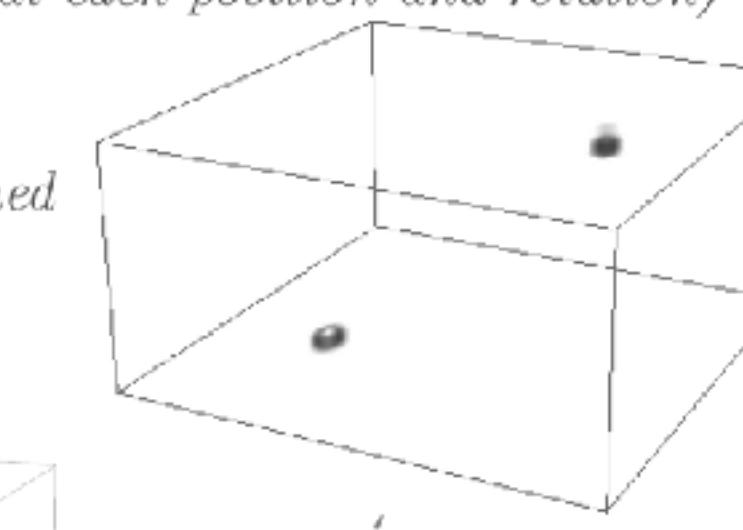
$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$

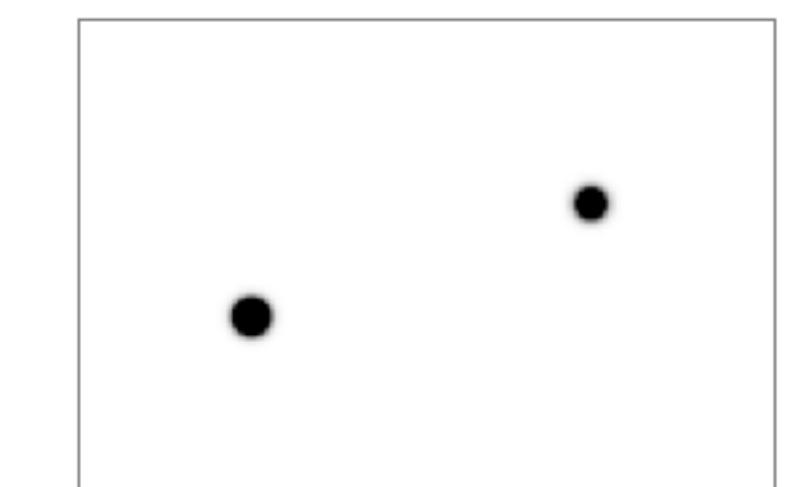
:



*G feature map (activation for faces at each position and rotation)*

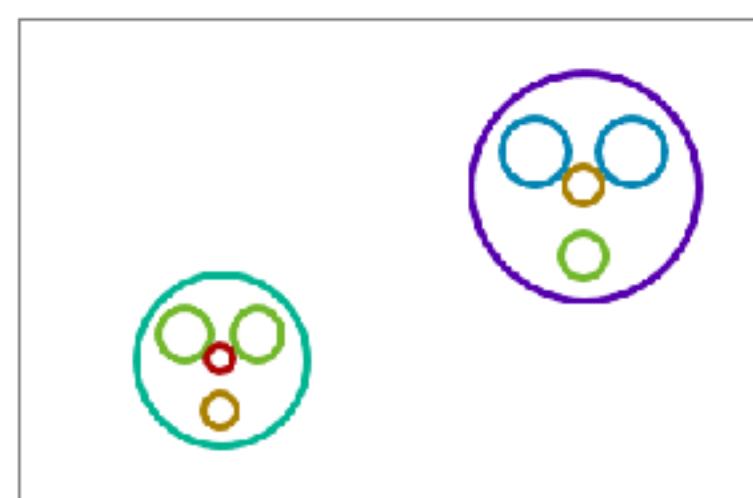


2D feature map



*Projection over sub-group  $H$  guarantees local invariance*

Scale-translation group  $\mathbb{R}^2 \rtimes \mathbb{R}^+$



$$s = 2$$



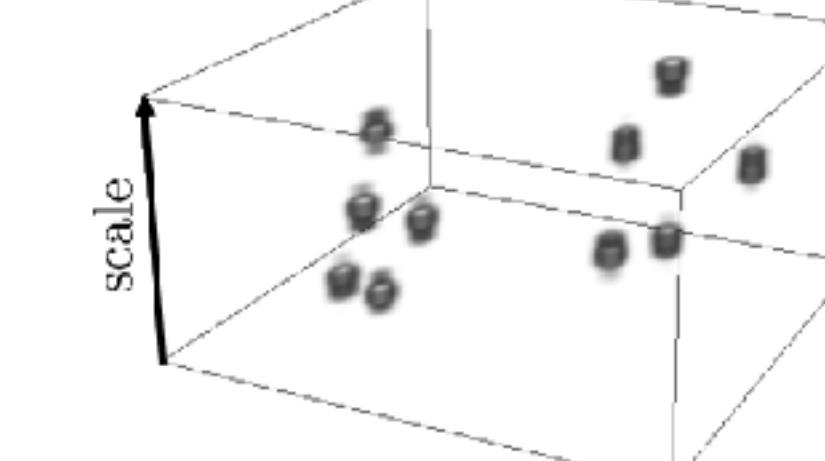
$$s = 1.4$$



$$s = 1$$

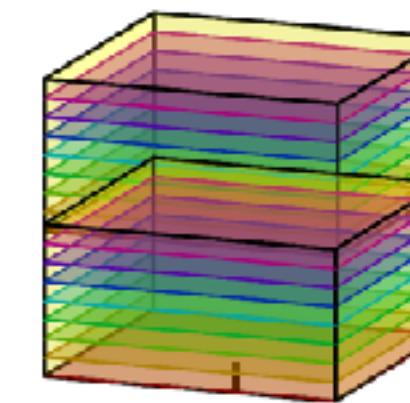


*G-feature maps are equivariant w.r.t. translation and scaling of the input*



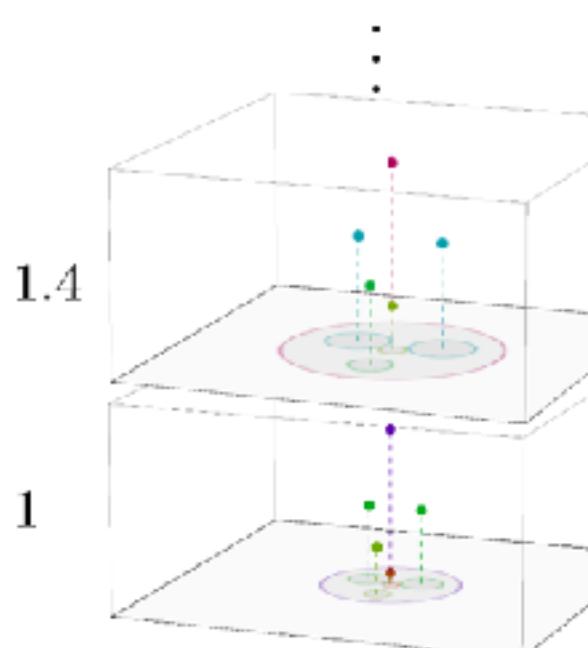
*Activation for circles at each position and scale*

Group conv layer



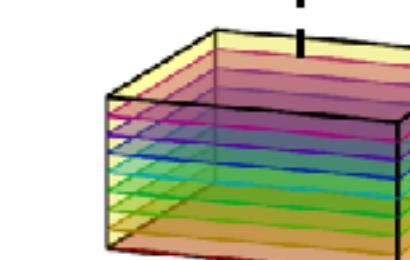
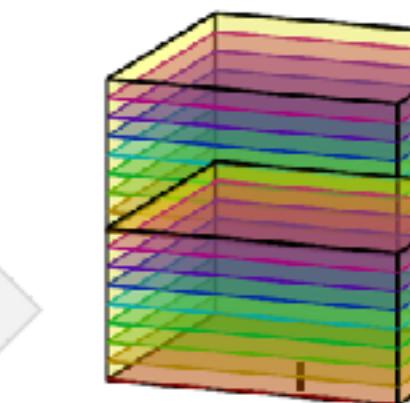
$$s = 1.4$$

$$s = 1$$



*G-conv kernels assign weights to activations in a pattern of relative poses*

Projection layer

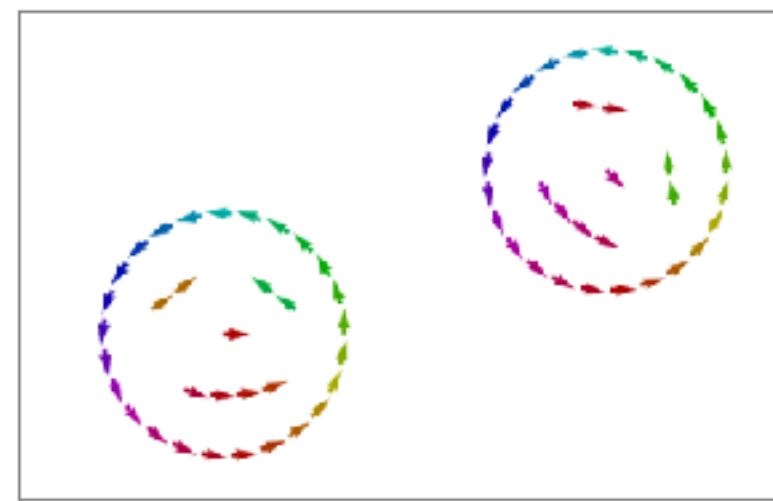


*Activation for faces at each position and scale*

Lifting layer

# Roto-translation group $SE(2) = \mathbb{R}^2 \rtimes SO(2)$

2D feature map



*Using a set of transformed 2D conv kernels*

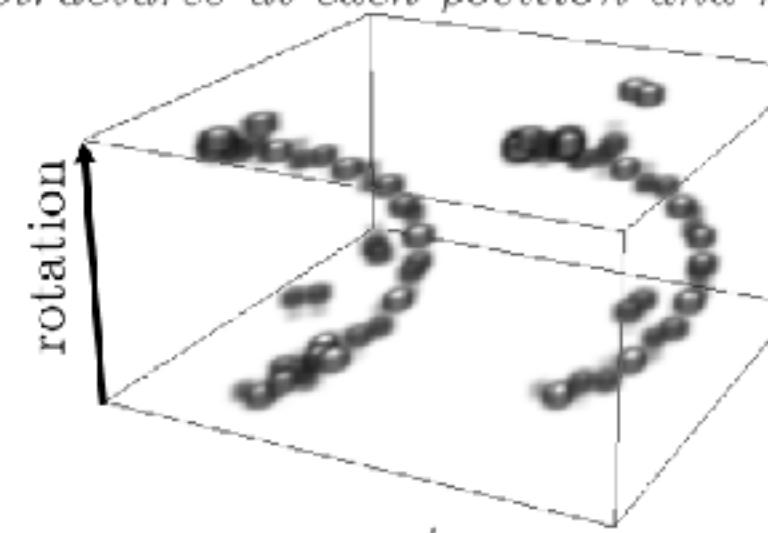
$$\theta = \frac{\pi}{2}$$

$$\theta = \frac{\pi}{4}$$

$$\theta = 0$$



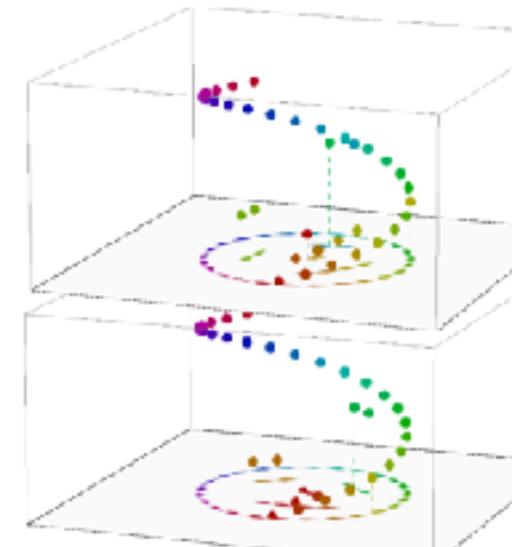
*G feature map (activation for oriented structures at each position and rotation)*



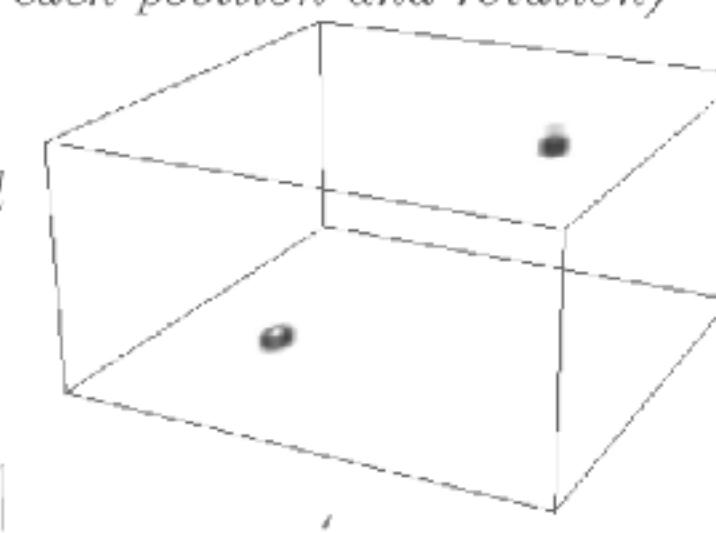
*Using a set of transformed G-conv kernels*

$$\theta = \frac{\pi}{4}$$

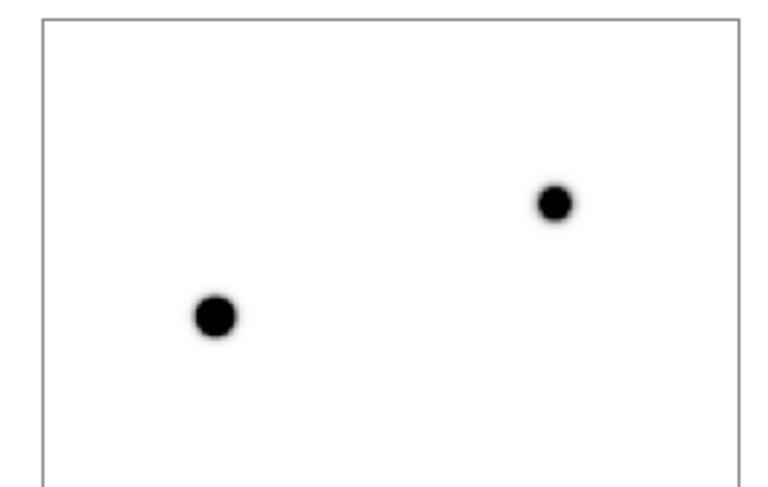
$$\theta = 0$$



*G feature map (activation for faces at each position and rotation)*

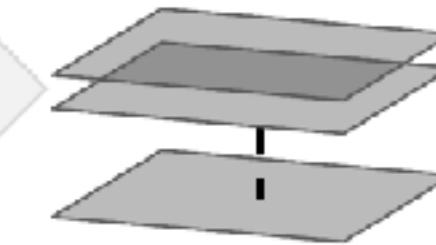


2D feature map

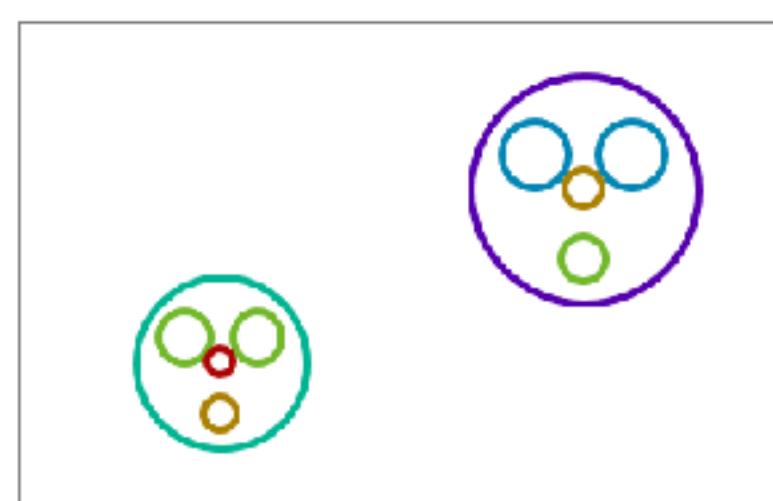


*Projection over sub-group  $H$  guarantees local invariance*

Projection layer



Scale-translation group  $\mathbb{R}^2 \times \mathbb{R}^+$



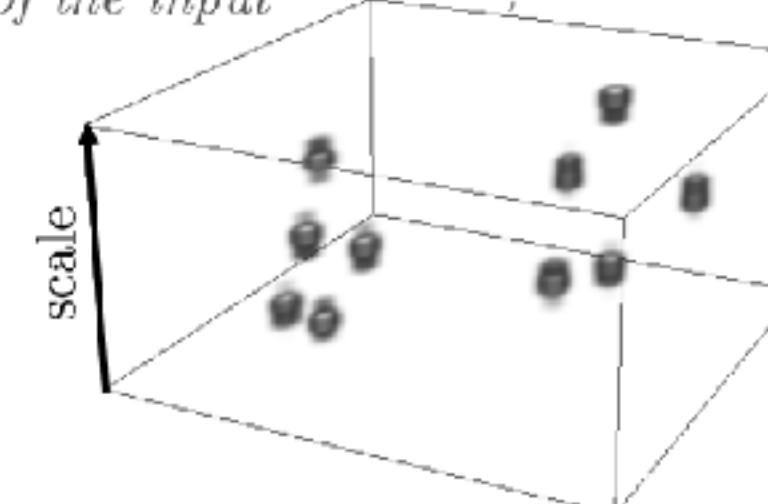
$$s = 2$$

$$s = 1.4$$

$$s = 1$$



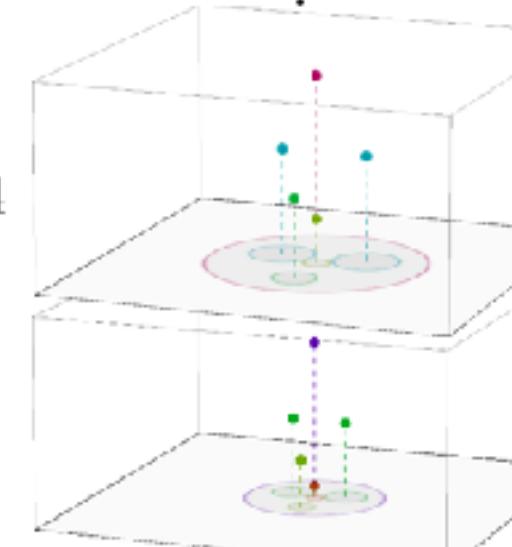
*G-feature maps are equivariant w.r.t. translation and scaling of the input*



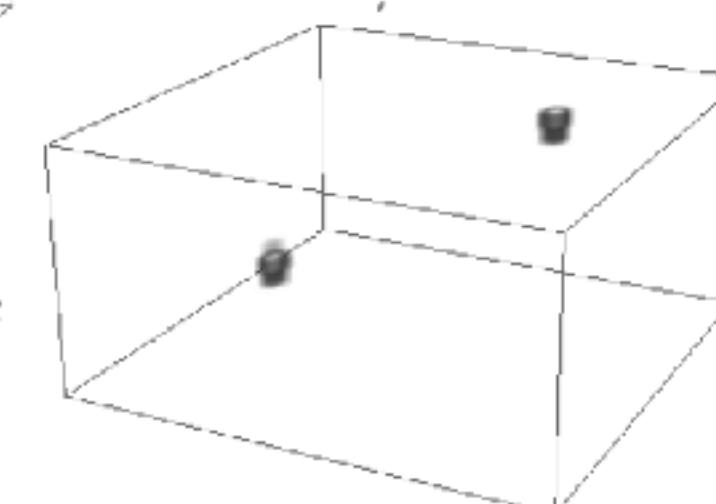
*Activation for circles at each position and scale*

$$s = 1.4$$

$$s = 1$$



*G-conv kernels assign weights to activations in a pattern of relative poses*



*Activation for faces at each position and scale*

Lifting layer

Group conv layer

Projection layer

**1. Motivation**

Equivariance → weight-sharing and generalization

**2. Pattern matching using group theory**

Group theory: symmetries & recognition by components  
(features have “poses”)

**3. Group convolutions**

Template matching over groups

**4. Example**

**5. G-convs are all you need!**

**6. Steerable group convolutions**

**7. Feature fields and escnn library**

**8. Equivariant tensor product layers**

**9. Equivariant graph NNs**

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

## 4. Example

5. G-convs are all you need!

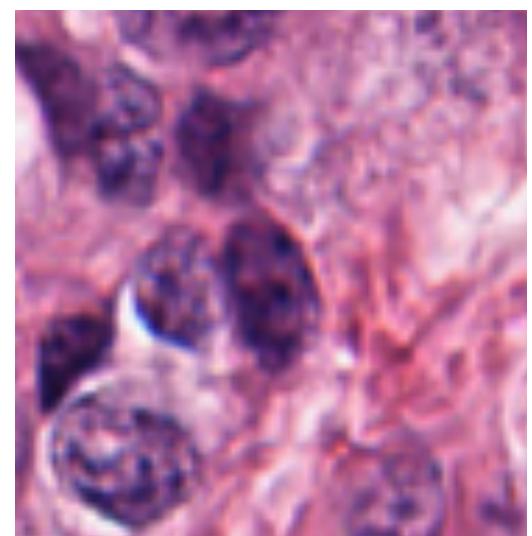
6. Steerable group convolutions

7. Feature fields and escnn library

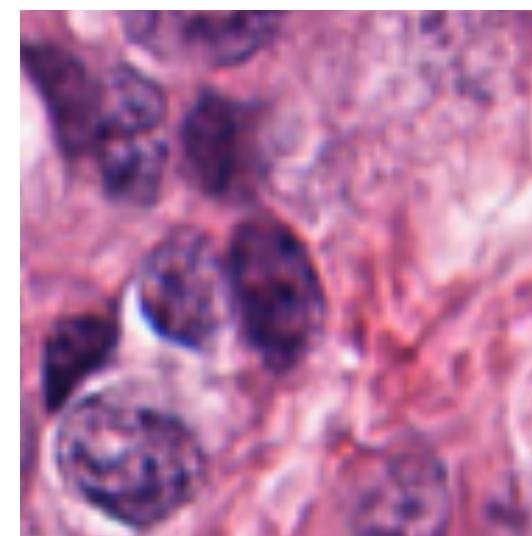
8. Equivariant tensor product layers

9. Equivariant graph NNs

# Architecture for rotation invariant mitotic cell detection

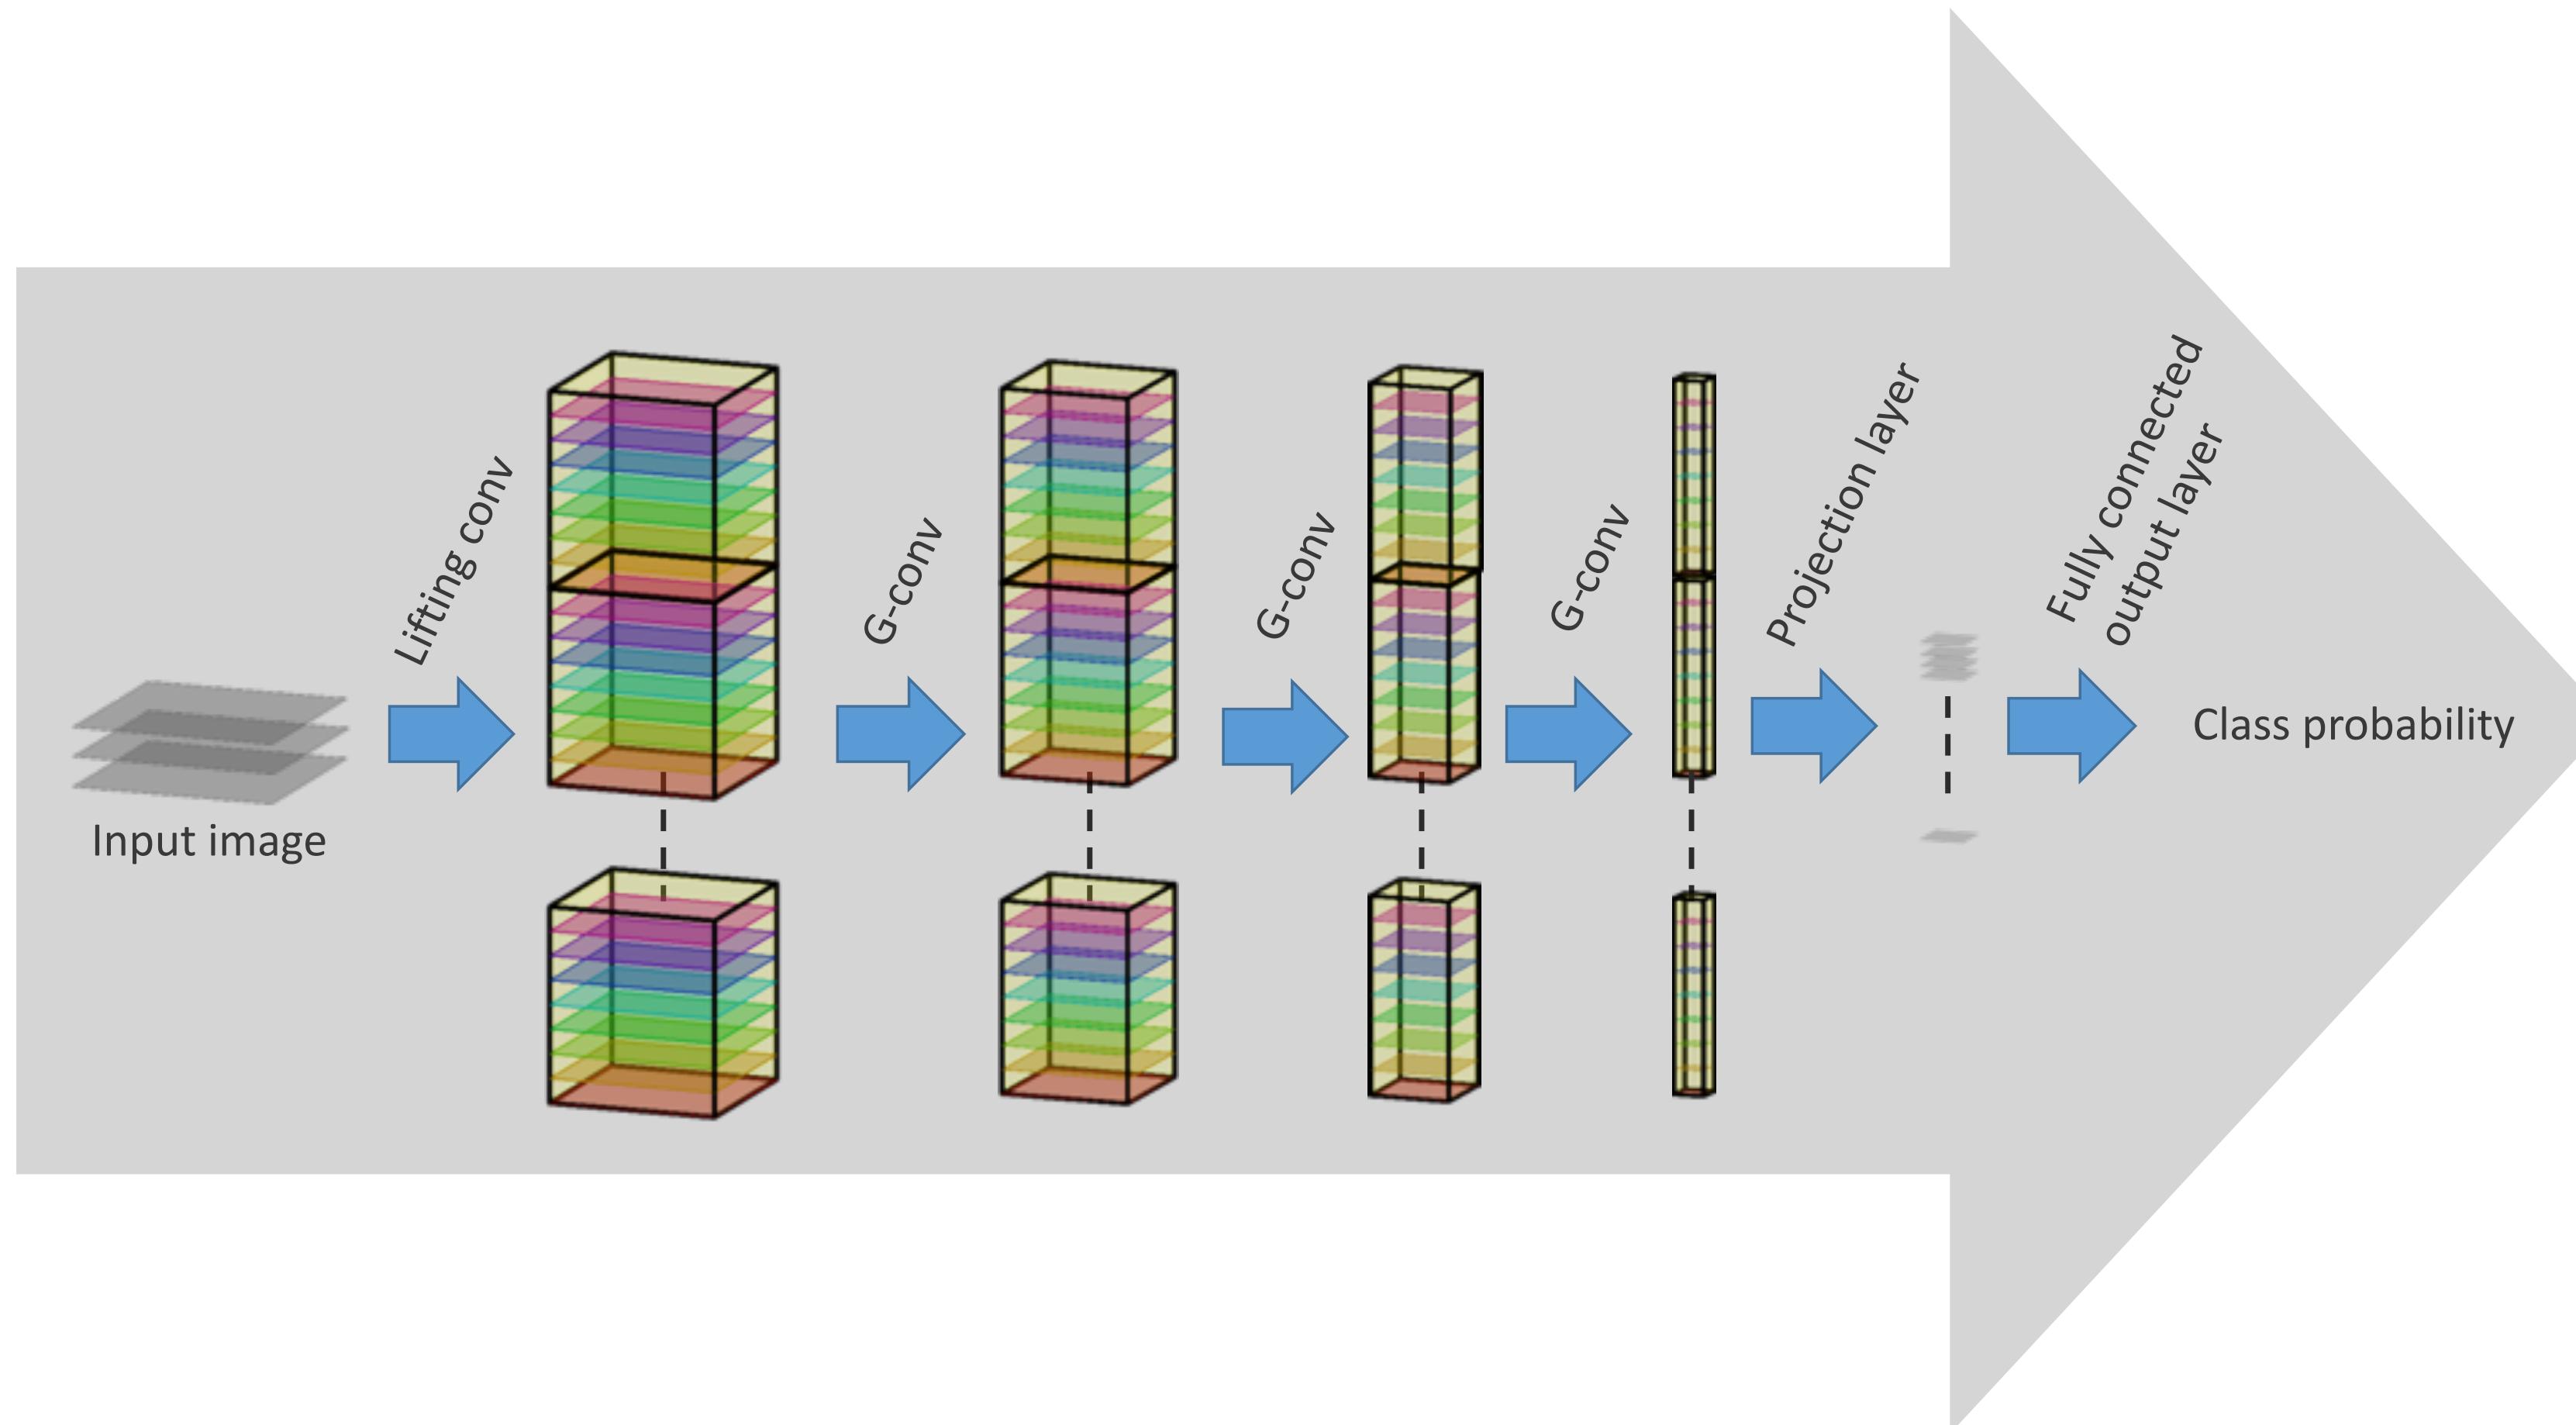
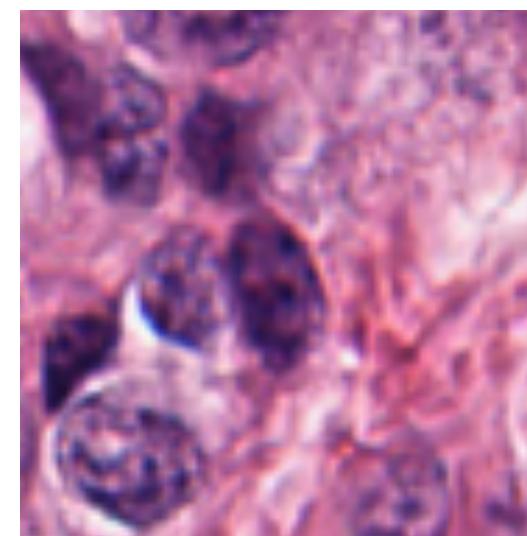


# Architecture for rotation invariant mitotic cell detection



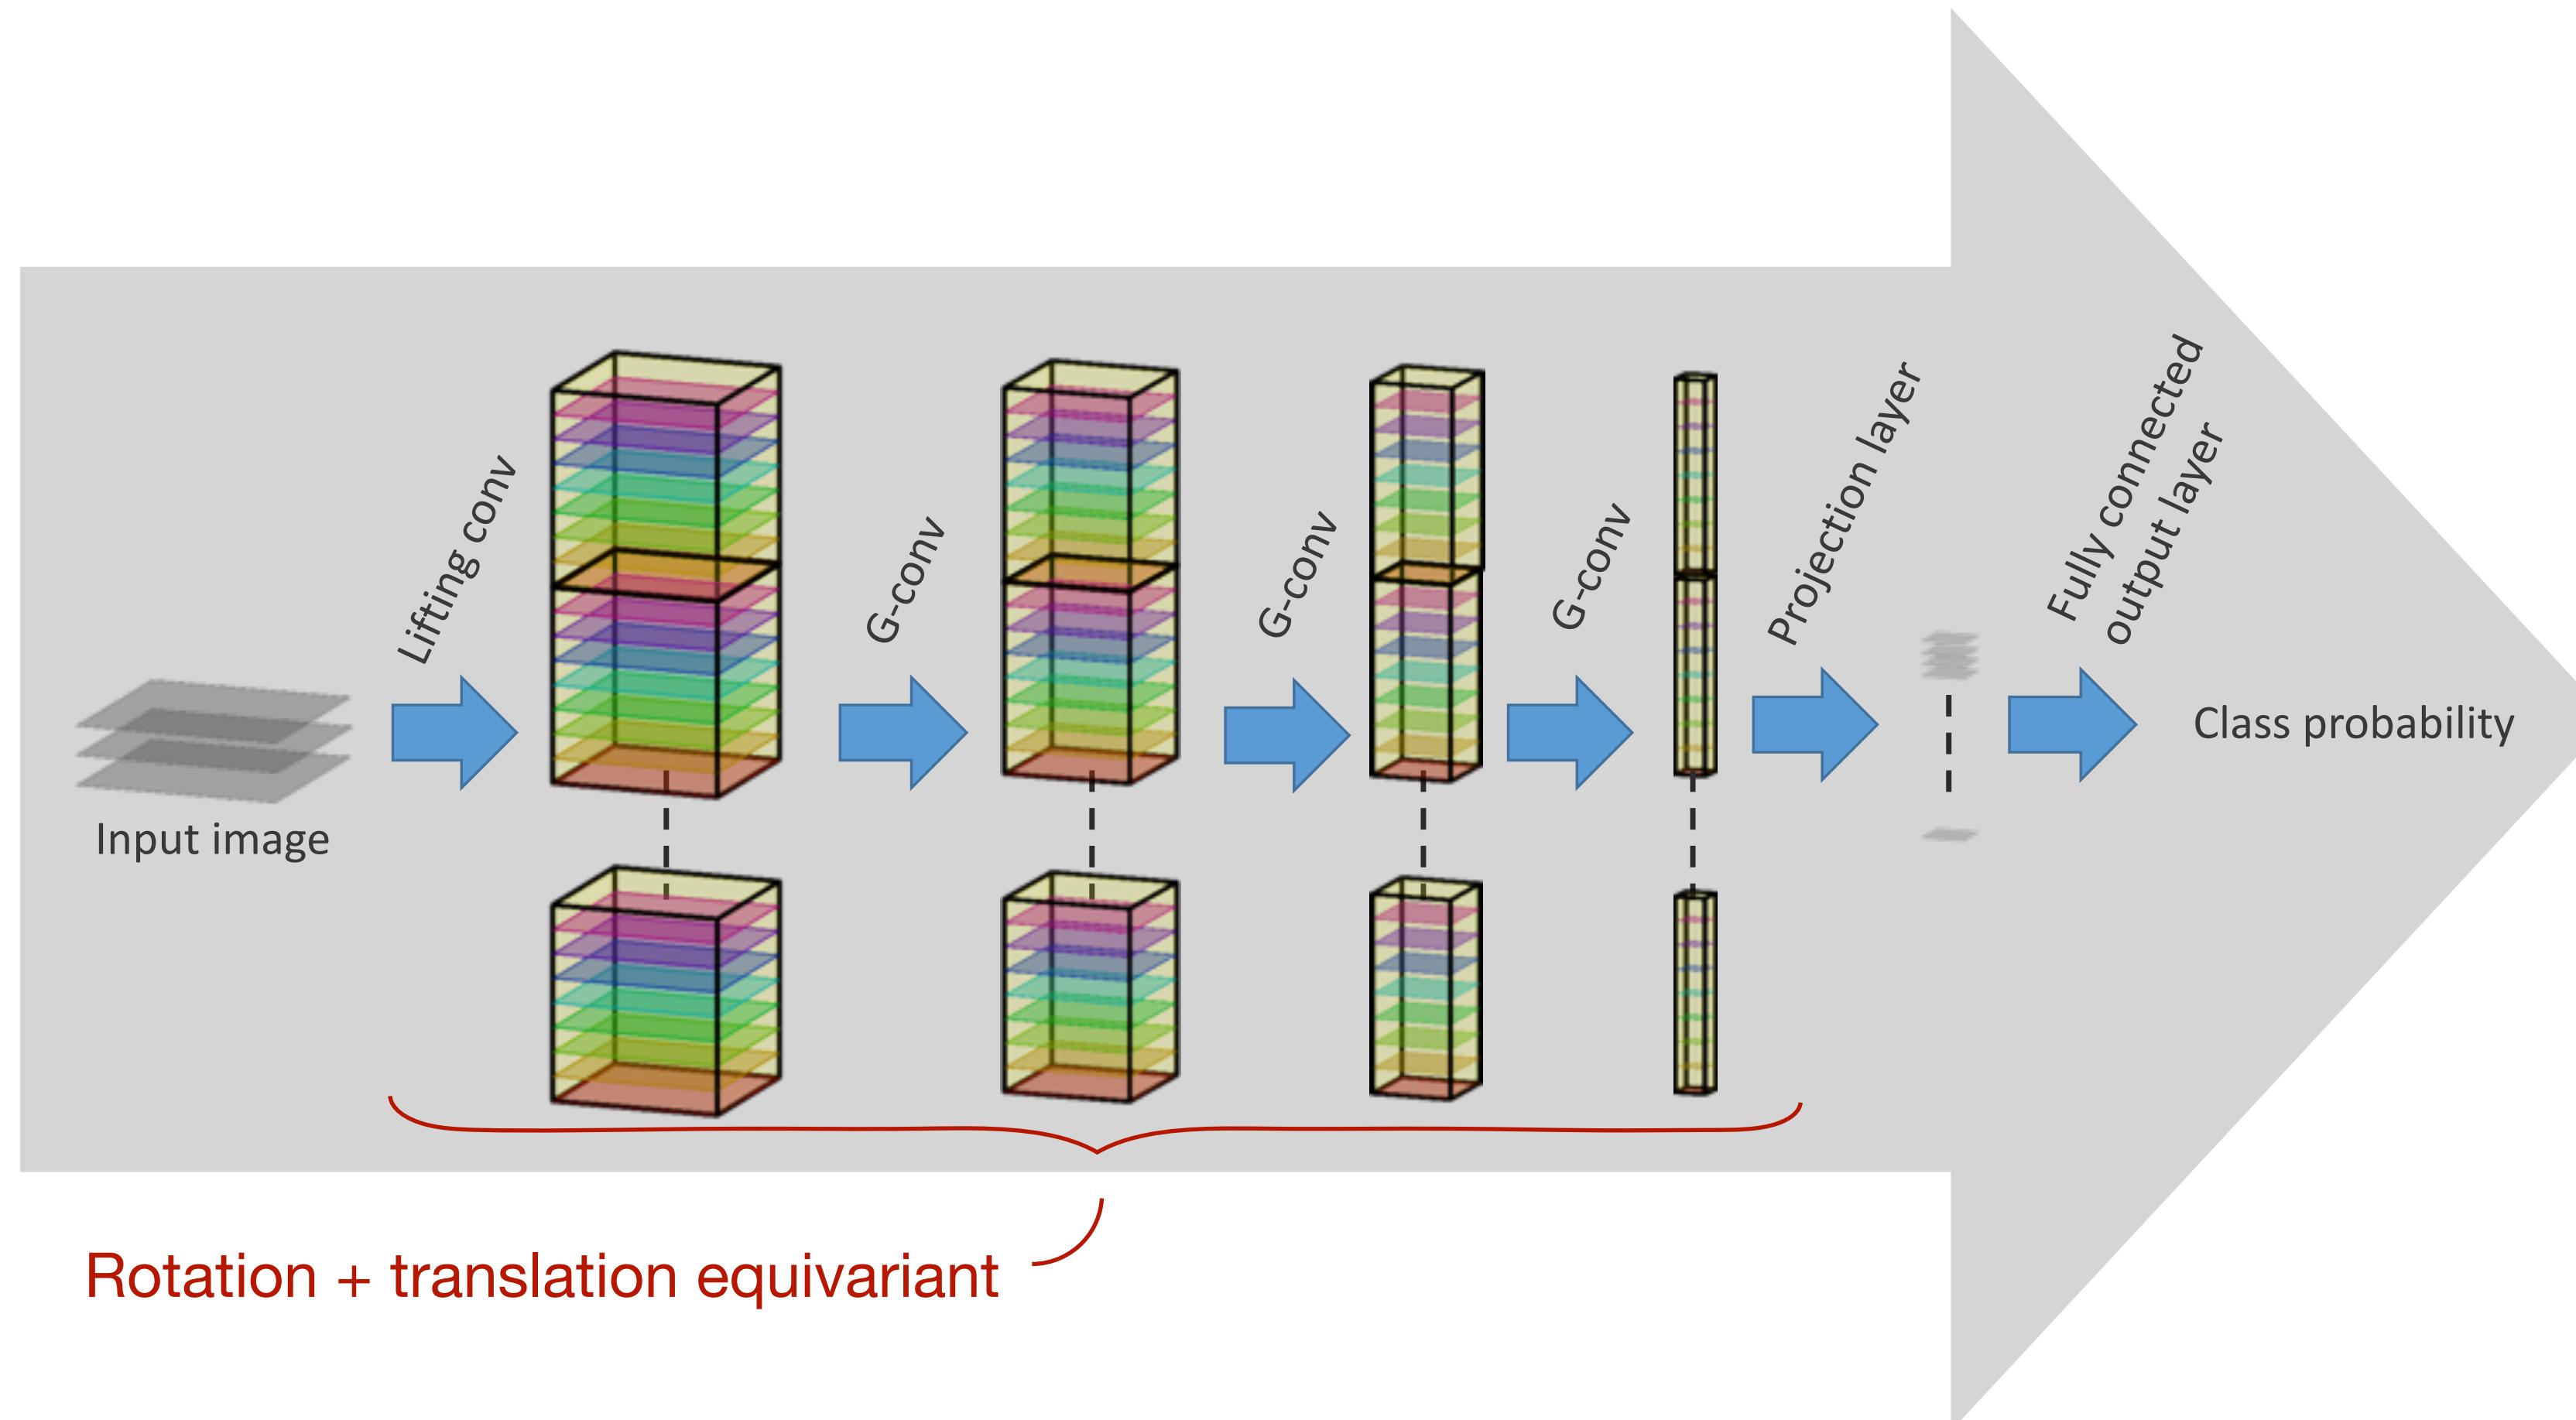
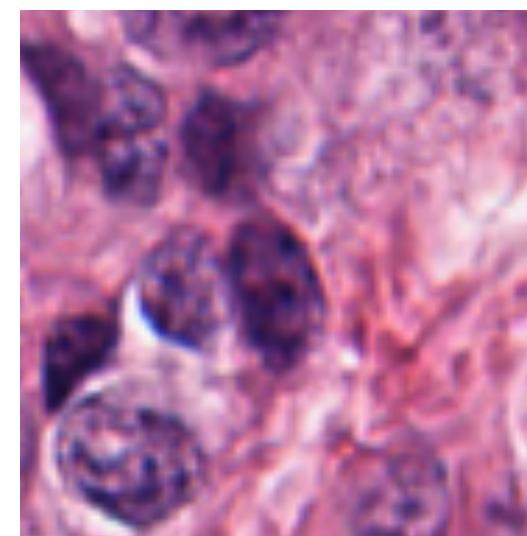
“normal” (0)  
vs  
“mitotic” (1)

# Architecture for rotation invariant mitotic cell detection



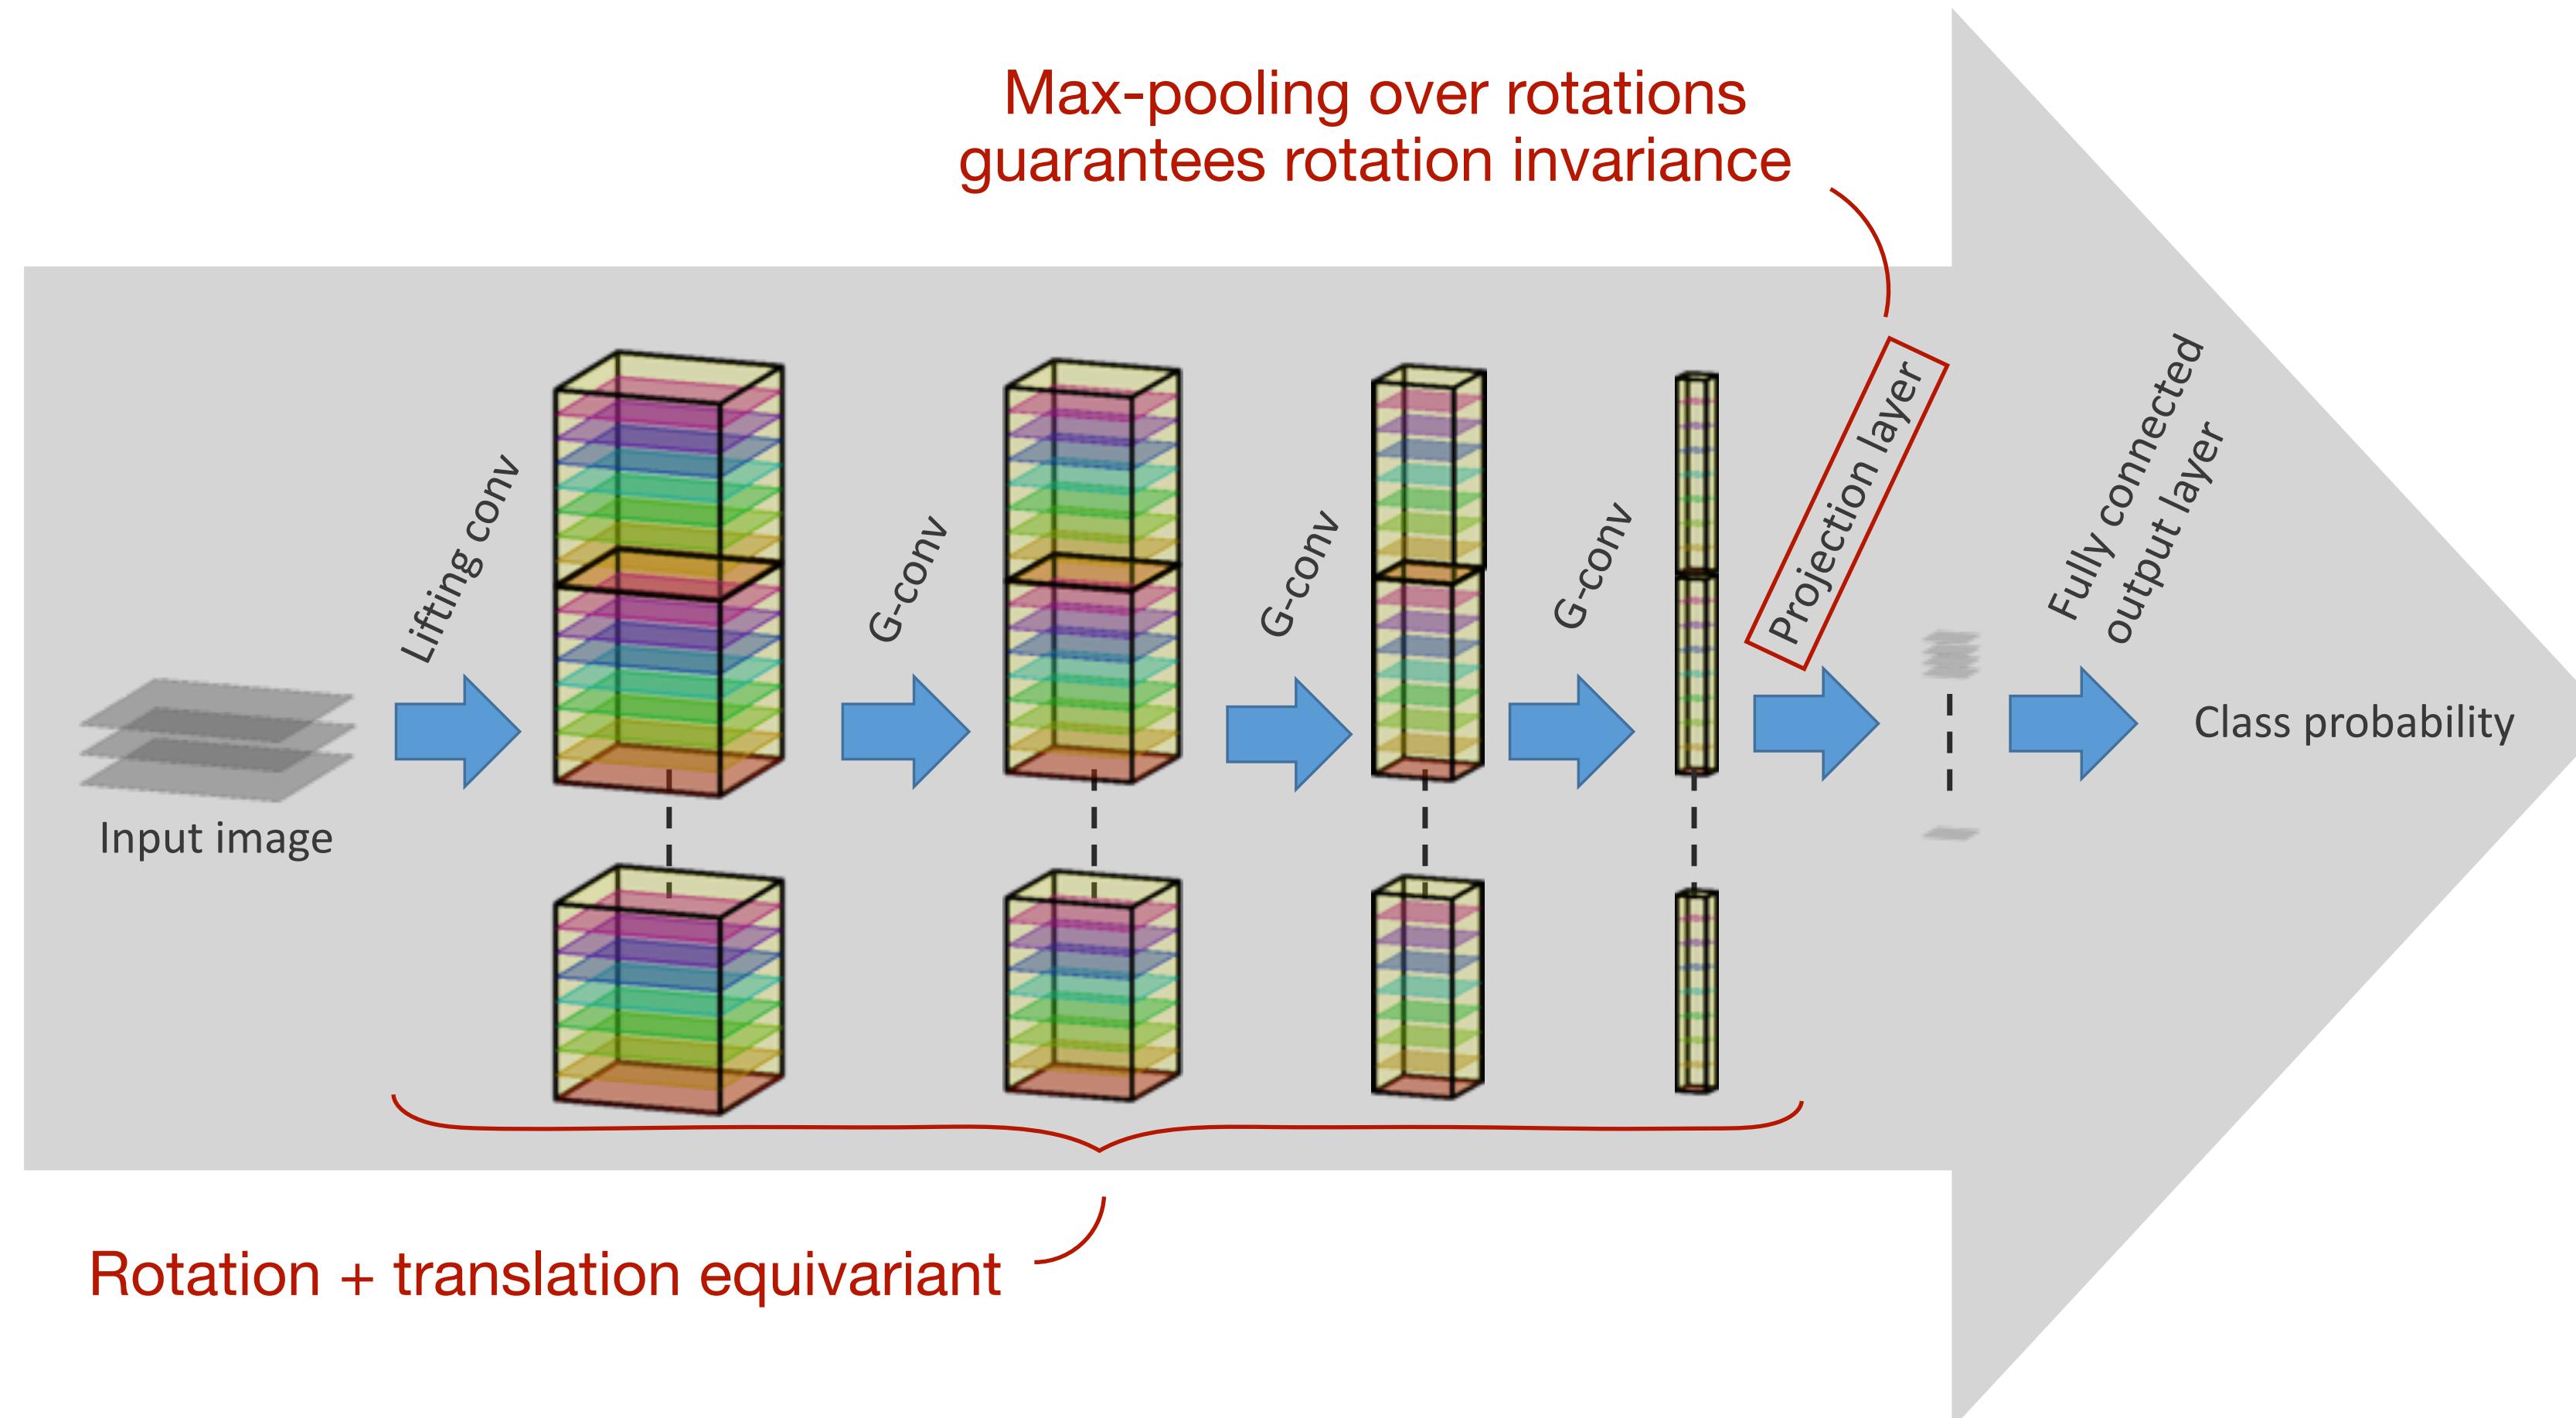
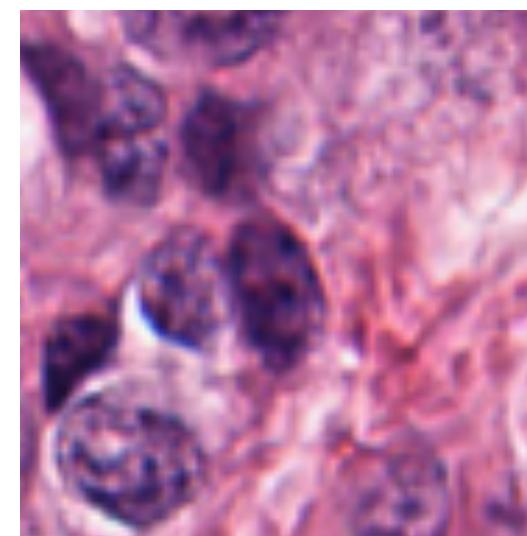
“normal” (0)  
vs  
“mitotic” (1)

# Architecture for rotation invariant mitotic cell detection



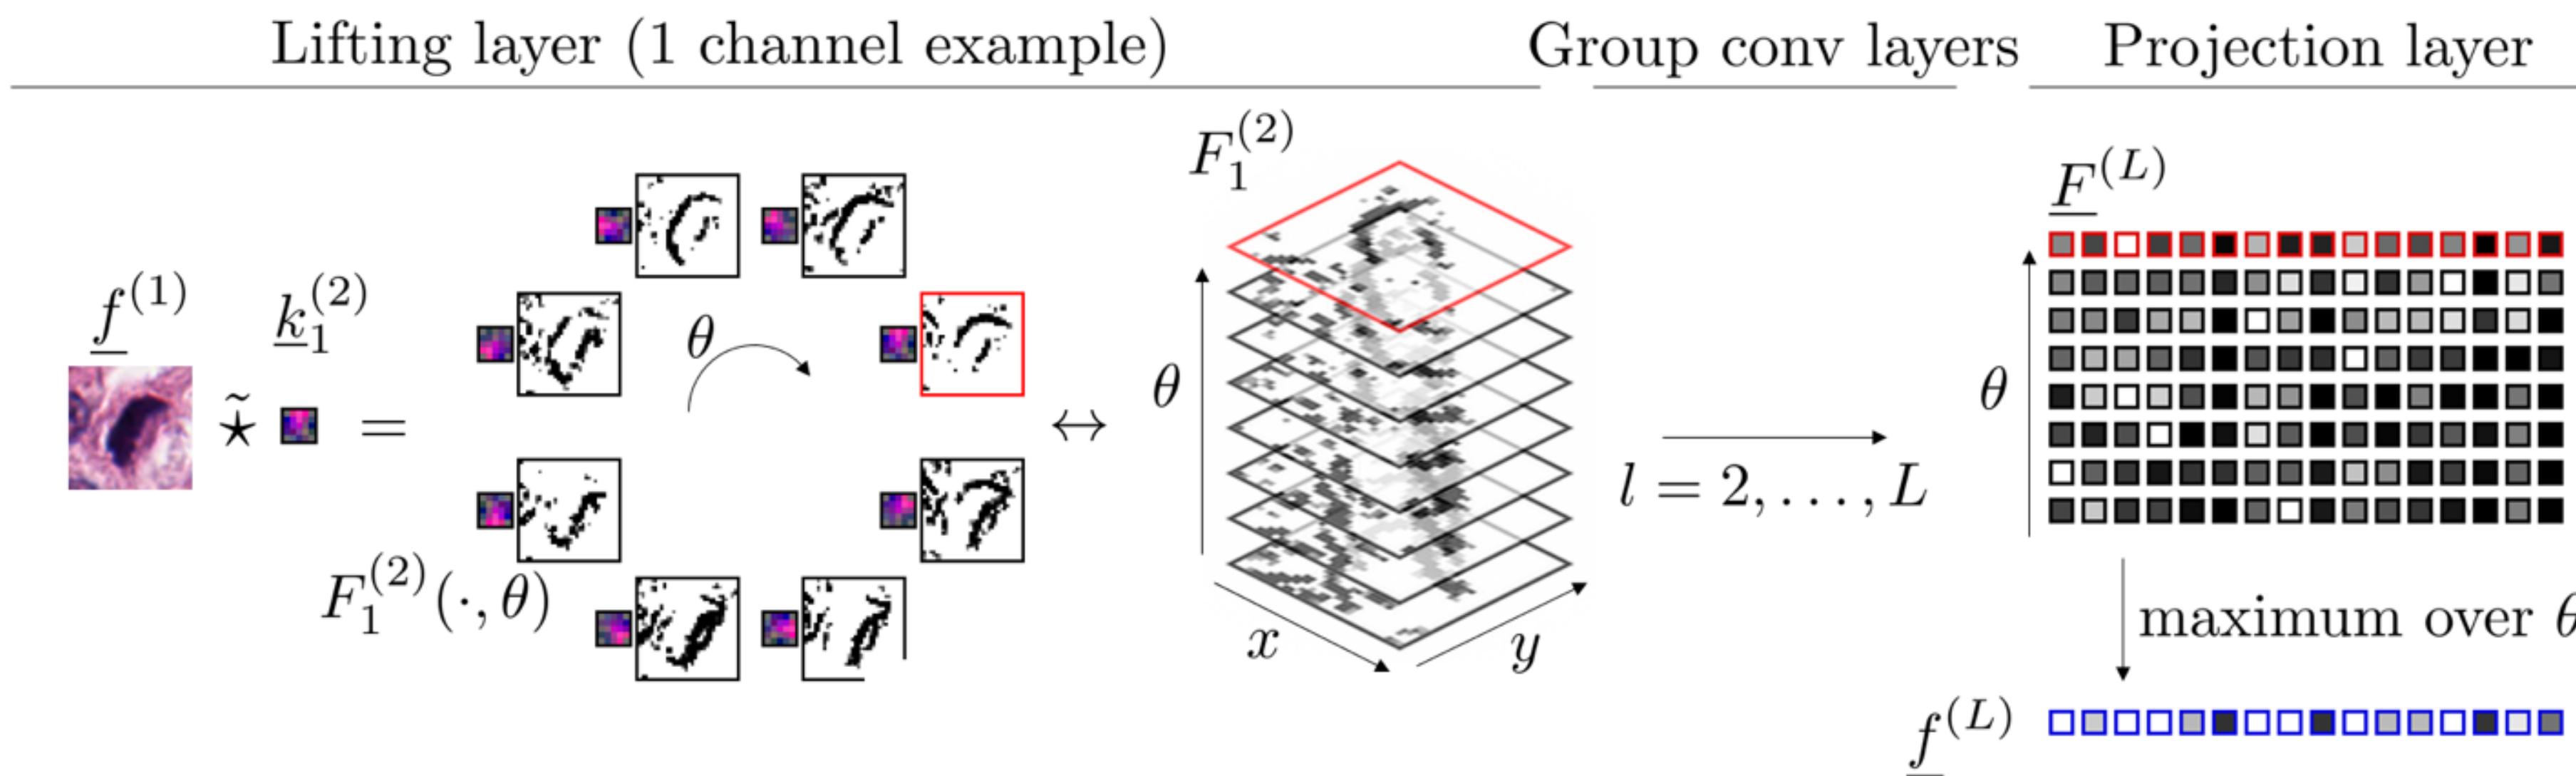
“normal” (0)  
vs  
“mitotic” (1)

# Architecture for rotation invariant mitotic cell detection

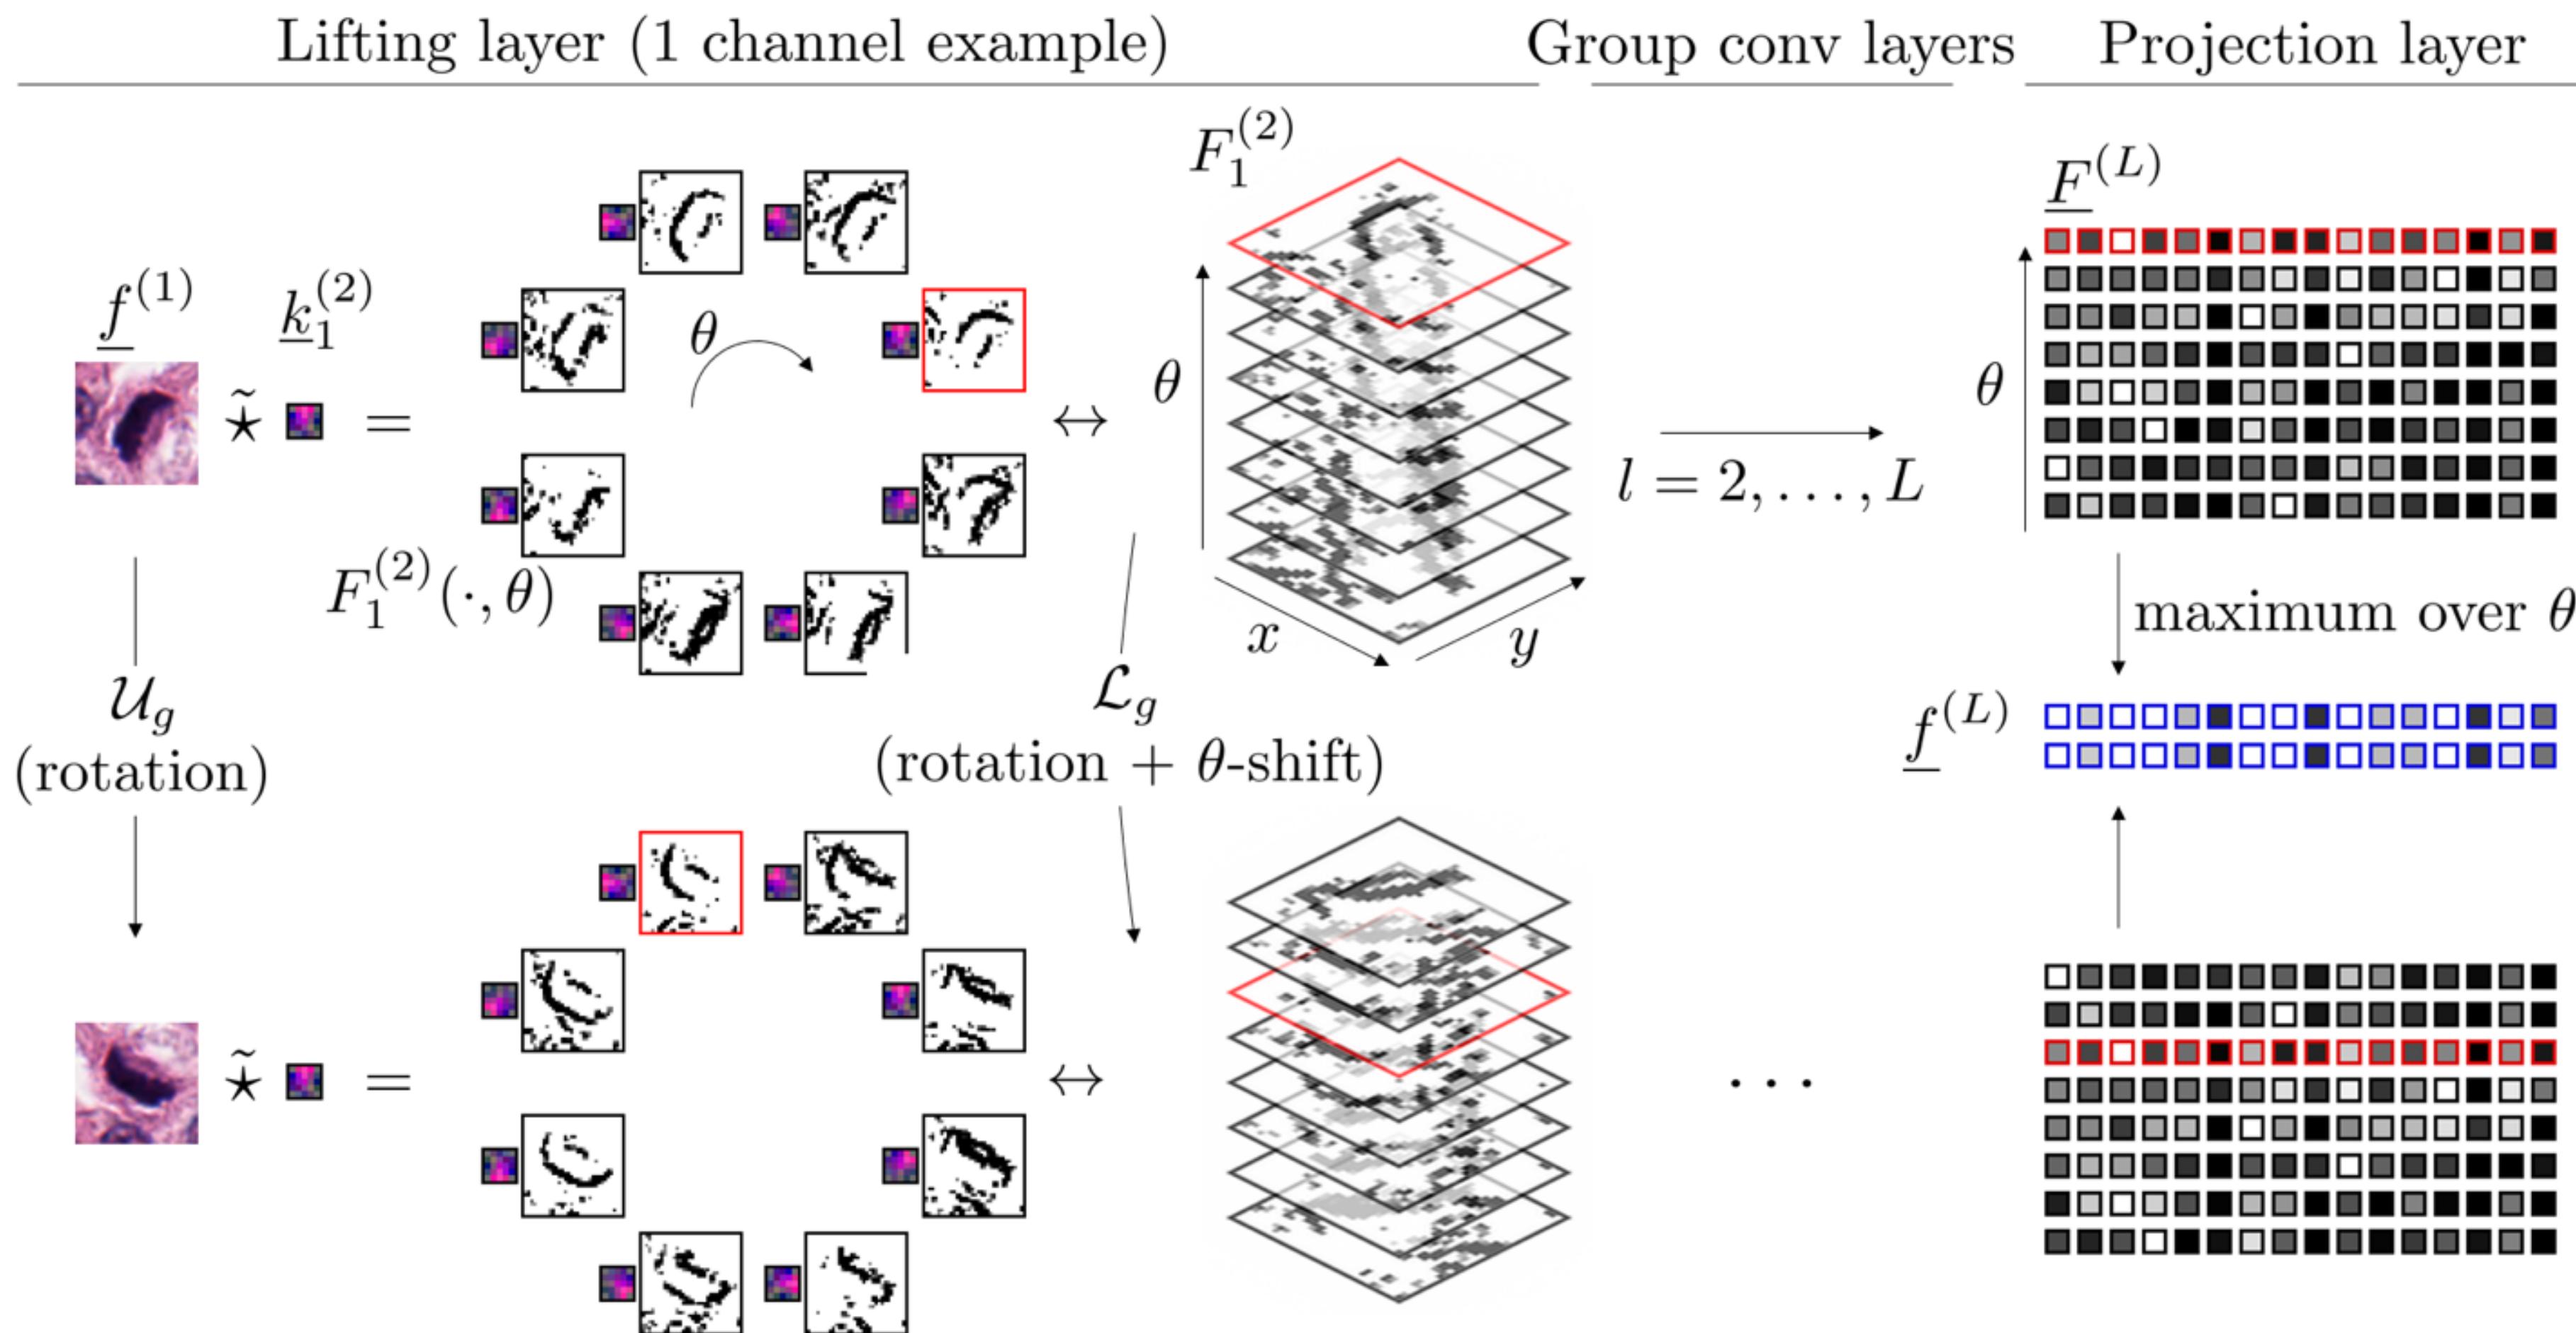


“normal” (0)  
vs  
“mitotic” (1)

# Architecture for rotation invariant mitotic cell detection

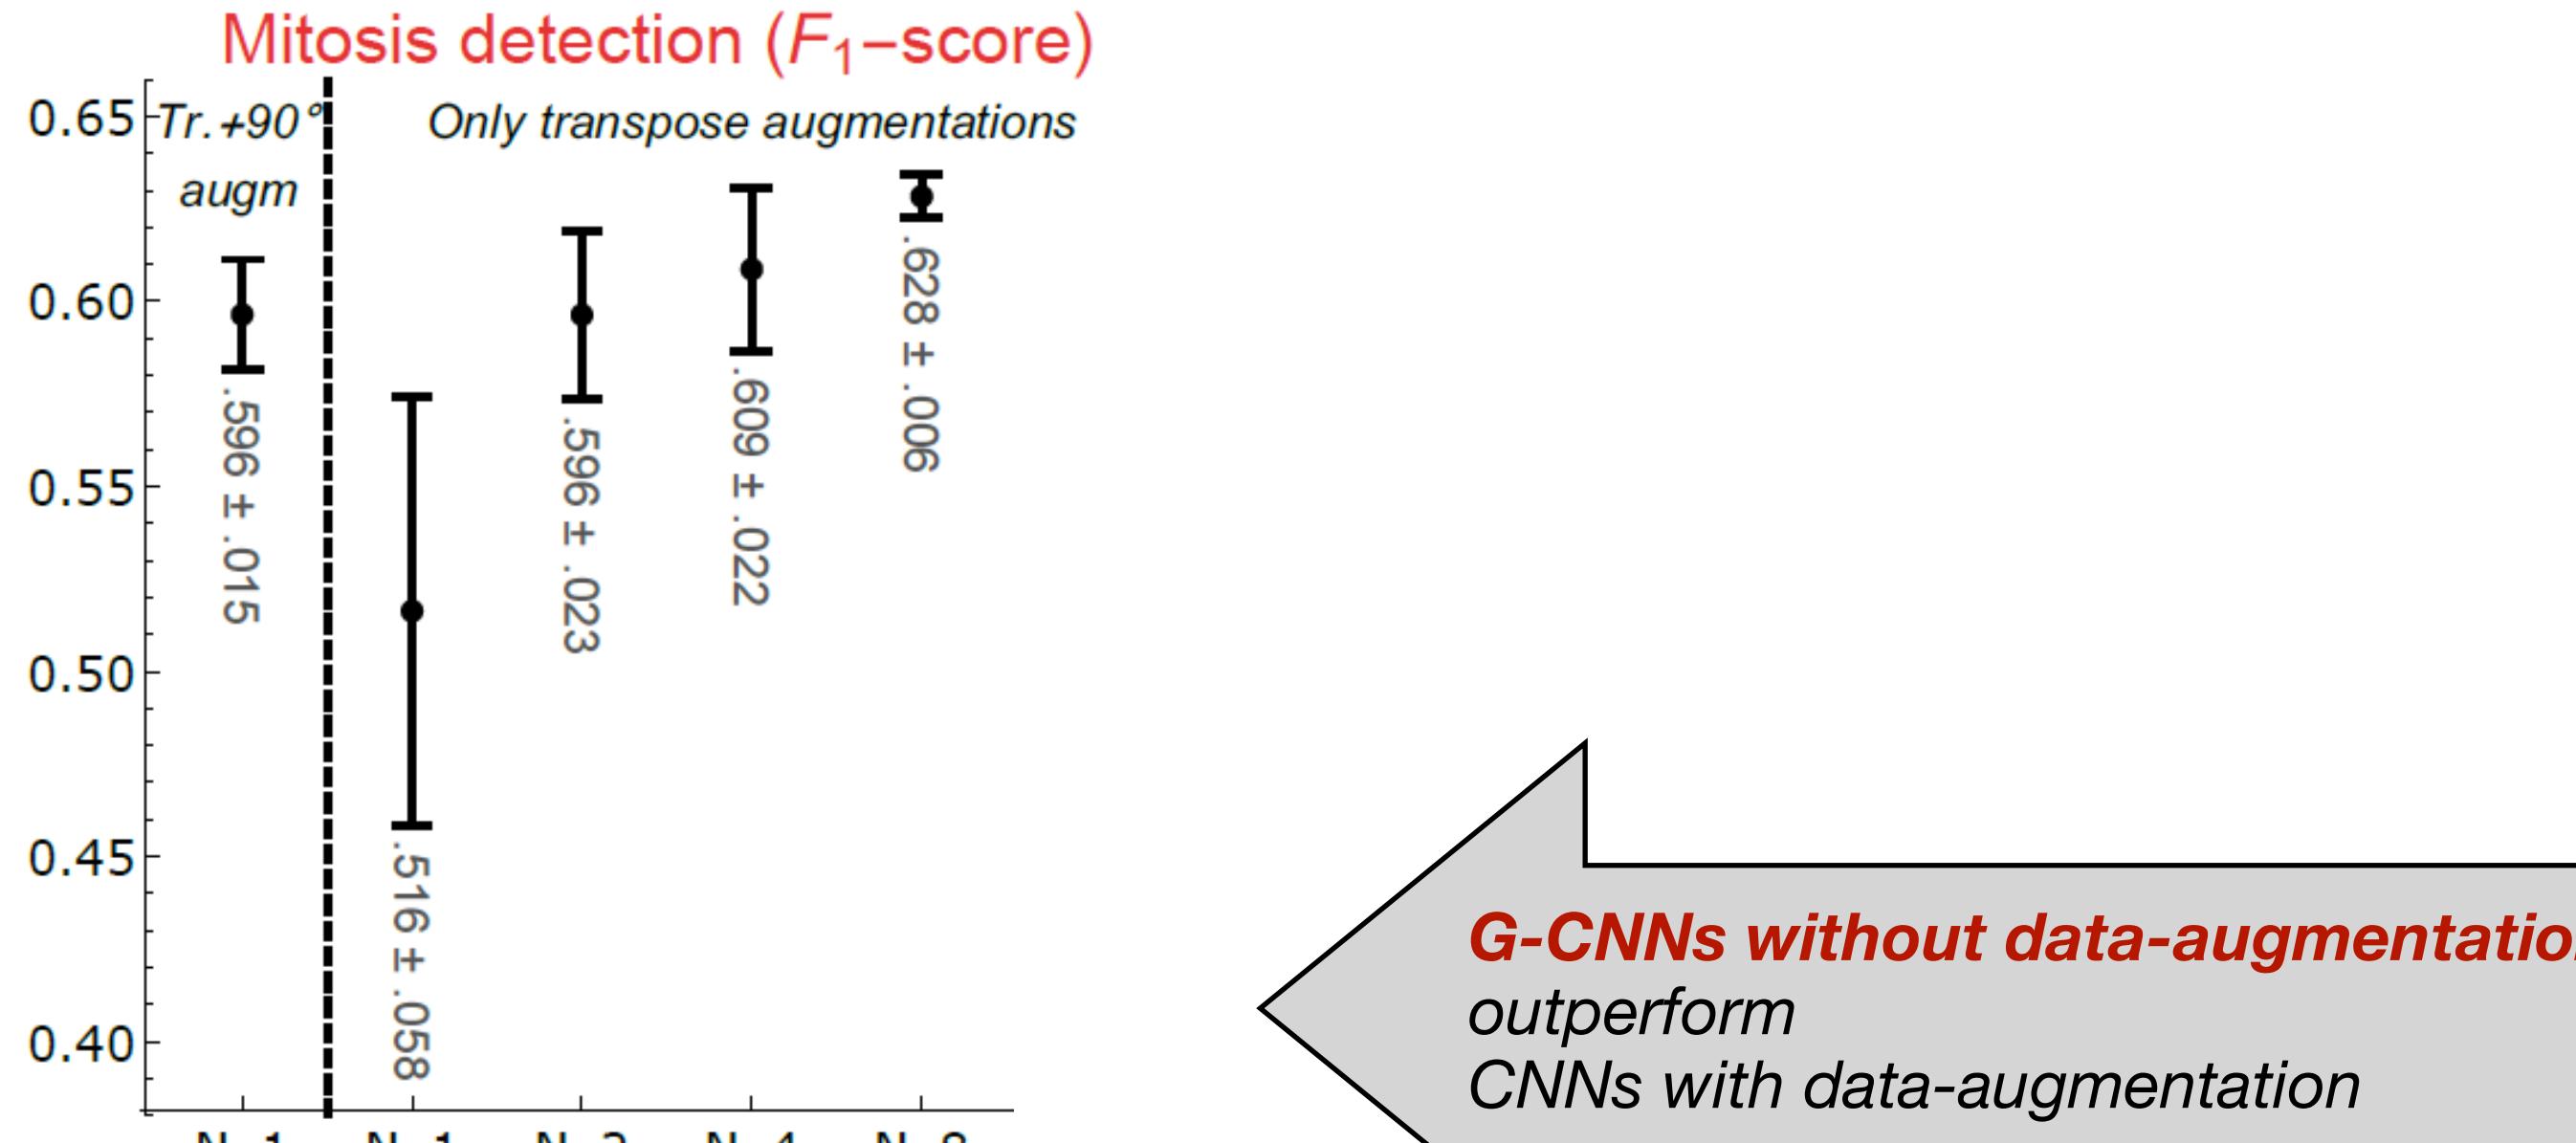
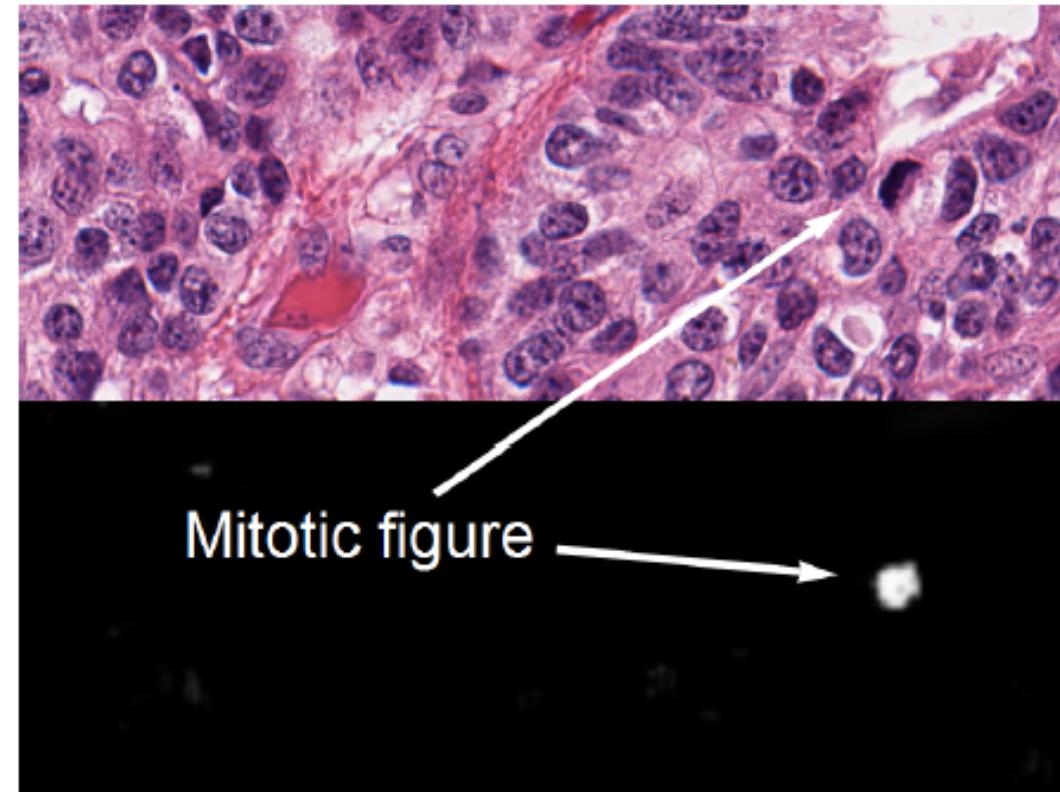


# Architecture for rotation invariant mitotic cell detection



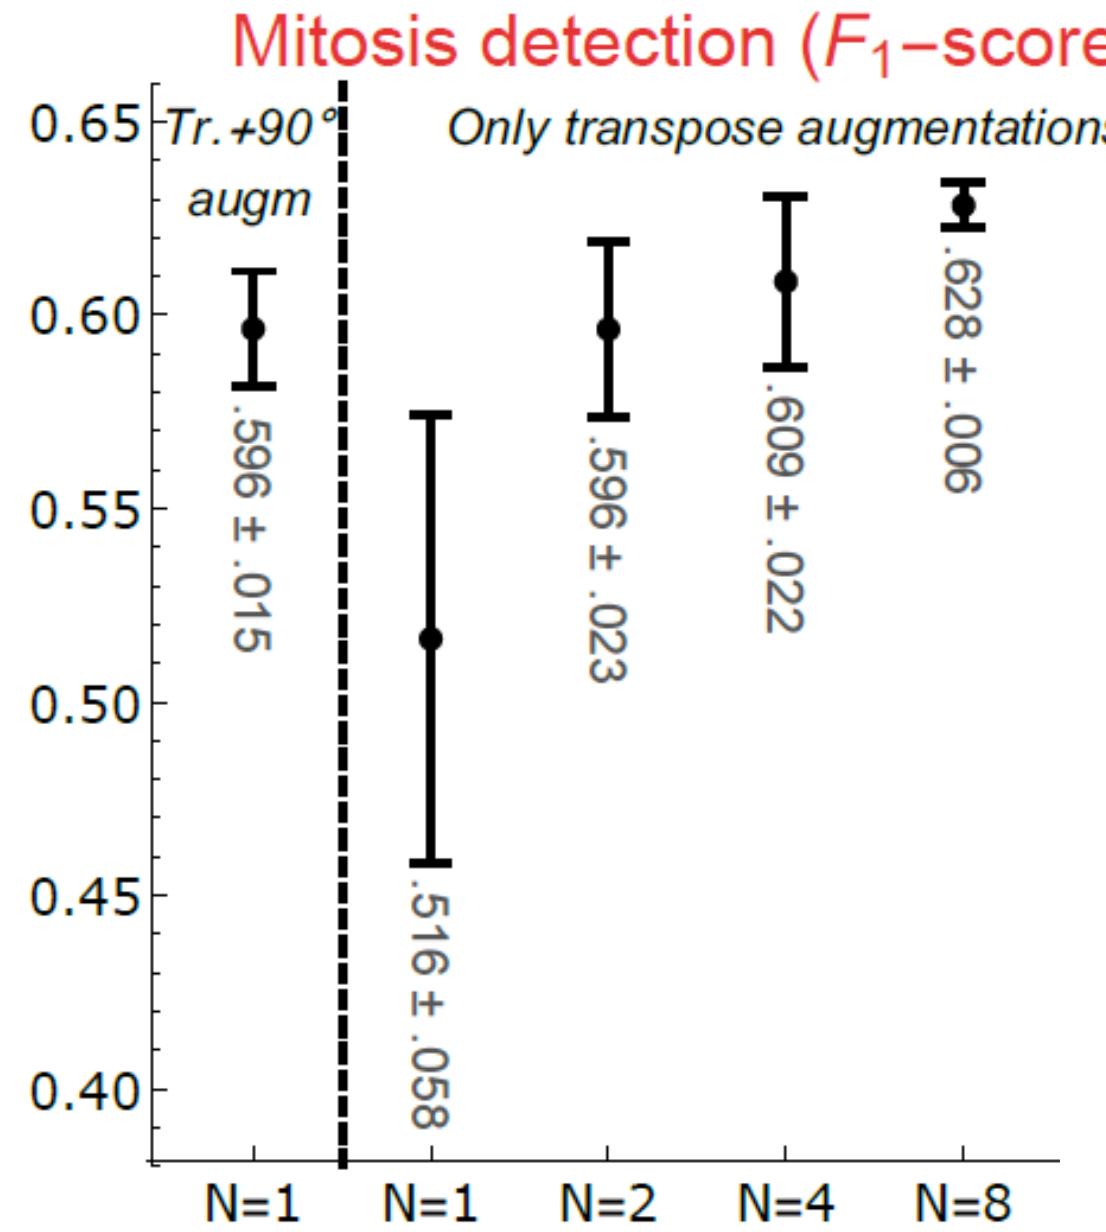
# Architecture for rotation invariant mitotic cell detection

Bekkers & Lafarge et al. MICCAI 2018

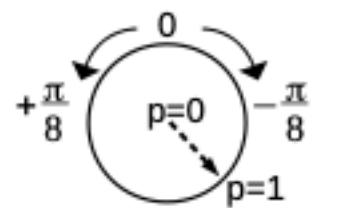
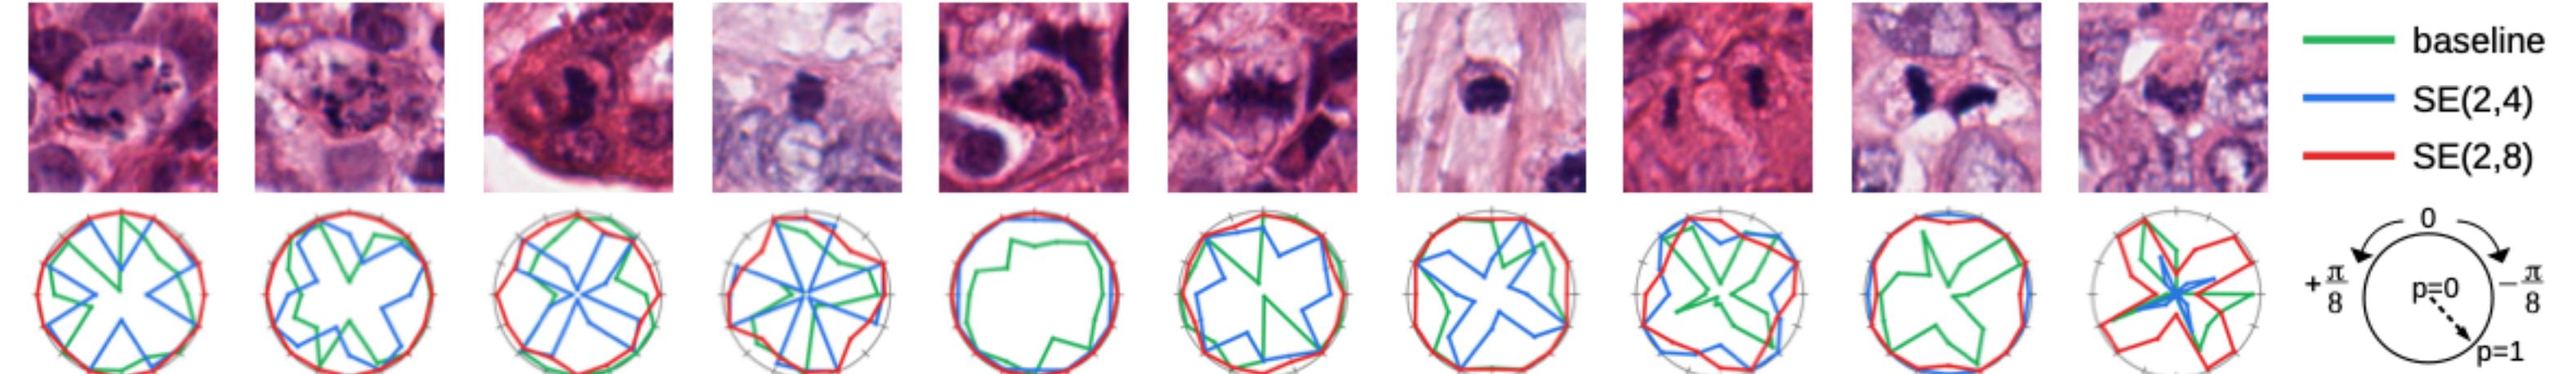


# Architecture for rotation invariant mitotic cell detection

Bekkers & Lafarge et al. MICCAI 2018



Lafarge et al. MedIA 2020

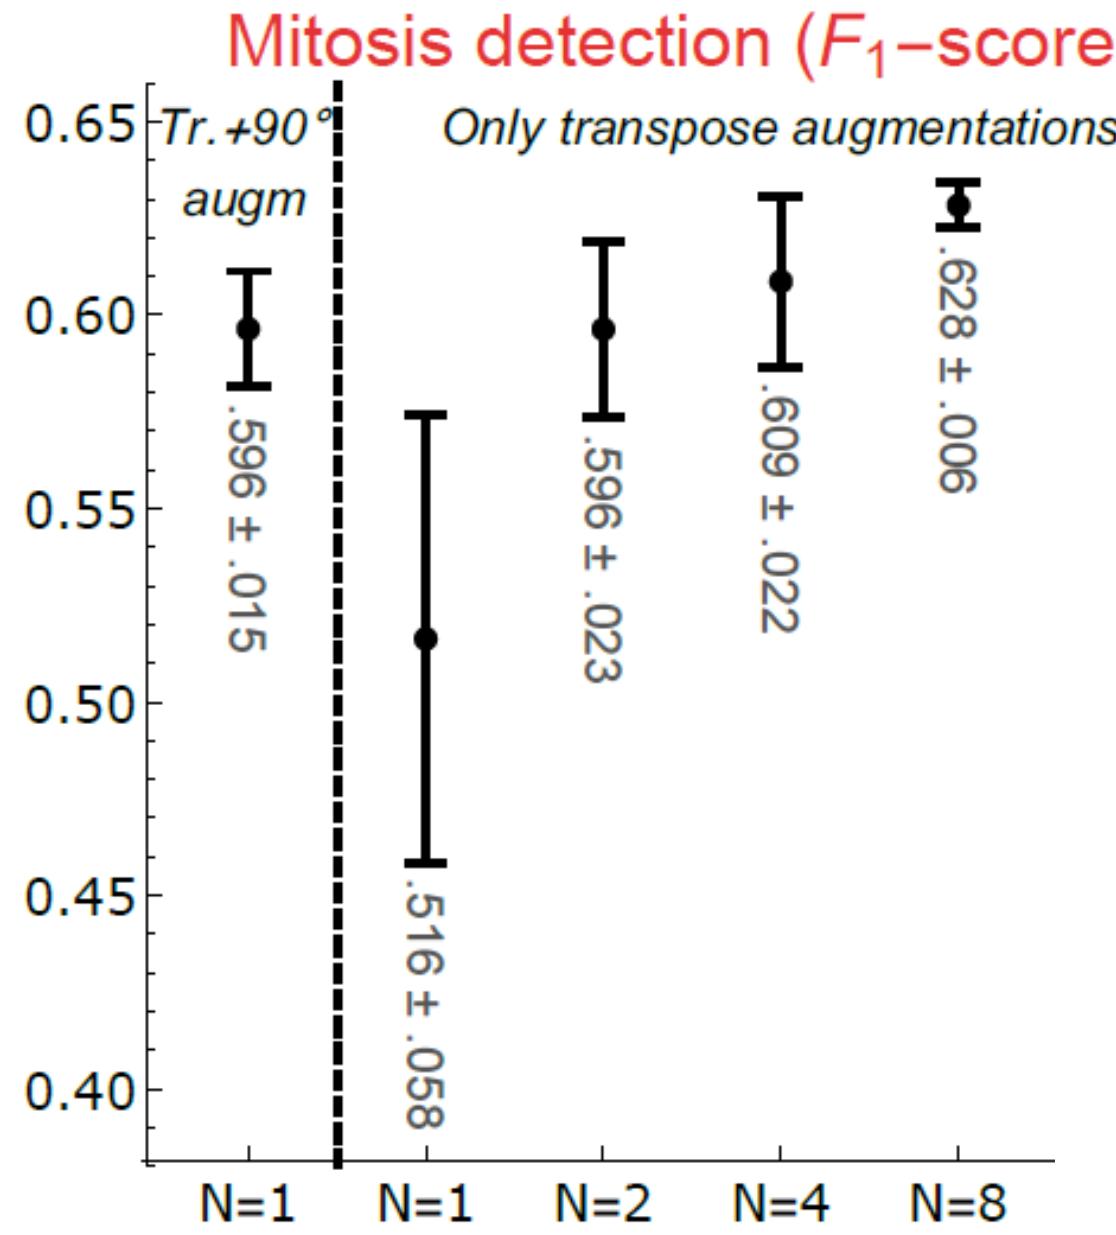
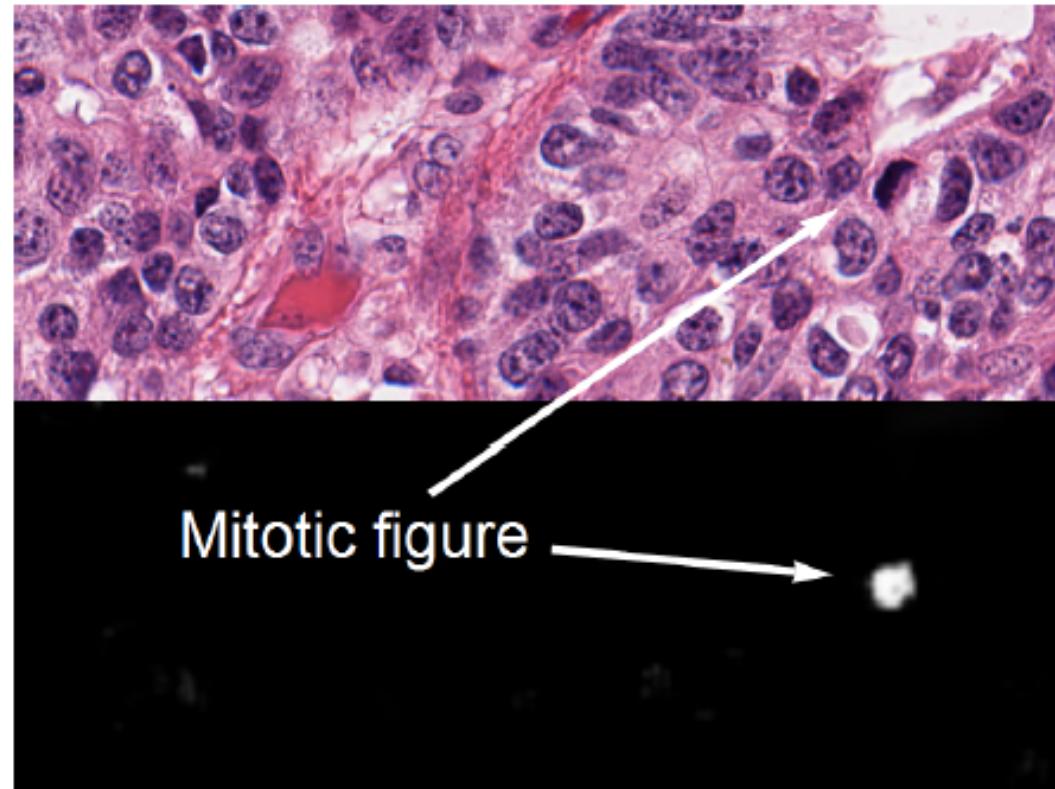


G-CNNs guarantee geometric stability.  
They are **robust to input distortions**,  
regular CNNs aren't...

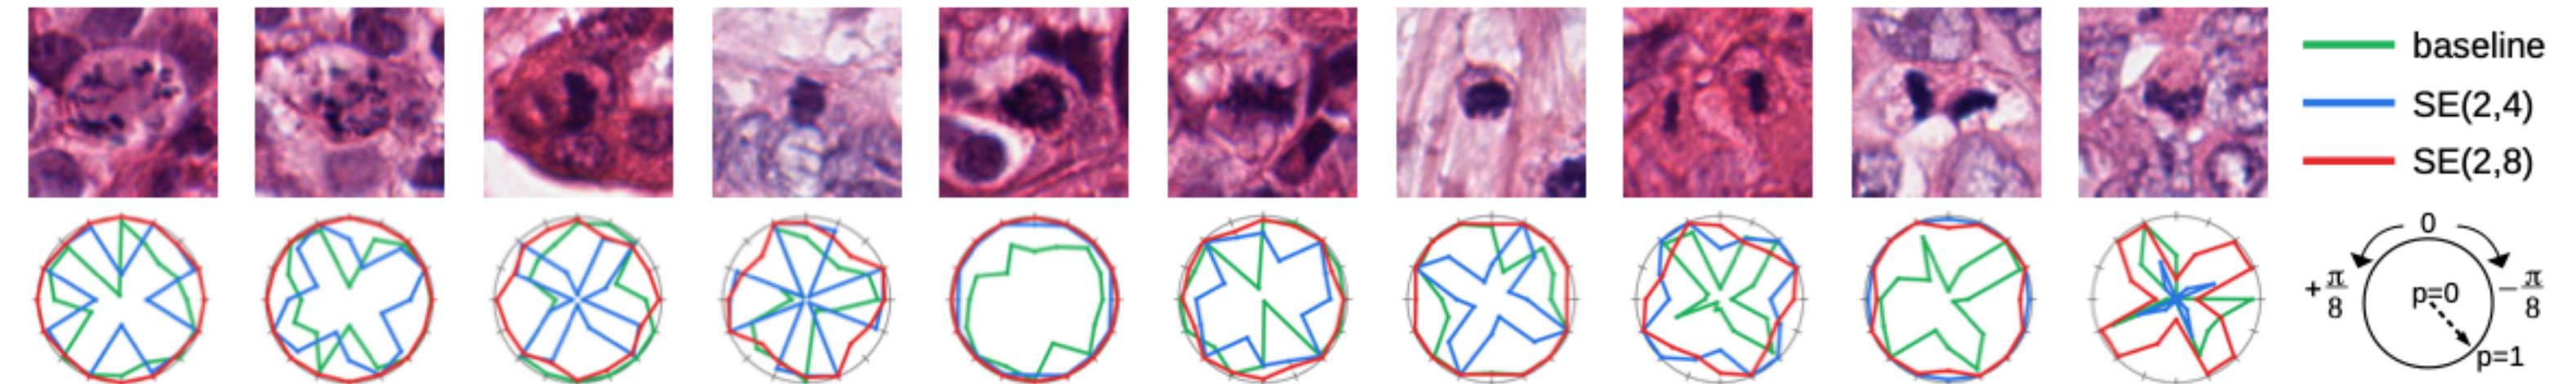
**G-CNNs without data-augmentation**  
outperform  
CNNs with data-augmentation

# Architecture for rotation invariant mitotic cell detection

Bekkers & Lafarge et al. MICCAI 2018



Lafarge et al. MedIA 2020



Lafarge et al. ArXiv/Media 2020

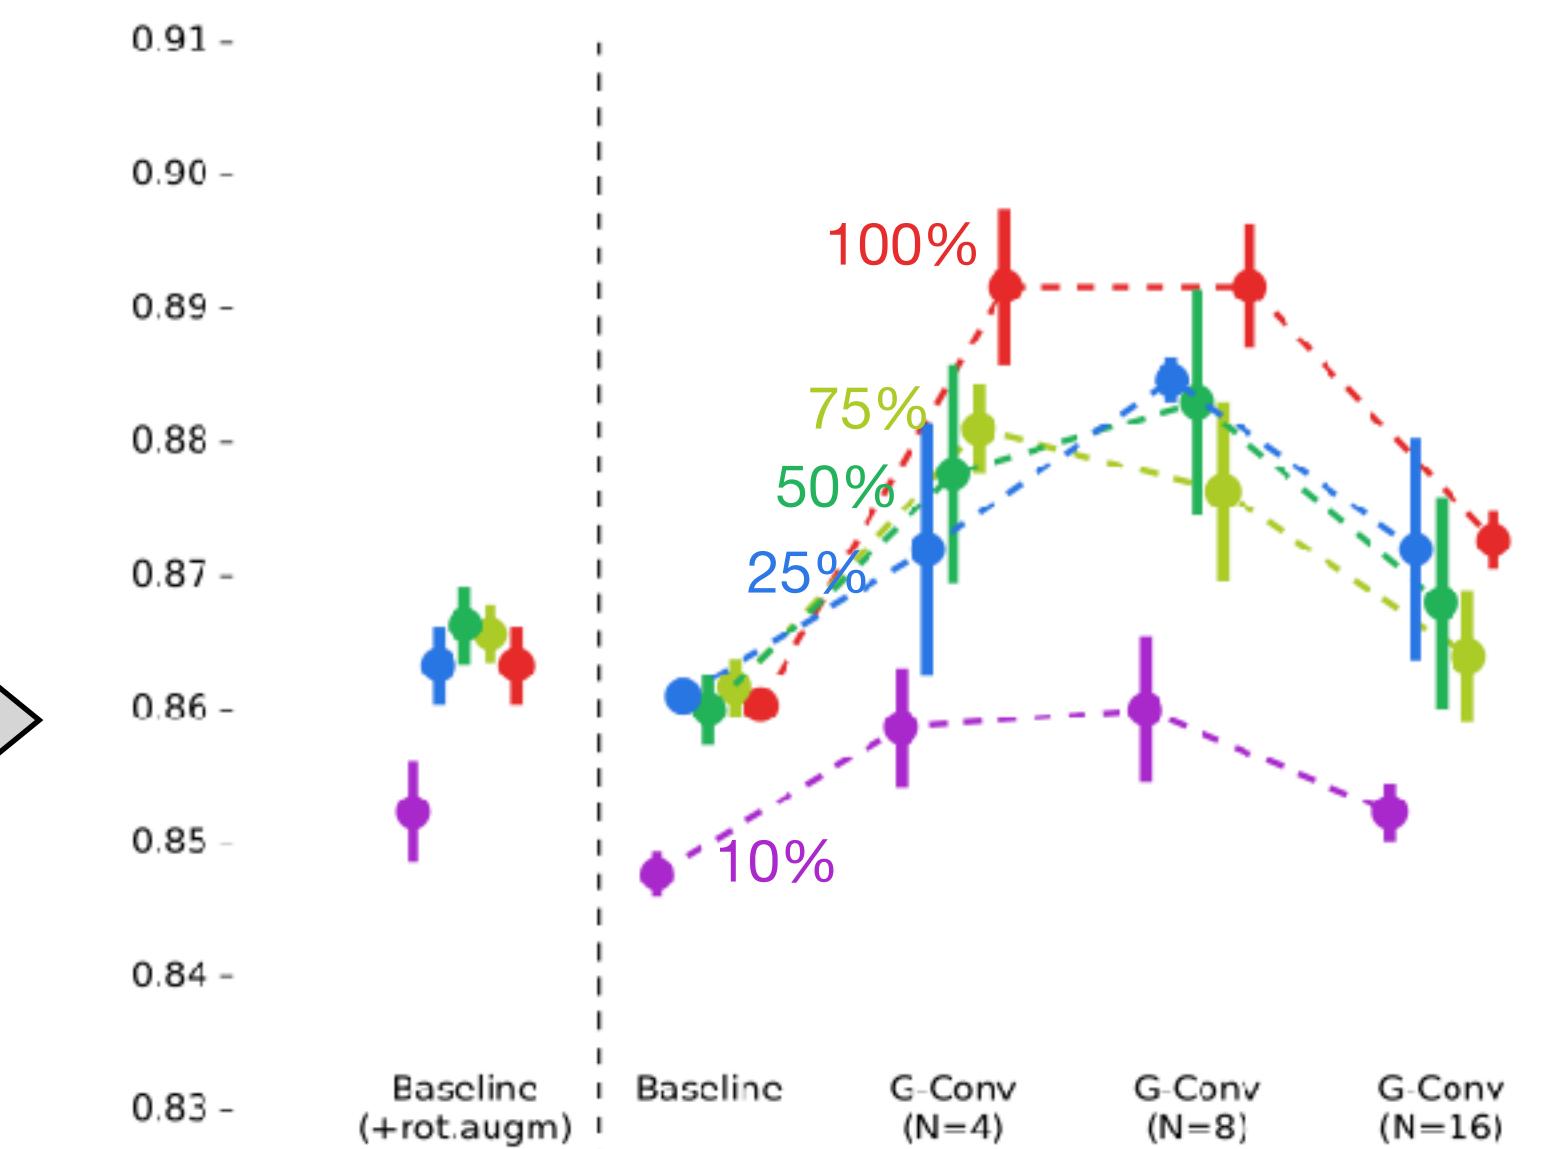
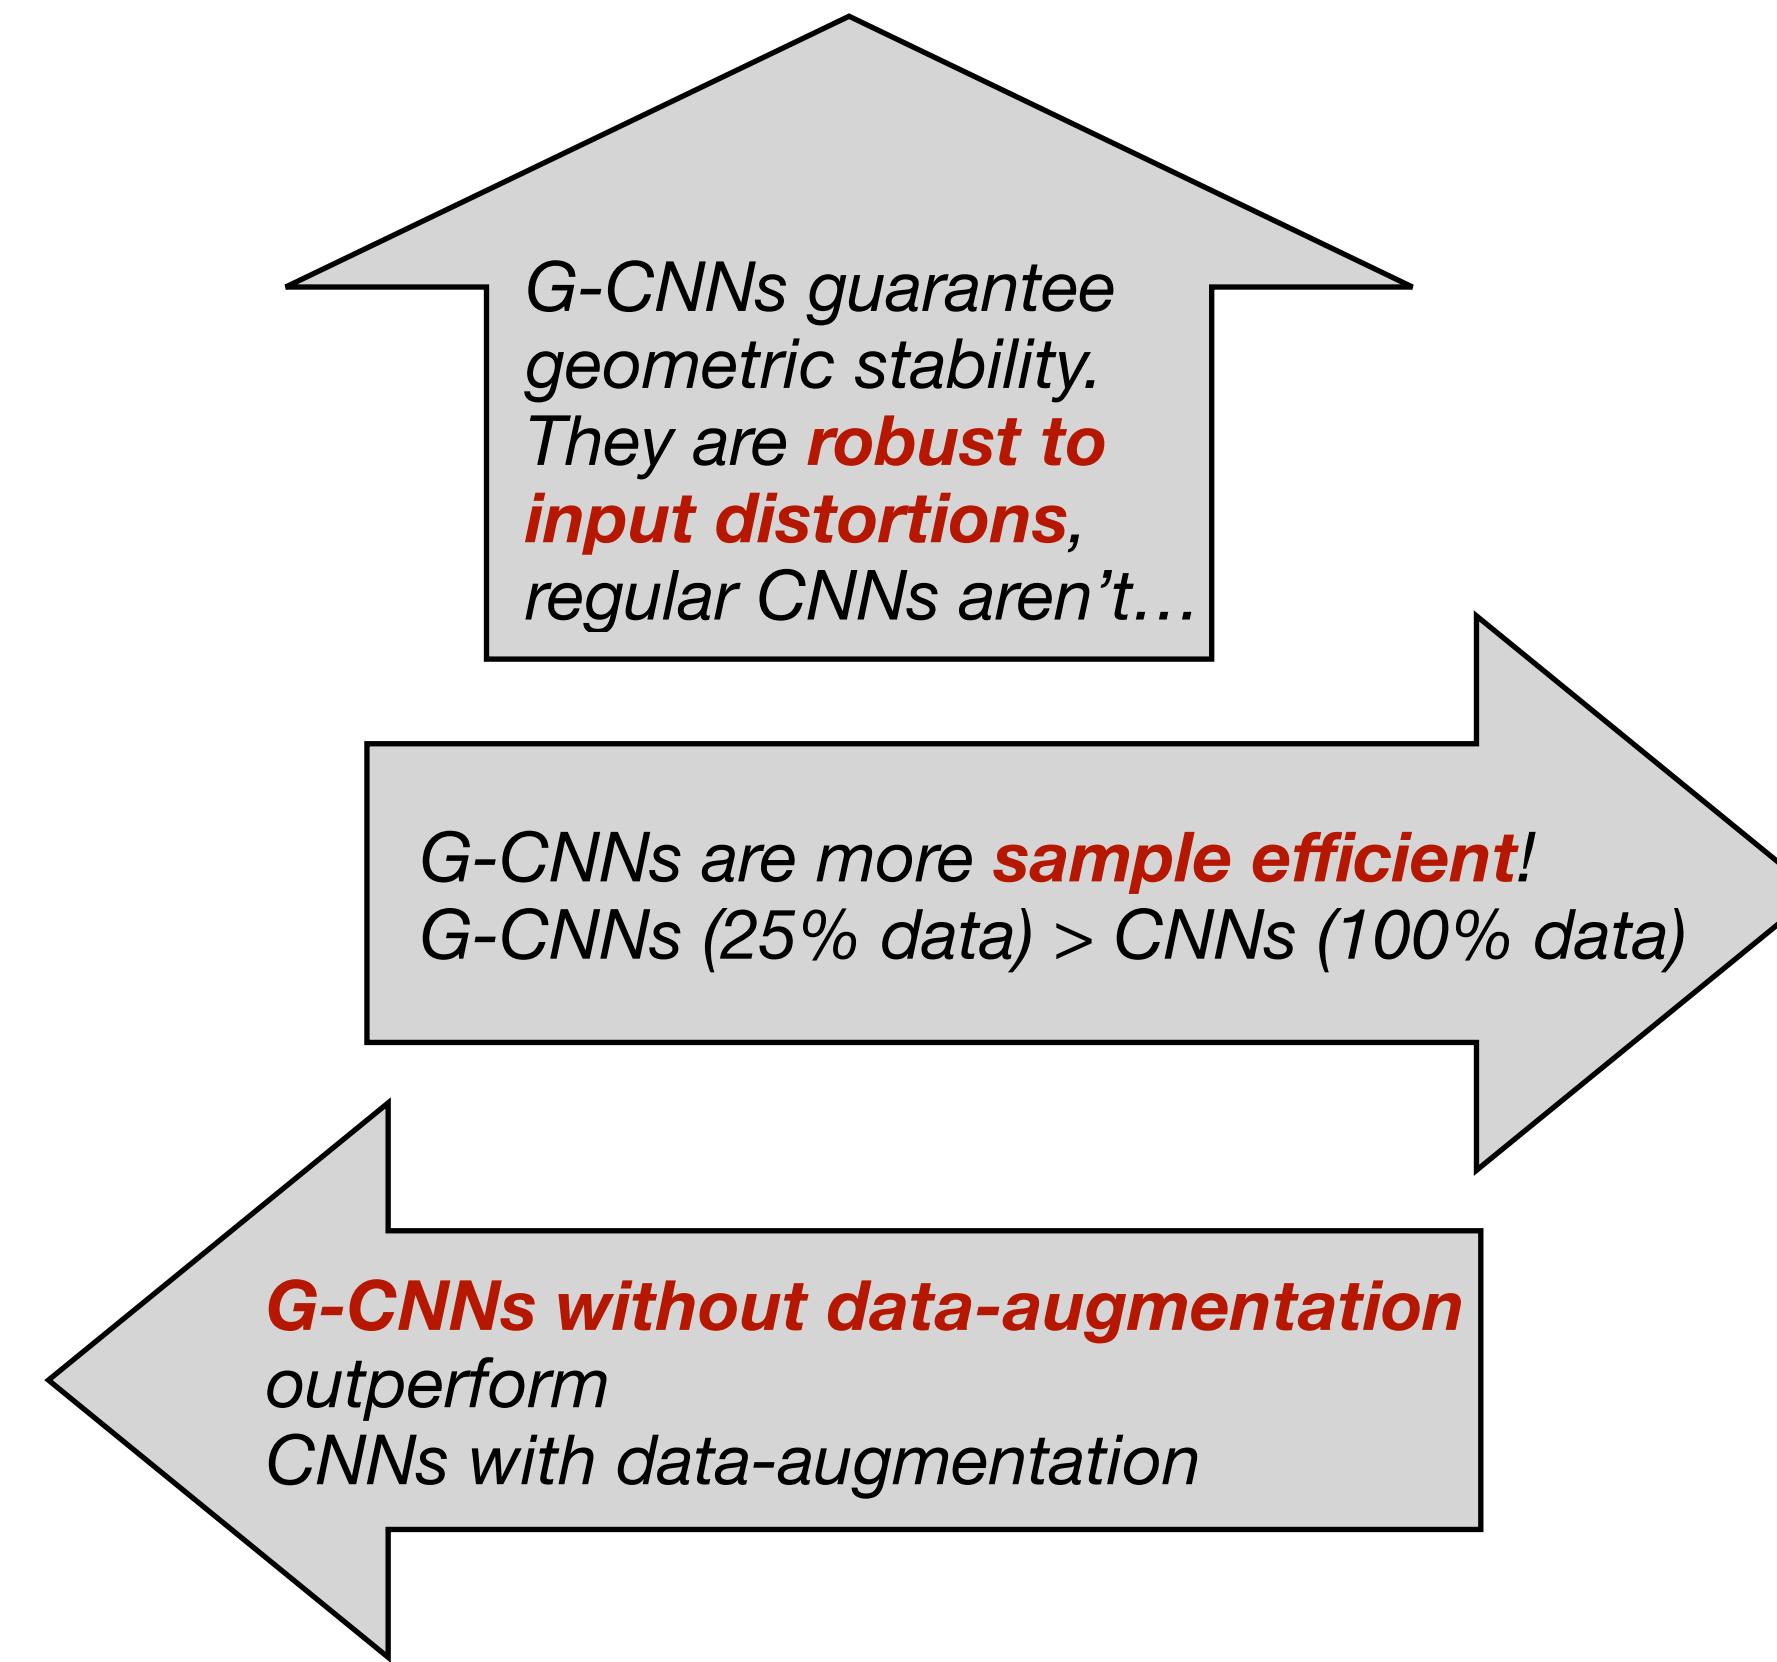


Figure 7: Mean and Standard Deviation plots summarizing the accuracy of the tumor classification models. Mean  $\pm$  standard deviation is indicated. Color identifies the different data regime (red: 100%; lime: 75%; green: 50%; blue: 25%; purple: 10%).

# G-CNNs rule!

- The right inductive bias: guaranteed equivariance  
(no loss of information)
- Performance gains that can't be obtained by data-augmentation alone  
(both local and global equivariance/invariance)
- Increased sample efficiency  
(increased weight sharing, no geometric augmentation necessary)

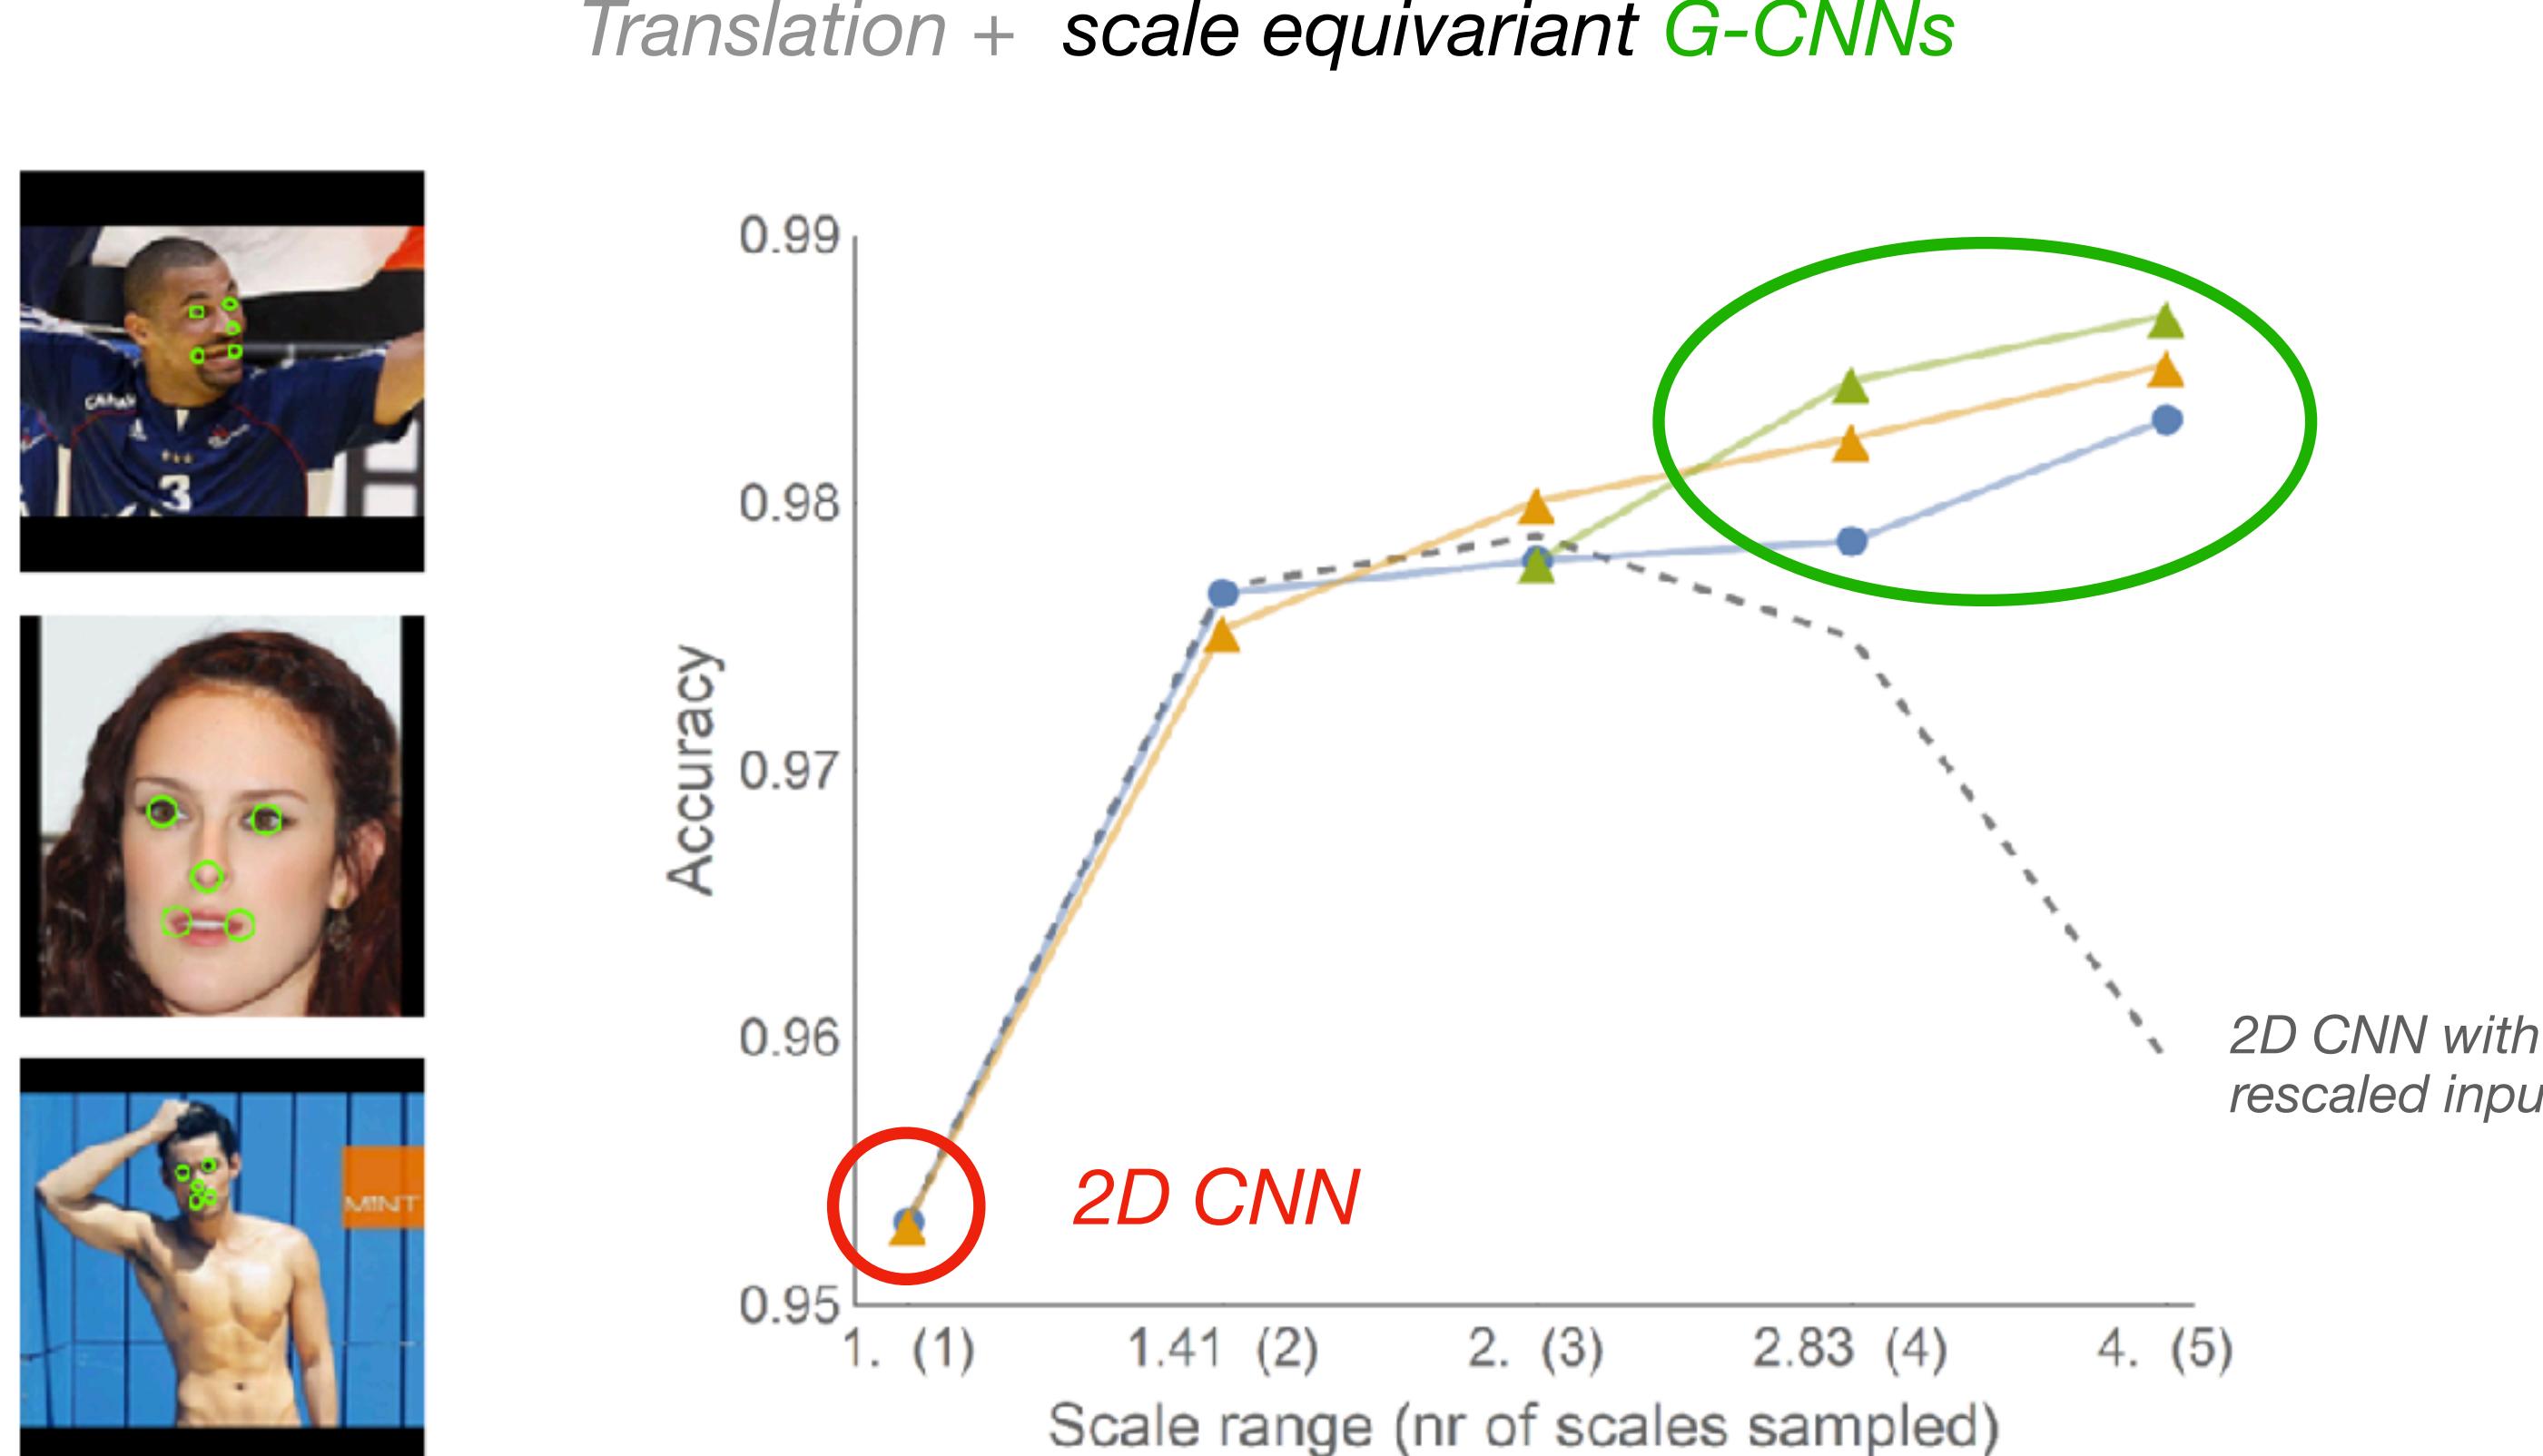
# G-CNNs rule!

- The right inductive bias: guaranteed equivariance (no loss of information)
  - Performance gains that can't be obtained by data-augmentation alone (both local and global equivariance/invariance)
  - Increased sample efficiency  
(increased weight sharing, no geometric augmentation necessary)



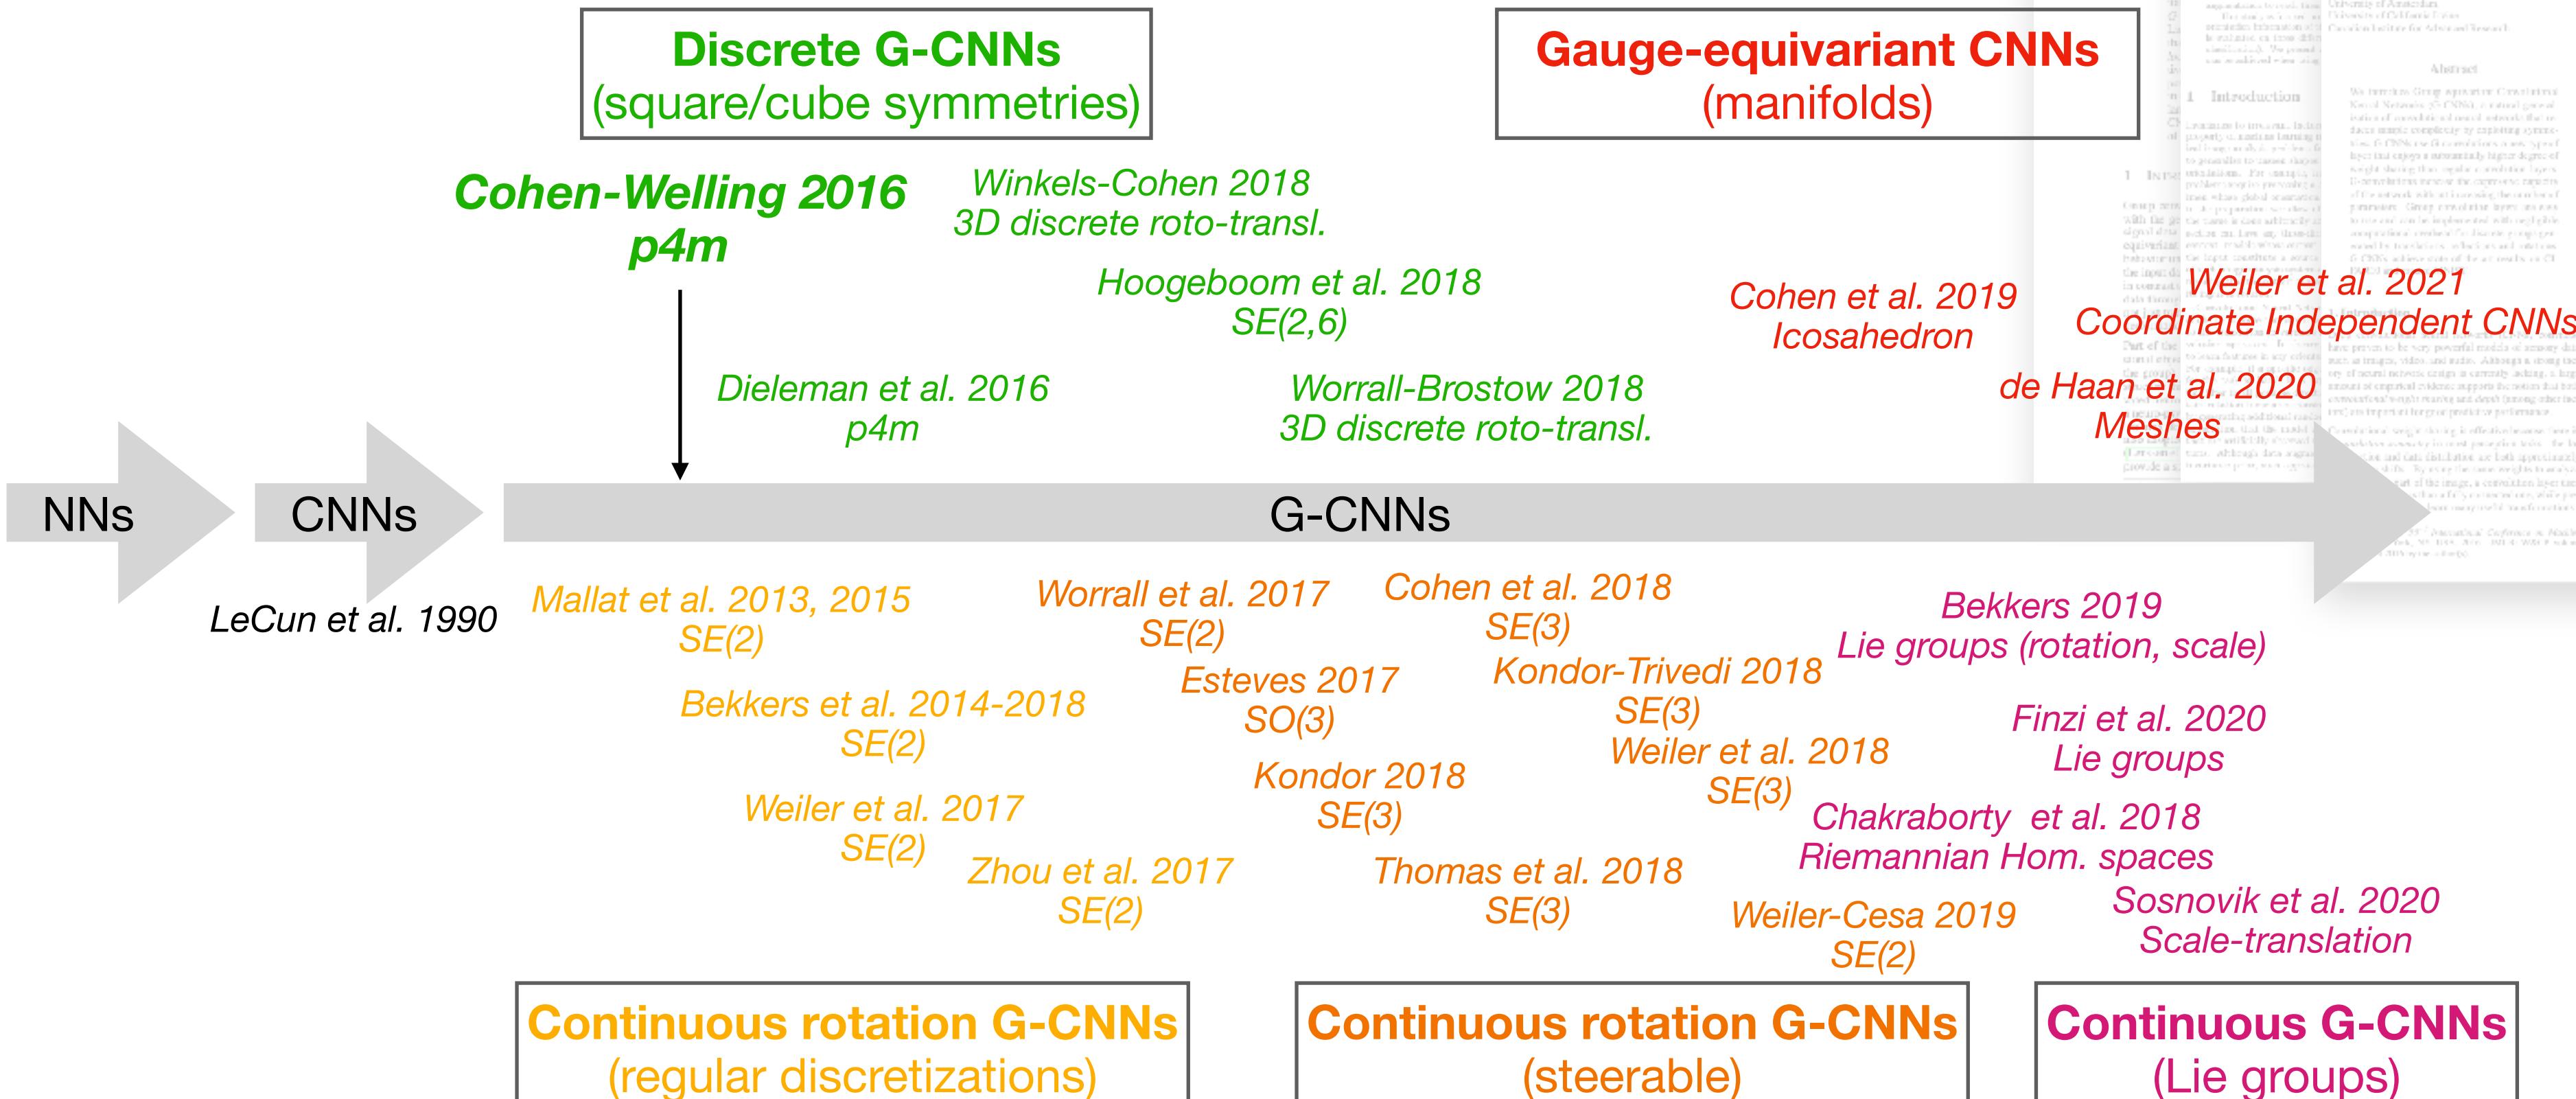
# From rotation to scale equivariant CNNs

Bekkers ICLR 2020



# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



# Cesa-Lang-Weiler 2022

## $\mathbb{R}^d \times H$ with $H$ compact

<https://quva-lab.github.io/escnn/>

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

6. Steerable group convolutions

7. Feature fields and escnn library

8. Equivariant tensor product layers

9. Equivariant graph NNs

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

**5. G-convs are all you need!**

6. Steerable group convolutions

7. Feature fields and escnn library

8. Equivariant tensor product layers

9. Equivariant graph NNs

# Theorem (G-convs are all you need)

Bekkers ICLR 2020, Thm. 1\*

Let  $\mathcal{K} : \mathbb{L}_2(X) \rightarrow \mathbb{L}_2(Y)$  map between signals on homogeneous spaces of  $G$ .

Let homogeneous space  $Y \equiv G/H$  such that  $H = \text{Stab}_G(y_0)$  for some chosen origin  $y_0 \in Y$  and let  $g_y \in G$  such that  $\forall_{y \in Y} : y = g_y y_0$ .

Then  $\mathcal{K}$  is equivariant to group  $G$  if and only if:

1. It is a group convolution:  $[\mathcal{K}f](y) = \int_X \frac{1}{|g_y|} k(g_y^{-1}x)f(x)dx$

2. The kernel satisfies a symmetry constraint:  $\forall_{h \in H} : \frac{1}{|g_y|} k(hx) = k(x)$

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

7. Feature fields and escnn library

8. Equivariant tensor product layers

9. Equivariant graph NNs

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

7. Feature fields and escnn library

8. Equivariant tensor product layers

9. Equivariant graph NNs

# Steerable basis

A vector  $Y(x) = \begin{pmatrix} \vdots \\ Y_l(x) \\ \vdots \end{pmatrix} \in \mathbb{K}^L$  with (basis) functions  $Y_l \in \mathbb{L}_2(X)$  is steerable if

$$\forall_{g \in G} : \quad Y(gx) = \rho(g)Y(x),$$

where  $gx$  denotes the action of  $G$  on  $X$  and  $\rho(g) \in \mathbb{K}^{L \times L}$  is a representation of  $G$ .

*I.e., we can transform all basis functions simply by taking a linear combination of the original basis functions.*

# Function in steerable basis

Let

$$f(x | \hat{\mathbf{w}}) = \hat{\mathbf{w}}^\dagger Y(x) \quad (Y(x) \text{ a steerable basis})$$

Then we can **steer**/shift this function by transforming the weights  $\hat{\mathbf{w}}$

$$f(h^{-1}x | \hat{\mathbf{w}}) = f(x | \rho(h)\hat{\mathbf{w}})$$

---

$$\begin{aligned} (\rho(h)\hat{\mathbf{w}})^\dagger & \left( \begin{array}{c} \text{3D Box Image} \\ \parallel \\ Y(x) \end{array} \right) & = & \text{2D Grayscale Image} \\ & \parallel & & \parallel \\ & f(x | \rho(h)\hat{\mathbf{w}}) & & \end{aligned}$$

# Function in steerable basis

Let

$$f(x | \hat{\mathbf{w}}) = \hat{\mathbf{w}}^\dagger Y(x) \quad (Y(x) \text{ a steerable basis})$$

Then we can **steer**/shift this function by transforming the weights  $\hat{\mathbf{w}}$

$$f(h^{-1}x | \hat{\mathbf{w}}) = f(x | \rho(h)\hat{\mathbf{w}})$$

---

$$\begin{aligned} (\rho(h)\hat{\mathbf{w}})^\dagger & \left( \begin{array}{c} \text{3D Plot of a function} \\ \parallel \\ Y(x) \end{array} \right) & = & \text{2D Grayscale Image} \\ & \parallel & & \parallel \\ & f(x | \rho(h)\hat{\mathbf{w}}) & & \end{aligned}$$

# Example: Steerable basis on $S^1$ (circular harmonics)

Basis functions (for  $\mathbb{L}_2(S^1)$ ):

$$Y_l(\alpha) = e^{il\alpha}$$

Are steered by representations:

$$\rho_l(\theta) = e^{il\theta}$$

Proof:

$$\begin{aligned} Y_l(\alpha - \theta) &= e^{il(\alpha - \theta)} \\ &= e^{-il\theta} e^{ila} \\ &= \rho_l(-\theta) Y_l(\alpha) \end{aligned}$$

# Example: Steerable basis on $S^1$ (circular harmonics)

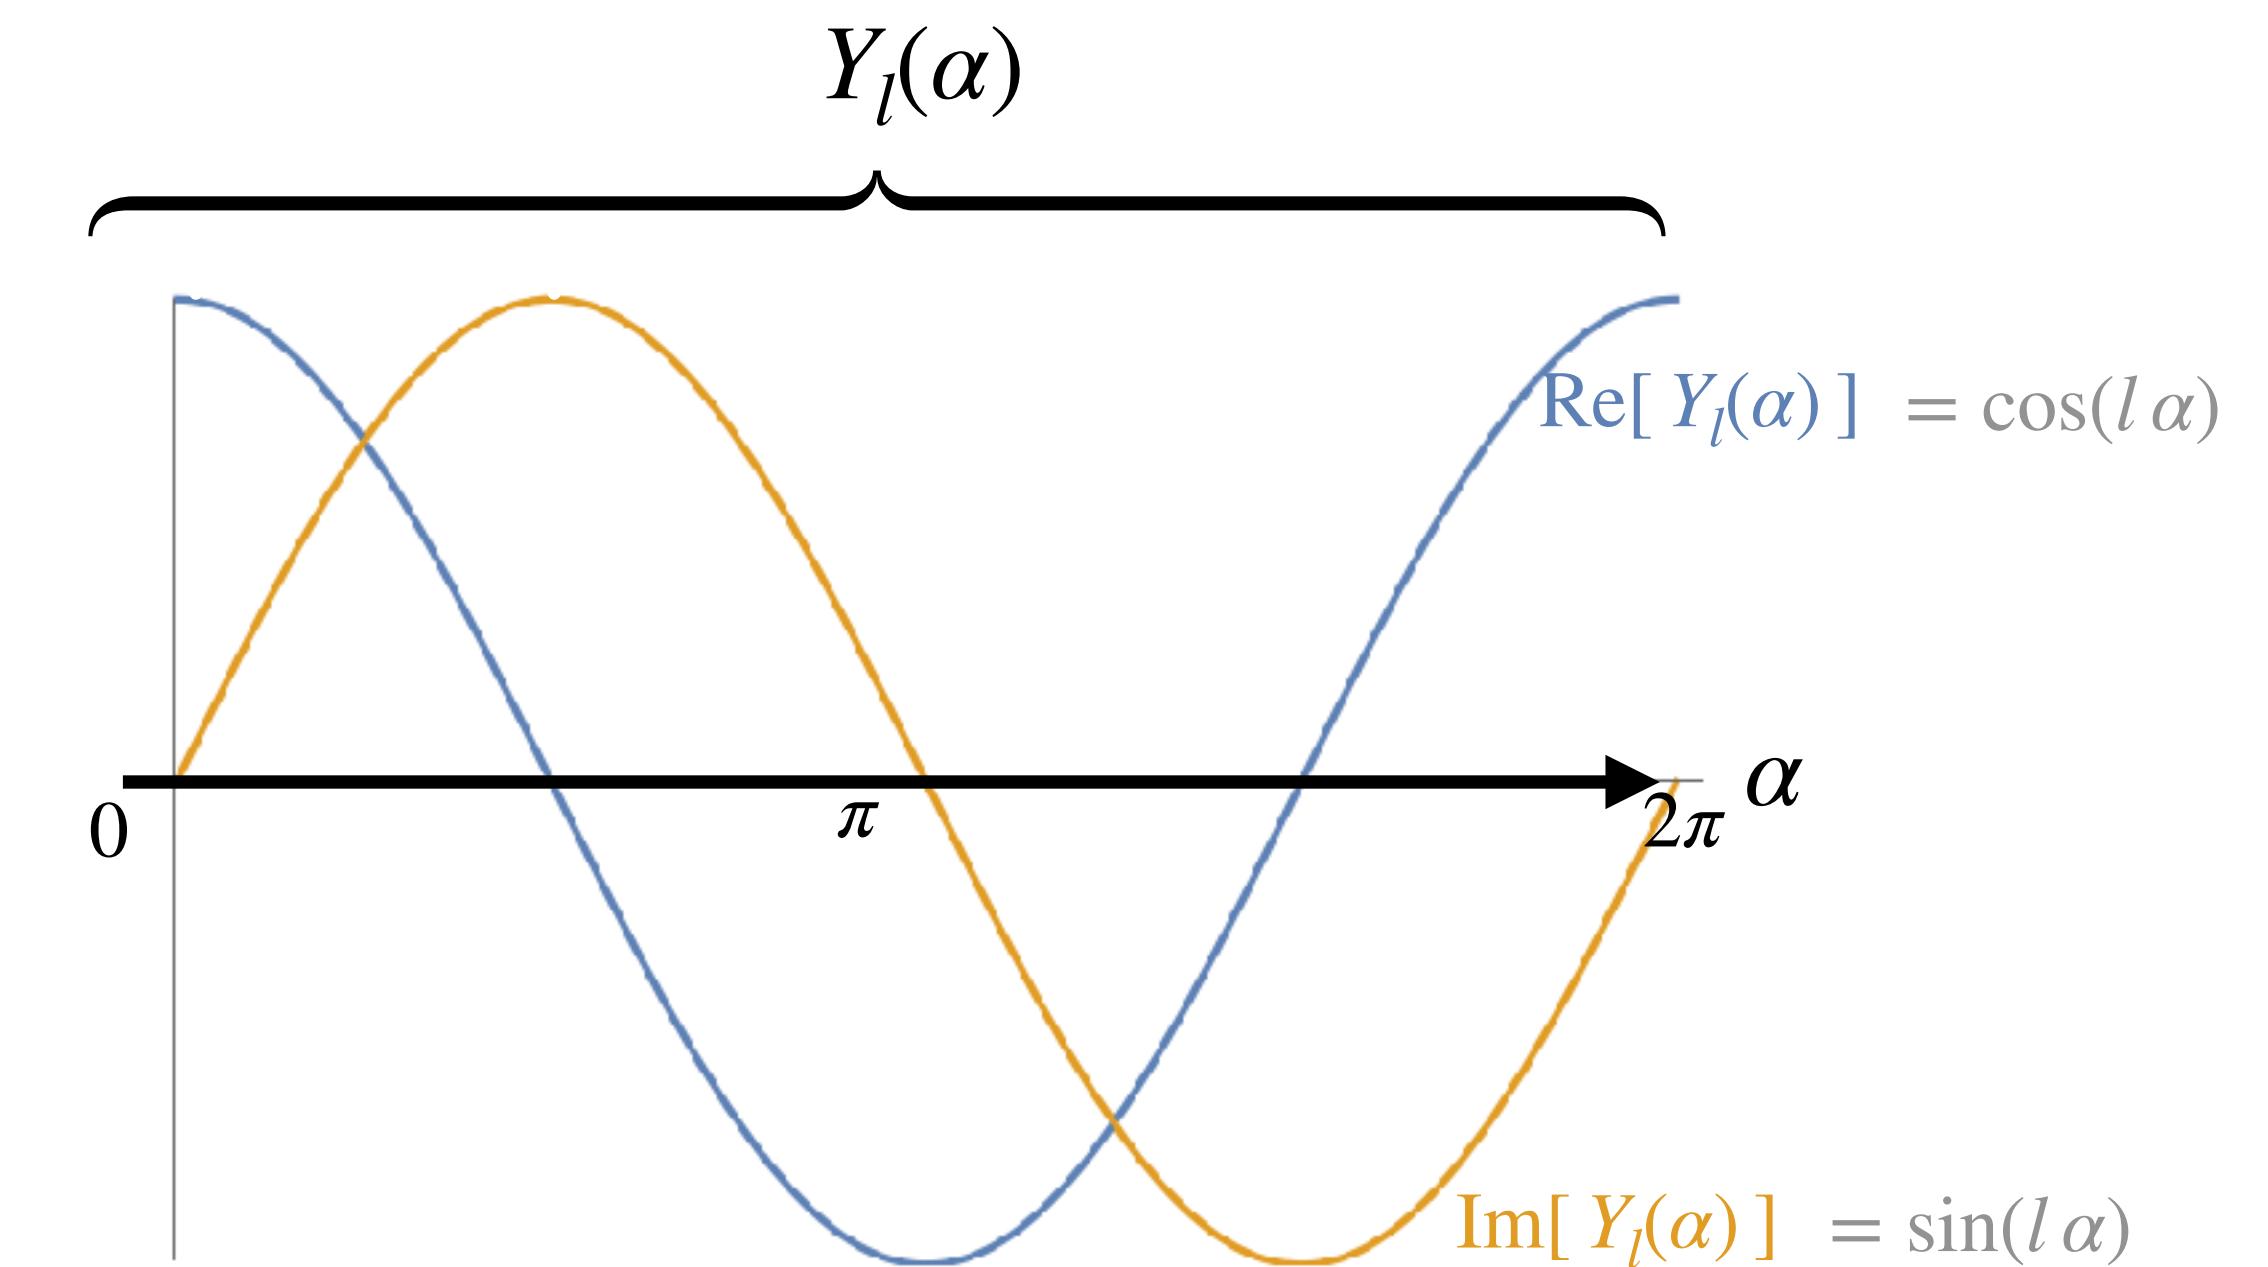
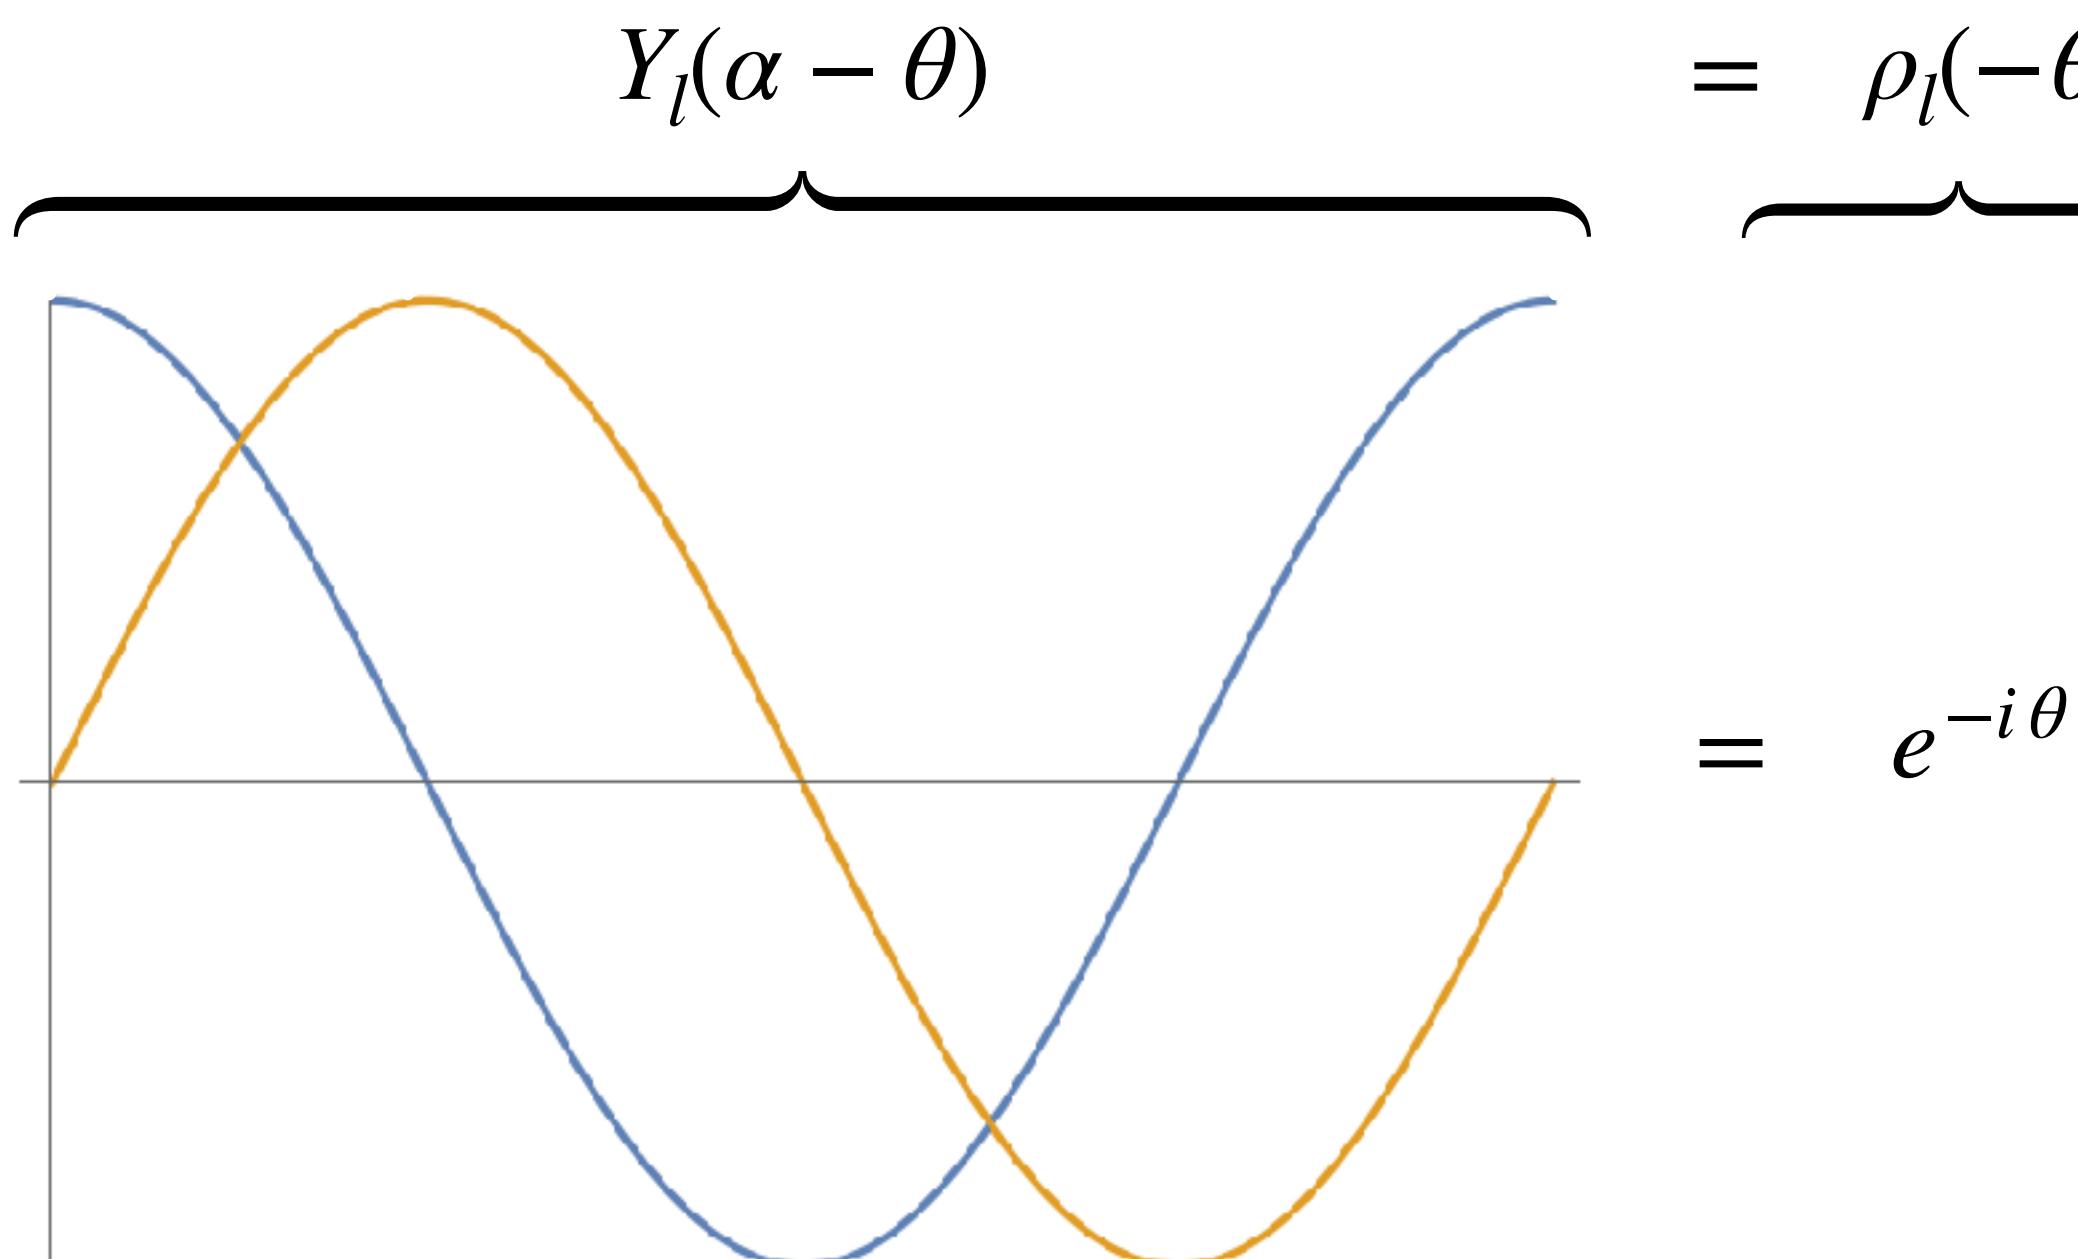
Basis functions (for  $\mathbb{L}_2(S^1)$ ):

Are steered by representations:

$$Y_l(\alpha) = e^{il\alpha}$$

$$\rho_l(\theta) = e^{il\theta}$$

Proof:  $Y_l(\alpha - \theta) = e^{il(\alpha-\theta)}$   
 $= e^{-il\theta} e^{ila}$   
 $= \rho_l(-\theta) Y_l(\alpha)$



# Example: Steerable basis on $S^1$ (circular harmonics)

$$\begin{array}{c}
 Y(\alpha - \theta) \\
 \left( \begin{array}{c} \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \end{array} \right) \\
 = \left( \begin{array}{cccccc} e^{i3\theta} & 0 & 0 & 0 & 0 & 0 \\ 0 & e^{i2\theta} & 0 & 0 & 0 & 0 \\ 0 & 0 & e^{i1\theta} & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & e^{-i1\theta} & 0 \\ 0 & 0 & 0 & 0 & 0 & e^{-i2\theta} \\ 0 & 0 & 0 & 0 & 0 & 0 \end{array} \right) \left( \begin{array}{c} \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \end{array} \right) \\
 Y(\alpha)
 \end{array}$$

$\rho(-\theta) = \bigoplus_{l=-L}^L \rho_l(-\theta)$

# Example: Steerable basis on $S^1$ (circular harmonics)

$$\begin{array}{c}
 Y(\alpha - \theta) \\
 \left( \begin{array}{c} \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \end{array} \right) \\
 = \left( \begin{array}{cccccc} e^{i3\theta} & 0 & 0 & 0 & 0 & 0 \\ 0 & e^{i2\theta} & 0 & 0 & 0 & 0 \\ 0 & 0 & e^{i1\theta} & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & e^{-i1\theta} & 0 \\ 0 & 0 & 0 & 0 & 0 & e^{-i2\theta} \\ 0 & 0 & 0 & 0 & 0 & 0 \end{array} \right) \left( \begin{array}{c} \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \\ \vdots \\ \text{blue wave} \\ \text{orange wave} \end{array} \right) \\
 Y(\alpha)
 \end{array}$$

$\rho(-\theta) = \bigoplus_{l=-L}^L \rho_l(-\theta)$

# Example: Steerable basis on $S^1$ (circular harmonics)

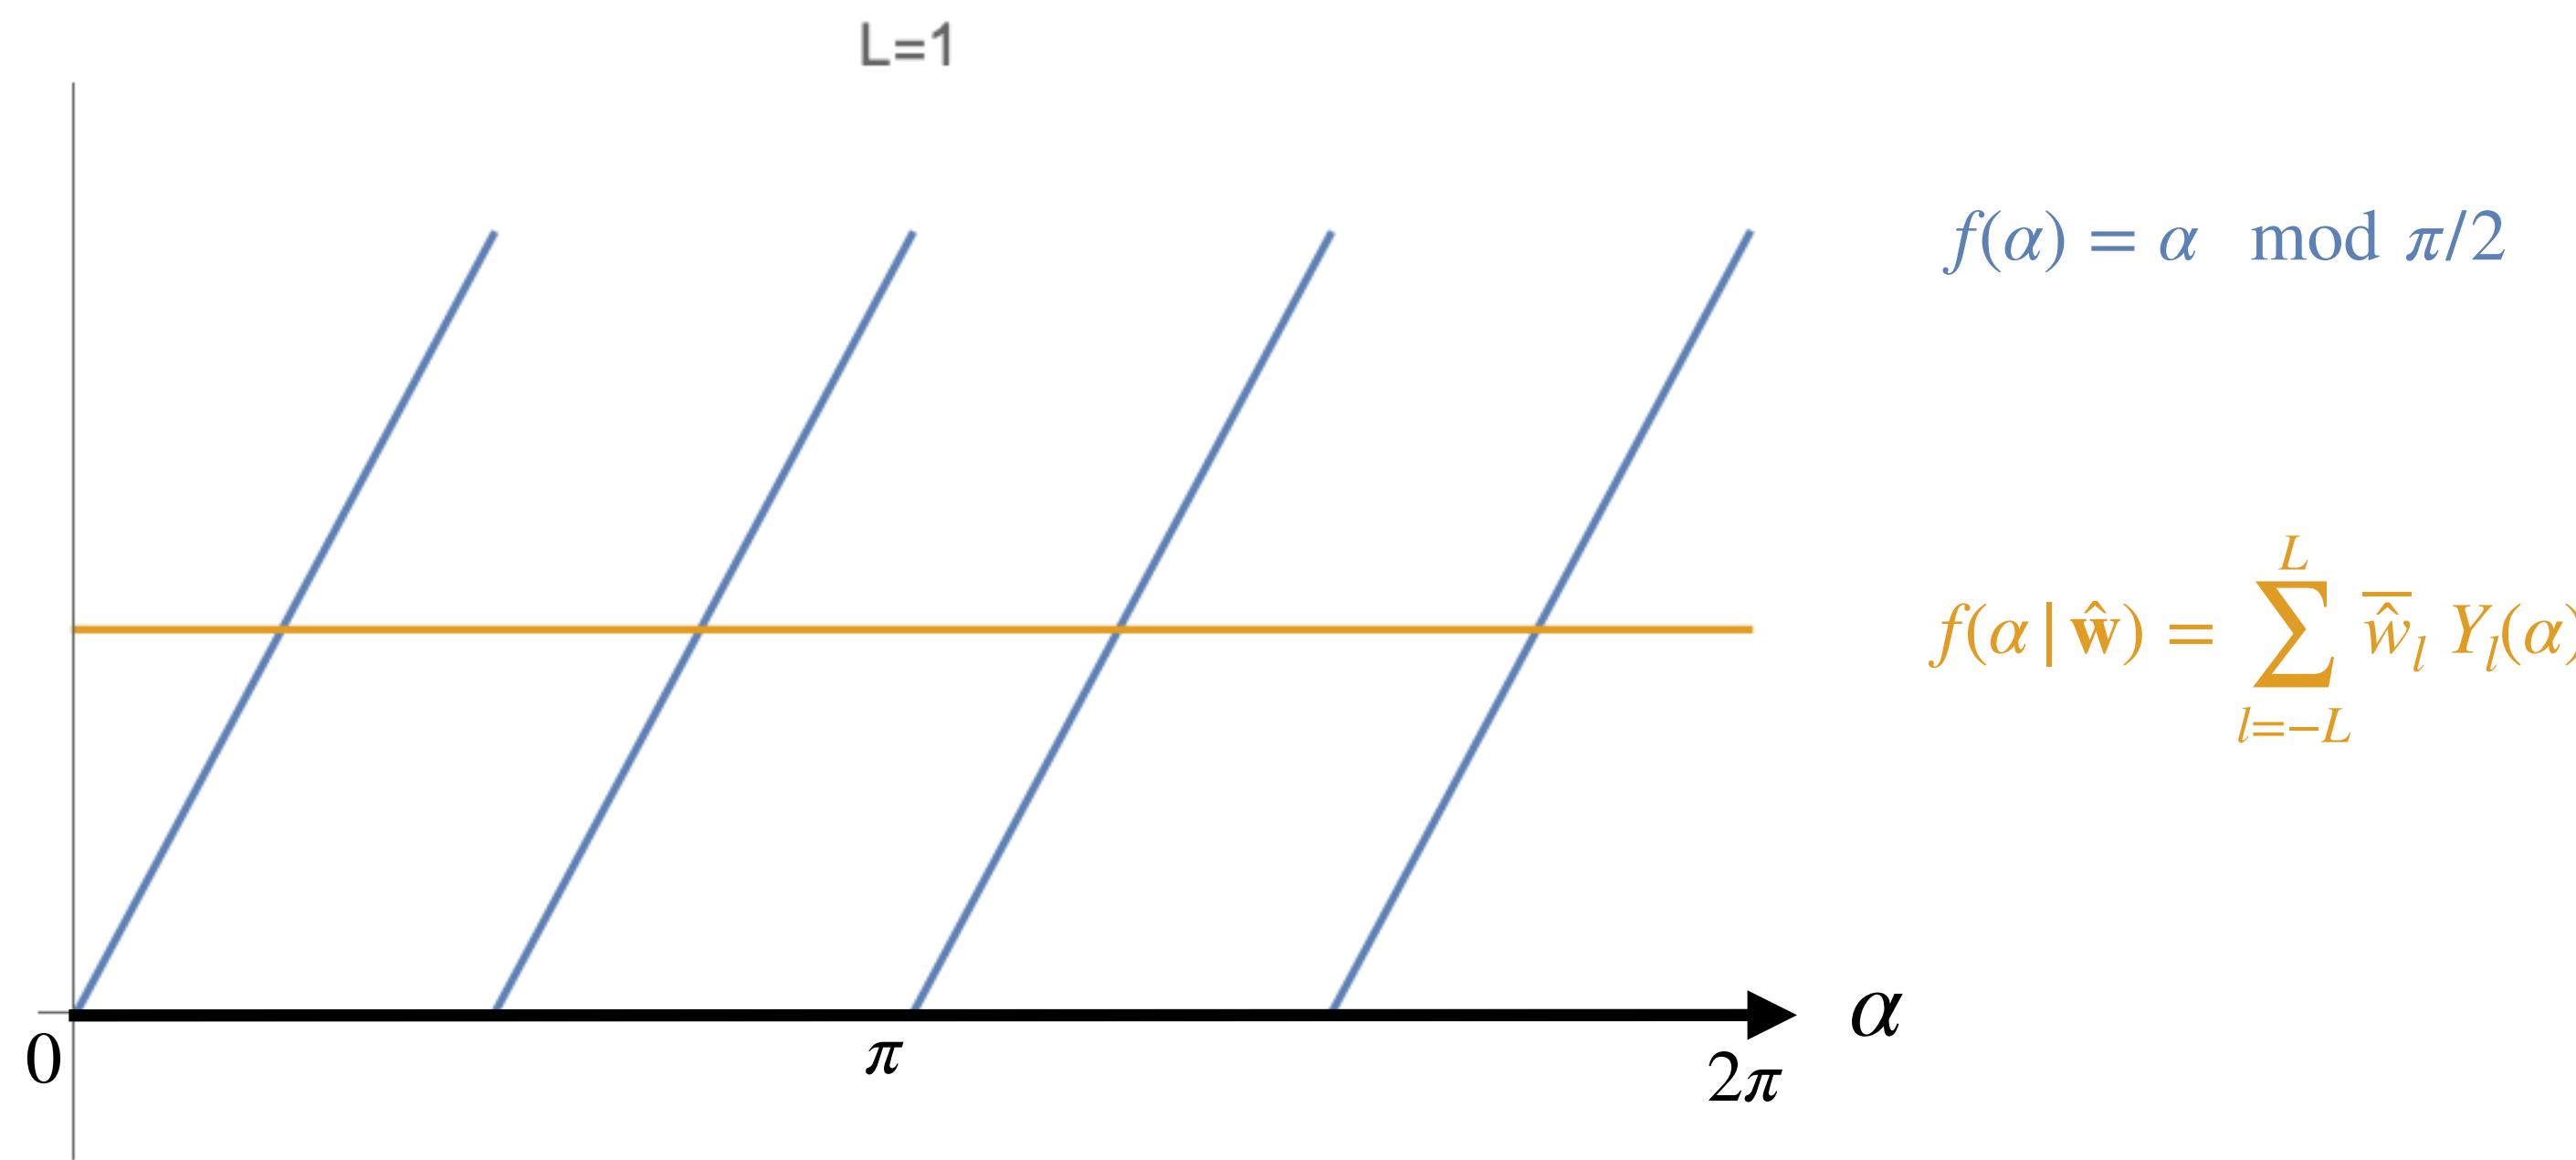
Basis functions (for  $\mathbb{L}_2(S^1)$ ):

$$Y_l(\alpha) = e^{il\alpha}$$

**Form a complete orthonormal (Fourier) basis:**

$Y_l$  are given by the irreps of  $SO(2)$  and hence form orthogonal basis (Peter-Weyl Theorem)

$$f(\alpha | \hat{\mathbf{w}}) = \sum_{l=-\infty}^{\infty} \overline{\hat{w}_l} Y_l(\alpha)$$



# Example: Steerable basis on $S^1$ (circular harmonics)

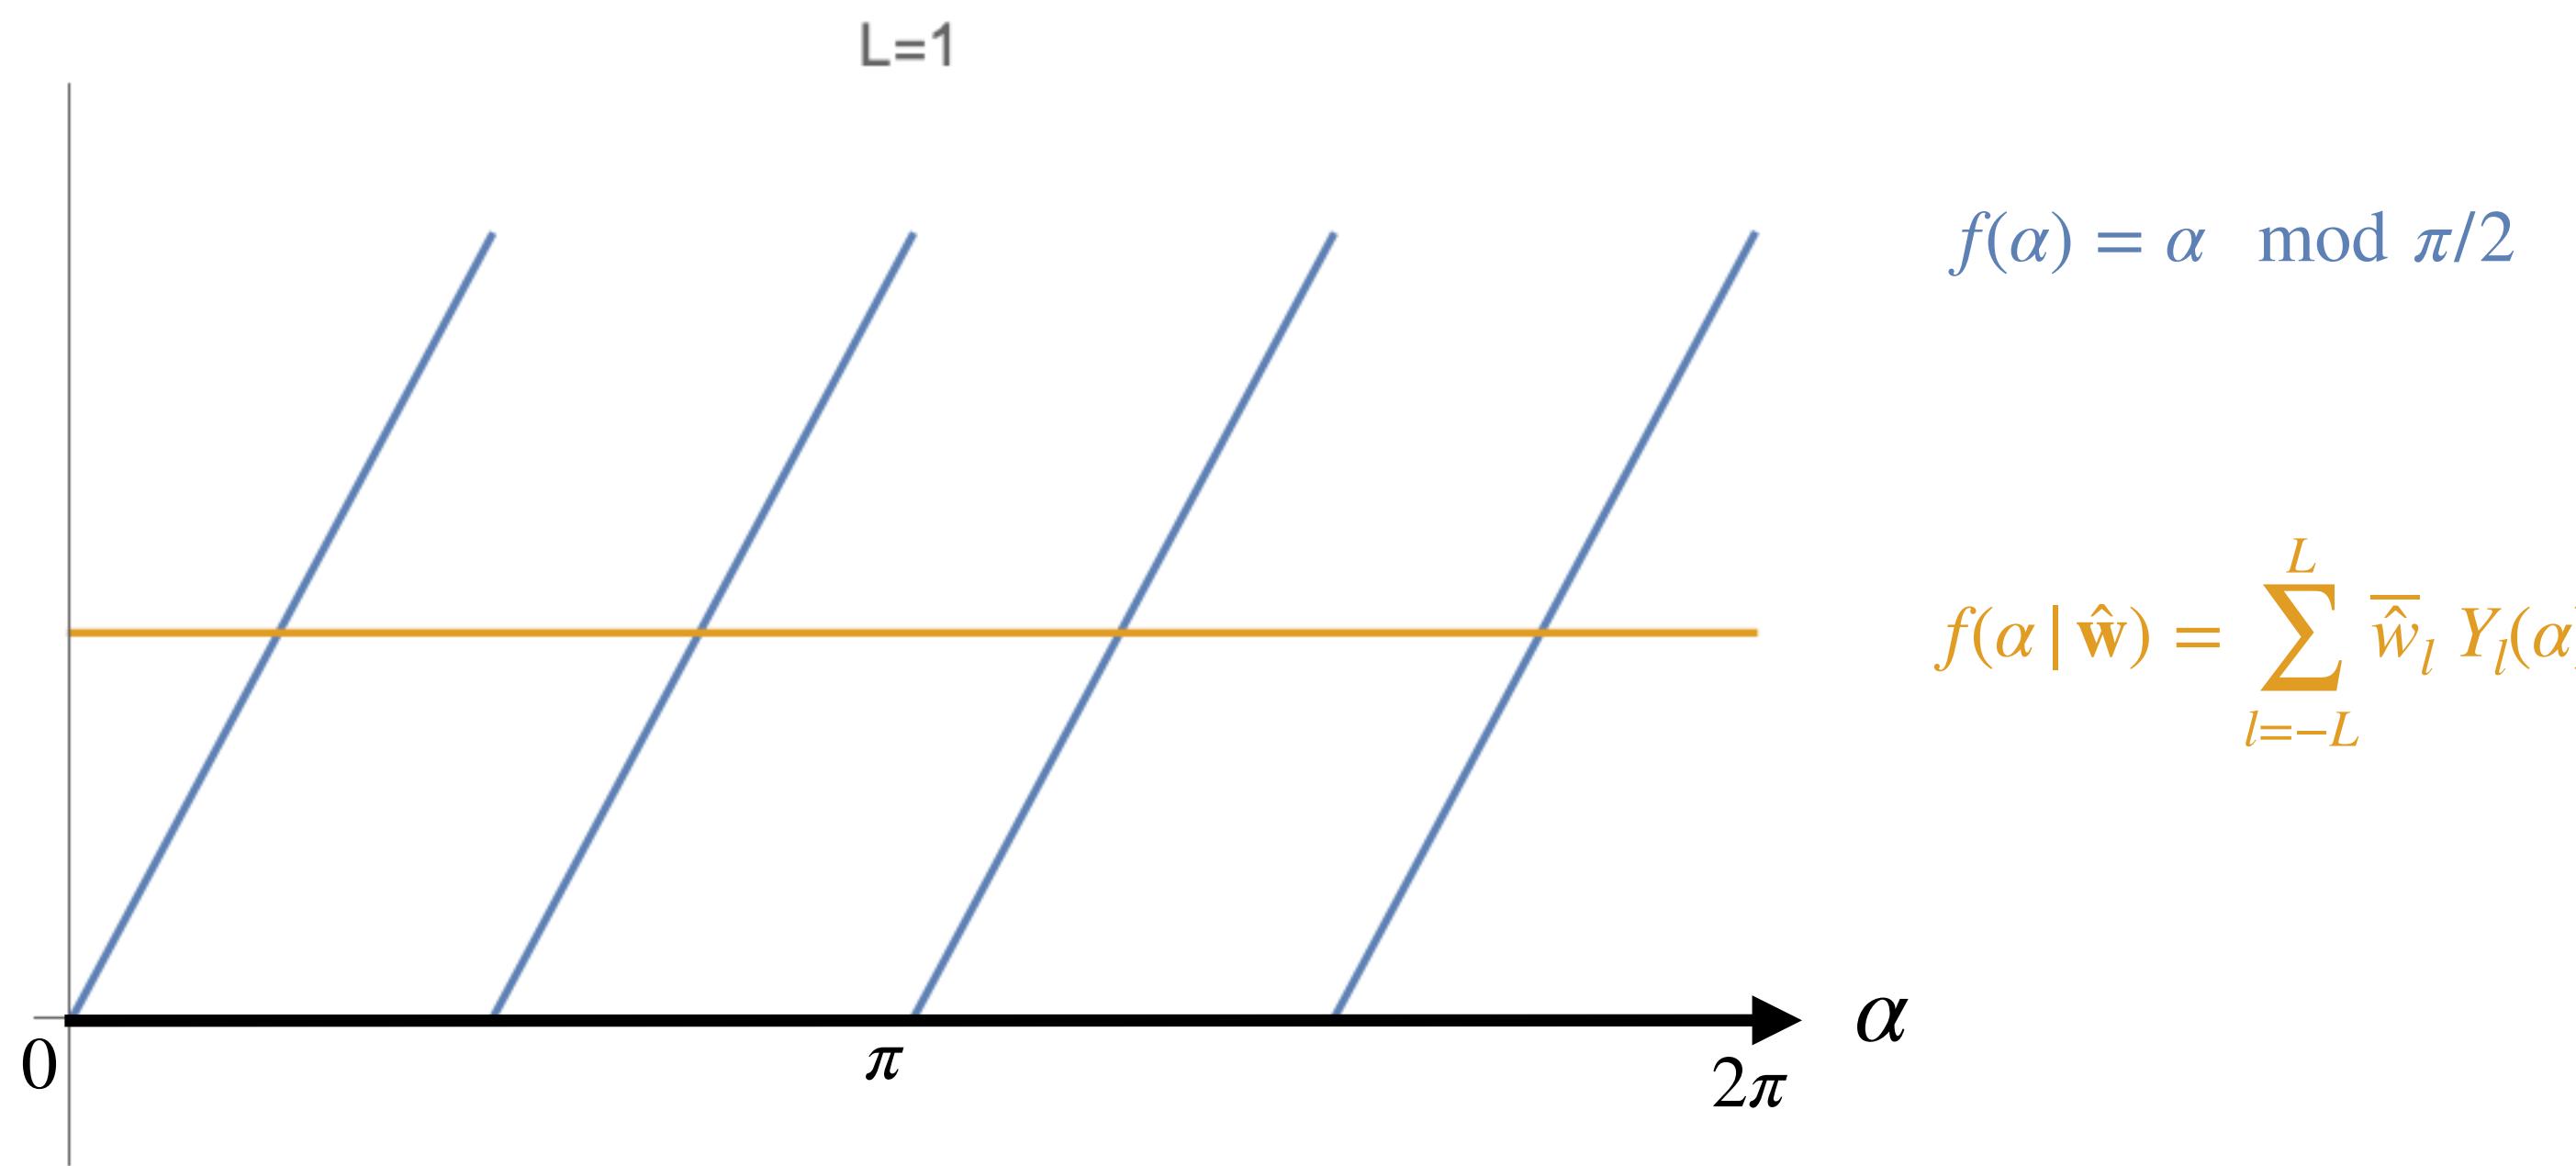
Basis functions (for  $\mathbb{L}_2(S^1)$ ):

$$Y_l(\alpha) = e^{il\alpha}$$

**Form a complete orthonormal (Fourier) basis:**

$Y_l$  are given by the irreps of  $SO(2)$  and hence form orthogonal basis (Peter-Weyl Theorem)

$$f(\alpha | \hat{\mathbf{w}}) = \sum_{l=-\infty}^{\infty} \hat{w}_l \overline{Y_l}(\alpha)$$



# Example: Steerable basis on $S^1$ (circular harmonics)

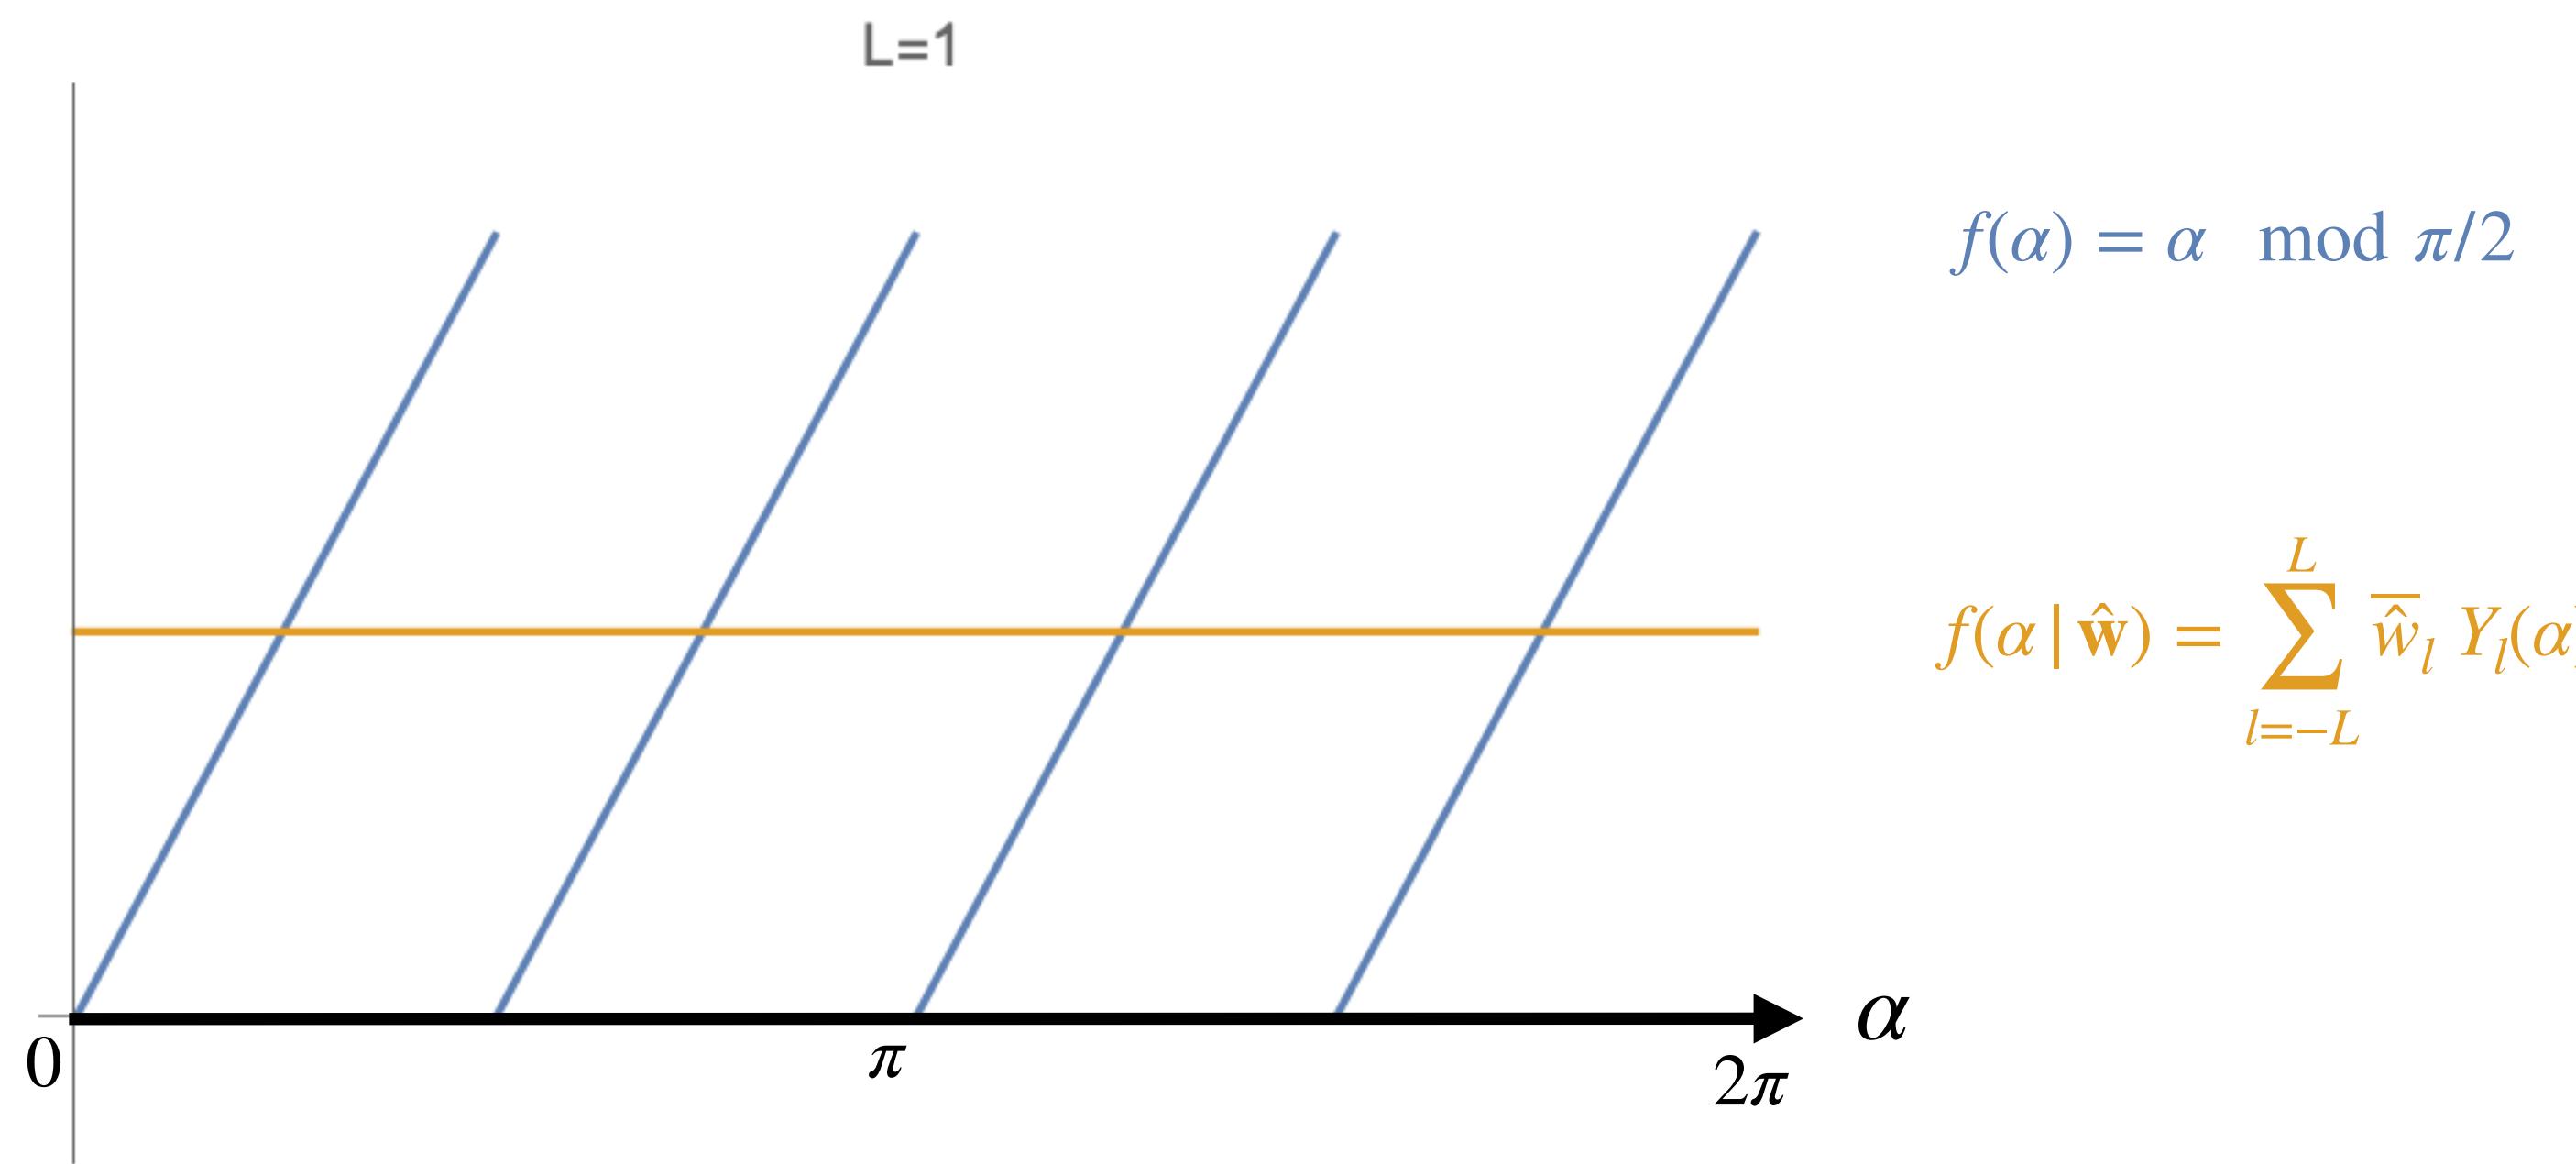
Basis functions (for  $\mathbb{L}_2(S^1)$ ):

$$Y_l(\alpha) = e^{il\alpha}$$

**Form a complete orthonormal (Fourier) basis:**

$$f(\alpha | \hat{\mathbf{w}}) = \sum_{l=-\infty}^{\infty} \hat{w}_l Y_l(-\alpha)$$

$Y_l$  are given by the irreps of  $SO(2)$  and hence form orthogonal basis (Peter-Weyl Theorem)



# Example: Steerable basis on $S^1$ (circular harmonics)

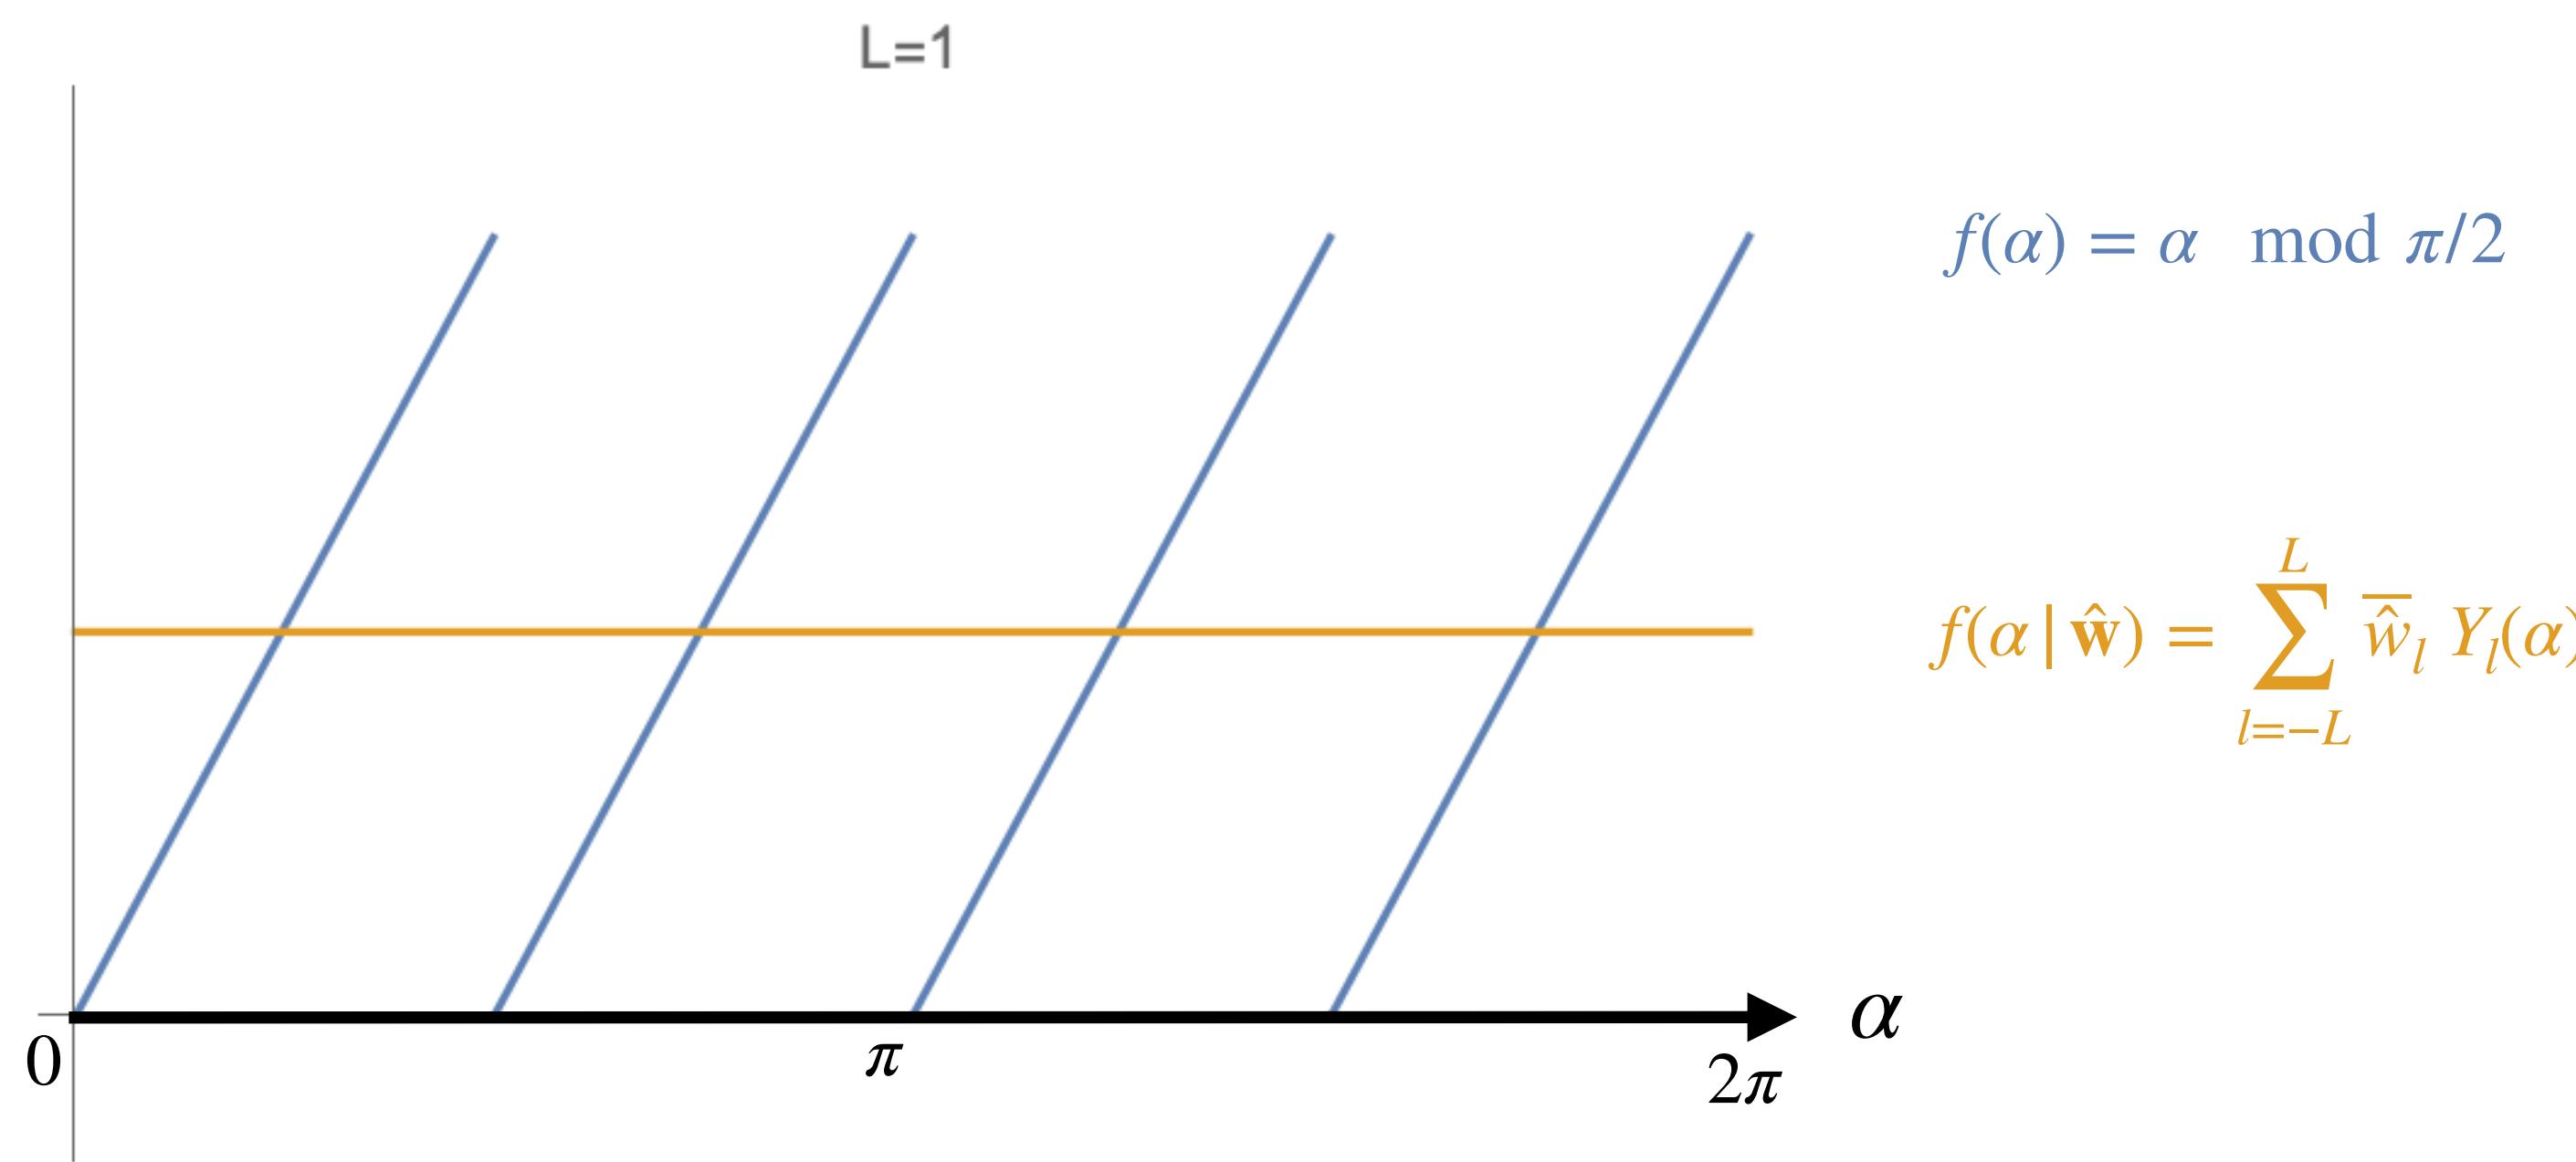
Basis functions (for  $\mathbb{L}_2(S^1)$ ):

$$Y_l(\alpha) = e^{il\alpha}$$

**Form a complete orthonormal (Fourier) basis:**

$Y_l$  are given by the irreps of  $SO(2)$  and hence form orthogonal basis (Peter-Weyl Theorem)

$$f(\alpha | \hat{\mathbf{w}}) = \sum_{l=-\infty}^{\infty} \overline{\hat{w}_l} Y_l(\alpha)$$



# Example: Steerable basis on $S^1$ (circular harmonics)

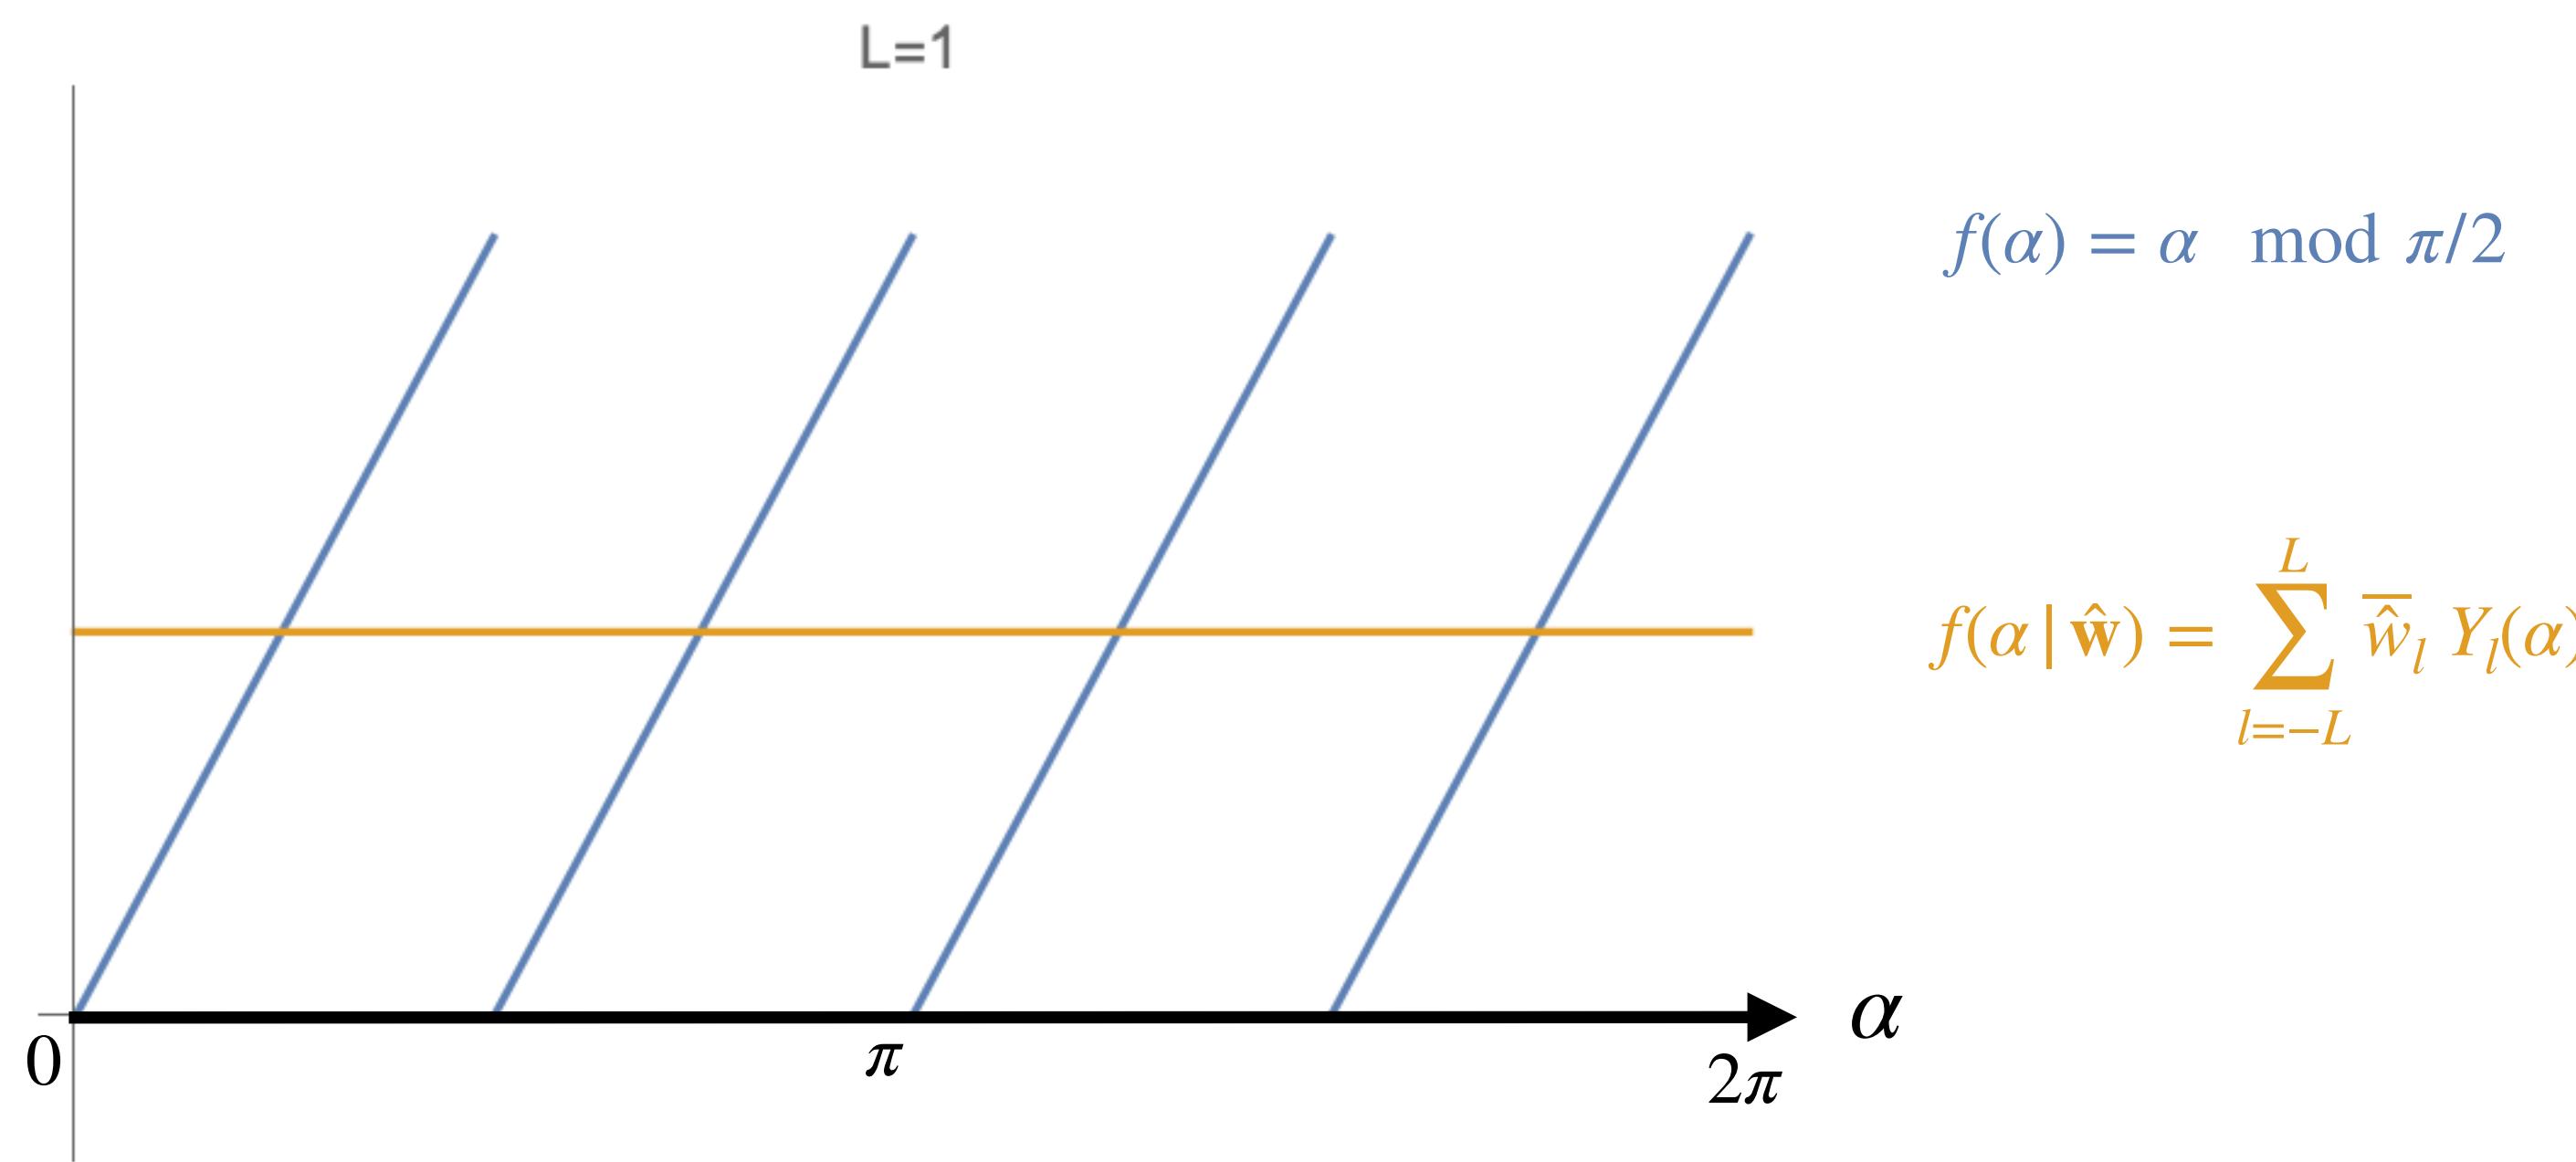
Basis functions (for  $\mathbb{L}_2(S^1)$ ):

$$Y_l(\alpha) = e^{il\alpha}$$

**Form a complete orthonormal (Fourier) basis:**

$Y_l$  are given by the irreps of  $SO(2)$  and hence form orthogonal basis (Peter-Weyl Theorem)

$$f(\alpha | \hat{\mathbf{w}}) = \sum_{l=-\infty}^{\infty} \overline{\hat{w}_l} Y_l(\alpha)$$



# Example: Steerable basis on $S^1$ (circular harmonics)

Let

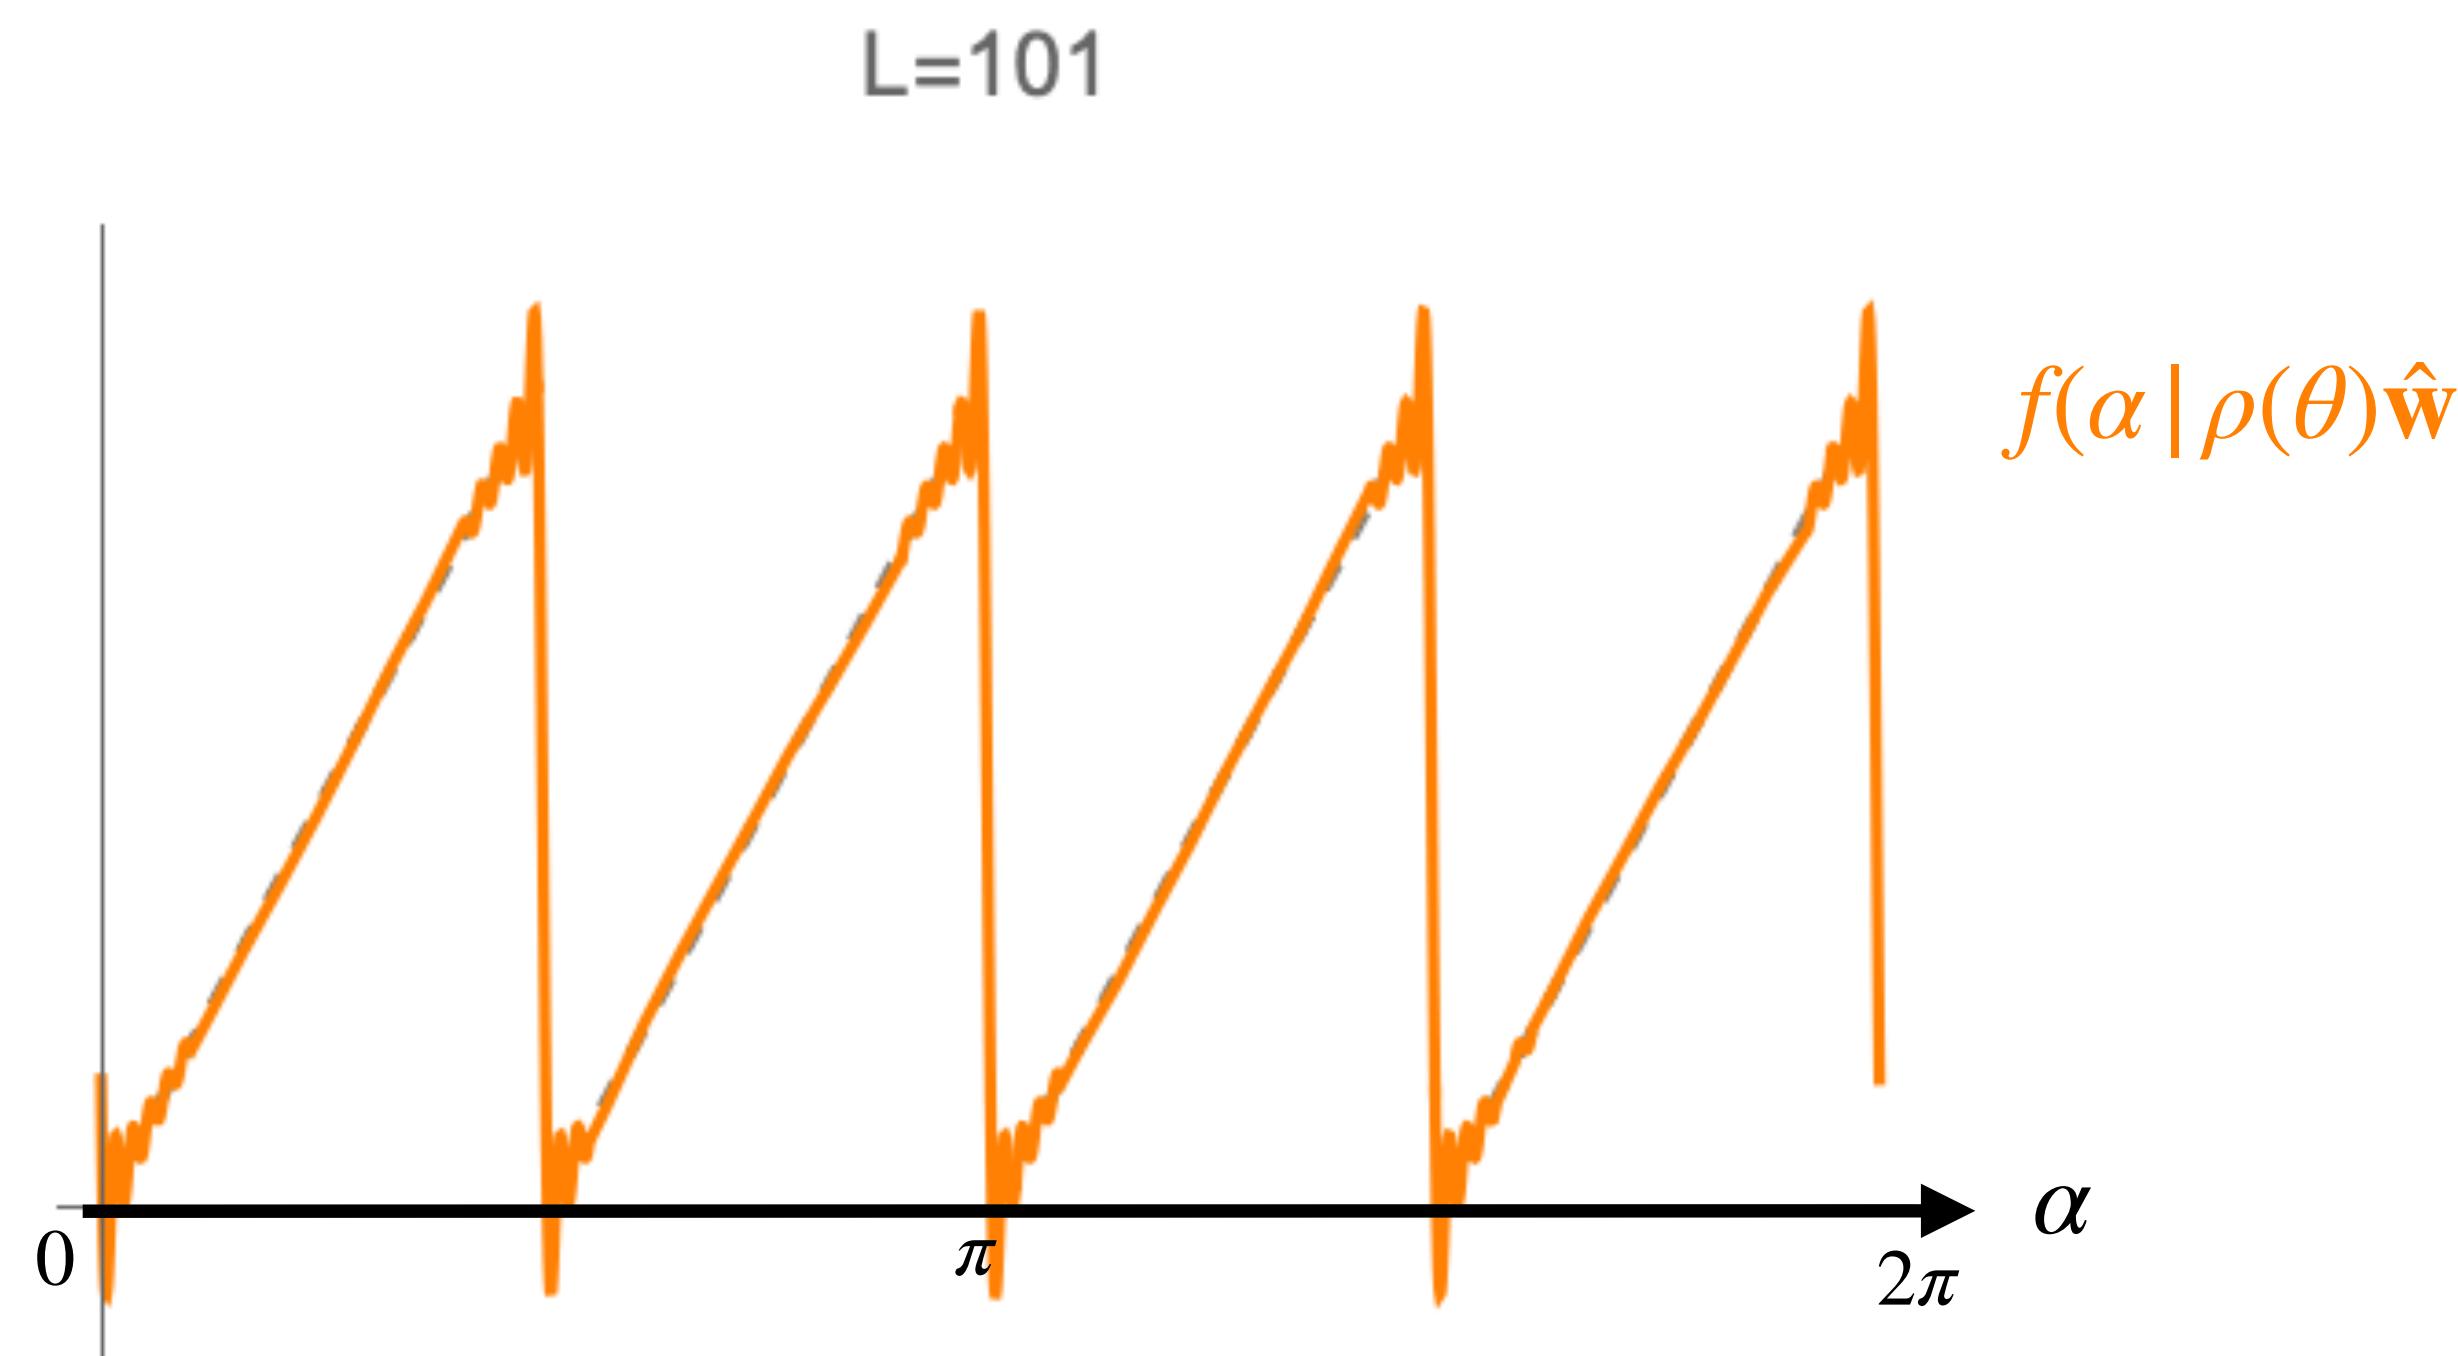
$$f(\alpha | \hat{\mathbf{w}}) = \hat{\mathbf{w}}^\dagger Y(\alpha)$$

Then we can **steer**/shift this function by transforming the weights  $\hat{\mathbf{w}}$

$$f(\alpha - \theta | \hat{\mathbf{w}}) = f(\alpha | \rho(\theta)\hat{\mathbf{w}})$$

Proof:

$$\begin{aligned} f(\alpha - \theta | \hat{\mathbf{w}}) &= \hat{\mathbf{w}}^\dagger Y(\alpha - \theta) \\ &= \hat{\mathbf{w}}^\dagger \rho(-\theta) Y(\alpha) \\ &= \hat{\mathbf{w}}^\dagger \rho(\theta)^\dagger Y(\alpha) \\ &= (\rho(\theta)\hat{\mathbf{w}})^\dagger Y(\alpha) \\ &= f(\alpha | \rho(\theta)\hat{\mathbf{w}}) \end{aligned}$$



# Example: Steerable basis on $S^1$ (circular harmonics)

Let

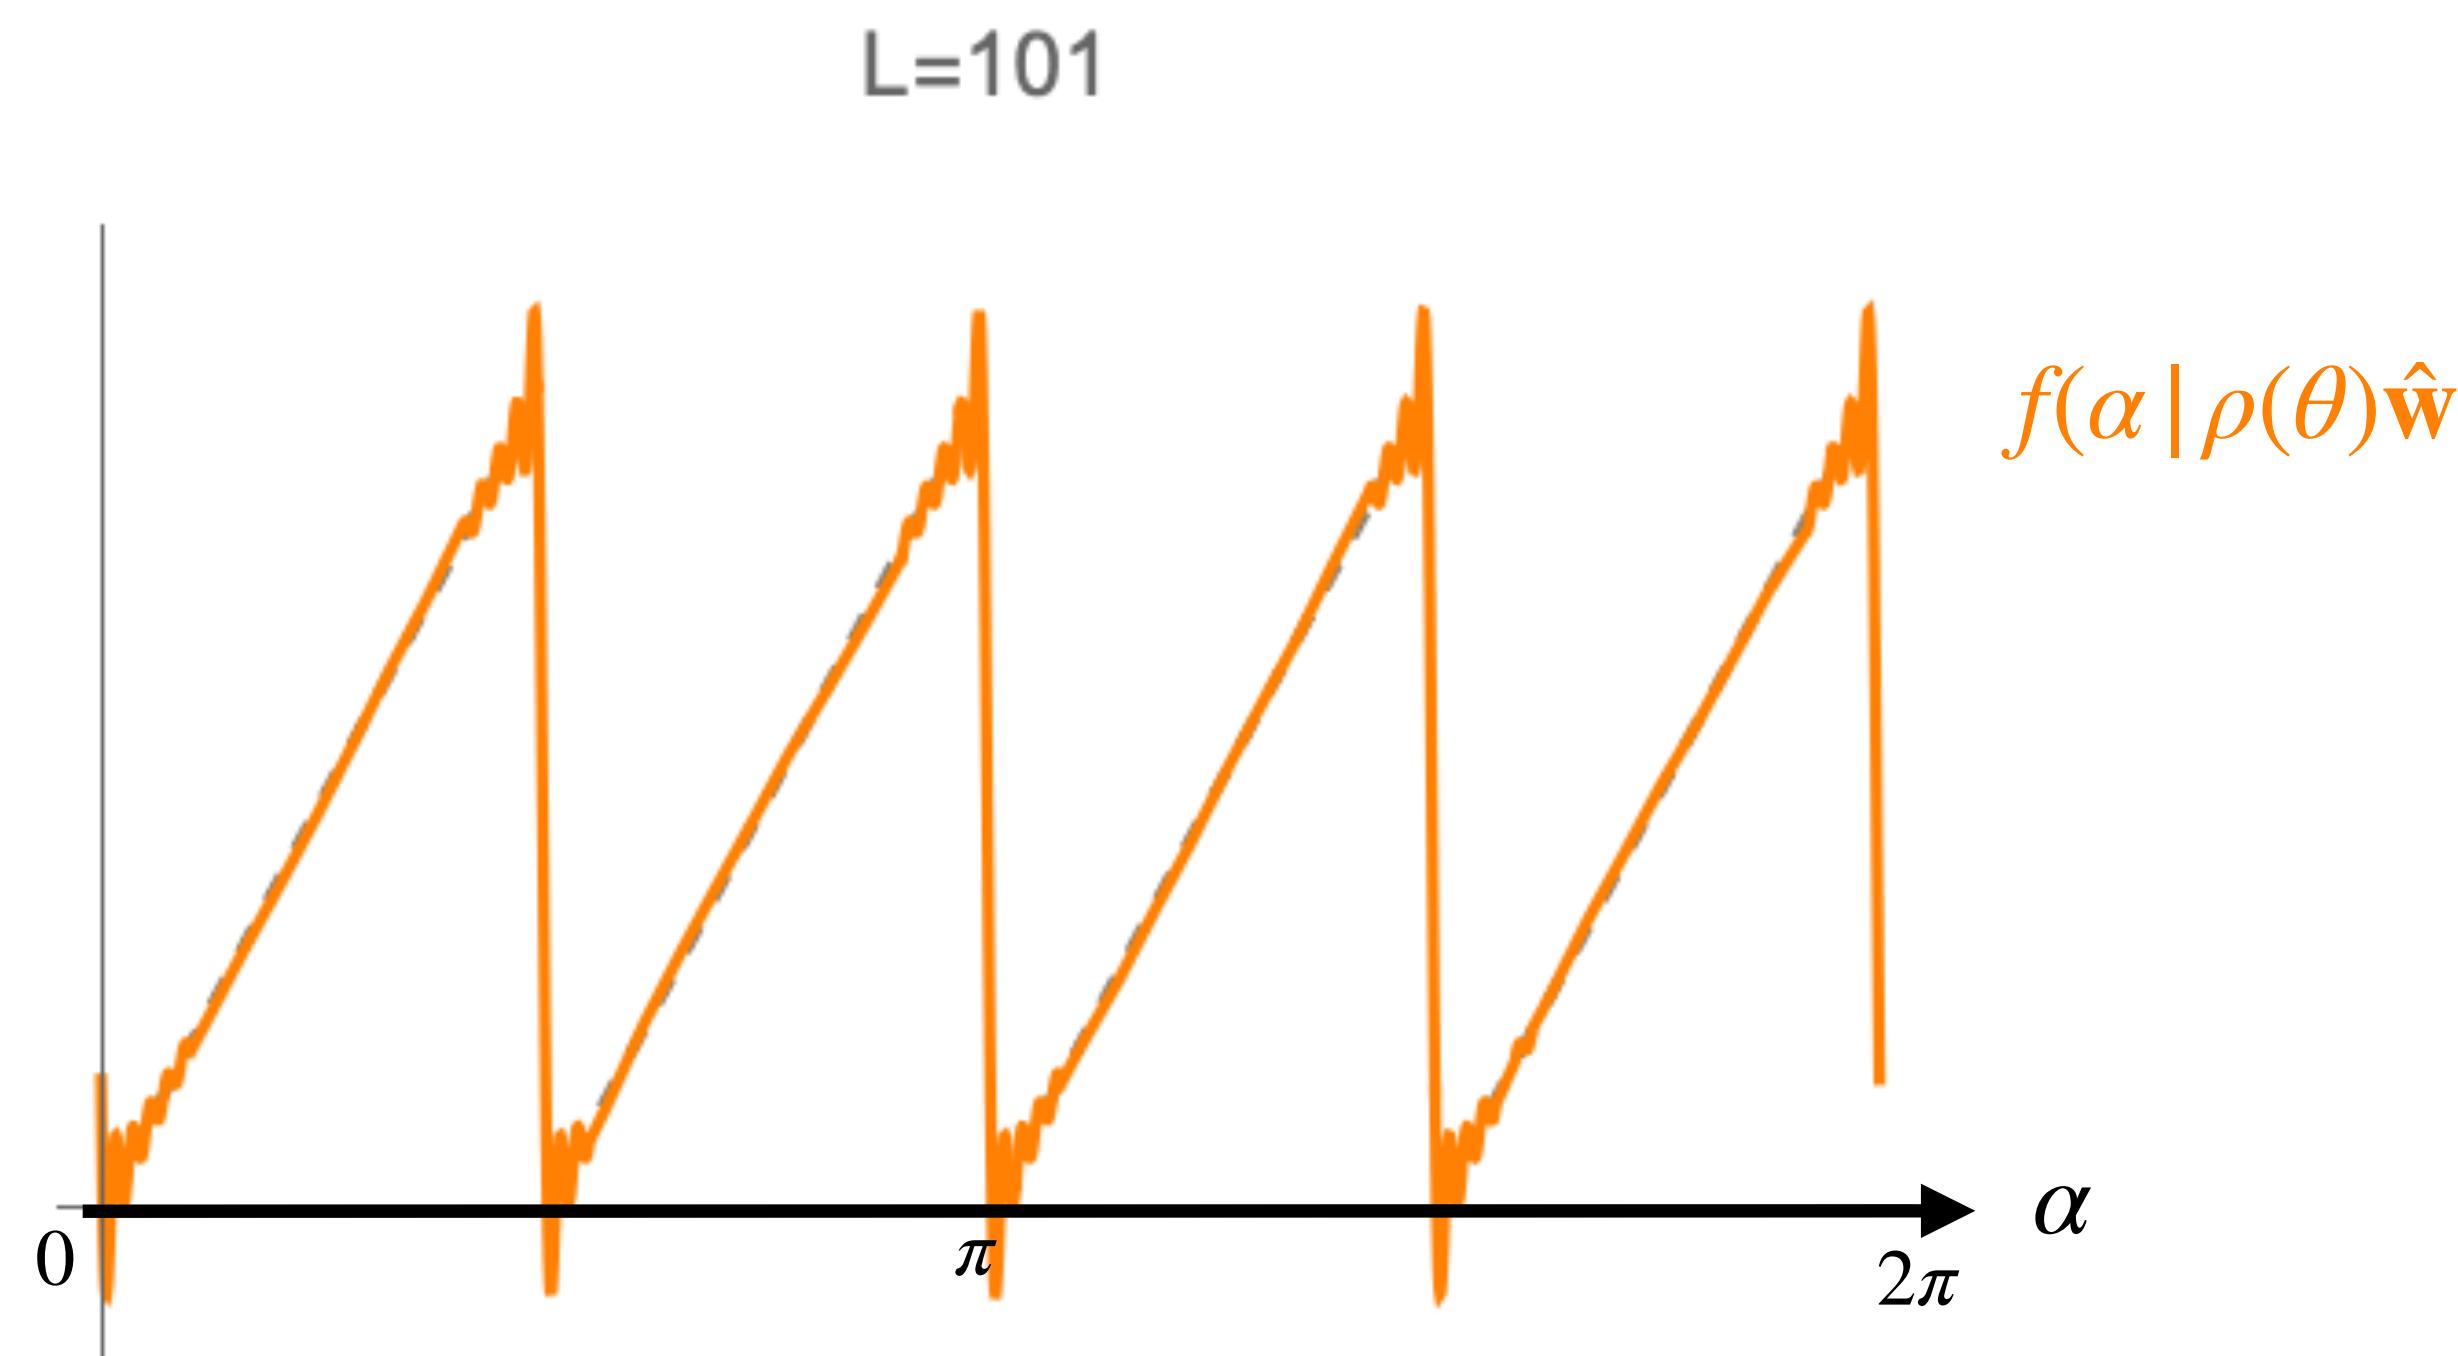
$$f(\alpha | \hat{\mathbf{w}}) = \hat{\mathbf{w}}^\dagger Y(\alpha)$$

Then we can **steer**/shift this function by transforming the weights  $\hat{\mathbf{w}}$

$$f(\alpha - \theta | \hat{\mathbf{w}}) = f(\alpha | \rho(\theta)\hat{\mathbf{w}})$$

Proof:

$$\begin{aligned} f(\alpha - \theta | \hat{\mathbf{w}}) &= \hat{\mathbf{w}}^\dagger Y(\alpha - \theta) \\ &= \hat{\mathbf{w}}^\dagger \rho(-\theta) Y(\alpha) \\ &= \hat{\mathbf{w}}^\dagger \rho(\theta)^\dagger Y(\alpha) \\ &= (\rho(\theta)\hat{\mathbf{w}})^\dagger Y(\alpha) \\ &= f(\alpha | \rho(\theta)\hat{\mathbf{w}}) \end{aligned}$$



# Two dimensional rotation-steerable functions

- The previous functions  $\rho_l(\theta) = e^{il\theta}$  are (irreducible) representations of  $SO(2)$

Recall lecture 1.6 (Group Theory | Homogeneous/quotient spaces)

## Transitive action

**Transitive action:** An action  $\odot : G \times X \rightarrow X$  such that

$$\forall_{x_0, x \in X} \exists_{g \in G} : x = g \odot x_0$$

$(\mathbb{R}^2, +)$  acts transitively on  $\mathbb{R}^2$

$SE(2)$  acts transitively on  $\mathbb{R}^2$

$SO(2)$  does not ...

6

- Though not transitively...
- It does act transitively on  $S^1$  though
- Use polar coordinates  $\mathbb{R}^2 \ni \mathbf{x} \leftrightarrow (r, \alpha) \in \mathbb{R}^+ \times S^1$  to come up with a rotation-steerable basis for  $\mathbb{L}_2(\mathbb{R}^2)$ !

# Two dimensional rotation-steerable functions

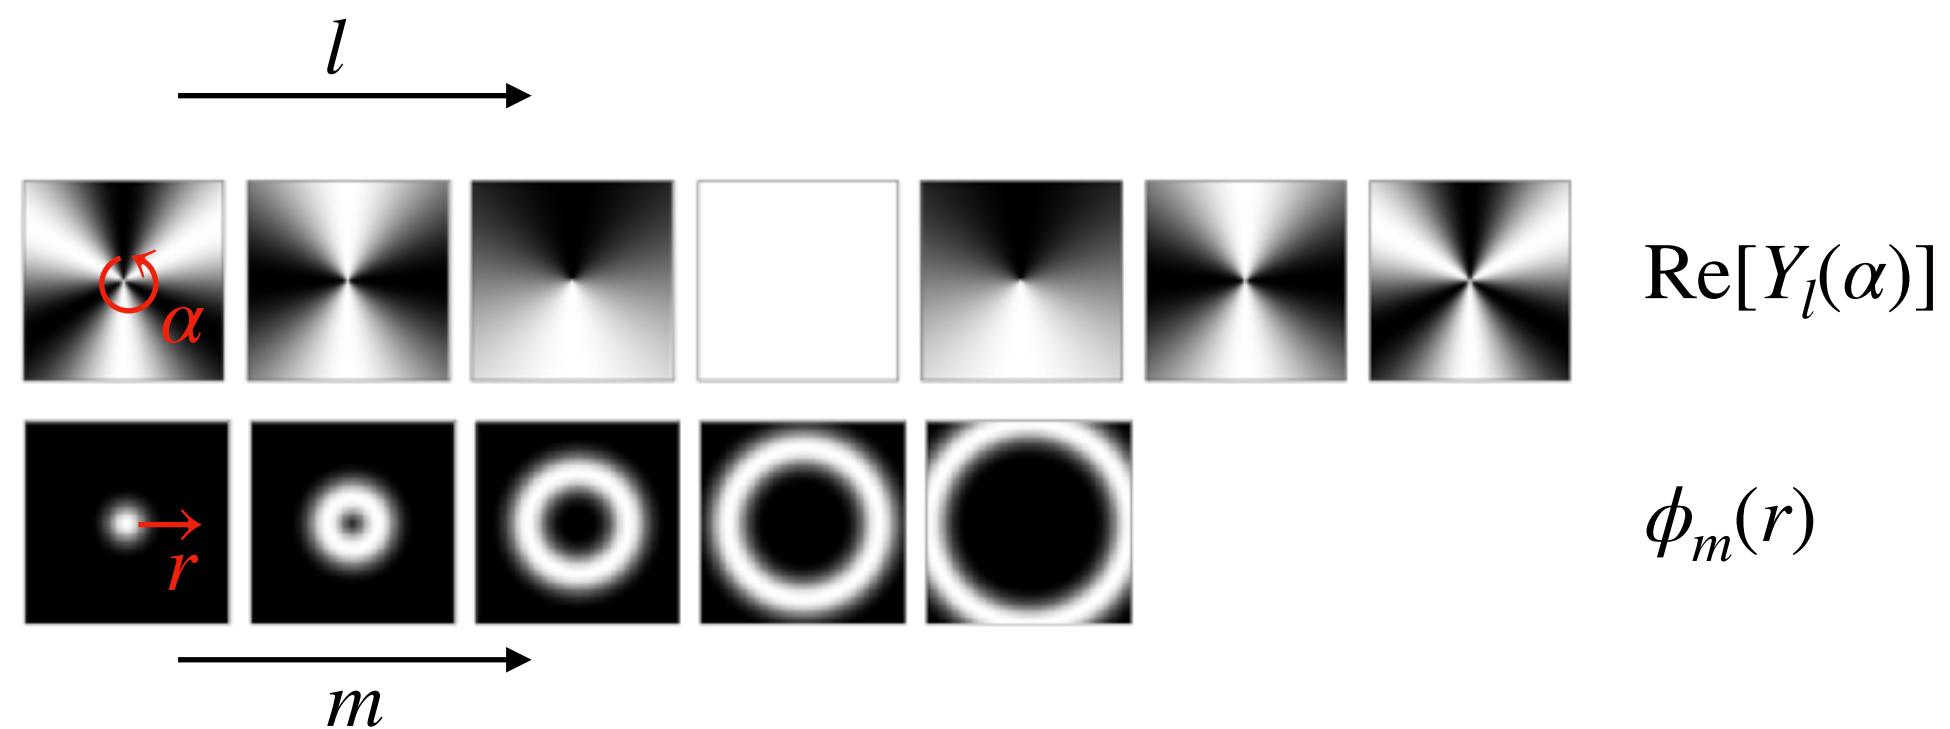
- Consider polar-separable convolution kernel:

$$k(\mathbf{x} | \mathbf{w}) = k^\rightarrow(r | \mathbf{w}) k^\circlearrowleft(\alpha | \mathbf{w}),$$

- with  $k^\circlearrowleft$  in an  $SO(2)$  steerable basis, and  $k^\rightarrow$  in some radial basis:

$$k^\circlearrowleft(\alpha | \mathbf{w}) = \sum_l \bar{w}_l Y_l(\alpha), \quad \text{e.g., with} \quad Y_l(\alpha) = e^{il\alpha},$$

$$k^\rightarrow(r | \mathbf{w}) = \sum_m w_m \phi_m(r)$$



- Then we may as well write it as

$$\begin{aligned} k(\mathbf{x} | \mathbf{w}) &= \sum_l \sum_m w_m \bar{w}_l \phi_m(r) Y_l(\alpha) \\ &= \sum_l \sum_m \bar{w}_{ml} \phi_m(r) Y_l(\alpha) \quad (\text{"absorb" weights}) \\ &= \sum_l \hat{w}_l(r) Y_l(\alpha) \end{aligned}$$

with radius dependent weights  $\hat{w}_l(r) = \sum_m w_{ml} \phi_m(r)$

- Then such kernel is clearly rotation steerable!

$$k(\mathbf{R}_\theta^{-1}\mathbf{x} | \hat{\mathbf{w}}(r)) = k(\mathbf{x} | \rho(\theta)\hat{\mathbf{w}}(r))$$

# Two dimensional rotation-steerable functions

- Consider polar-separable convolution kernel:

$$k(\mathbf{x} \mid \mathbf{w}) = k^\rightarrow(r \mid \mathbf{w}) k^\circlearrowleft(\alpha \mid \mathbf{w}),$$

- with  $k^\circlearrowleft$  in an  $SO(2)$  steerable basis, and  $k^\rightarrow$  in some radial basis:

$$k^\circlearrowleft(\alpha \mid \mathbf{w}) = \sum_l \bar{w}_l Y_l(\alpha), \quad \text{e.g., with} \quad Y_l(\alpha) = e^{il\alpha},$$

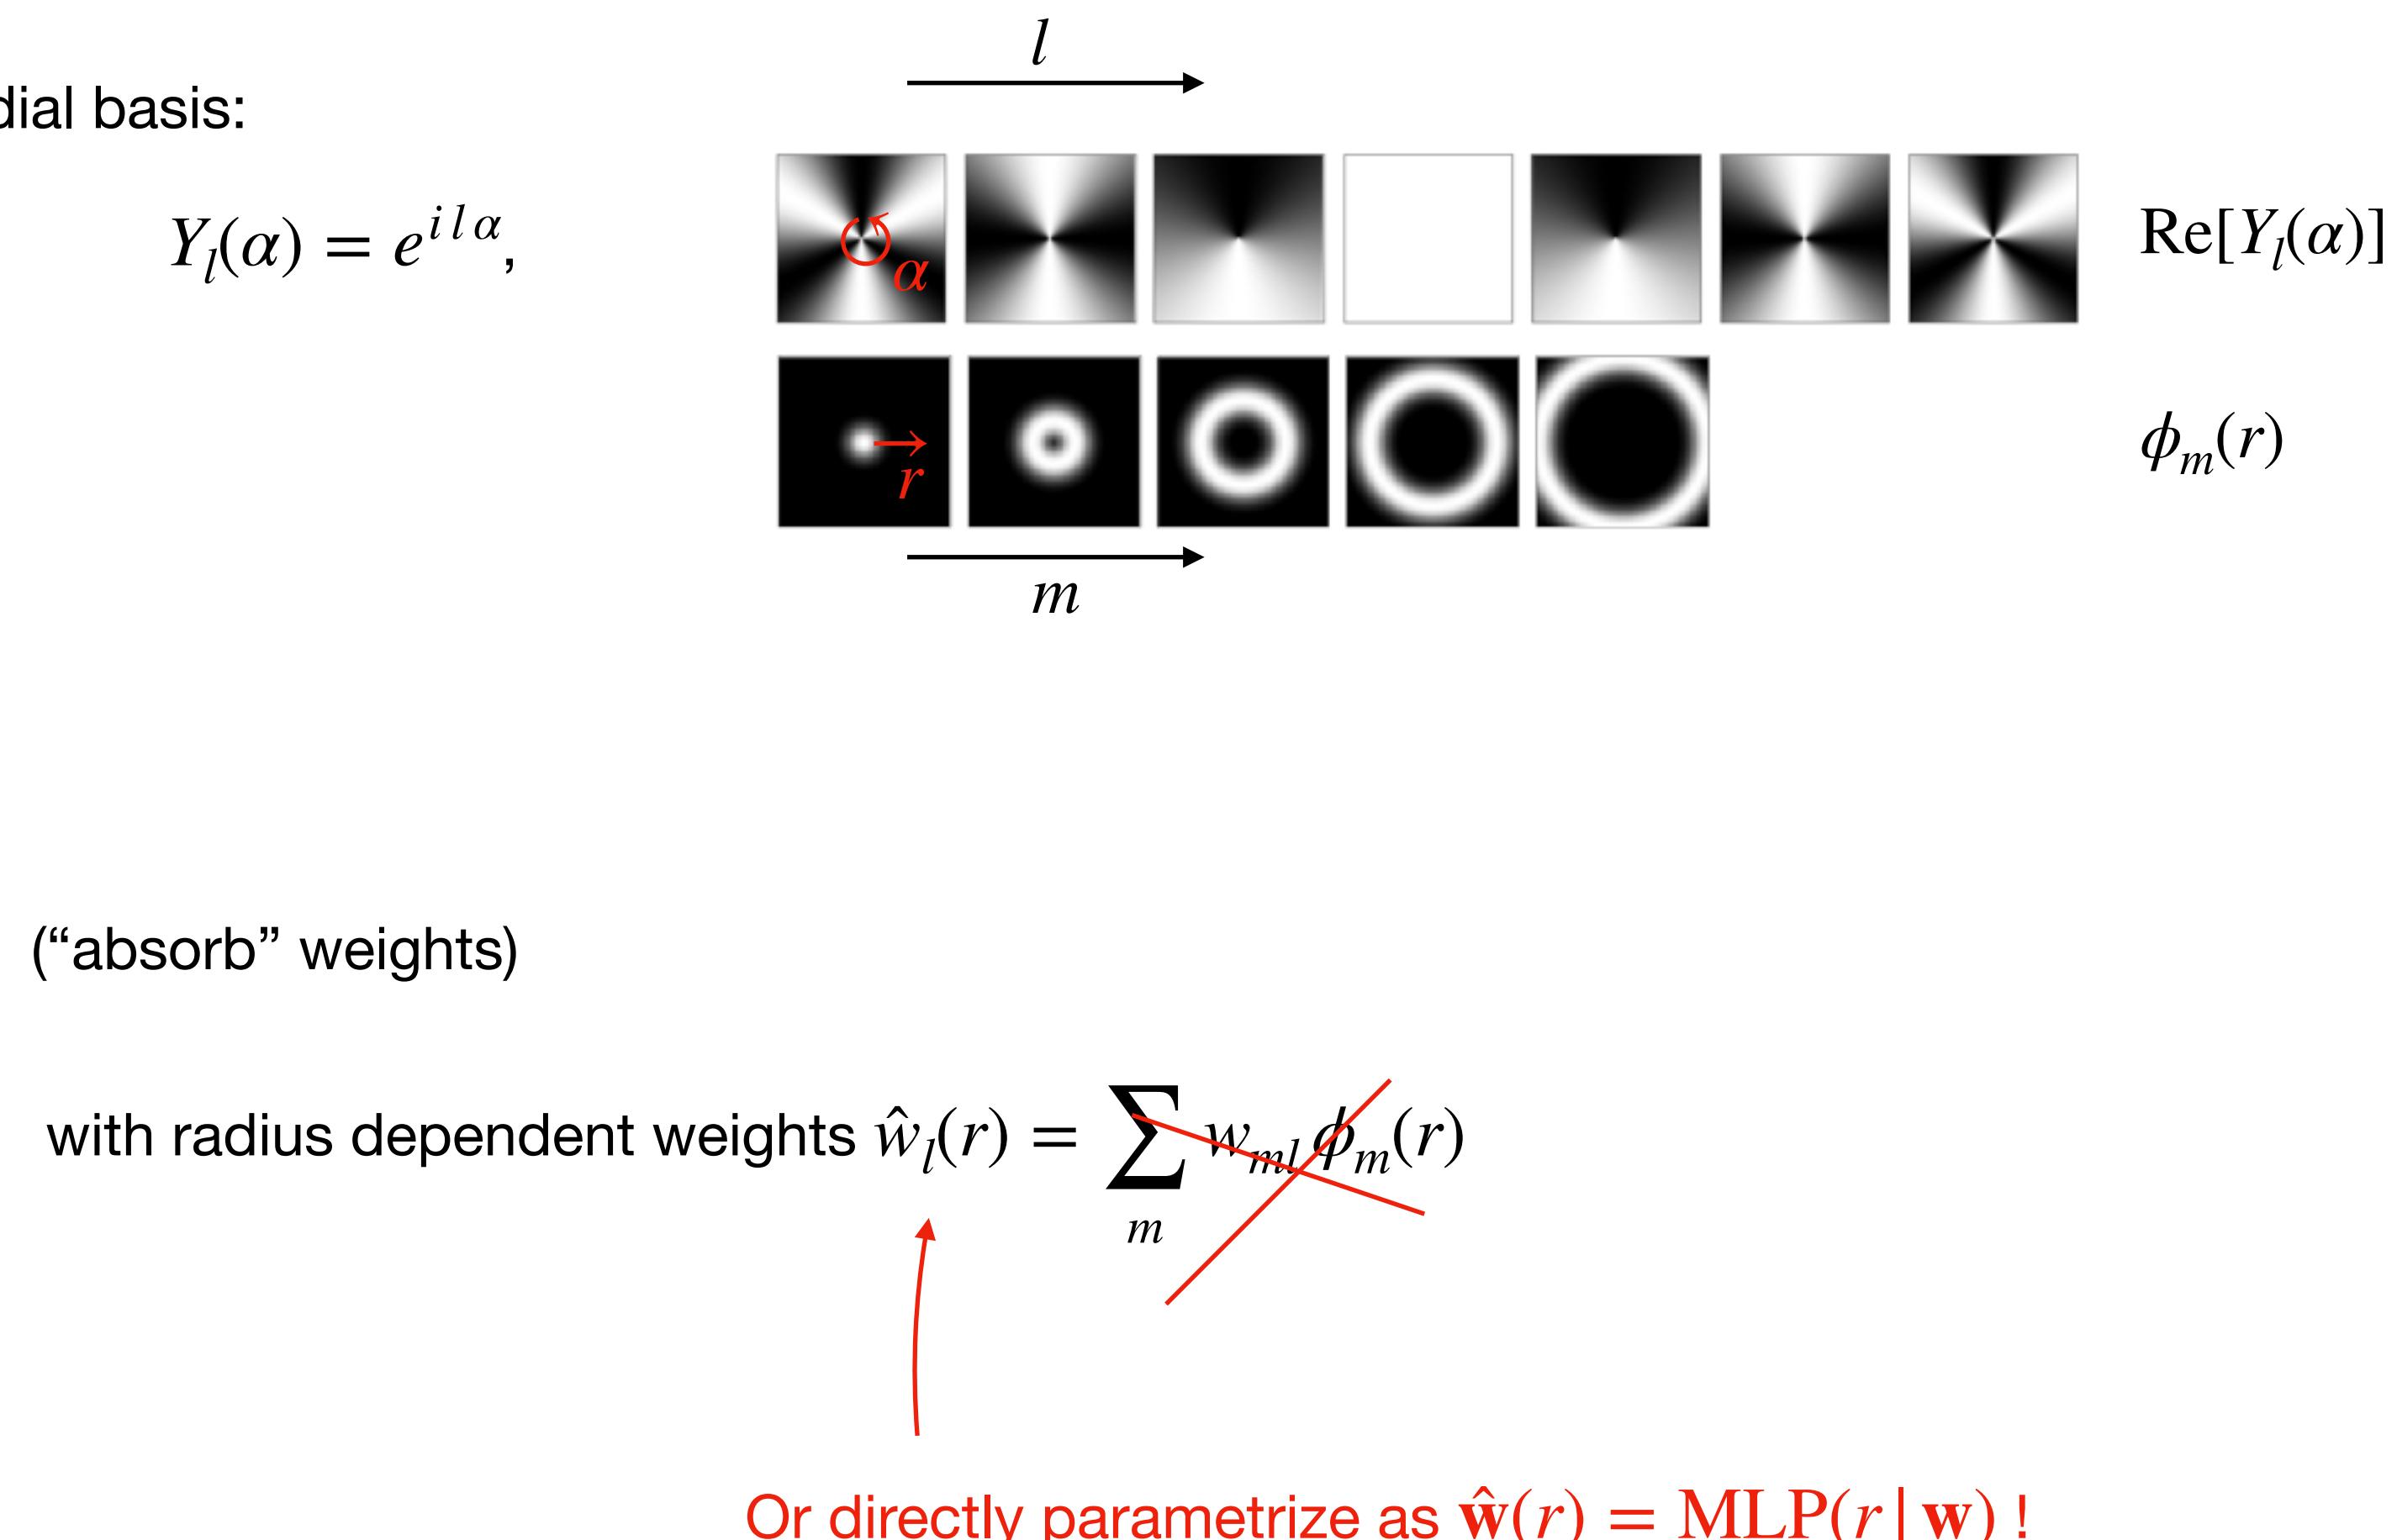
$$k^\rightarrow(r \mid \mathbf{w}) = \sum_m w_m \phi_m(r)$$

- Then we may as well write it as

$$\begin{aligned} k(\mathbf{x} \mid \mathbf{w}) &= \sum_l \sum_m w_m \bar{w}_l \phi_m(r) Y_l(\alpha) \\ &= \sum_l \sum_m \bar{w}_{ml} \phi_m(r) Y_l(\alpha) \quad (\text{"absorb" weights}) \\ &= \sum_l \hat{w}_l(r) Y_l(\alpha) \end{aligned}$$

- Then such kernel is clearly rotation steerable!

$$k(\mathbf{R}_\theta^{-1}\mathbf{x} \mid \hat{\mathbf{w}}(r)) = k(\mathbf{x} \mid \rho(\theta)\hat{\mathbf{w}}(r))$$



# Complex (irreducible) representations

$$\begin{array}{ccc}
 Y(\mathbf{R}_\theta^{-1} \mathbf{x}) & = & \rho(\mathbf{R}_\theta^{-1}) \\
 \begin{array}{cc}
 \text{Re} & \text{Im} \\
 \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) & \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) \\
 \longrightarrow & \longrightarrow
 \end{array} & = & \begin{pmatrix} e^{3i\theta} & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & e^{2i\theta} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & e^{1i\theta} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & e^{-1i\theta} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & e^{-2i\theta} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & e^{-3i\theta} \end{pmatrix} \\
 \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) & & \begin{array}{cc}
 \text{Re} & \text{Im} \\
 \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) & \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) \\
 \longrightarrow & \longrightarrow
 \end{array}
 \end{array}
 \end{array}$$

# Complex (irreducible) representations

$$\begin{array}{ccc}
 Y(\mathbf{R}_\theta^{-1} \mathbf{x}) & = & \rho(\mathbf{R}_\theta^{-1}) \\
 \begin{array}{cc}
 \text{Re} & \text{Im} \\
 \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) & \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) \\
 \longrightarrow & \longrightarrow
 \end{array} & = & \begin{pmatrix} e^{3i\theta} & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & e^{2i\theta} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & e^{1i\theta} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & e^{-1i\theta} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & e^{-2i\theta} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & e^{-3i\theta} \end{pmatrix} \\
 \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) & & \begin{array}{cc}
 \text{Re} & \text{Im} \\
 \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) & \left( \begin{array}{c} \text{Re} \\ \text{Im} \\ \vdots \\ \text{Re} \\ \text{Im} \end{array} \right) \\
 \longrightarrow & \longrightarrow
 \end{array}
 \end{array}
 \end{array}$$

# Complex (irreducible) representations

$$\begin{array}{ccc}
 Y(\mathbf{R}_\theta^{-1} \mathbf{x}) & = & \rho(\mathbf{R}_\theta^{-1}) \\
 \left( \begin{array}{cc} \text{Re} & \text{Im} \\ \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} & \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} \end{array} \right) & = & \left( \begin{array}{ccccccc} e^{3i\theta} & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & e^{2i\theta} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & e^{1i\theta} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & e^{-1i\theta} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & e^{-2i\theta} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & e^{-3i\theta} \end{array} \right) \\
 \left( \begin{array}{cc} \text{Re} & \text{Im} \\ \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} & \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} \end{array} \right) & = & \left( \begin{array}{cc} \cos(l\alpha) & \sin(l\alpha) \\ \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} & \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} \end{array} \right)
 \end{array}$$

The diagram illustrates the decomposition of a complex representation  $Y(\mathbf{R}_\theta^{-1} \mathbf{x})$  into its real and imaginary parts. On the left, two 4x4 matrices are shown: one for the real part (labeled "Re") and one for the imaginary part (labeled "Im"). Arrows point from these matrices to a central 4x4 matrix representing the rotation  $\rho(\mathbf{R}_\theta^{-1})$ . This central matrix is a diagonal block matrix with entries  $e^{3i\theta}$ ,  $e^{2i\theta}$ ,  $e^{1i\theta}$ , 1,  $e^{-1i\theta}$ ,  $e^{-2i\theta}$ , and  $e^{-3i\theta}$ . To the right, the decomposition is shown as a product of the rotation matrix and two 4x4 matrices representing the real and imaginary components, labeled "Re" and "Im". The "Re" matrix contains the cosine and sine terms, while the "Im" matrix contains the remaining terms. A red box highlights the cosine term  $\cos(l\alpha)$ .

# Real (irreducible) representations

$$Y(\mathbf{R}_\theta^{-1} \mathbf{x}) = \rho(\mathbf{R}_\theta^{-1}) Y(\mathbf{x})$$

$$\begin{pmatrix} \text{Image 1} \\ \text{Image 2} \\ \text{Image 3} \\ \text{Image 4} \\ \text{Image 5} \\ \text{Image 6} \\ \text{Image 7} \\ \text{Image 8} \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & \cos \theta & \sin \theta & 0 & 0 & 0 & 0 \\ 0 & -\sin \theta & \cos \theta & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & \cos 2\theta & \sin 2\theta & 0 & 0 \\ 0 & 0 & 0 & -\sin 2\theta & \cos 2\theta & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & \cos 3\theta & \sin 3\theta \\ 0 & 0 & 0 & 0 & -\sin 3\theta & \cos 3\theta \end{pmatrix} \begin{pmatrix} \text{Image 1} \\ \text{Image 2} \\ \text{Image 3} \\ \text{Image 4} \\ \text{Image 5} \\ \text{Image 6} \\ \text{Image 7} \\ \text{Image 8} \end{pmatrix}$$

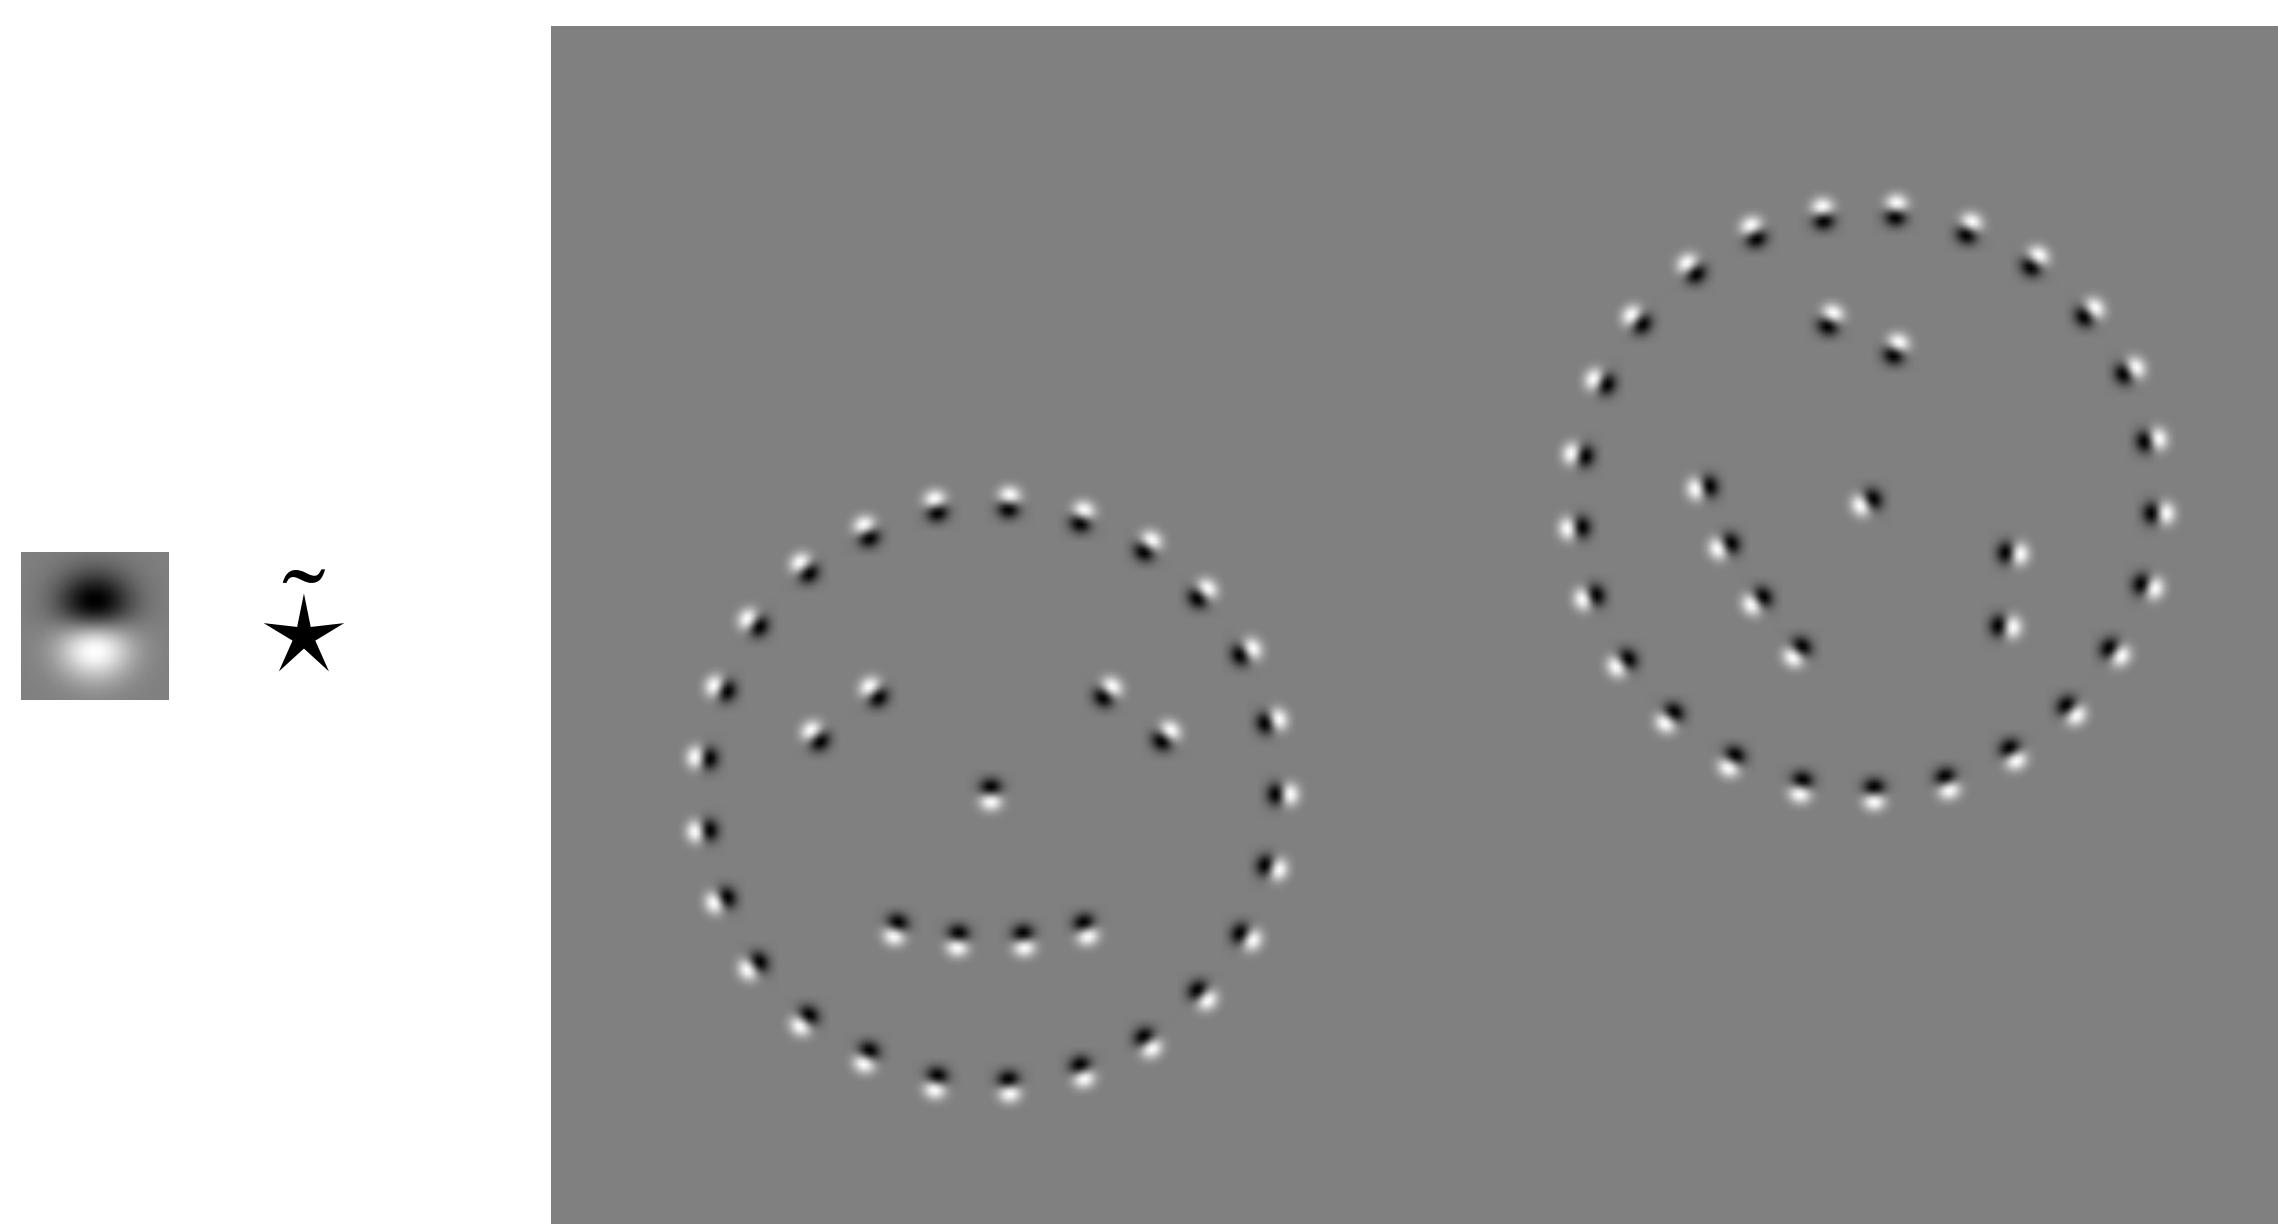
# Real (irreducible) representations

$$Y(\mathbf{R}_\theta^{-1} \mathbf{x}) = \rho(\mathbf{R}_\theta^{-1}) Y(\mathbf{x})$$
$$\left( \begin{array}{c} \text{Image 1} \\ \text{Image 2} \\ \text{Image 3} \\ \text{Image 4} \\ \text{Image 5} \\ \text{Image 6} \\ \text{Image 7} \\ \text{Image 8} \end{array} \right) = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & \cos \theta & \sin \theta & 0 & 0 & 0 & 0 \\ 0 & -\sin \theta & \cos \theta & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & \cos 2\theta & \sin 2\theta & 0 & 0 \\ 0 & 0 & 0 & -\sin 2\theta & \cos 2\theta & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & \cos 3\theta & \sin 3\theta \\ 0 & 0 & 0 & 0 & -\sin 3\theta & \cos 3\theta \end{pmatrix} \left( \begin{array}{c} \text{Image 1} \\ \text{Image 2} \\ \text{Image 3} \\ \text{Image 4} \\ \text{Image 5} \\ \text{Image 6} \\ \text{Image 7} \\ \text{Image 8} \end{array} \right)$$

# Regular group convolutions revisited

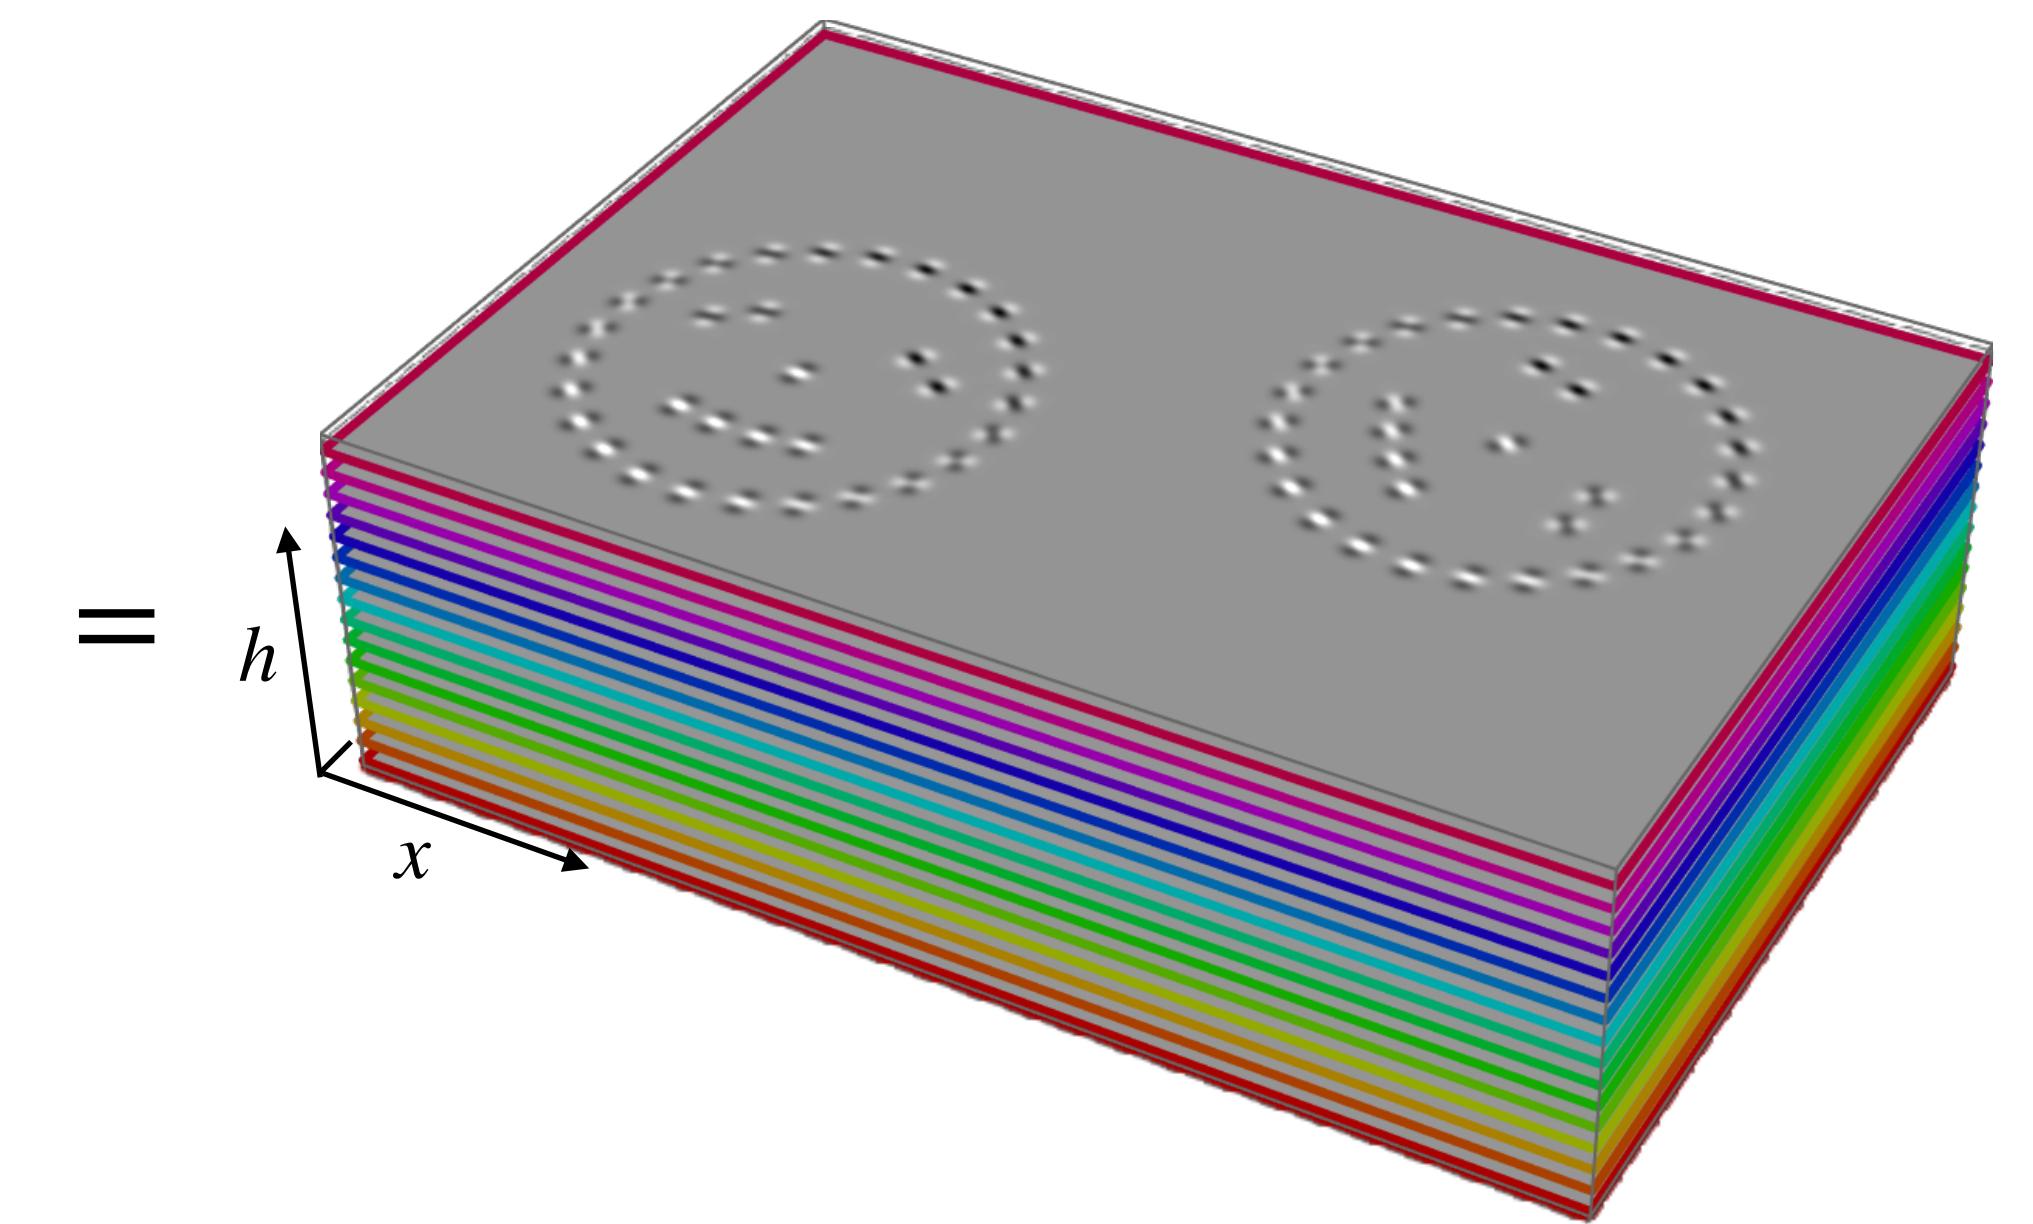
Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)}$   
(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

---



2D convolution kernel

2D input feature map



$SE(2)$  output feature map

# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

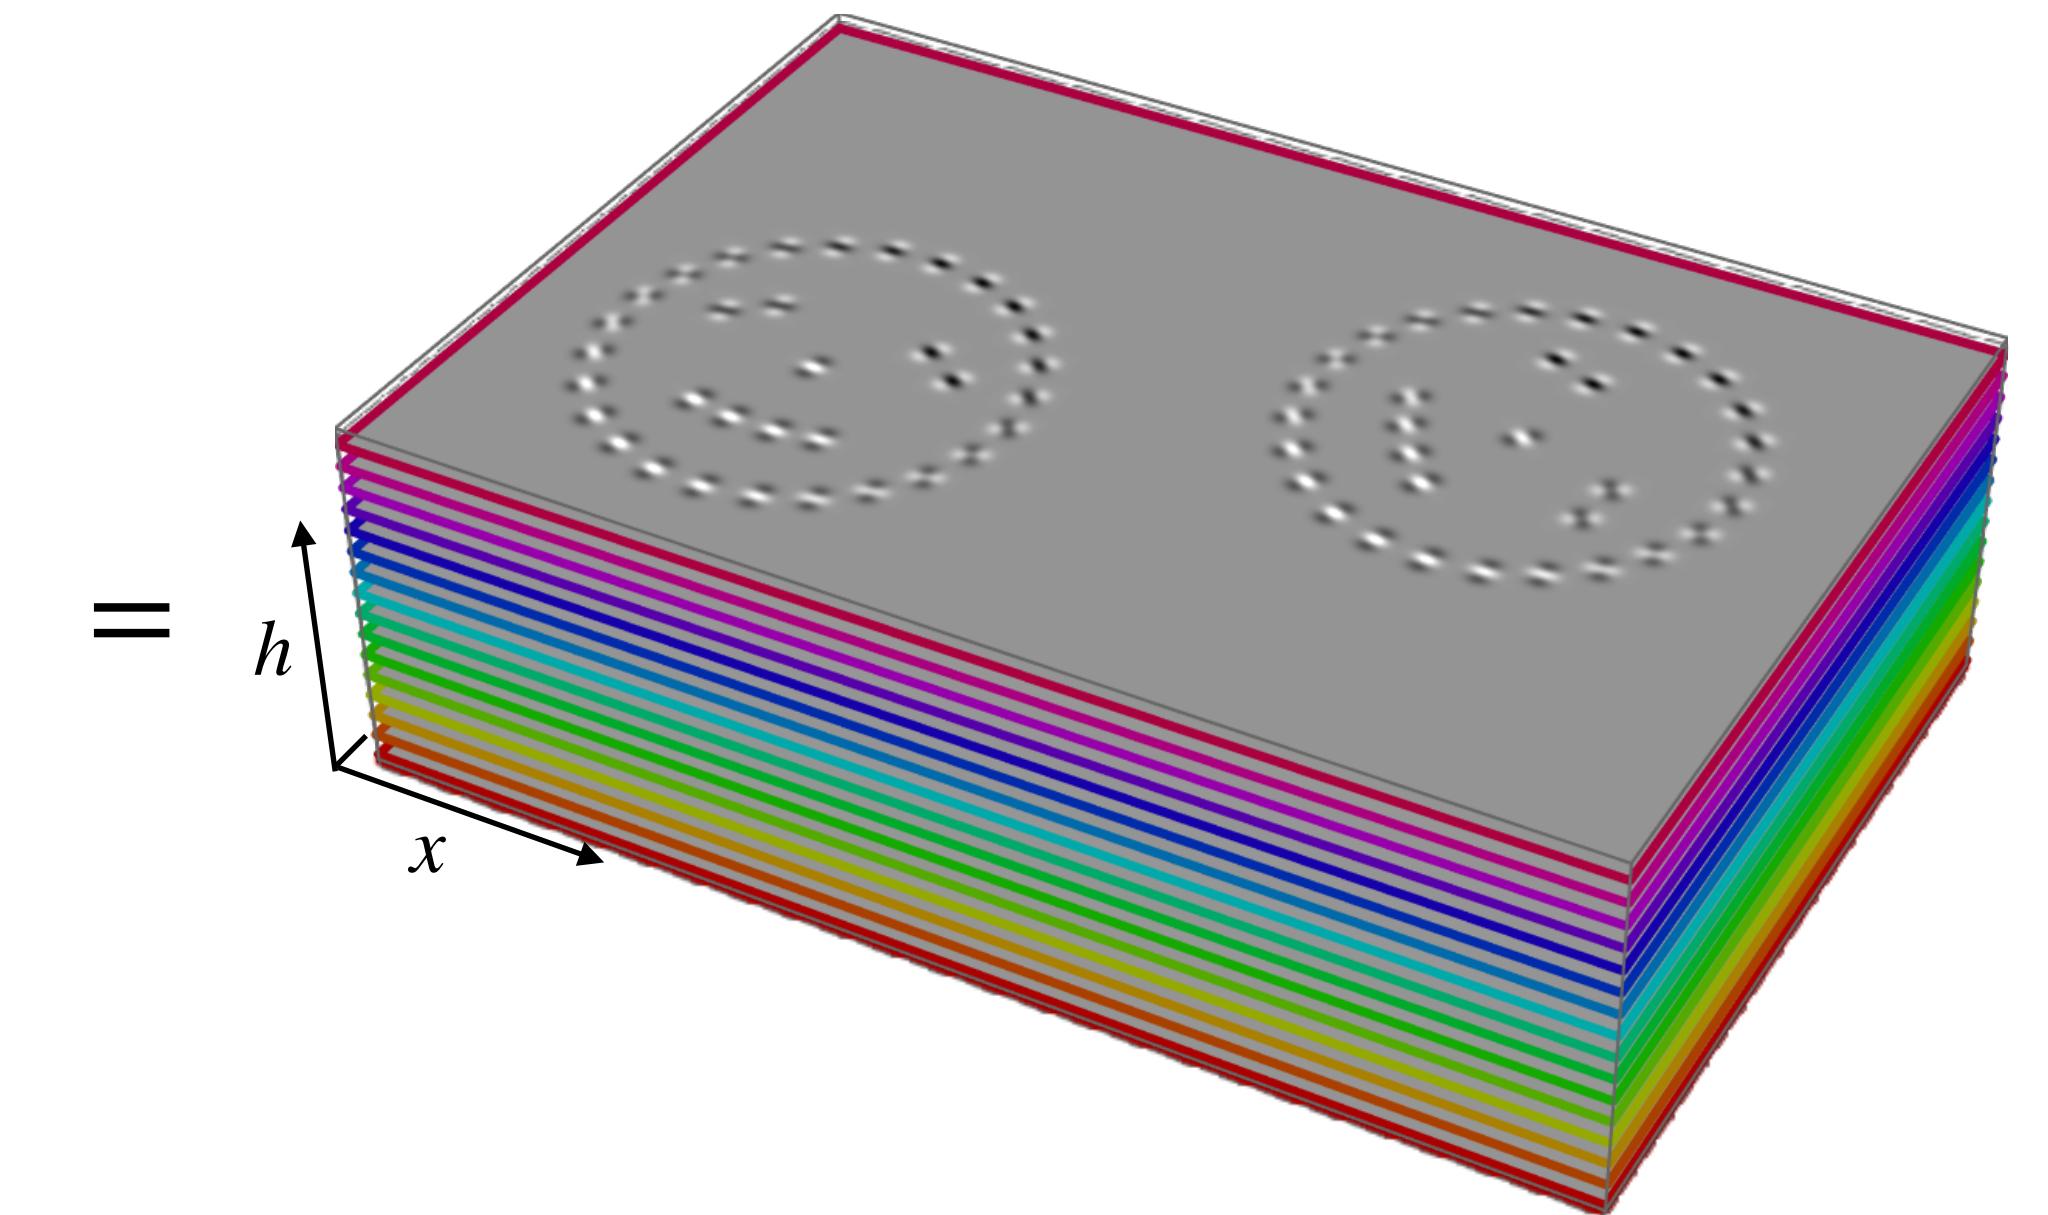
$$(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)}$$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}')f(\mathbf{x}')d\mathbf{x}'$$



2D convolution kernel

2D input feature map



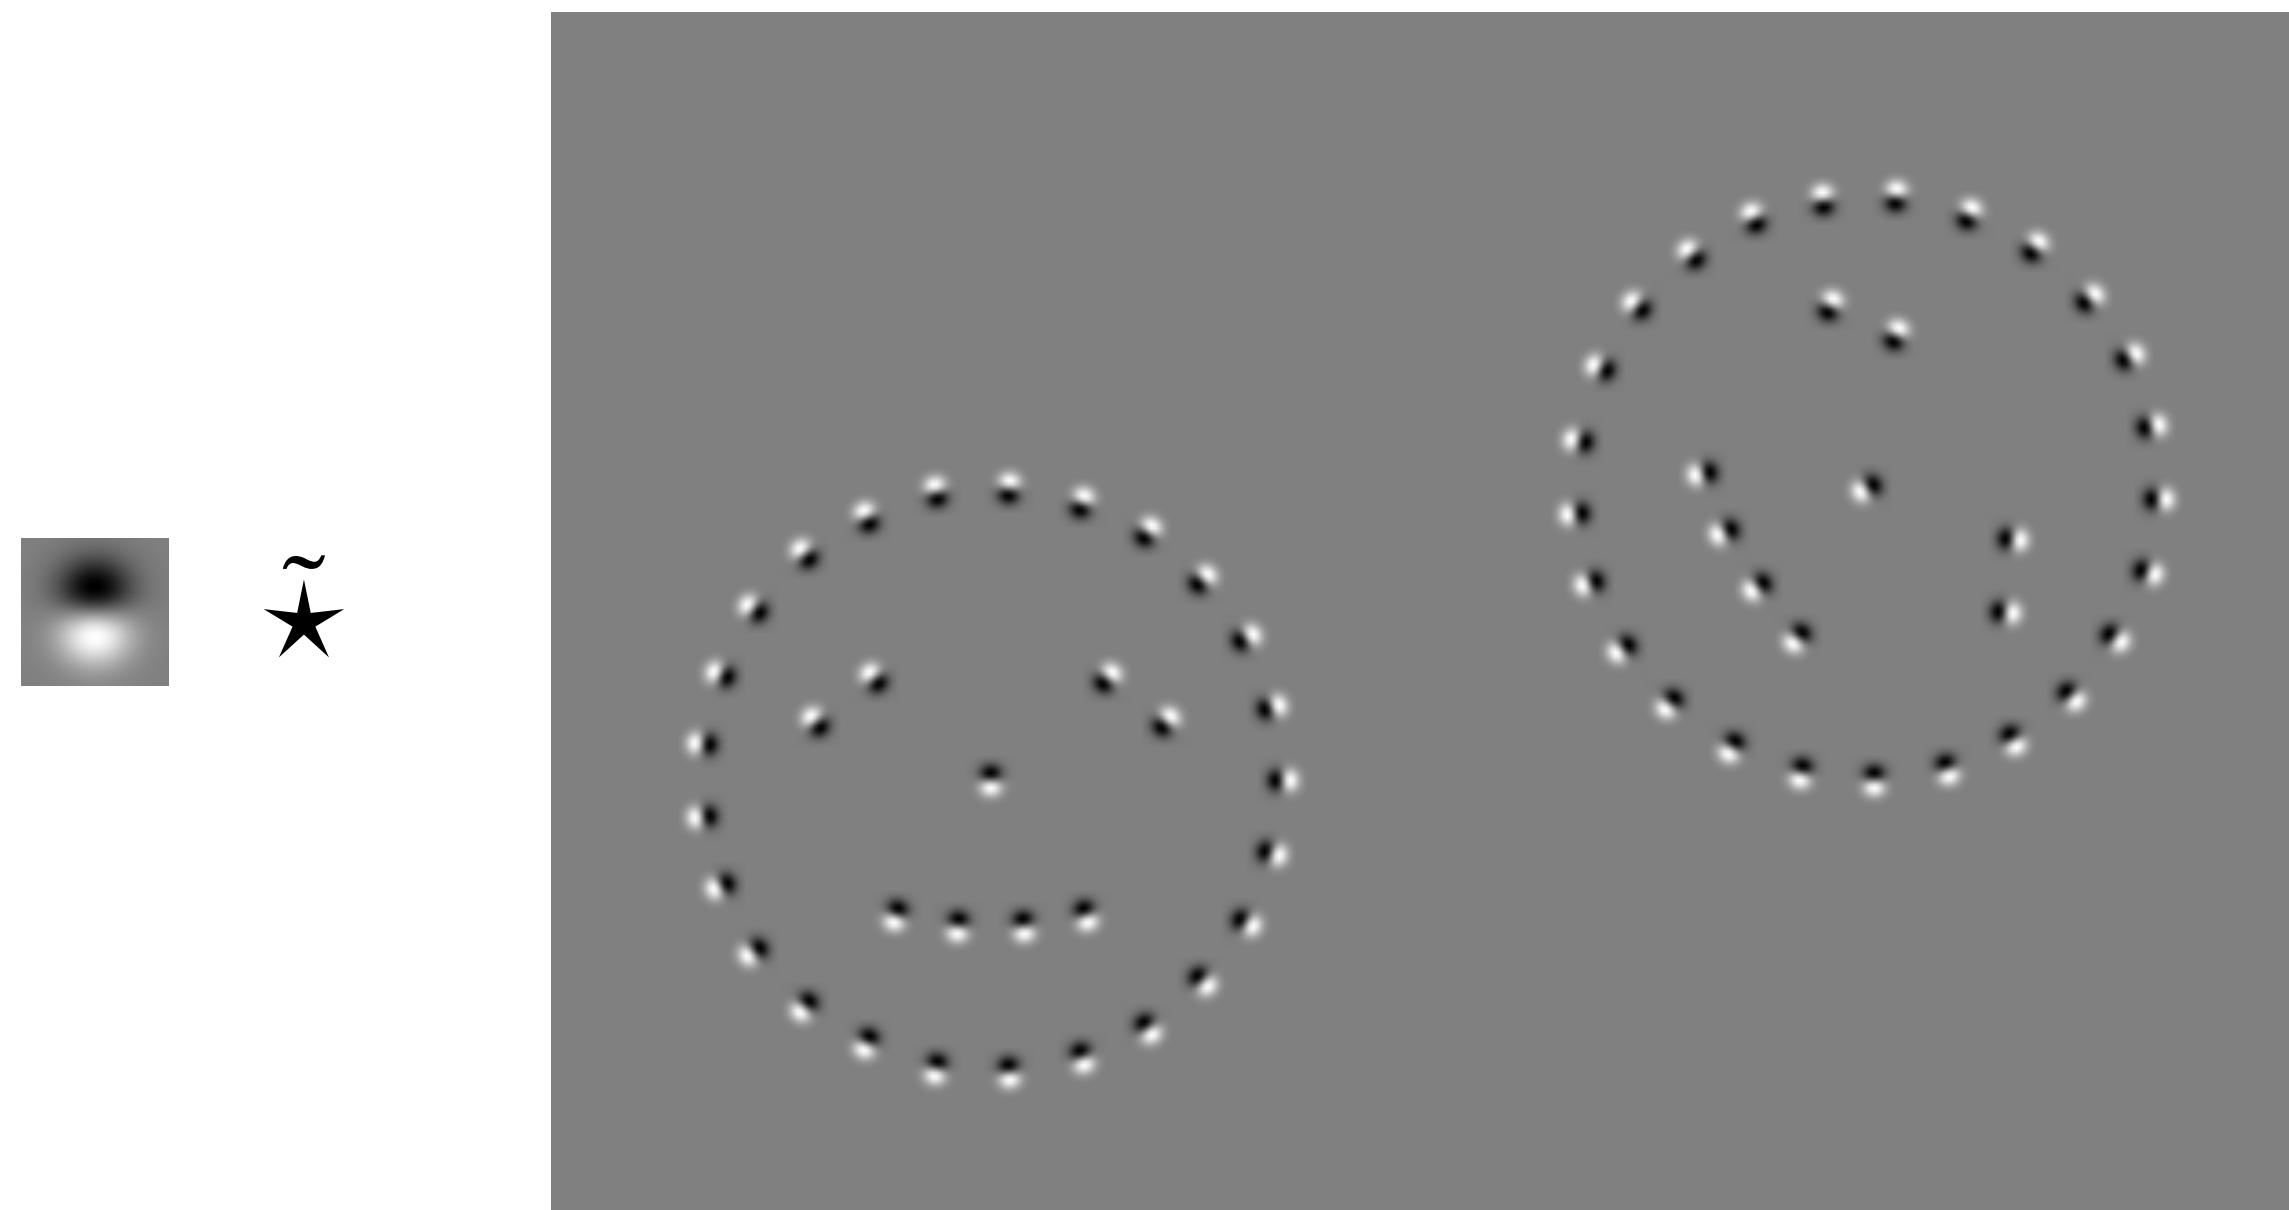
$SE(2)$  output feature map

# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)} = (\mathcal{L}_{\mathbf{x}}^{(\mathbb{R}^d,+)\rightarrow \mathbb{L}_2(\mathbb{R}^d)} \mathcal{L}_h^{H \rightarrow \mathbb{L}_2(\mathbb{R}^d)} k, f)_{\mathbb{L}_2(\mathbb{R}^d)}$

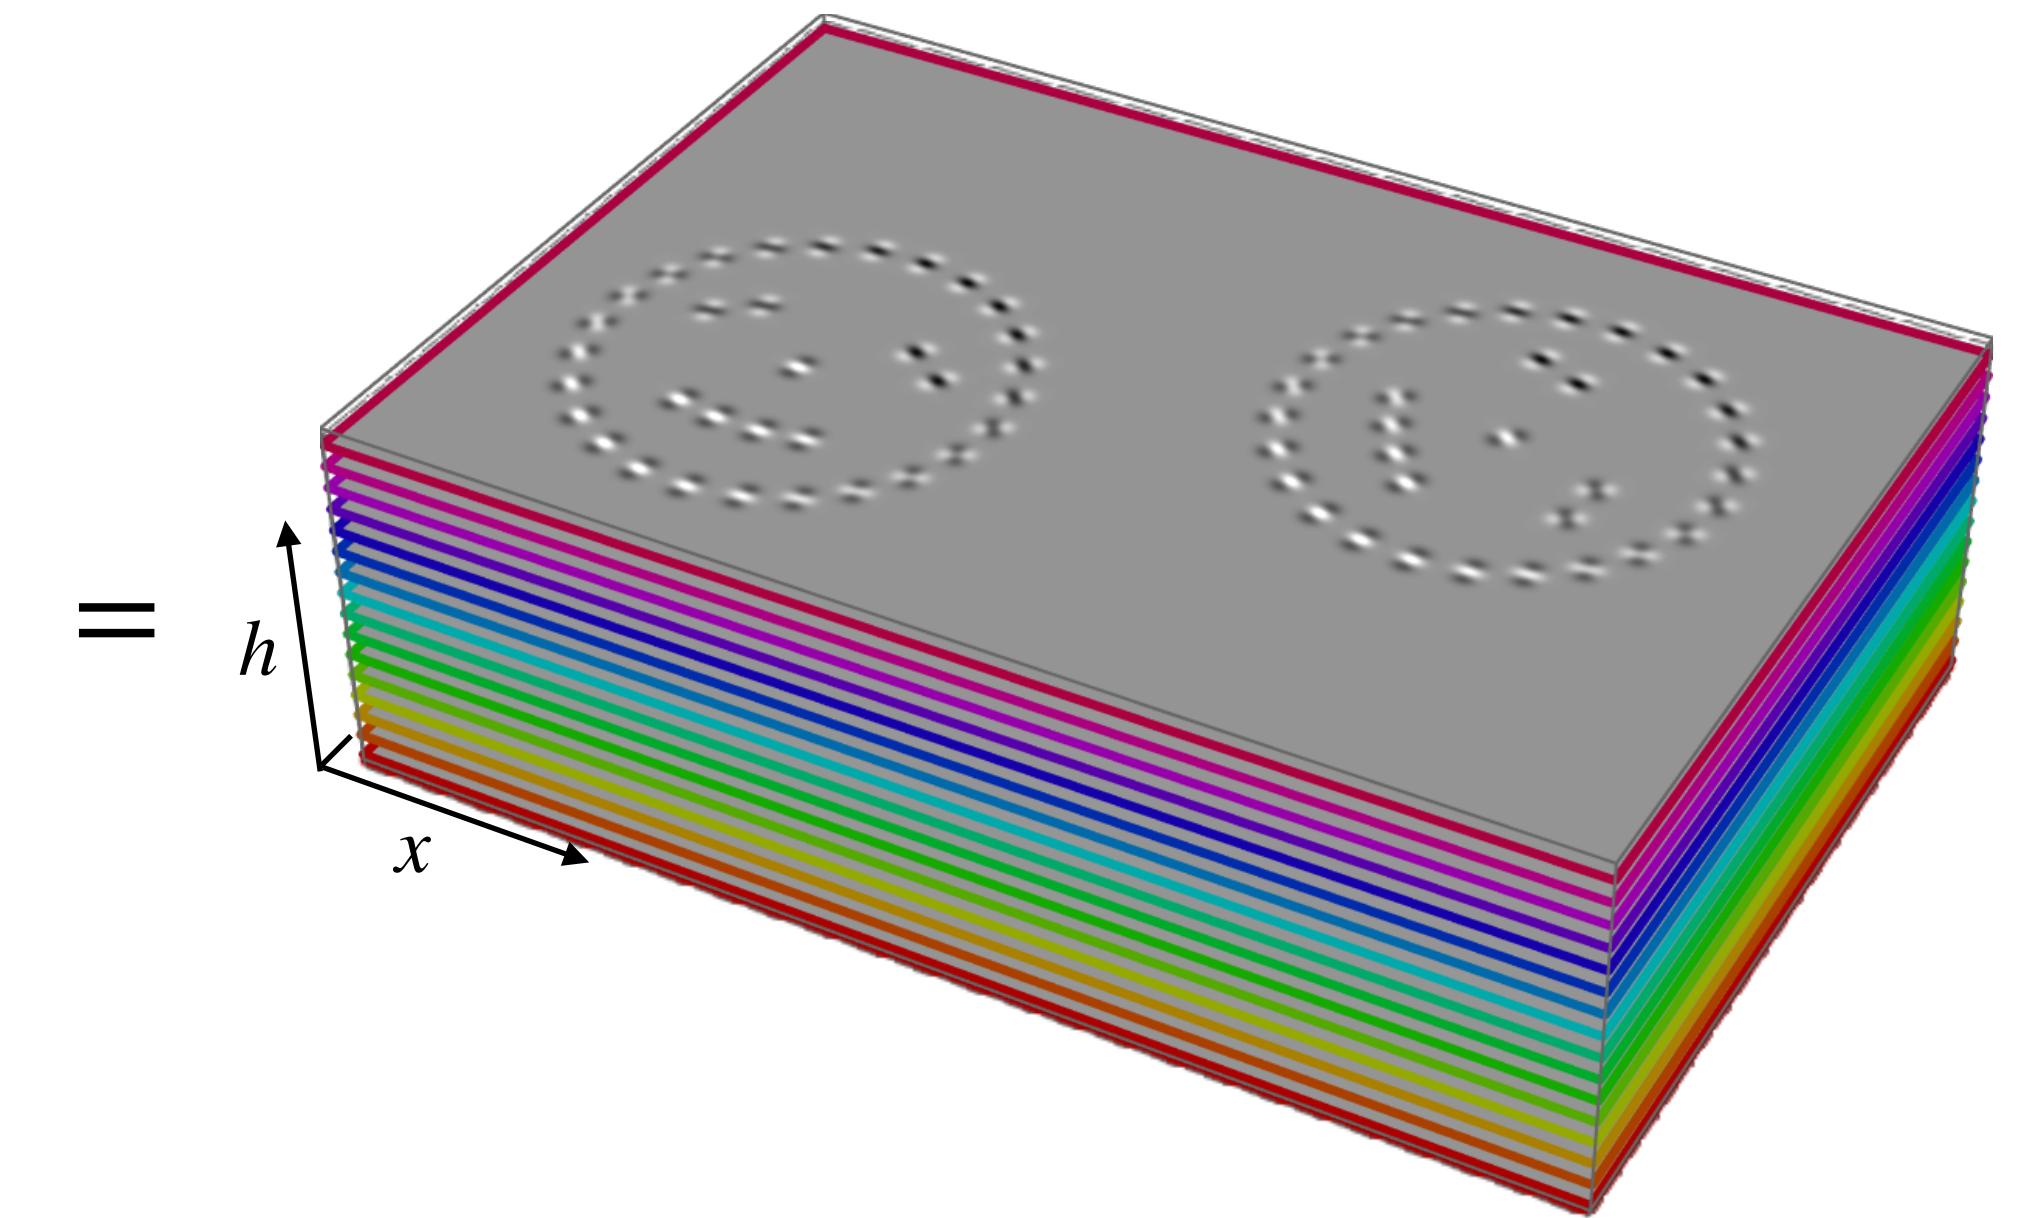
(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}')f(\mathbf{x}')d\mathbf{x}'$$



2D convolution kernel

2D input feature map



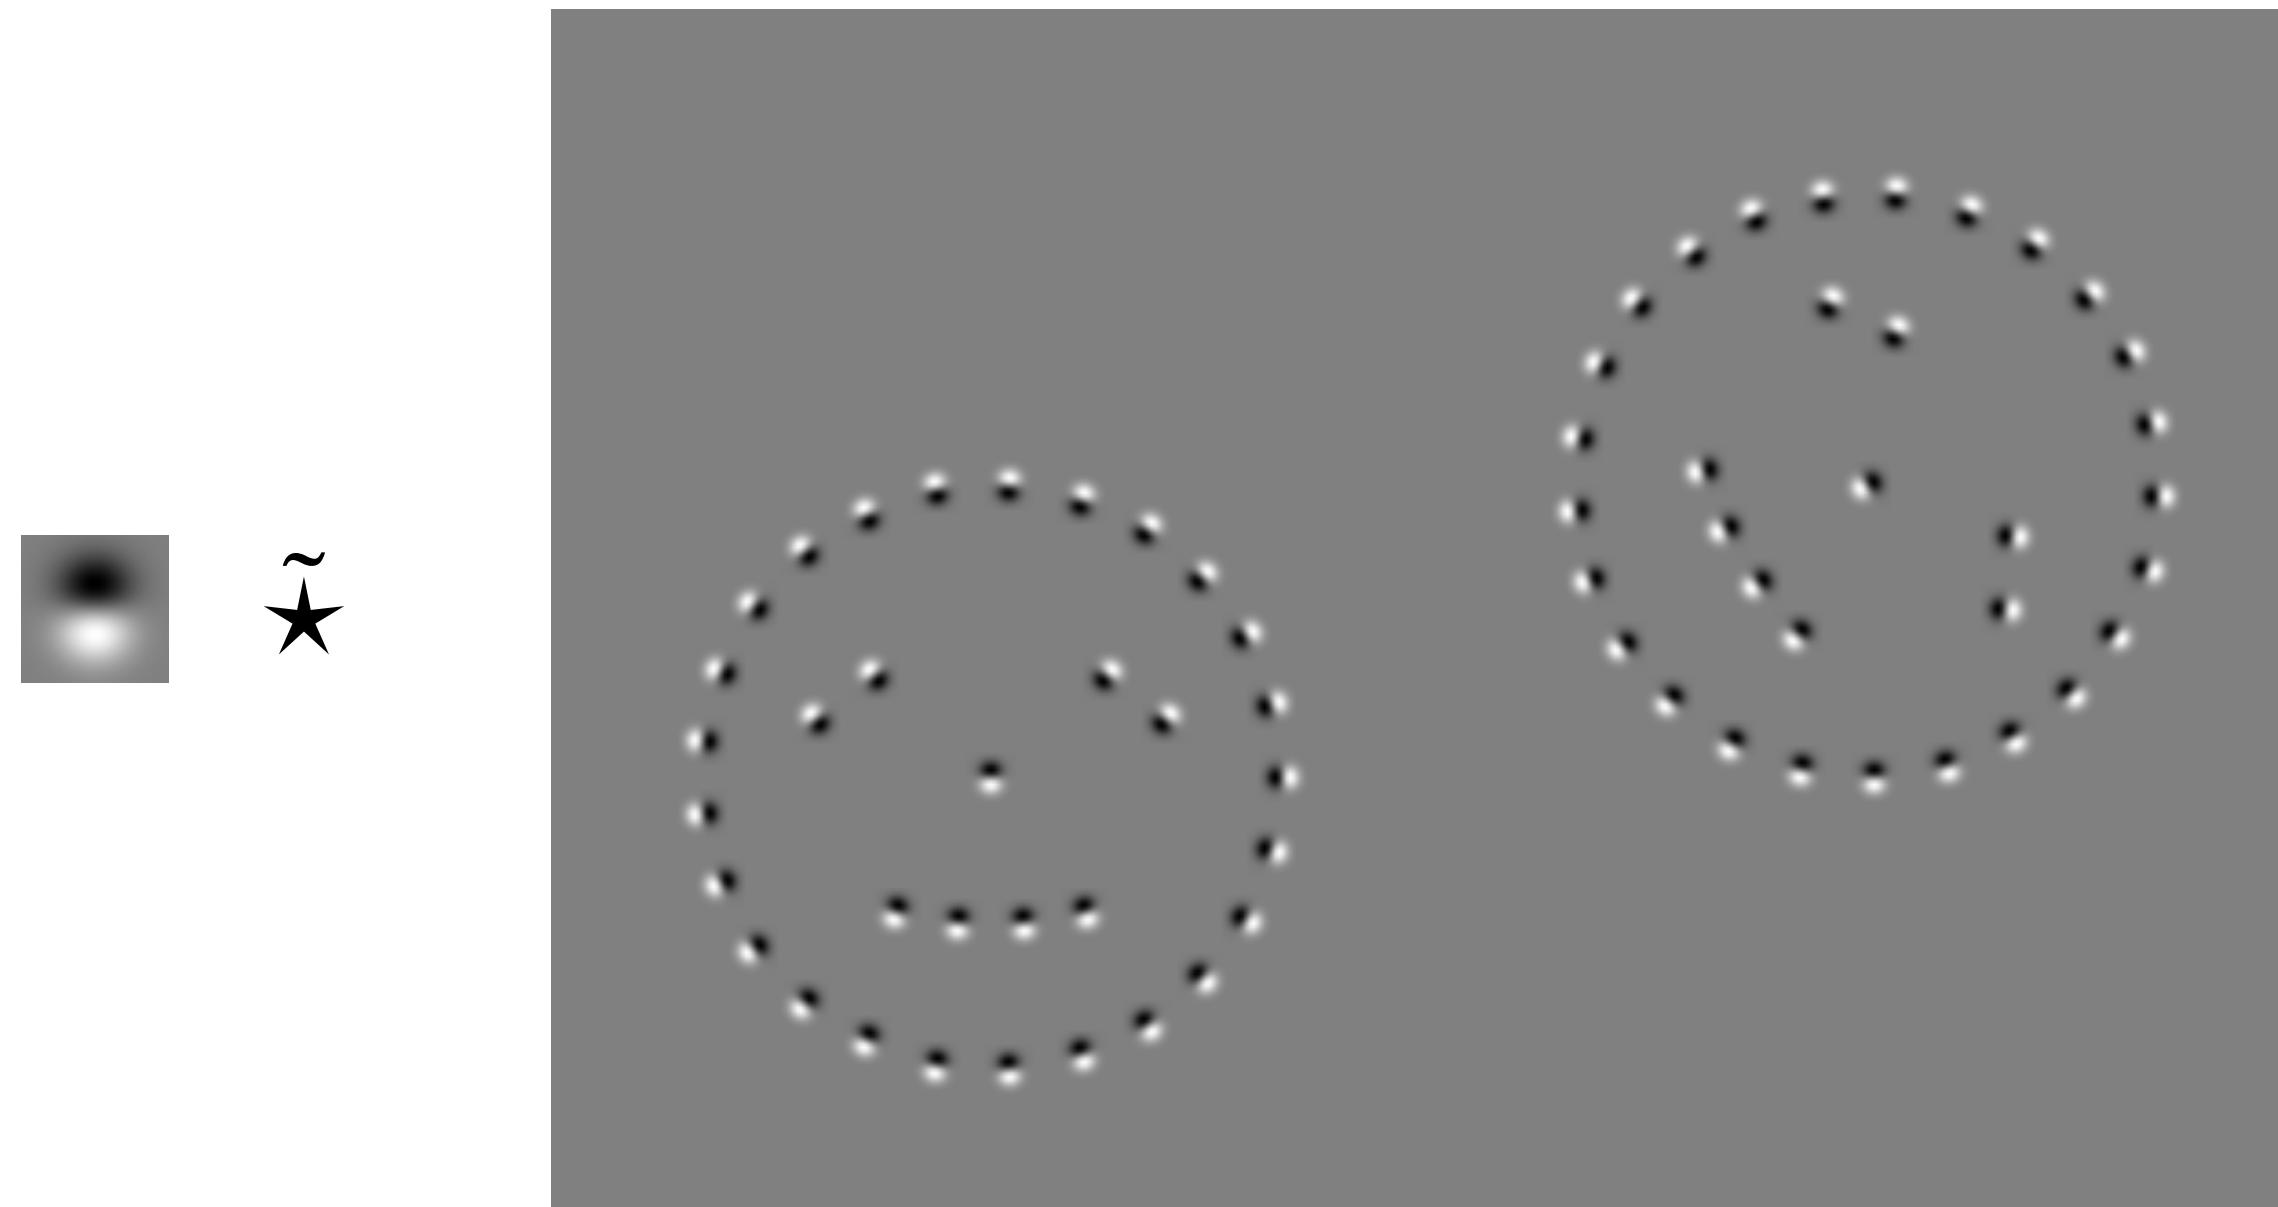
$SE(2)$  output feature map

# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)} = (\mathcal{L}_{\mathbf{x}}^{(\mathbb{R}^d,+)\rightarrow \mathbb{L}_2(\mathbb{R}^d)} \mathcal{L}_h^{H \rightarrow \mathbb{L}_2(\mathbb{R}^d)} k, f)_{\mathbb{L}_2(\mathbb{R}^d)}$

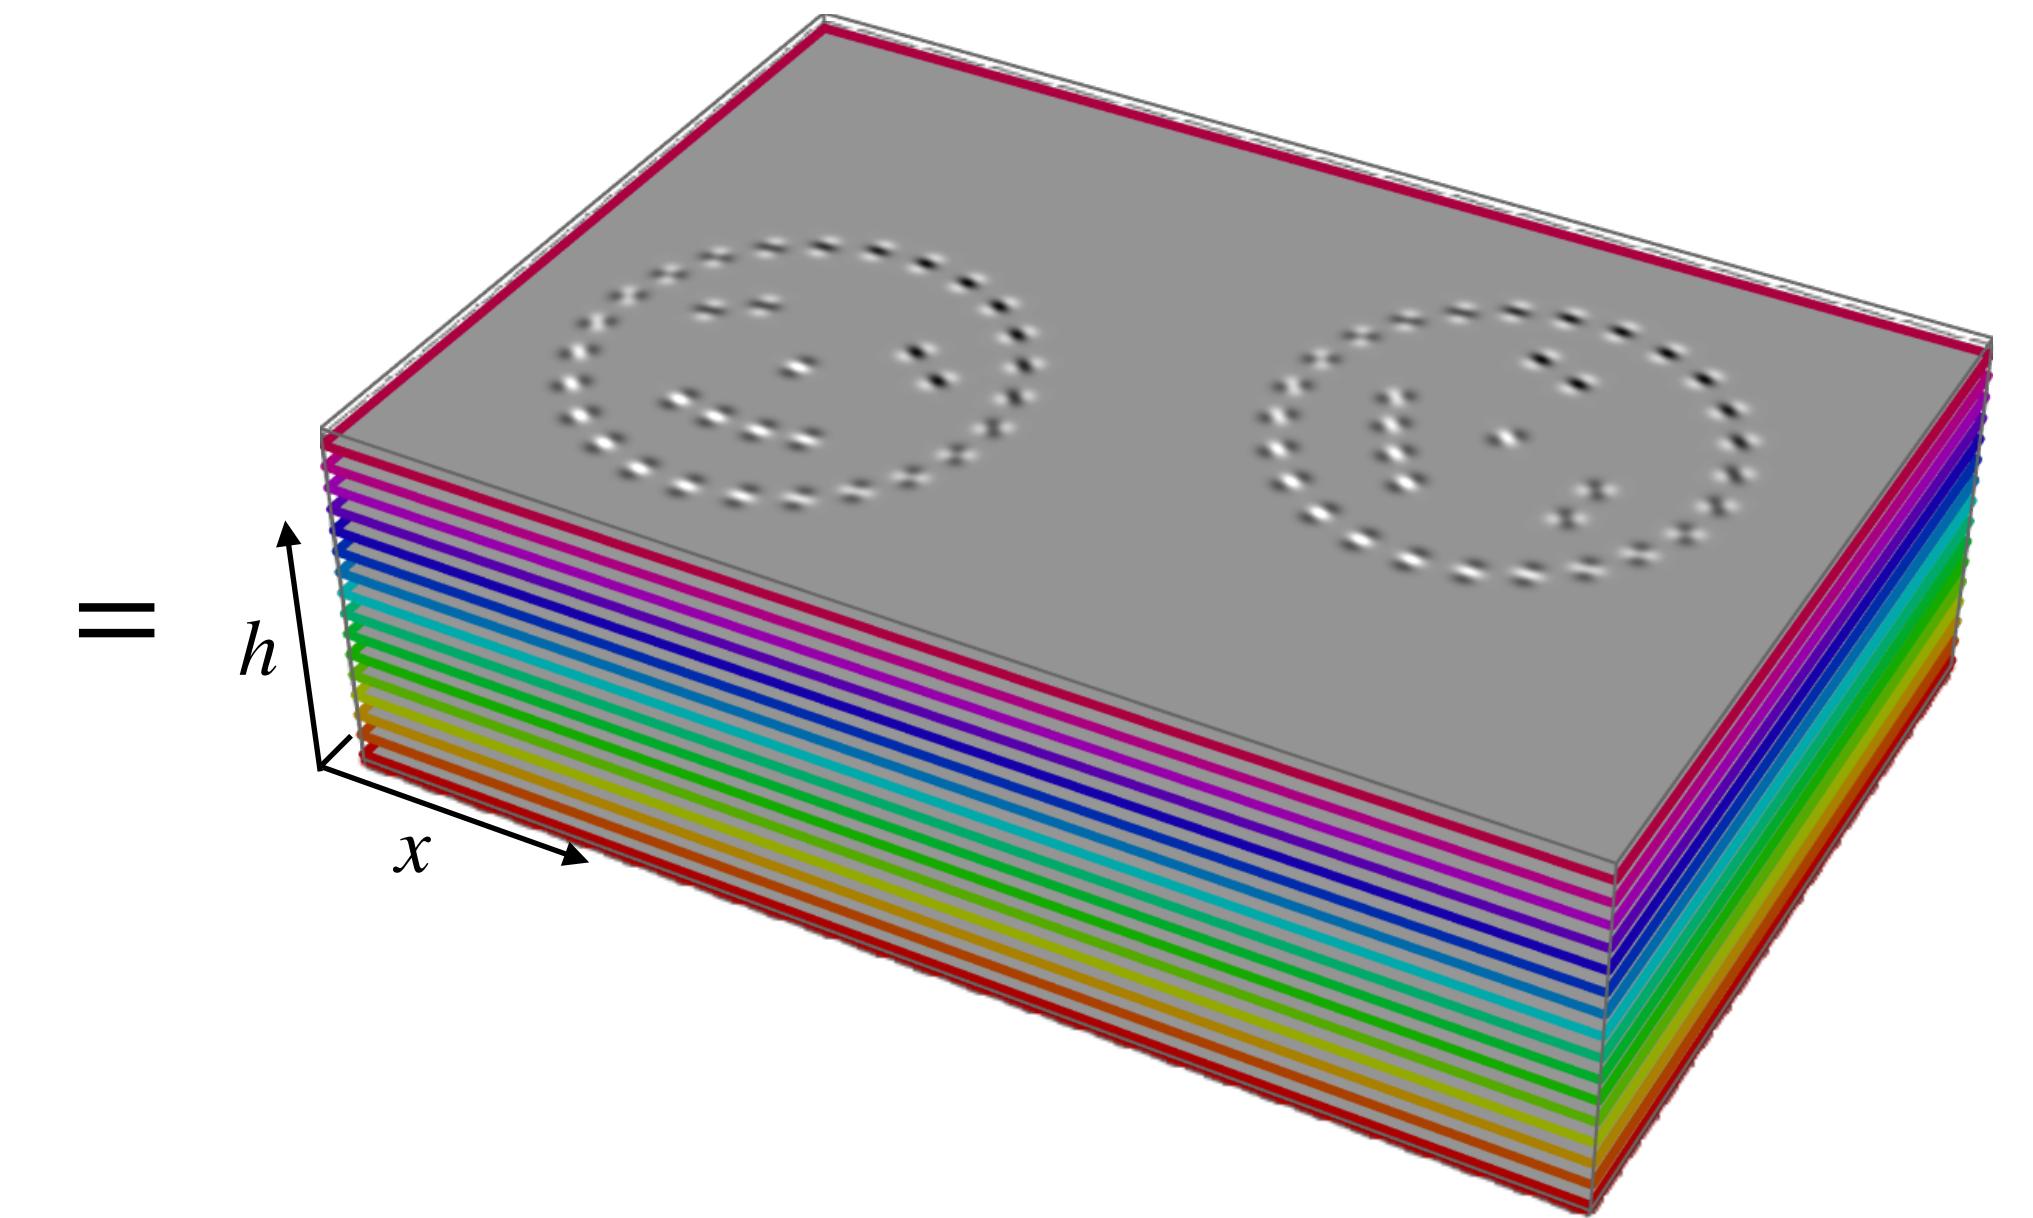
(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}')f(\mathbf{x}')d\mathbf{x}' = \int_{\mathbb{R}^2} k(h^{-1}(\mathbf{x}' - \mathbf{x}))f(\mathbf{x}')d\mathbf{x}'$$



2D convolution kernel

2D input feature map



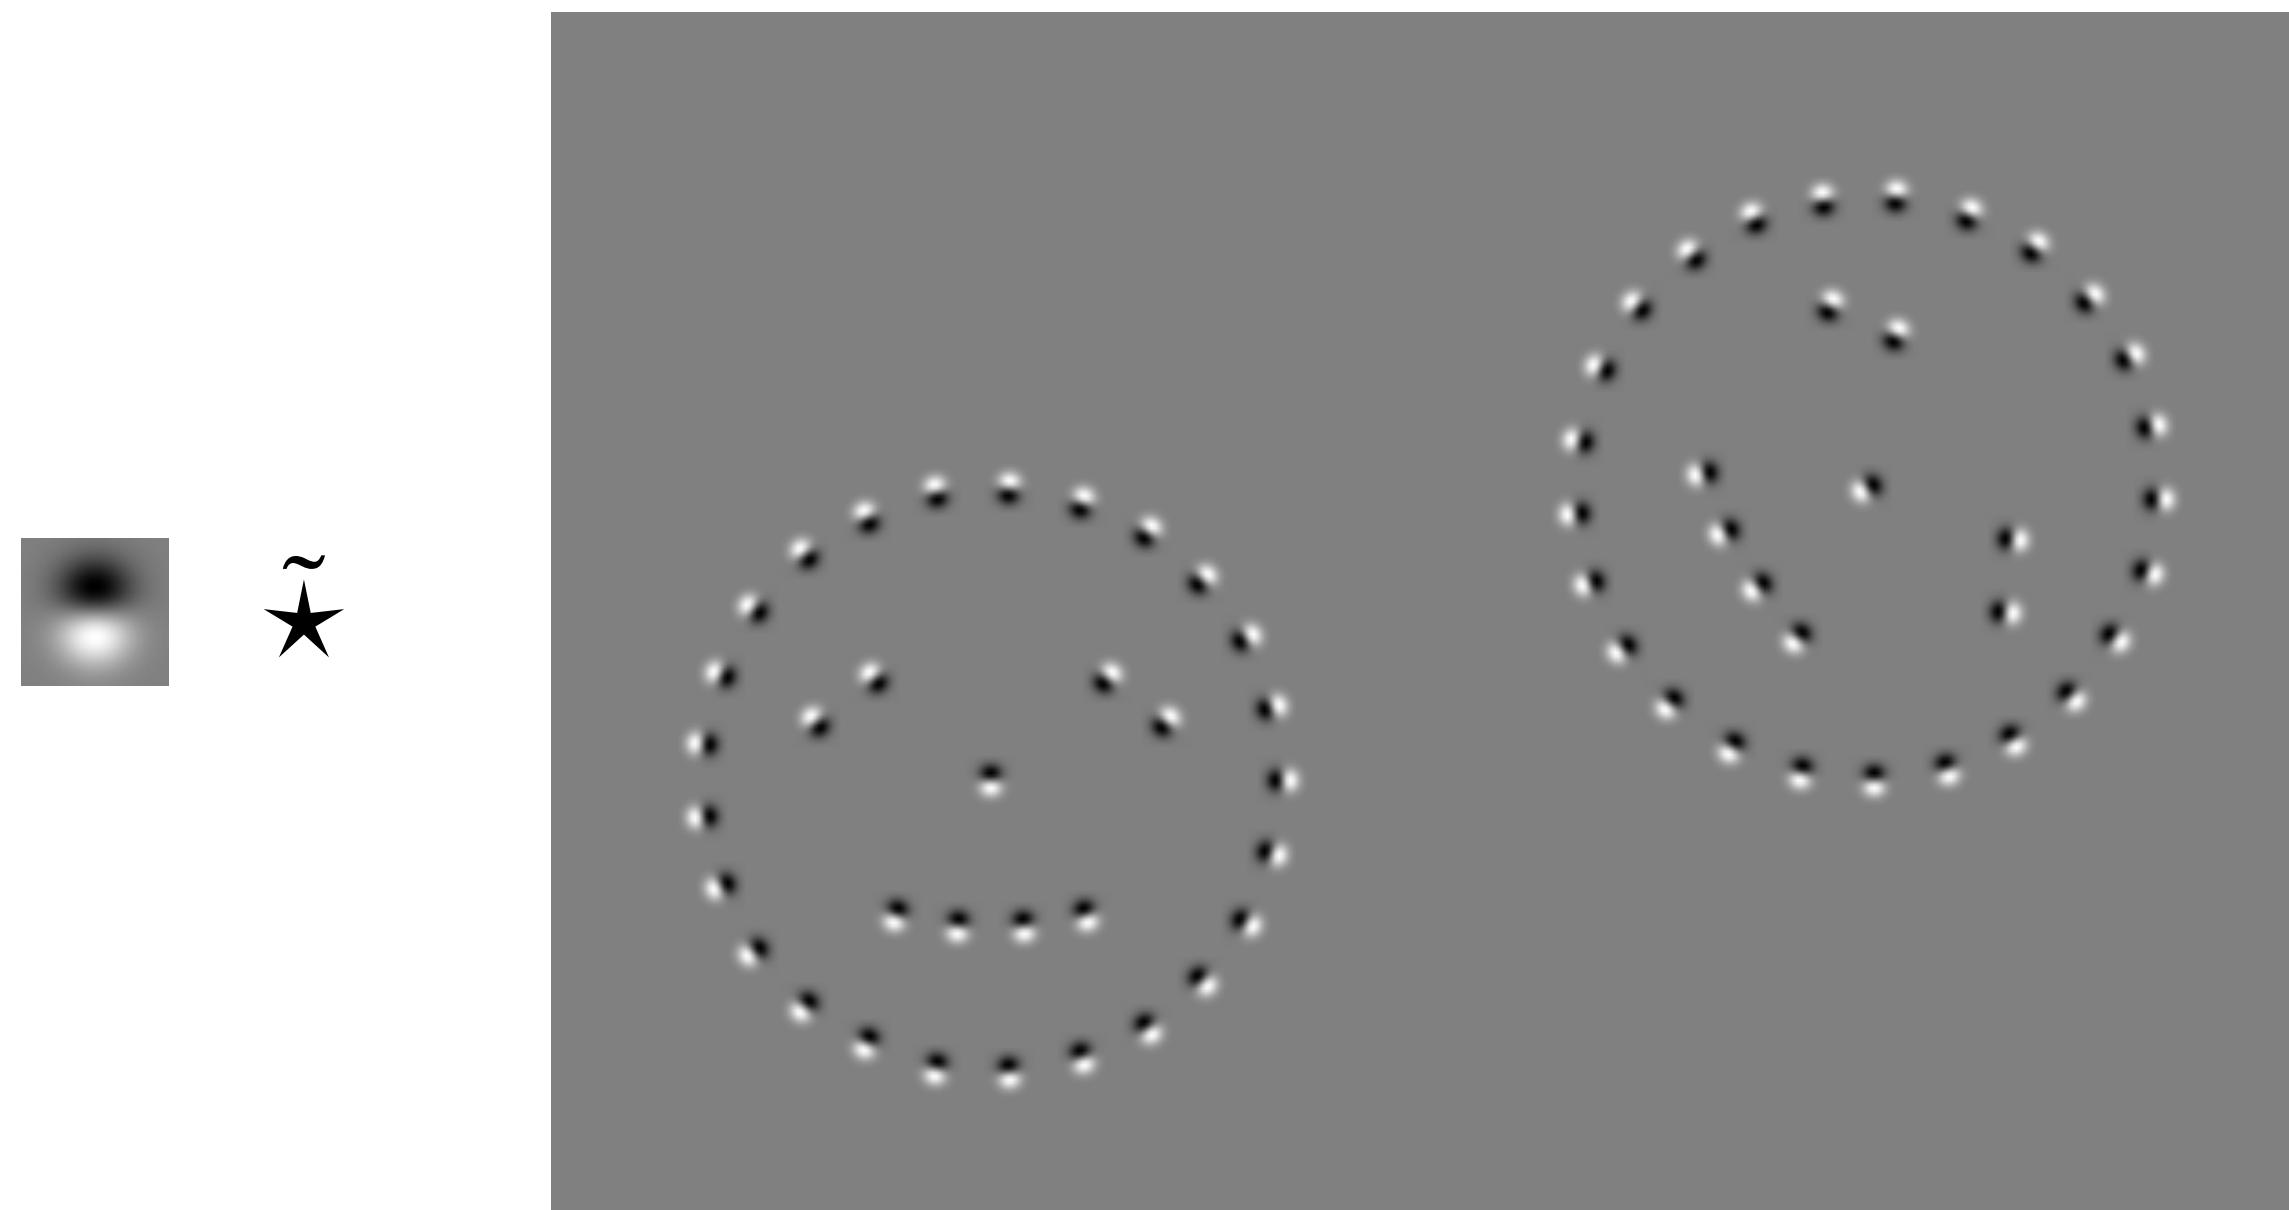
$SE(2)$  output feature map

# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)} = (\mathcal{L}_{\mathbf{x}}^{(\mathbb{R}^d,+)\rightarrow \mathbb{L}_2(\mathbb{R}^d)} \mathcal{L}_h^{H \rightarrow \mathbb{L}_2(\mathbb{R}^d)} k, f)_{\mathbb{L}_2(\mathbb{R}^d)}$

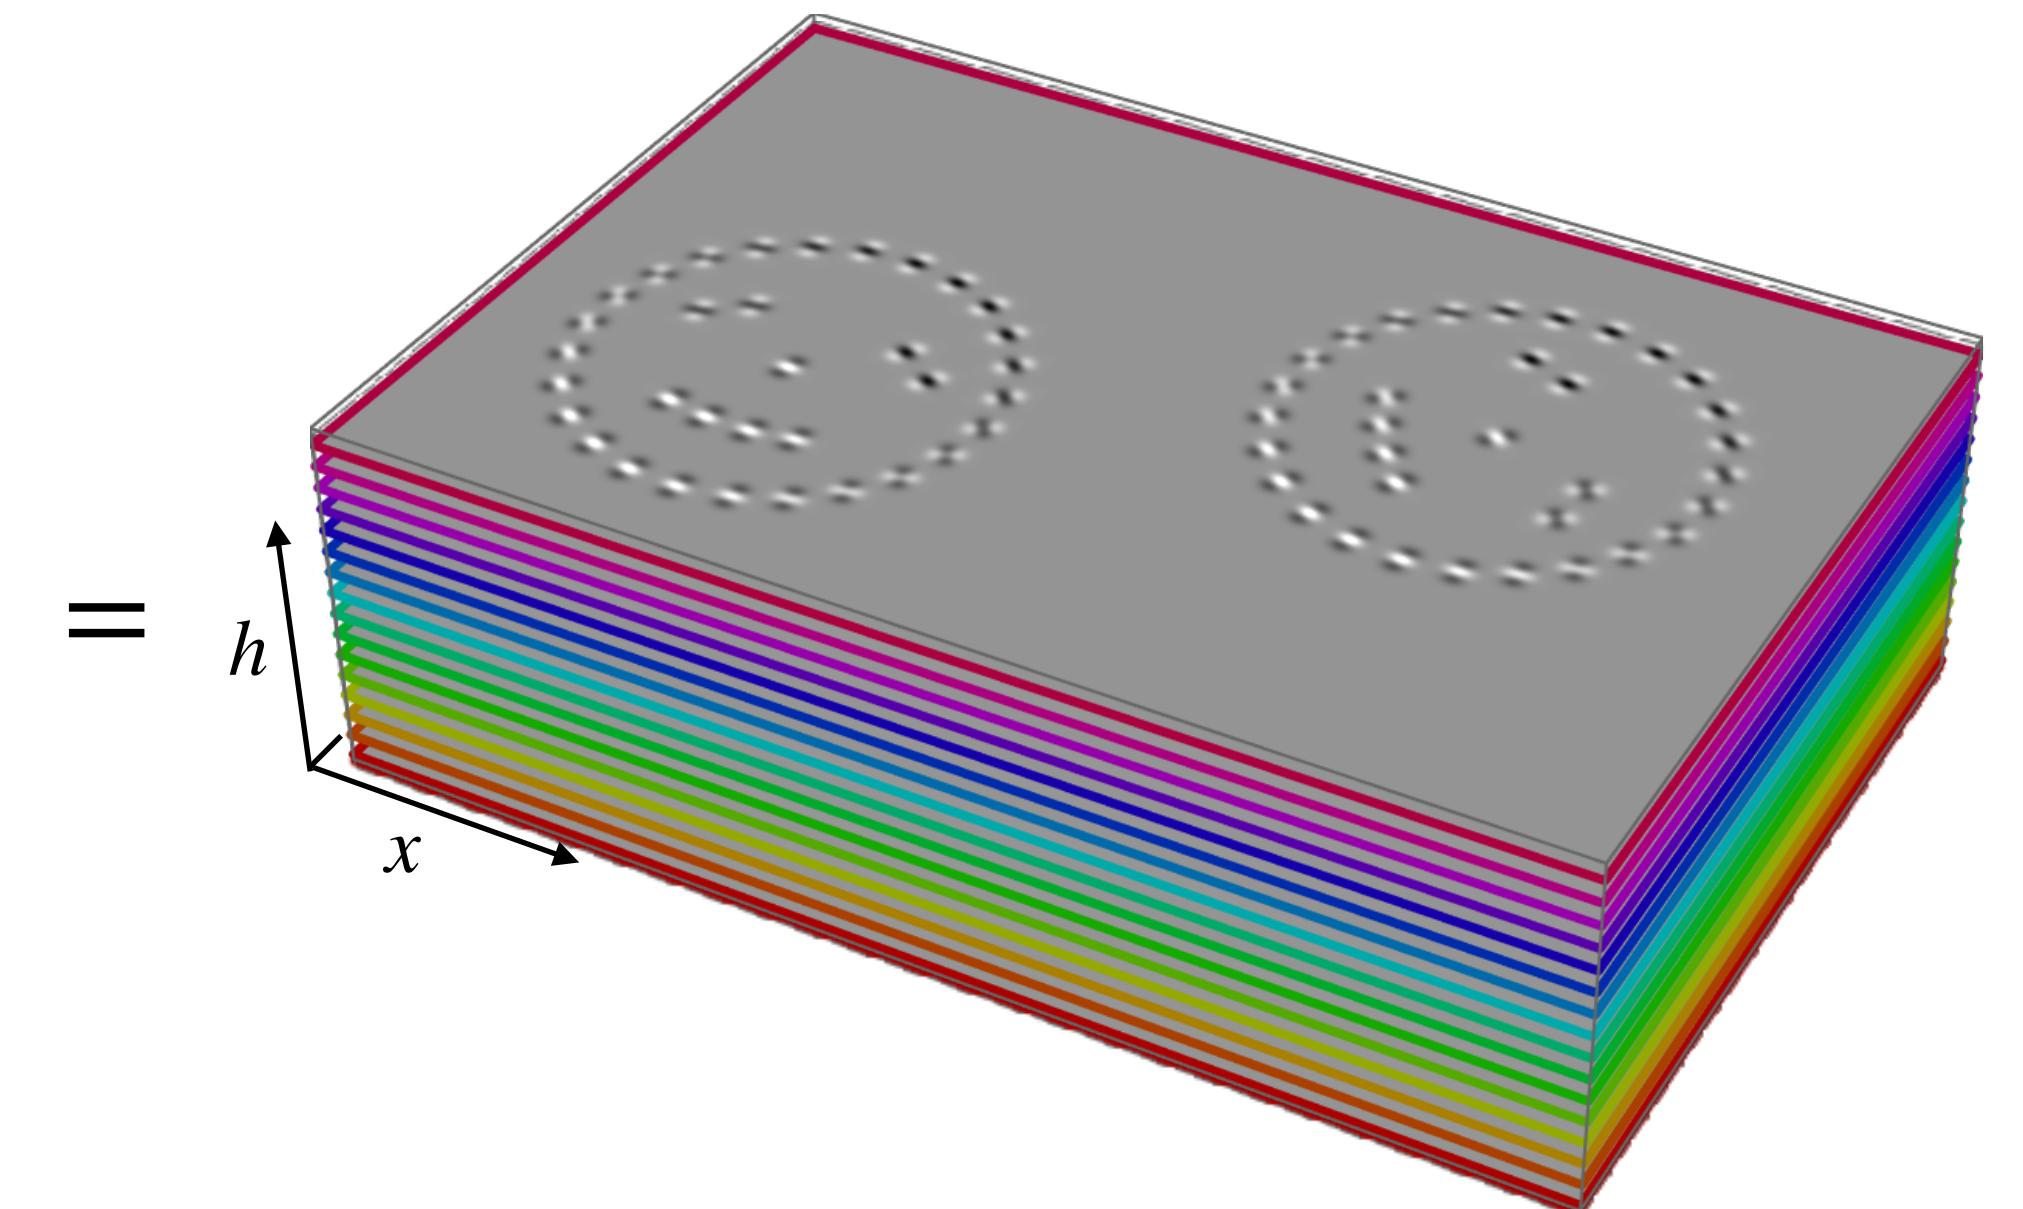
(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}')f(\mathbf{x}')d\mathbf{x}' = \int_{\mathbb{R}^2} k_h(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$



2D convolution kernel

2D input feature map



$SE(2)$  output feature map

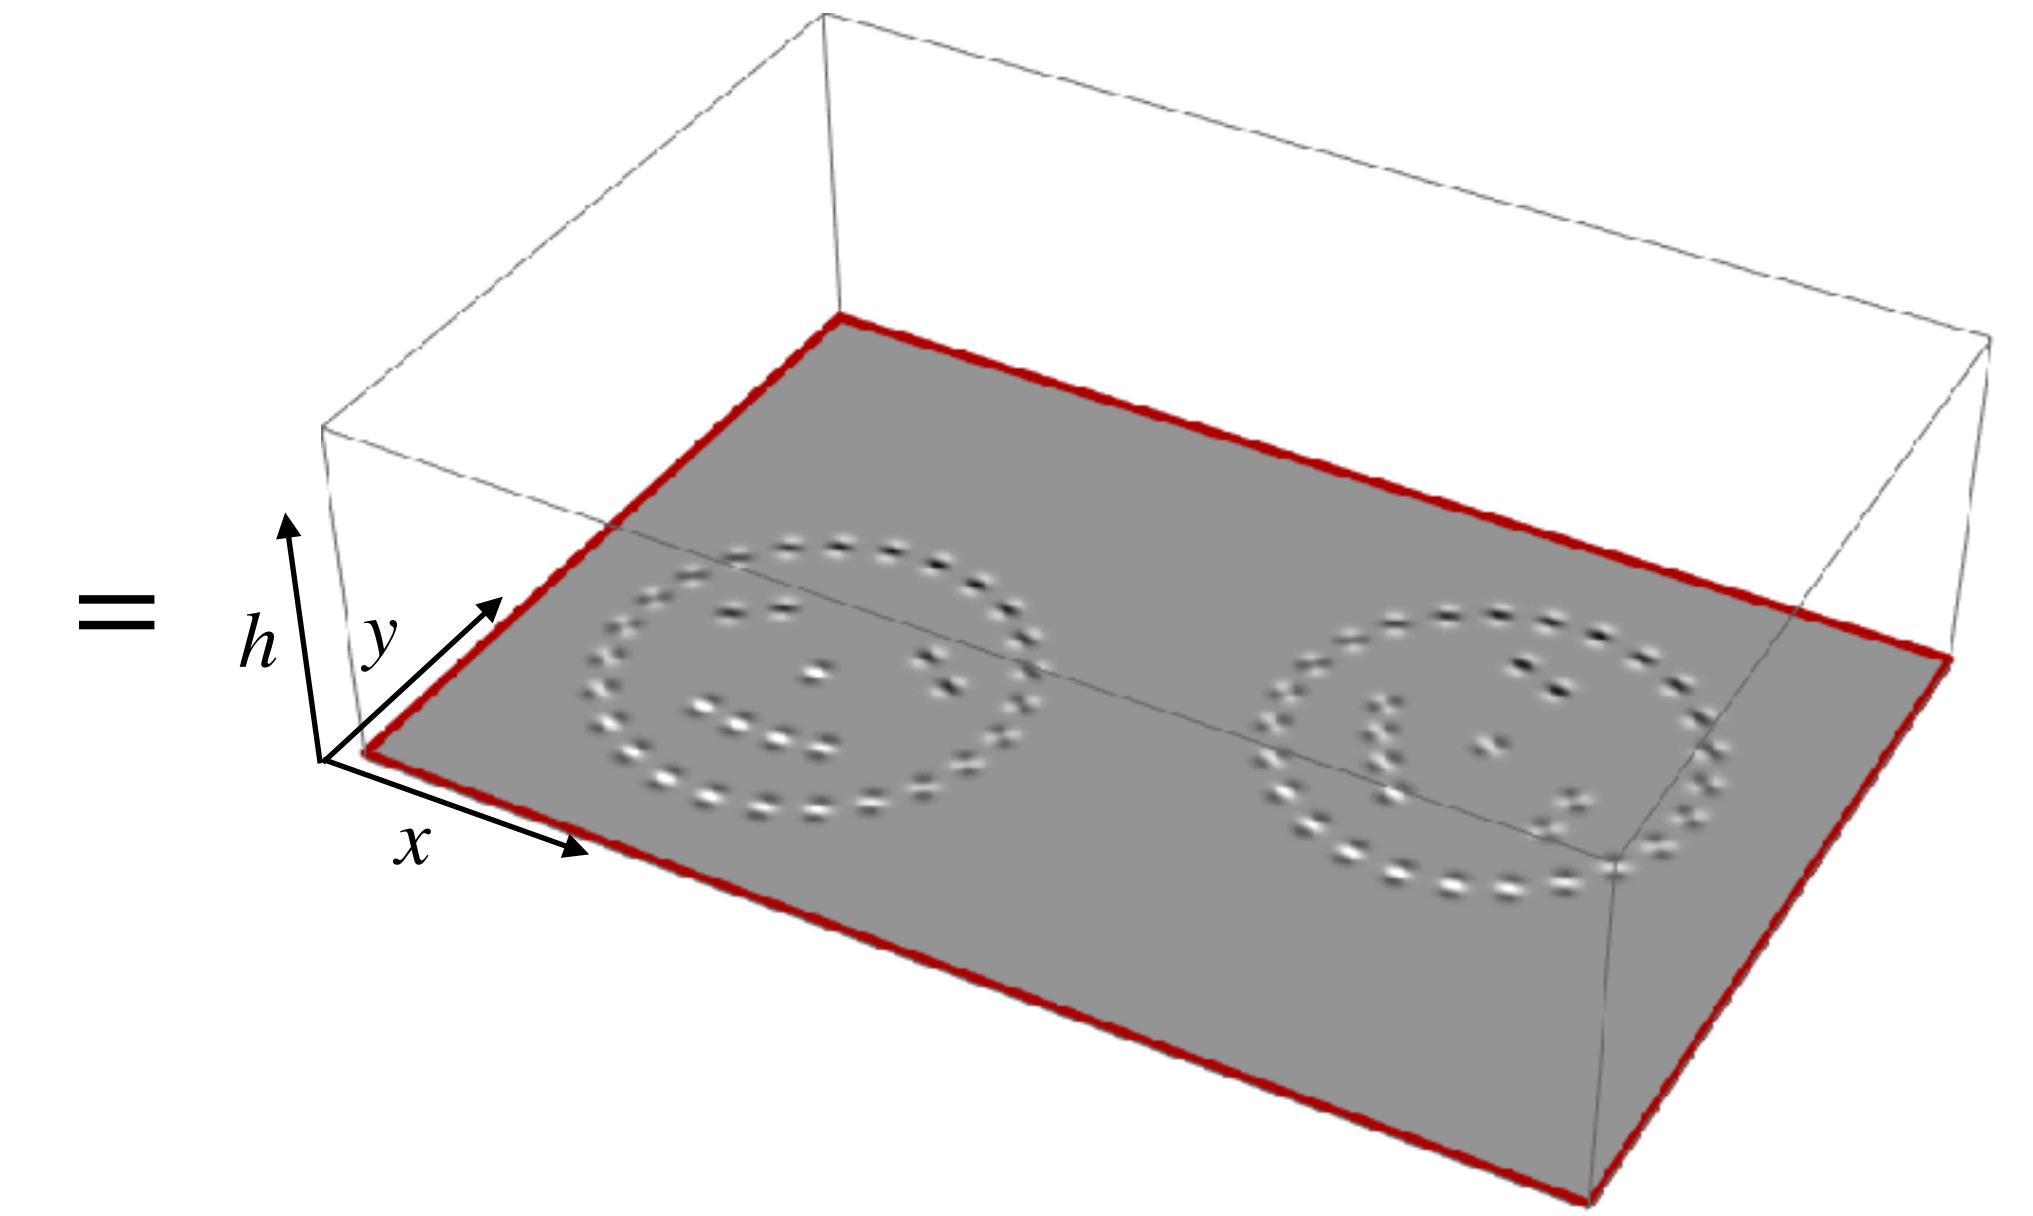
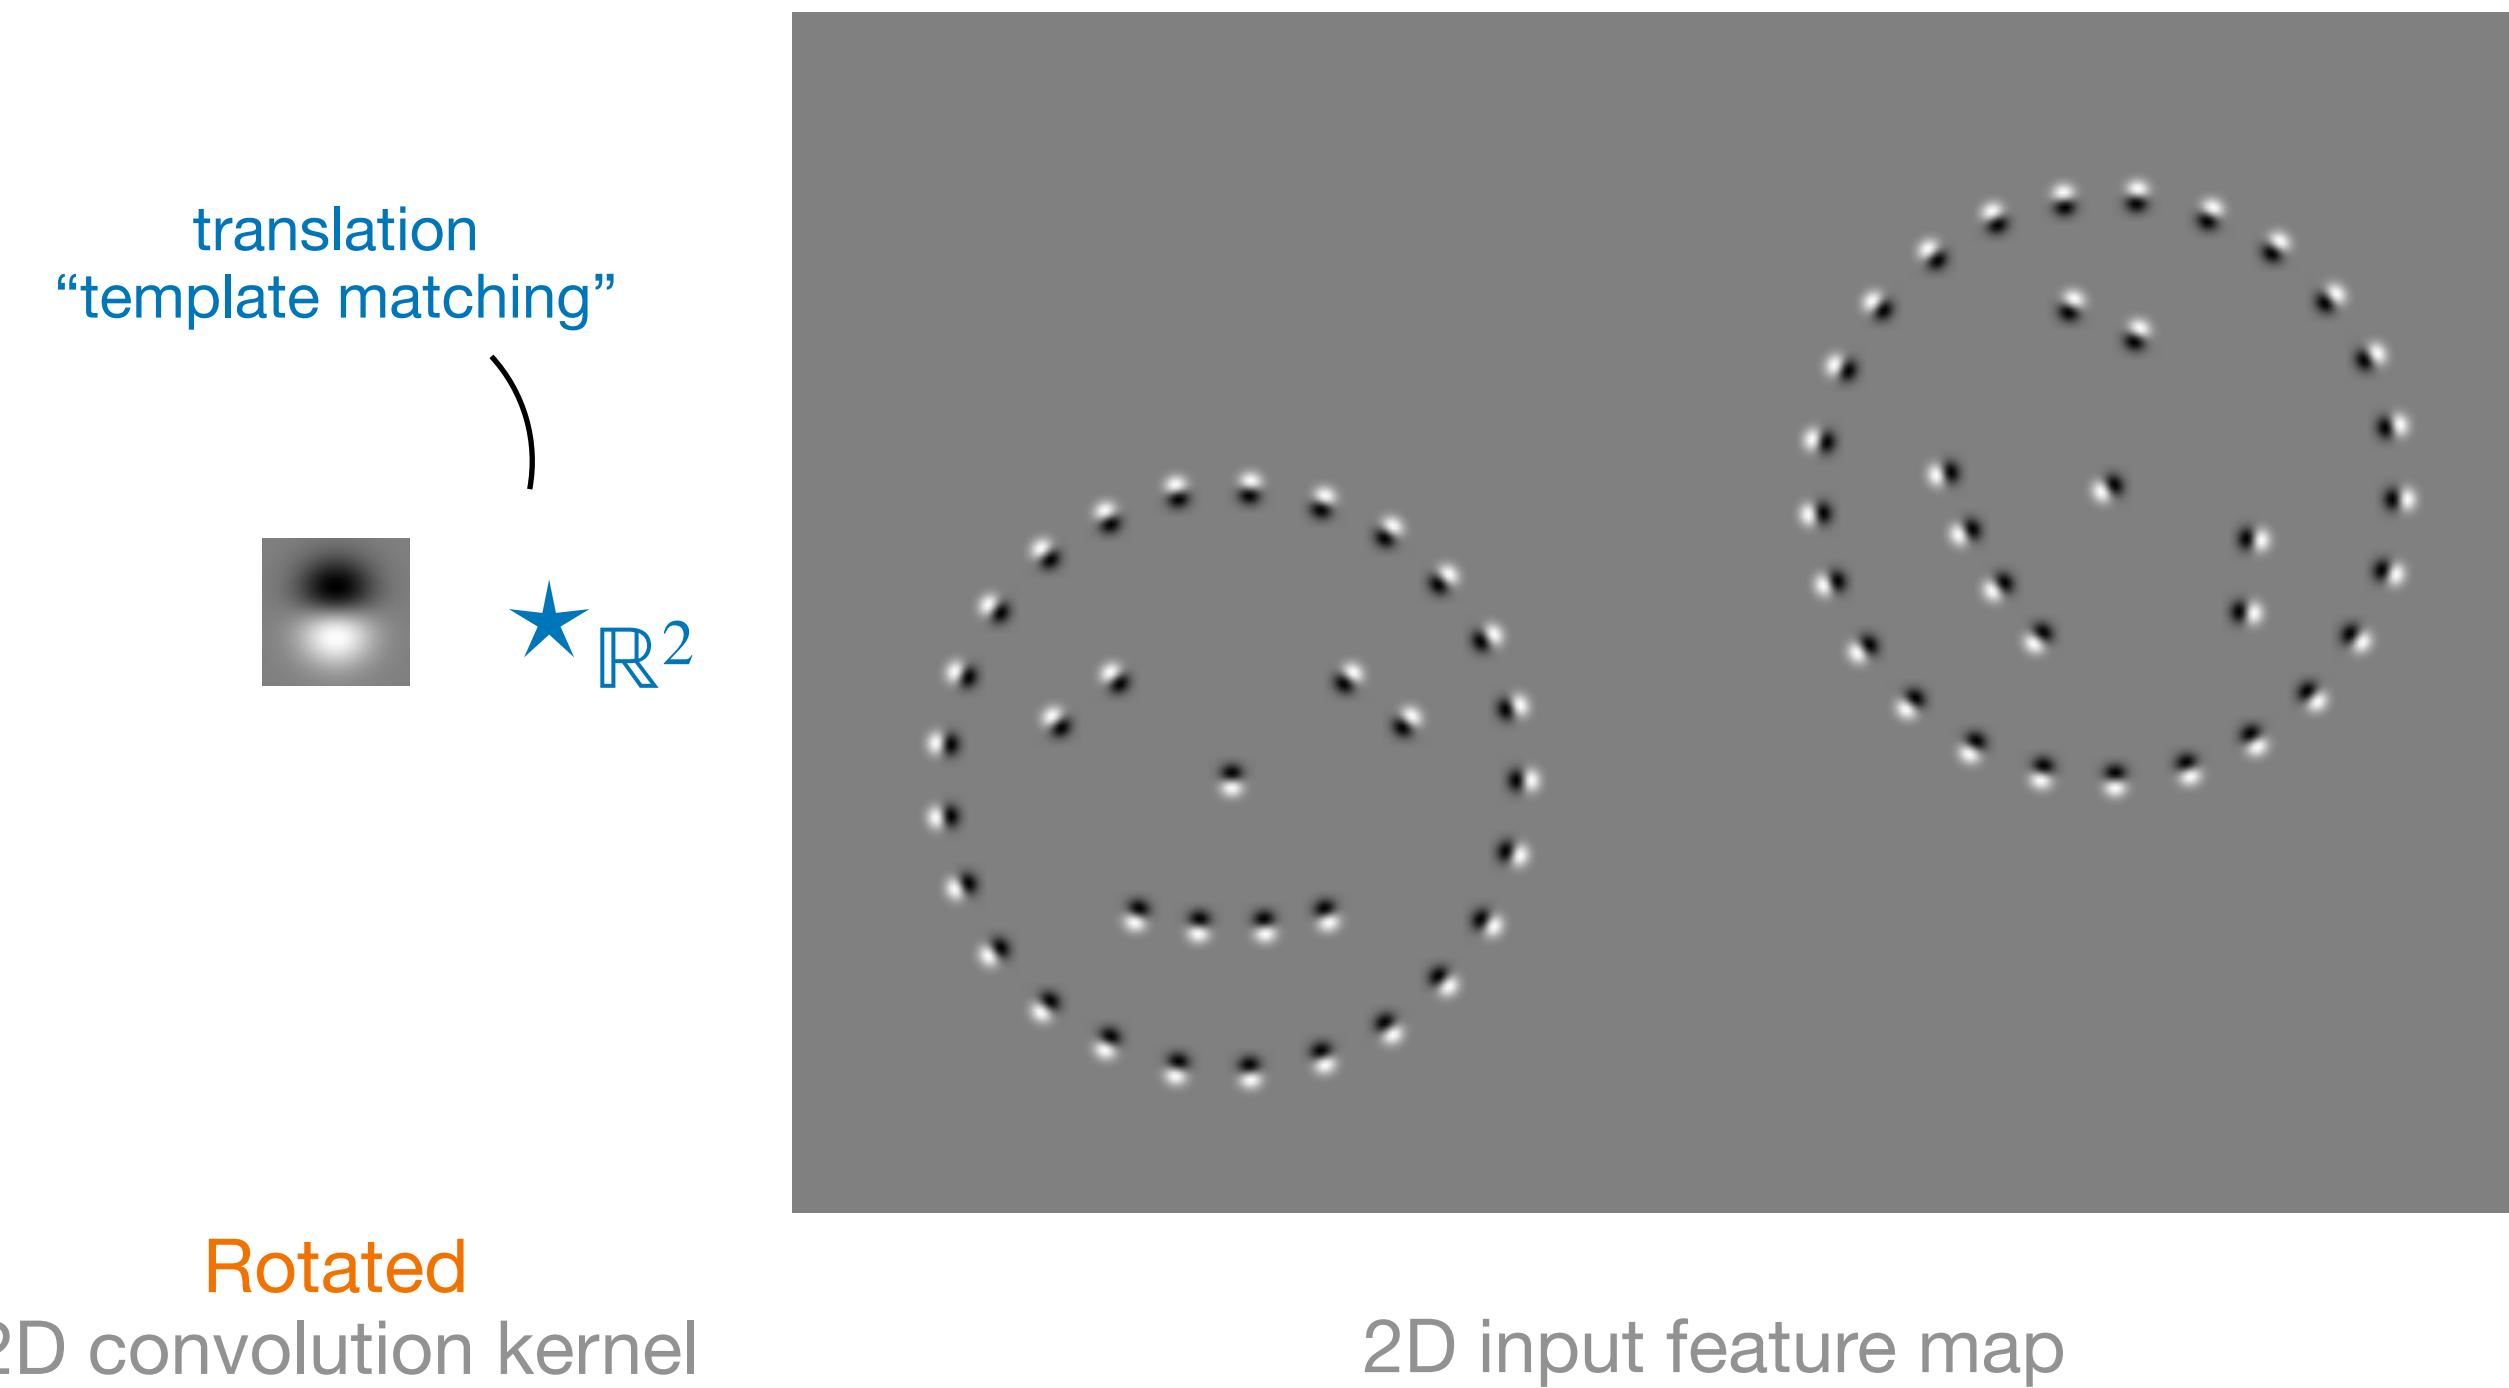
# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)} = (\mathcal{L}_{\mathbf{x}}^{(\mathbb{R}^d,+)\rightarrow \mathbb{L}_2(\mathbb{R}^d)} \mathcal{L}_h^{H \rightarrow \mathbb{L}_2(\mathbb{R}^d)} k, f)_{\mathbb{L}_2(\mathbb{R}^d)}$

(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}')f(\mathbf{x}')d\mathbf{x}' = \int_{\mathbb{R}^2} k_h(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$


---



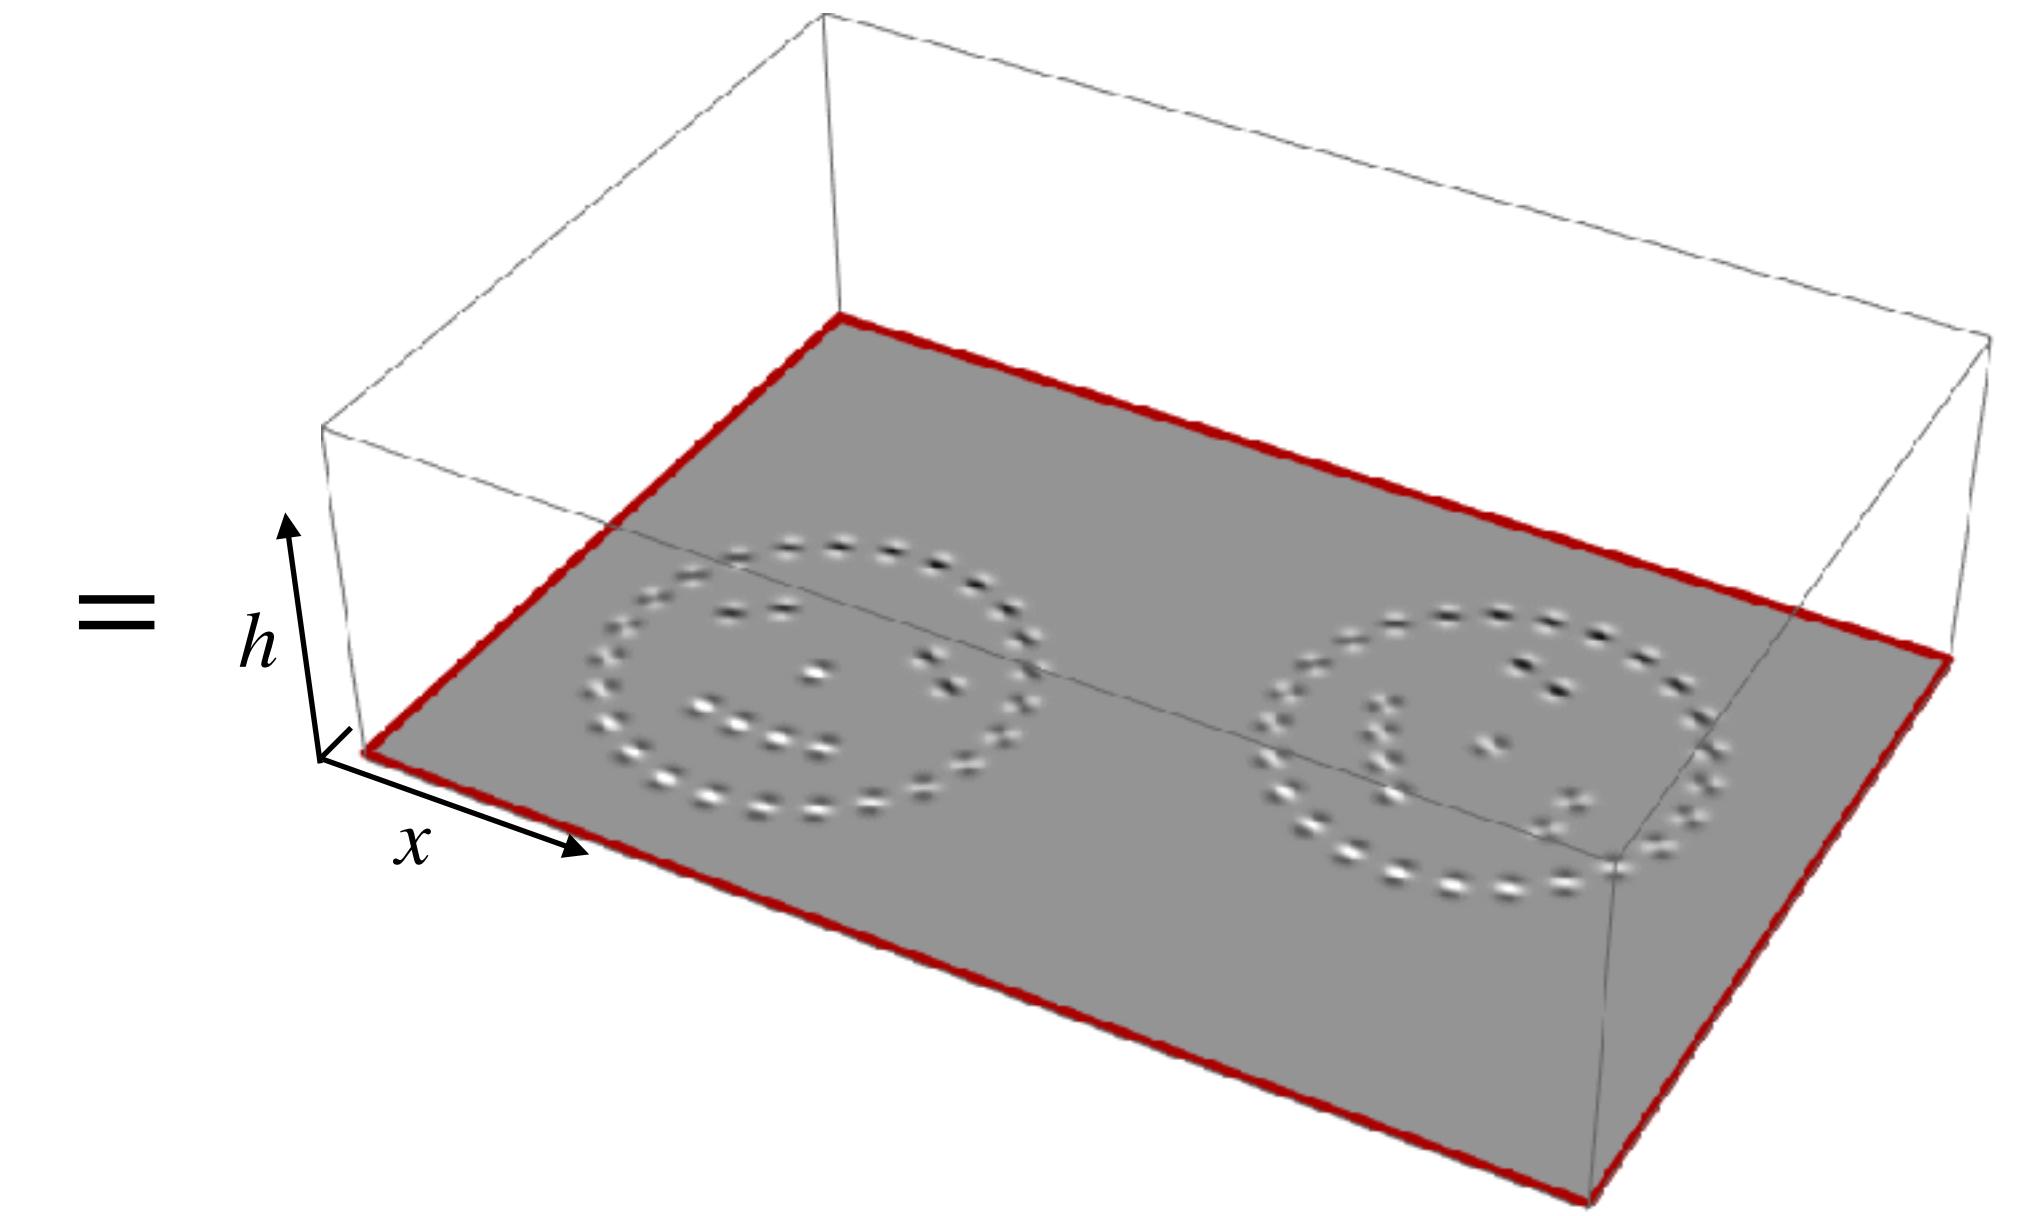
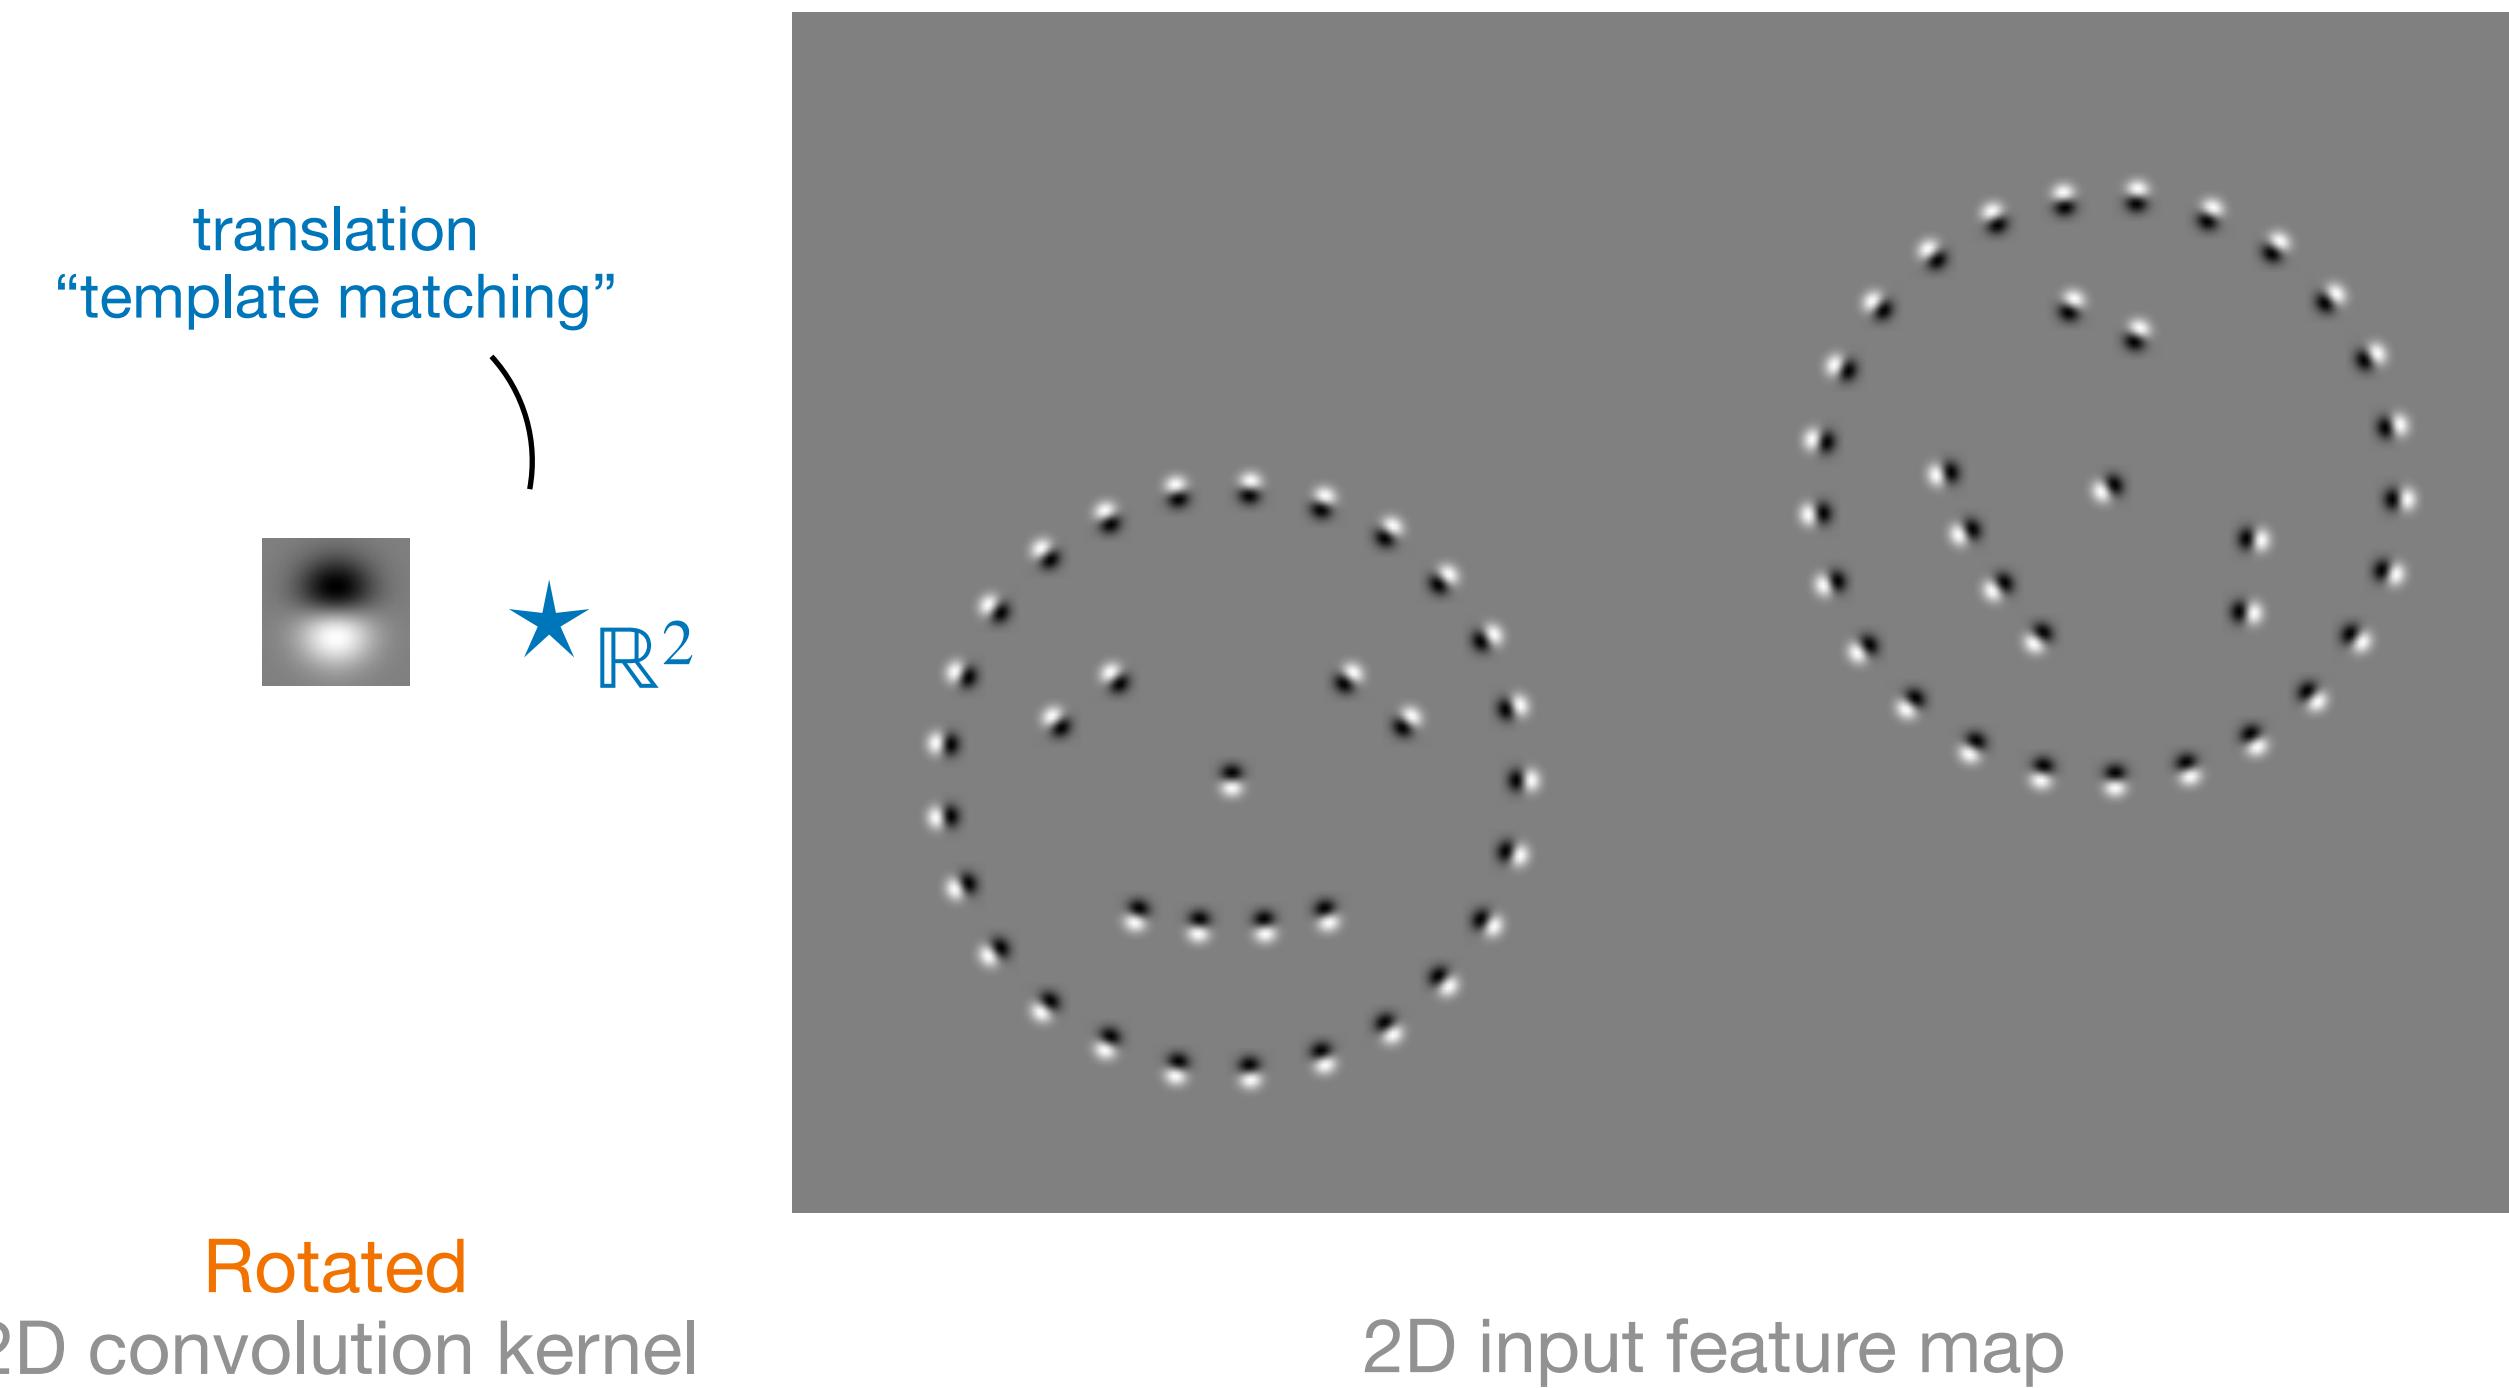
# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)} = (\mathcal{L}_{\mathbf{x}}^{(\mathbb{R}^d,+)\rightarrow \mathbb{L}_2(\mathbb{R}^d)} \mathcal{L}_h^{H \rightarrow \mathbb{L}_2(\mathbb{R}^d)} k, f)_{\mathbb{L}_2(\mathbb{R}^d)}$

(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}')f(\mathbf{x}')d\mathbf{x}' = \int_{\mathbb{R}^2} k_h(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$


---



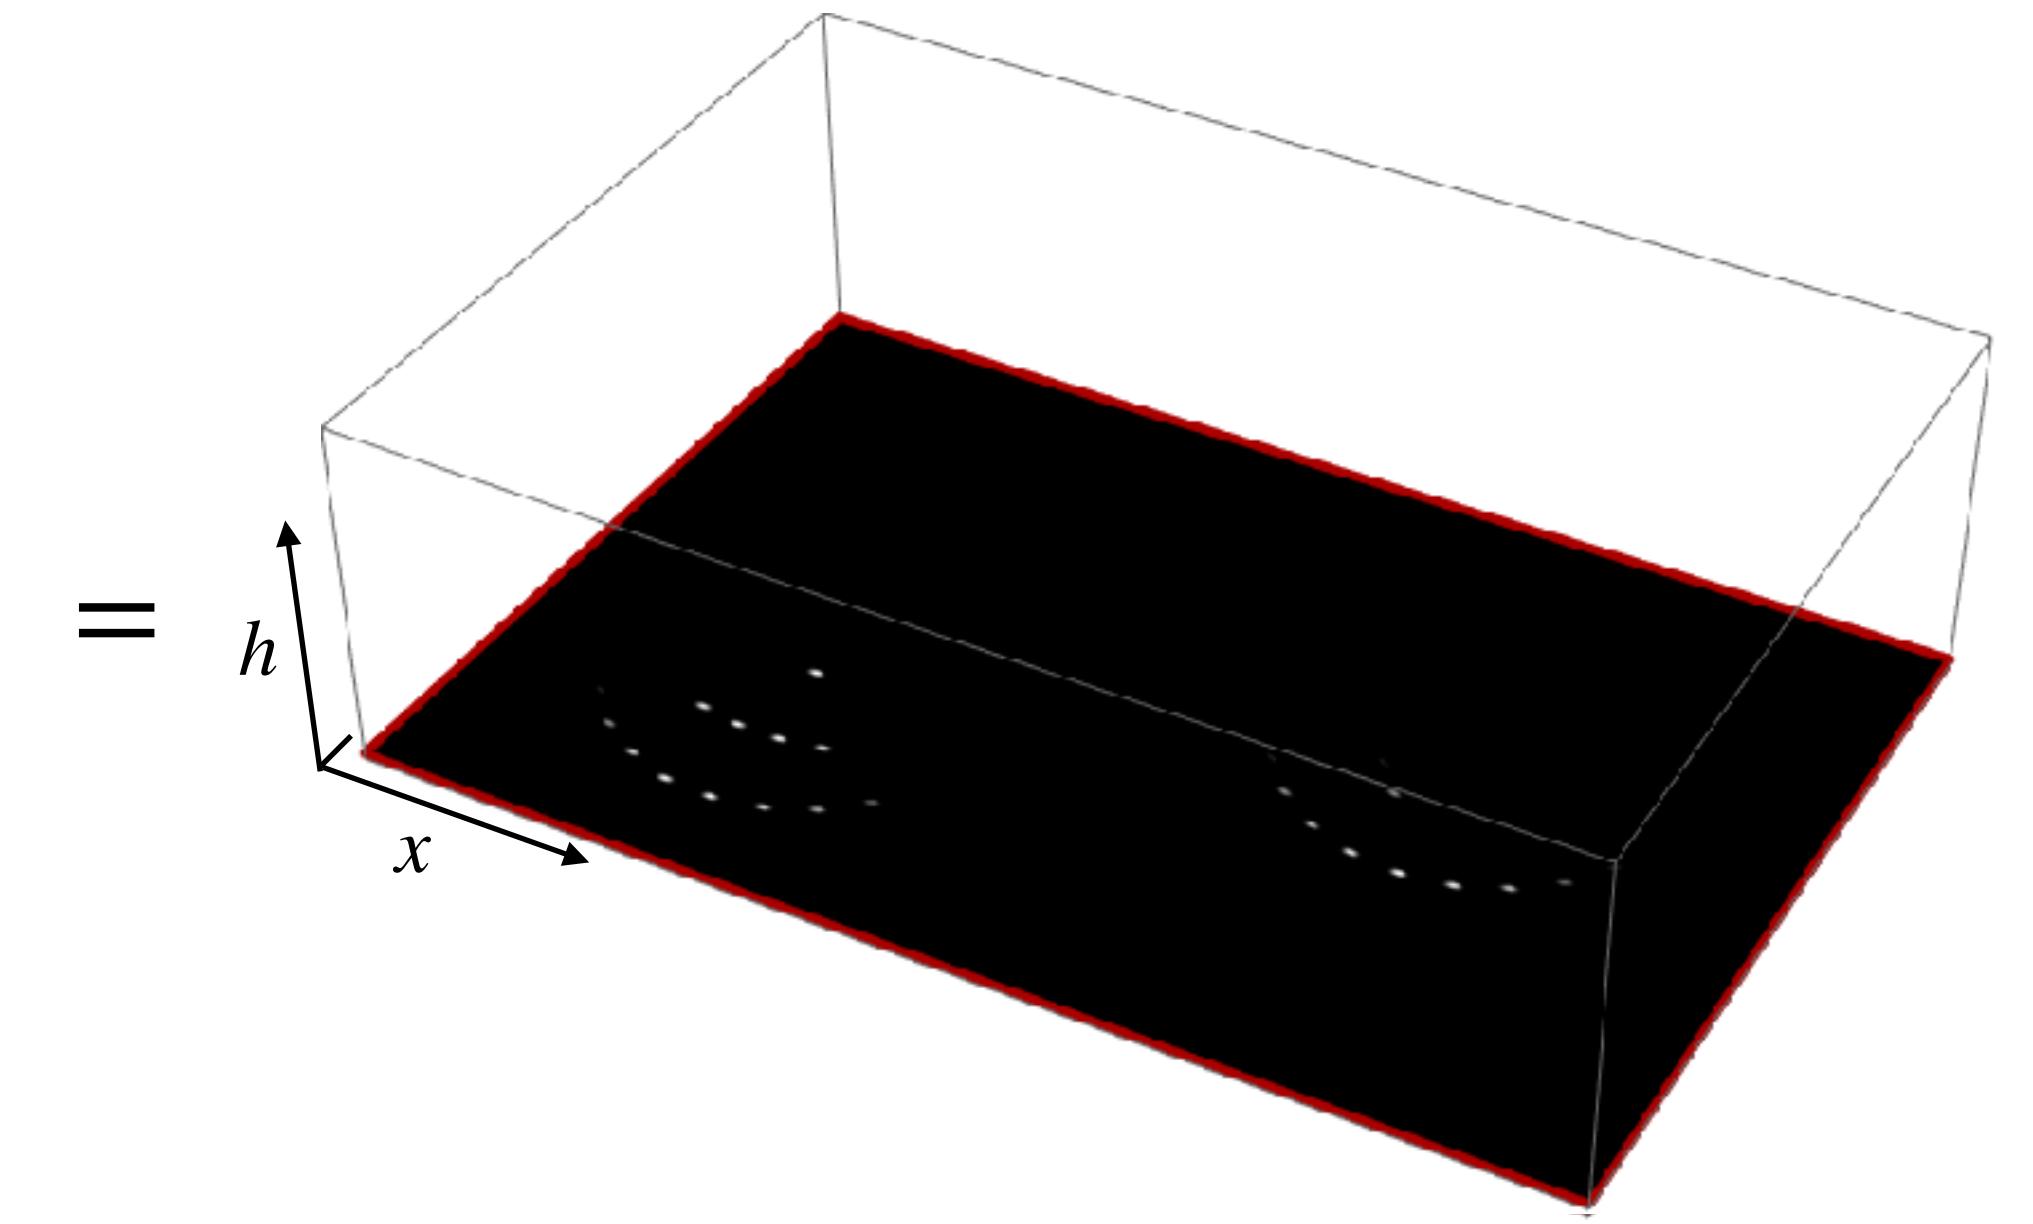
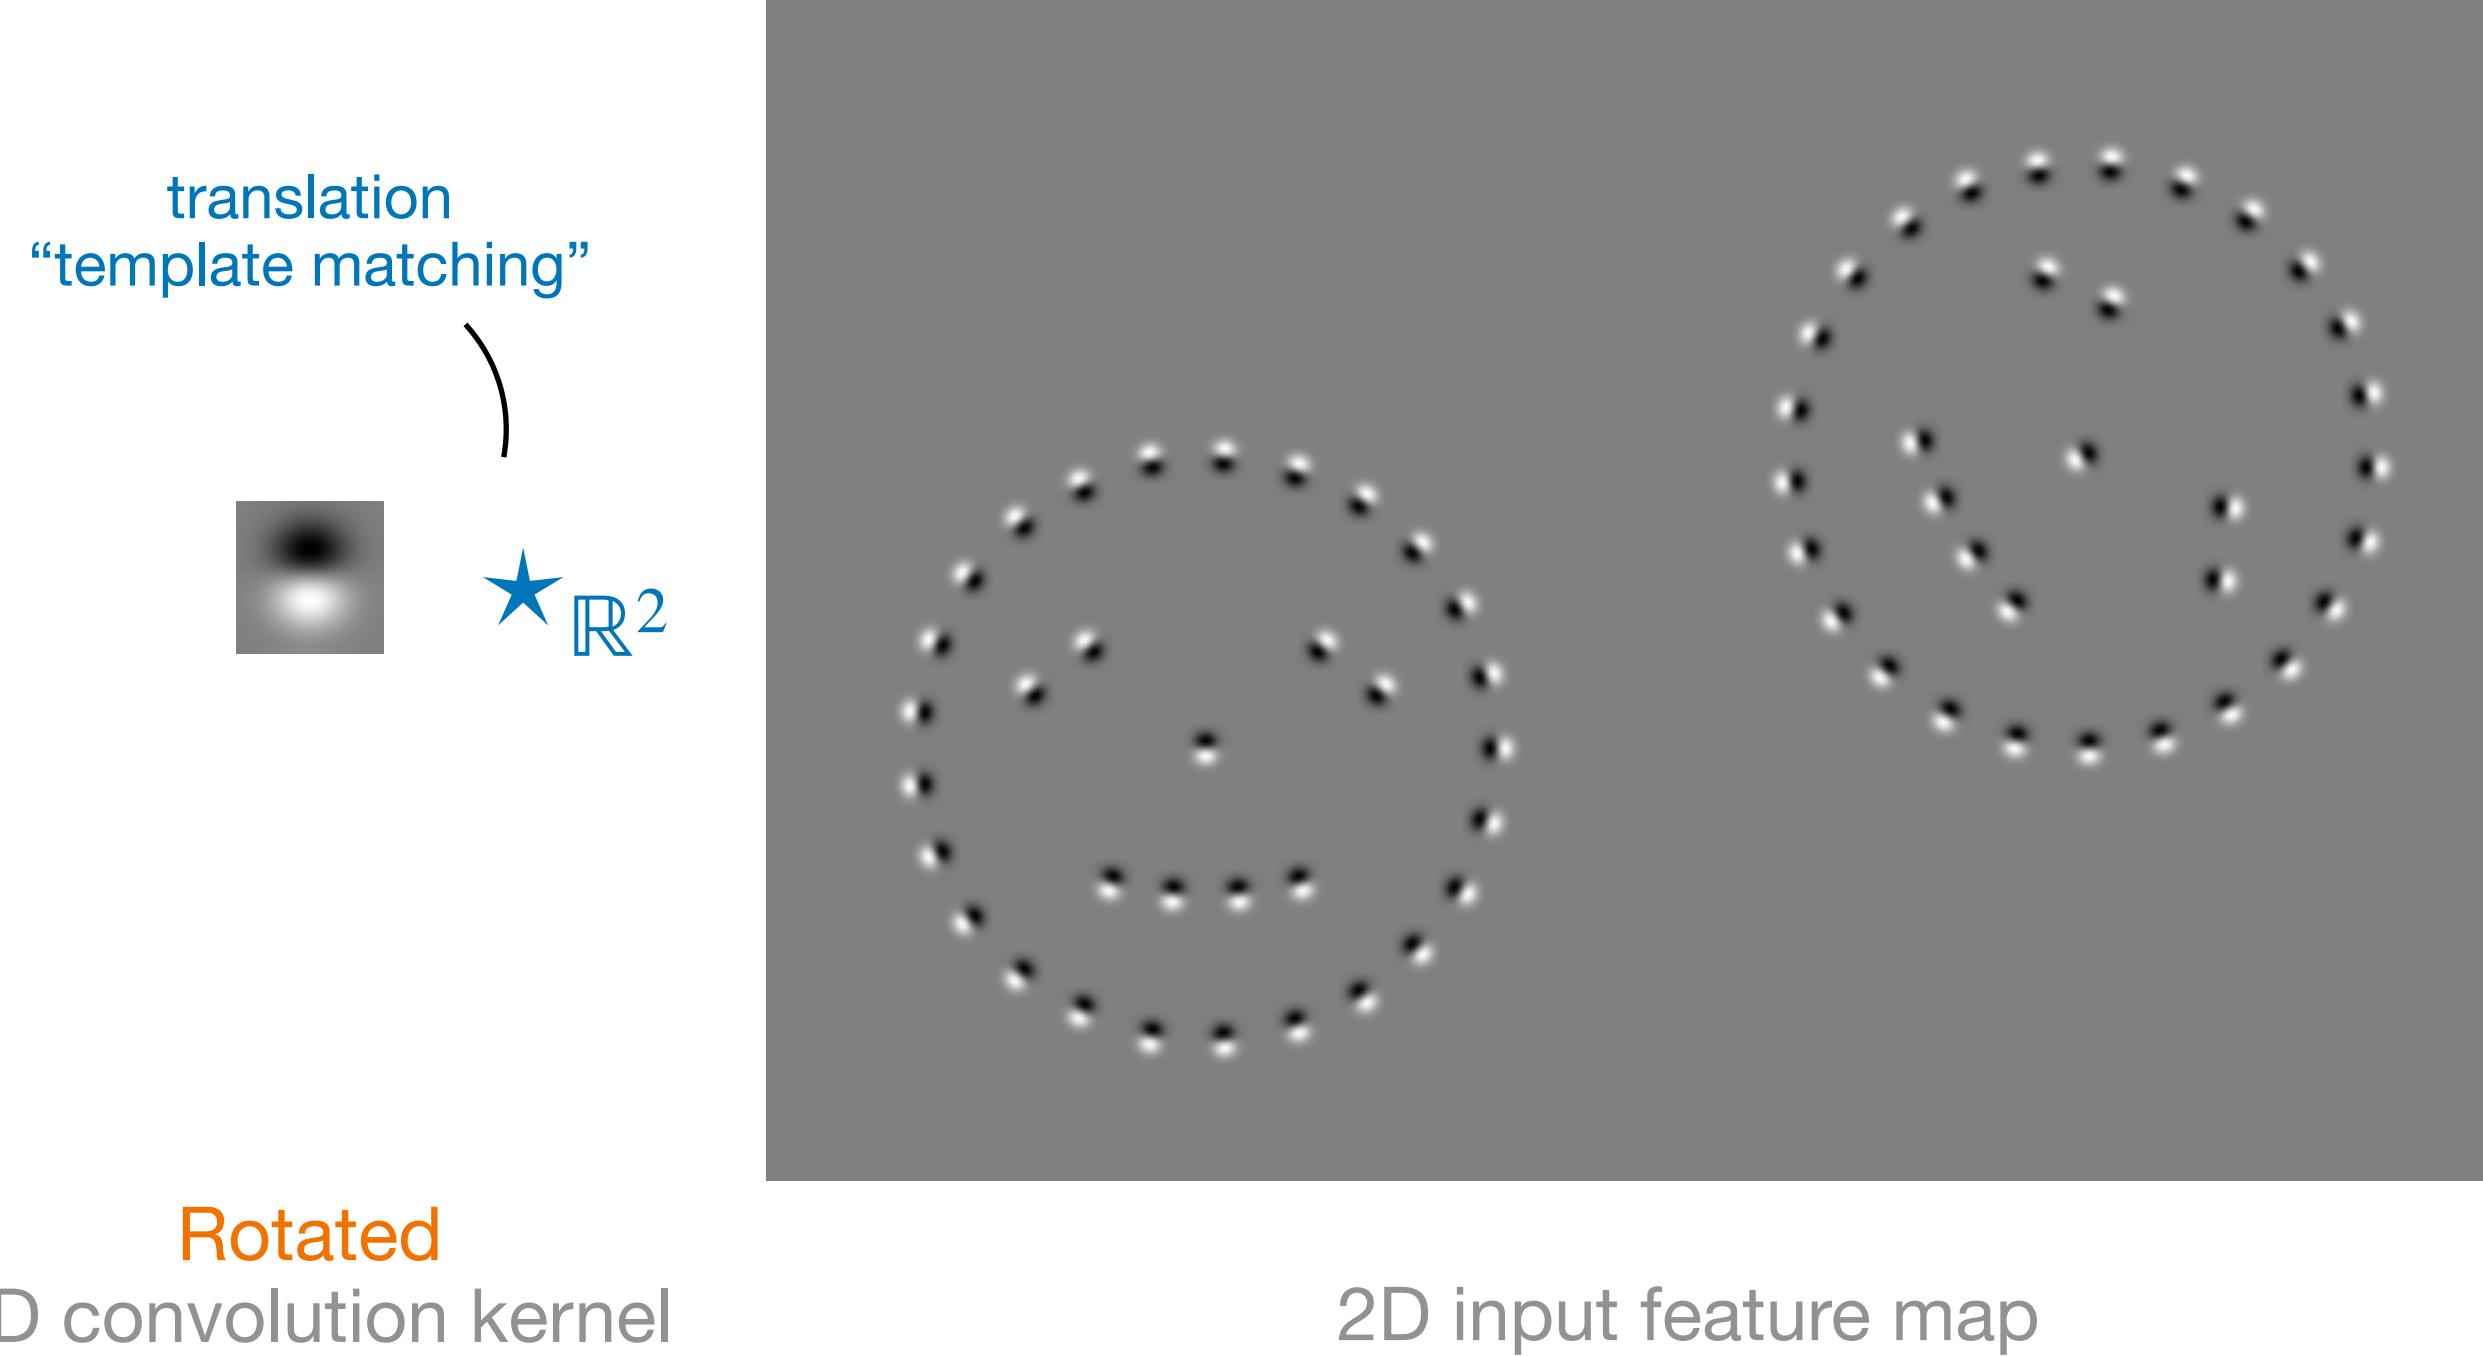
**$SE(2)$  output feature map**

# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  
 (e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

$$(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)} = (\mathcal{L}_{\mathbf{x}}^{(\mathbb{R}^d,+)\rightarrow \mathbb{L}_2(\mathbb{R}^d)} \mathcal{L}_h^{H \rightarrow \mathbb{L}_2(\mathbb{R}^d)} k, f)_{\mathbb{L}_2(\mathbb{R}^d)}$$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}') f(\mathbf{x}') d\mathbf{x}' = \int_{\mathbb{R}^2} k_h(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

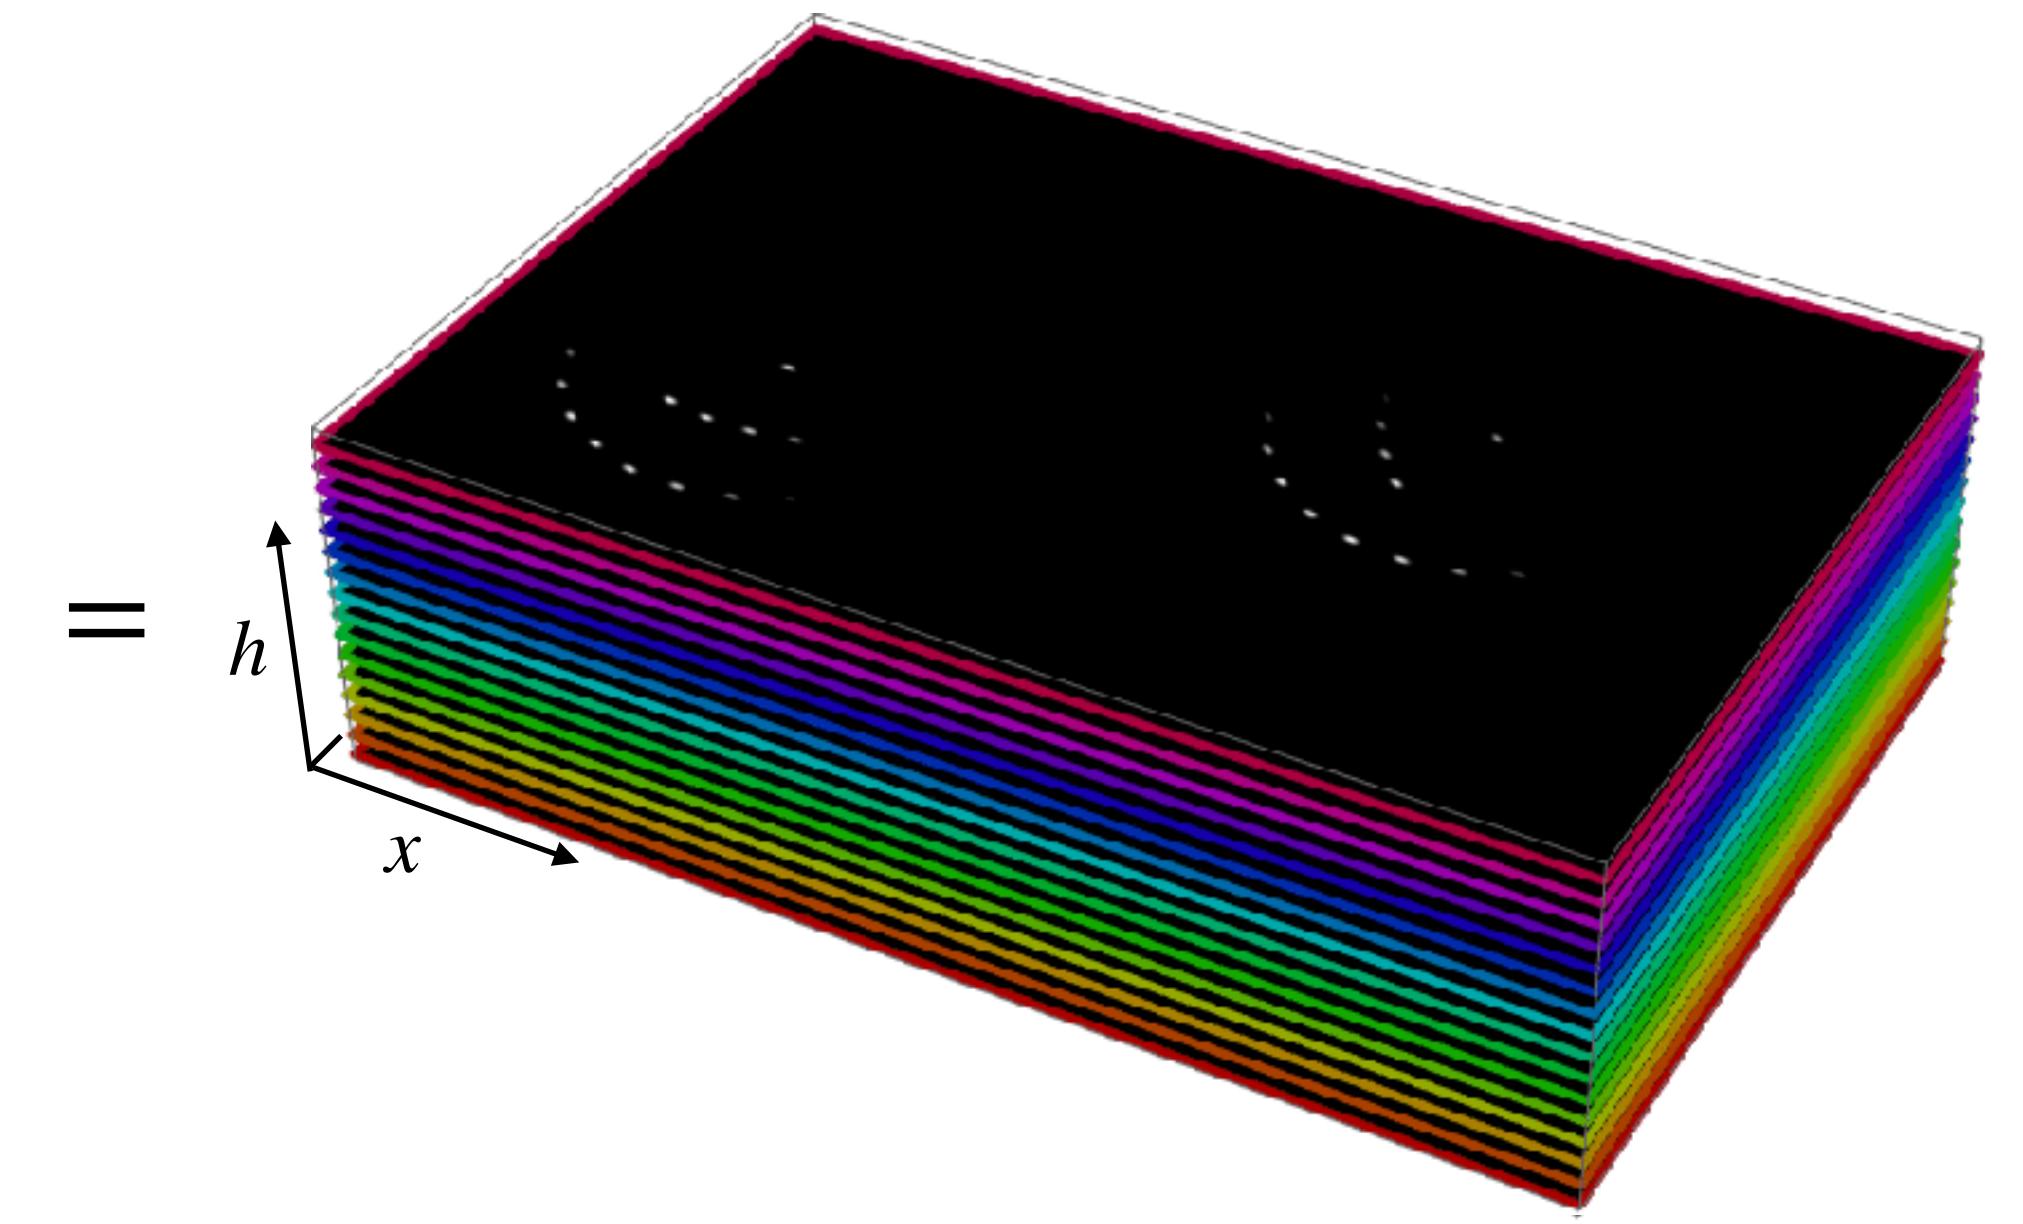
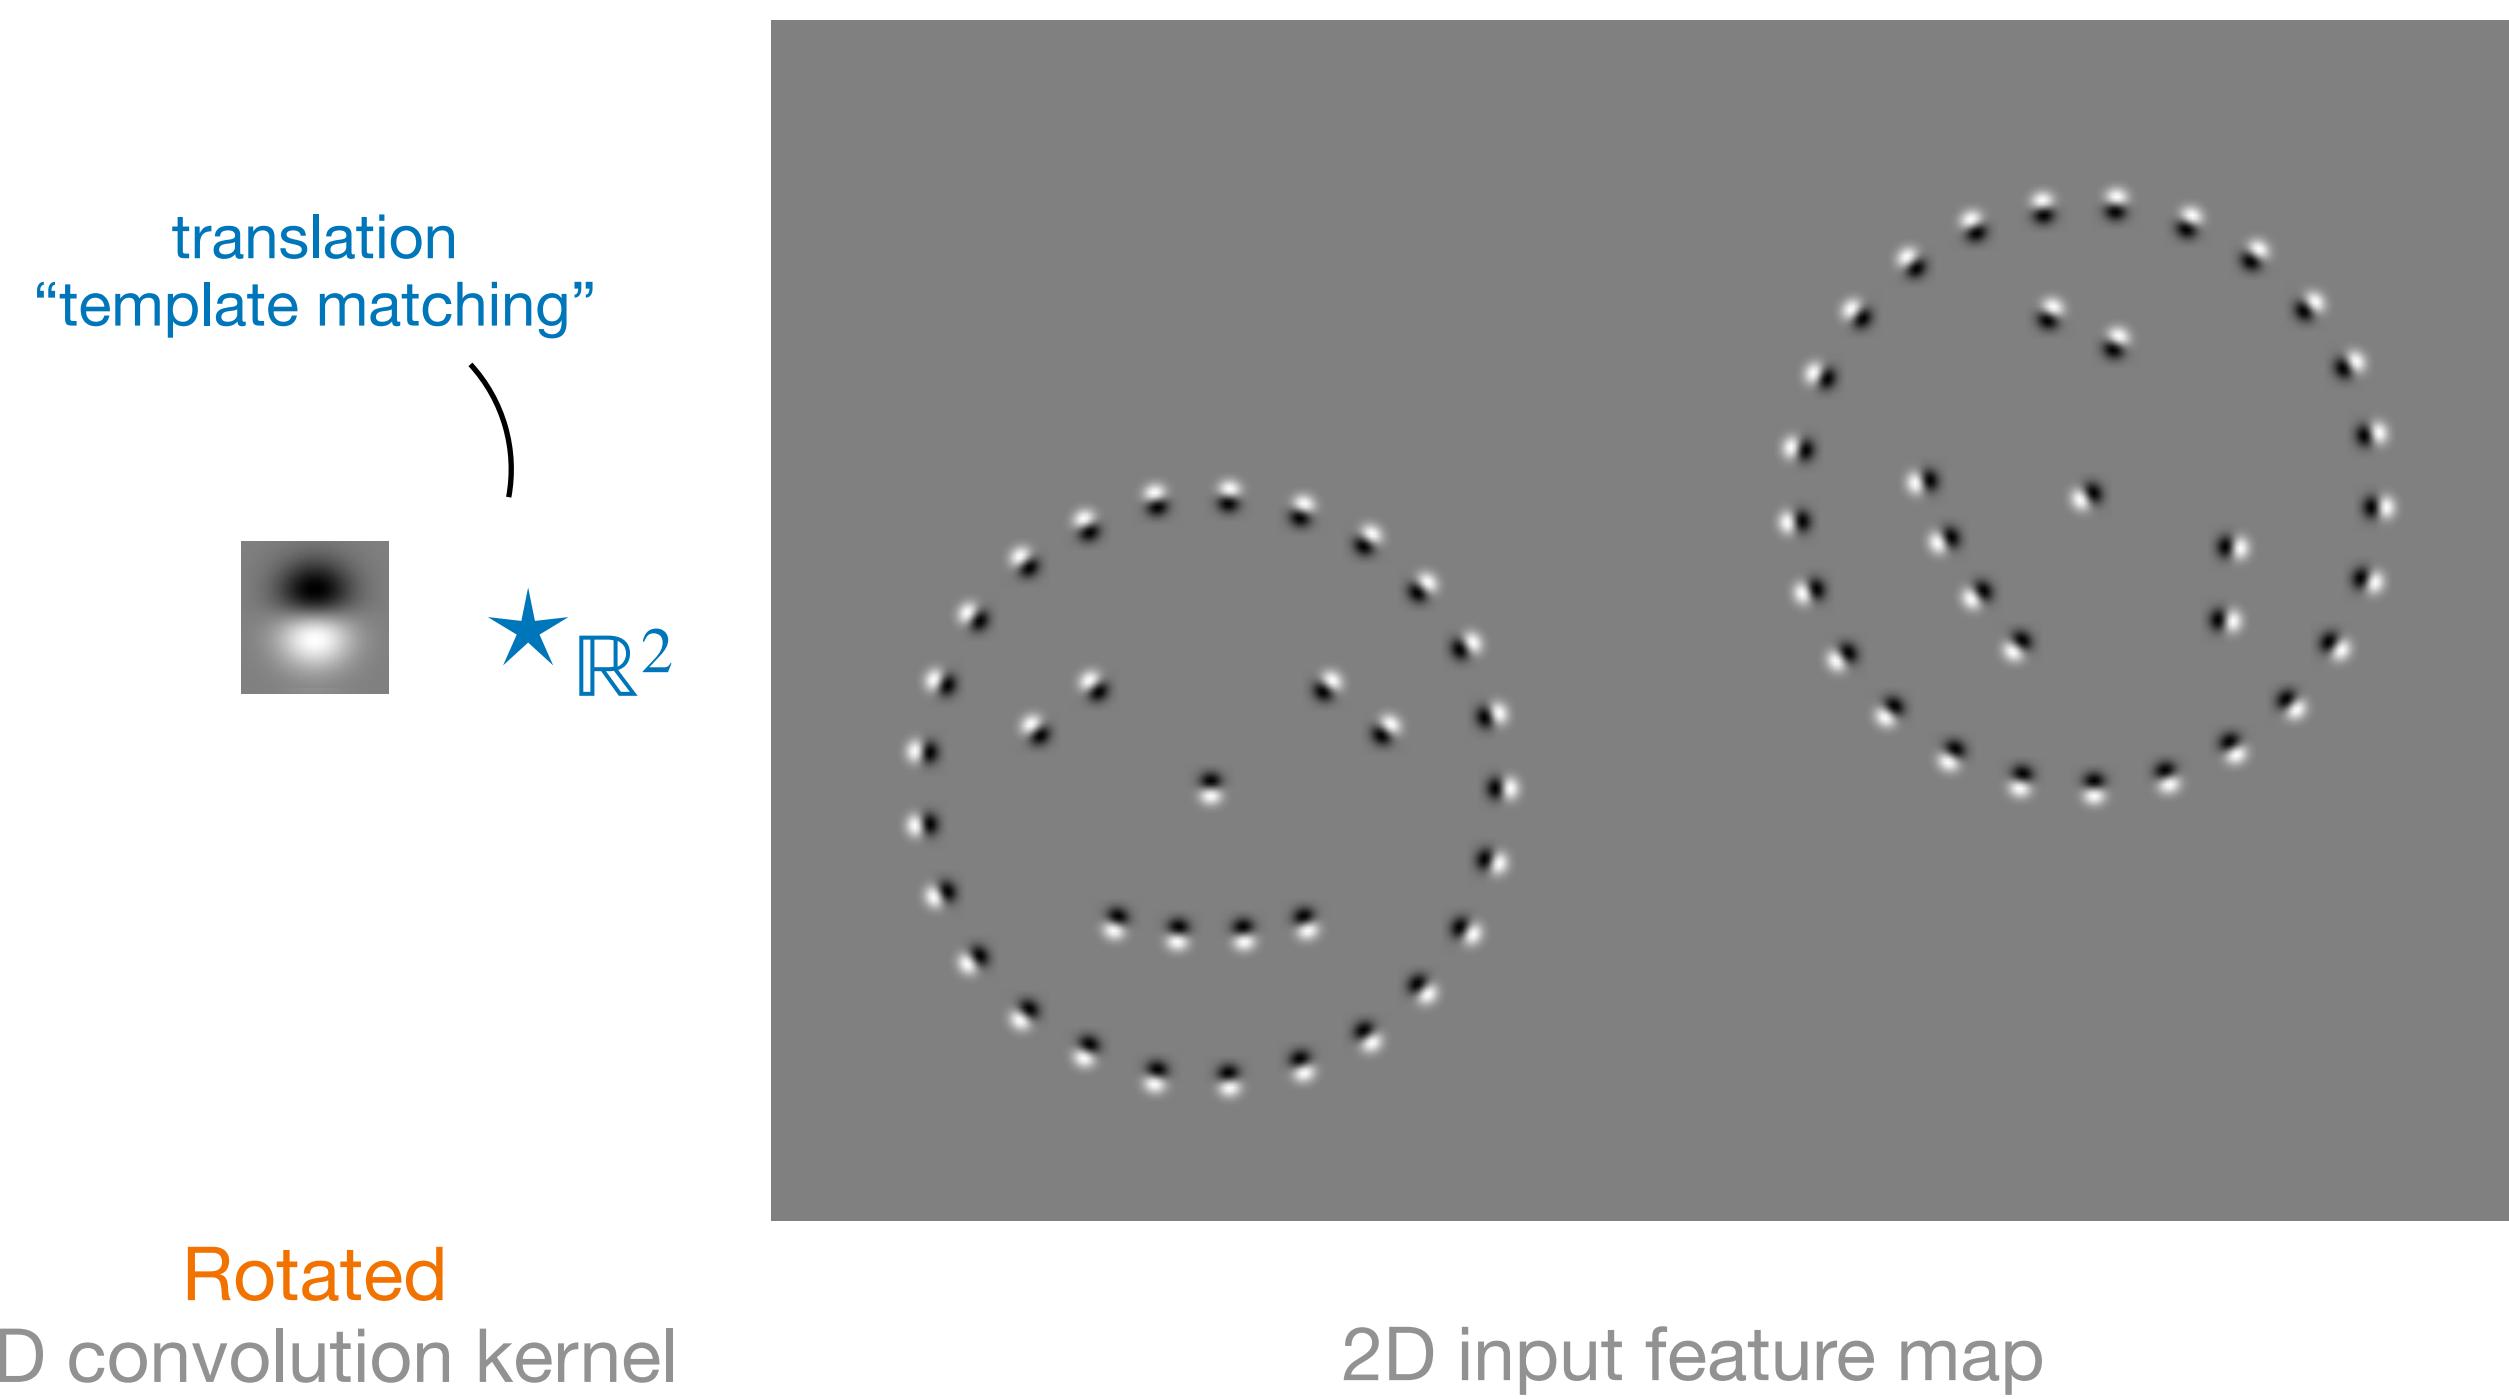


# Regular group convolutions revisited

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(g) = (\mathcal{L}_g^{G \rightarrow \mathbb{L}_2(X)} k, f)_{\mathbb{L}_2(X)} = (\mathcal{L}_{\mathbf{x}}^{(\mathbb{R}^d,+)} \rightarrow \mathbb{L}_2(\mathbb{R}^d) \mathcal{L}_h^{H \rightarrow \mathbb{L}_2(\mathbb{R}^d)} k, f)_{\mathbb{L}_2(\mathbb{R}^d)}$

(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

$$= \int_{\mathbb{R}^d} k(g^{-1}\mathbf{x}')f(\mathbf{x}')d\mathbf{x}' = \int_{\mathbb{R}^2} k_h(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$



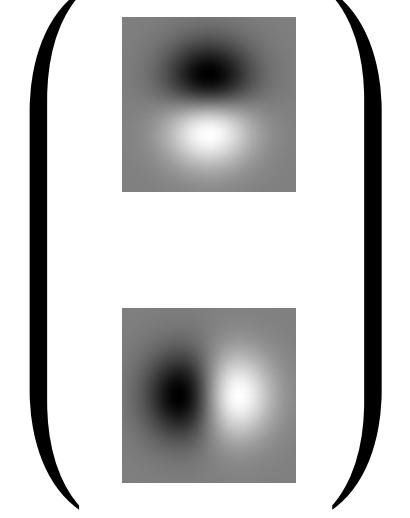
$SE(2)$  output feature map (after ReLU)

# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):       $(k \tilde{\star} f)(\mathbf{x}, \textcolor{orange}{h}) = \int_{\mathbb{R}^d} k(\textcolor{orange}{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$

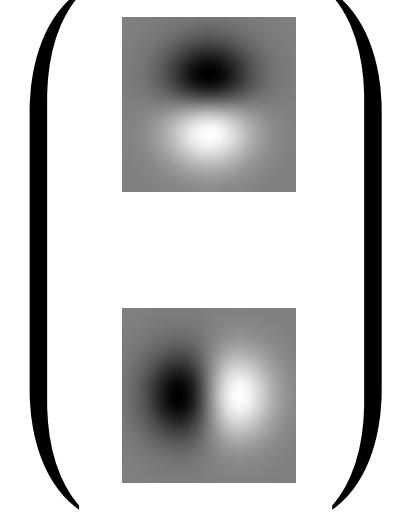
# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$

$$k(\mathbf{x} \mid \hat{\mathbf{w}}) = \hat{\mathbf{w}}^\dagger Y(\mathbf{x})$$


# Lifting convolution with steerable kernel

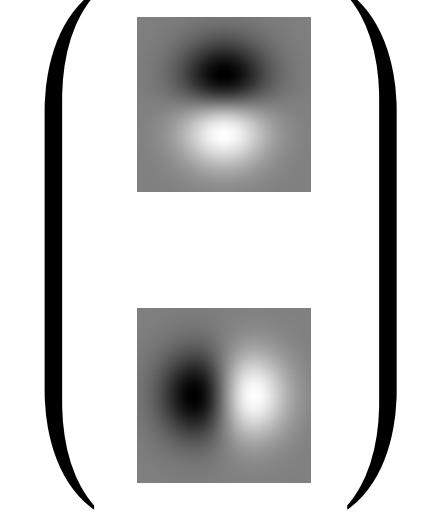
Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x})$$


# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

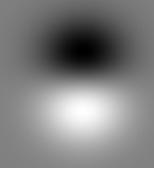
$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$
$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \quad Y(\mathbf{x})$$


# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$
$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

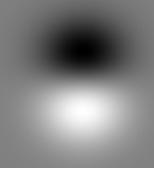
$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \begin{pmatrix} Y(\mathbf{x}) \\ \vdots \\ Y(\mathbf{x}) \end{pmatrix}$$


$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$
$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \begin{pmatrix} Y(\mathbf{x}) \\ \vdots \\ Y(\mathbf{x}) \end{pmatrix}$$


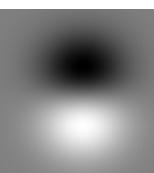
$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$
$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$

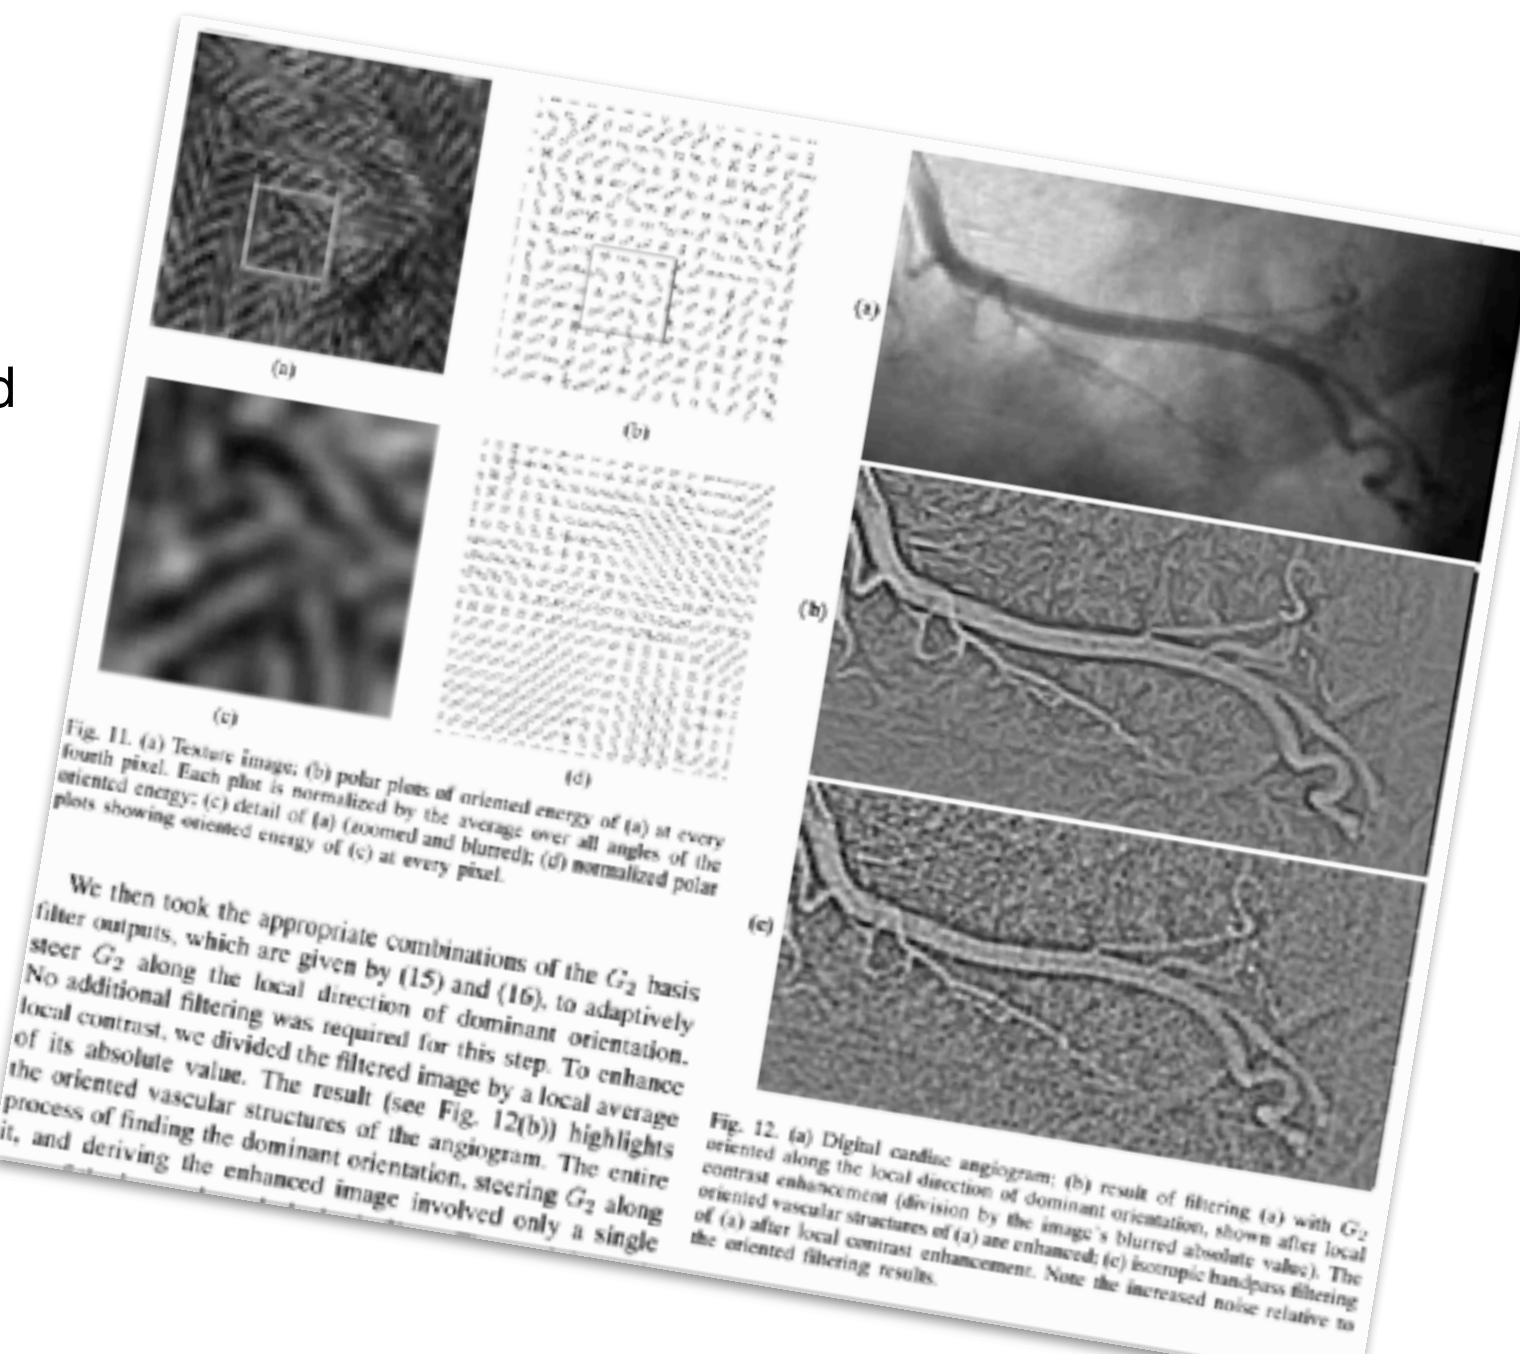
$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \begin{pmatrix} Y(\mathbf{x}) \\ \vdots \\ Y(\mathbf{x}) \end{pmatrix}$$


$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

- ! Freeman, W. T., & Adelson, E. H. (1991). The design and
- use of steerable filters. IEEE Transactions on Pattern analysis and machine intelligence, 13(9), 891-906.



# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ): 
$$(k \tilde{\star} f)(\mathbf{x}, \theta) = (\rho(\mathbf{R}_\theta)\hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$
  
(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

---

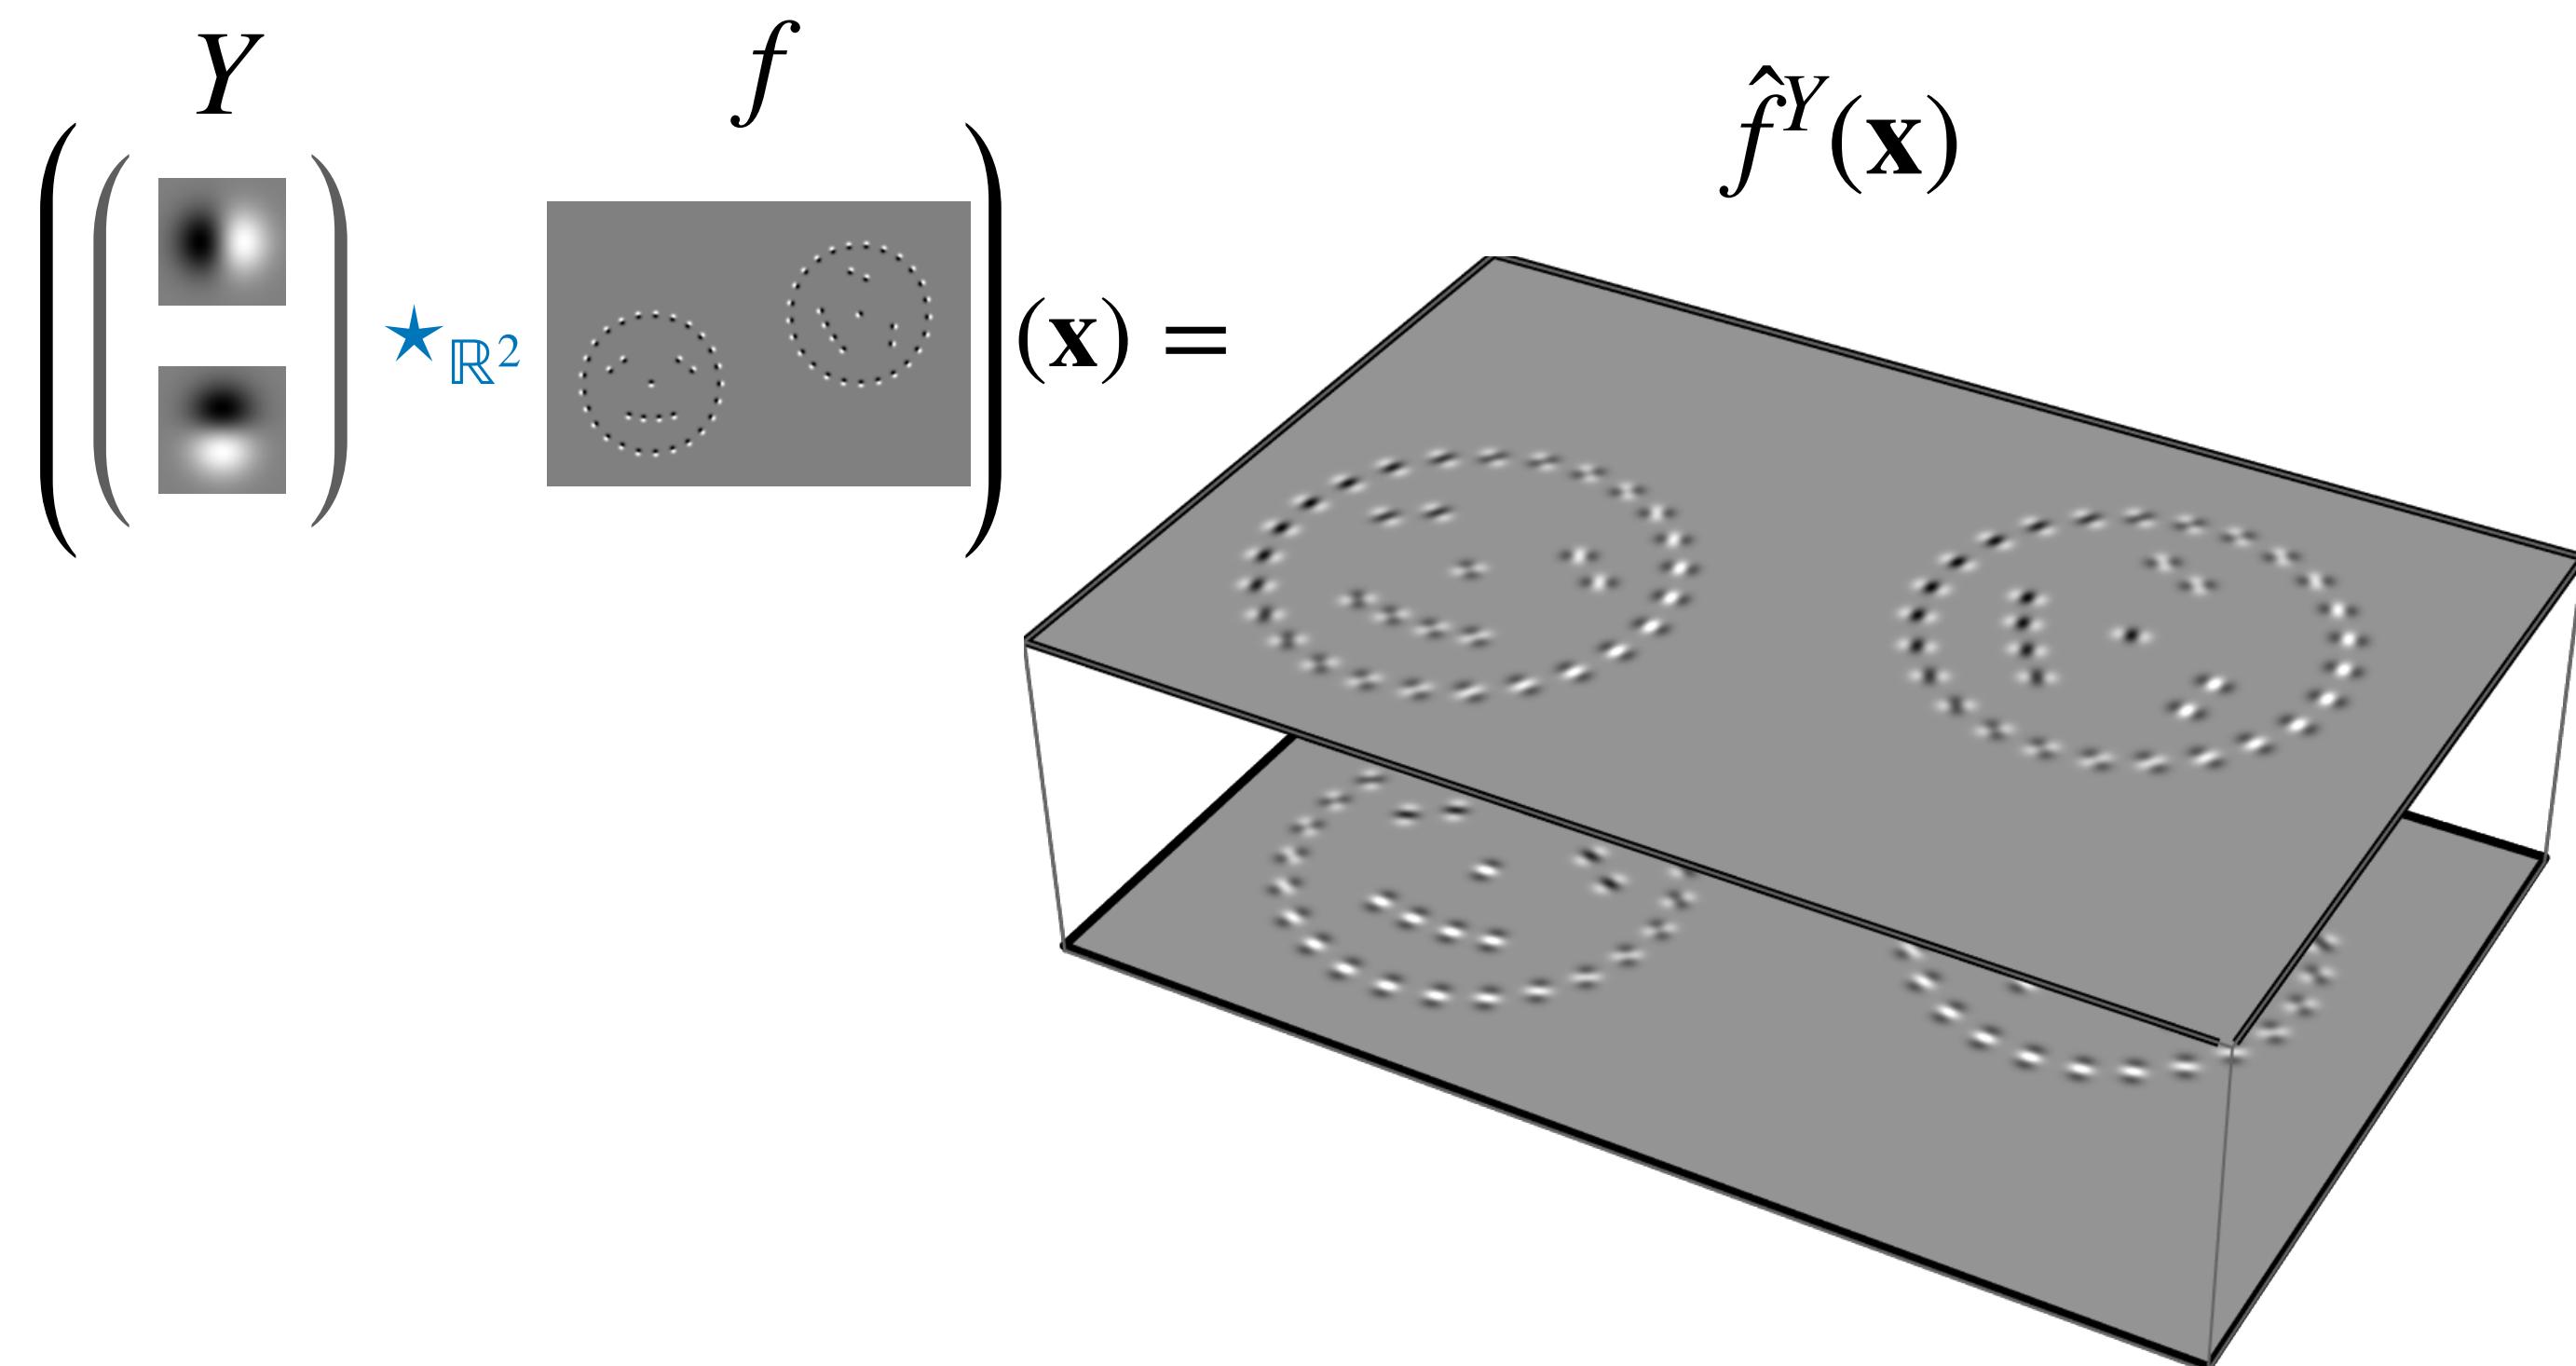
# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \theta) = (\rho(\mathbf{R}_\theta) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

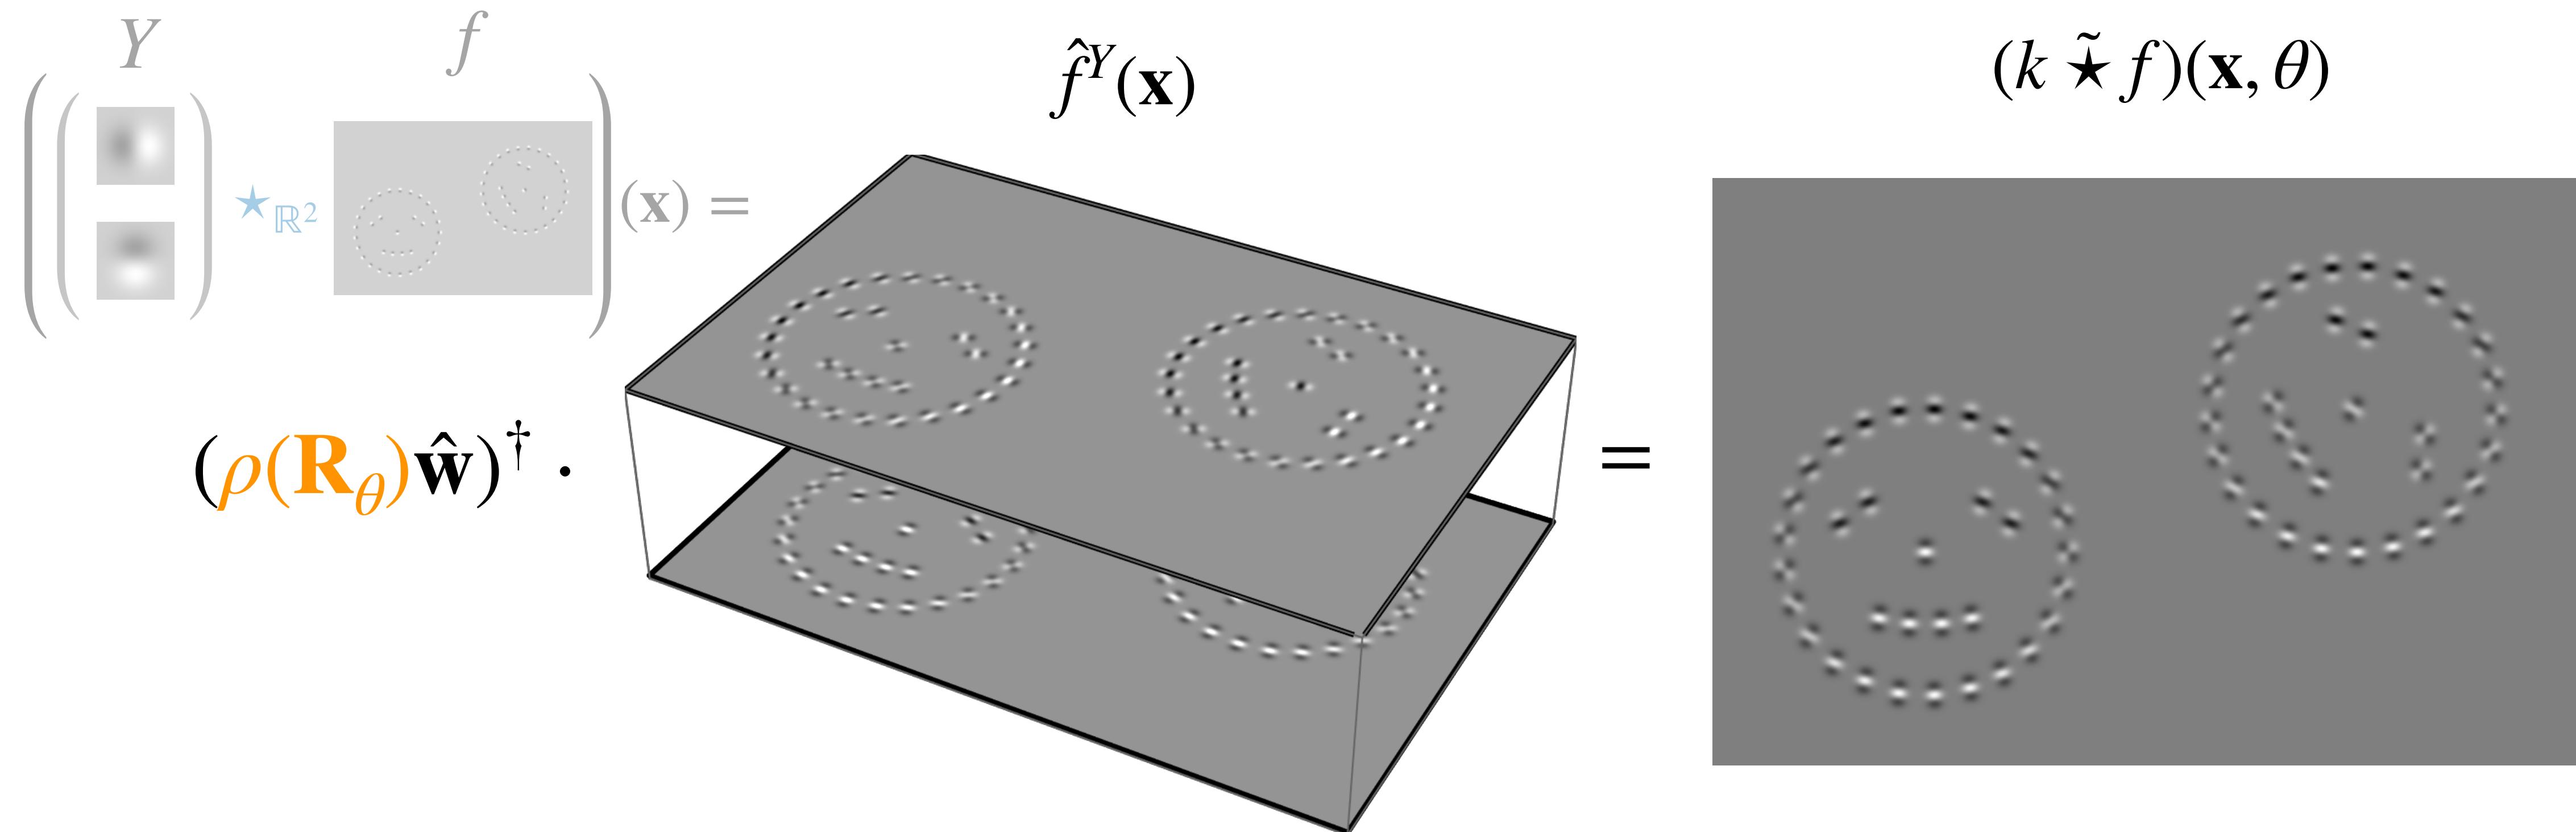
---



# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):  $(k \tilde{\star} f)(\mathbf{x}, \theta) = (\rho(\mathbf{R}_\theta)\hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$   
(e.g.  $G = SE(2) = \mathbb{R}^2 \rtimes SO(2)$ )  
 $X = \mathbb{R}^2$

---



# Lifting convolution with steerable kernel

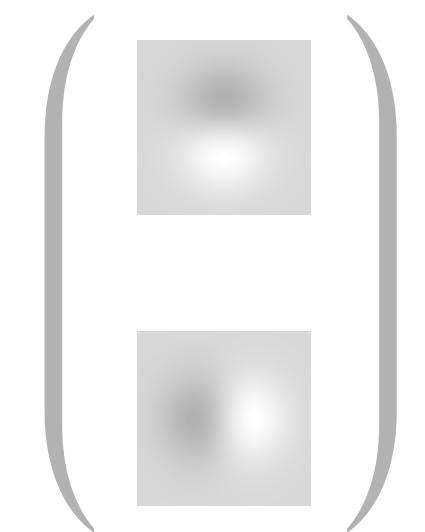
Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$

$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}^{-1}) \hat{\mathbf{w}})^T Y(\mathbf{x})$$


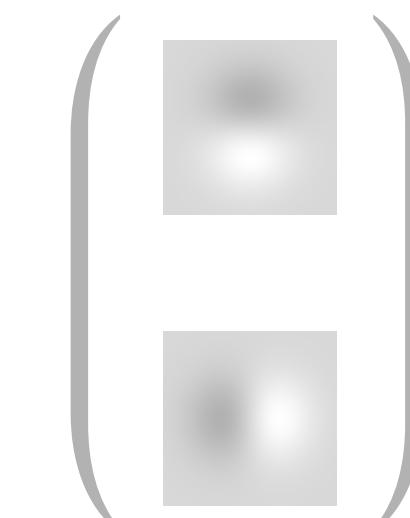
- ! Freeman, W. T., & Adelson, E. H. (1991). The design and use of steerable filters. IEEE Transactions on Pattern analysis and machine intelligence, 13(9), 891-906.

# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$

$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}^{-1}) \hat{\mathbf{w}})^T Y(\mathbf{x})$$


$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

- ! Freeman, W. T., & Adelson, E. H. (1991). The design and use of steerable filters. IEEE Transactions on Pattern analysis and machine intelligence, 13(9), 891-906.

# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$

$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

$$= \text{tr}( \hat{f}^Y(\mathbf{x}) \hat{\mathbf{w}}^\dagger \rho(\mathbf{h}^{-1}) )$$

$$\mathbf{a}^T \mathbf{b} = \text{tr}(\mathbf{b} \mathbf{a}^T) \quad \text{and} \quad \rho(h)^\dagger = \rho(h^{-1})$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}^{-1}) \hat{\mathbf{w}})^T Y(\mathbf{x})$$


- ! Freeman, W. T., & Adelson, E. H. (1991). The design and use of steerable filters. IEEE Transactions on Pattern analysis and machine intelligence, 13(9), 891-906.

# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$

$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}^{-1}) \hat{\mathbf{w}})^T Y(\mathbf{x})$$


$$\begin{pmatrix} \text{blurred gray square} \\ \text{sharper gray square} \\ \text{sharp gray square} \end{pmatrix}$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

$$= \text{tr}( \hat{f}^Y(\mathbf{x}) \hat{\mathbf{w}}^\dagger \rho(\mathbf{h}^{-1}) )$$

$$= \text{tr}( \hat{f}(\mathbf{x}) \rho(\mathbf{h}^{-1}) )$$

- ! Freeman, W. T., & Adelson, E. H. (1991). The design and use of steerable filters. IEEE Transactions on Pattern analysis and machine intelligence, 13(9), 891-906.

$$\mathbf{a}^T \mathbf{b} = \text{tr}(\mathbf{b} \mathbf{a}^T) \quad \text{and} \quad \rho(h)^\dagger = \rho(h^{-1})$$

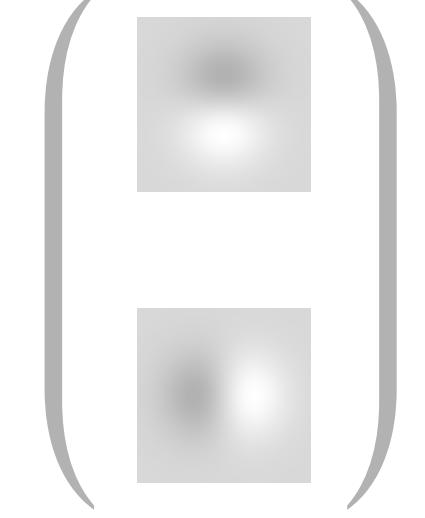
$$\hat{f}(\mathbf{x}) = \hat{f}^Y(\mathbf{x}) \hat{\mathbf{w}}^\dagger$$

# Lifting convolution with steerable kernel

Group convolution ( $G = \mathbb{R}^d \rtimes H$ ):

$$(k \tilde{\star} f)(\mathbf{x}, \mathbf{h}) = \int_{\mathbb{R}^d} k(\mathbf{h}^{-1}(\mathbf{x}' - \mathbf{x}) \mid \hat{\mathbf{w}}) f(\mathbf{x}') d\mathbf{x}'$$

$$= \int_{\mathbb{R}^d} (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$k(\mathbf{h}^{-1} \mathbf{x} \mid \hat{\mathbf{w}}) = (\rho(\mathbf{h}^{-1}) \hat{\mathbf{w}})^T Y(\mathbf{x})$$


$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \int_{\mathbb{R}^d} Y(\mathbf{x}' - \mathbf{x}) f(\mathbf{x}') d\mathbf{x}'$$

$$= (\rho(\mathbf{h}) \hat{\mathbf{w}})^\dagger \hat{f}^Y(\mathbf{x})$$

$$= \text{tr}( \hat{f}^Y(\mathbf{x}) \hat{\mathbf{w}}^\dagger \rho(\mathbf{h}^{-1}) )$$

$$= \text{tr}( \hat{f}(\mathbf{x}) \rho(\mathbf{h}^{-1}) )$$

- ! Freeman, W. T., & Adelson, E. H. (1991). The design and use of steerable filters. IEEE Transactions on Pattern analysis and machine intelligence, 13(9), 891-906.

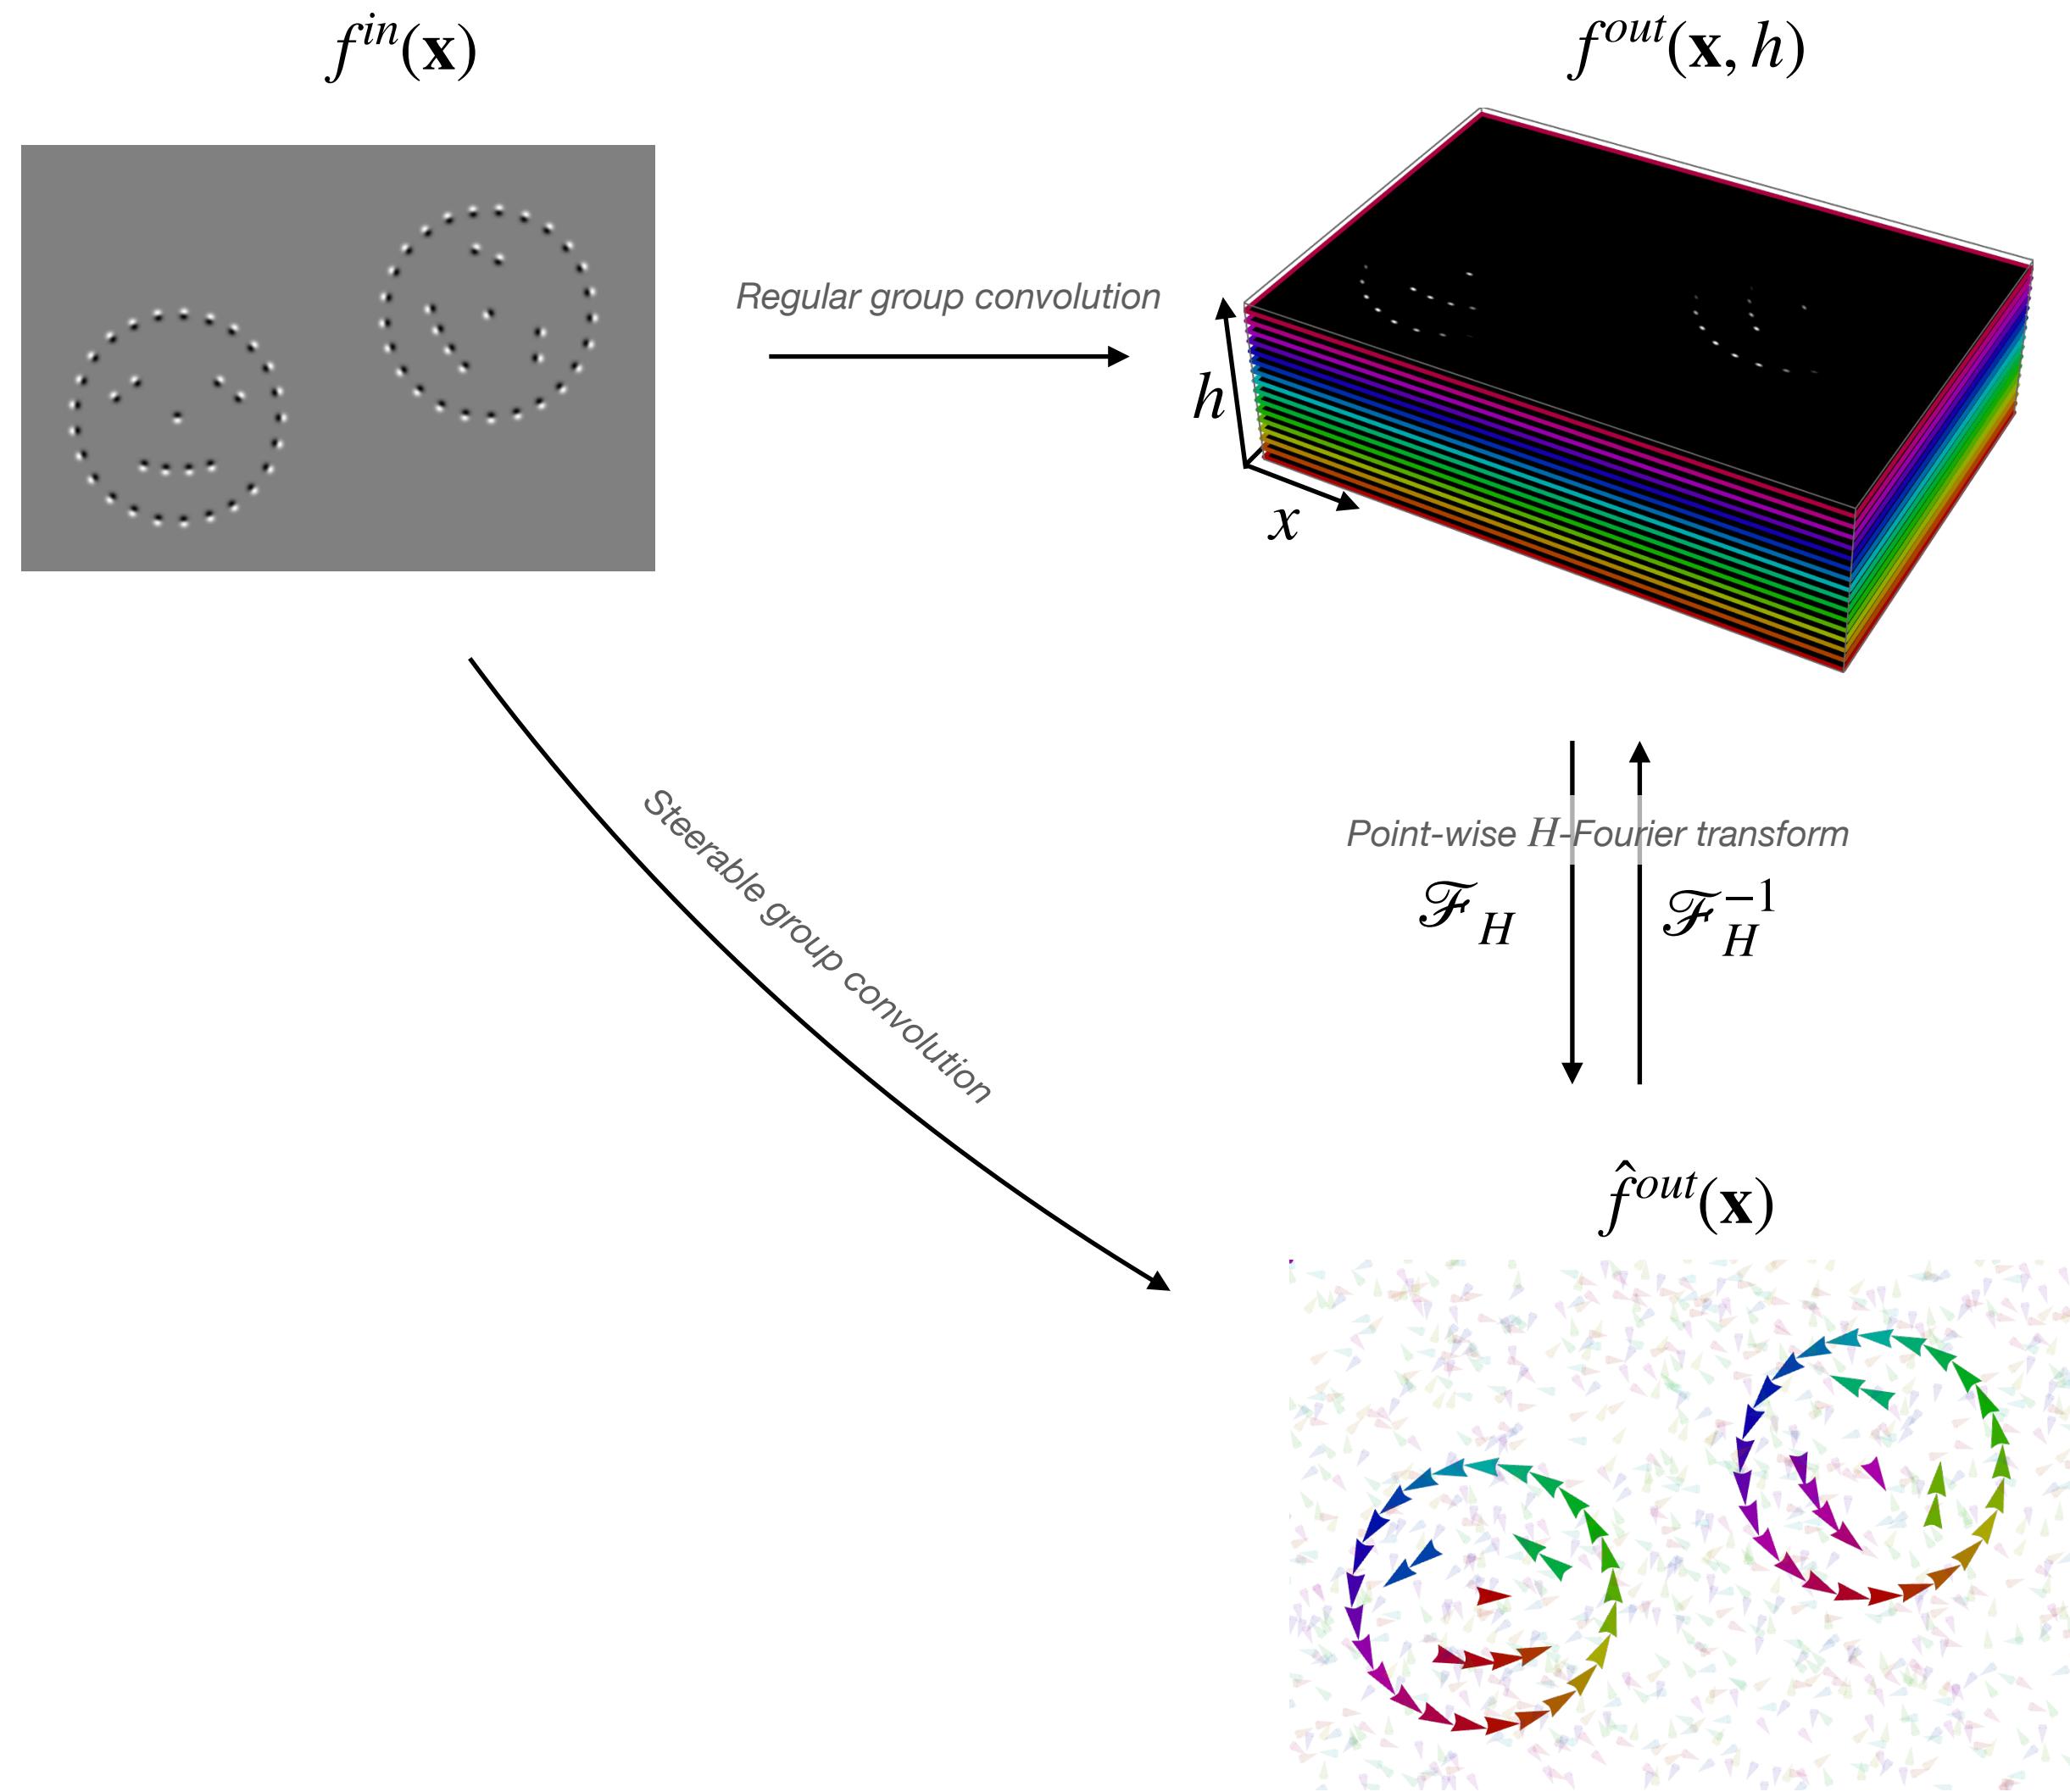
$$\mathbf{a}^T \mathbf{b} = \text{tr}(\mathbf{b} \mathbf{a}^T) \quad \text{and} \quad \rho(h)^\dagger = \rho(h^{-1})$$

$$\hat{f}(\mathbf{x}) = \hat{f}^Y(\mathbf{x}) \hat{\mathbf{w}}^\dagger$$

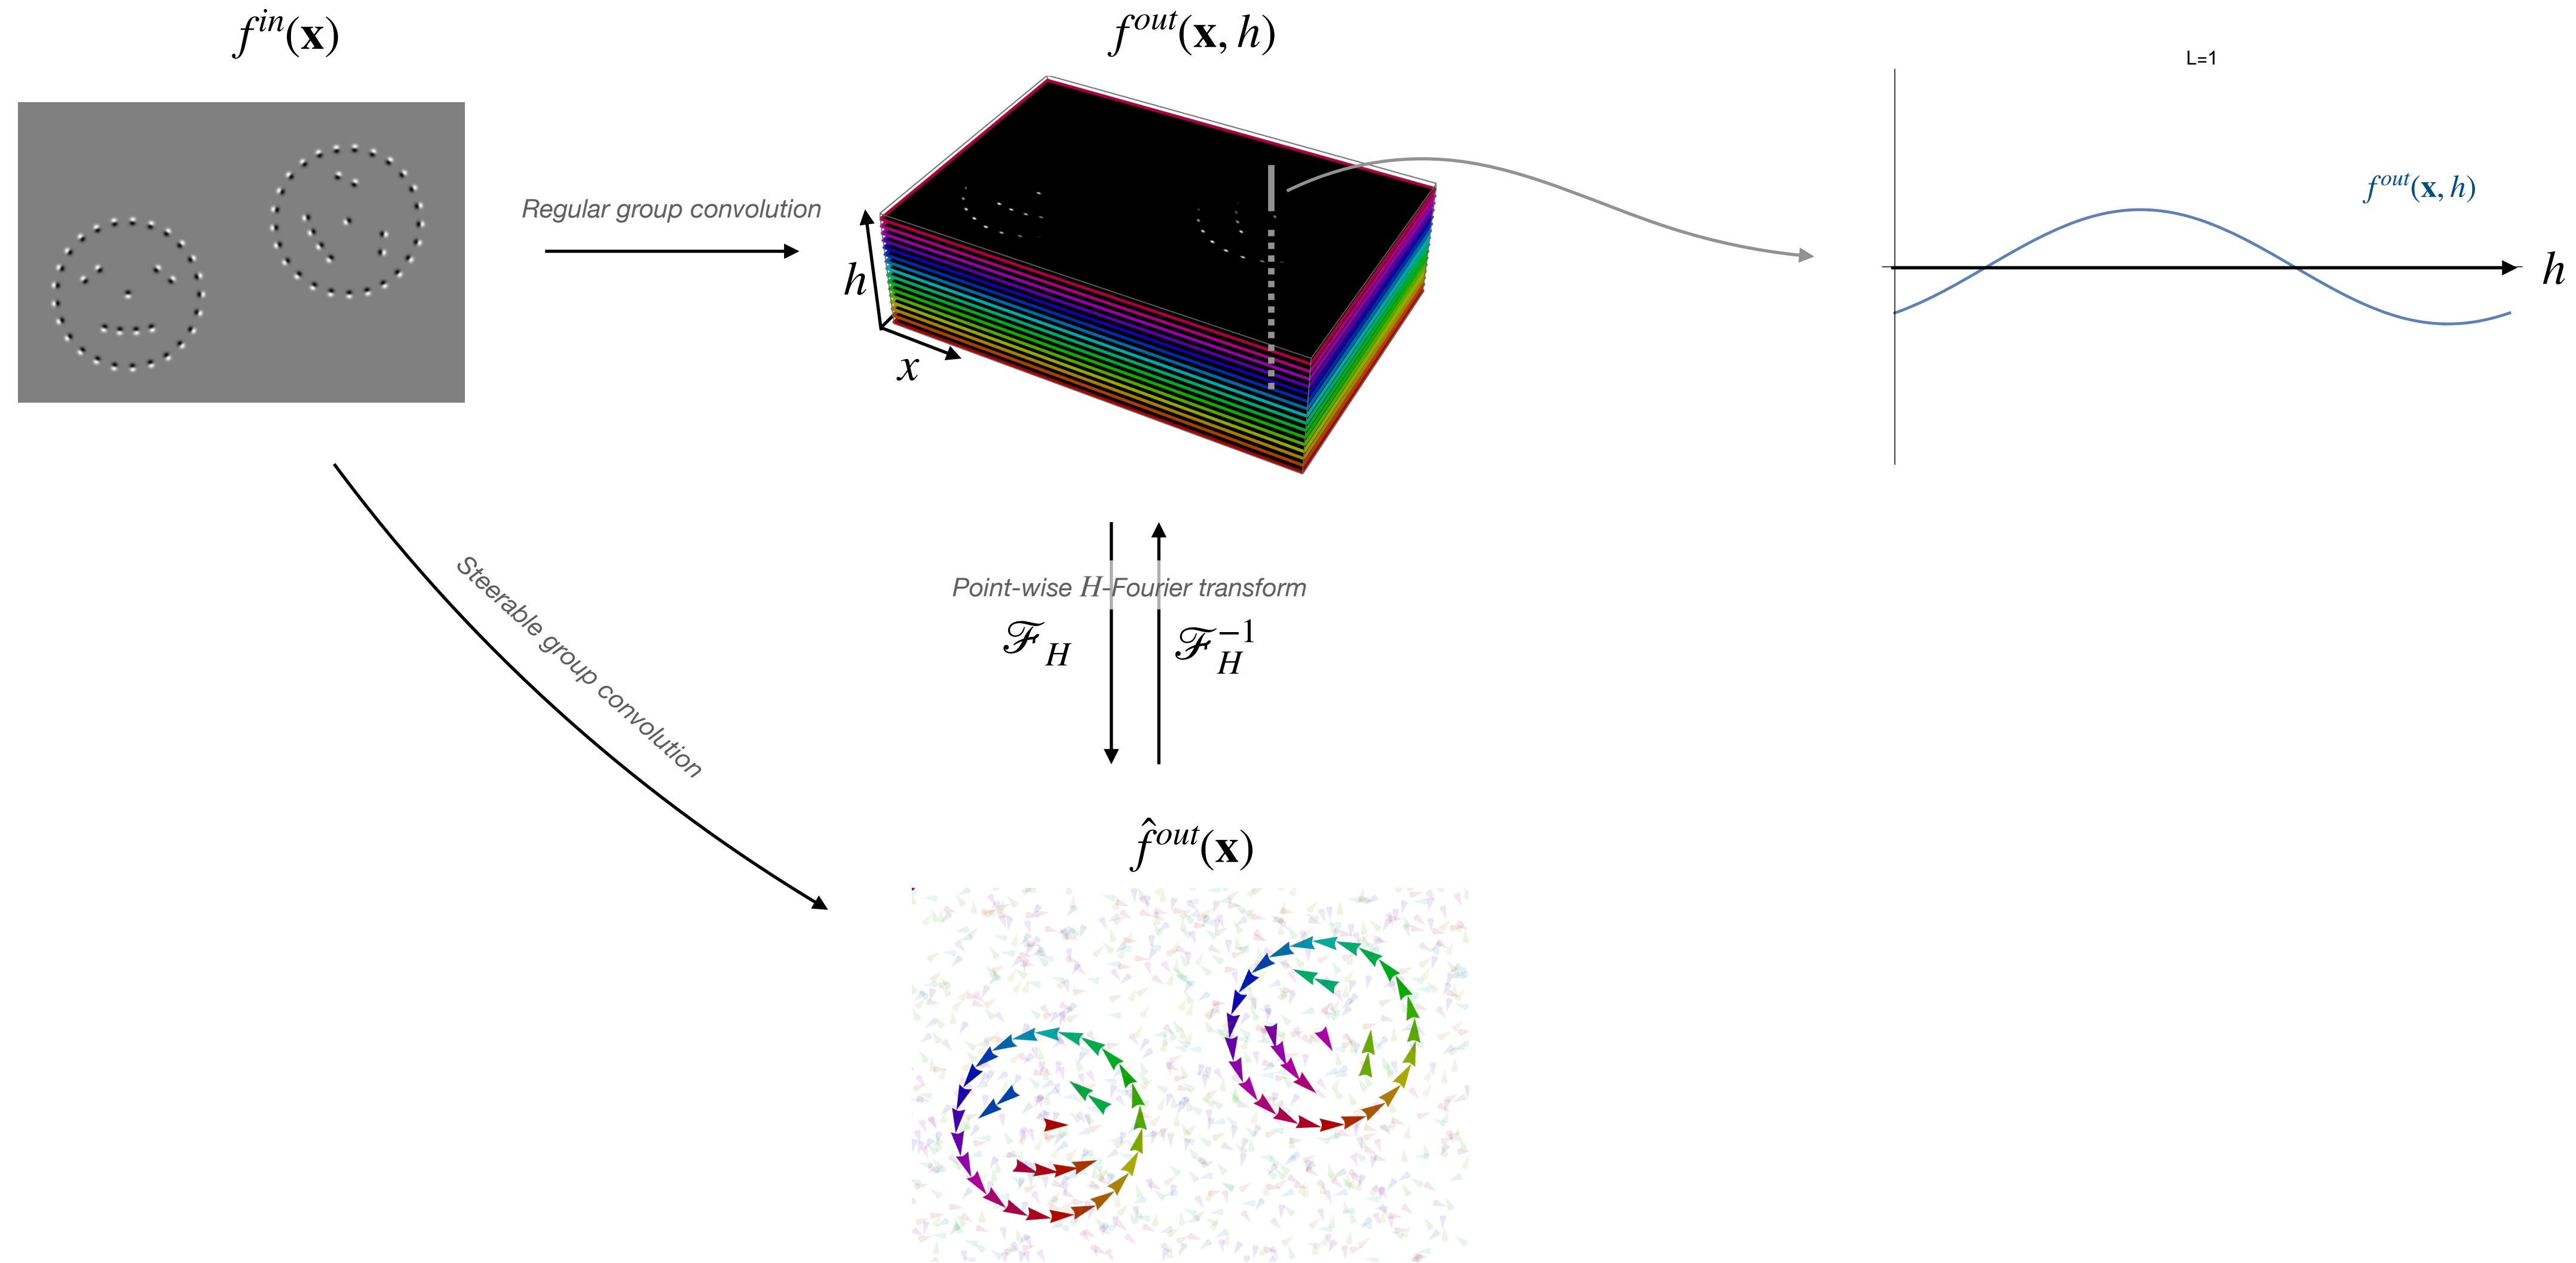
$$= \mathcal{F}_H^{-1}[ \hat{f}(\mathbf{x}) ](\mathbf{h})$$

Inverse  $H$ -Fourier transform!

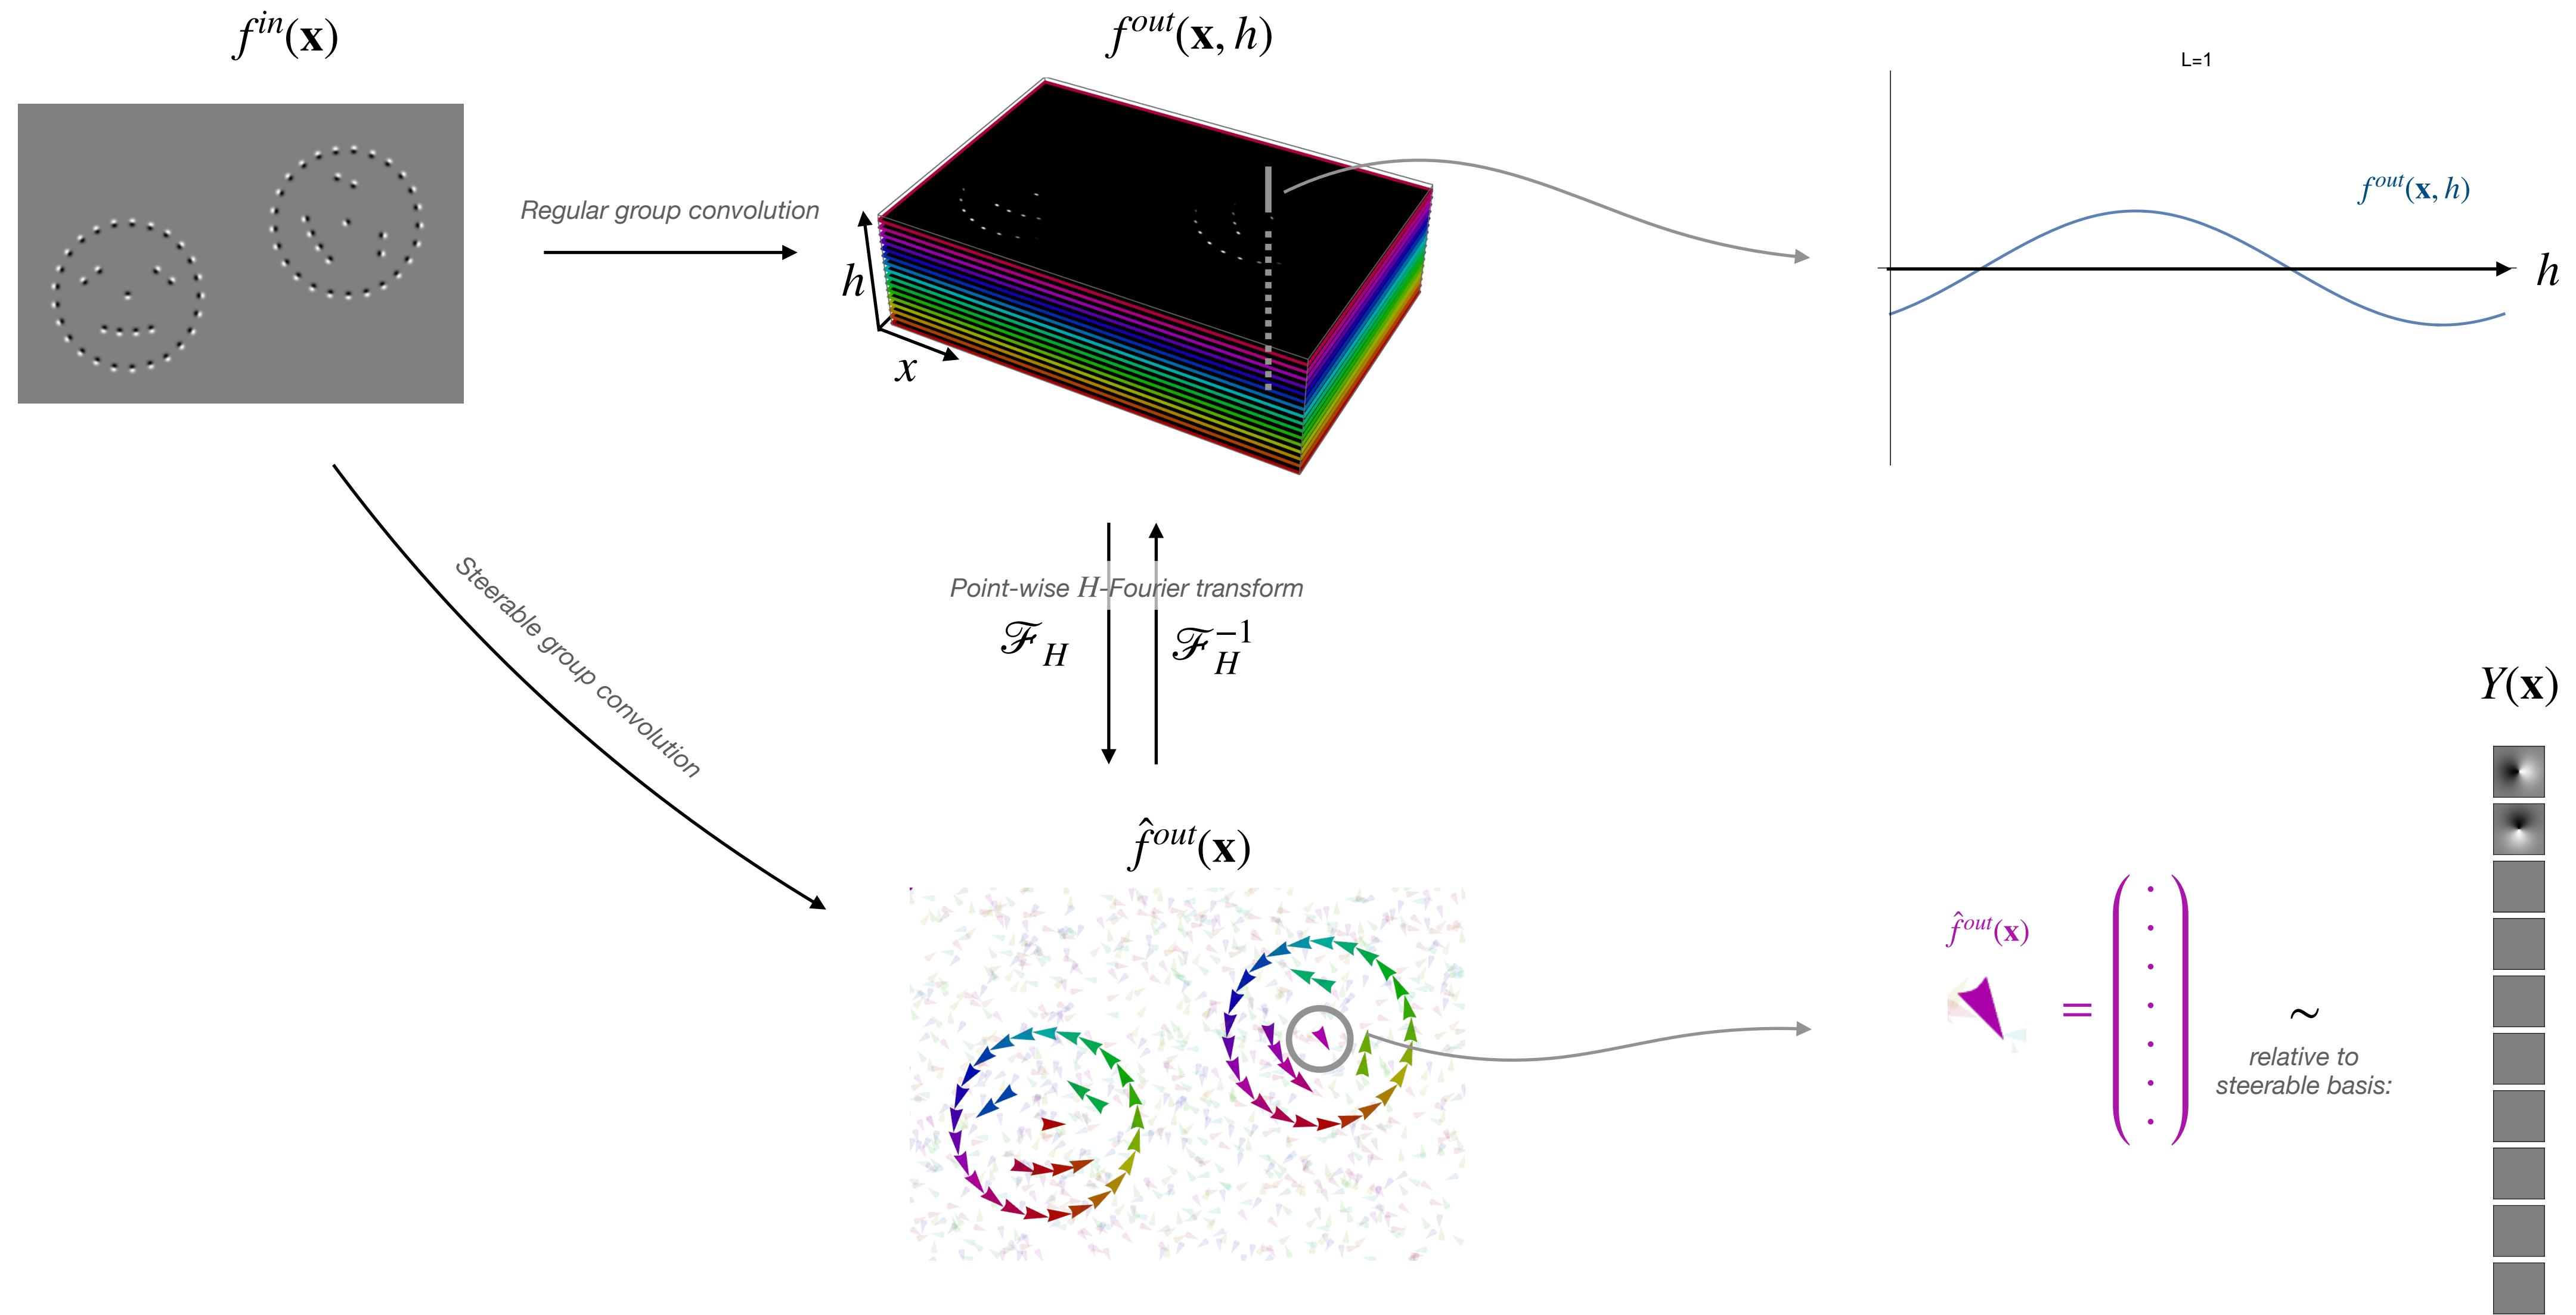
# From regular to steerable via a Fourier transform



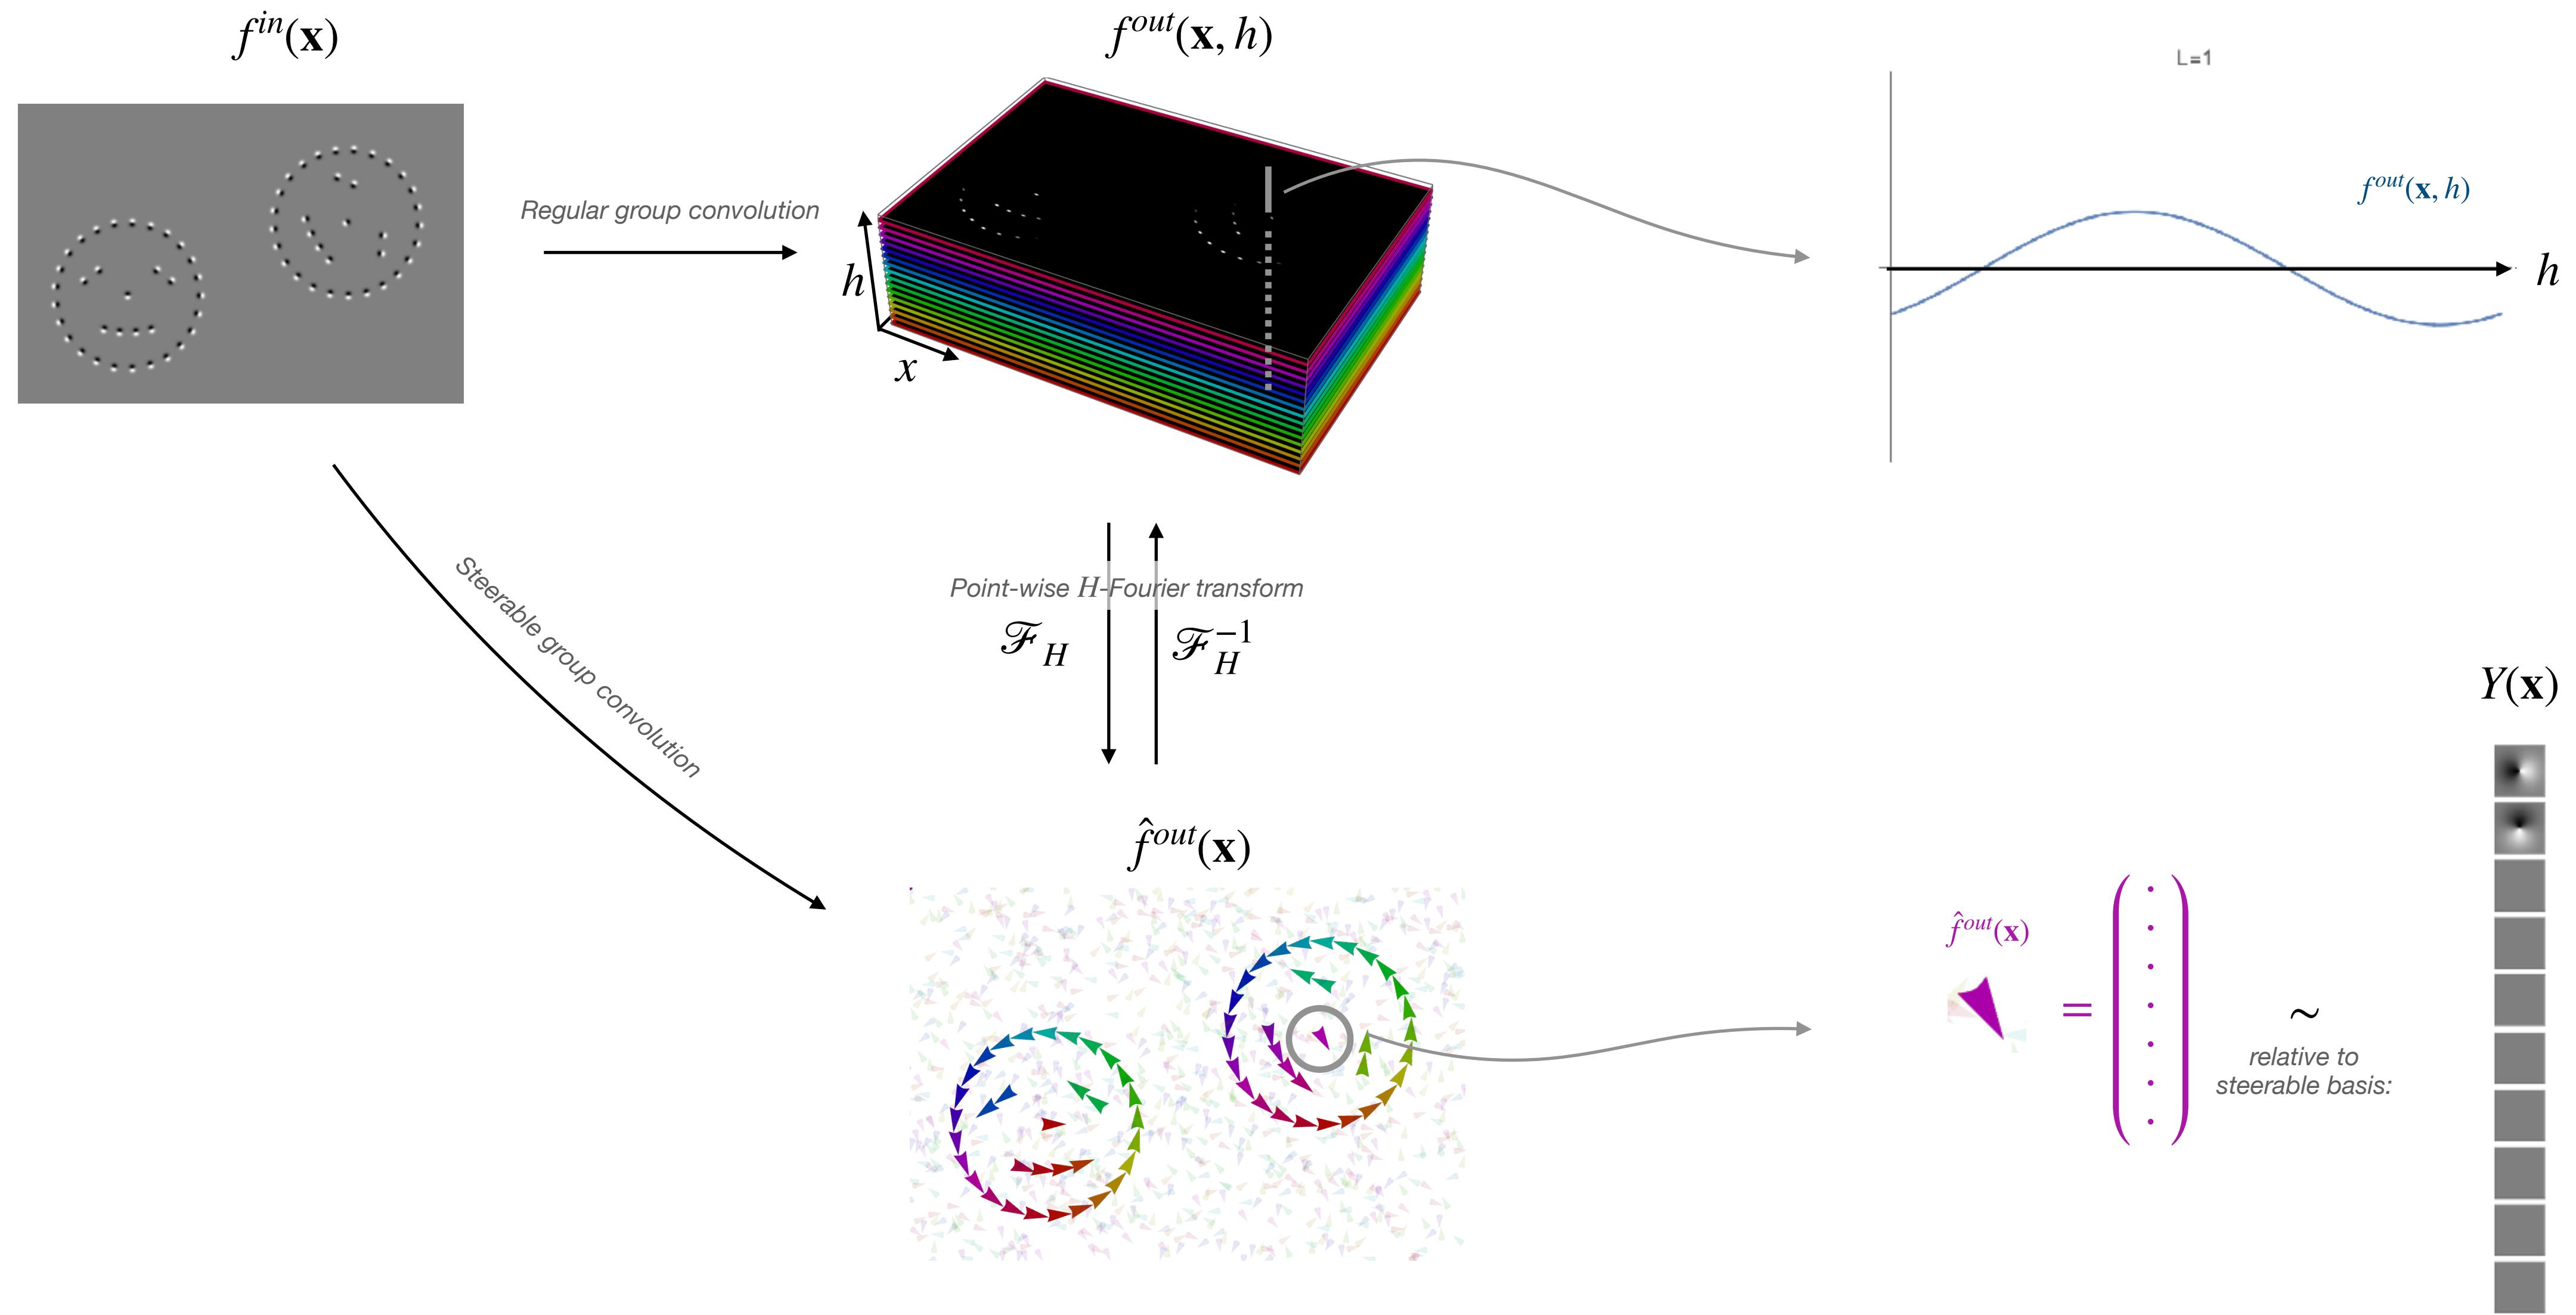
# From regular to steerable via a Fourier transform



# From regular to steerable via a Fourier transform



# From regular to steerable via a Fourier transform



# From regular to steerable via a Fourier transform

**Regular group convolutions:**  
Domain expanded feature maps

$$f^{(l)} : \mathbb{R}^d \times \mathbf{H} \rightarrow \mathbb{R}$$

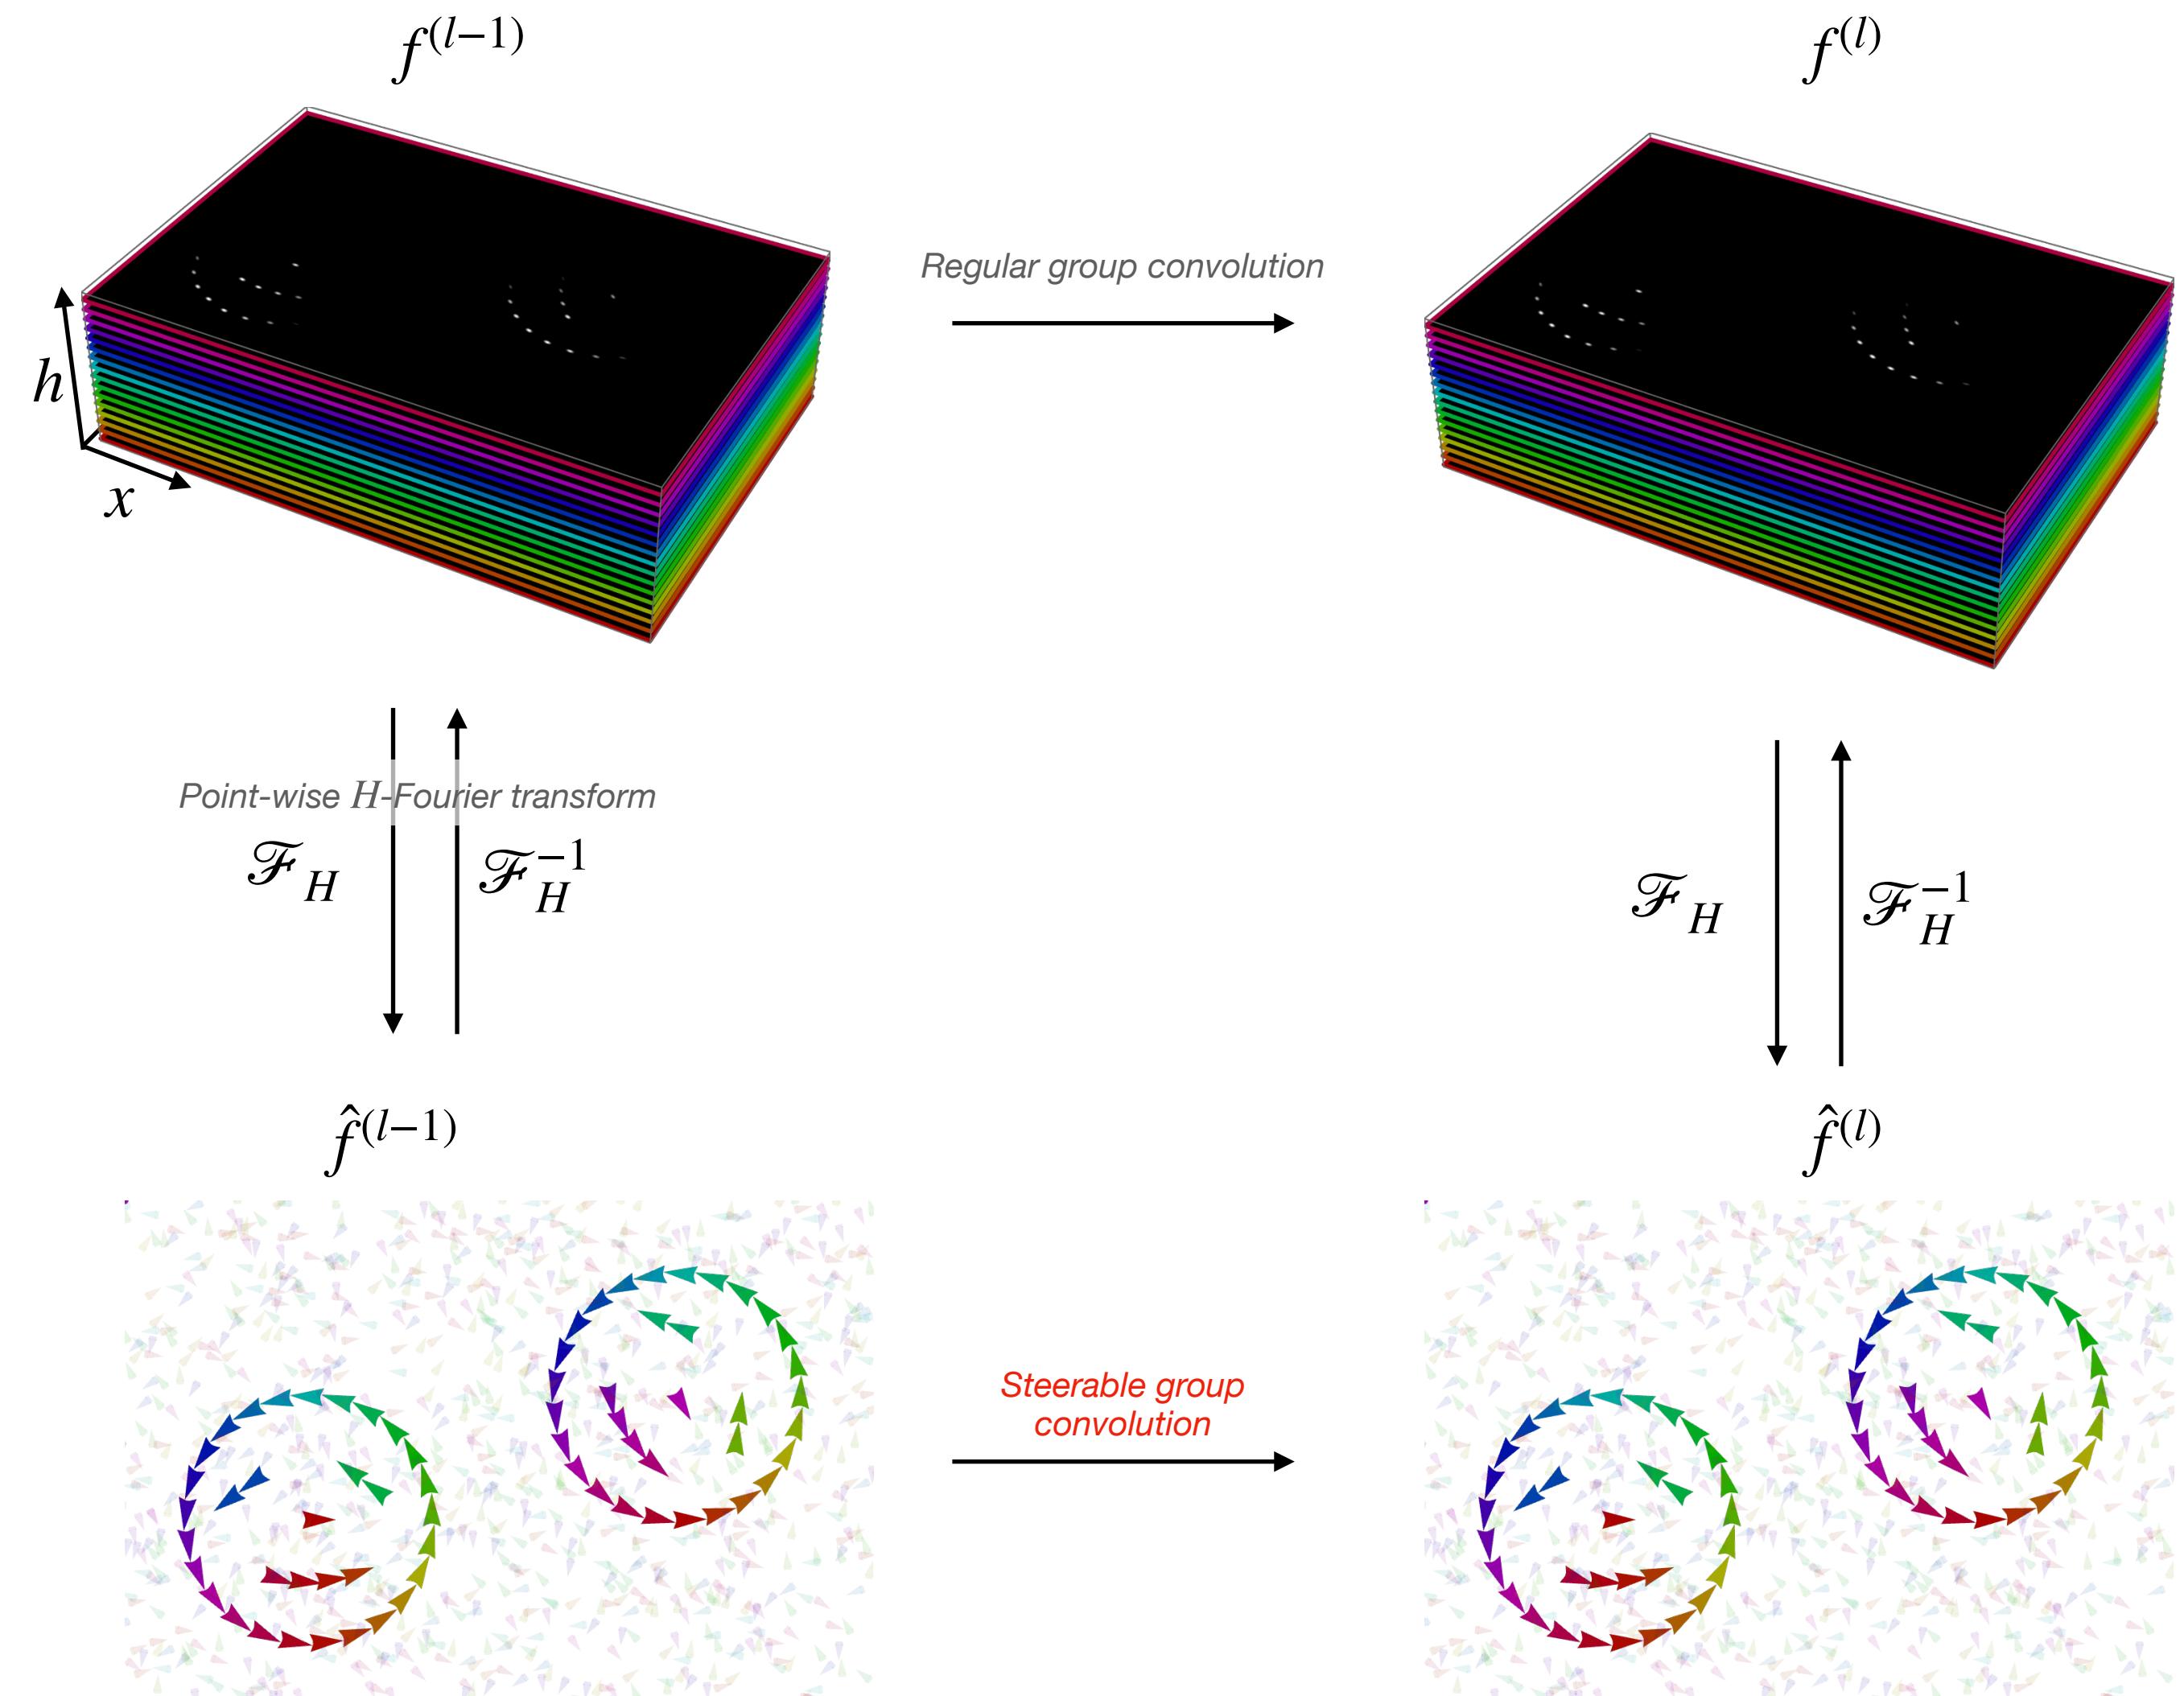
*added axis*

**Steerable group convolutions:**  
Co-domain expanded feature  
maps (feature fields)

$$\hat{f}^{(l)} : \mathbb{R}^d \rightarrow \mathbf{V}_\mathbf{H}$$

*vector field instead of scalar field*

*(vectors in  $\mathbf{V}_\mathbf{H}$  transform via group  $\mathbf{H}$  representations)*



1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

**Efficient (band-limited) grid-free g-convs**

7. Feature fields and escnn library

8. Equivariant tensor product layers

9. Equivariant graph NNs

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

Efficient (band-limited) **grid-free** g-convs

7. Feature fields and escnn library

8. Equivariant tensor product layers

9. Equivariant graph NNs

# Feature field and induced representation

We call  $\hat{f} : \mathbb{R}^d \rightarrow \mathbb{R}^{d_\rho}$  a feature vector field, or simply a **feature field**, if its

<i>codomain</i>	transforms via a representation	$\rho(h)$	of $H$
<i>domain</i>	transforms via the action	$g^{-1}$	of $G = (\mathbb{R}^d, +) \rtimes H$

Representation  $\rho$  defines the **type** of the field, and together with the group action of  $G = (\mathbb{R}^d, +) \rtimes H$  defines the **induced representation**

$$(\text{Ind}_H^G[\rho](\mathbf{x}, h)\hat{f})(\mathbf{x}') := \rho(h)\hat{f}(h^{-1}(\mathbf{x}' - \mathbf{x}))$$

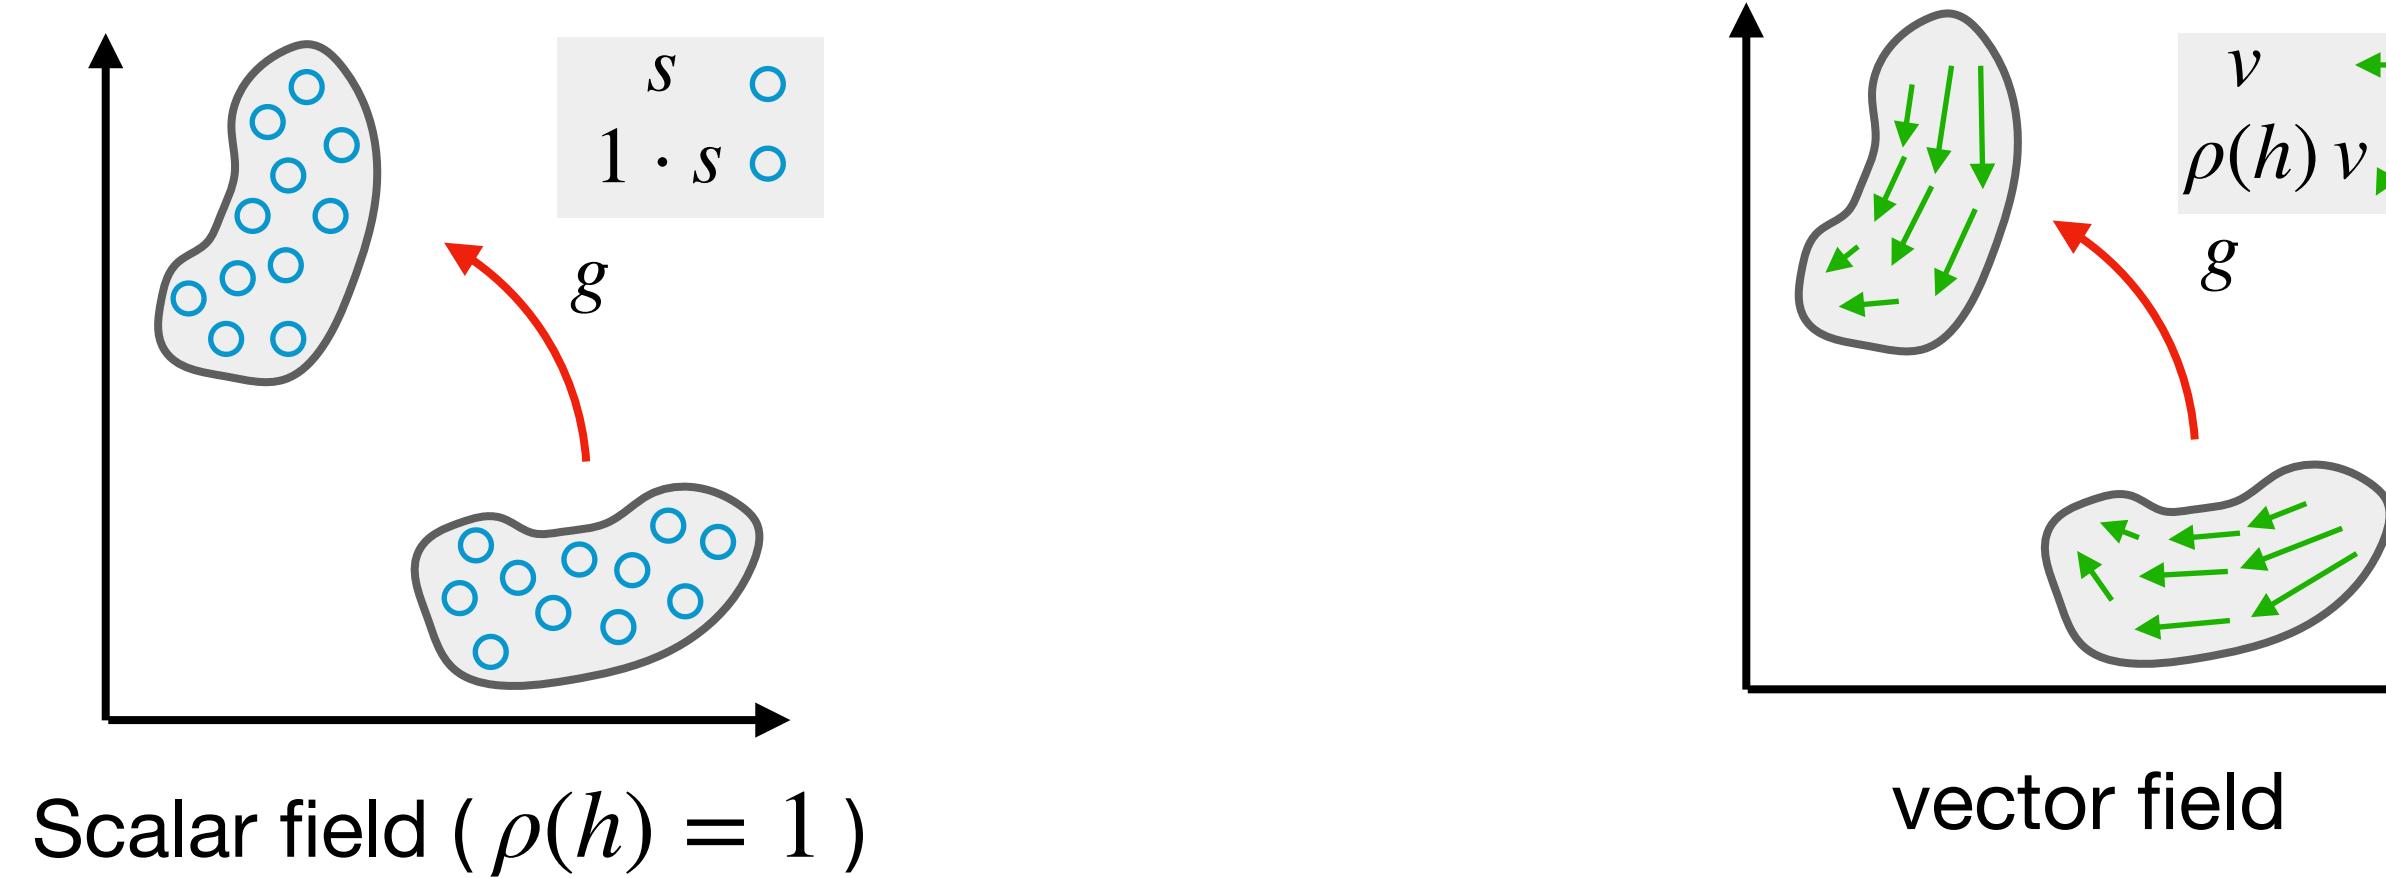
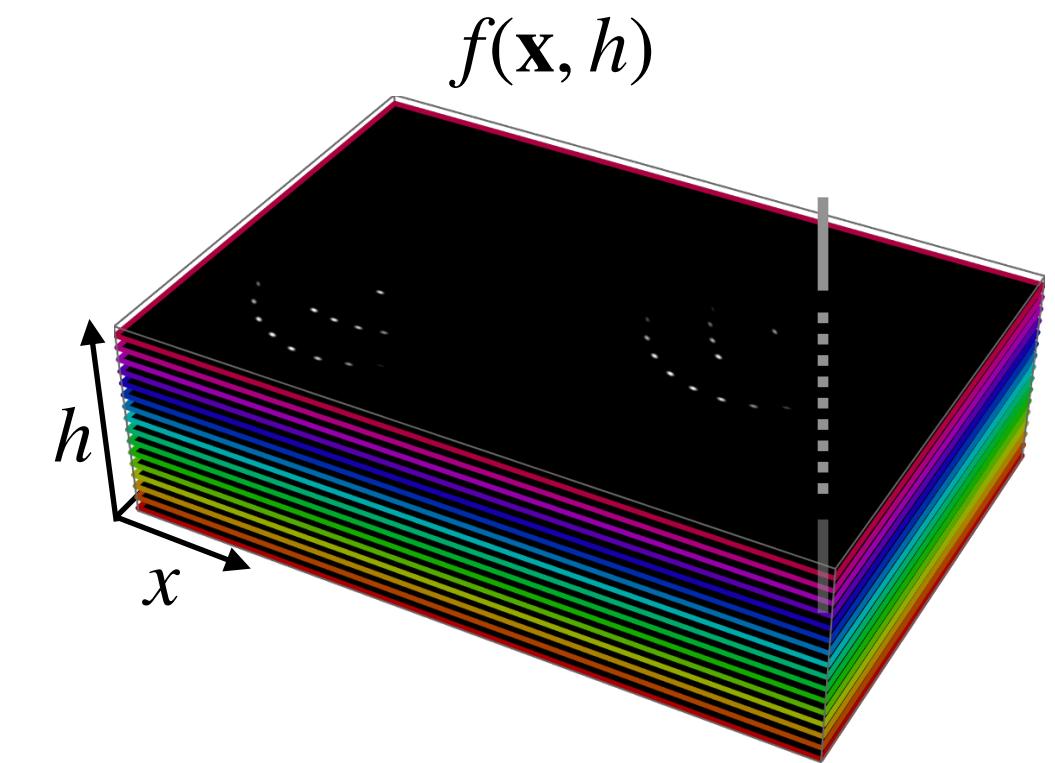


Figure adapted from: Weiler, M., & Cesa, G. (2019). General e (2)-equivariant steerable cnns. NeurIPS  
See also <https://github.com/QUVA-Lab/e2cnn>

# Feature field and induced representation

**Regular  $G$  feature maps:**  $f(\mathbf{x}, h)$  considered so far can be considered feature fields.

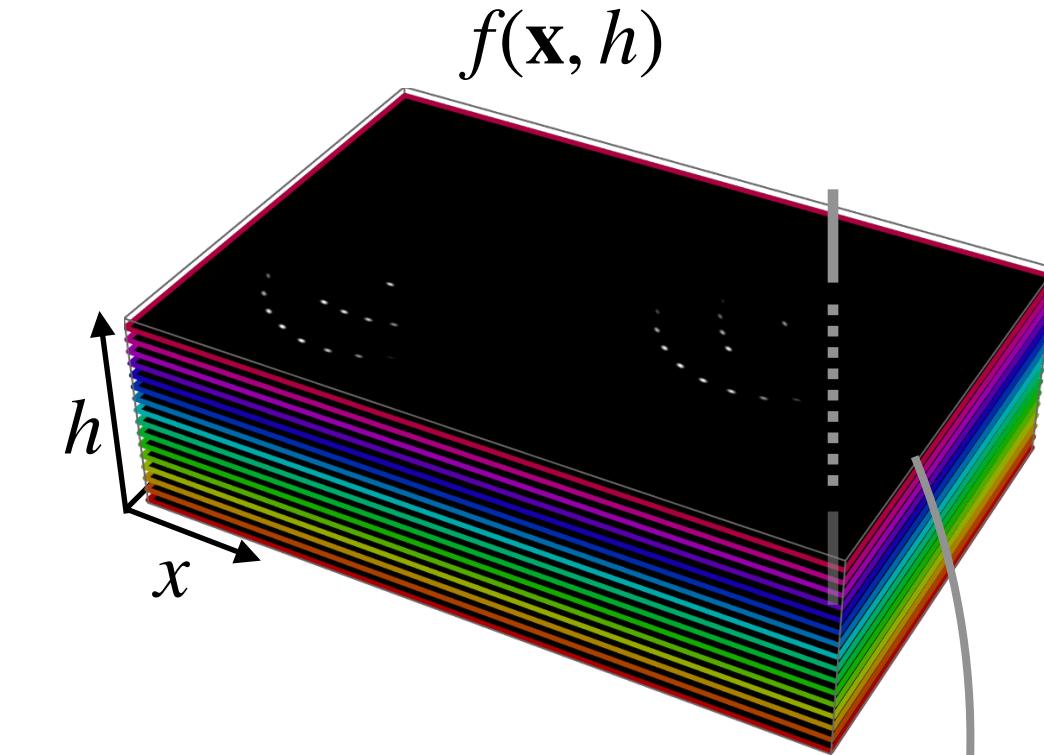
$$(\mathcal{L}_g f)(\mathbf{x}', h') = f(h^{-1}(\mathbf{x}' - \mathbf{x}), h^{-1}h)$$



# Feature field and induced representation

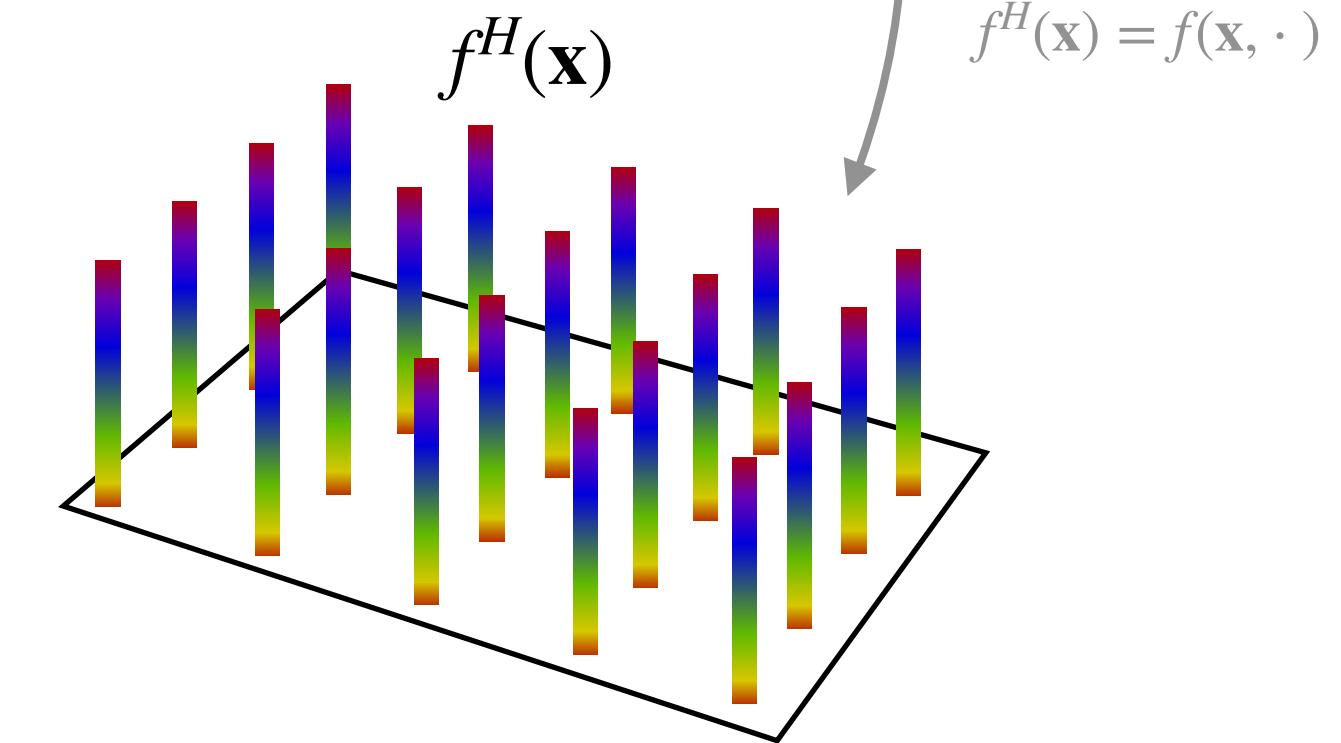
**Regular  $G$  feature maps:**  $f(\mathbf{x}, h)$  considered so far can be considered feature fields.

$$(\mathcal{L}_g f)(\mathbf{x}', h') = f(h^{-1}(\mathbf{x}' - \mathbf{x}), h^{-1}h)$$



**Regular  $H$  feature fields:** Let  $f^H(\mathbf{x}) = f(\mathbf{x}, \cdot)$  be the field of functions  $f^H(\mathbf{x}) : H \rightarrow \mathbb{R}$  on the subgroup  $H$ , then the functions (**fibers**) transform via the regular representation  $\mathcal{L}_h^H$  ( recall.  $\mathcal{L}_h^H f(h') = f(h^{-1}h')$  )

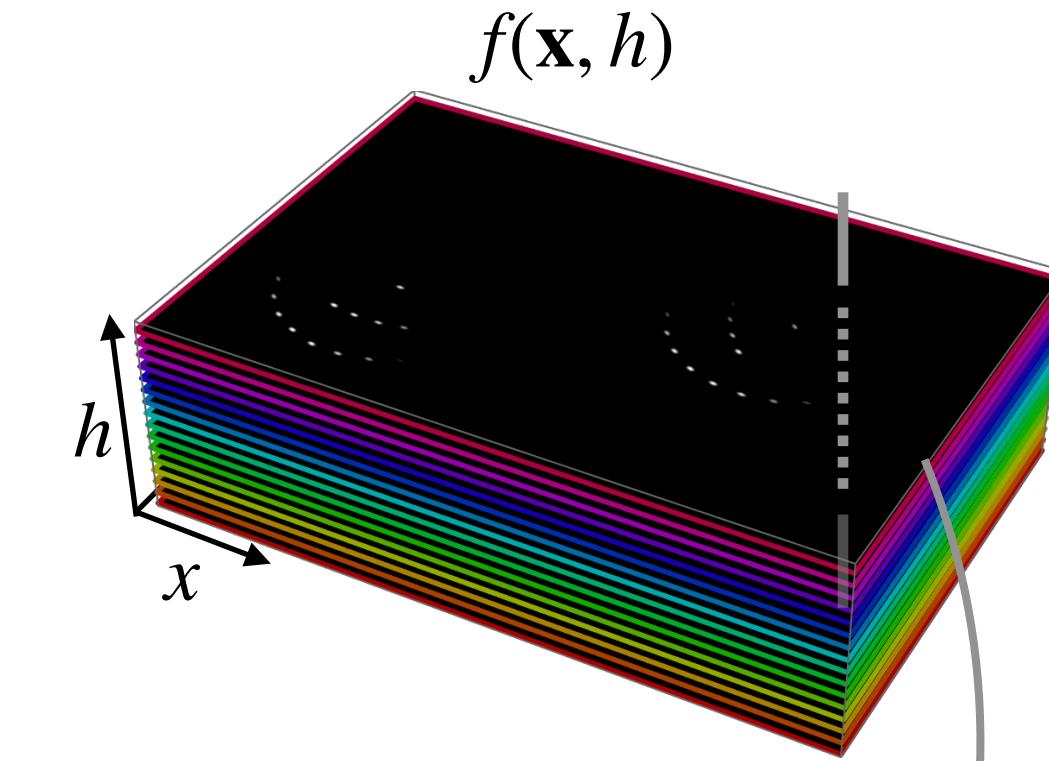
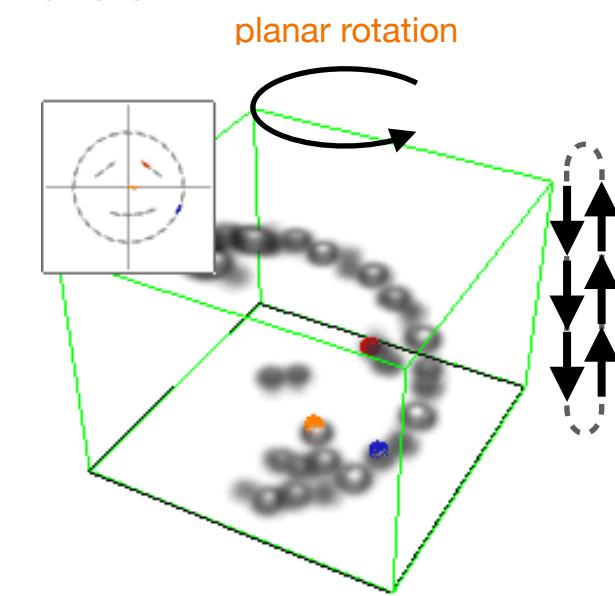
$$(\mathcal{L}_g f)(\mathbf{x}', h') \iff (\text{Ind}_H^G[\mathcal{L}_h^H](\mathbf{x}, h)f^H)(\mathbf{x}')$$



# Feature field and induced representation

**Regular  $G$  feature maps:**  $f(\mathbf{x}, h)$  considered so far can be considered feature fields.

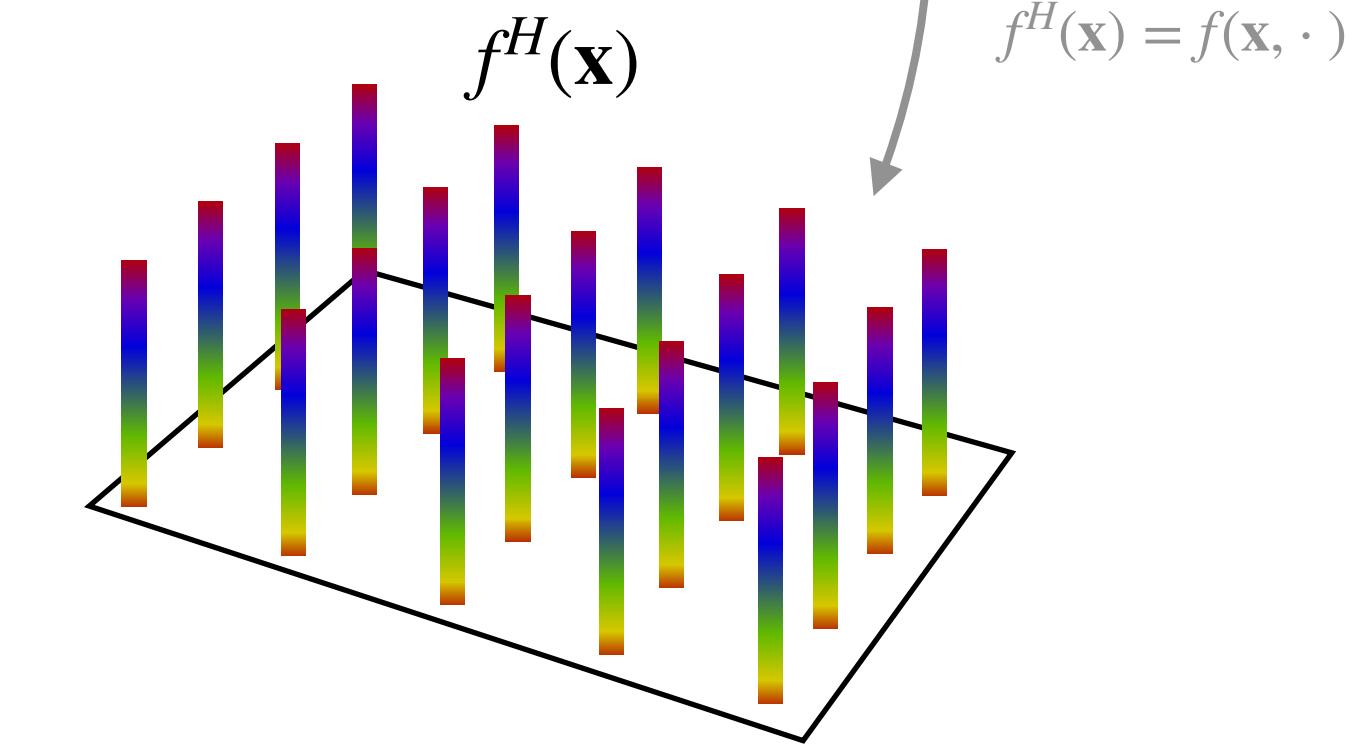
$$(\mathcal{L}_g f)(\mathbf{x}', h') = f(h^{-1}(\mathbf{x}' - \mathbf{x}), h^{-1}h)$$



**Regular  $H$  feature fields:** Let  $f^H(\mathbf{x}) = f(\mathbf{x}, \cdot)$  be the field of functions  $f^H(\mathbf{x}) : H \rightarrow \mathbb{R}$  on the subgroup  $H$ , then the functions (**fibers**) transform via the regular representation  $\mathcal{L}_h^H$

( recall.  $\mathcal{L}_h^H f(h') = f(h^{-1}h')$  )

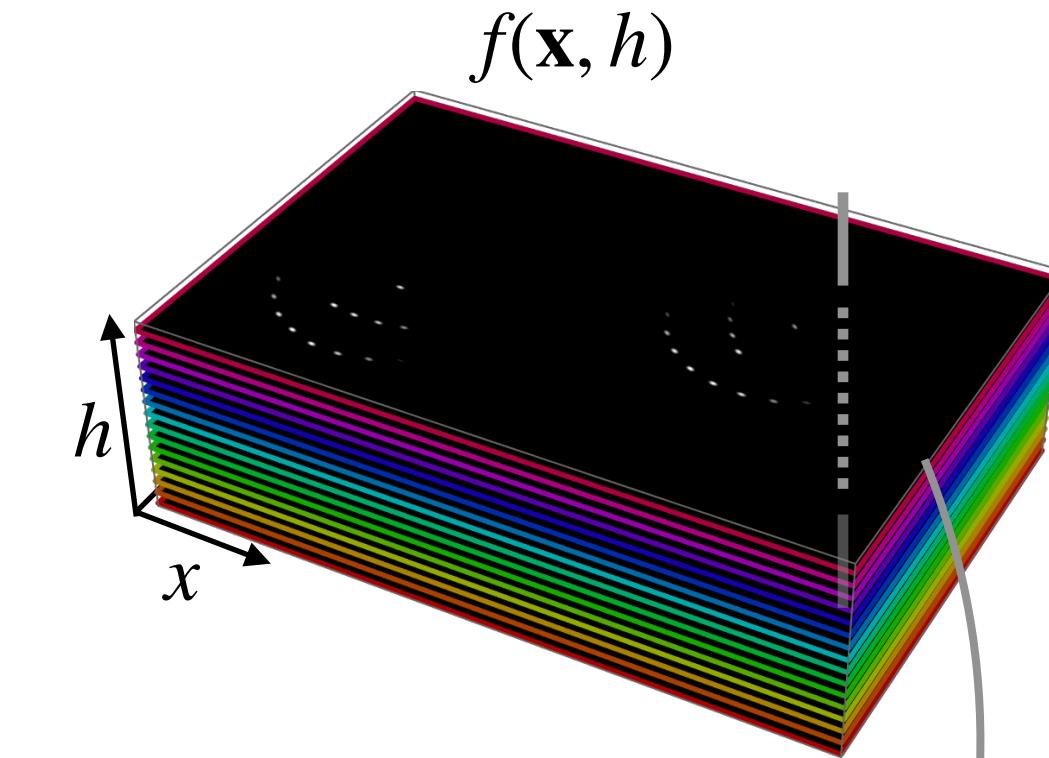
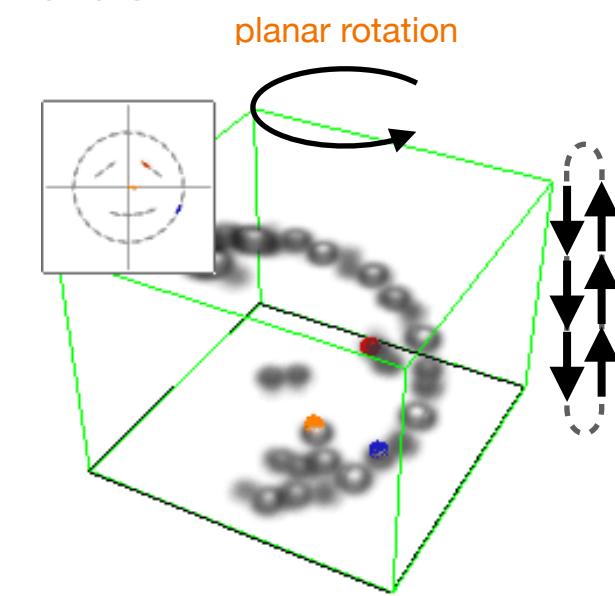
$$(\mathcal{L}_g f)(\mathbf{x}', h') \iff (\text{Ind}_H^G[\mathcal{L}_h^H](\mathbf{x}, h)f^H)(\mathbf{x}')$$



# Feature field and induced representation

**Regular  $G$  feature maps:**  $f(\mathbf{x}, h)$  considered so far can be considered feature fields.

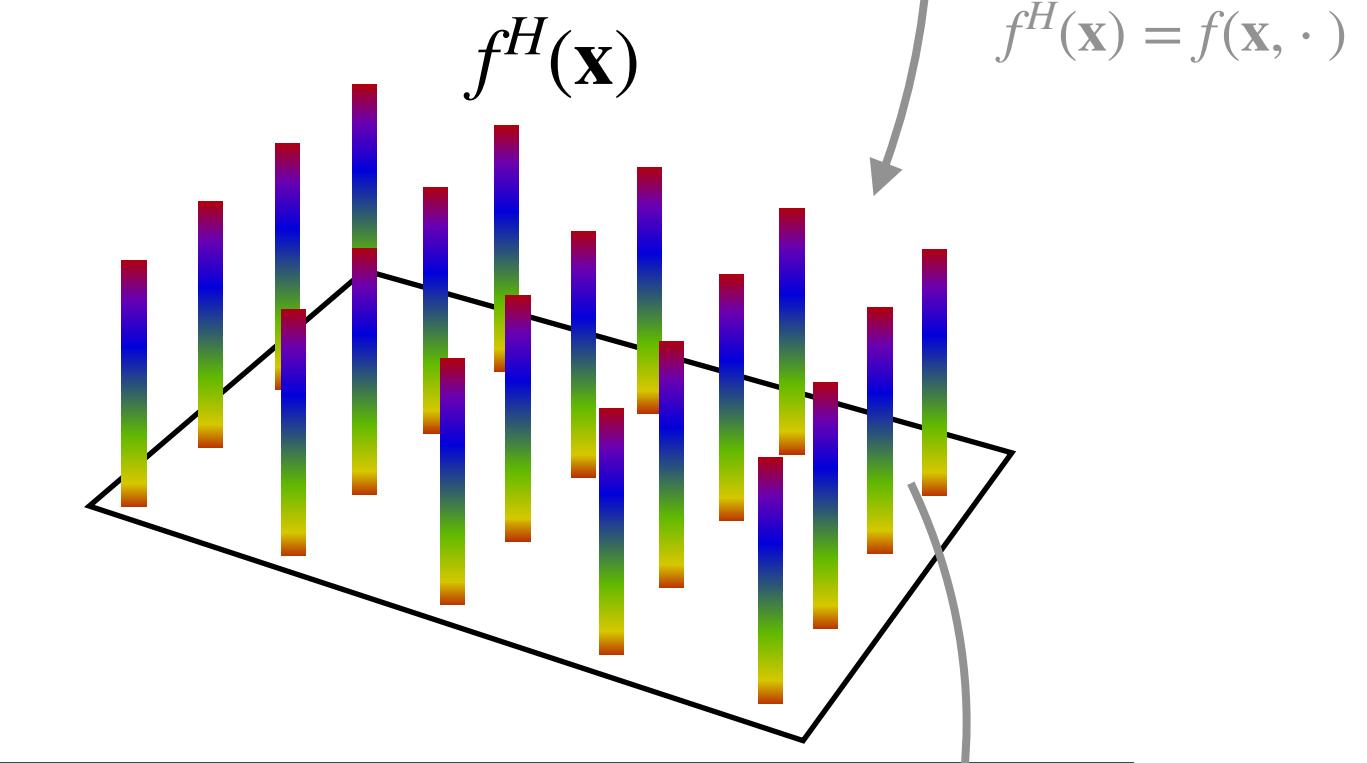
$$(\mathcal{L}_g f)(\mathbf{x}', h') = f(h^{-1}(\mathbf{x}' - \mathbf{x}), h^{-1}h)$$



**Regular  $H$  feature fields:** Let  $f^H(\mathbf{x}) = f(\mathbf{x}, \cdot)$  be the field of functions  $f^H(\mathbf{x}) : H \rightarrow \mathbb{R}$  on the subgroup  $H$ , then the functions (**fibers**) transform via the regular representation  $\mathcal{L}_h^H$

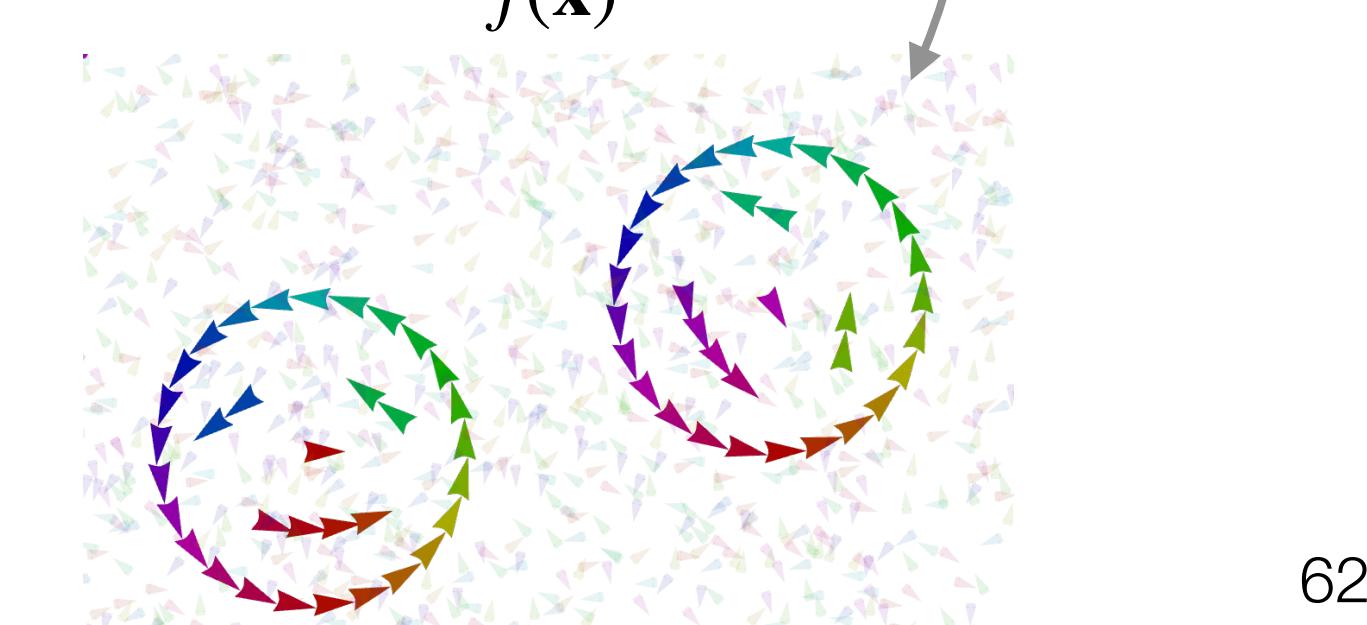
(recall.  $\mathcal{L}_h^H f(h') = f(h^{-1}h')$ )

$$(\mathcal{L}_g f)(\mathbf{x}', h') \iff (\text{Ind}_H^G[\mathcal{L}_h^H](\mathbf{x}, h)f^H)(\mathbf{x}')$$



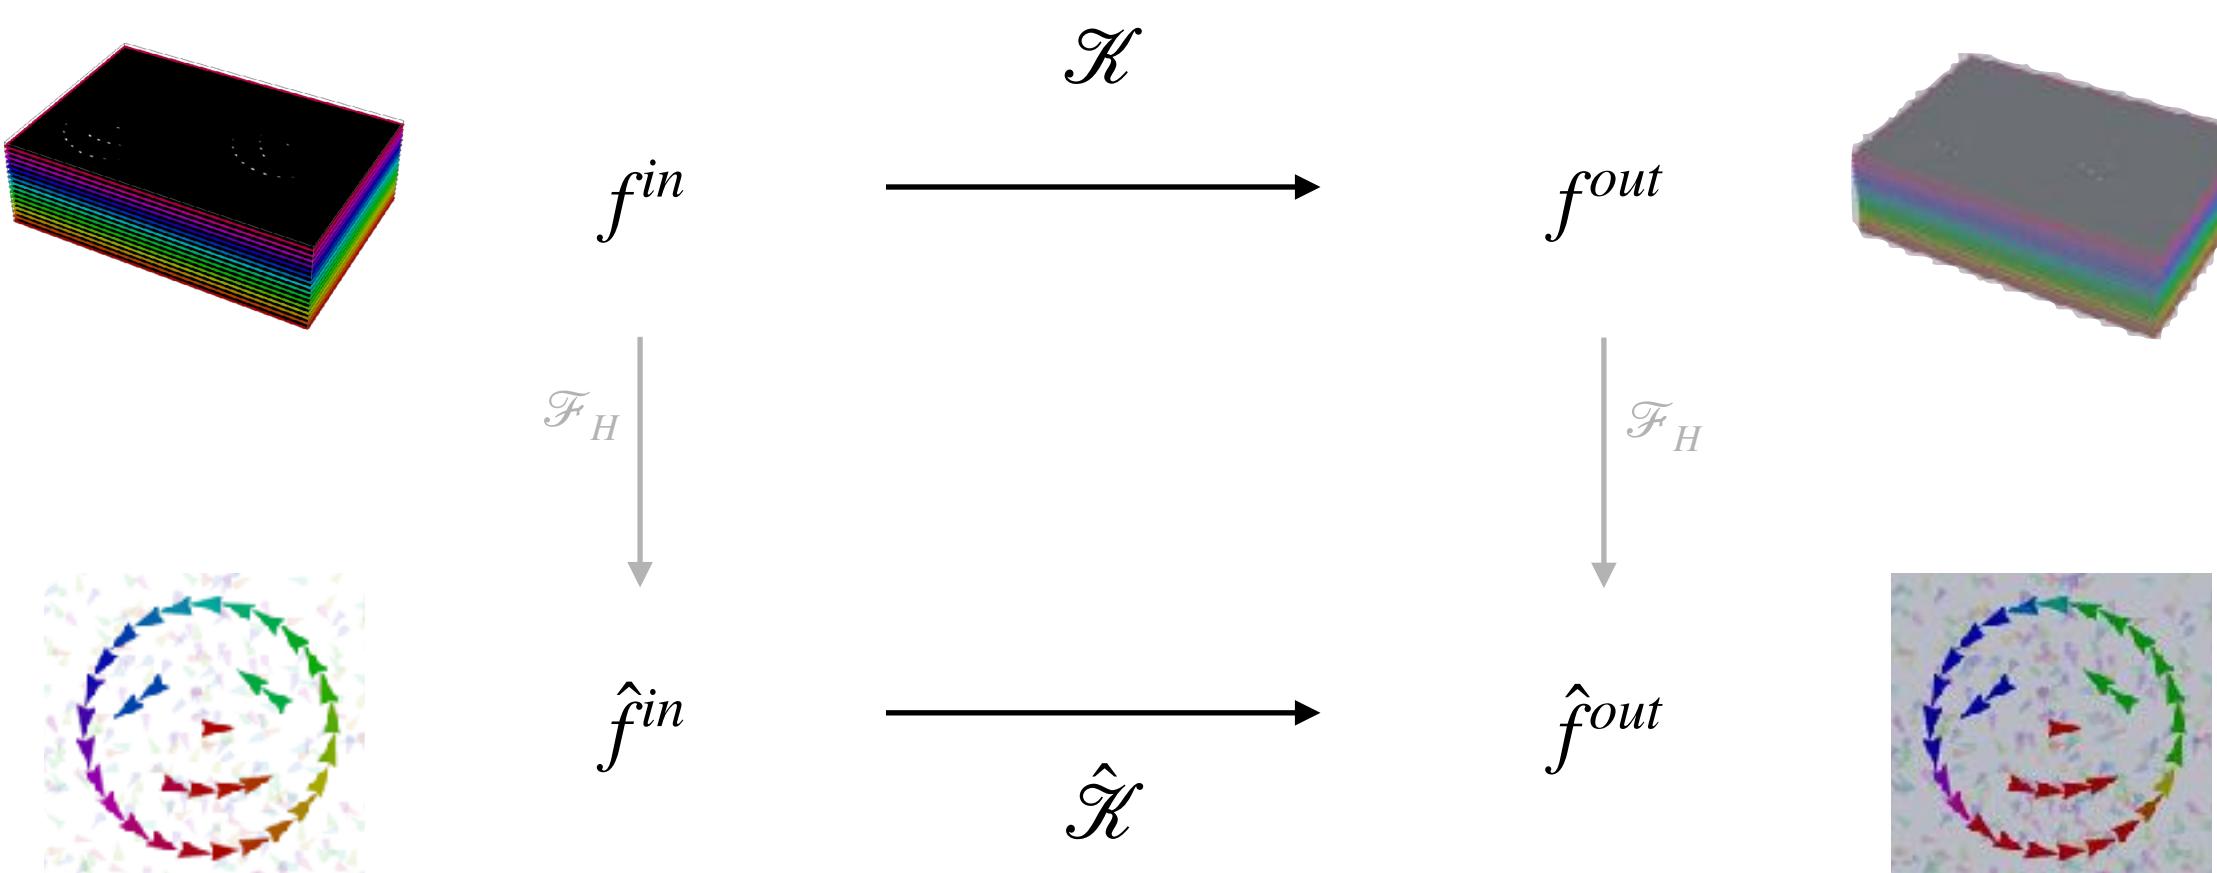
**Steerable  $H$  feature fields:** Since the fibers  $f^H(\mathbf{x})$  are functions on  $H$  we can represent them via their Fourier coefficients  $\hat{f}(\mathbf{x}) = \mathcal{F}_H[f^H(\mathbf{x})]$ . These vectors of coefficients transform via irreps  $\rho(h) = \bigoplus_l \rho_l(h)$

$$(\mathcal{L}_g f)(\mathbf{x}', h') \iff (\text{Ind}_H^G[\mathcal{L}_h^H](\mathbf{x}, h)\hat{f})(\mathbf{x}') \iff (\text{Ind}_H^G[\rho(h)](\mathbf{x}, h)\hat{f})(\mathbf{x}')$$



# Steerable group convolutions

$$\text{Group convolution } \mathcal{K}[f](g) = \int_G k(g^{-1}g')f(g)d\mathbf{x}'dg$$



$$\text{Normal convolution } \mathcal{K}[\hat{f}](\mathbf{x}) = \int_{\mathbb{R}^d} k(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$

**but with kernel  $k : \mathbb{R}^d \rightarrow \mathbb{R}^{d_Y \times d_X}$  satisfying constraint**

$$\forall_{h \in H} \forall_{\mathbf{x} \in \mathbb{R}^d} : k(g \mathbf{x}) = \rho_Y(h)k(\mathbf{x})\rho_X(h^{-1})$$

---

## 3D Steerable CNNs: Learning Rotationally Equivariant Features in Volumetric Data

---

Maurice Weiler\*  
University of Amsterdam  
[m.weiler@uva.nl](mailto:m.weiler@uva.nl)

Mario Geiger\*  
EPFL  
[mario.geiger@epfl.ch](mailto:mario.geiger@epfl.ch)

Max Welling  
University of Amsterdam, CIFAR,  
Qualcomm AI Research  
[m.welling@uva.nl](mailto:m.welling@uva.nl)

Wouter Boomsma  
University of Copenhagen  
[wb@di.ku.dk](mailto:wb@di.ku.dk)

Taco Cohen  
Qualcomm AI Research  
[taco.cohen@gmail.com](mailto:taco.cohen@gmail.com)

### Abstract

We present a convolutional network that is equivariant to rigid body motions. The model uses scalar-, vector-, and tensor fields over 3D Euclidean space to represent data, and equivariant convolutions to map between such representations. These SE(3)-equivariant convolutions utilize equivariant kernels which are parameterized as a linear combination of a complete steerable kernel basis, which is derived analytically in this paper. We prove that equivariant convolutions are the most general equivariant linear maps between fields over  $\mathbb{R}^3$ . Our experimental results confirm the effectiveness of 3D Steerable CNNs for the problem of amino acid propensity prediction and protein structure classification, both of which have inherent SE(3) symmetry.

### 1 Introduction

Increasingly, machine learning techniques are being applied in the natural sciences. Many problems in this domain, such as the analysis of protein structure, exhibit exact or approximate symmetries. It has long been understood that the equations that define a model or natural law should respect the symmetries of the system under study, and that knowledge of symmetries provides a powerful constraint on the space of admissible models. Indeed, in theoretical physics, this idea is enshrined as a fundamental principle, known as Einstein's principle of general covariance. Machine learning, which is, like physics, concerned with the induction of predictive models, is no different: our models must respect known symmetries in order to produce physically meaningful results.

A lot of recent work, reviewed in Sec. 2, has focused on the problem of developing equivariant networks, which respect some known symmetry. In this paper, we develop the theory of SE(3)-equivariant networks. This is far from trivial, because SE(3) is both non-commutative and non-compact. Nevertheless, at run-time, all that is required to make a 3D convolution equivariant using our method, is to parameterize the convolution kernel as a linear combination of pre-computed steerable basis kernels. Hence, the 3D Steerable CNN incorporates equivariance to symmetry transformations without deviating far from current engineering best practices.

The architectures presented here fall within the framework of Steerable G-CNNs [8] [10] [40] [45], which represent their input as fields over a homogeneous space ( $\mathbb{R}^3$  in this case), and use steerable

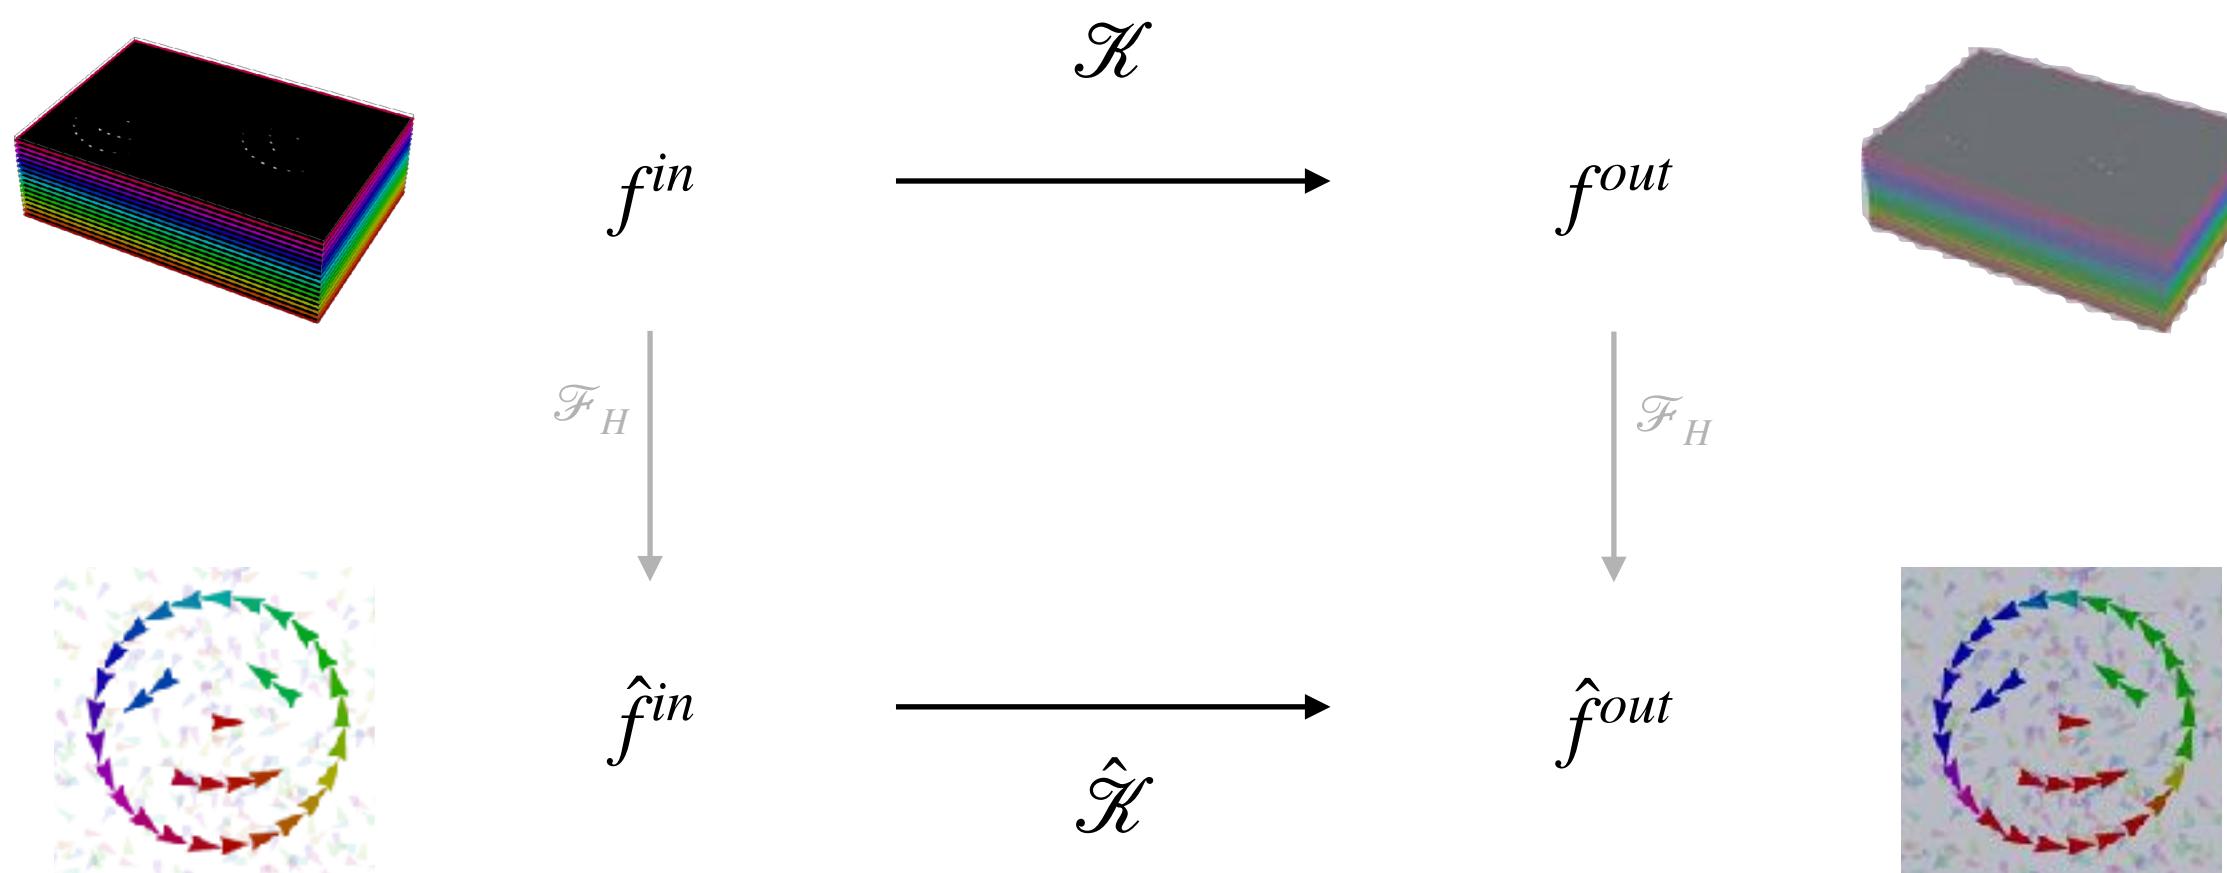
\* Equal Contribution. MG initiated the project, derived the kernel space constraint, wrote the first network implementation and ran the Shrec17 experiment. MW solved the kernel constraint analytically, designed the anti-aliased kernel sampling in discrete space and coded / ran many of the CATH experiments.

Source code is available at <https://github.com/mariogeiger/se3cnn>

32nd Conference on Neural Information Processing Systems (NeurIPS 2018), Montréal, Canada.

# Steerable group convolutions

$$\text{Group convolution } \mathcal{K}[f](g) = \int_G k(g^{-1}g')f(g)d\mathbf{x}'dg$$



$$\text{Normal convolution } \mathcal{K}[\hat{f}](\mathbf{x}) = \int_{\mathbb{R}^d} k(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$

**but with kernel  $k : \mathbb{R}^d \rightarrow \mathbb{R}^{d_Y \times d_X}$  satisfying constraint**

$$\forall_{h \in H} \forall_{\mathbf{x} \in \mathbb{R}^d} : k(g \mathbf{x}) = \rho_Y(h)k(\mathbf{x})\rho_X(h^{-1})$$

---

## General E(2) - Equivariant Steerable CNNs

---

Maurice Weiler\*  
University of Amsterdam, QUVA Lab  
m.weiler@uva.nl

Gabriele Cesa†  
University of Amsterdam  
cesa.gabriele@gmail.com

### Abstract

The big empirical success of group equivariant networks has led in recent years to the sprouting of a great variety of equivariant network architectures. A particular focus has thereby been on rotation and reflection equivariant CNNs for planar images. Here we give a general description of E(2)-equivariant convolutions in the framework of *Steerable CNNs*. The theory of Steerable CNNs thereby yields constraints on the convolution kernels which depend on group representations describing the transformation laws of feature spaces. We show that these constraints for arbitrary group representations can be reduced to constraints under irreducible representations. A general solution of the kernel space constraint is given for arbitrary representations of the Euclidean group E(2) and its subgroups. We implement a wide range of previously proposed and entirely new equivariant network architectures and extensively compare their performances. E(2)-steerable convolutions are further shown to yield remarkable gains on CIFAR-10, CIFAR-100 and STL-10 when used as drop-in replacement for non-equivariant convolutions.

### 1 Introduction

The equivariance of neural networks under symmetry group actions has in the recent years proven to be a fruitful prior in network design. By guaranteeing a desired transformation behavior of convolutional features under transformations of the network input, equivariant networks achieve improved generalization capabilities and sample complexities compared to their non-equivariant counterparts. Due to their great practical relevance, a big pool of rotation- and reflection-equivariant models for planar images has been proposed by now. Unfortunately, an empirical survey, reproducing and comparing all these different approaches, is still missing.

An important step in this direction is given by the theory of *Steerable CNNs* [1, 2, 3, 4, 5] which defines a very general notion of equivariant convolutions on homogeneous spaces. In particular, steerable CNNs describe E(2)-equivariant (i.e. rotation- and reflection-equivariant) convolutions on the image plane  $\mathbb{R}^2$ . The feature spaces of steerable CNNs are thereby defined as spaces of *feature fields*, characterized by a group representation which determines their transformation behavior under transformations of the input. In order to preserve the specified transformation law of feature spaces, the convolutional kernels are subject to a linear constraint, depending on the corresponding group representations. While this constraint has been solved for specific groups and representations [1, 2], no general solution strategy has been proposed so far. In this work we give a general strategy which reduces the solution of the kernel space constraint under arbitrary representations to much simpler constraints under single, *irreducible* representations.

Specifically for the Euclidean group E(2) and its subgroups, we give a general solution of this kernel space constraint. As a result, we are able to implement a wide range of equivariant models, covering regular GCNNs [6, 7, 8, 9, 10, 11], classical Steerable CNNs [1], Harmonic Networks [12], gated Harmonic Networks [2], Vector Field Networks [13], Scattering Transforms [14, 15, 16, 17, 18] and entirely new architectures, in one unified framework. In addition, we are able to build hybrid models, mixing different field types (representations) of these networks both over layers and within layers.

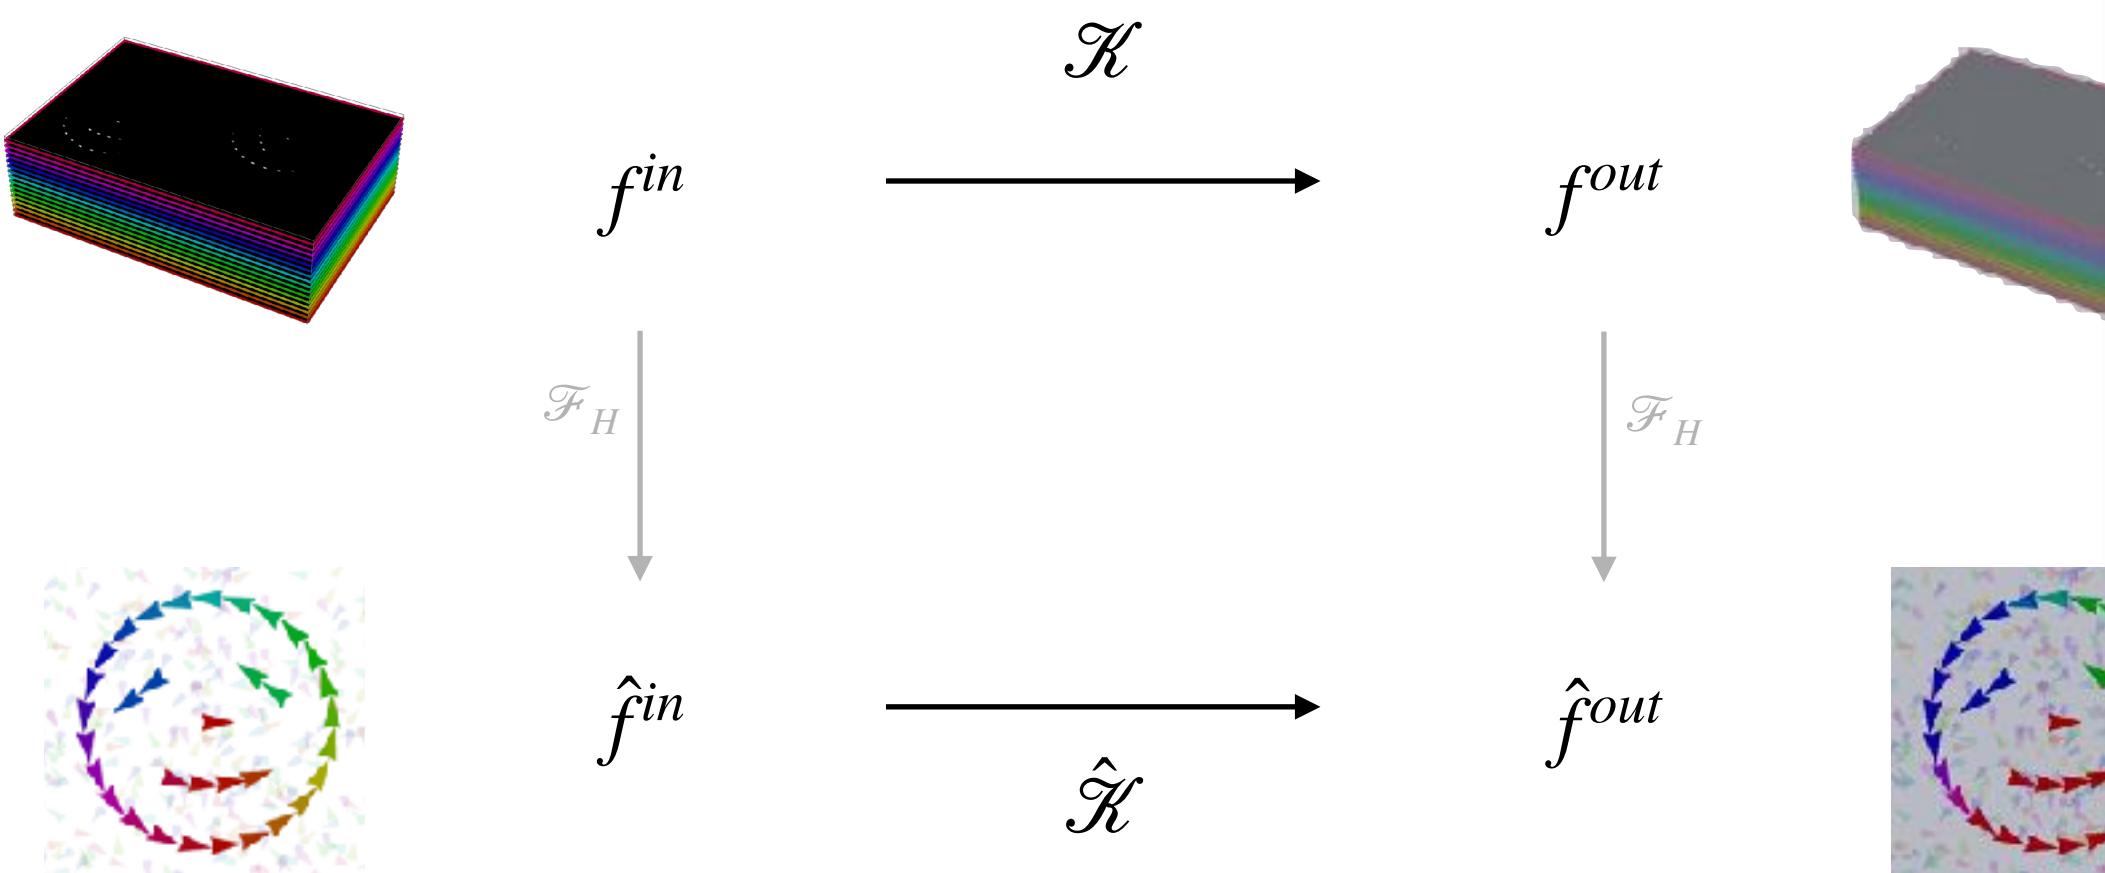
\* Equal contribution, author ordering determined by random number generator.

† This research has been conducted during an internship at QUVA lab, University of Amsterdam.

# Steerable group convolutions

Published as a conference paper at ICLR 2022

$$\text{Group convolution } \mathcal{K}[f](g) = \int_G k(g^{-1}g')f(g)d\mathbf{x}'dg$$



$$\text{Normal convolution } \mathcal{K}[\hat{f}](\mathbf{x}) = \int_{\mathbb{R}^d} k(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$

**but with kernel  $k : \mathbb{R}^d \rightarrow \mathbb{R}^{d_Y \times d_X}$  satisfying constraint**

$$\forall_{h \in H} \forall_{\mathbf{x} \in \mathbb{R}^d} : k(g \mathbf{x}) = \rho_Y(h)k(\mathbf{x})\rho_X(h^{-1})$$

## A PROGRAM TO BUILD $E(n)$ -EQUIVARIANT STEERABLE CNNS

Gabriele Cesa  
Qualcomm AI Research\*  
University of Amsterdam  
gcesa@qti.qualcomm.com

Leon Lang  
University of Amsterdam  
l.lang@uva.nl

Maurice Weiler  
University of Amsterdam  
m.weiler.ml@gmail.com

### ABSTRACT

Equivariance is becoming an increasingly popular design choice to build data efficient neural networks by exploiting prior knowledge about the symmetries of the problem at hand. Euclidean steerable CNNs are one of the most common classes of equivariant networks. While the constraints these architectures need to satisfy are understood, existing approaches are tailored to specific (classes of) groups. No generally applicable method that is *practical* for implementation has been described so far. In this work, we generalize the Wigner-Eckart theorem proposed in [Lang & Weiler (2020)], which characterizes general  $G$ -steerable kernel spaces for compact groups  $G$  over their homogeneous spaces, to arbitrary  $G$ -spaces. This enables us to directly parameterize filters in terms of a band-limited basis on the whole space rather than on  $G$ 's orbits, but also to easily implement steerable CNNs equivariant to a large number of groups. To demonstrate its generality, we instantiate our method on a variety of isometry groups acting on the Euclidean space  $\mathbb{R}^3$ . Our framework allows us to build  $E(3)$  and  $SE(3)$ -steerable CNNs like previous works, but also CNNs with arbitrary  $G \leq O(3)$ -steerable kernels. For example, we build 3D CNNs equivariant to the symmetries of platonic solids or choose  $G = SO(2)$  when working with 3D data having only azimuthal symmetries. We compare these models on 3D shapes and molecular datasets, observing improved performance by matching the model's symmetries to the ones of the data.

### 1 INTRODUCTION

In machine learning, it is common for learning tasks to present a number of *symmetries*. A symmetry in the data occurs, for example, when some property (e.g., the label) does not change if a set of transformations is applied to the data itself, e.g. translations or rotations of images. Symmetries are algebraically described by *groups*. If prior knowledge about the symmetries of a task is available, it is usually beneficial to encode them in the models used (Shawe-Taylor (1989); Cohen & Welling (2016a)). The property of such models is referred to as *equivariance* and is obtained by introducing some *equivariance constraints* in the architecture (see Eq. 2). A classical example are convolutional neural networks (CNNs), which achieve translation equivariance by constraining linear layers to be convolution operators. A wider class of equivariant models are Euclidean steerable CNNs (Cohen & Welling (2016b); Weiler et al. (2018a); Weiler & Cesa (2019); Jenner & Weiler (2022)), which guarantee equivariance to isometries  $\mathbb{R}^n \times G$  of a Euclidean space  $\mathbb{R}^n$ , i.e., to translations and a group  $G$  of origin-preserving transformations, such as rotations and reflections. As proven in Weiler et al. (2018a; 2021); Jenner & Weiler (2022), this requires convolutions with  $G$ -steerable (equivariant) kernels.

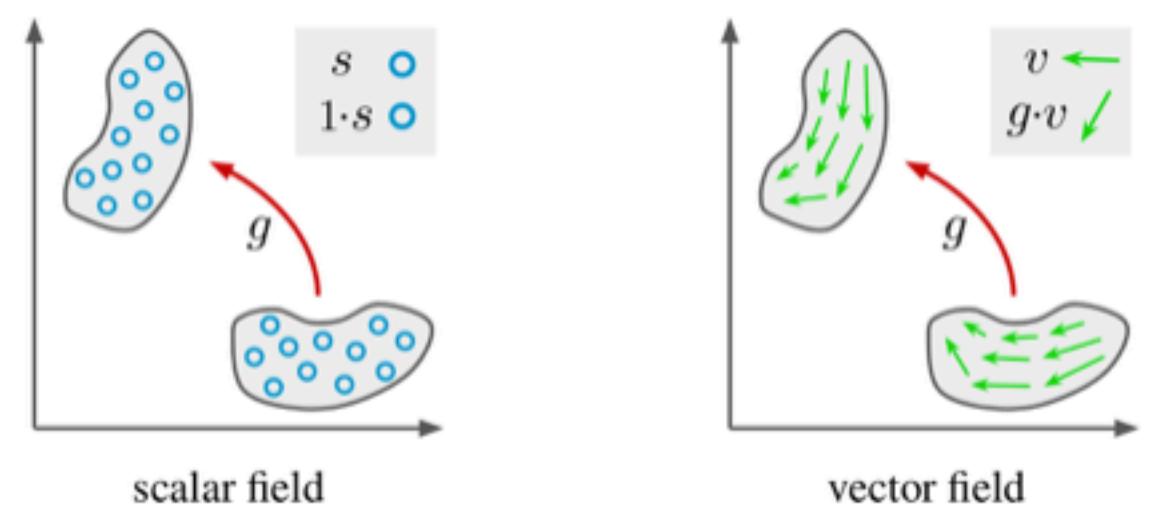
Our goal is developing a program to parameterize with minimal requirements arbitrary  $G$ -steerable kernel spaces, with compact  $G$ , which are required to implement  $\mathbb{R}^n \times G$  equivariant CNNs. [Lang & Weiler (2020)] provides a first step in this direction by generalizing the *Wigner-Eckart theorem* from quantum mechanics to obtain a general technique to parametrize  $G$ -steerable kernel spaces over *orbits* of a compact  $G$ . The theorem reduces the task of building steerable kernel bases to that of finding some pure representation theoretic ingredients. Since the equivariance constraint only relates points  $g.x \in \mathbb{R}^n$  in the same *orbit*  $G.x \subset \mathbb{R}^n$ , a kernel can take independent values on different orbits. Fig. 1 shows

\*Qualcomm AI Research is an initiative of Qualcomm Technologies, Inc.

escnn is a [PyTorch](#) extension for equivariant deep learning. escnn is the successor of the [e2cnn](#) library, which only supported planar isometries.

*Equivariant neural networks* guarantee a specified transformation behavior of their feature spaces under transformations of their input. For instance, classical convolutional neural networks (CNNs) are by design equivariant to translations of their input. This means that a translation of an image leads to a corresponding translation of the network's feature maps. This package provides implementations of neural network modules which are equivariant under all *isometries*  $E(2)$  of the image plane  $\mathbb{R}^2$  and all *isometries*  $E(3)$  of the 3D space  $\mathbb{R}^3$ , that is, under *translations, rotations and reflections* (and can, potentially, be extended to all isometries  $E(n)$  of  $\mathbb{R}^n$ ). In contrast to conventional CNNs, E(n)-equivariant models are guaranteed to generalize over such transformations, and are therefore more data efficient.

The feature spaces of E(n)-Equivariant Steerable CNNs are defined as spaces of *feature fields*, being characterized by their transformation law under rotations and reflections. Typical examples are scalar fields (e.g. gray-scale images or temperature fields) or vector fields (e.g. optical flow or electromagnetic fields).



Instead of a number of channels, the user has to specify the field types and their *multiplicities* in order to define a feature space. Given a specified input- and output feature space, our `R2conv` and `R3conv` modules instantiate

$$\text{Normal convolution } \mathcal{K}[\hat{f}](\mathbf{x}) = \int_{\mathbb{R}^d} k(\mathbf{x}' - \mathbf{x})f(\mathbf{x}')d\mathbf{x}'$$

but with kernel  $k : \mathbb{R}^d \rightarrow \mathbb{R}^{d_Y \times d_X}$  satisfying constraint

$$\forall_{h \in H} \forall_{\mathbf{x} \in \mathbb{R}^d} : k(g \mathbf{x}) = \rho_Y(h)k(\mathbf{x})\rho_X(h^{-1})$$

# up convolutions

Published as a conference paper at ICLR 2022

## A PROGRAM TO BUILD E( $n$ )-EQUIVARIANT STEERABLE CNNS

Gabriele Cesa  
Qualcomm AI Research\*  
University of Amsterdam  
gcesa@qti.qualcomm.com

Leon Lang  
University of Amsterdam  
l.lang@uva.nl

Maurice Weiler  
University of Amsterdam  
m.weiler.ml@gmail.com

### ABSTRACT

Equivariance is becoming an increasingly popular design choice to build data efficient neural networks by exploiting prior knowledge about the symmetries of the problem at hand. Euclidean steerable CNNs are one of the most common classes of equivariant networks. While the constraints these architectures need to satisfy are understood, existing approaches are tailored to specific (classes of) groups. No generally applicable method that is *practical* for implementation has been described so far. In this work, we generalize the Wigner-Eckart theorem proposed in [Lang & Weiler \(2020\)](#), which characterizes general  $G$ -steerable kernel spaces for compact groups  $G$  over their homogeneous spaces, to arbitrary  $G$ -spaces. This enables us to directly parameterize filters in terms of a band-limited basis on the whole space rather than on  $G$ 's orbits, but also to easily implement steerable CNNs equivariant to a large number of groups. To demonstrate its generality, we instantiate our method on a variety of isometry groups acting on the Euclidean space  $\mathbb{R}^3$ . Our framework allows us to build  $E(3)$  and  $SE(3)$ -steerable CNNs like previous works, but also CNNs with arbitrary  $G \leq O(3)$ -steerable kernels. For example, we build 3D CNNs equivariant to the symmetries of platonic solids or choose  $G = SO(2)$  when working with 3D data having only azimuthal symmetries. We compare these models on 3D shapes and molecular datasets, observing improved performance by matching the model's symmetries to the ones of the data.

### 1 INTRODUCTION

In machine learning, it is common for learning tasks to present a number of *symmetries*. A symmetry in the data occurs, for example, when some property (e.g., the label) does not change if a set of transformations is applied to the data itself, e.g. translations or rotations of images. Symmetries are algebraically described by *groups*. If prior knowledge about the symmetries of a task is available, it is usually beneficial to encode them in the models used ([Shawe-Taylor 1989](#) [Cohen & Welling 2016a](#)). The property of such models is referred to as *equivariance* and is obtained by introducing some *equivariance constraints* in the architecture (see Eq. 2). A classical example are convolutional neural networks (CNNs), which achieve translation equivariance by constraining linear layers to be convolution operators. A wider class of equivariant models are Euclidean steerable CNNs ([Cohen & Welling 2016b](#); [Weiler et al. 2018a](#); [Weiler & Cesa 2019](#); [Jenner & Weiler 2022](#)), which guarantee equivariance to isometries  $\mathbb{R}^n \times G$  of a Euclidean space  $\mathbb{R}^n$ , i.e., to translations and a group  $G$  of origin-preserving transformations, such as rotations and reflections. As proven in [Weiler et al. \(2018a; 2021\)](#); [Jenner & Weiler \(2022\)](#), this requires convolutions with  $G$ -steerable (equivariant) kernels.

Our goal is developing a program to parameterize with minimal requirements arbitrary  $G$ -steerable kernel spaces, with compact  $G$ , which are required to implement  $\mathbb{R}^n \times G$  equivariant CNNs. [Lang & Weiler \(2020\)](#) provides a first step in this direction by generalizing the *Wigner-Eckart theorem* from quantum mechanics to obtain a general technique to parametrize  $G$ -steerable kernel spaces over *orbits* of a compact  $G$ . The theorem reduces the task of building steerable kernel bases to that of finding some pure representation theoretic ingredients. Since the equivariance constraint only relates points  $g \cdot x \in \mathbb{R}^n$  in the same *orbit*  $G \cdot x \subset \mathbb{R}^n$ , a kernel can take independent values on different orbits. Fig. 1 shows

\*Qualcomm AI Research is an initiative of Qualcomm Technologies, Inc.

`escnn` is a [PyTorch](#) extension for equivariant deep learning. `escnn` is the successor of the `e2cnn` library, which only supported planar isometries.

*Equivariant neural networks* guarantee a specified transformation behavior of their feature spaces under transformations of their input. For instance, classical convolutional neural networks (CNNs) are by design equivariant to translations of their input. This means that a translation of an image leads to a corresponding translation of the network's feature maps. This package provides implementations of neural network modules which are equivariant under all *isometries*  $E(2)$  of the image plane  $\mathbb{R}^2$  and all *isometries*  $E(3)$  of the 3D space  $\mathbb{R}^3$ , that is, under *translations, rotations and reflections* (and can, potentially, be extended to all isometries  $E(n)$  of  $\mathbb{R}^n$ ). In contrast to conventional CNNs,  $E(n)$ -equivariant models are guaranteed to generalize over such transformations.

The feature characterizes gray-scale

Instead of a feature spa

## Getting Started

`escnn` is easy to use since it provides a high level user interface which abstracts most intricacies of group and representation theory away. The following code snippet shows how to perform an equivariant convolution from an RGB-image to 10 *regular* feature fields (corresponding to a [group convolution](#)).

```
from escnn import gspaces
from escnn import nn
import torch

r2_act = gspaces.rot2dOnR2(N=8)
feat_type_in = nn.FieldType(r2_act, 3*[r2_act.trivial_repr])
feat_type_out = nn.FieldType(r2_act, 10*[r2_act.regular_repr])

conv = nn.R2Conv(feat_type_in, feat_type_out, kernel_size=5)
relu = nn.ReLU(feat_type_out)

x = torch.randn(16, 3, 32, 32)
x = feat_type_in(x)

y = relu(conv(x))
```

Line 5 specifies the symmetry group action on the image plane  $\mathbb{R}^2$  under which the network should be equivariant. We choose the [cyclic group](#)  $C_8$ , which describes discrete rotations by multiples of  $2\pi/8$ . Line 6 specifies the input feature field types. The three color channels of an RGB image are thereby to be identified as three independent scalar fields, which transform under the [trivial representation](#) of  $C_8$ . Similarly, the output feature space is in line 7 specified to consist of 10 feature fields which transform under the [regular representation](#) of  $C_8$ . The  $C_8$ -equivariant convolution is then instantiated by passing the input and output type as well as the kernel size to the constructor (line 9). Line 10 instantiates an equivariant ReLU nonlinearity which will operate on the output field and is therefore passed the output field type.

# up convolutions

Published as a conference paper at ICLR 2022

## A PROGRAM TO BUILD $E(n)$ -EQUIVARIANT STEERABLE CNNS

Gabriele Cesa  
Qualcomm AI Research\*  
University of Amsterdam  
gcesa@qti.qualcomm.com

Leon Lang  
University of Amsterdam  
l.lang@uva.nl

Maurice Weiler  
University of Amsterdam  
m.weiler.ml@gmail.com

### ABSTRACT

Equivariance is becoming an increasingly popular design choice to build data efficient neural networks by exploiting prior knowledge about the symmetries of the problem at hand. Euclidean steerable CNNs are one of the most common classes of equivariant networks. While the constraints these architectures need to satisfy are understood, existing approaches are tailored to specific (classes of) groups. No generally applicable method that is *practical* for implementation has been described so far. In this work, we generalize the Wigner-Eckart theorem proposed in [Lang & Weiler \(2020\)](#), which characterizes general  $G$ -steerable kernel spaces for compact groups  $G$  over their homogeneous spaces, to arbitrary  $G$ -spaces. This enables us to directly parameterize filters in terms of a band-limited basis on the whole space rather than on  $G$ 's orbits, but also to easily implement steerable CNNs equivariant to a large number of groups. To demonstrate its generality, we instantiate our method on a variety of isometry groups acting on the Euclidean space  $\mathbb{R}^3$ . Our framework allows us to build  $E(3)$  and  $SE(3)$ -steerable CNNs like previous works, but also CNNs with arbitrary  $G \leq O(3)$ -steerable kernels. For example, we build 3D CNNs equivariant to the symmetries of platonic solids or choose  $G = SO(2)$  when working with 3D data having only azimuthal symmetries. We compare these models on 3D shapes and molecular datasets, observing improved performance by matching the model's symmetries to the ones of the data.

### 1 INTRODUCTION

In machine learning, it is common for learning tasks to present a number of *symmetries*. A symmetry in the data occurs, for example, when some property (e.g., the label) does not change if a set of transformations is applied to the data itself, e.g. translations or rotations of images. Symmetries are algebraically described by *groups*. If prior knowledge about the symmetries of a task is available, it is usually beneficial to encode them in the models used ([Shawe-Taylor 1989](#) [Cohen & Welling 2016a](#)). The property of such models is referred to as *equivariance* and is obtained by introducing some *equivariance constraints* in the architecture (see Eq. 2). A classical example are convolutional neural networks (CNNs), which achieve translation equivariance by constraining linear layers to be convolution operators. A wider class of equivariant models are Euclidean steerable CNNs ([Cohen & Welling 2016b](#); [Weiler et al. 2018a](#); [Weiler & Cesa 2019](#); [Jenner & Weiler 2022](#)), which guarantee equivariance to isometries  $\mathbb{R}^n \times G$  of a Euclidean space  $\mathbb{R}^n$ , i.e., to translations and a group  $G$  of origin-preserving transformations, such as rotations and reflections. As proven in [Weiler et al. \(2018a; 2021\)](#); [Jenner & Weiler \(2022\)](#), this requires convolutions with  $G$ -steerable (equivariant) kernels.

Our goal is developing a program to parameterize with minimal requirements arbitrary  $G$ -steerable kernel spaces, with compact  $G$ , which are required to implement  $\mathbb{R}^n \times G$  equivariant CNNs. [Lang & Weiler \(2020\)](#) provides a first step in this direction by generalizing the *Wigner-Eckart theorem* from quantum mechanics to obtain a general technique to parametrize  $G$ -steerable kernel spaces over *orbits* of a compact  $G$ . The theorem reduces the task of building steerable kernel bases to that of finding some pure representation theoretic ingredients. Since the equivariance constraint only relates points  $g.x \in \mathbb{R}^n$  in the same orbit  $G.x \subset \mathbb{R}^n$ , a kernel can take independent values on different orbits. Fig. 1 shows

\*Qualcomm AI Research is an initiative of Qualcomm Technologies, Inc.

# Steerable group convolutions

E(n)-Equivariant Steerable CNNs

[Documentation](#) | [Paper ICLR 22](#) | [Paper NeurIPS 19](#) | [e2cnn library](#) | [e2cnn experiments](#) | [Thesis](#)

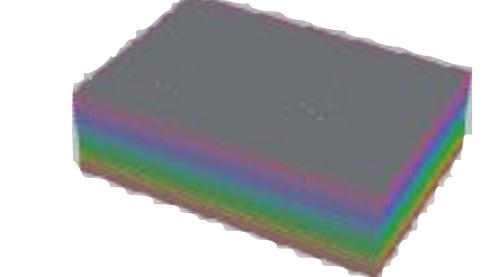
`escnn` is a [PyTorch](#) extension for equivariant deep learning. `escnn` is the successor of the `e2cnn` library, which only supported planar isometries.

<https://github.com/QUVA-Lab/escnn>

Type-0 field usual 2D feature map  
(e.g. for segmentation)



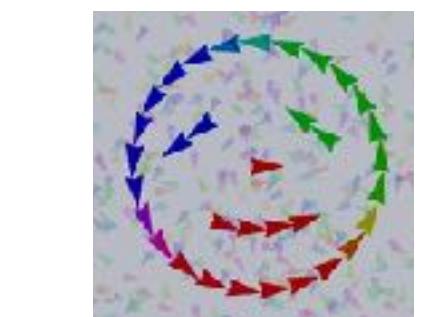
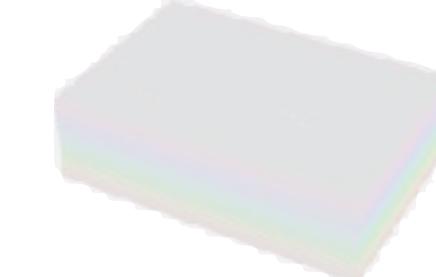
Regular feature types



$\mathcal{K}$

$f$

$f^{out}$



$\mathcal{F}_H$

$f^{out}$

$\hat{\mathcal{K}}$

$\hat{f}^{out}$

$\hat{\mathcal{K}}$

$\hat{\mathcal{K}}$

$\hat{f}^{out}$

$\mathcal{F}_H$

$\hat{f}^{out}$

$\hat{\mathcal{K}}$

$\hat{\mathcal{K}}$

$\hat{f}^{out}$

$\mathcal{K}$

$\hat{\mathcal{K}}$

&lt;

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

Efficient (band-limited) **grid-free** g-convs

7. Feature fields and escnn library

Flexible framework for equivariant layers

8. Equivariant tensor product layers

9. Equivariant graph NNs

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

Efficient (band-limited) **grid-free** g-convs

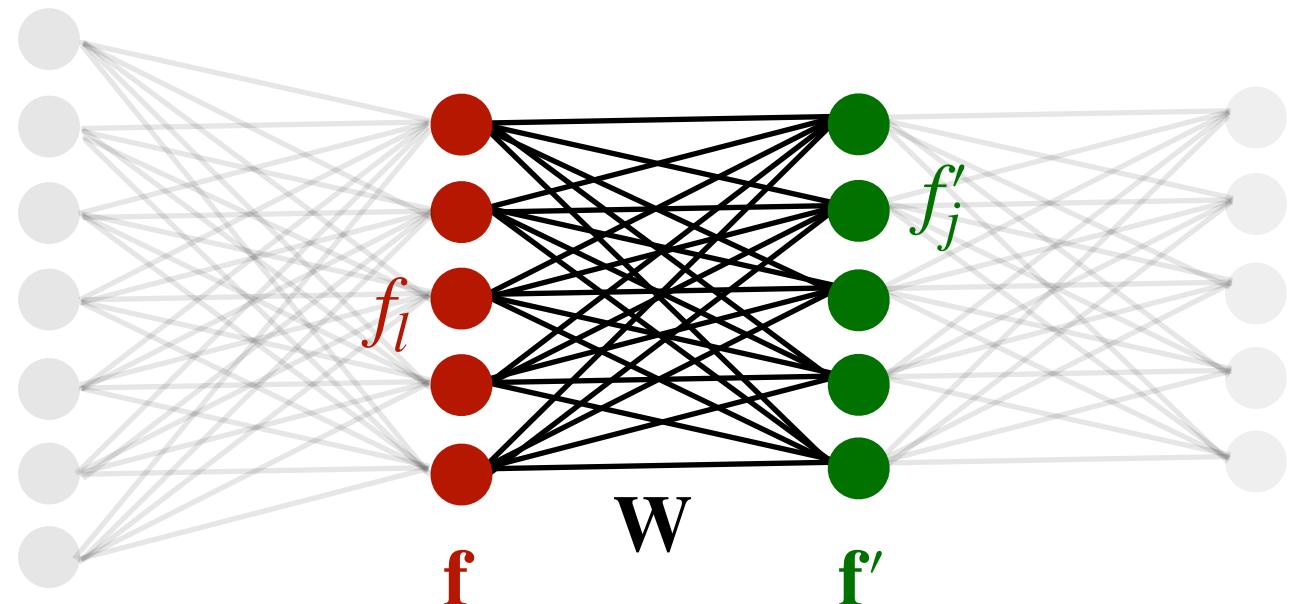
7. Feature fields and escnn library

Flexible framework for equivariant layers

8. Equivariant tensor product layers

9. Equivariant graph NNs

# Equivariant MLP



$$\begin{array}{c} \mathbf{f} \\ \text{linear layer} \end{array} \mapsto \mathbf{f}' = \mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) \begin{array}{c} \mathbf{f} \\ \text{activation} \end{array} \mapsto \mathbf{f}'' = \sigma(\mathbf{f}')$$

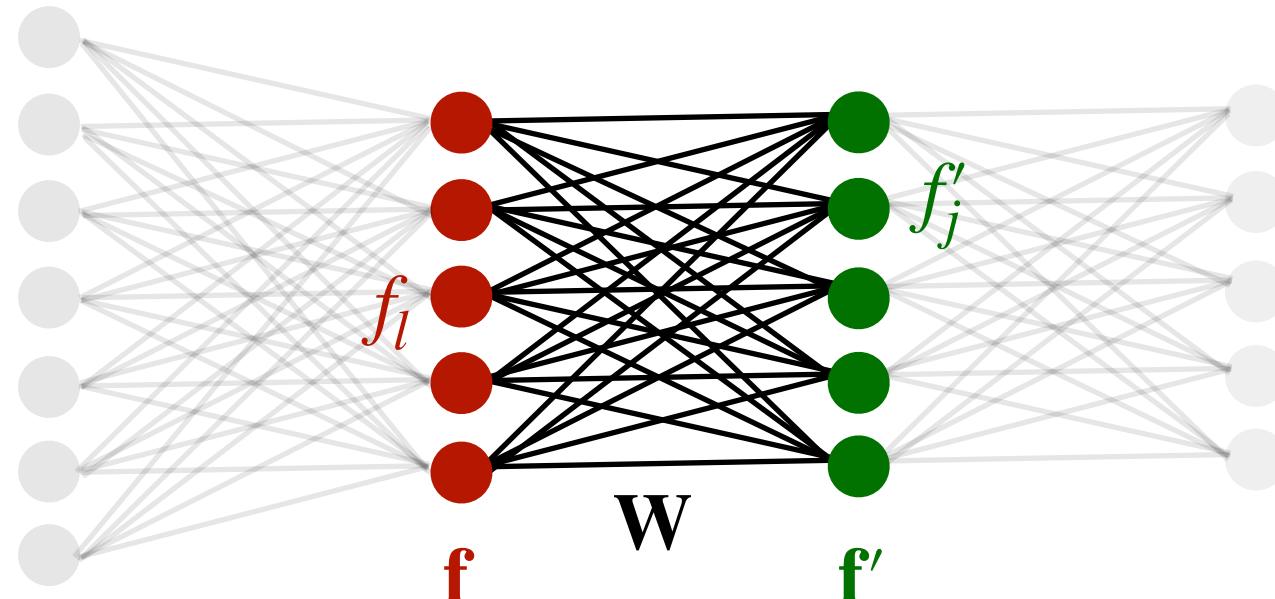
( Repeat  $L$  times )

**Linear layer** (matrix-vector multiplication)

$$\mathbf{f}' = \mathbf{W} \mathbf{f}$$

$$f'_j = \sum_l w_l^j f_l$$

# Equivariant MLP



$\mathbf{f} \xrightarrow{\text{linear layer}} \mathbf{f}' = \mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) \mathbf{f} \xrightarrow{\text{activation}} \mathbf{f}'' = \sigma(\mathbf{f}')$

( Repeat  $L$  times )

**Linear layer** (matrix-vector multiplication)

$$\mathbf{f}' = \mathbf{W} \mathbf{f}$$

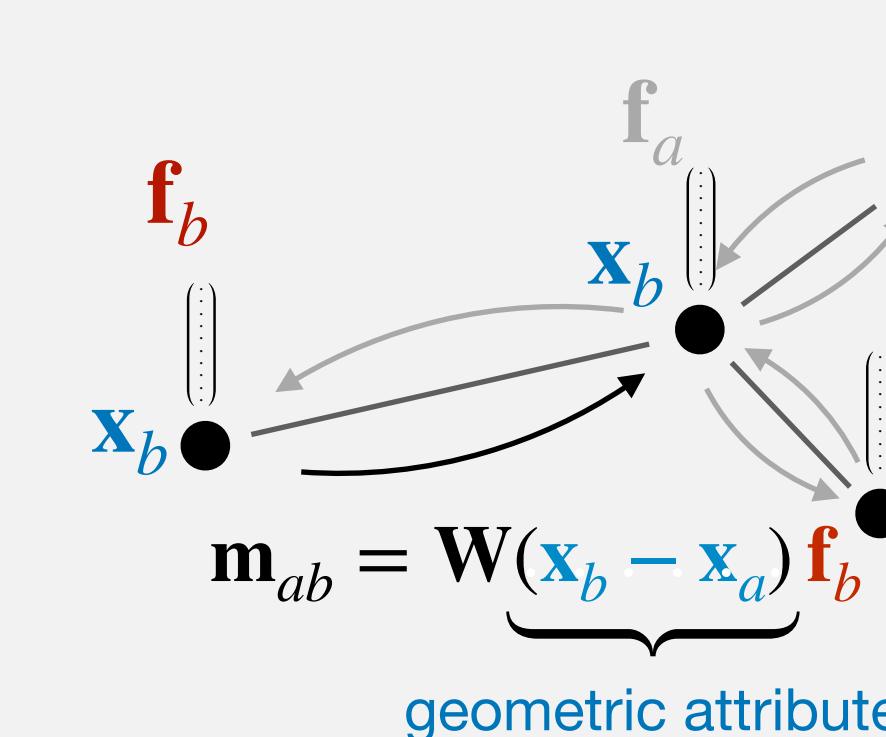
$$f'_j = \sum_l w_l^j f_l$$

**Conditional linear layer** (weight matrix depends on  $\mathbf{x}_b - \mathbf{x}_a$ )

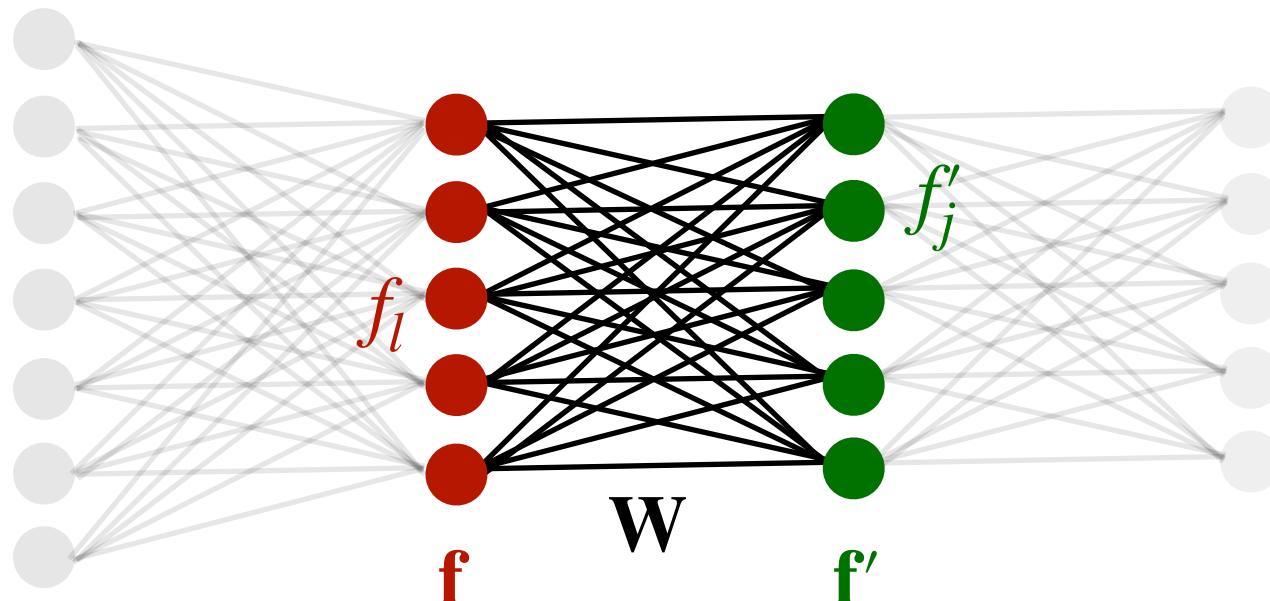
$$\mathbf{f}' = \mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) \mathbf{f}$$

$$f'_j = \sum_l w_l^j (\mathbf{x}_b - \mathbf{x}_a) f_l$$

Convolutional message passing



# Equivariant MLP



$$\begin{array}{c} \mathbf{f} \\ \text{linear layer} \end{array} \mapsto \mathbf{f}' = \mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) \mathbf{f} \mapsto \begin{array}{c} \mathbf{f}'' = \sigma(\mathbf{f}') \\ \text{activation} \end{array}$$

( Repeat  $L$  times )

**Linear layer** (matrix-vector multiplication)

$$\mathbf{f}' = \mathbf{W} \mathbf{f}$$

$$f'_j = \sum_l w_l^j f_l$$

**Conditional linear layer** (weight matrix depends on  $\mathbf{x}_b - \mathbf{x}_a$ )

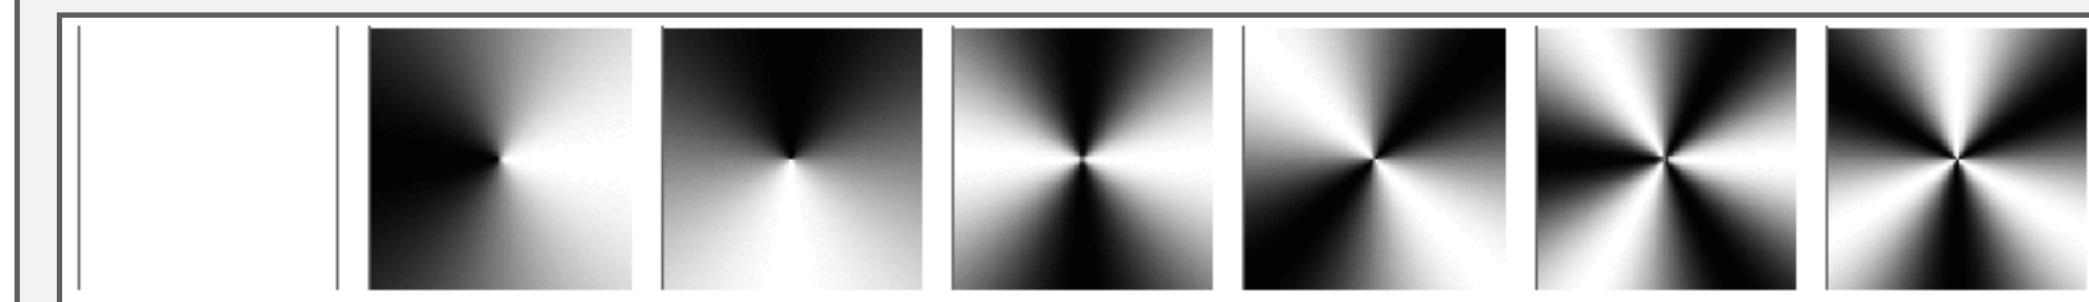
$$\mathbf{f}' = \mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) \mathbf{f}$$

$$f'_j = \sum_l w_l^j (\mathbf{x}_b - \mathbf{x}_a) f_l$$

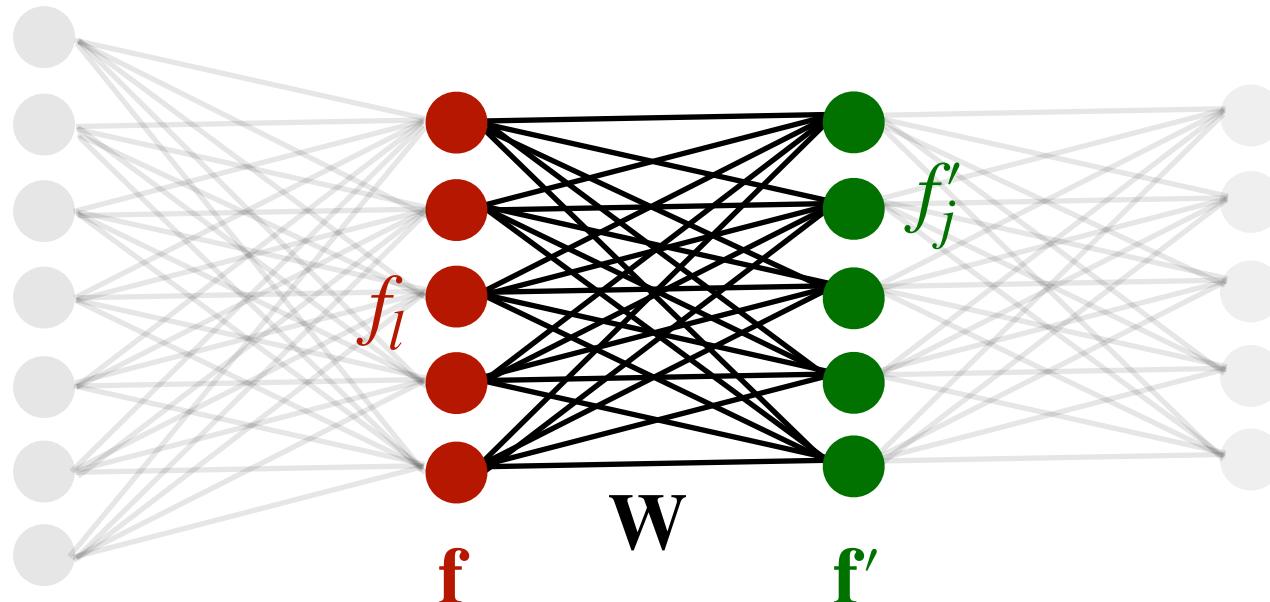
Let  $\mathbf{W} : \mathbb{R}^3 \rightarrow \mathbb{R}^{C \times C}$  a matrix valued function (conv-kernel)

- Expanded in a basis  $\mathbf{Y}(\mathbf{x}) = \begin{pmatrix} \vdots \\ Y_J(\mathbf{x}) \\ \vdots \end{pmatrix}$
- Basis (coordinate embedding) functions  $Y_J : \mathbb{R}^3 \rightarrow \mathbb{R}$
- Matrix-valued weights  $\mathbf{W}_J$  with elements  $w_{Jl}^j$

$$\mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) = \sum_J \mathbf{W}_J Y_J(\mathbf{x}_b - \mathbf{x}_a)$$



# Equivariant MLP



$$\begin{array}{c} \mathbf{f} \\ \text{linear layer} \end{array} \mapsto \mathbf{f}' = \mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) \mathbf{f} \mapsto \mathbf{f}'' = \sigma(\mathbf{f}') \quad \text{(Repeat } L \text{ times)}$$

activation

**Linear layer** (matrix-vector multiplication)

$$\mathbf{f}' = \mathbf{W} \mathbf{f}$$

$$f'_j = \sum_l w_l^j f_l$$

**Conditional linear layer** (weight matrix depends on  $\mathbf{x}_b - \mathbf{x}_a$ )

$$\mathbf{f}' = \mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) \mathbf{f}$$

$$f'_j = \sum_l w_l^j (\mathbf{x}_b - \mathbf{x}_a) f_l$$

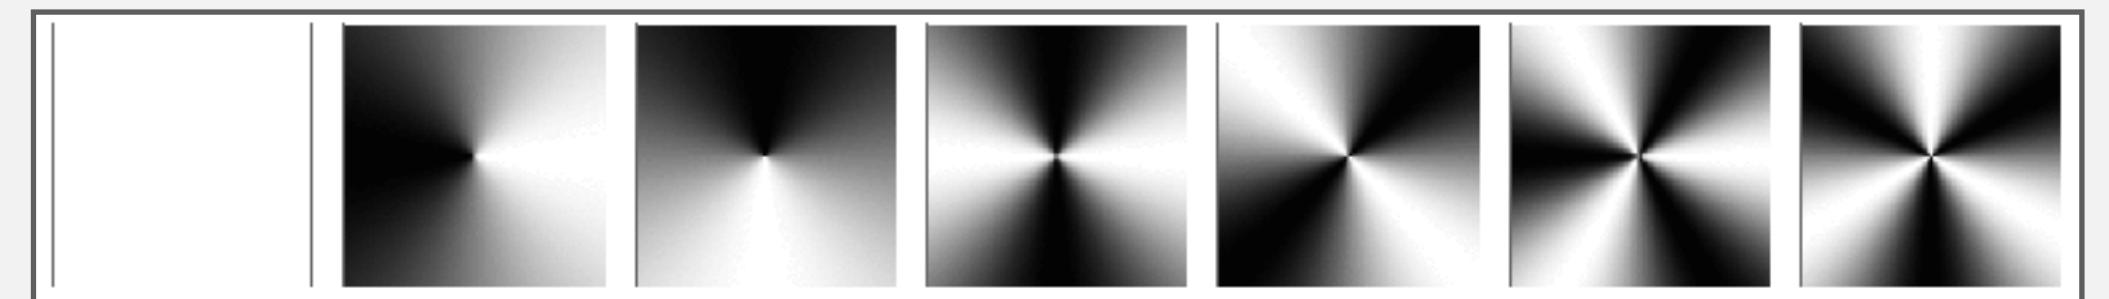
$$\mathbf{f}' = \mathbf{f} \underset{\text{bilinear}}{W} Y_J(\mathbf{x}_b - \mathbf{x}_a)$$

$$f'_j = \sum_l \sum_J w_{Jl}^j Y_J(\mathbf{x}_b - \mathbf{x}_a) f_l$$

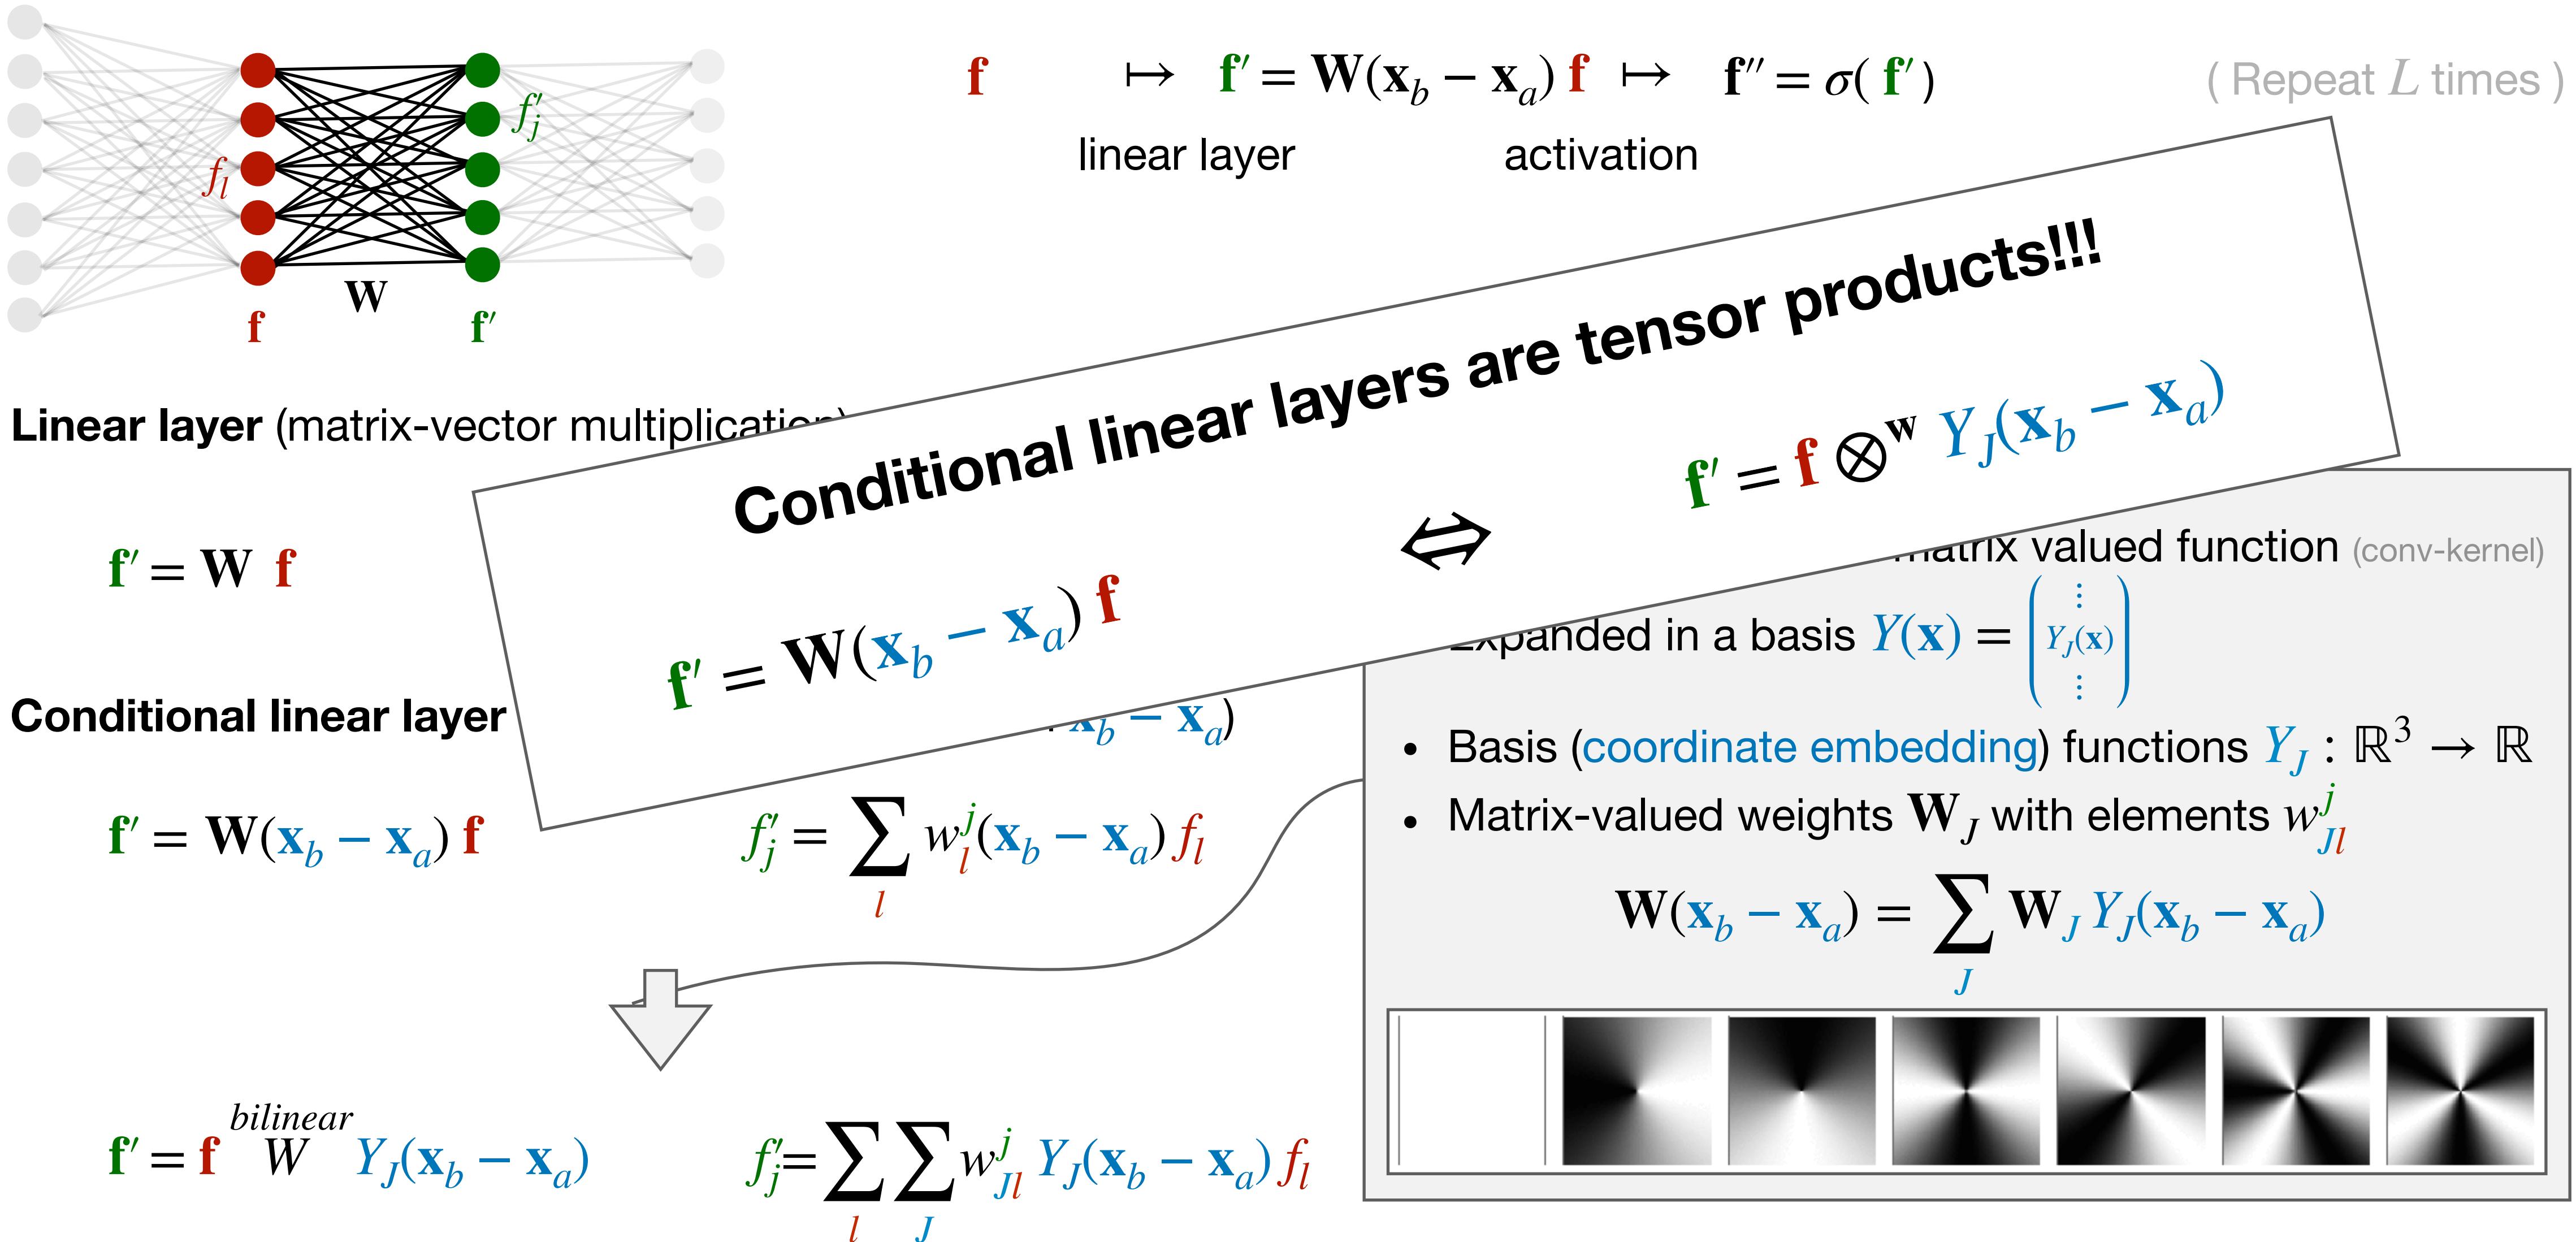
Let  $\mathbf{W} : \mathbb{R}^3 \rightarrow \mathbb{R}^{C \times C}$  a matrix valued function (conv-kernel)

- Expanded in a basis  $\mathbf{Y}(\mathbf{x}) = \begin{pmatrix} \vdots \\ Y_J(\mathbf{x}) \\ \vdots \end{pmatrix}$
- Basis (coordinate embedding) functions  $Y_J : \mathbb{R}^3 \rightarrow \mathbb{R}$
- Matrix-valued weights  $\mathbf{W}_J$  with elements  $w_{Jl}^j$

$$\mathbf{W}(\mathbf{x}_b - \mathbf{x}_a) = \sum_J \mathbf{W}_J Y_J(\mathbf{x}_b - \mathbf{x}_a)$$



# Equivariant MLP



1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

Efficient (band-limited) grid-free g-convs

7. Feature fields and escnn library

Flexible framework for equivariant layers

8. Equivariant tensor product layers

**Conv layers  $\leftrightarrow$  TPs with coordinate embeddings**  
(Clebsch-Gordan: equivariant TP)

9. Equivariant graph NNs

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

Efficient (band-limited) **grid-free** g-convs

7. Feature fields and escnn library

Flexible framework for equivariant layers

8. Equivariant tensor product layers

Conv layers  $\leftrightarrow$  TPs with coordinate embeddings  
(Clebsch-Gordan: equivariant TP)

9. Equivariant graph NNs

# Neural Message Passing for Quantum Chemistry

Justin Gilmer<sup>1</sup> Samuel S. Schoenholz<sup>1</sup> Patrick F. Riley<sup>2</sup> Oriol Vinyals<sup>3</sup> George E. Dahl<sup>1</sup>

## Abstract

Supervised learning on molecules has incredible potential to be useful in chemistry, drug discovery, and materials science. Luckily, several promising and closely related neural network models invariant to molecular symmetries have already been described in the literature. These models learn a message passing algorithm and aggregation procedure to compute a function of their entire input graph. At this point, the next step is to find a particularly effective variant of this general approach and apply it to chemical prediction benchmarks until we either solve them or reach the limits of the approach. In this paper, we reformulate existing models into a single common framework we call Message Passing Neural Networks (MPNNs) and explore additional novel variations within this framework. Using MPNNs we demonstrate state of the art results on an important molecular property prediction benchmark; these results are strong enough that we believe future work should focus on datasets with larger molecules or more accurate ground truth labels.

## 1. Introduction

The past decade has seen remarkable success in the use of deep neural networks to understand and translate natural language (Wu et al., 2016), generate and decode complex audio signals (Hinton et al., 2012), and infer features from real-world images and videos (Krizhevsky et al., 2012). Although chemists have applied machine learning to many problems over the years, predicting the properties of molecules and materials using machine learning (and especially deep learning) is still in its infancy. To date, most research applying machine learning to chemistry tasks (Hansen et al., 2015; Huang & von Lilienfeld, 2016;

<sup>1</sup>Google Brain <sup>2</sup>Google <sup>3</sup>Google DeepMind. Correspondence to: Justin Gilmer <[gilmer@google.com](mailto:gilmer@google.com)>, George E. Dahl <[gdahl@google.com](mailto:gdahl@google.com)>.

*Proceedings of the 34<sup>th</sup> International Conference on Machine Learning*, Sydney, Australia, PMLR 70, 2017. Copyright 2017 by the author(s).

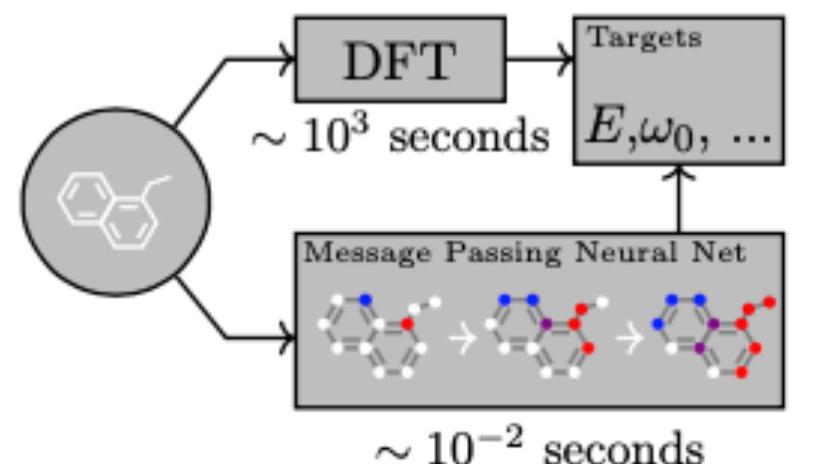


Figure 1. A Message Passing Neural Network predicts quantum properties of an organic molecule by modeling a computationally expensive DFT calculation.

Rupp et al., 2012; Rogers & Hahn, 2010; Montavon et al., 2012; Behler & Parrinello, 2007; Schoenholz et al., 2016) has revolved around feature engineering. While neural networks have been applied in a variety of situations (Merkwirth & Lengauer, 2005; Micheli, 2009; Lusci et al., 2013; Duvenaud et al., 2015), they have yet to become widely adopted. This situation is reminiscent of the state of image models before the broad adoption of convolutional neural networks and is due, in part, to a dearth of empirical evidence that neural architectures with the appropriate inductive bias can be successful in this domain.

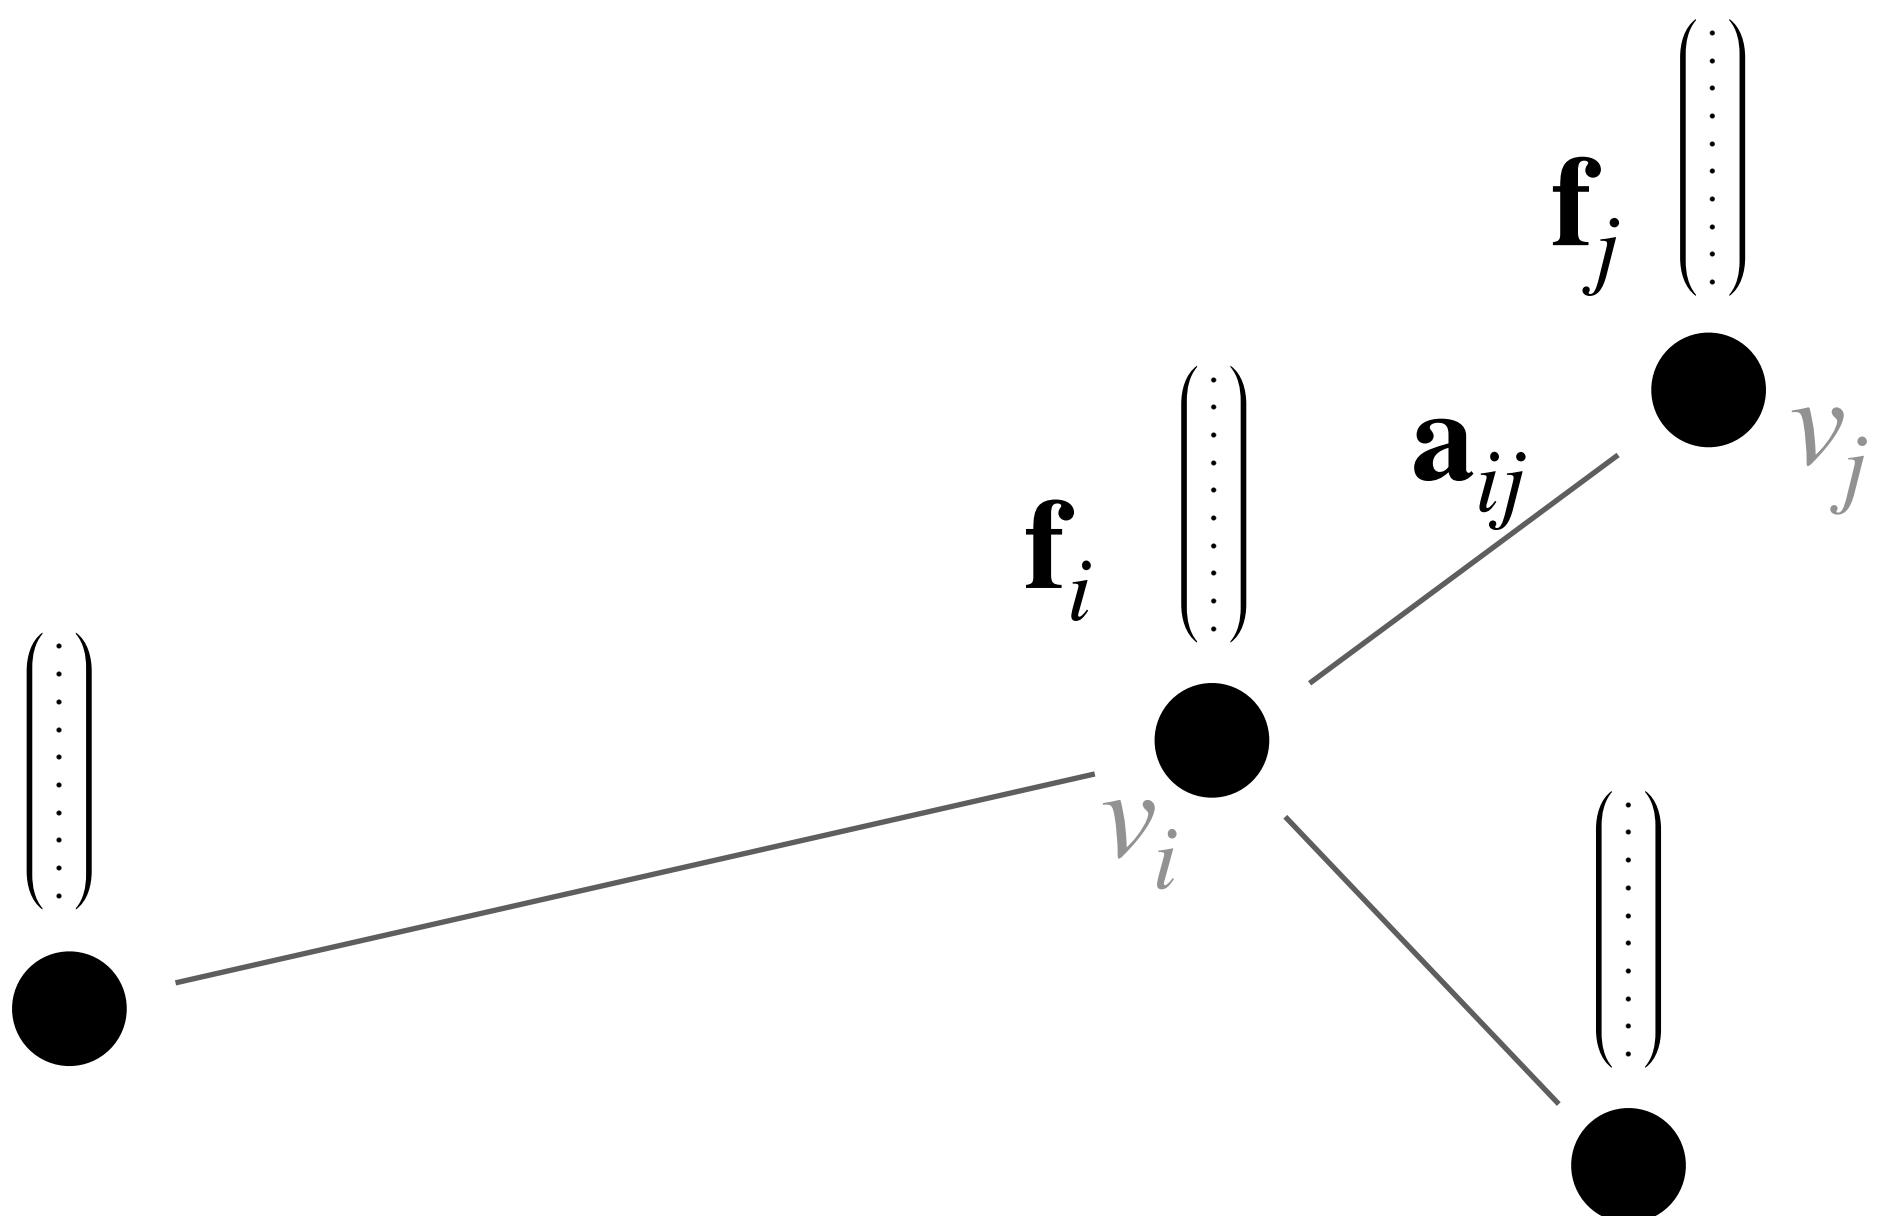
Recently, large scale quantum chemistry calculation and molecular dynamics simulations coupled with advances in high throughput experiments have begun to generate data at an unprecedented rate. Most classical techniques do not make effective use of the larger amounts of data that are now available. The time is ripe to apply more powerful and flexible machine learning methods to these problems, assuming we can find models with suitable inductive biases. The symmetries of atomic systems suggest neural networks that operate on graph structured data and are invariant to graph isomorphism might also be appropriate for molecules. Sufficiently successful models could someday help automate challenging chemical search problems in drug discovery or materials science.

In this paper, our goal is to demonstrate effective machine learning models for chemical prediction problems

# The Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

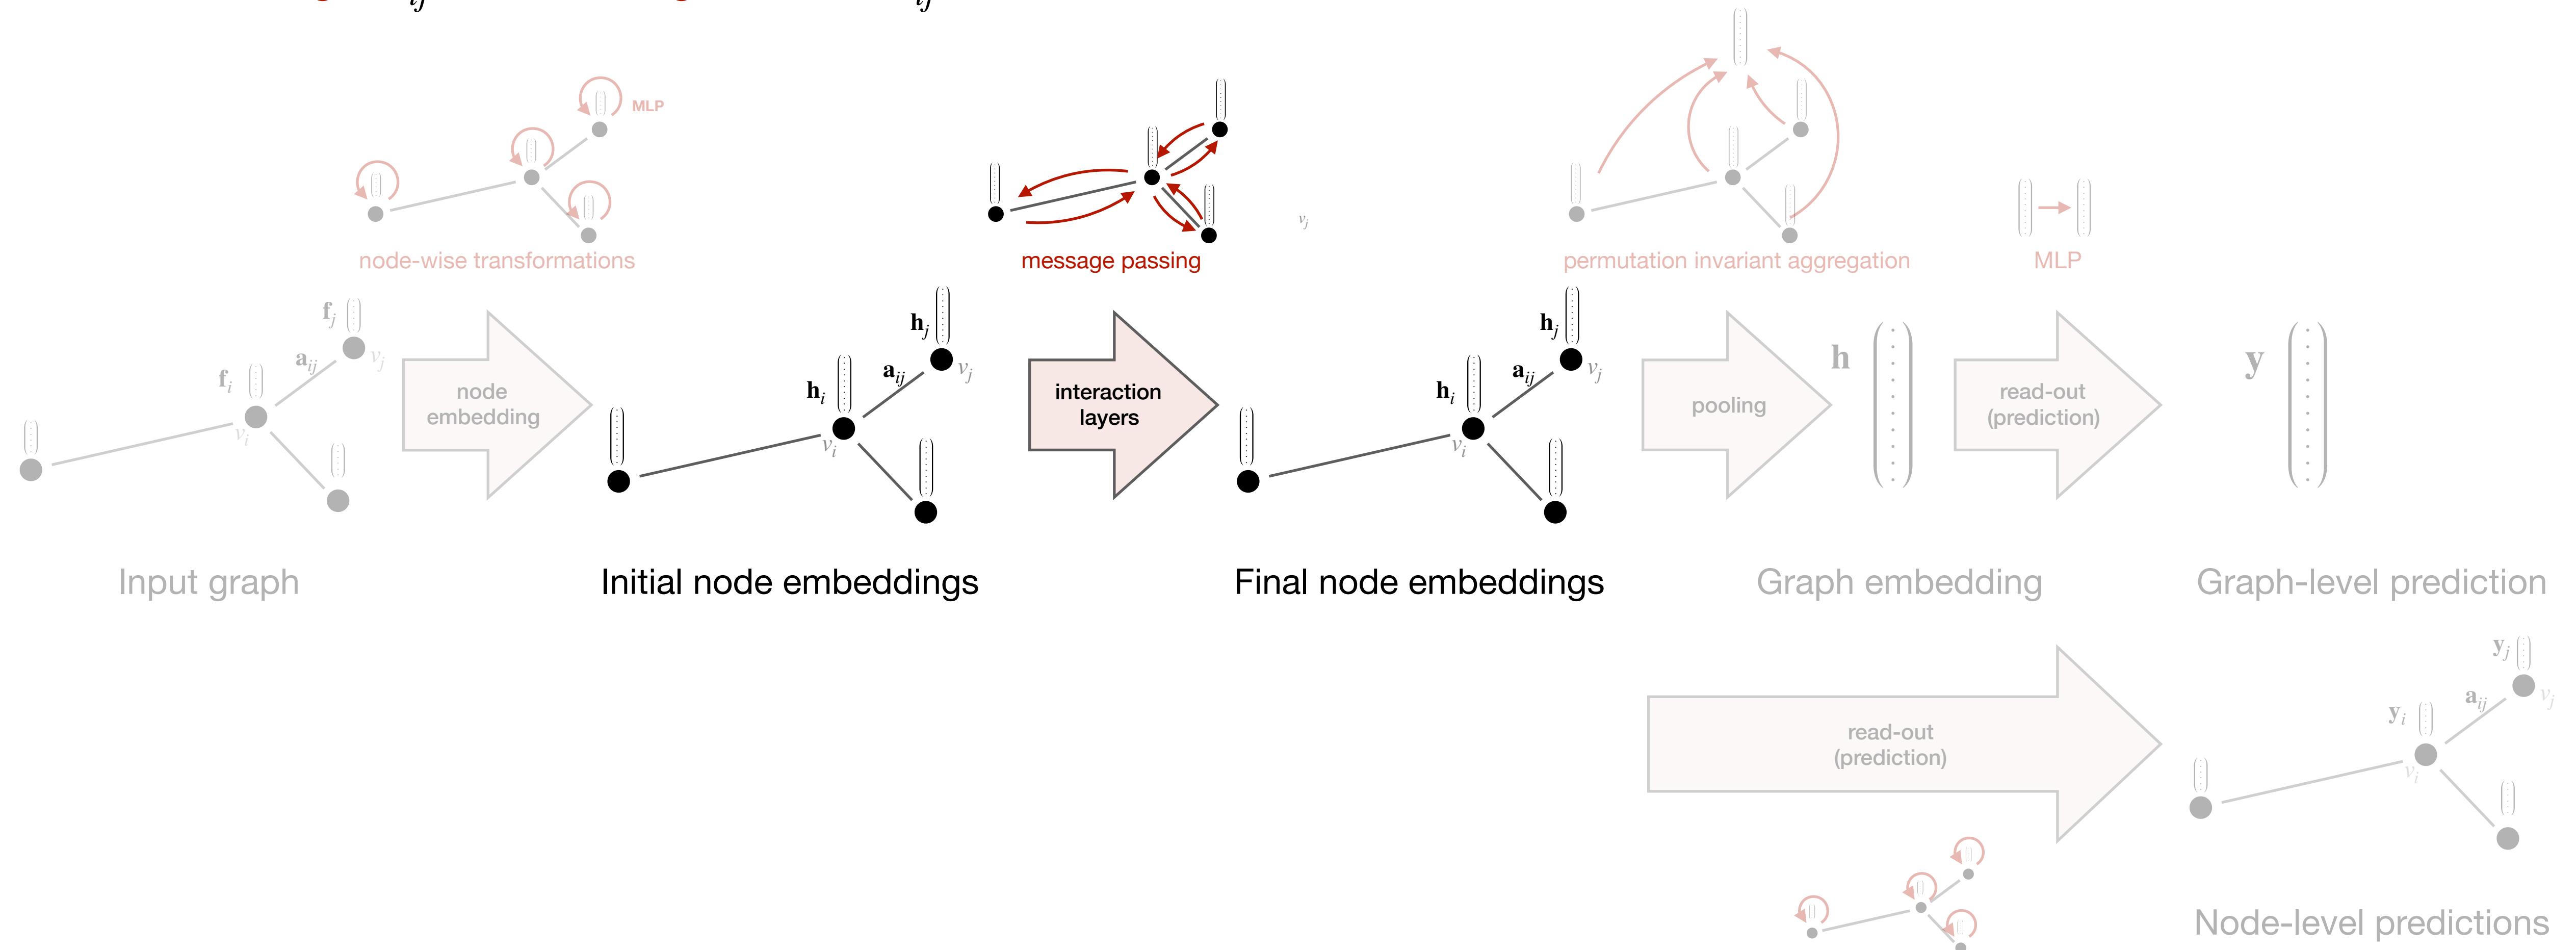


# The Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

**Goal:** iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

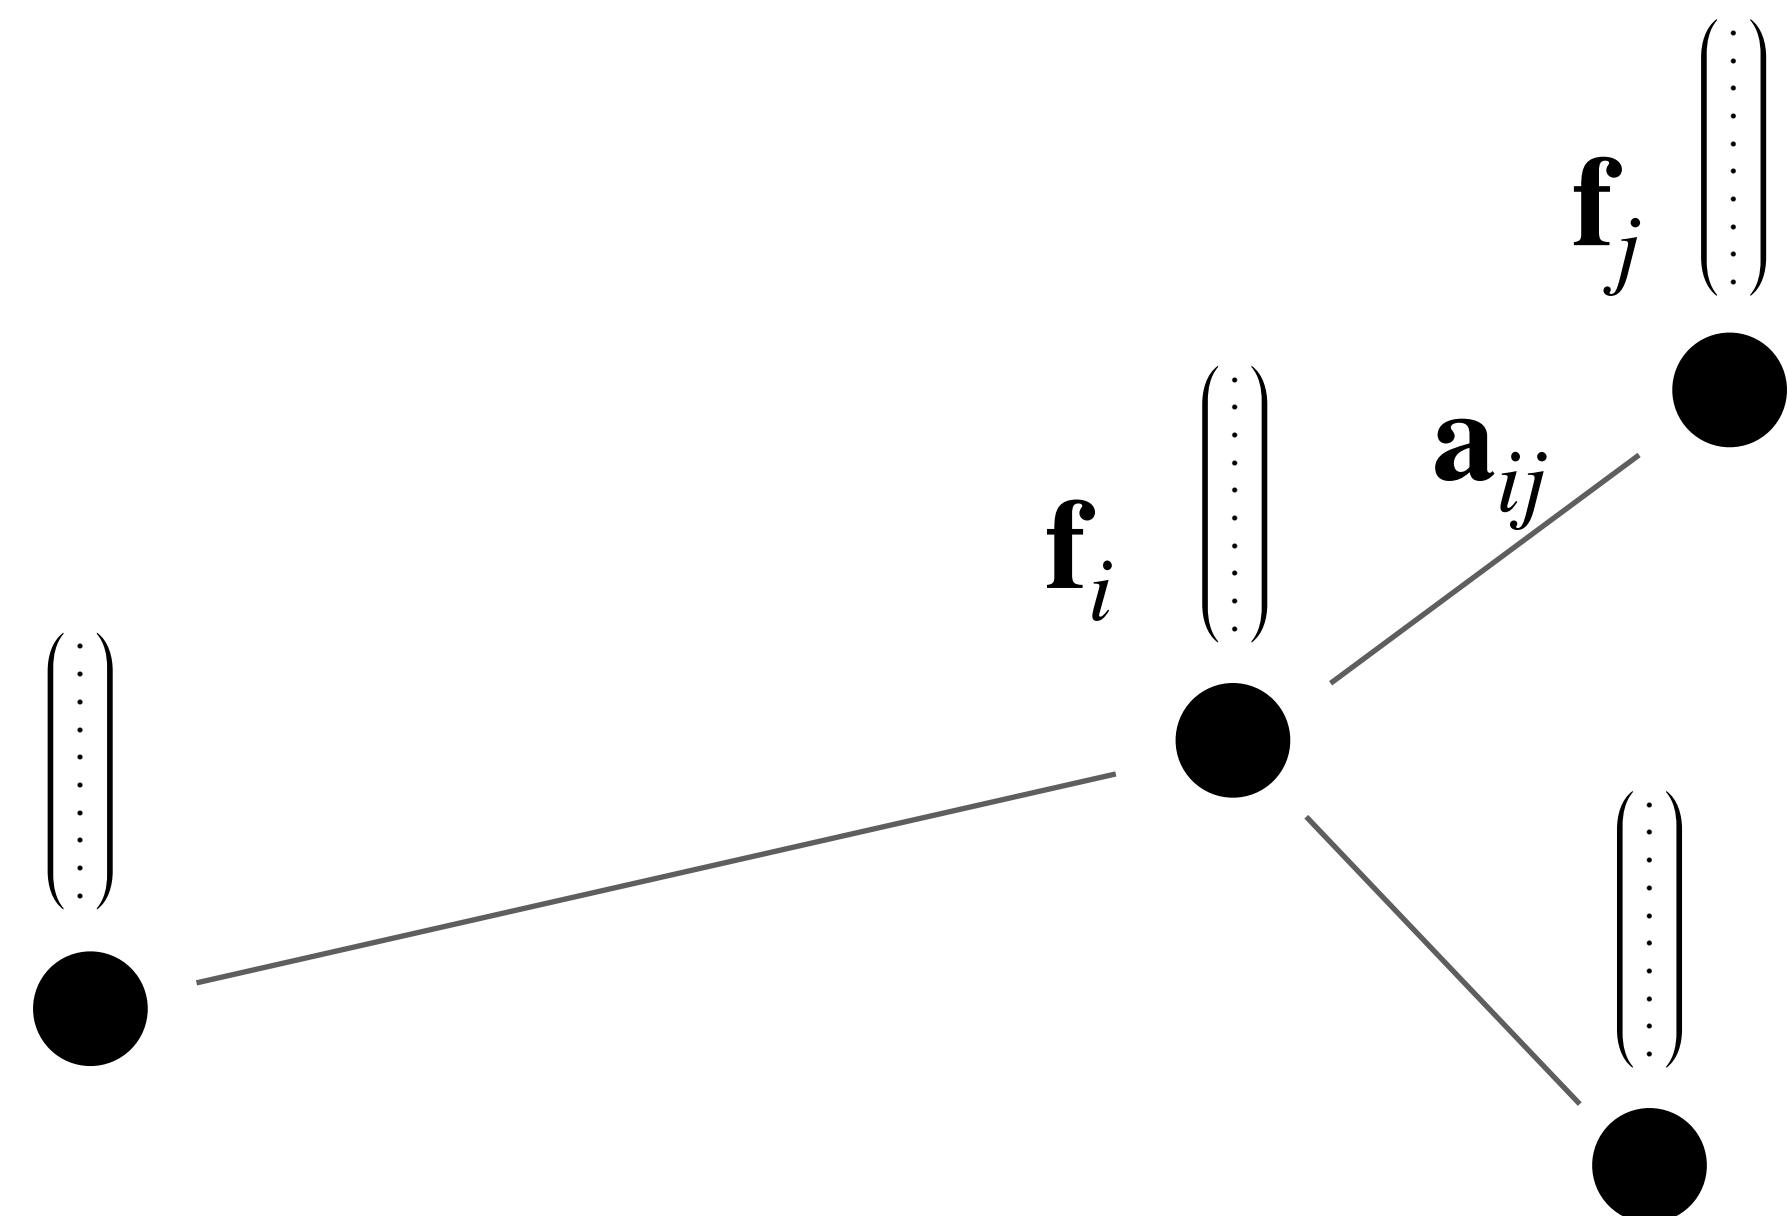


# The Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$



Message passing layer:

- Messages

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

- Aggregate + node updates

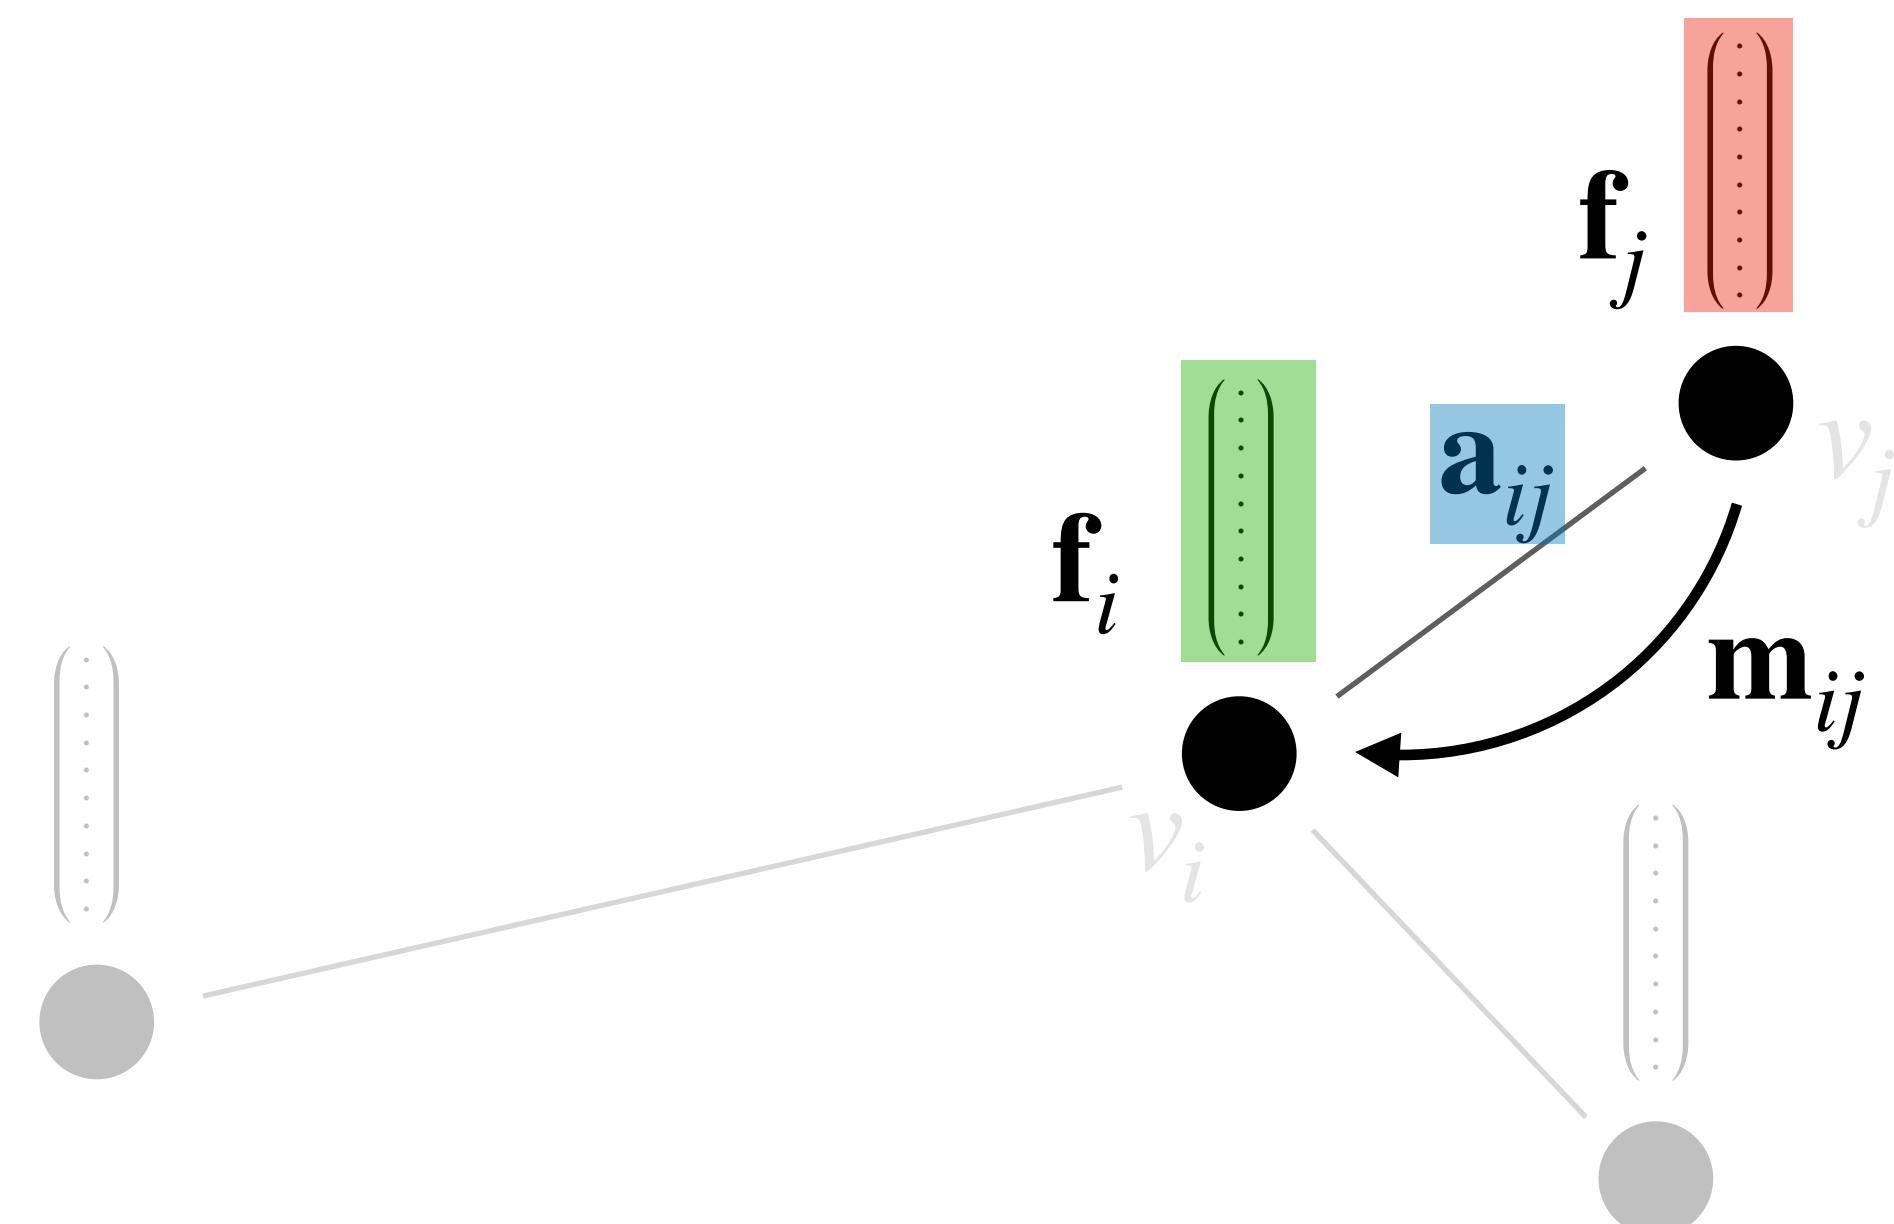
$$\mathbf{f}'_i = \phi_f(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij})$$

# The Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$



Message passing layer:

- Messages

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

- Aggregate + node updates

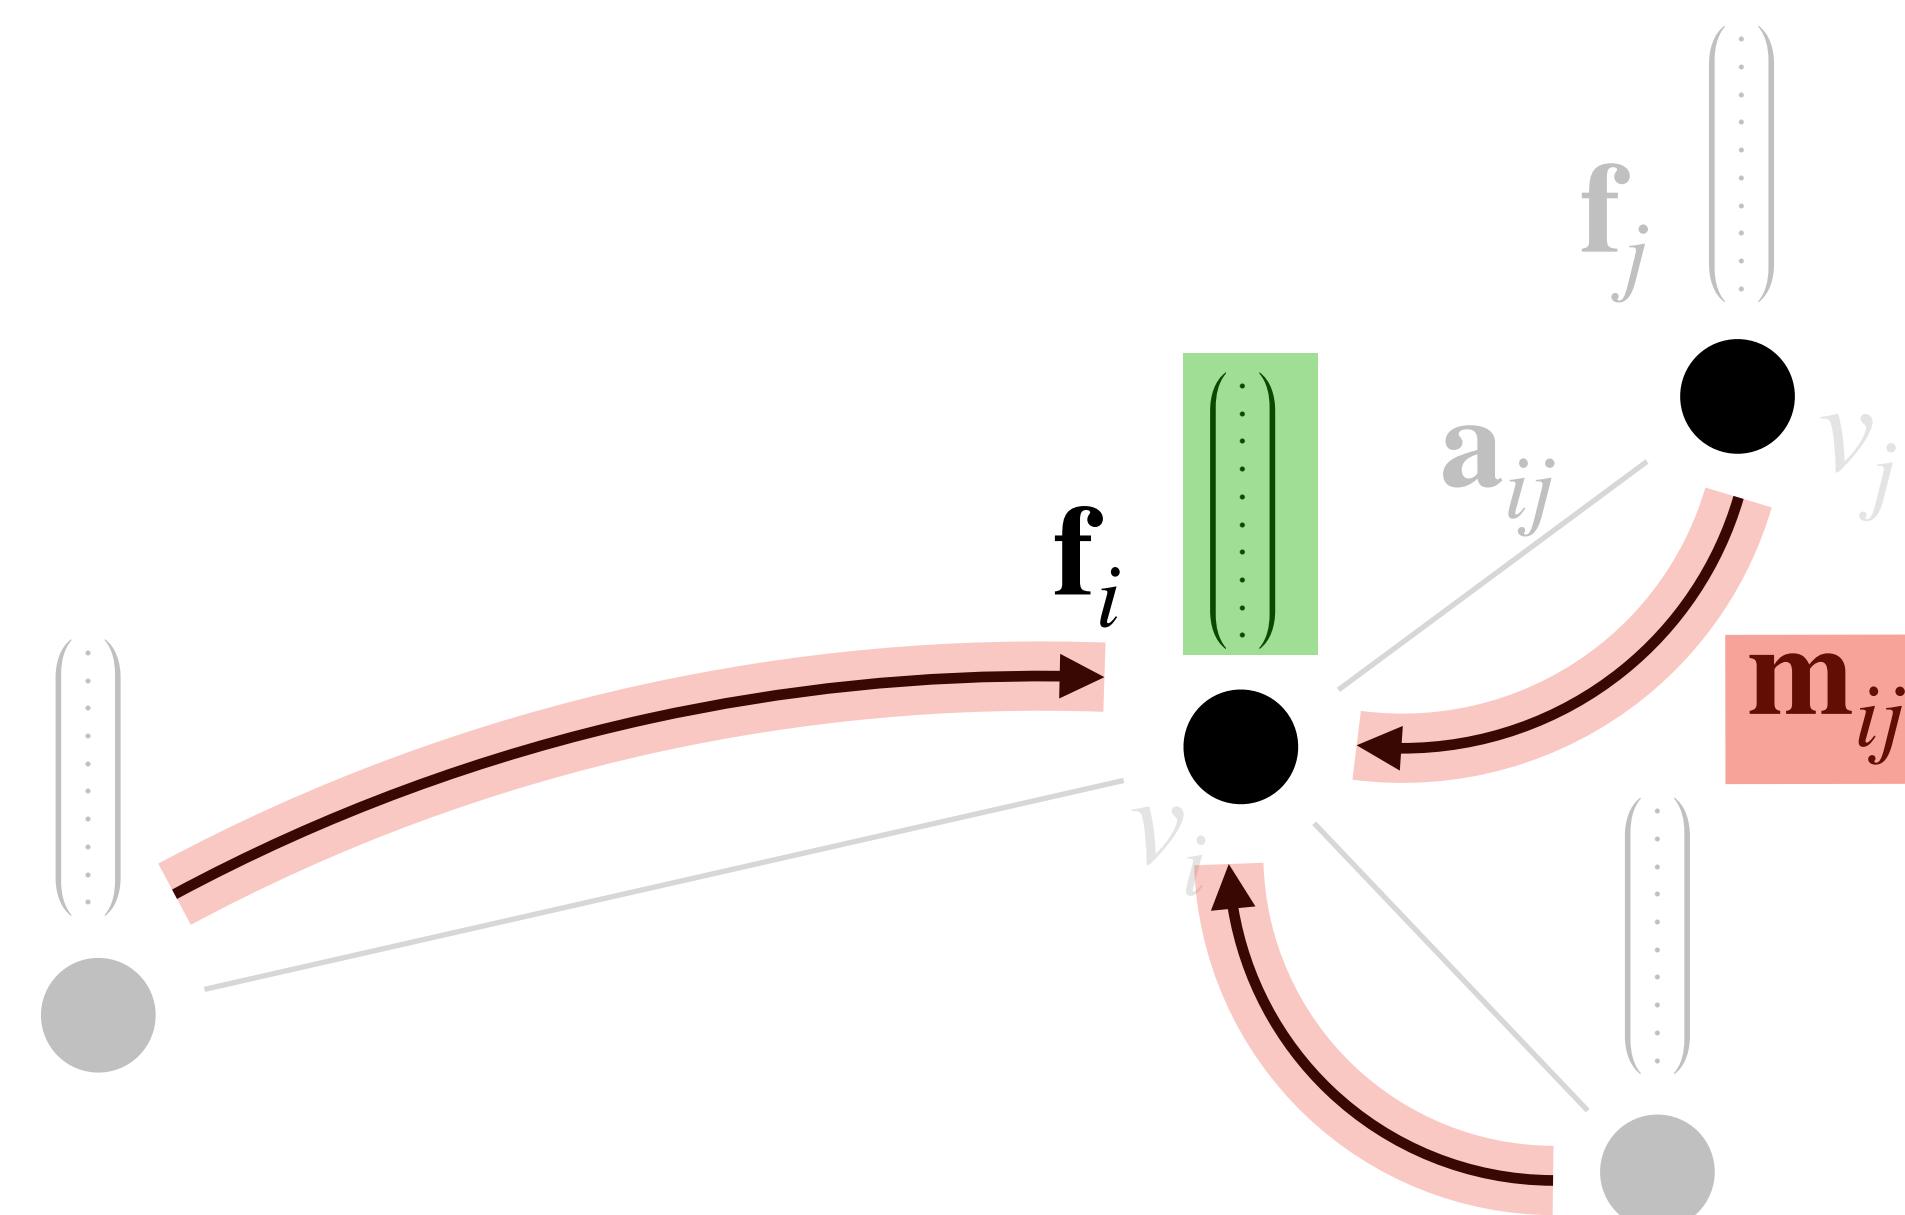
$$\mathbf{f}'_i = \phi_f(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij})$$

# The Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$



Message passing layer:

- Messages

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

- Aggregate + node updates

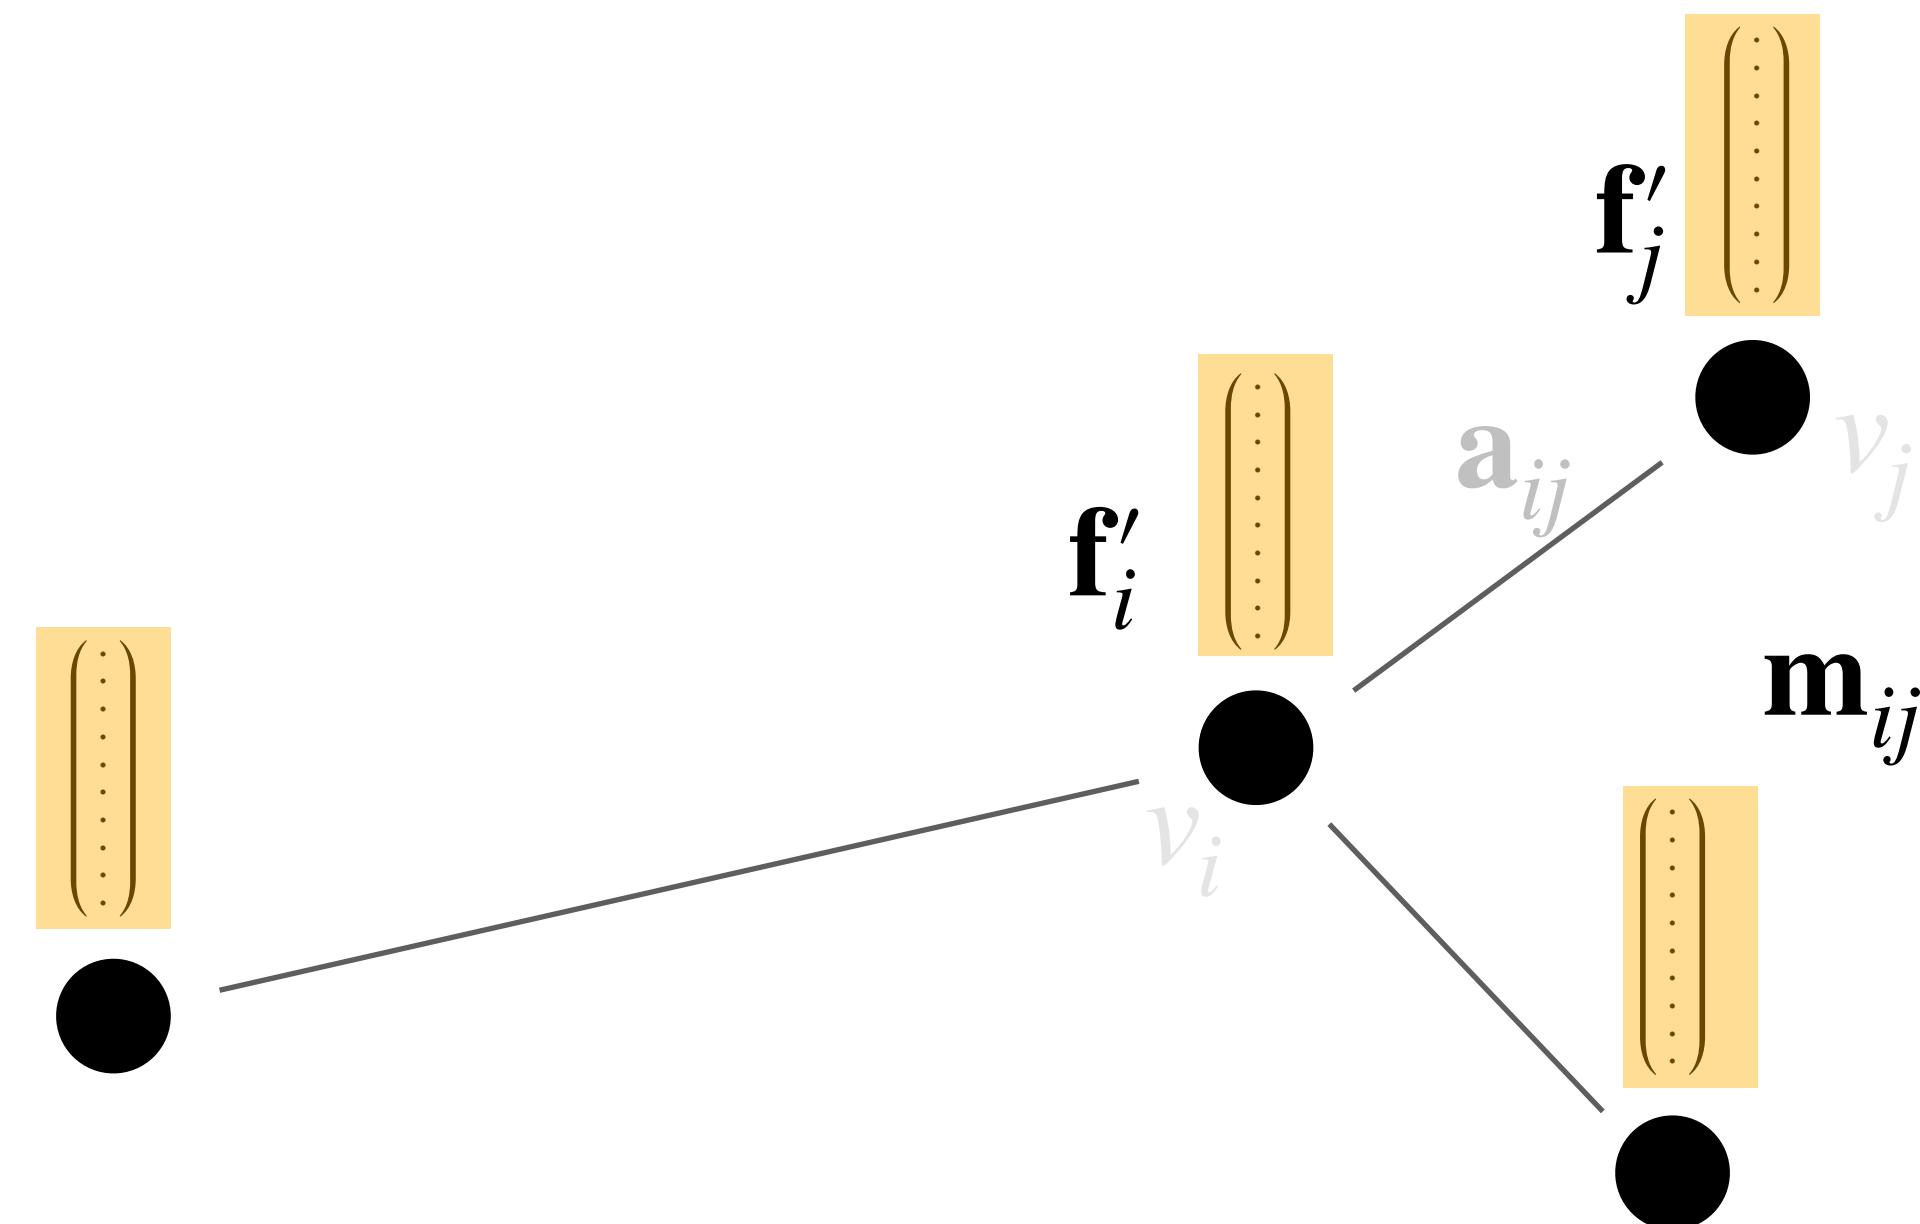
$$\mathbf{f}'_i = \phi_f(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij})$$

# The Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$



Message passing layer:

- Messages

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

- Aggregate + node updates

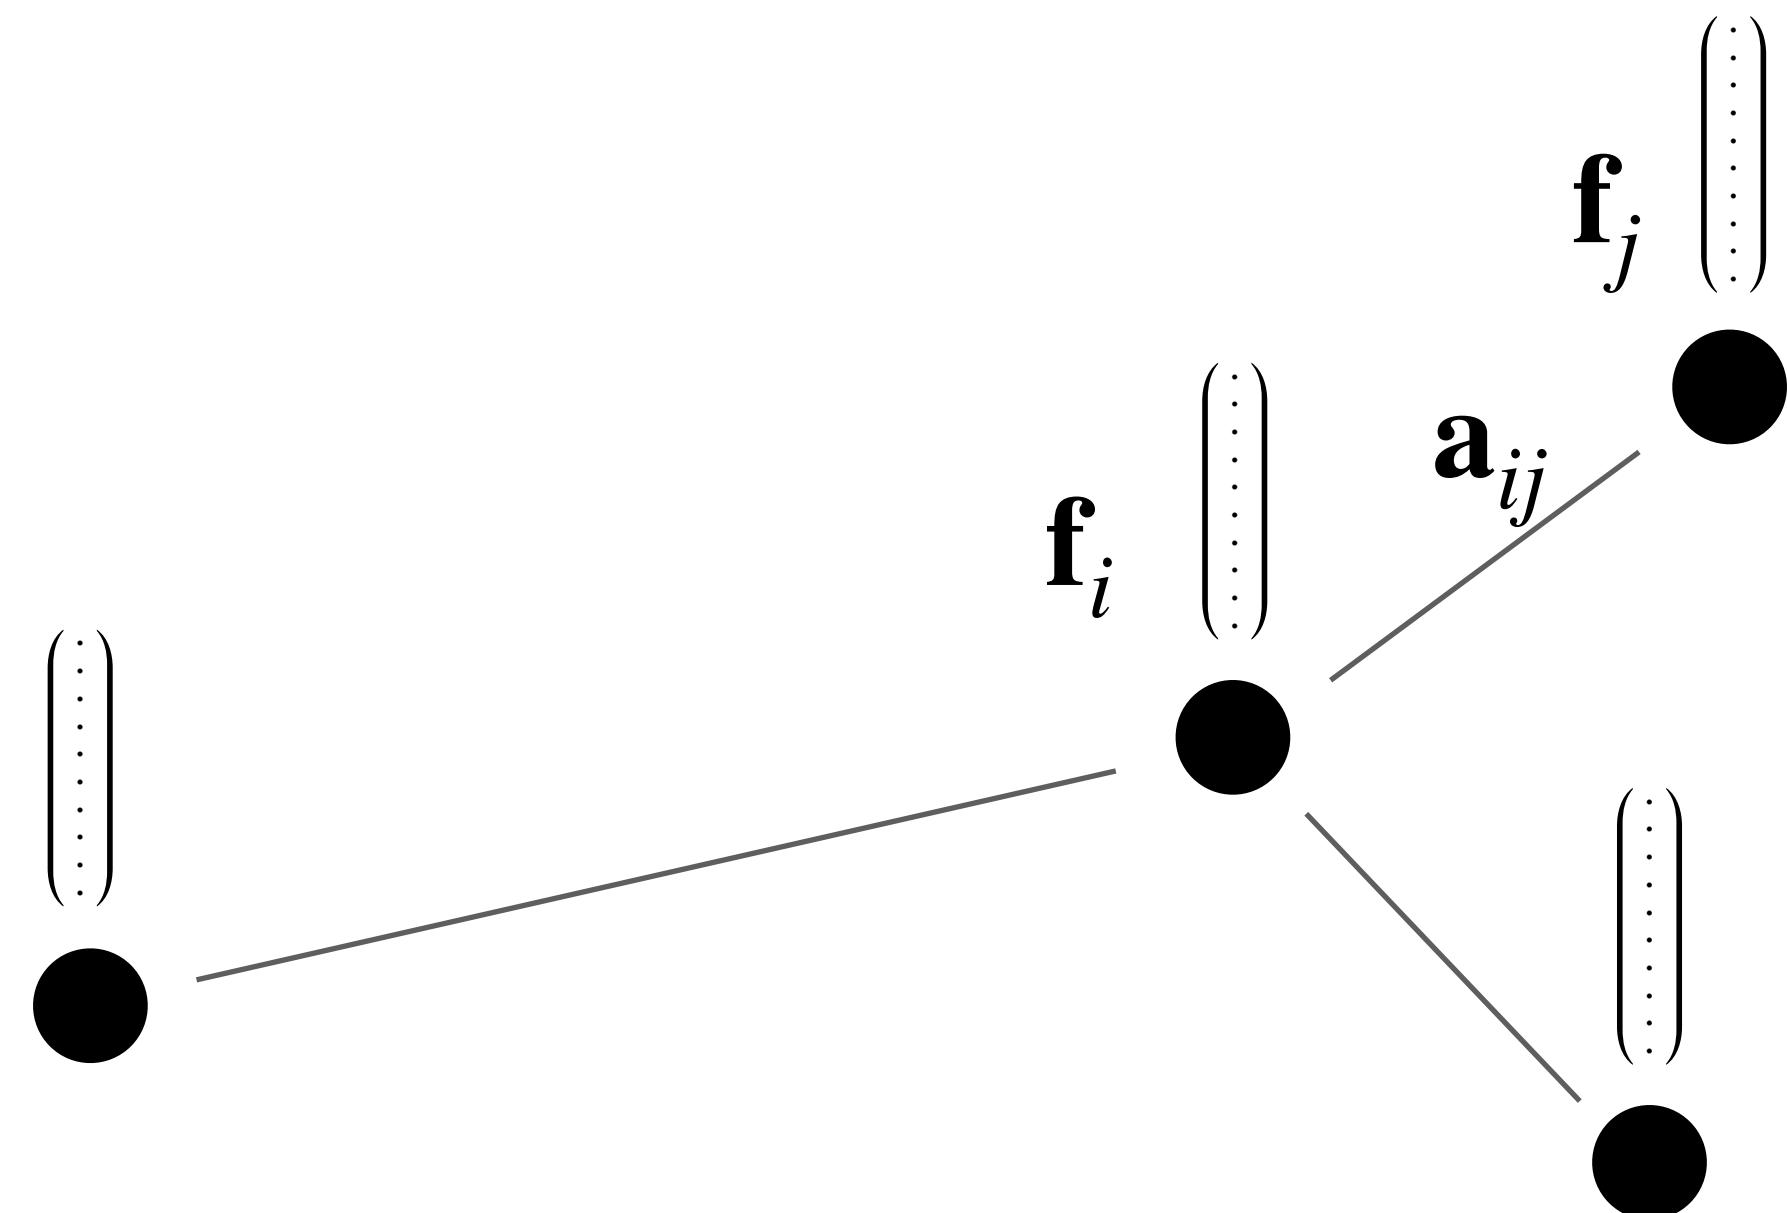
$$\mathbf{f}'_i = \phi_f(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij})$$

# The Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$



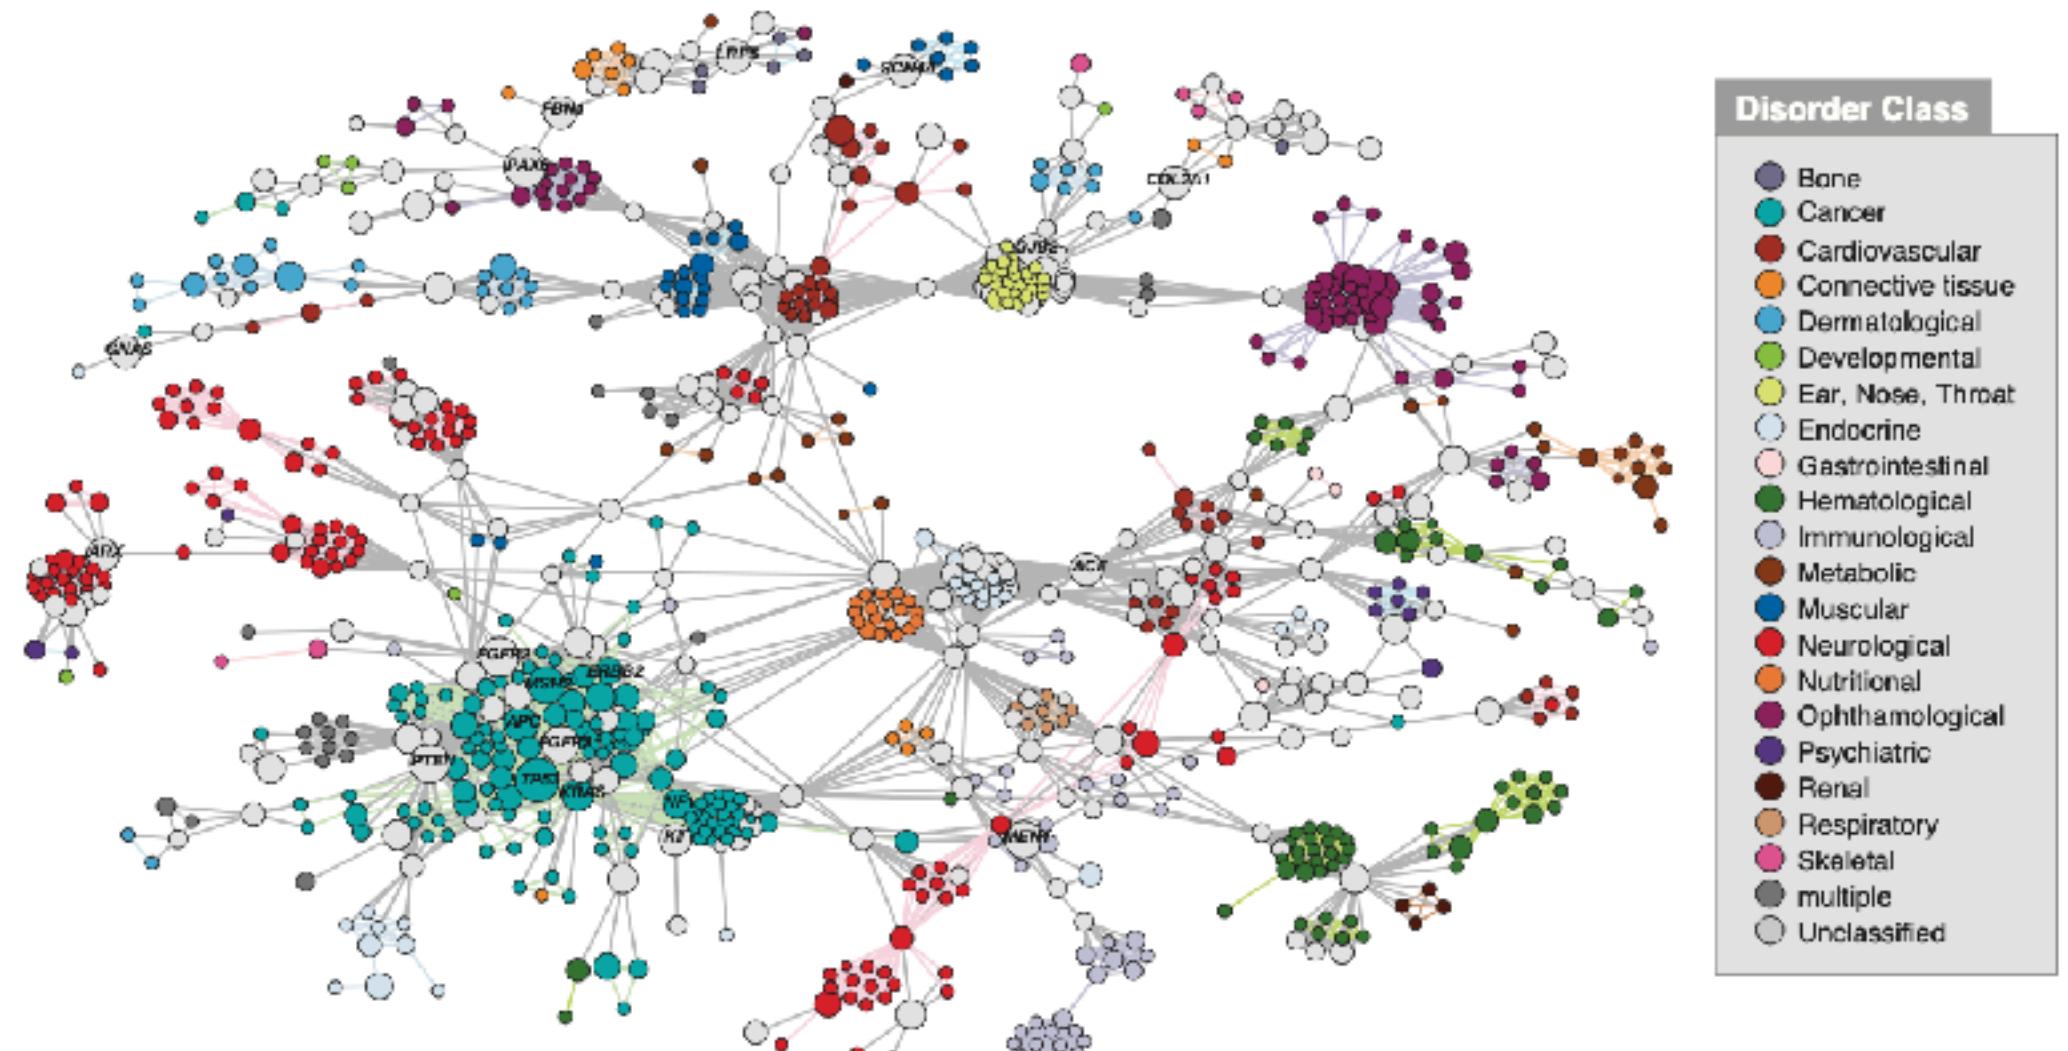
Message passing layer:

- Messages

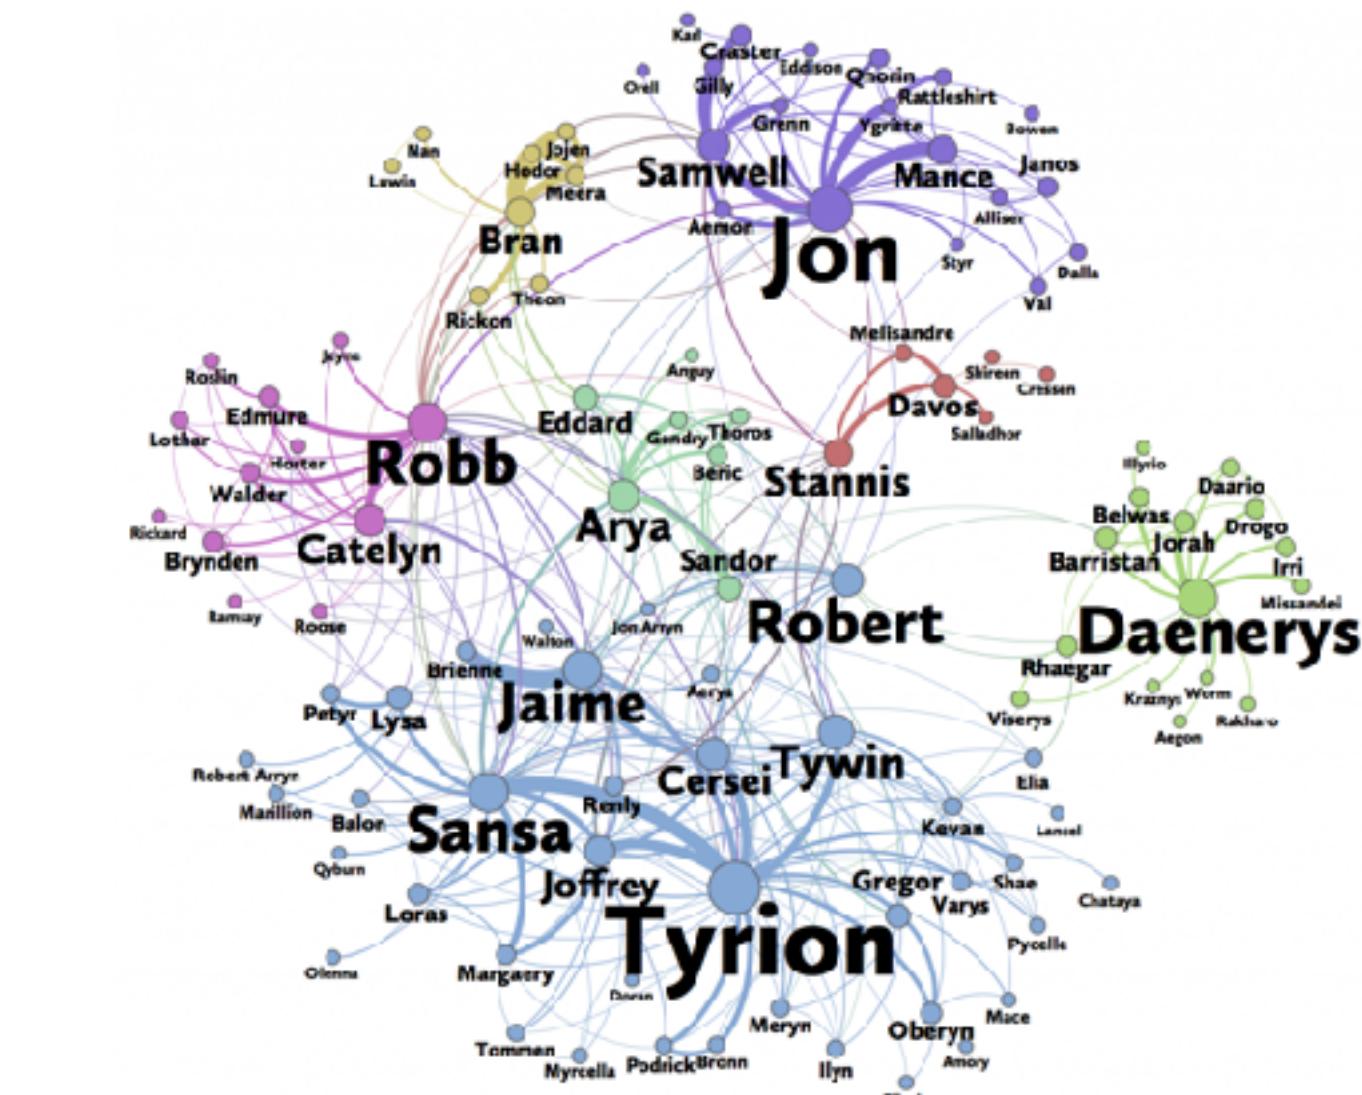
$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

- Aggregate + node updates

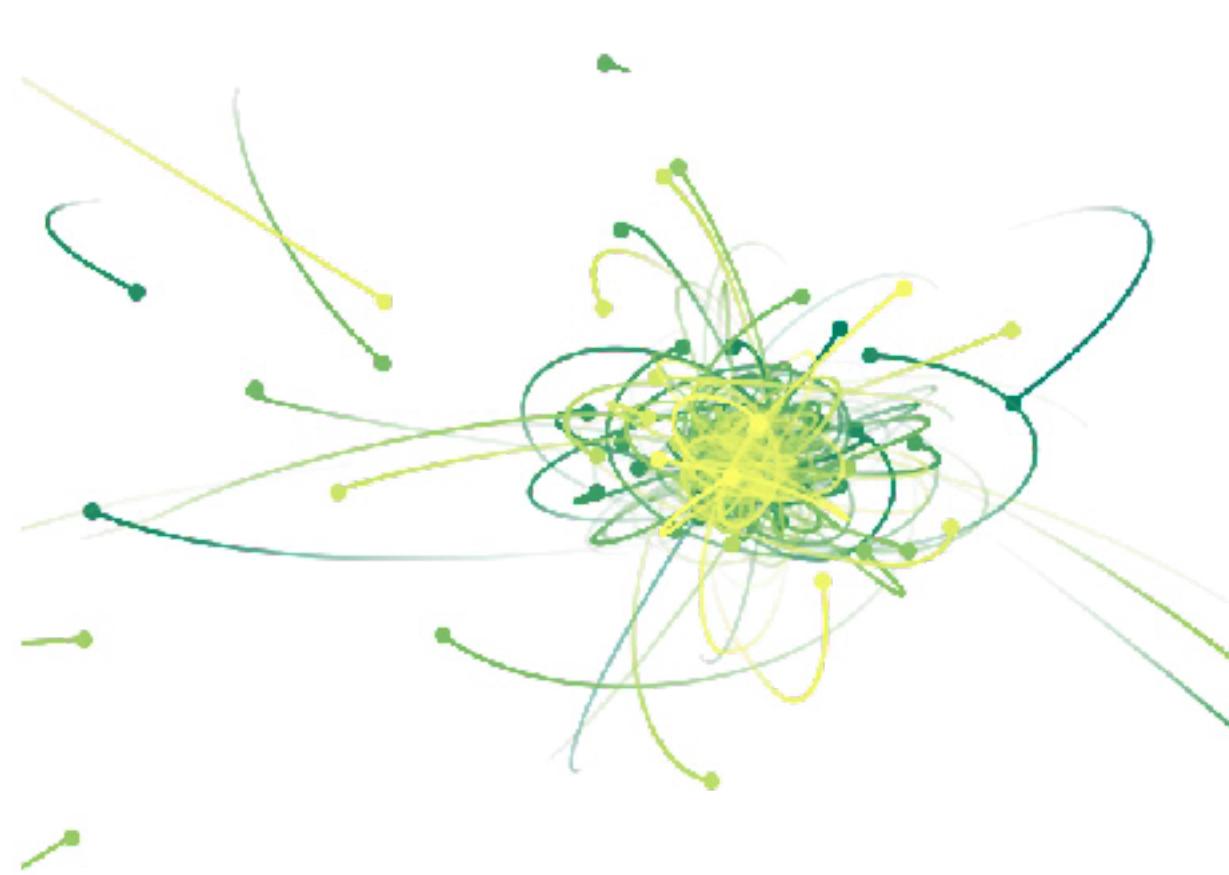
$$\mathbf{f}'_i = \phi_f(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij})$$



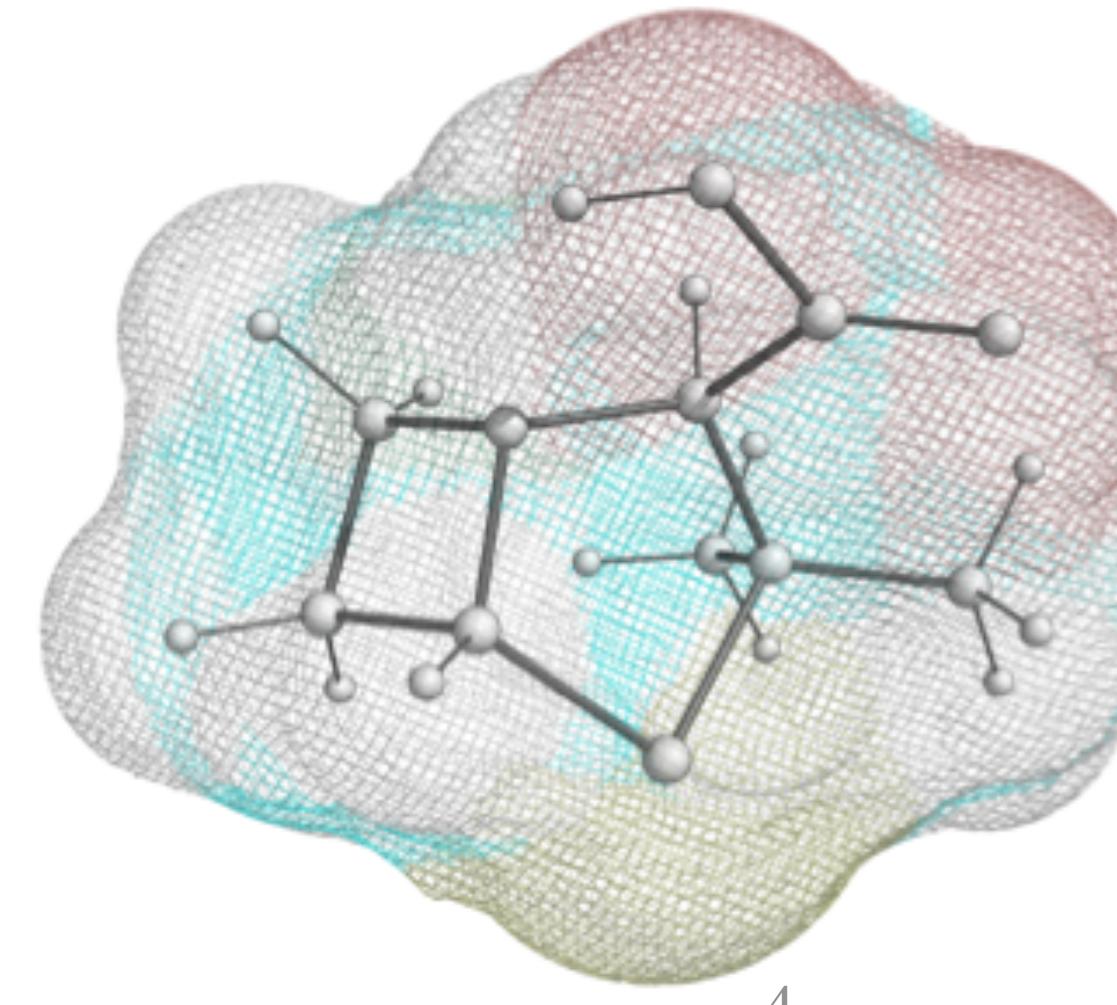
**Gene/Protein interaction graphs<sup>1</sup>**



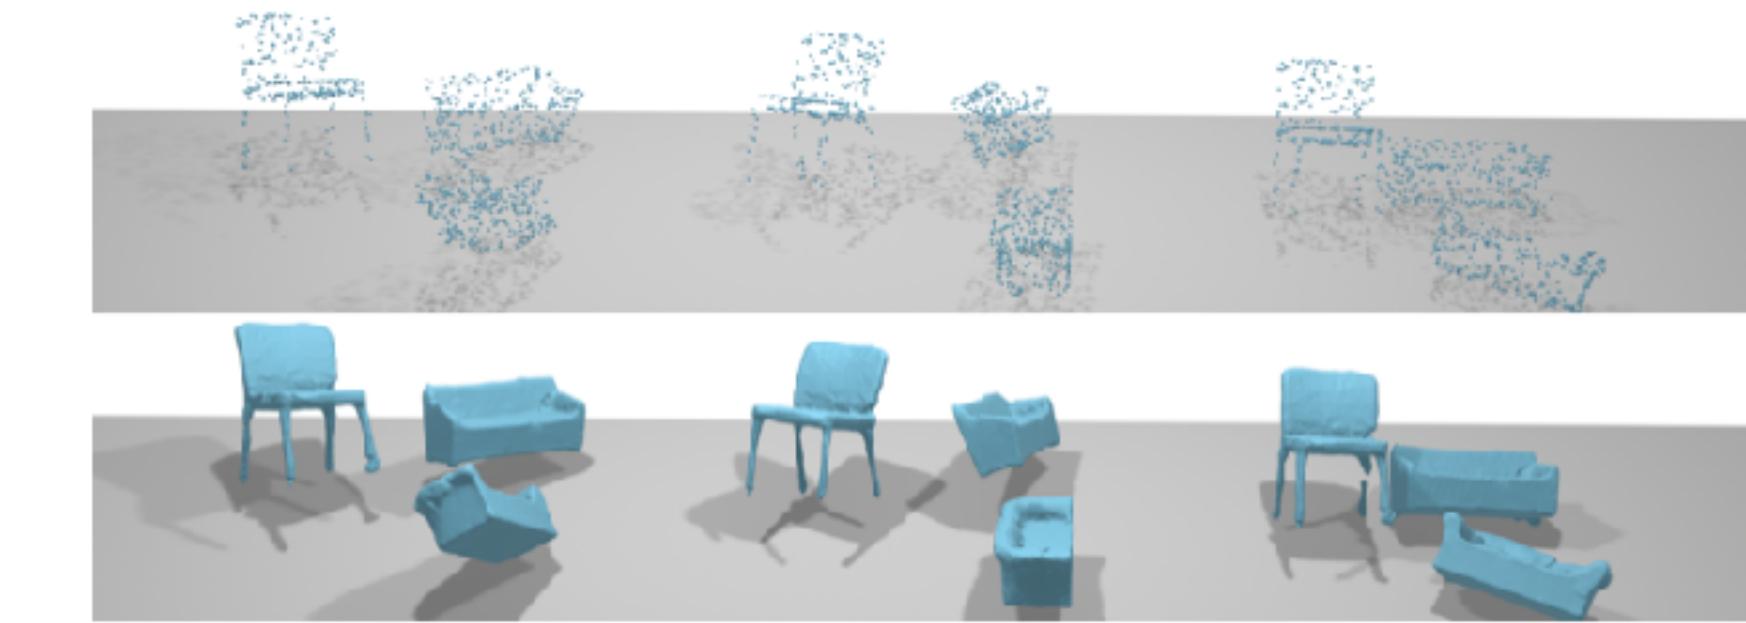
**Social network graph<sup>2</sup>**



**Physical system<sup>3</sup>**



**Molecule<sup>4</sup>**



**Point cloud/shapes<sup>5</sup>**

Figures from: <sup>1</sup>Goh, K. I., Cusick, M. E., Valle, D., Childs, B., Vidal, M., & Barabási, A. L. (2007). The human disease network. *Proceedings of the National Academy of Sciences*, 104(21), 8685-8690.

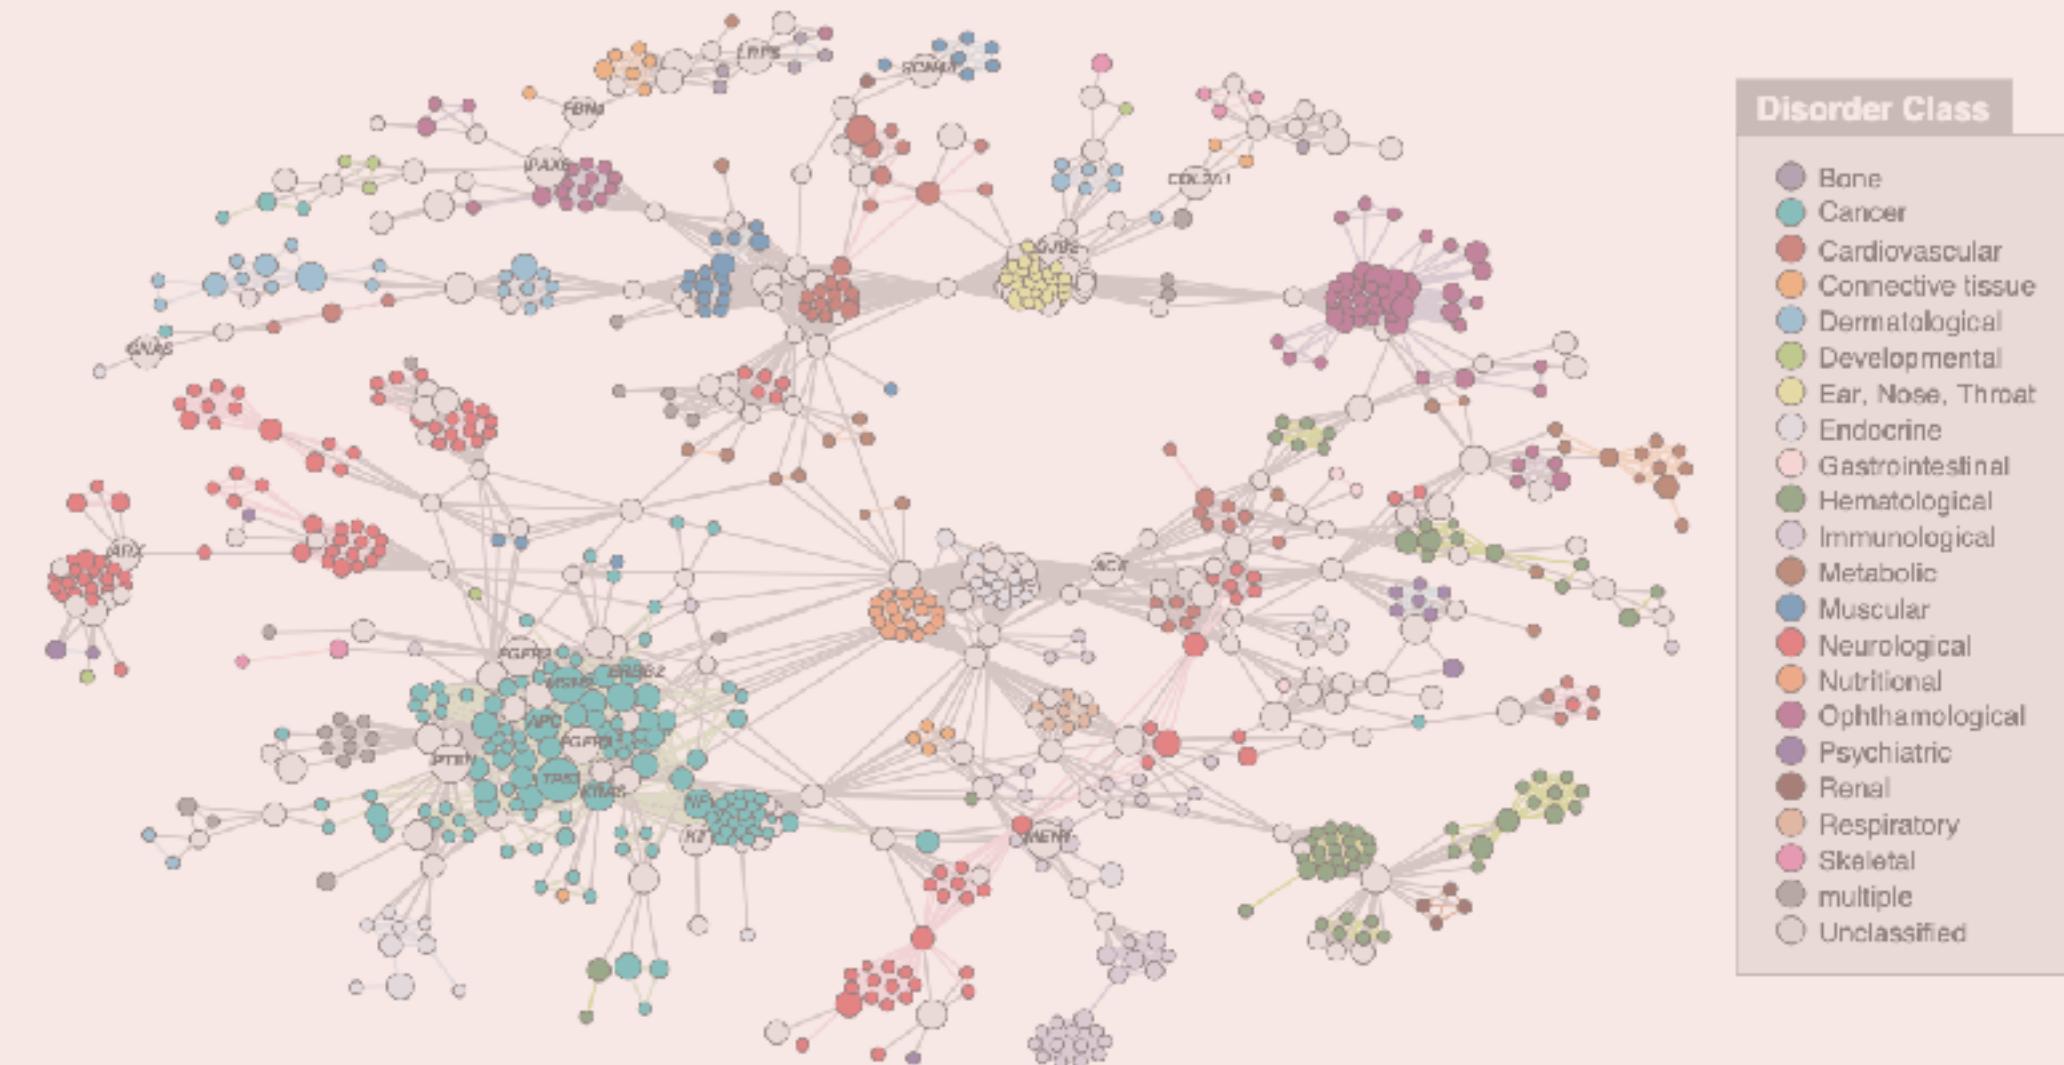
<sup>2</sup><https://predictivehacks.com/social-network-analysis-of-game-of-thrones/>

<sup>3</sup>Brandstetter, J., Hesselink, R., van der Pol, E., Bekkers, E., & Welling, M. (2021). Geometric and Physical Quantities improve E (3) Equivariant Message Passing. In ICLR 2022

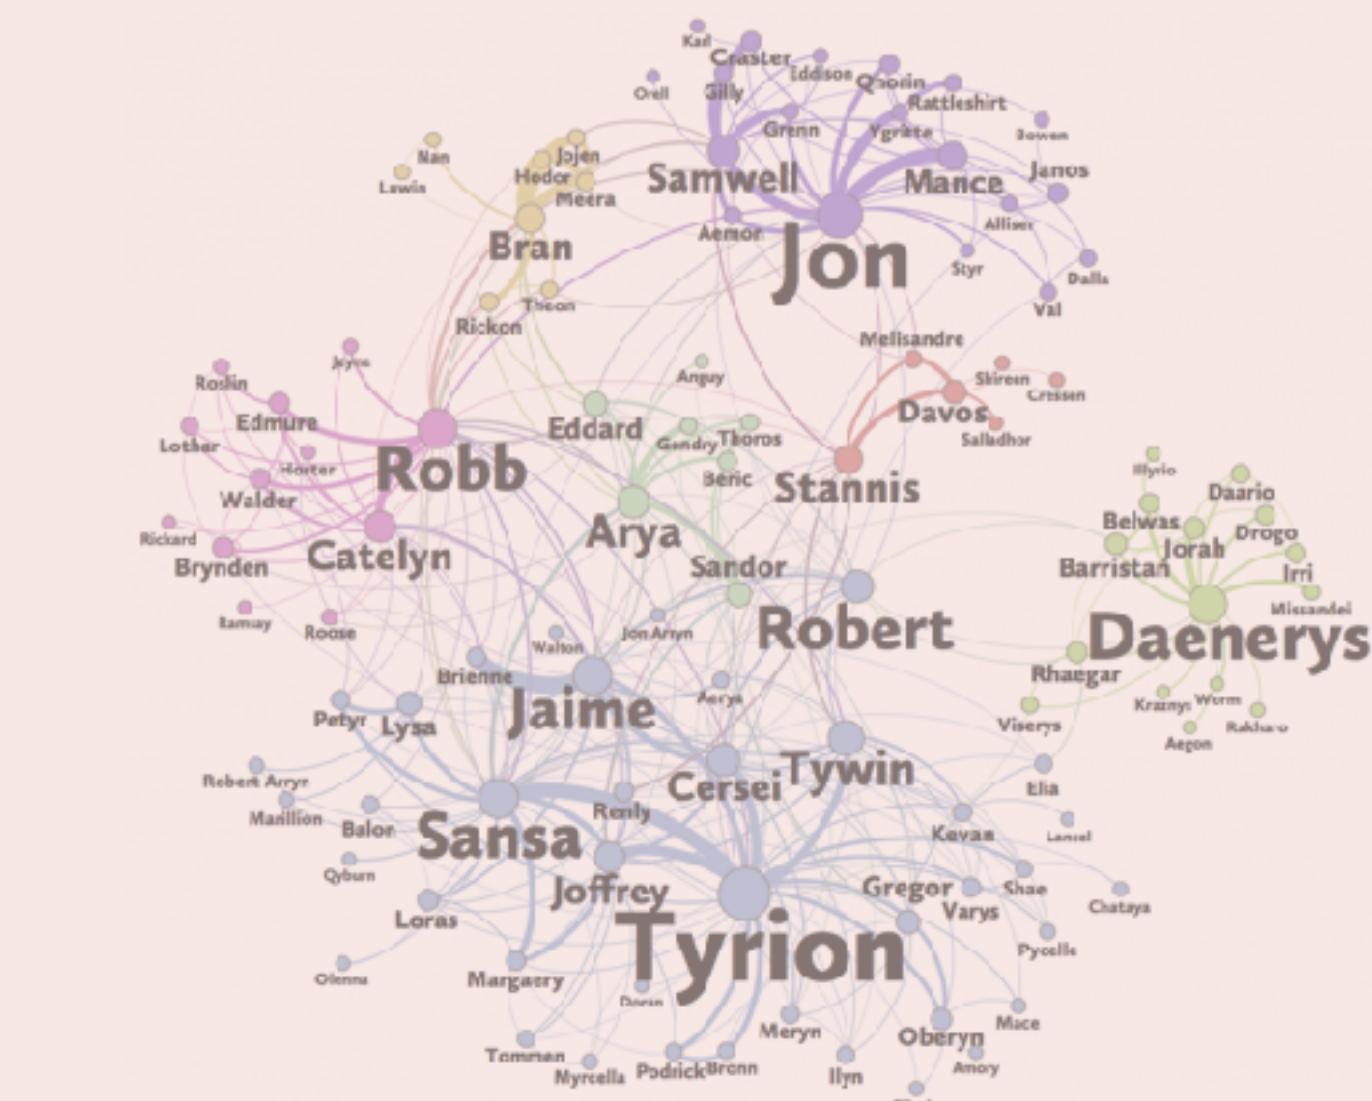
<sup>4</sup>Atz, K., Grisoni, F., & Schneider, G. (2021). Geometric deep learning on molecular representations. *Nature Machine Intelligence*, 1-10.

<sup>5</sup>Chatzipantazis, E., Pertigkiozoglou, S., Dobriban, E., & Daniilidis, K. (2022). SE (3)-Equivariant Attention Networks for Shape Reconstruction in Function Space. arXiv preprint arXiv:2204.02394.

## General graphs

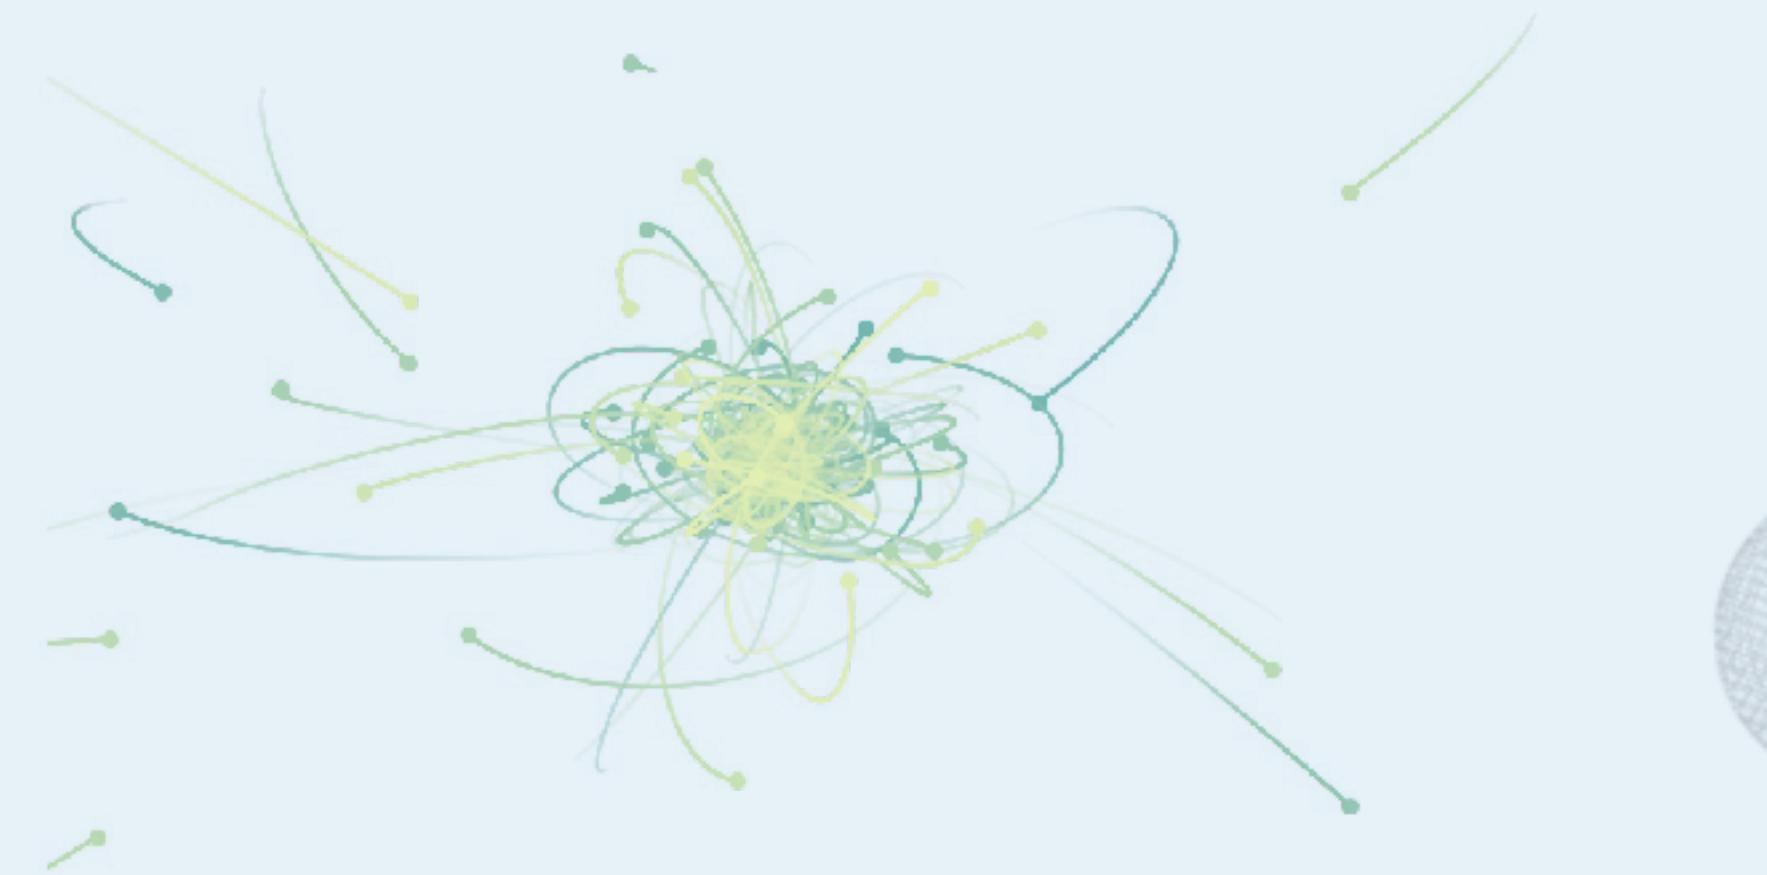


Gene/Protein interaction graphs<sup>1</sup>

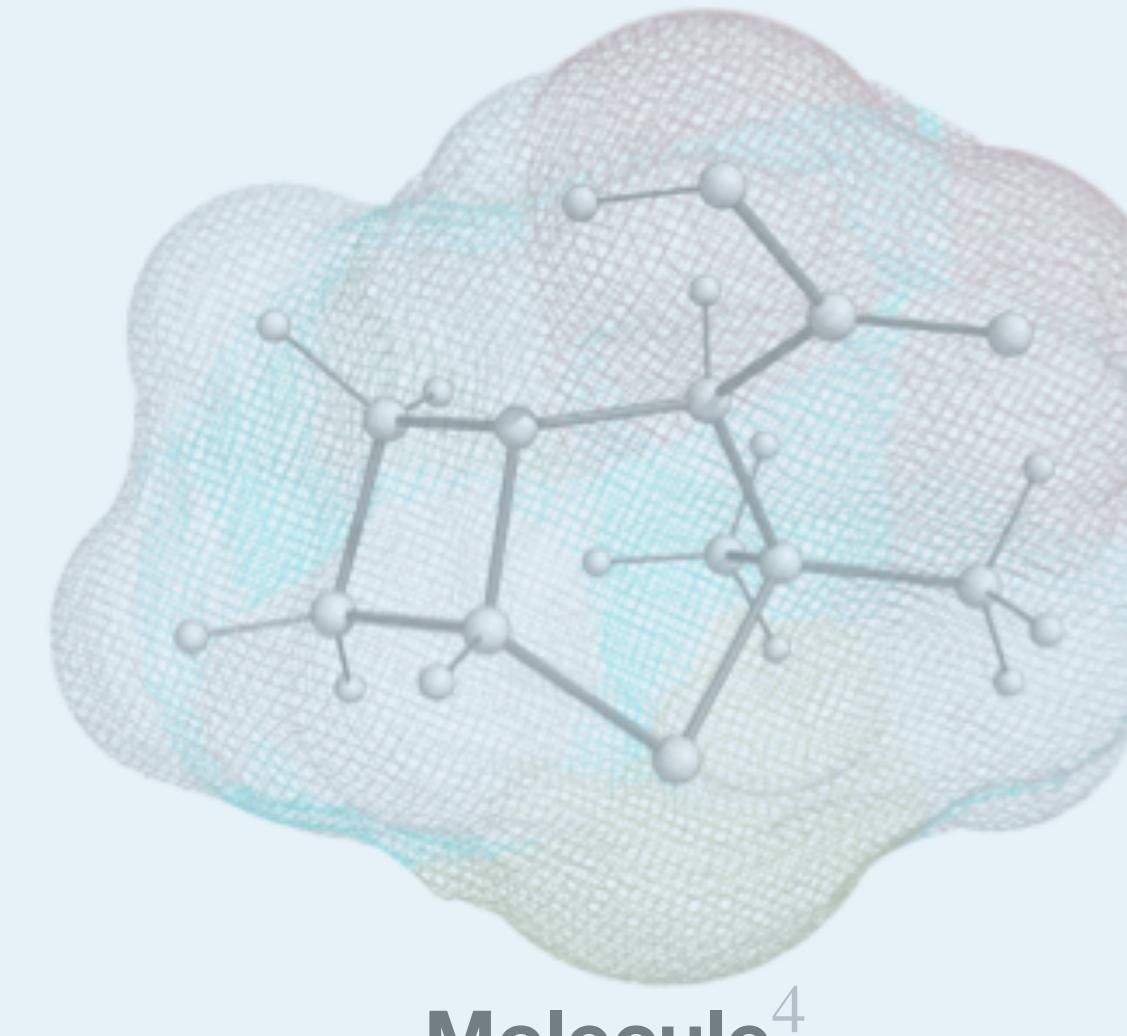


Social network graph<sup>2</sup>

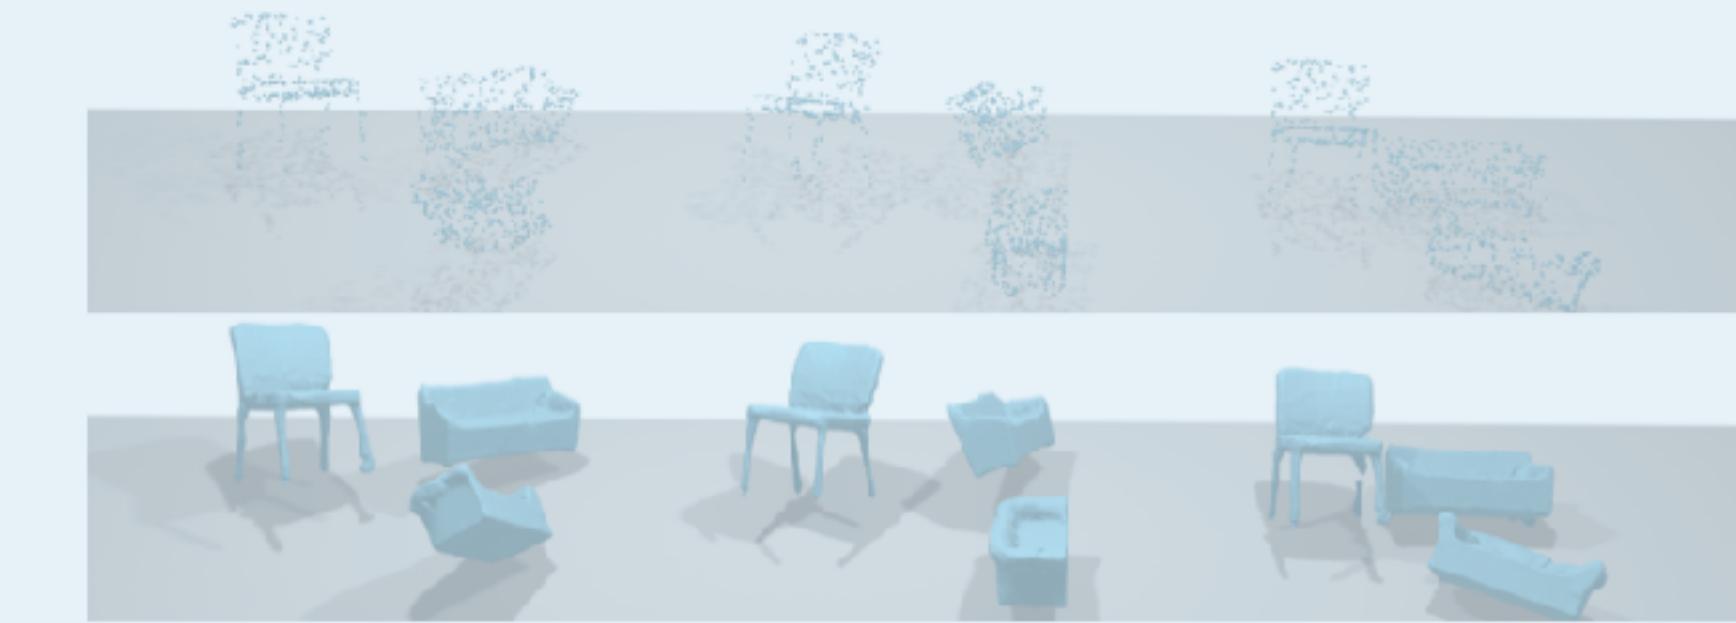
## Geometric graphs (nodes correspond to points in a manifold)



Physical system<sup>3</sup>



Molecule<sup>4</sup>



Point cloud/shapes<sup>5</sup>

Figures from: <sup>1</sup>Goh, K. I., Cusick, M. E., Valle, D., Childs, B., Vidal, M., & Barabási, A. L. (2007). The human disease network. *Proceedings of the National Academy of Sciences*, 104(21), 8685-8690.

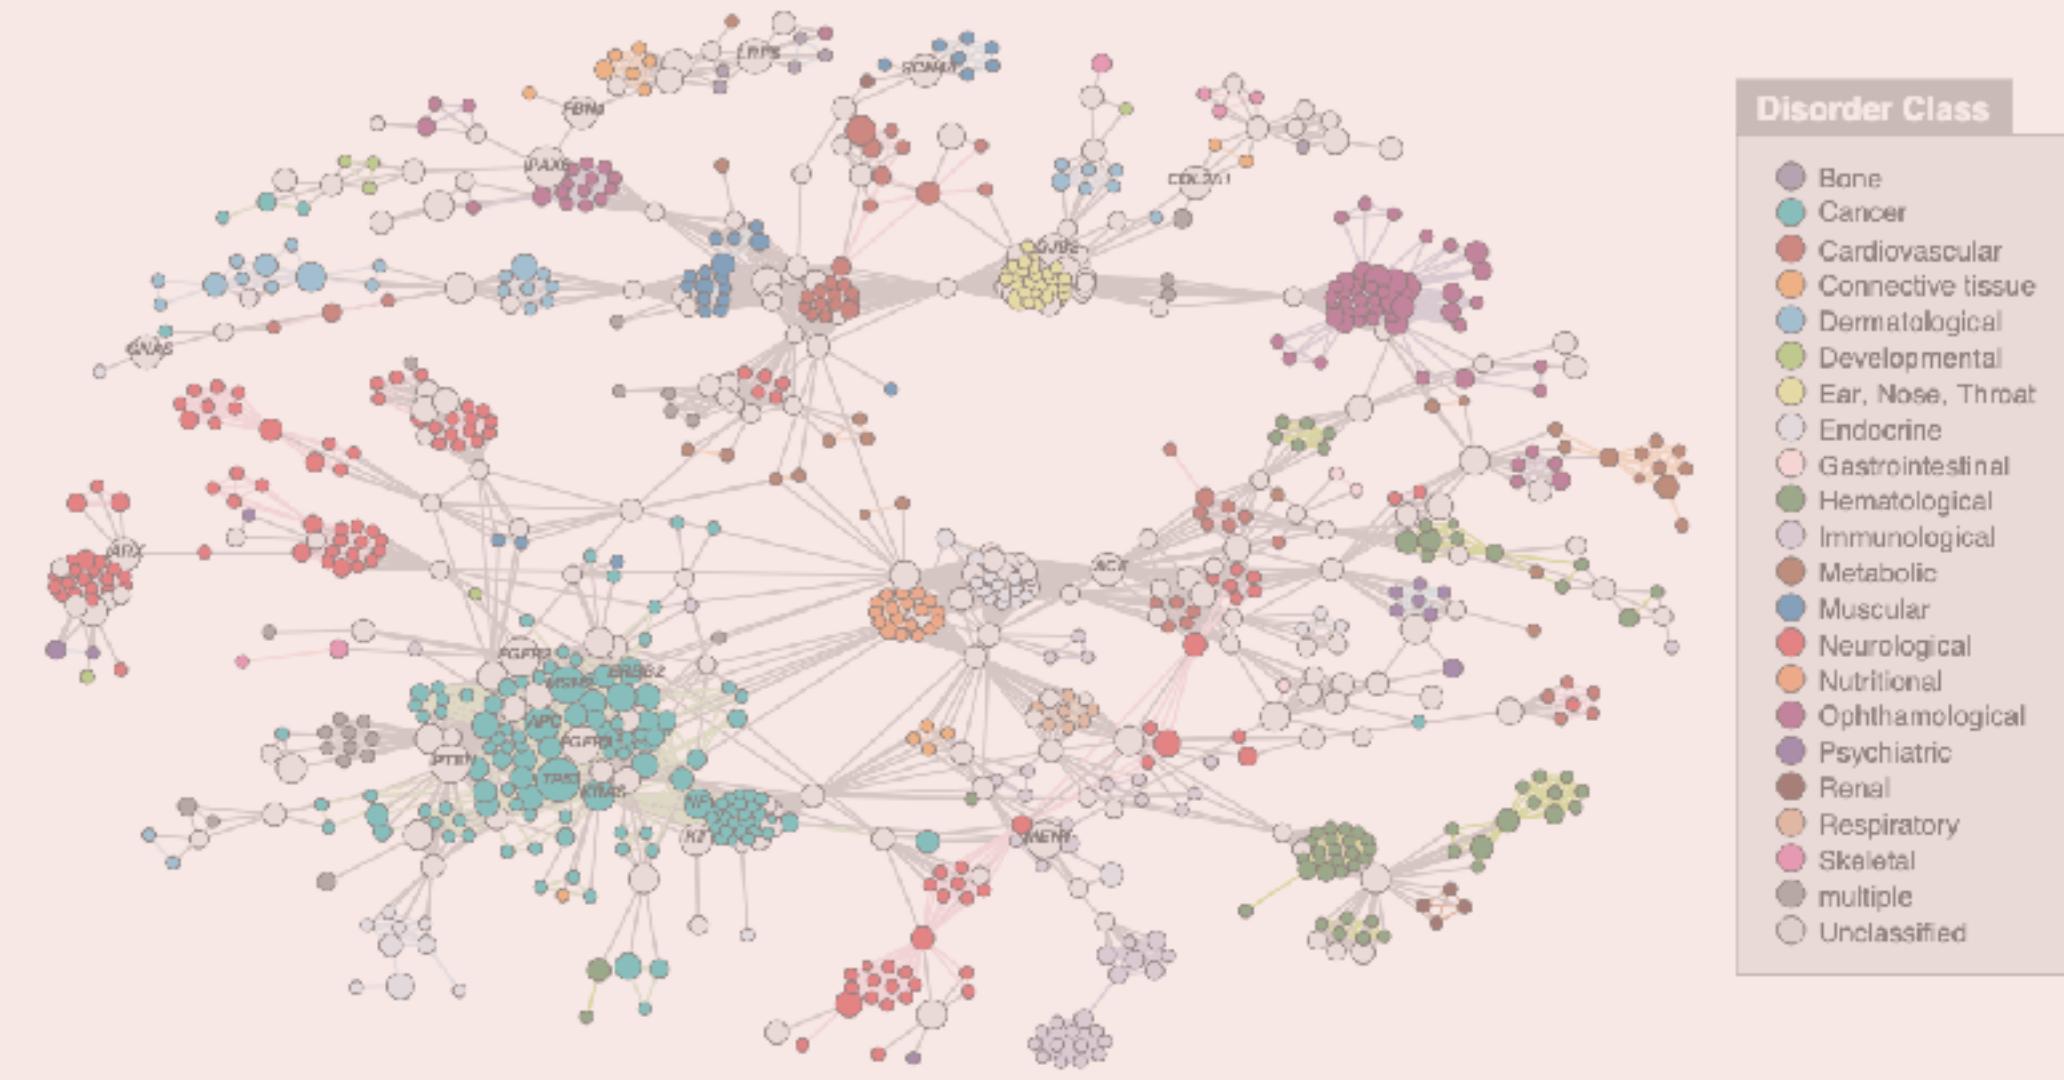
<sup>2</sup><https://predictivehacks.com/social-network-analysis-of-game-of-thrones/>

<sup>3</sup>Brandstetter, J., Hesselink, R., van der Pol, E., Bekkers, E., & Welling, M. (2021). Geometric and Physical Quantities improve E (3) Equivariant Message Passing. In ICLR 2022

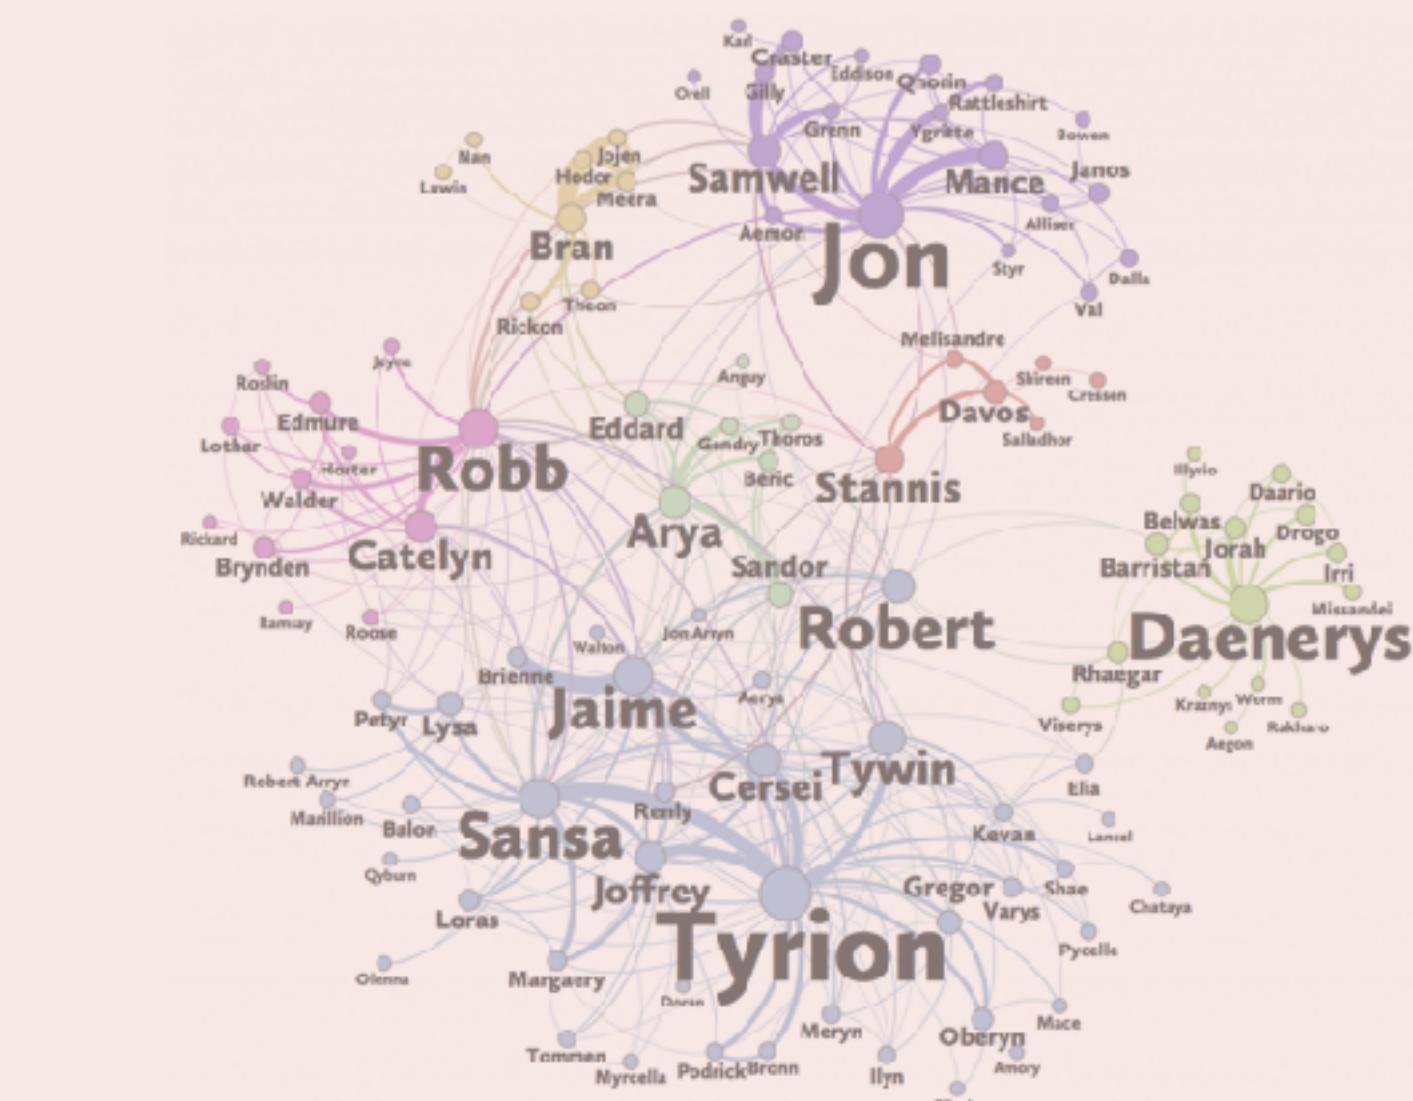
<sup>4</sup>Atz, K., Grisoni, F., & Schneider, G. (2021). Geometric deep learning on molecular representations. *Nature Machine Intelligence*, 1-10.

<sup>5</sup>Chatzipantazis, E., Pertigkiozoglou, S., Dobriban, E., & Daniilidis, K. (2022). SE (3)-Equivariant Attention Networks for Shape Reconstruction in Function Space. arXiv preprint arXiv:2204.02394. 81

## General graphs



Gene/Protein interaction graphs<sup>1</sup>



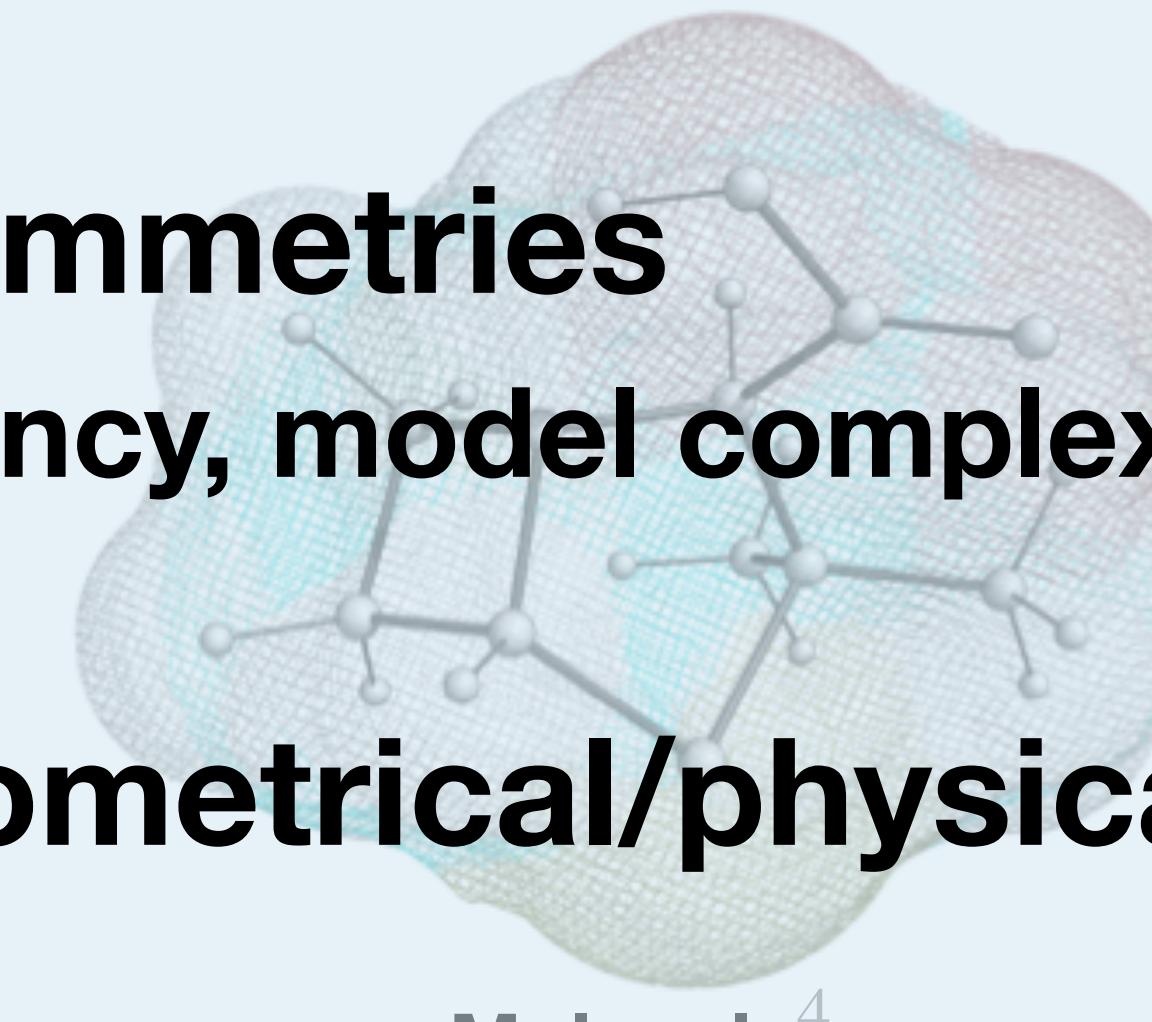
Social network graph<sup>2</sup>

## Geometric graphs (nodes correspond to points in a manifold)

**1. Leverage symmetries  
(sample efficiency, model complexity, generalizability)**



Physical system<sup>3</sup>



Molecule<sup>4</sup>



Point cloud/shapes<sup>5</sup>

Figures from: <sup>1</sup>Goh, K. I., Cusick, M. E., Valle, D., Childs, B., Vidal, M., & Barabási, A. L. (2007). The human disease network. *Proceedings of the National Academy of Sciences*, 104(21), 8685-8690.

<sup>2</sup><https://predictivehacks.com/social-network-analysis-of-game-of-thrones/>

<sup>3</sup>Brandstetter, J., Hesselink, R., van der Pol, E., Bekkers, E., & Welling, M. (2021). Geometric and Physical Quantities improve E (3) Equivariant Message Passing. In ICLR 2022

<sup>4</sup>Atz, K., Grisoni, F., & Schneider, G. (2021). Geometric deep learning on molecular representations. *Nature Machine Intelligence*, 1-10.

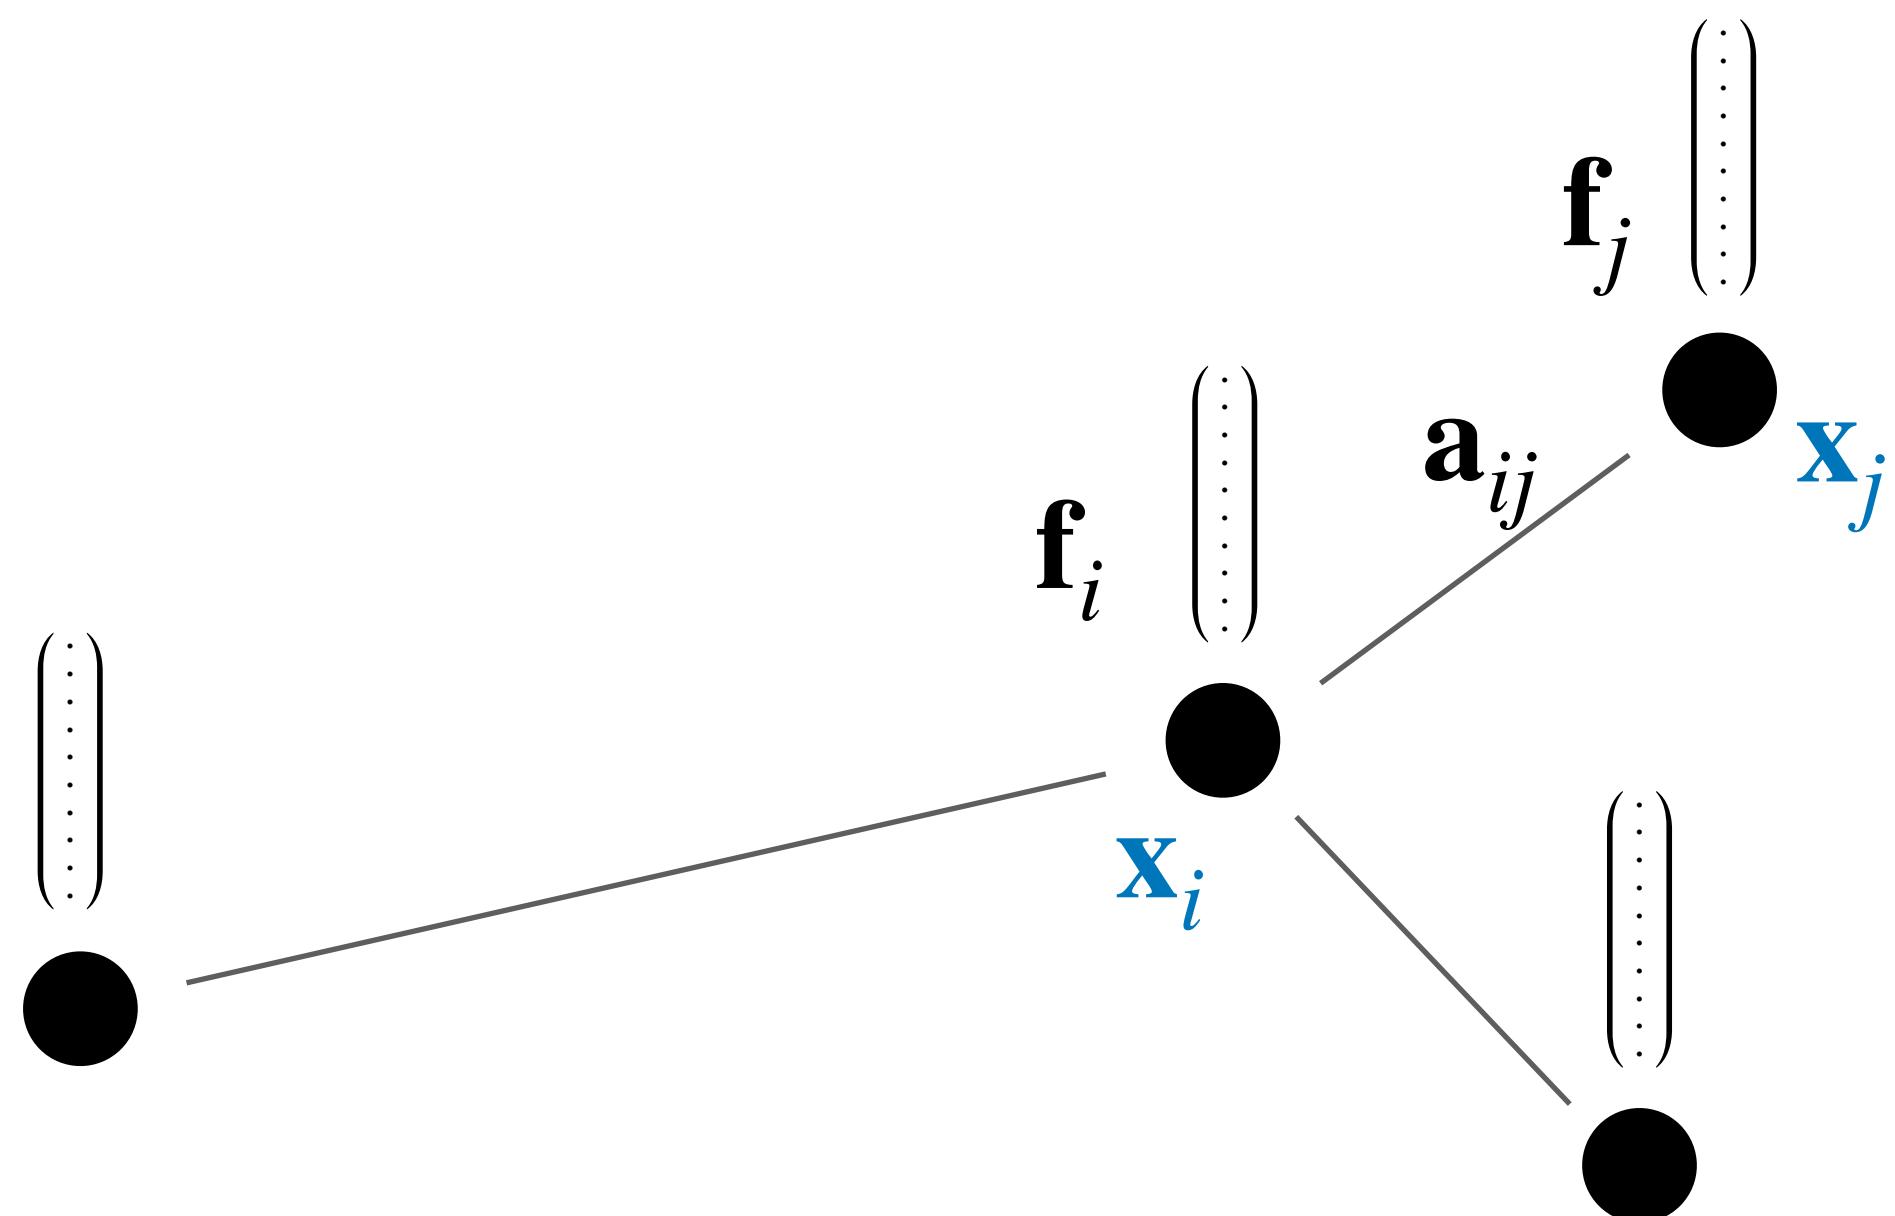
<sup>5</sup>Chatzipantazis, E., Pertigkiozoglou, S., Dobriban, E., & Daniilidis, K. (2022). SE (3)-Equivariant Attention Networks for Shape Reconstruction in Function Space. arXiv preprint arXiv:2204.02394. 81

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$



Message passing layer:

- Messages

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

- Aggregate + node updates

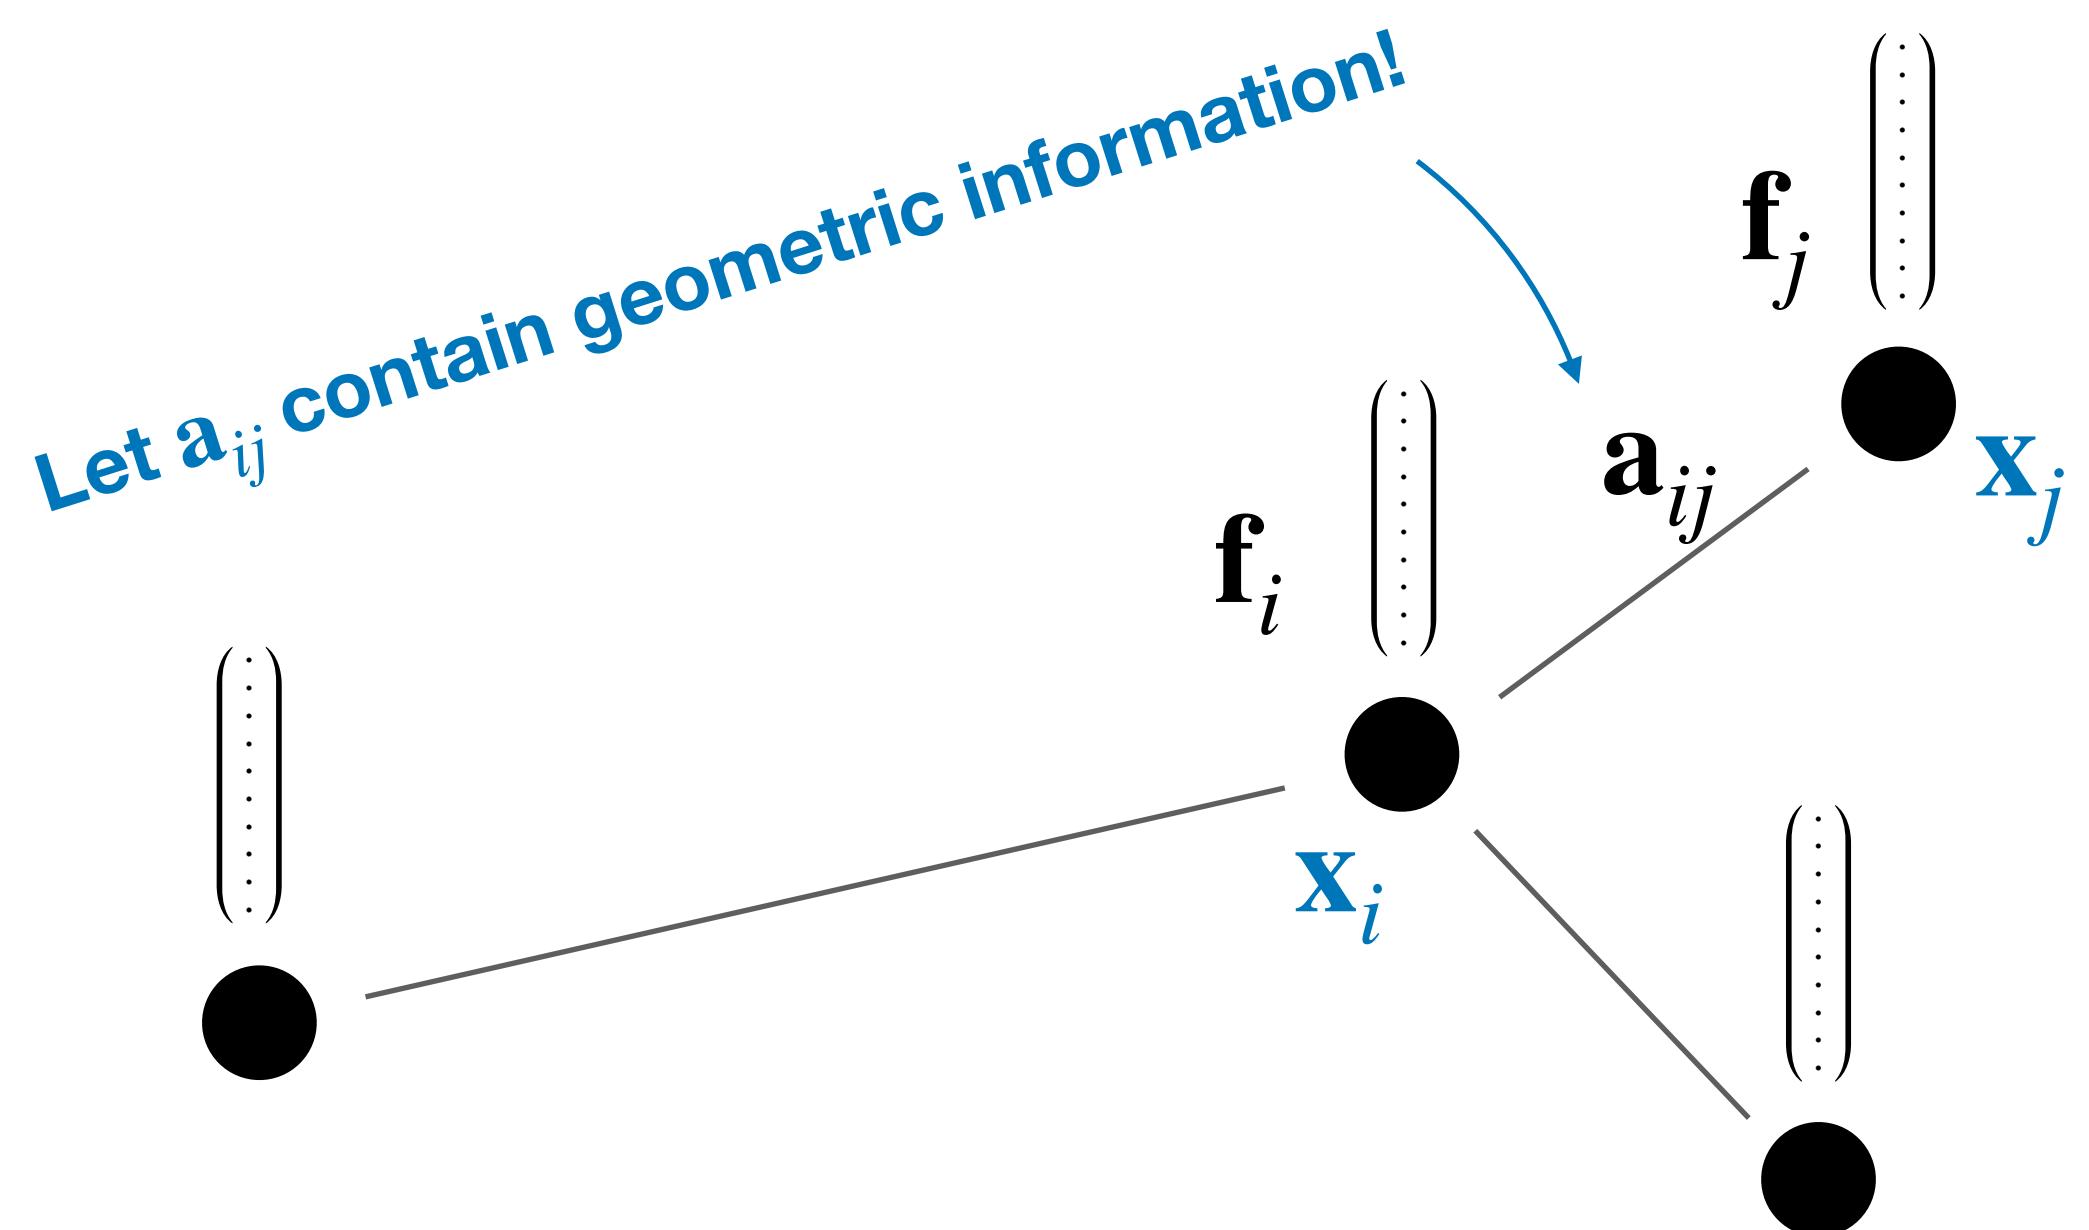
$$\mathbf{f}'_i = \phi_f(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij})$$

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$



Message passing layer:

- Messages

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

- Aggregate + node updates

$$\mathbf{f}'_i = \phi_f(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij})$$

“Condition” messages on geometry

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$

Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

Message passing layer:

- Messages

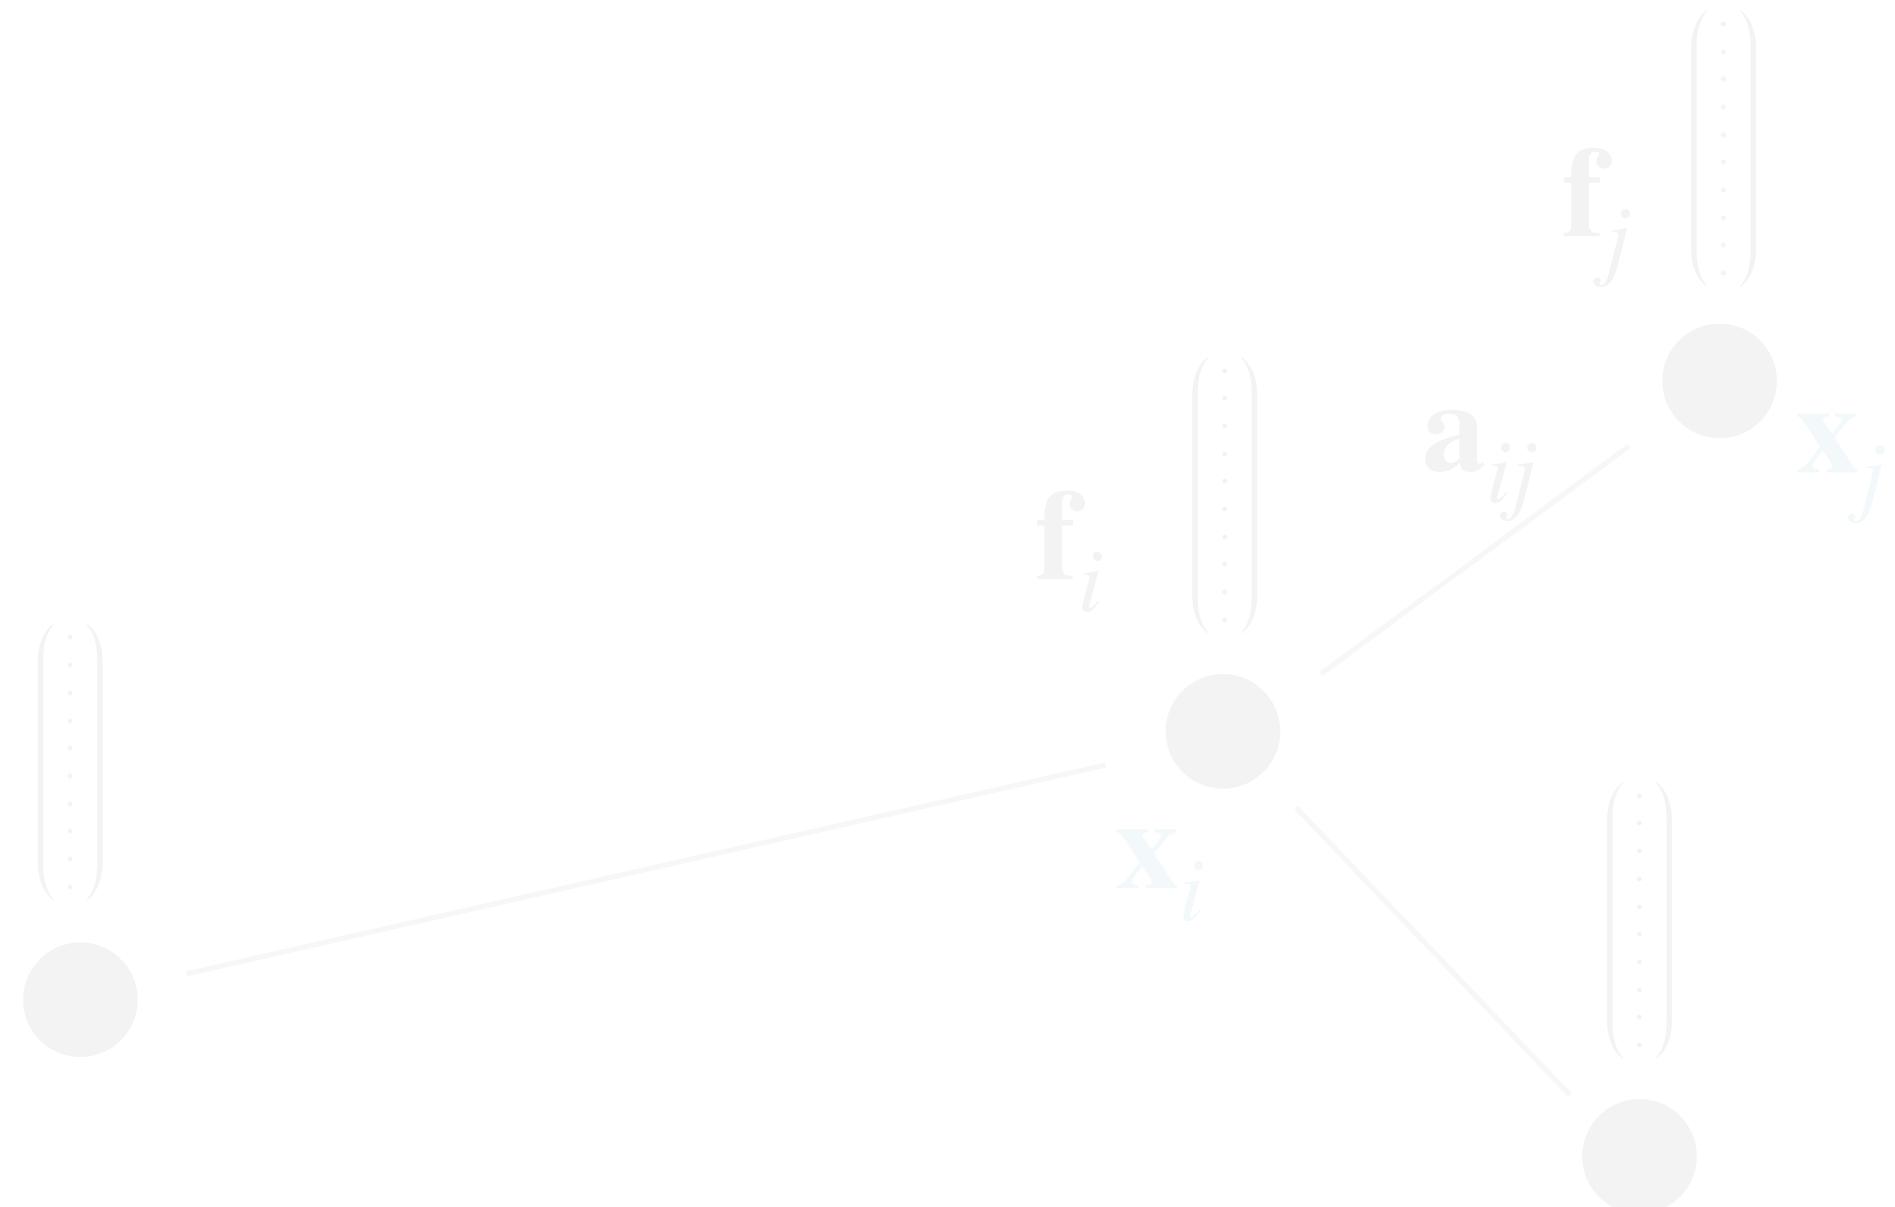
$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{x}_j - \mathbf{x}_i)$$



# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$



Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

Message passing layer:

- Messages

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{x}_j - \mathbf{x}_i)$$

Only equivariant to translations...

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$



Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

Message passing layer:

- Messages

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{x}_j - \mathbf{x}_i)$$

Only equivariant to translations...

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$



Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

Message passing layer:

- Messages

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{x}_j - \mathbf{x}_i)$$

Only equivariant to translations...

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

Full  $E(3)$  equivariance, but a bit restrictive...

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$



Goal: iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

Message passing layer:

- Messages

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{x}_j - \mathbf{x}_i)$$

Only equivariant to translations...

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

Full  $E(3)$  equivariance, but a bit restrictive...

$$(X = G) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, g_j^{-1}g_i)$$

Solution 1: Lift to the group!

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$



**Goal:** iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

Message passing layer:

- Messages

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{x}_j - \mathbf{x}_i)$$

Only equivariant to translations...

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

Full  $E(3)$  equivariance, but a bit restrictive...

$$(X = G) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, g_j^{-1}g_i)$$

Solution 1: Lift to the group!

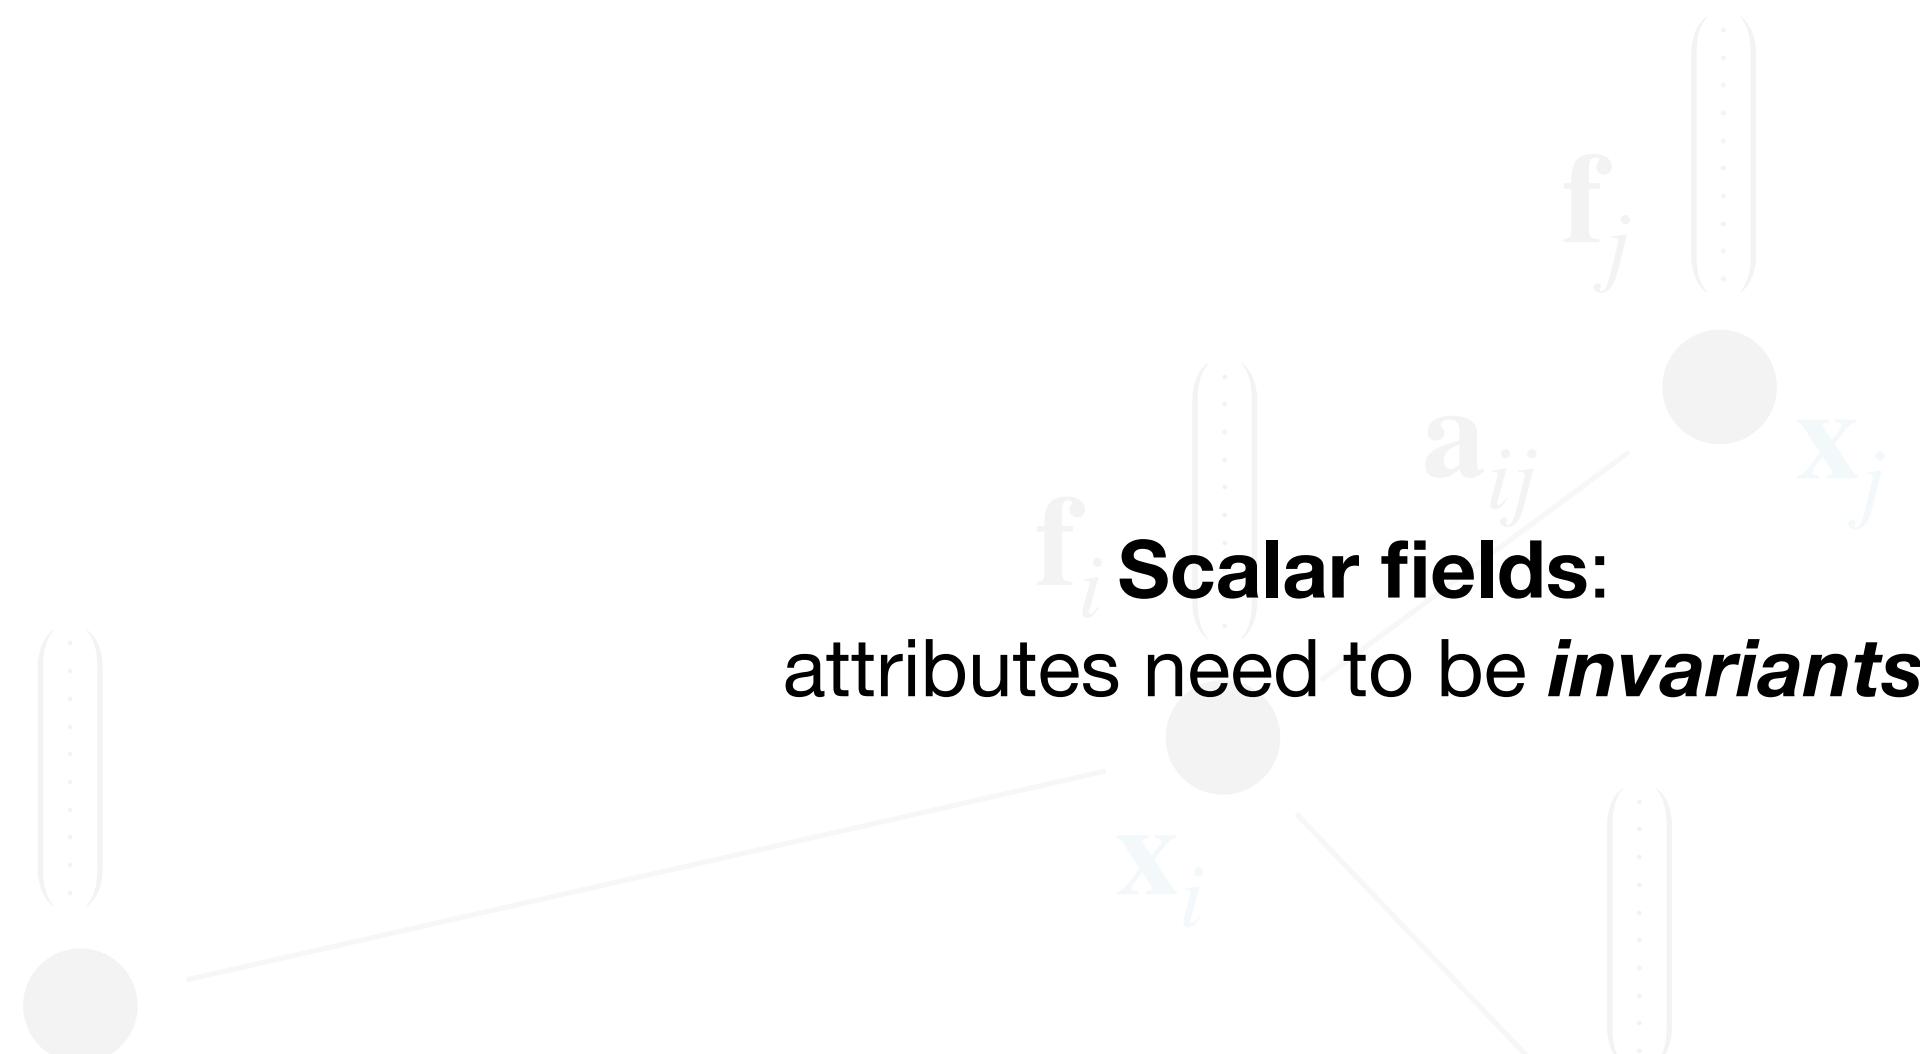
$$(X = \mathbb{R}^d) \quad \hat{\mathbf{m}}_{ij} = \hat{\phi}_m(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, Y(\mathbf{x}_j - \mathbf{x}_i))$$

Solution 2: work with steerable feature fields!

# The Geometric Message Passing Framework

Graph  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

- nodes  $v_i \in \mathcal{V}$  with node feature  $\mathbf{f}_i \in \mathbb{R}^{C_v}$  and position  $x_i \in X$
- edges  $e_{ij} \in \mathcal{E}$  with edge attribute  $\mathbf{a}_{ij} \in \mathbb{R}^{C_e}$



**Scalar fields:**  
attributes need to be *invariants*!

**Steerable feature fields:**  
attributes can be *covariants*!

**Goal:** iteratively update node features to obtain useful hidden representations  $\mathbf{h} \in \mathbb{R}^{C_h}$

Message passing layer:

- Messages

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{x}_j - \mathbf{x}_i)$$

Only equivariant to translations...

$$(X = \mathbb{R}^d) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

Full  $E(3)$  equivariance, but a bit restrictive...

$$(X = G) \quad \mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, g_j^{-1}g_i)$$

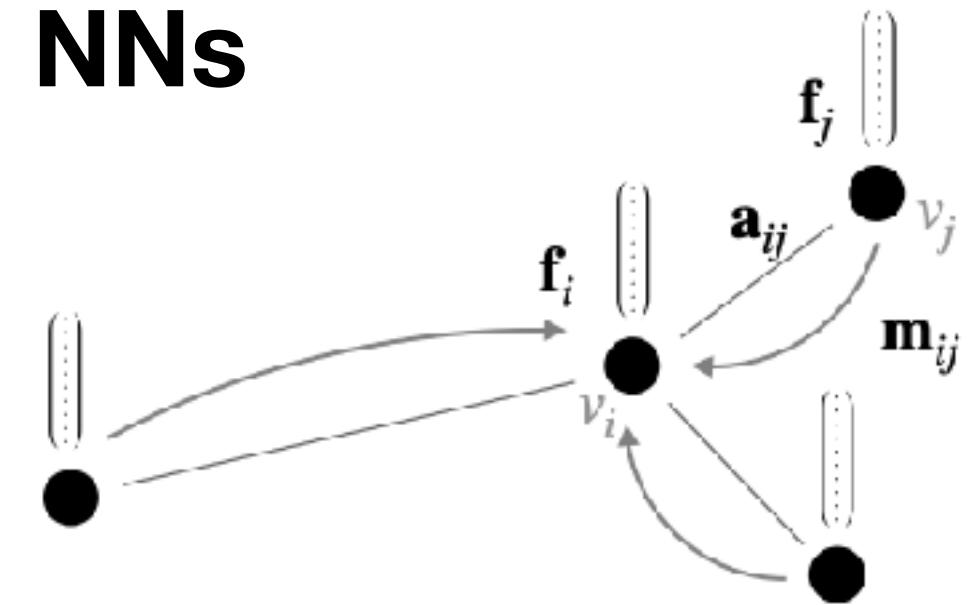
Solution 1: Lift to the group!

$$(X = \mathbb{R}^d) \quad \hat{\mathbf{m}}_{ij} = \hat{\phi}_m(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, Y(\mathbf{x}_j - \mathbf{x}_i))$$

Solution 2: work with steerable feature fields!

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

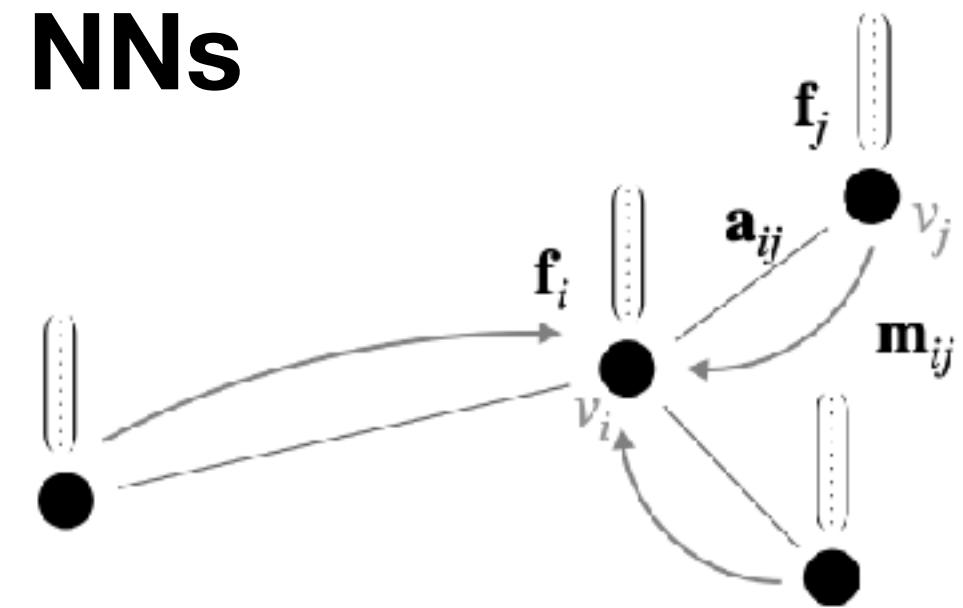
$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

## Classic point convolutions

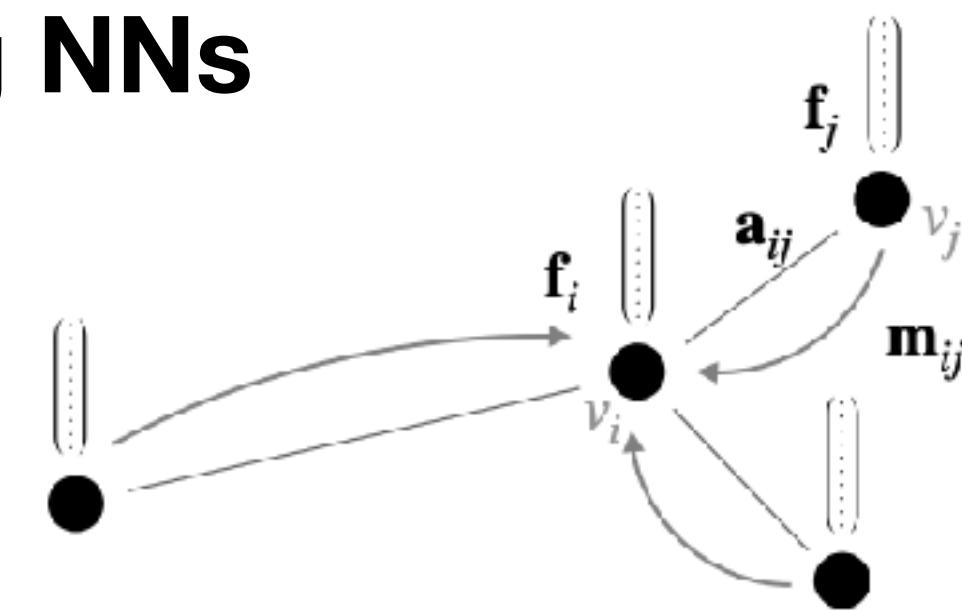
(Lecture 1.7: regular g-convs on homogeneous spaces)

$$\mathbf{m}_{ij} = \mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)\mathbf{f}_j$$

$$\mathbf{m}_{ij} = \mathbf{W}(g_i^{-1}g_j)\mathbf{f}_j$$

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

## Classic point convolutions

(Lecture 1.7: regular g-convs on homogeneous spaces)

$$\mathbf{m}_{ij} = \mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)\mathbf{f}_j$$

$$\mathbf{m}_{ij} = \mathbf{W}(g_i^{-1}g_j)\mathbf{f}_j$$

## Steerable G-CNNs

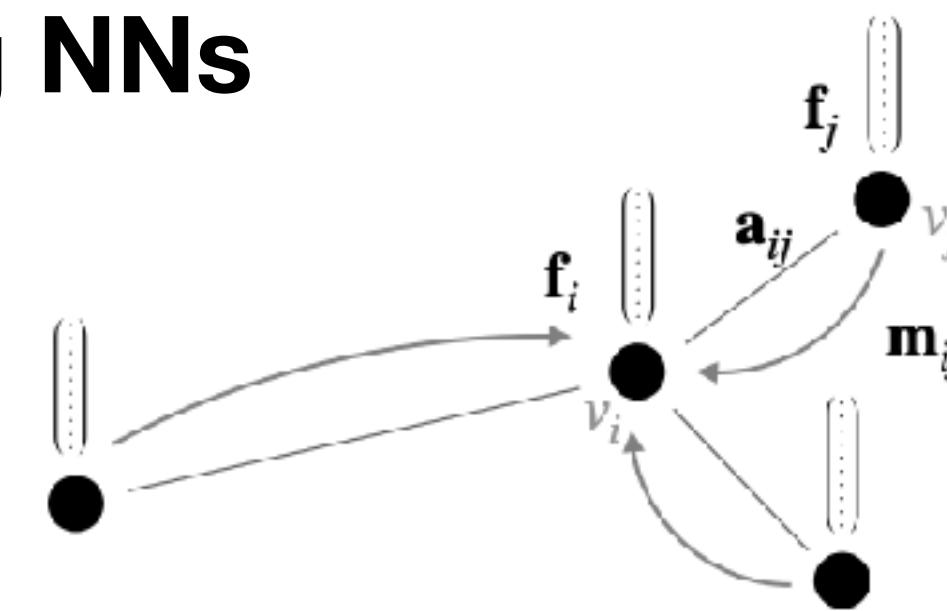
(Lecture 2: steerable g-convs)

$$\mathbf{m}_{ij} = \mathbf{W}_{\hat{\mathbf{a}}_{ij}}(\|\mathbf{x}_j - \mathbf{x}_i\|)\hat{\mathbf{f}}_j$$

$$:= \hat{\mathbf{f}}_j \otimes_{cg}^{\mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)} \hat{\mathbf{a}}_{ij}$$

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

## Classic point convolutions

(Lecture 1.7: regular g-convs on homogeneous spaces)

$$\mathbf{m}_{ij} = \mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)\mathbf{f}_j$$

$$\mathbf{m}_{ij} = \mathbf{W}(g_i^{-1}g_j)\mathbf{f}_j$$

## Steerable G-CNNs

(Lecture 2: steerable g-convs)

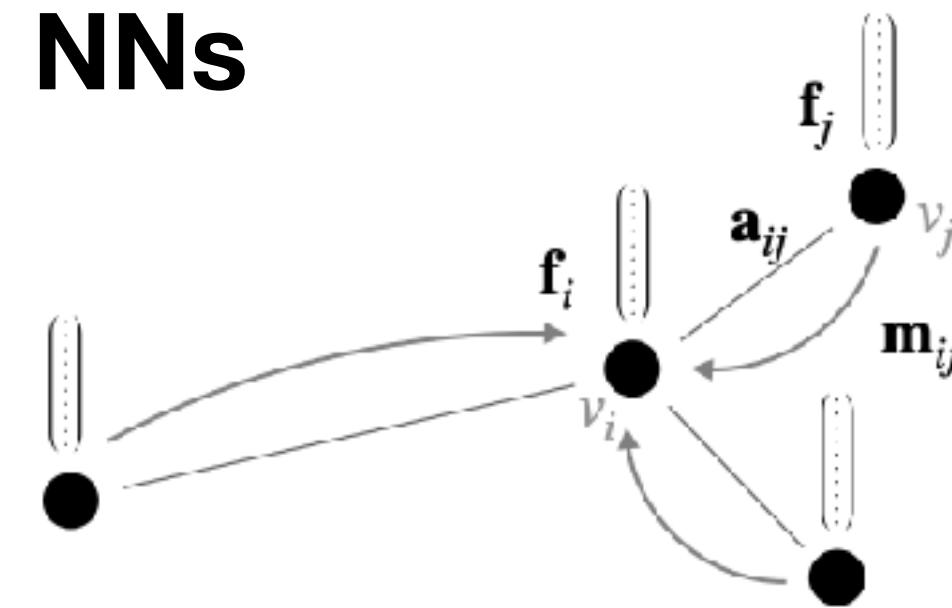
$$\mathbf{m}_{ij} = \mathbf{W}_{\hat{\mathbf{a}}_{ij}}(\|\mathbf{x}_j - \mathbf{x}_i\|)\hat{\mathbf{f}}_j$$

$$:= \hat{\mathbf{f}}_j \otimes_{cg}^{\mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)} \hat{\mathbf{a}}_{ij}$$

$\mathcal{F}_H$

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

## Classic point convolutions

(Lecture 1.7: regular g-convs on homogeneous spaces)

$$\mathbf{m}_{ij} = \mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)\mathbf{f}_j$$

$$\mathbf{m}_{ij} = \mathbf{W}(g_i^{-1}g_j)\mathbf{f}_j$$

## Steerable G-CNNs

(Lecture 2: steerable g-convs)

$$\mathbf{m}_{ij} = \mathbf{W}_{\hat{\mathbf{a}}_{ij}}(\|\mathbf{x}_j - \mathbf{x}_i\|)\hat{\mathbf{f}}_j$$

$$:= \hat{\mathbf{f}}_j \otimes_{cg}^{\mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)} \hat{\mathbf{a}}_{ij}$$

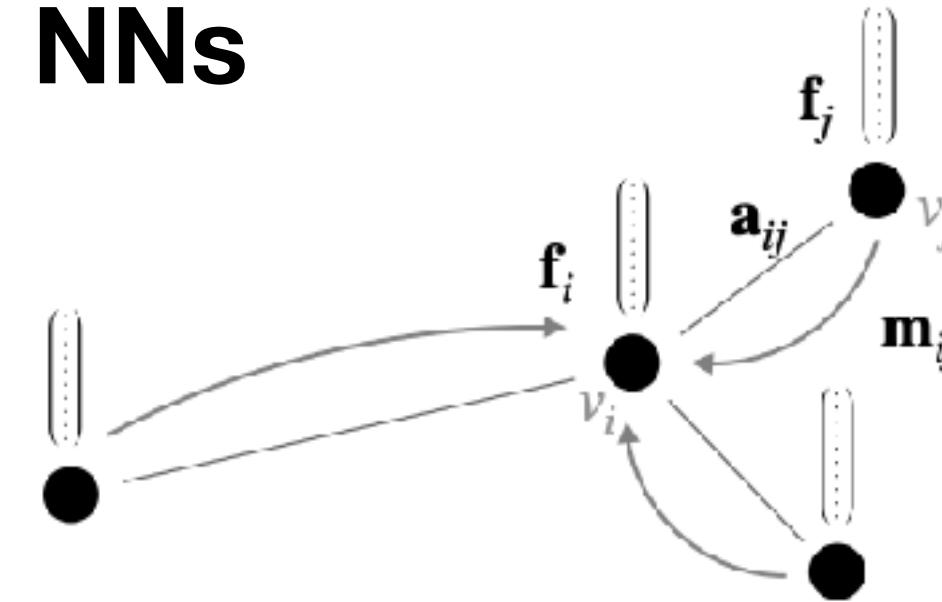
## Invariant Message Passing NNs

(Lecture 3)

$$\mathbf{m}_{ij} = \text{MLP}(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

## Classic point convolutions

(Lecture 1.7: regular g-convs on homogeneous spaces)

$$\mathbf{m}_{ij} = \mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)\mathbf{f}_j$$

$$\mathbf{m}_{ij} = \mathbf{W}(g_i^{-1}g_j)\mathbf{f}_j$$

## Steerable G-CNNs

(Lecture 2: steerable g-convs)

$$\mathbf{m}_{ij} = \mathbf{W}_{\hat{\mathbf{a}}_{ij}}(\|\mathbf{x}_j - \mathbf{x}_i\|)\hat{\mathbf{f}}_j$$

$$:= \hat{\mathbf{f}}_j \otimes_{cg}^{\mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)} \hat{\mathbf{a}}_{ij}$$

$\mathcal{F}_H$

## Invariant Message Passing NNs

(Lecture 3)

$$\mathbf{m}_{ij} = \text{MLP}(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

## Equivariant (Steerable) Message Passing NNs

(Lecture 3)

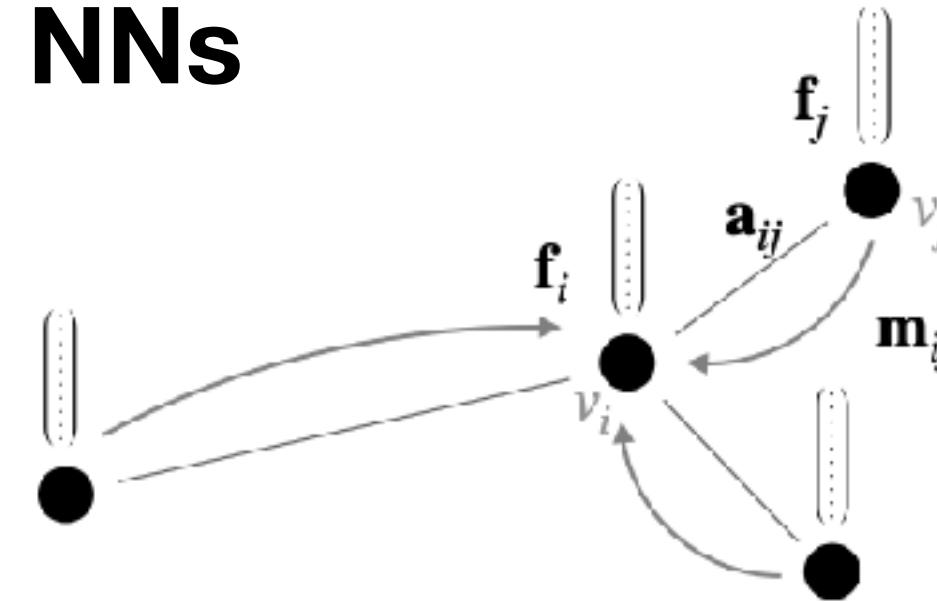
$$\hat{\mathbf{m}}_{ij} = \widehat{\text{MLP}}(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, \mathbf{x}_j - \mathbf{x}_i)$$

With steerable MLP:

$$\widehat{\text{MLP}}_{\hat{\mathbf{a}}_{ij}}(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, \|\mathbf{x}_j - \mathbf{x}_i\|) := \sigma(\mathbf{W}_{\hat{\mathbf{a}}_{ij}}^{(n)}(\dots(\sigma(\mathbf{W}_{\hat{\mathbf{a}}_{ij}}^{(1)}\hat{\mathbf{h}}_i))))$$

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

## Classic point convolutions

(Lecture 1.7: regular g-convs on homogeneous spaces)

$$\mathbf{m}_{ij} = \mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)\mathbf{f}_j$$

$$\mathbf{W}(g_i^{-1}g_j)\mathbf{f}_j$$

**Linear convolution**

## Steerable G-CNNs

(Lecture 2: steerable g-convs)

$$\mathbf{m}_{ij} = \mathbf{W}_{\hat{\mathbf{a}}_{ij}}(\|\mathbf{x}_j - \mathbf{x}_i\|)\hat{\mathbf{f}}_j$$

$$:= \hat{\mathbf{f}}_j \otimes_{cg}^{\mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)} \hat{\mathbf{a}}_{ij}$$

$$\mathcal{F}_H$$

## Invariant Message Passing NNs

(Lecture 3)

$$\mathbf{m}_{ij} = \text{MLP}(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

## Equivariant (Steerable) Message Passing NNs

(Lecture 3)

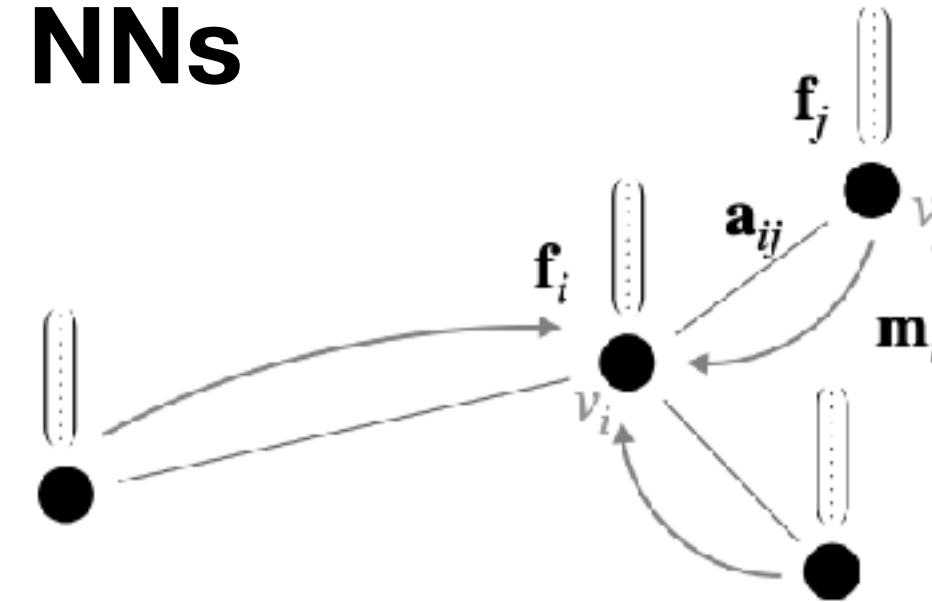
$$\hat{\mathbf{m}}_{ij} = \widehat{\text{MLP}}(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, \mathbf{x}_j - \mathbf{x}_i)$$

**With steerable MLP:**

$$\widehat{\text{MLP}}_{\hat{\mathbf{a}}_{ij}}(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, \|\mathbf{x}_j - \mathbf{x}_i\|) := \sigma(\mathbf{W}_{\hat{\mathbf{a}}_{ij}}^{(n)}(\dots(\sigma(\mathbf{W}_{\hat{\mathbf{a}}_{ij}}^{(1)}\hat{\mathbf{h}}_i))))$$

# Linear vs non-linear (group) convolutions

## Message passing NNs



Compute messages:

$$\mathbf{m}_{ij} = \phi_m(\mathbf{f}_i, \mathbf{f}_j, \mathbf{a}_{ij})$$

Aggregate and update:

$$\mathbf{f}'_i = \phi_f\left(\mathbf{f}_i, \sum_{j \in \mathcal{N}(i)} \mathbf{m}_{ij}\right)$$

## Classic point convolutions

(Lecture 1.7: regular g-convs on homogeneous spaces)

$$\mathbf{m}_{ij} = \mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)\mathbf{f}_j$$

$$\mathbf{W}(g_i^{-1}g_j)\mathbf{f}_j$$

**Linear convolution**

## Steerable G-CNNs

(Lecture 2: steerable g-convs)

$$\mathbf{m}_{ij} = \mathbf{W}_{\hat{\mathbf{a}}_{ij}}(\|\mathbf{x}_j - \mathbf{x}_i\|)\hat{\mathbf{f}}_j$$

$$:= \hat{\mathbf{f}}_j \otimes_{cg}^{\mathbf{W}(\|\mathbf{x}_j - \mathbf{x}_i\|)} \hat{\mathbf{a}}_{ij}$$

$$\mathcal{F}_H$$

## Invariant Message Passing NNs

(Lecture 3)

$$\mathbf{m}_{ij} = \text{MLP}(\mathbf{f}_i, \mathbf{f}_j, \|\mathbf{x}_j - \mathbf{x}_i\|)$$

**Non-linear “convolution”**

## Equivariant (Steerable) Message Passing NNs

(Lecture 3)

$$\hat{\mathbf{m}}_{ij} = \widehat{\text{MLP}}(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, \mathbf{x}_j - \mathbf{x}_i)$$

**With steerable MLP:**

$$\widehat{\text{MLP}}_{\hat{\mathbf{a}}_{ij}}(\hat{\mathbf{f}}_i, \hat{\mathbf{f}}_j, \|\mathbf{x}_j - \mathbf{x}_i\|) := \sigma(\mathbf{W}_{\hat{\mathbf{a}}_{ij}}^{(n)}(\dots(\sigma(\mathbf{W}_{\hat{\mathbf{a}}_{ij}}^{(1)}\hat{\mathbf{h}}_i))))$$

# GEOMETRIC AND PHYSICAL QUANTITIES IMPROVE E(3) EQUIVARIANT MESSAGE PASSING

Anonymous authors  
Paper under double-blind review

## ABSTRACT

Including covariant information, such as position, force, velocity or spin is important in many tasks in computational physics and chemistry. We introduce Steerable E(3) Equivariant Graph Neural Networks (SEGNNS) that generalise equivariant graph networks, such that node and edge attributes are not restricted to invariant scalars, but can contain covariant information, such as vectors or tensors. This model, composed of steerable MLPs, is able to incorporate geometric and physical information in both the message and update functions. Through the definition of steerable node attributes, the MLPs provide a new class of activation functions for general use with steerable feature fields. We discuss ours and related work through the lens of *equivariant non-linear convolutions*, which further allows us to pin-point the successful components of SEGNNS: *non-linear* message aggregation improves upon classic *linear* (steerable) point convolutions; *steerable messages* improve upon recent equivariant graph networks that send invariant messages. We demonstrate the effectiveness of our method on several tasks in computational physics and chemistry and provide extensive ablation studies.

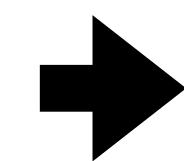
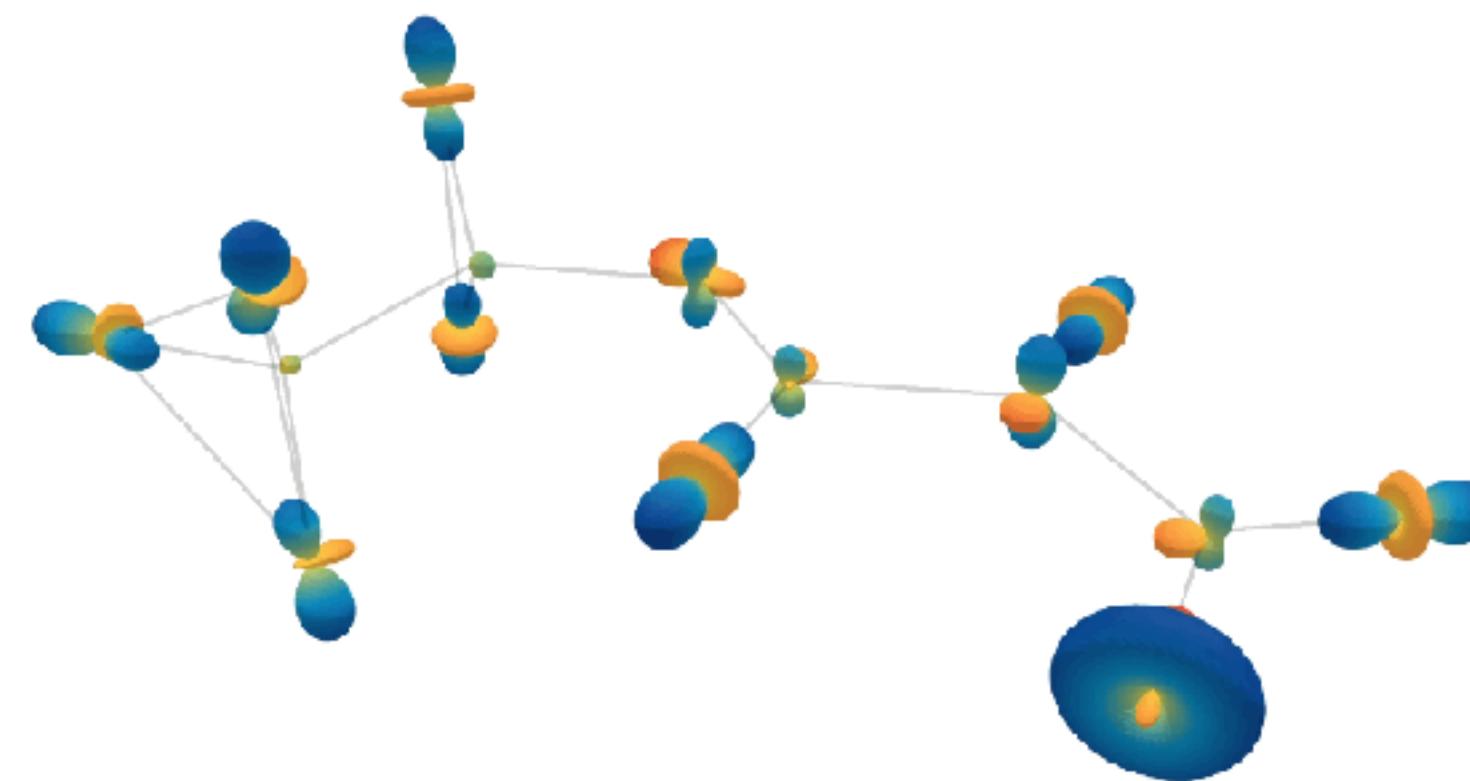
## 1 INTRODUCTION

The success of Convolutional Neural Networks (CNNs) (LeCun et al., 1998, 2015; Schmidhuber, 2015; Krizhevsky et al., 2012) is a key factor for the rise of deep learning, attributed to their capability of exploiting translation symmetries, hereby introducing a strong inductive bias. Recent work has shown that designing CNNs to exploit additional symmetries via group convolutions has even further increased their performance (Cohen & Welling, 2016, 2017; Worrall et al., 2017; Cohen et al., 2018; Kondor & Trivedi, 2018; Weiler et al., 2018; Bekkers et al., 2018; Bekkers, 2019; Weiler & Cesa, 2019). Graph neural networks (GNNs) and CNNs are closely related to each other via their aggregation of local information. More precisely, CNNs can be formulated as message passing layers (Gilmer et al., 2017) based on a sum aggregation of messages that are obtained by relative position-dependent *linear* transformations of neighbouring node features. The power of message passing layers is, however, that node features are transformed and propagated in a highly *non-linear* manner. Equivariant GNNs have been proposed before as either PointConv-type (Wu et al., 2019; Kristof et al., 2017) implementations of steerable (Thomas et al., 2018; Anderson et al., 2019; Fuchs et al., 2020) or regular group convolutions (Pinzi et al., 2020). The most important component in these methods are the convolution layers. Although powerful, such layers only (pseudo)<sup>1</sup> linearly transform the graphs and non-linearity is only obtained via point-wise activations.

In this paper, we propose non-linear E(3) *equivariant message passing* layers using the same principles that underlie steerable group convolutions, and view them as *non-linear group convolutions*. Central to our method is the use of steerable vectors and their equivariant transformations to represent and process node features; we present the underlying mathematics of both in Sec. 2 and illustrate it in Fig. 1 on a molecular graph. As a consequence, information at nodes and edges can now be rotationally invariant (scalar) or covariant (vector, tensor). In steerable message passing frameworks, the Clebsch-Gordan (CG) tensor product is used to steer the update and message functions by geometric information such as relative orientation (pose). Through a notion of steerable node attributes we provide a new class of equivariant activation functions for general use with steerable

<sup>1</sup>Methods such as SE(3)-transformers (Fuchs et al., 2020) and Cormorant (Anderson et al., 2019) include an input-dependent attention component that augments the convolutions.

# Task: Molecular property prediction

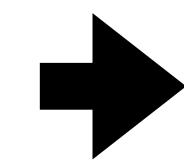
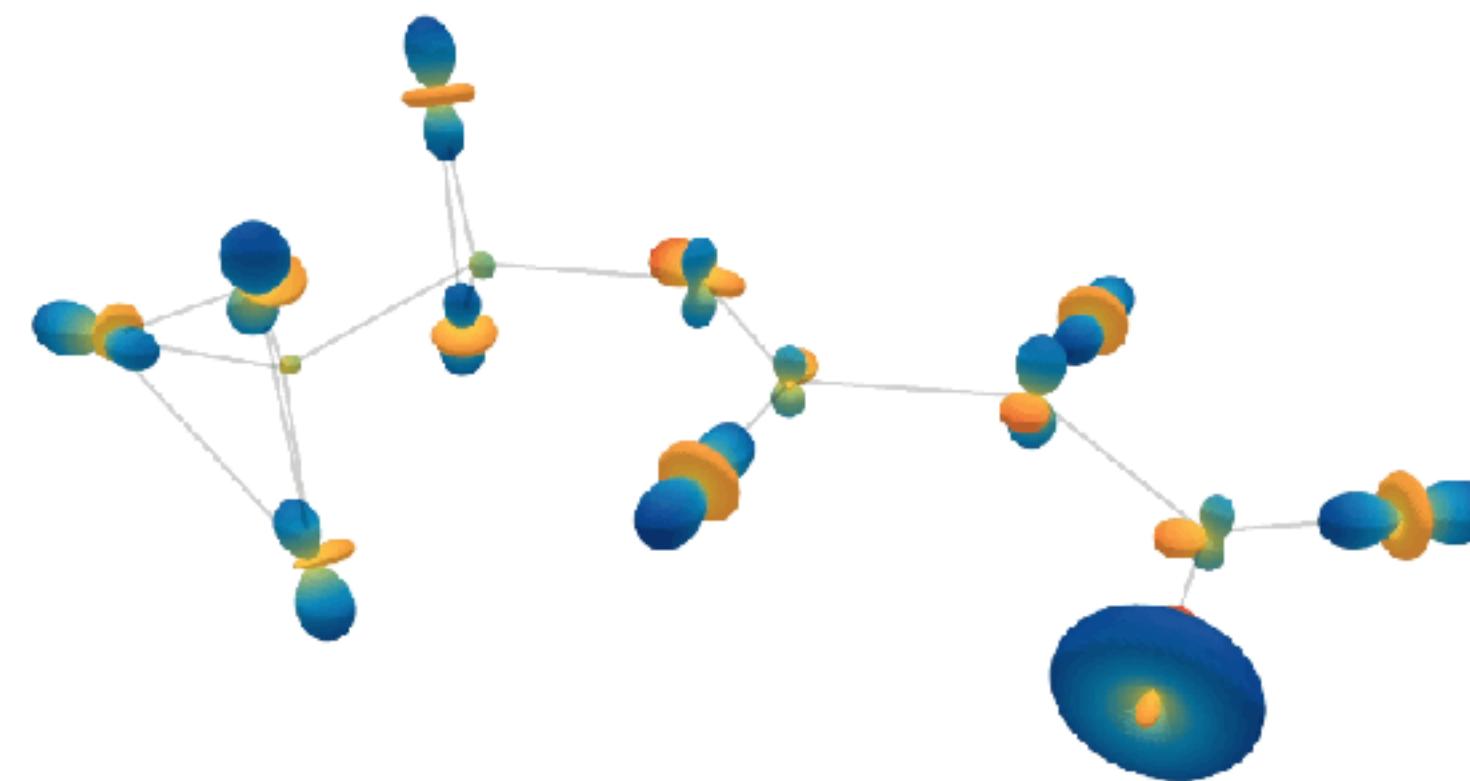


Property

	Task Units	Cutoff radius	$\alpha$ bohr <sup>3</sup>	$\Delta\epsilon$ meV	$\epsilon_{\text{HOMO}}$ meV	$\epsilon_{\text{LUMO}}$ meV	$\mu$ D	$C_\nu$ cal/mol K	Time [s]
<i>Isotropic (fully connected graph)</i>	(S)EGNN ( $l_f = 0, l_a = 0$ )	-	.091	53	34	28	.042	.043	0.016
<i>Isotropic (local)</i>	(S)EGNN ( $l_f = 0, l_a = 0$ )	2Å	.24	98	60	60	.34	.077	0.014
<i>Anisotropic (local)</i>	SEGNN ( $l_f = 1, l_a = 2$ )	2Å	.074	48	27	25	.031	.035	0.048
	SEGNN ( $l_f = 2, l_a = 3$ )	2Å	.060	42	24	21	.023	.031	0.097

**(Steerable) G-CNNs allow for local connectivity (Scales to large proteins!!!)**  
**isotropic convs require full connectivity in order to infer the geometry**

# Task: Molecular property prediction

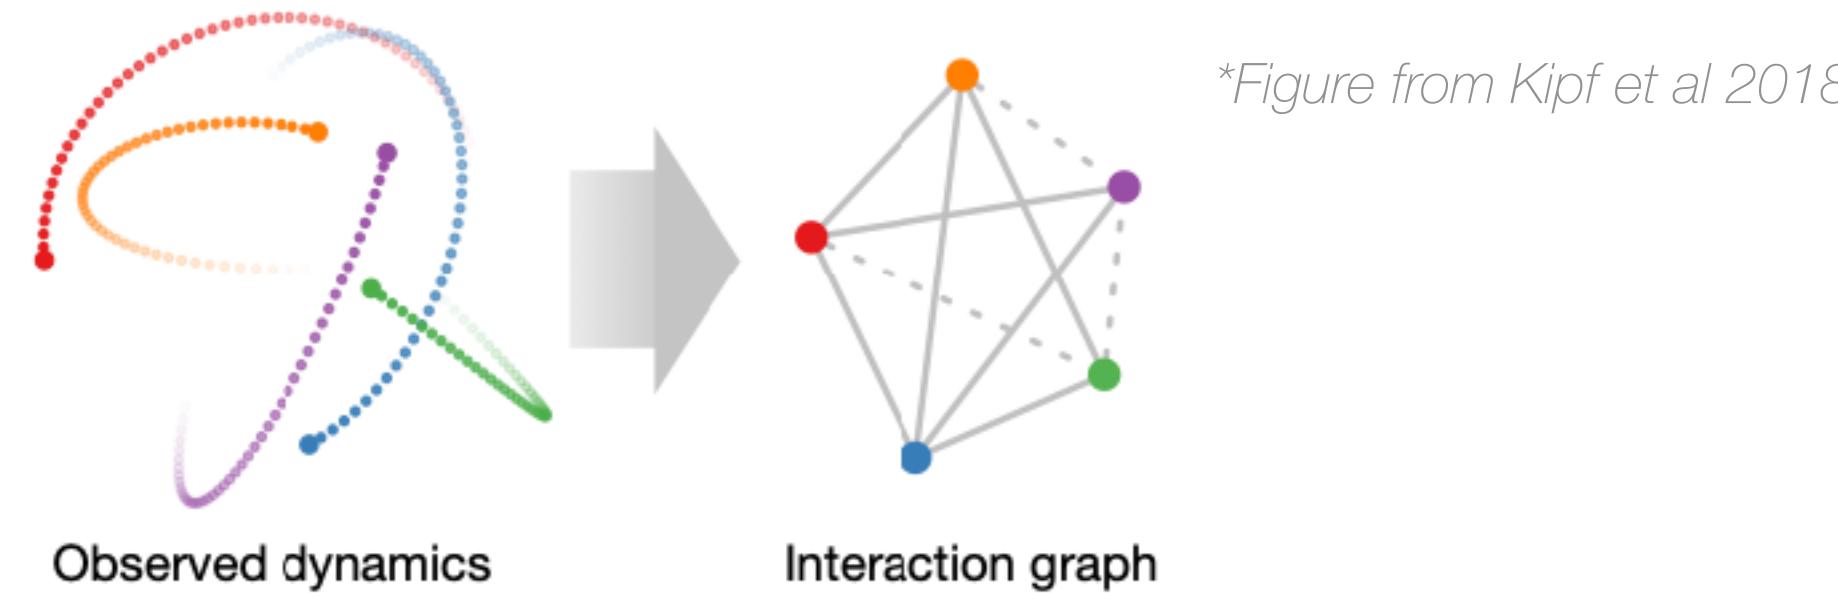


Property

	Task Units	Cutoff radius	$\alpha$ bohr <sup>3</sup>	$\Delta\epsilon$ meV	$\epsilon_{\text{HOMO}}$ meV	$\epsilon_{\text{LUMO}}$ meV	$\mu$ D	$C_\nu$ cal/mol K	Time [s]
<i>Isotropic (fully connected graph)</i>	(S)EGNN ( $l_f = 0, l_a = 0$ )	-	.091	53	34	28	.042	.043	0.016
<i>Isotropic (local)</i>	(S)EGNN ( $l_f = 0, l_a = 0$ )	2Å	.24	98	60	60	.34	.077	0.014
<i>Anisotropic (local)</i>	SEGNN ( $l_f = 1, l_a = 2$ )	2Å	.074	48	27	25	.031	.035	0.048
	SEGNN ( $l_f = 2, l_a = 3$ )	2Å	.060	42	24	21	.023	.031	0.097

**(Steerable) G-CNNs allow for local connectivity (Scales to large proteins!!!)**  
**isotropic convs require full connectivity in order to infer the geometry**

# Task: Trajectory prediction N-body problem

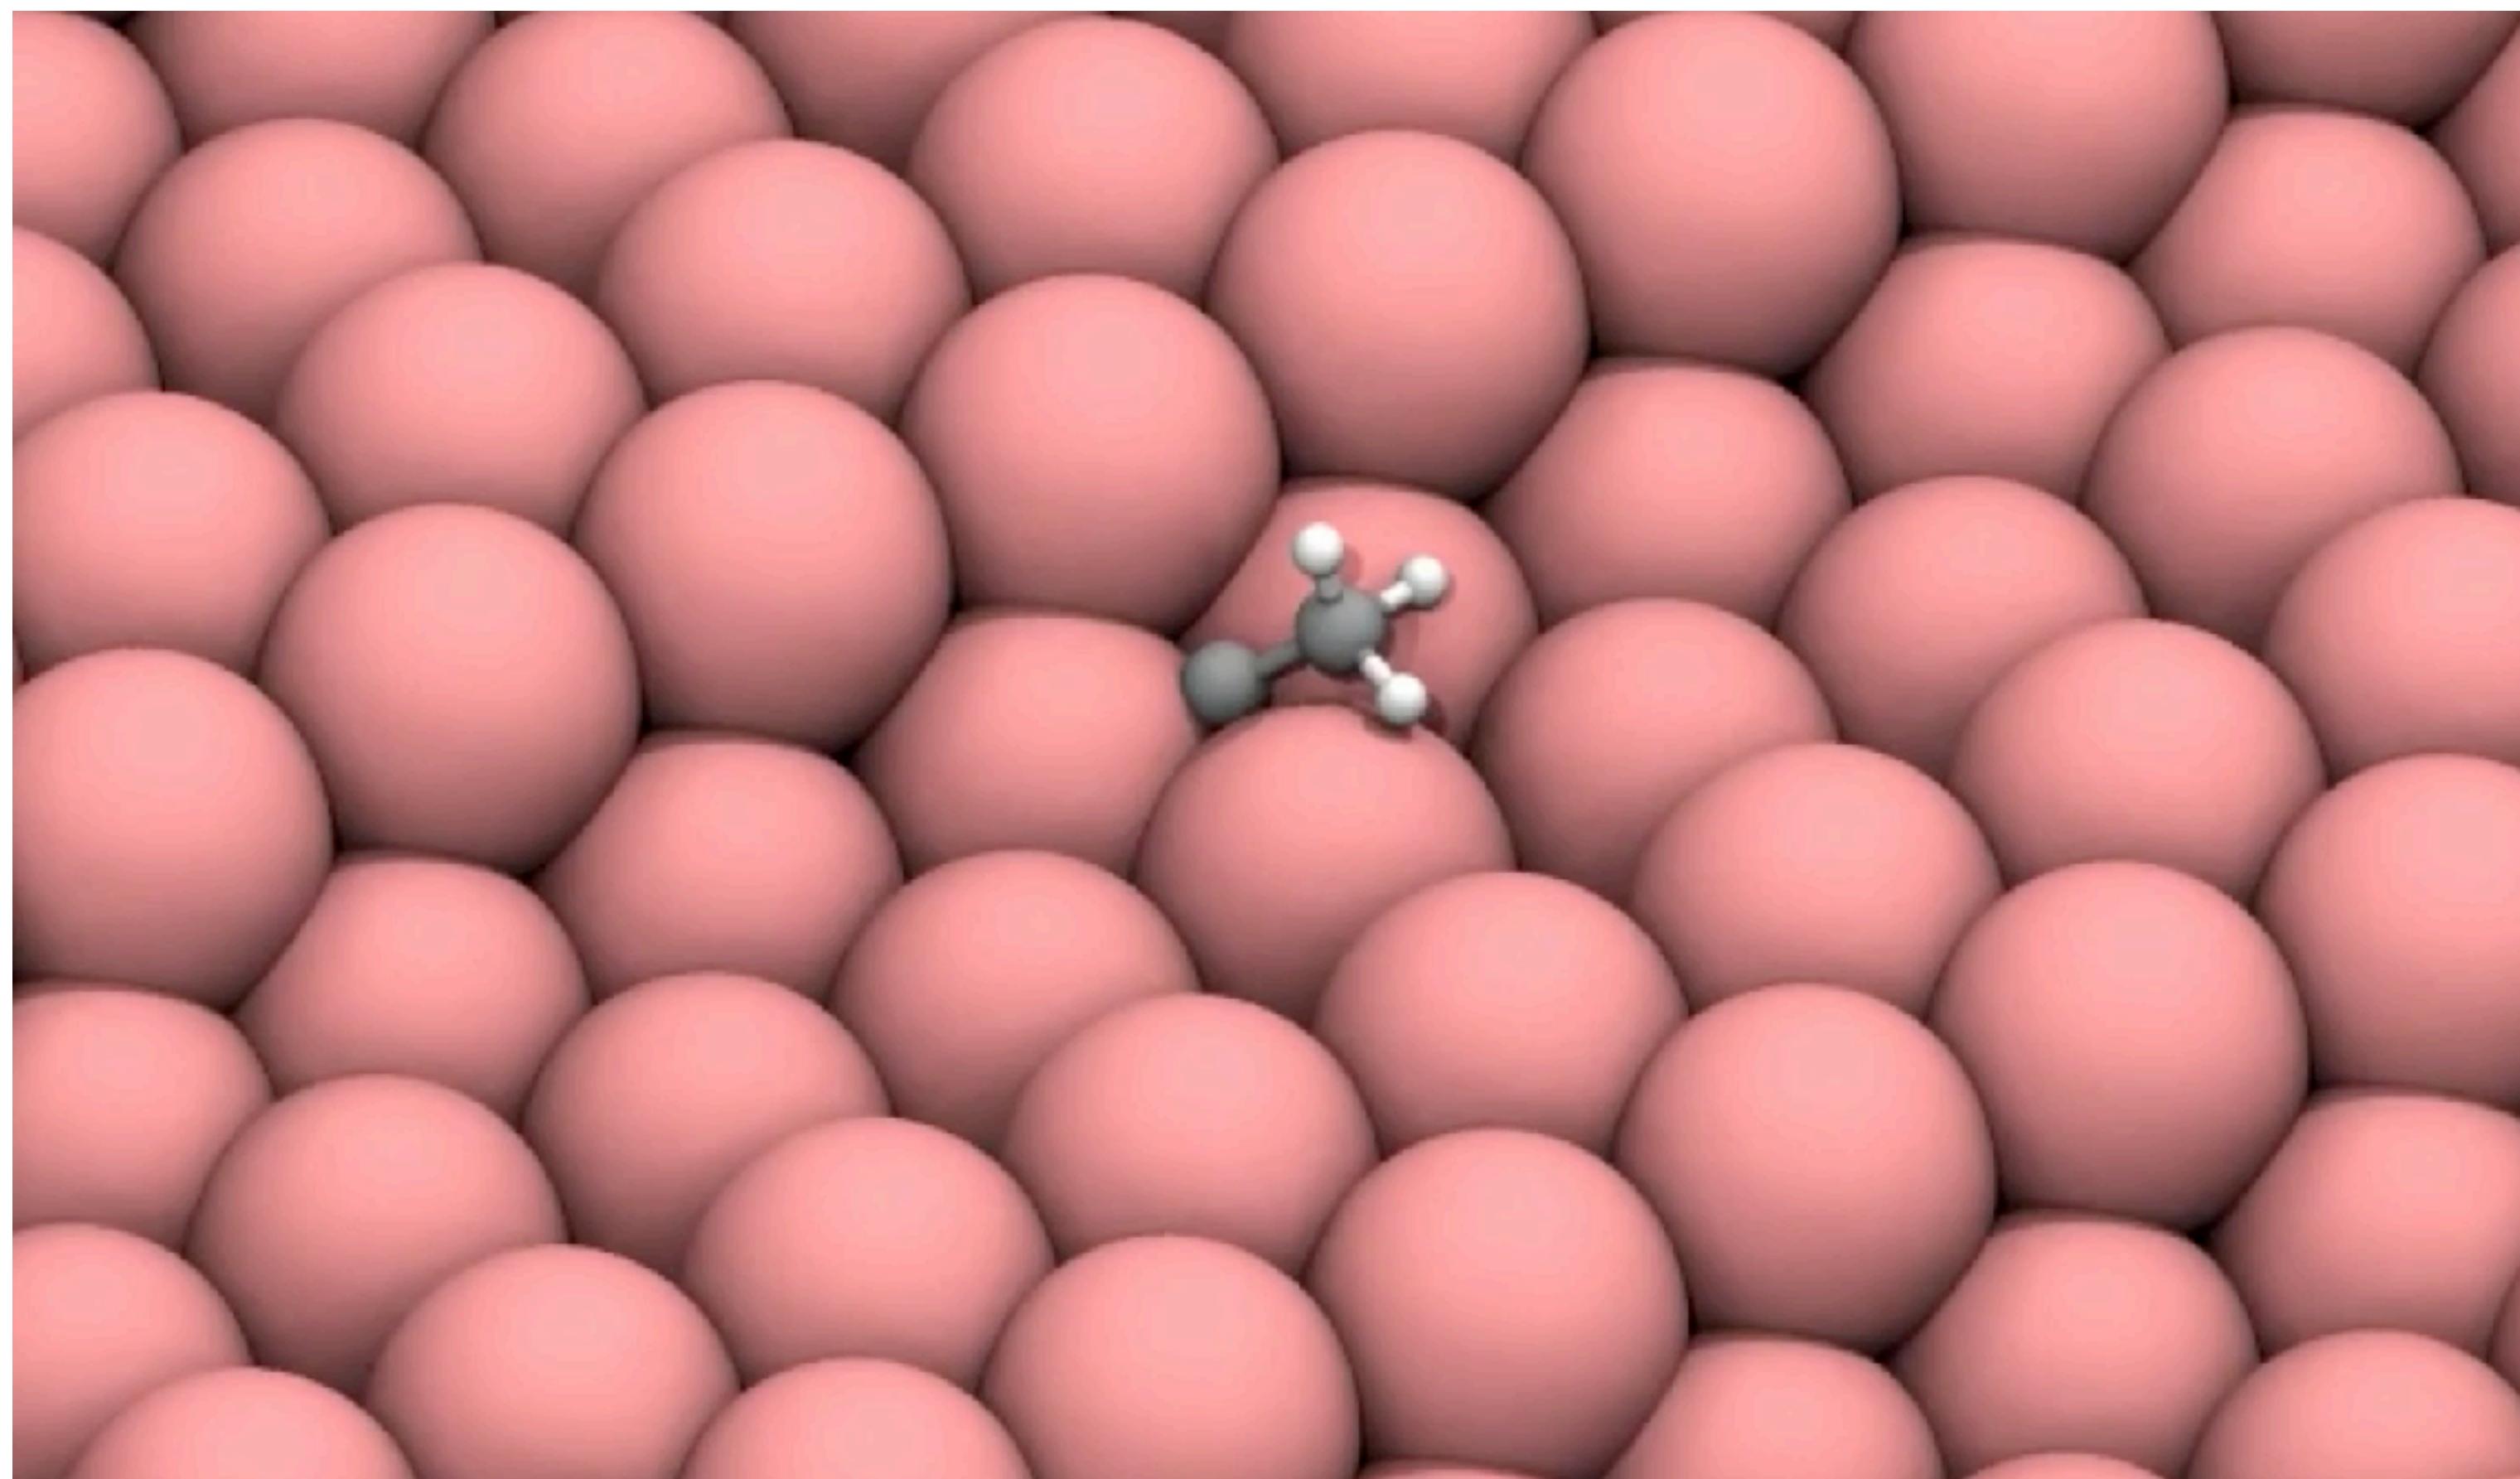


	Method	MSE	Time [s]
G-CNNs	Linear	.0819	.0001
	SE(3)-Tr. (Fuchs et al., 2020)	.0244	.0742
	TFN (Thomas et al., 2018)	.0155	.0182
	NMP (Gilmer et al., 2017)	.0107	.0017
	Radial Field (Köhler et al., 2019)	.0104	.0019
	EGNN (Satorras et al., 2021)	.0070± .00022	.0029
Isotropic “non-linear CNNs”	SE <sub>linear</sub> ( $l_f = 2, l_a = 2$ )	.0116± .00021	.064
	SE <sub>non-linear</sub> ( $l_f = 1, l_a = 1$ )	.0060± .00019	.031
	SEGNN <sub>G</sub> ( $l_f = 1, l_a = 1$ )	.0056± .00025	.025
	SEGNN <sub>G+P</sub> ( $l_f = 1, l_a = 1$ )	.0043± .00015	.026

G-CNNs outperform CNNs with isotropic kernels  
“Non-linear convolutions” outperform linear convolutions

# Steerable methods for computational chemistry

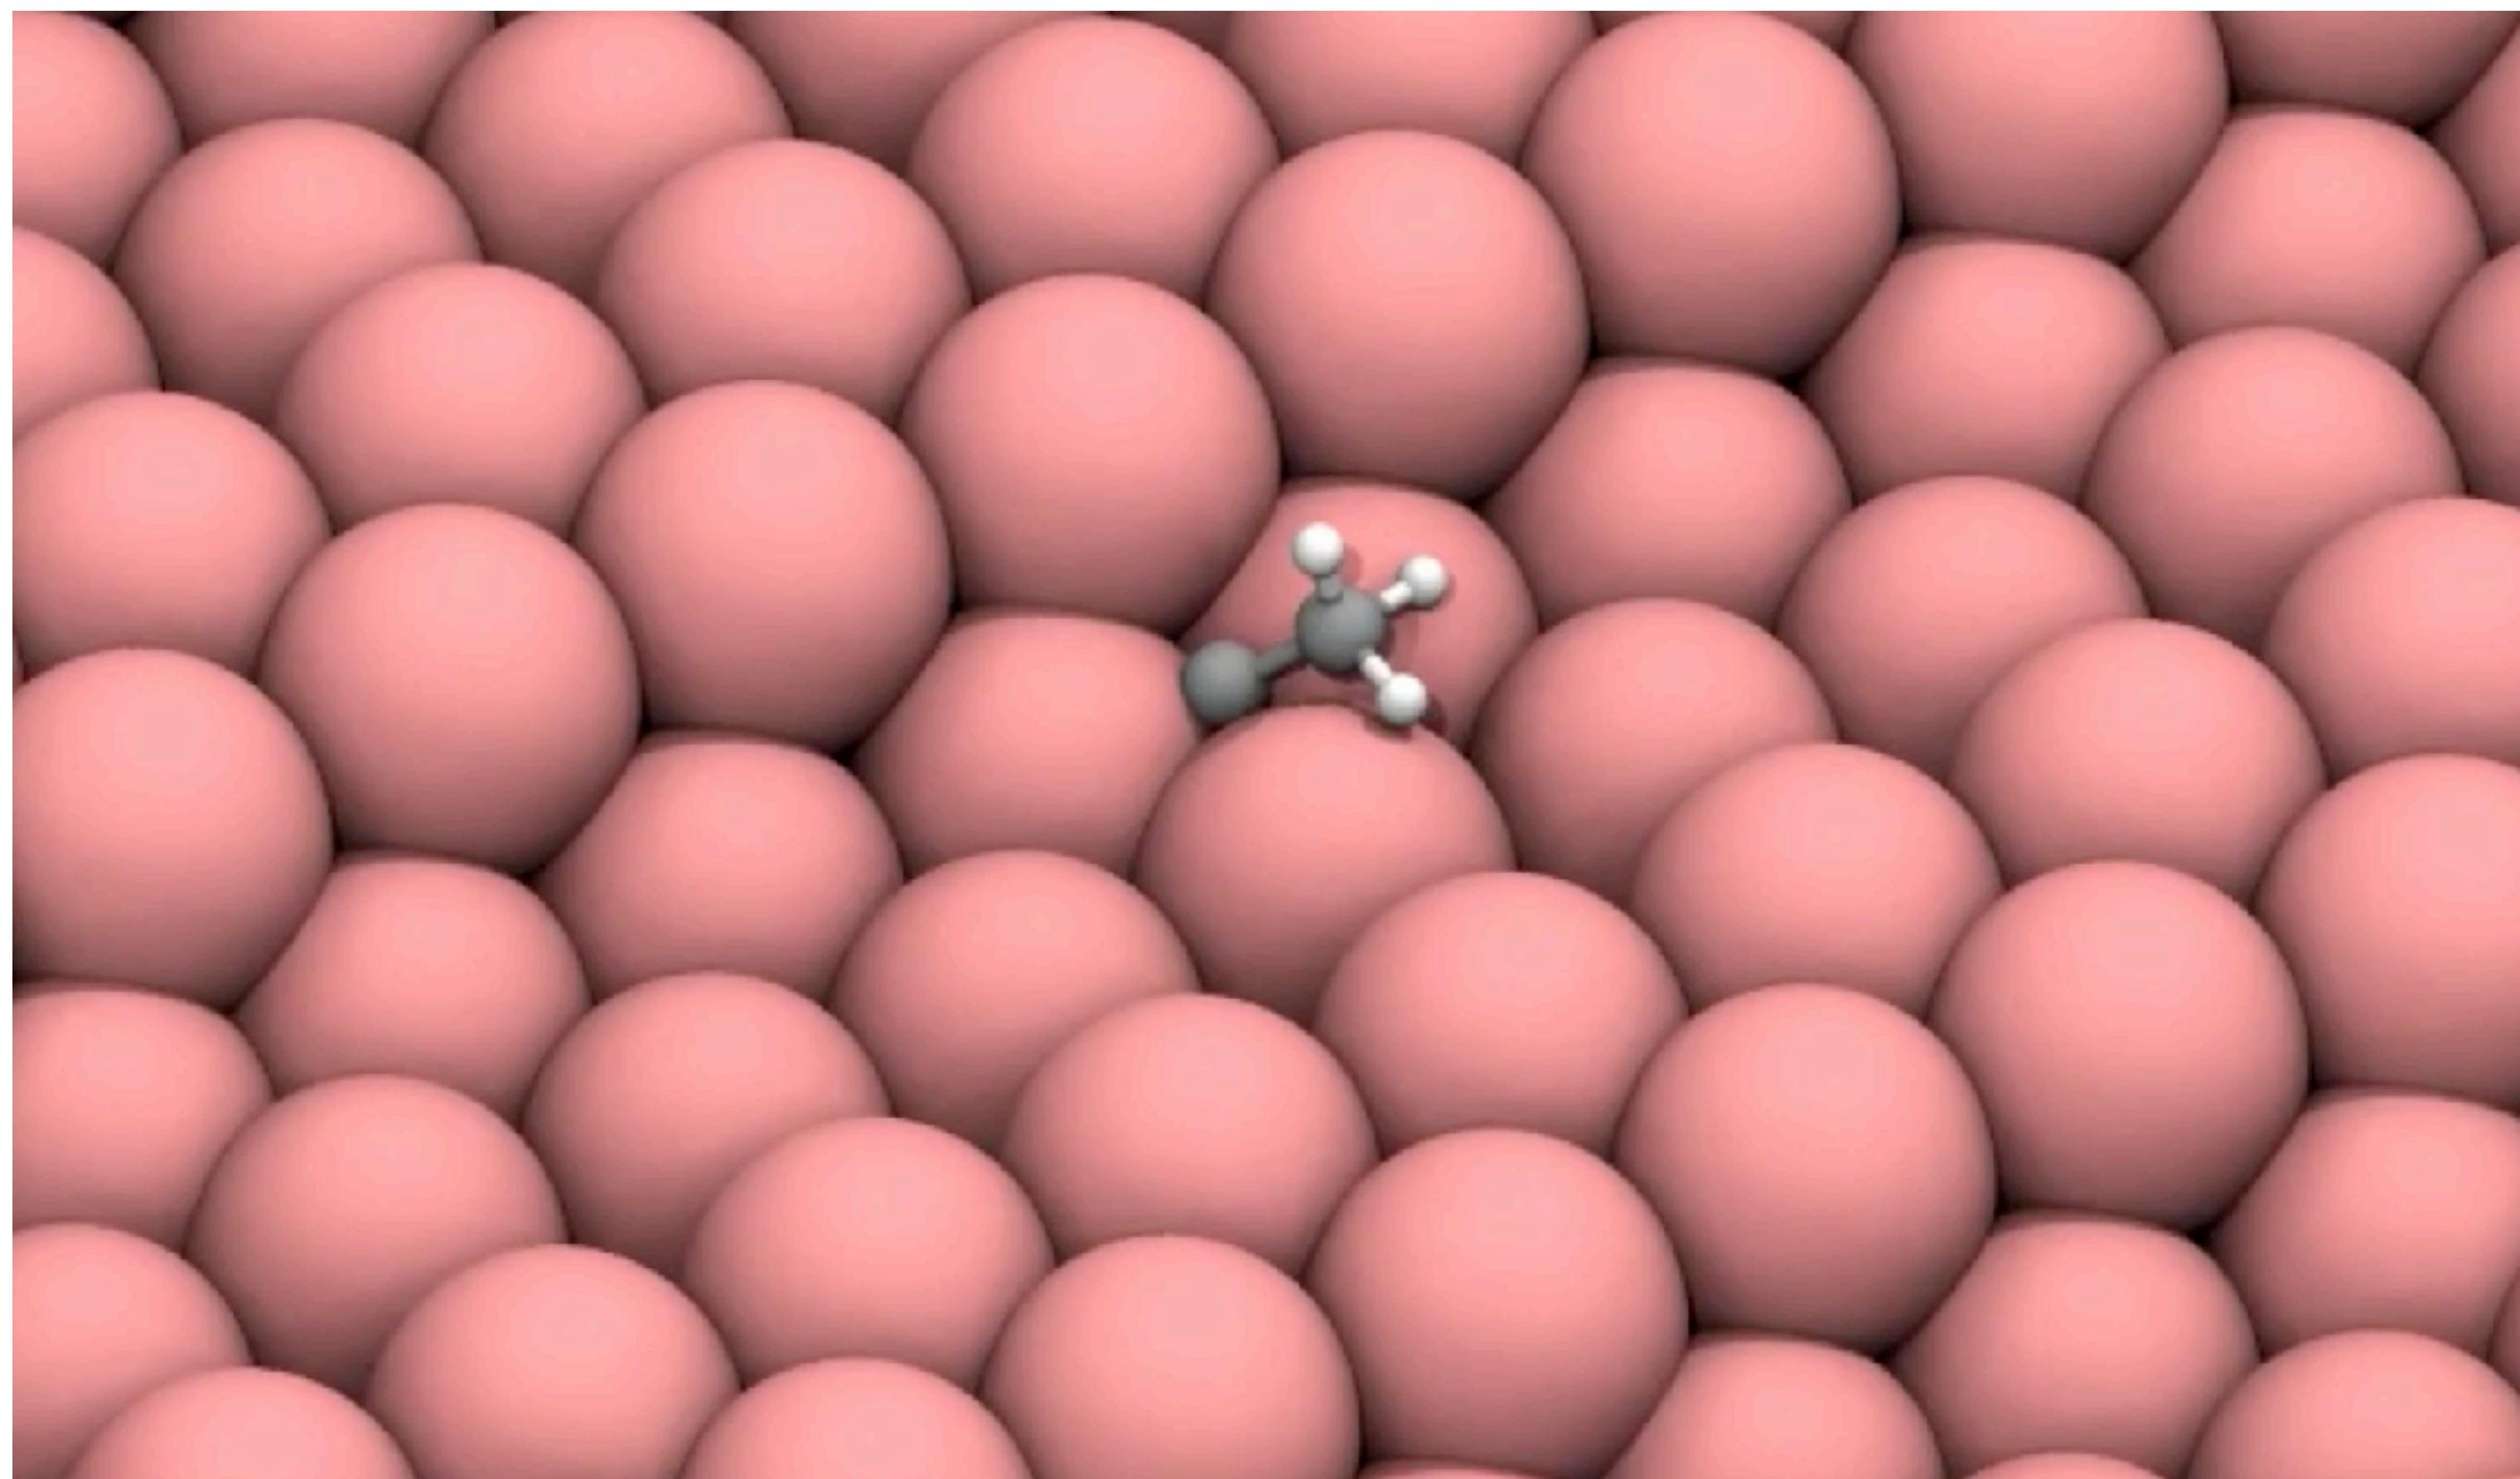
Brandstetter, Hesselink, van der Pol, Bekkers, Welling **Geometric and Physical Quantities Improve E(3) Equivariant Message Passing** - arXiv:2110.02905



Video: Open Catalyst Project

# Steerable methods for computational chemistry

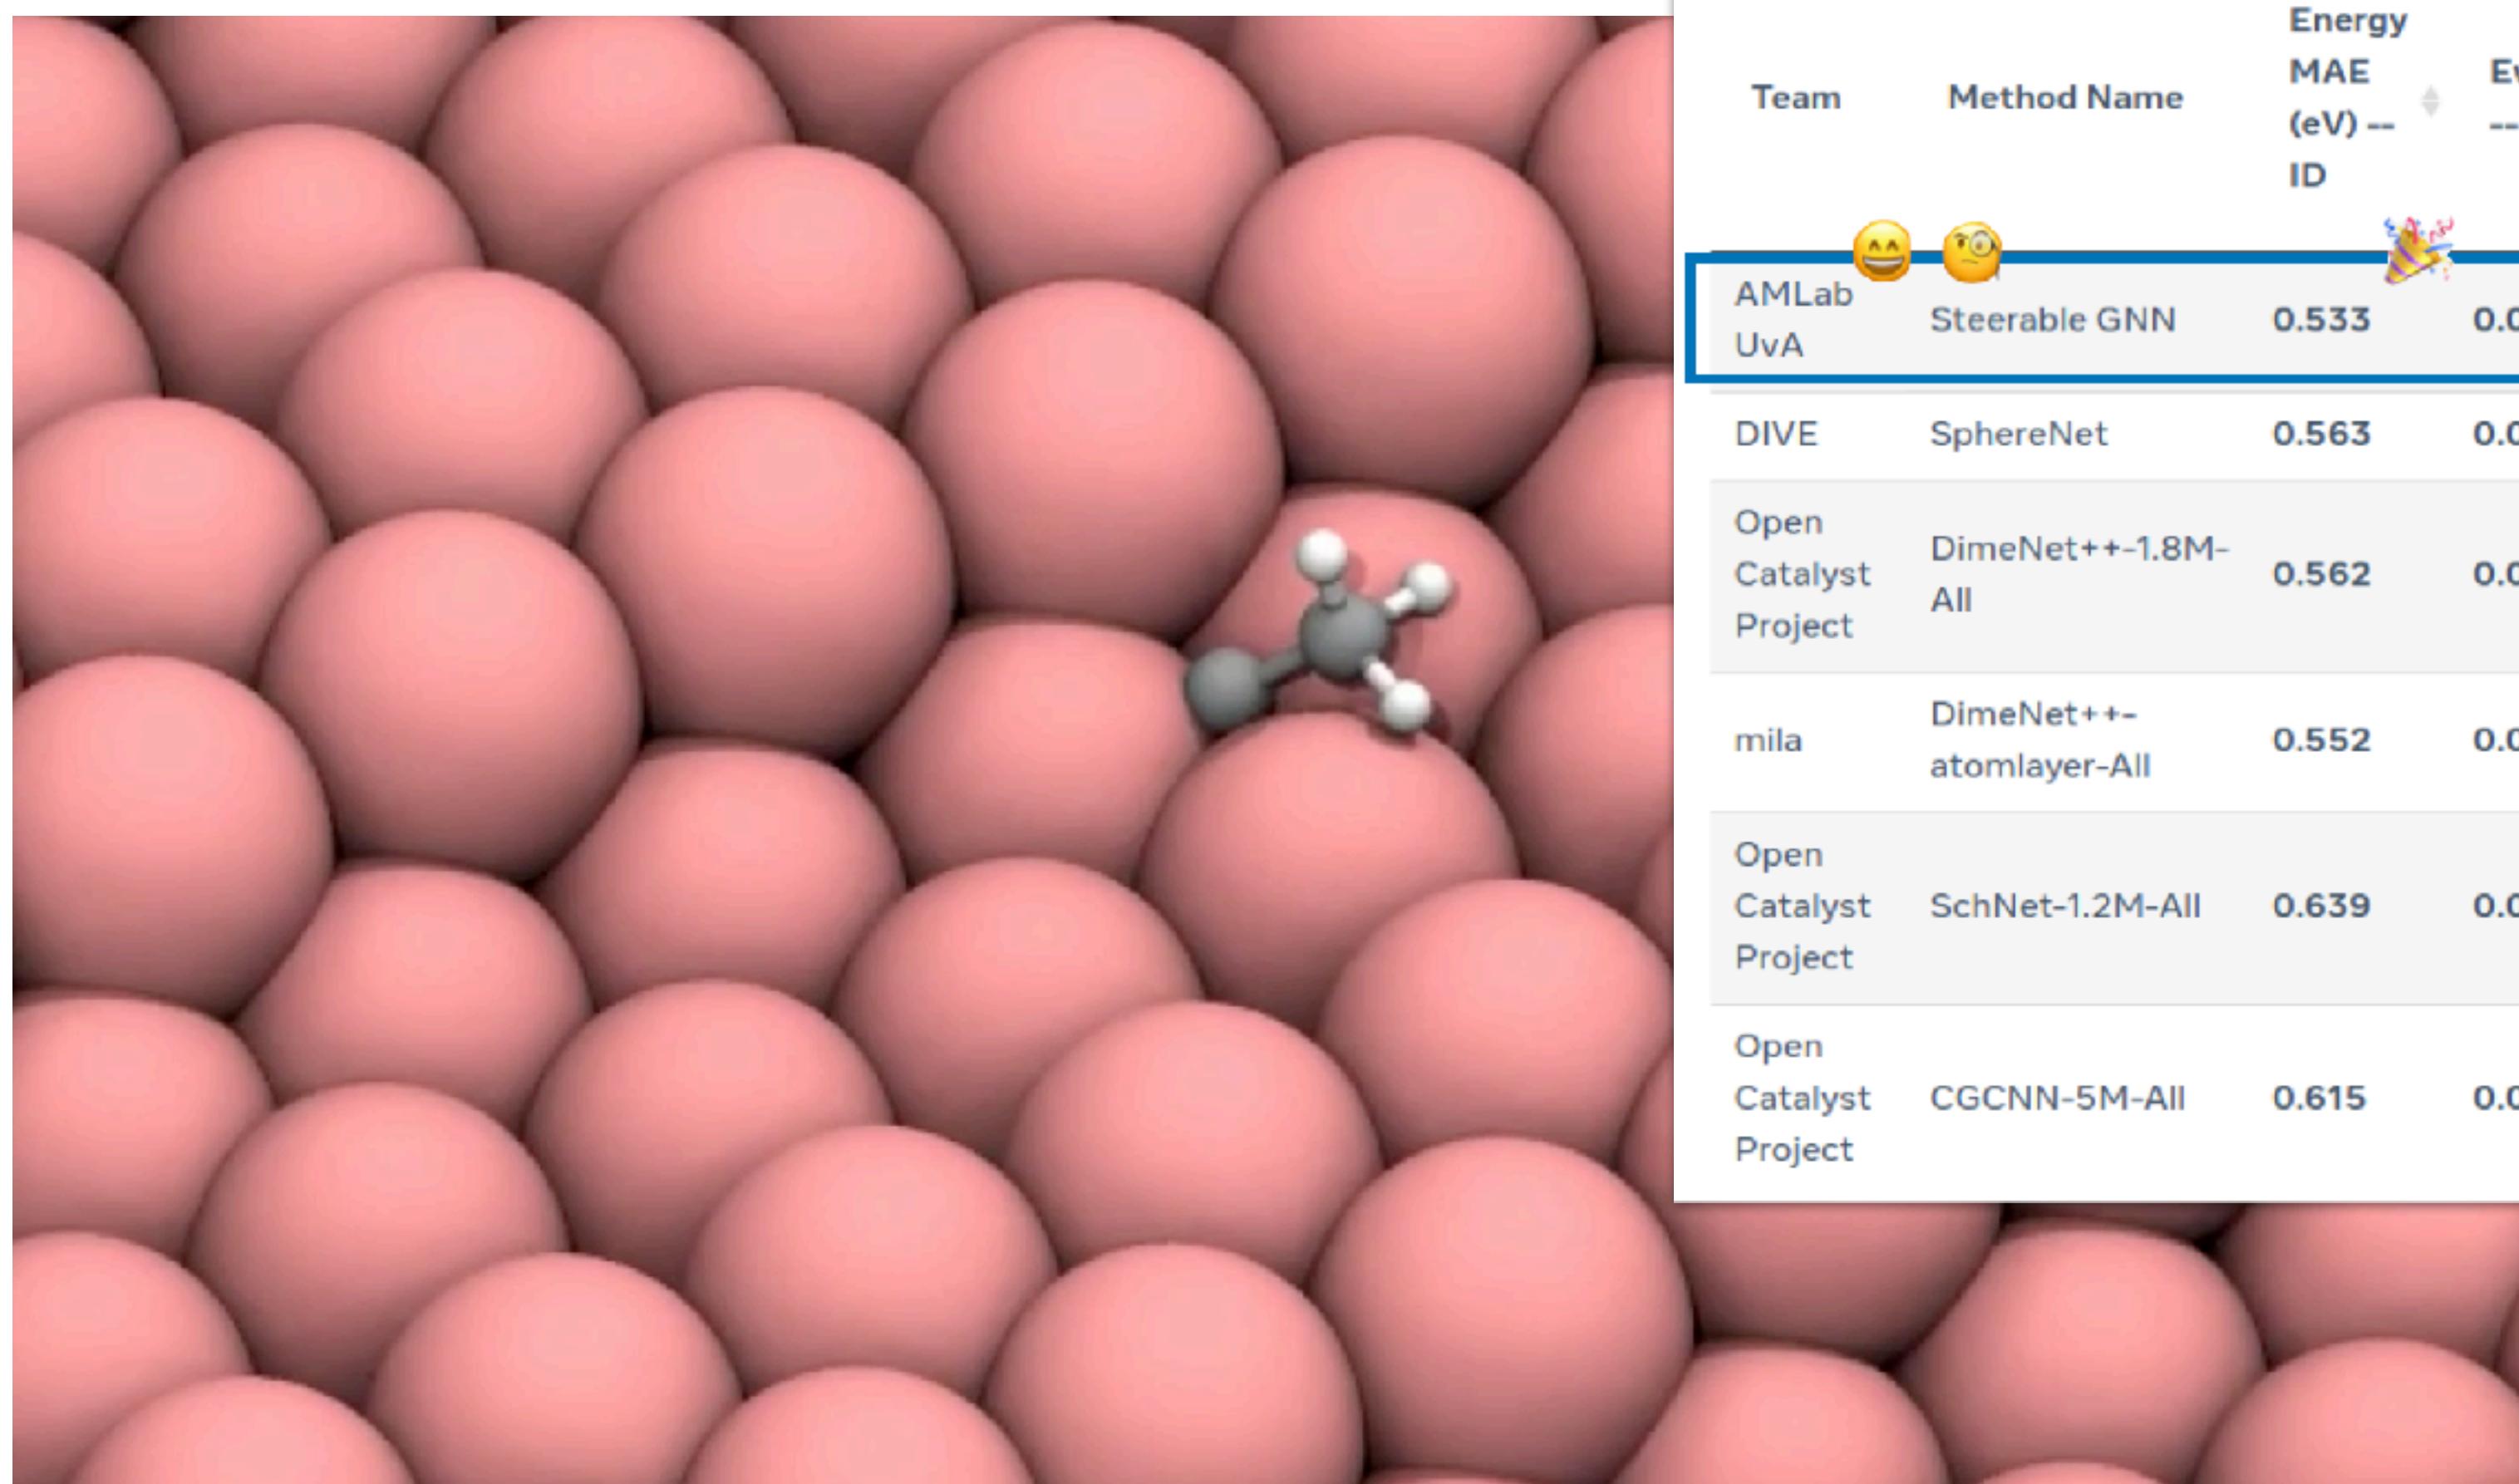
Brandstetter, Hesselink, van der Pol, Bekkers, Welling **Geometric and Physical Quantities Improve E(3) Equivariant Message Passing** - arXiv:2110.02905



Video: Open Catalyst Project

# Steerable methods for computational chemistry

Brandstetter, Hesselink, van der Pol, Bekkers, Welling **Geometric and Physical Quantities Improve E(3) Equivariant Message Passing** - arXiv:2110.02905



Team	Method Name	Energy MAE (eV) -- ID	EwT -- ID	Energy MAE (eV) -- OOD Ads	EwT -- OOD Ads	Energy MAE (eV) -- OOD Cat	EwT -- OOD Cat	Energy MAE (eV) -- OOD Both	EwT -- OOD Both	Submitted
AMLab UvA	Steerable GNN	0.533	0.0537	0.692	0.0246	0.537	0.0492	0.679	0.0263	2021/06/12
DIVE	SphereNet	0.563	0.0447	0.703	0.0229	0.571	0.0409	0.638	0.0241	2021/03/31
Open Catalyst Project	DimeNet++-1.8M-All	0.562	0.0425	0.725	0.0207	0.576	0.041	0.661	0.0241	2021/02/16
mila	DimeNet++-atomlayer-All	0.552	0.0489	0.747	0.0259	0.557	0.0459	0.688	0.0233	2021/06/07
Open Catalyst Project	SchNet-1.2M-All	0.639	0.0296	0.734	0.0233	0.662	0.0294	0.704	0.0221	2021/02/17
Open Catalyst Project	CGCNN-5M-All	0.615	0.034	0.915	0.0193	0.622	0.031	0.851	0.02	2021/02/18

Video: Open Catalyst Project

**Table 2: Comparison on QM9.****Table 2: Performance comparison on the QM9 dataset. Mean Error (MAE) between model predictions and ground truth**

		Task Units	$\alpha$ bohr <sup>3</sup>	$\Delta\epsilon$ meV	$\epsilon_{\text{HOMO}}$ meV	$\epsilon_{\text{LUMO}}$ meV	$\mu$ D	$C_v$ cal/mol l
<b>non-linear</b>	regular	no geometry	NMP	.092	69	43	38	.030
		$\mathbb{R}^3$	SchNet *	.235	63	41	34	.033
<b>pseudo-linear</b>	steerable	$\mathbb{R}^3$	Cormorant	.085	61	34	38	.038
	steerable	$SE(3)$	L1Net	.088	68	46	35	.043
<b>pseudo-linear</b>	regular	$G$	LieConv	.084	49	30	25	.032
	steerable	$SE(3)$	TFN	.223	58	40	38	.064
<b>non-linear</b>	steerable	$SE(3)$	SE(3)-Tr.	.142	53	35	33	.051
	regular	$\mathbb{R}^3 \times S^2 \times \mathbb{R}^+$	DimeNet++ *	<b>.043</b>	<b>32</b>	24	19	.029
<b>non-linear</b>	regular	$\mathbb{R}^3 \times S^2 \times \mathbb{R}^+$	SphereNet *	.046	<b>32</b>	<b>23</b>	<b>18</b>	.026
<b>non-linear</b>	regular?	$SE(3)$	PaiNN *	.045	45	27	20	.012
<b>non-linear</b>	regular	$\mathbb{R}^3$	EGNN	.071	48	29	25	.029
<b>non-linear</b>	steerable	$SE(3)$	SEGNN (Ours)	.060	42	24	21	.031

1. Motivation

Equivariance → weight-sharing and generalization

2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

3. Group convolutions

Template matching over groups

4. Example

Effective representation learning and generalization

5. G-convs are all you need!

Any equivariant linear layer is a group convolution

6. Steerable group convolutions

Efficient (band-limited) **grid-free** g-convs

7. Feature fields and escnn library

Flexible framework for equivariant layers

8. Equivariant tensor product layers

Conv layers  $\leftrightarrow$  TPs with coordinate embeddings  
(Clebsch-Gordan: equivariant TP)

9. Equivariant graph NNs

Equivariant MP via geometry conditioned layers

## 1. Motivation

Equivariance → weight-sharing and generalization

## 2. Pattern matching using group theory

Group theory: symmetries & recognition by components  
(features have “poses”)

## 3. Group convolutions

Template matching over groups

## 4. Example

Effective representation learning and generalization

## 5. G-convs are all you need!

Any equivariant linear layer is a group convolution

## 6. Steerable group convolutions

Efficient (band-limited) **grid-free g-convs**

## 7. Feature fields and escnn library

Flexible framework for equivariant layers

## 8. Equivariant tensor product layers

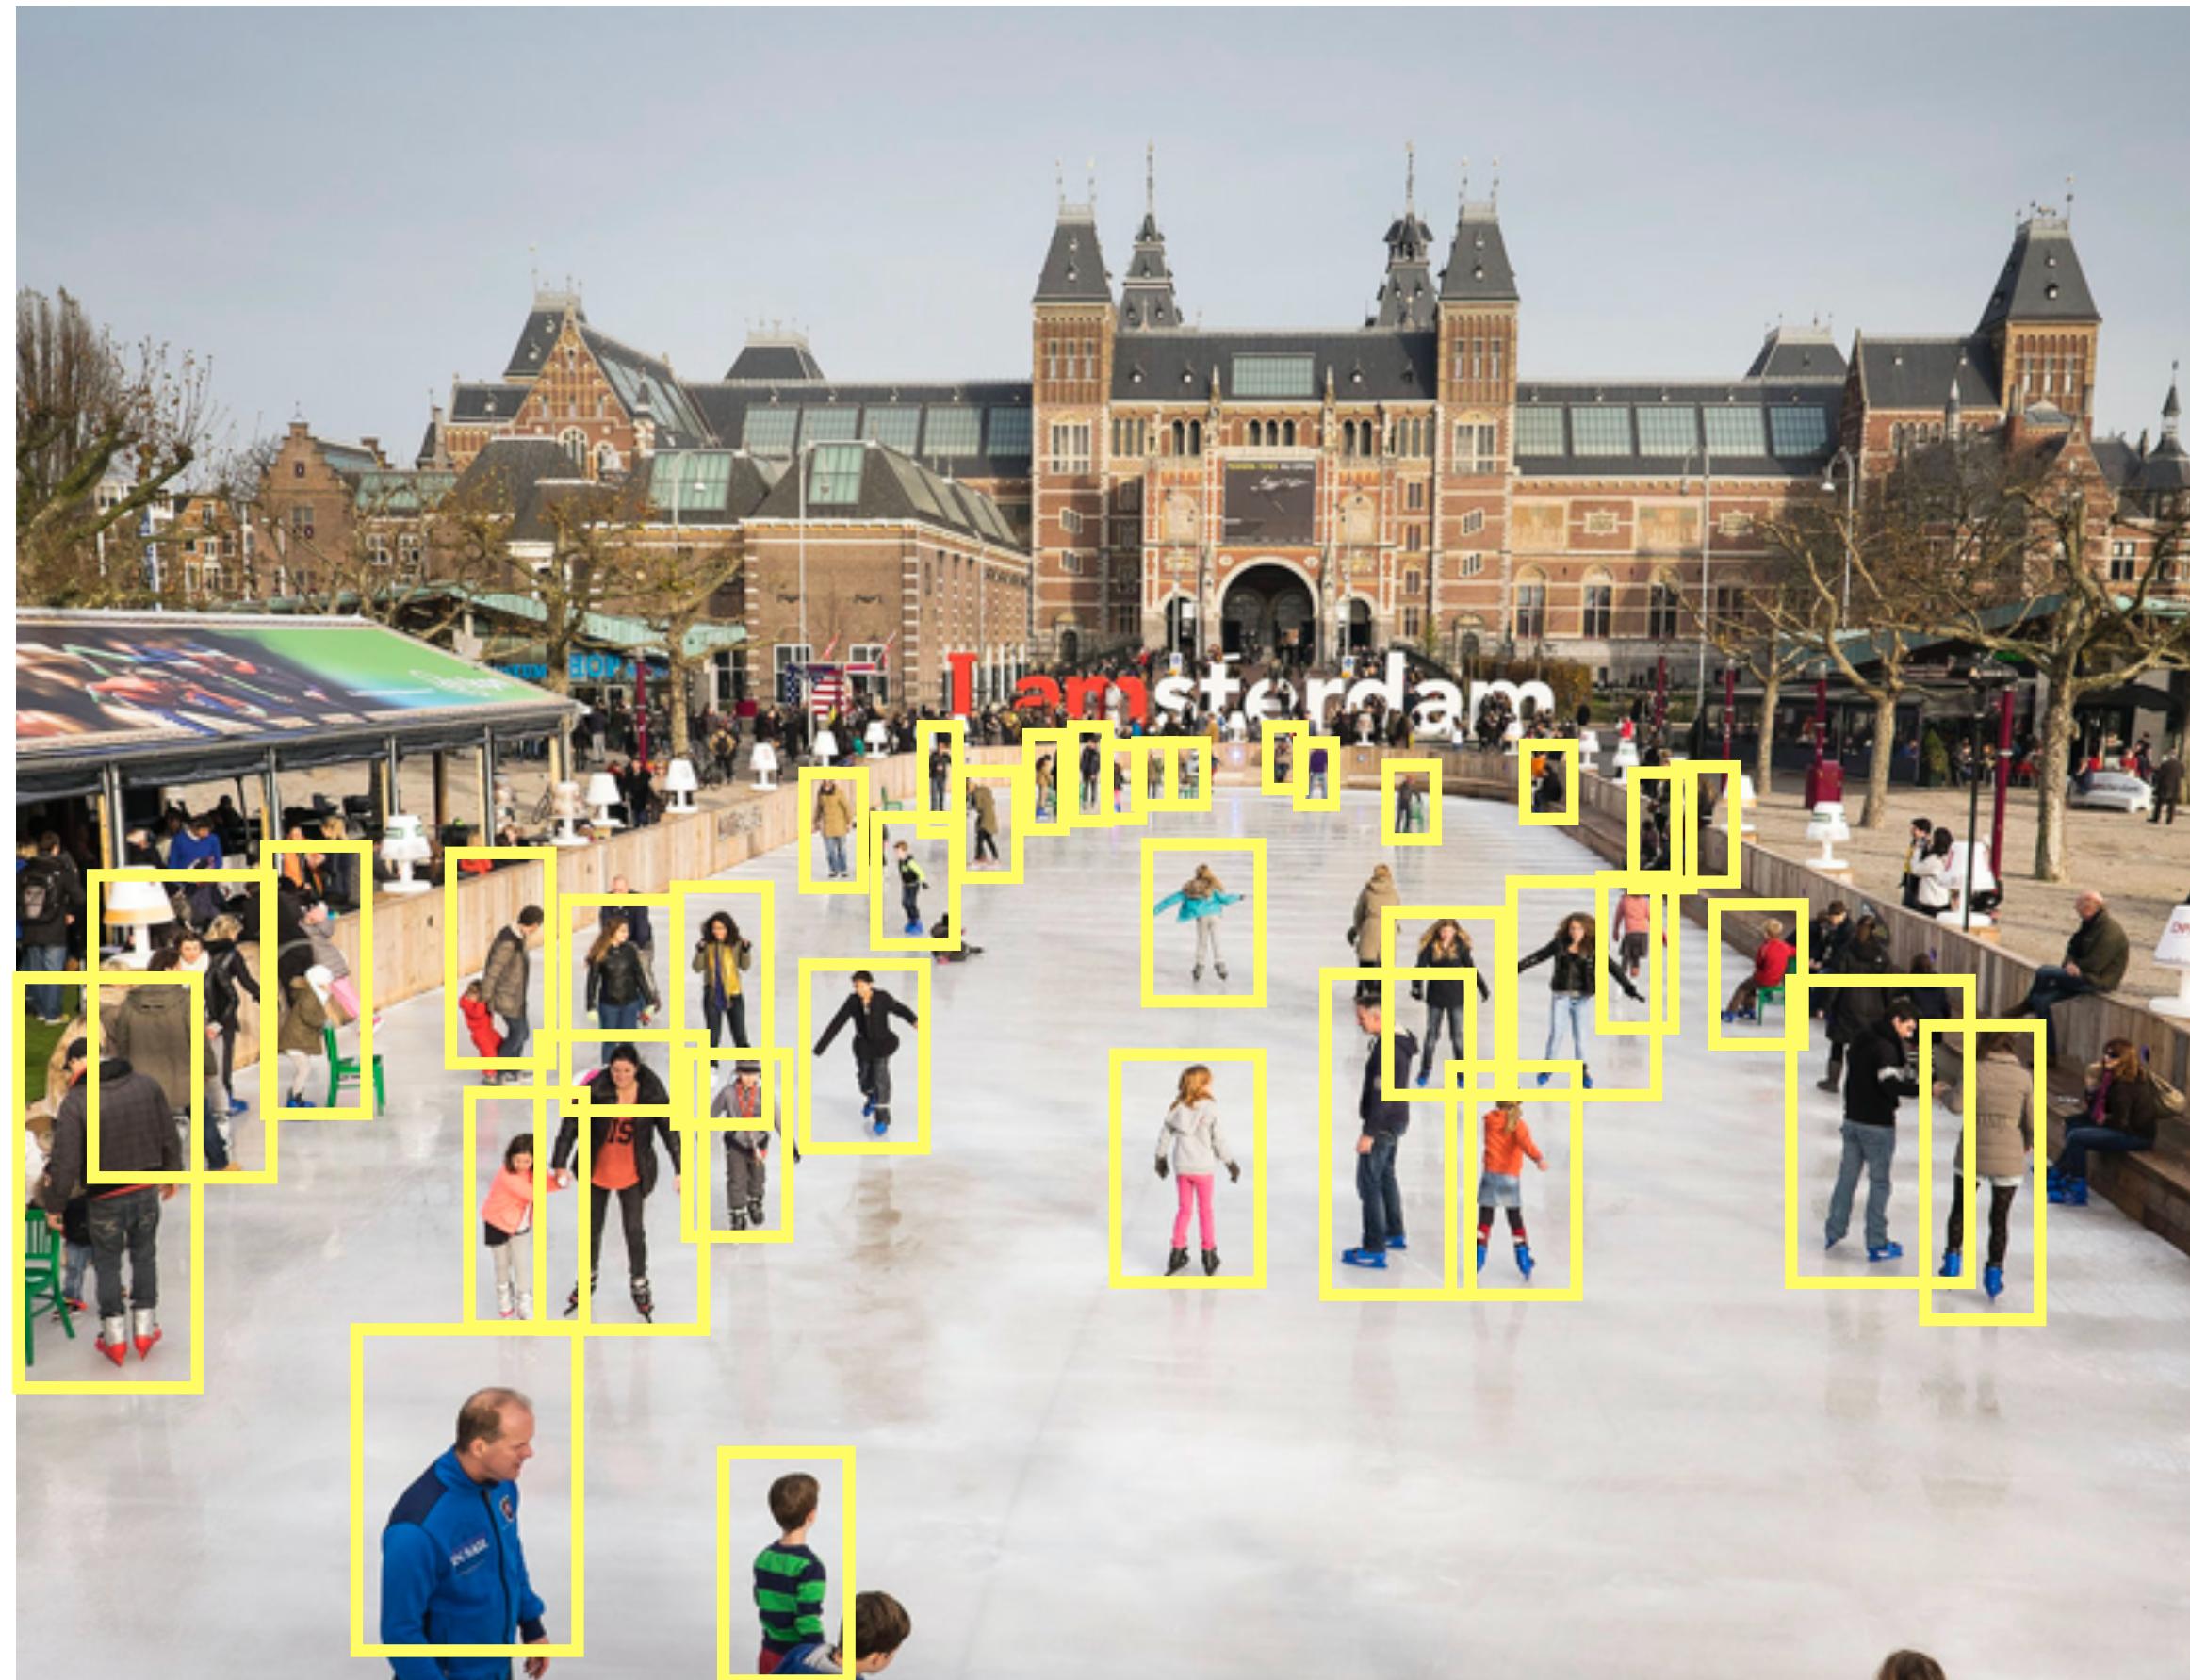
Conv layers  $\leftrightarrow$  TPs with coordinate embeddings  
(Clebsch-Gordan: equivariant TP)

## 9. Equivariant graph NNs

Equivariant MP via geometry conditioned layers

# Conclusion

# Geometric guarantees (equivariance)



Importance of equivariance:

- No information is lost when the input is transformed
- Guaranteed stability to (local + global) transformations

Group convolutions:

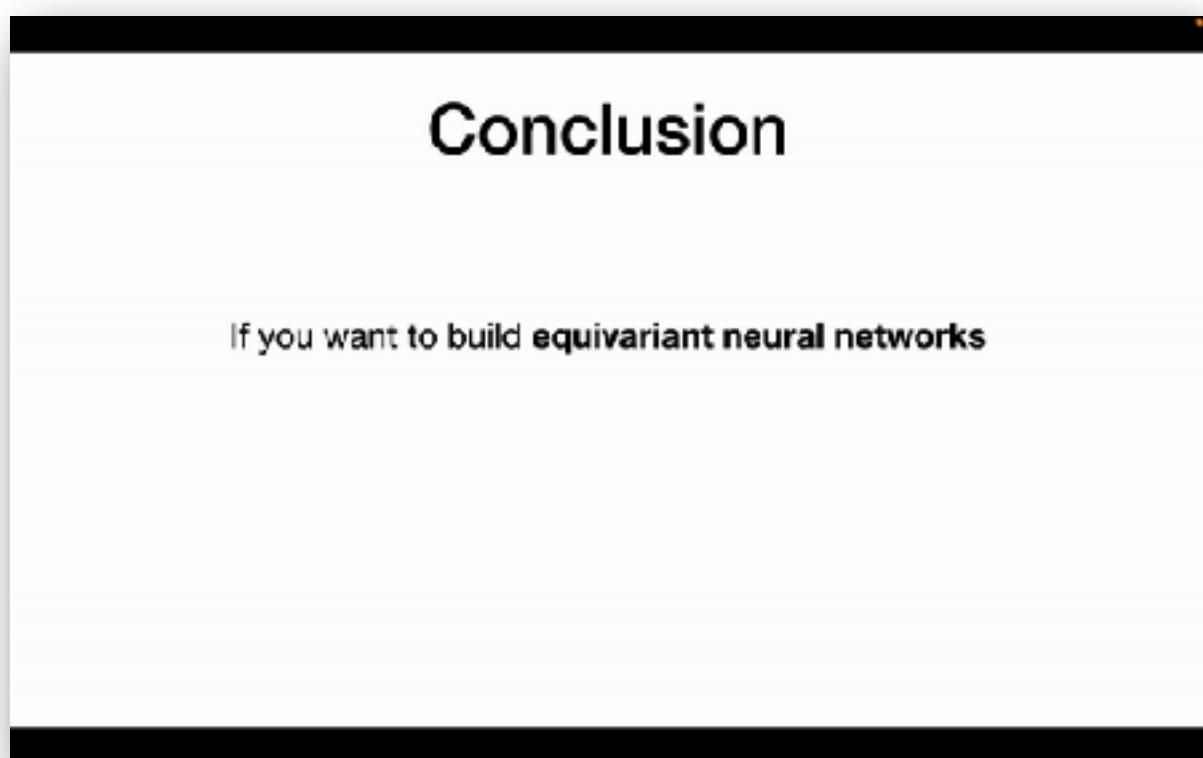
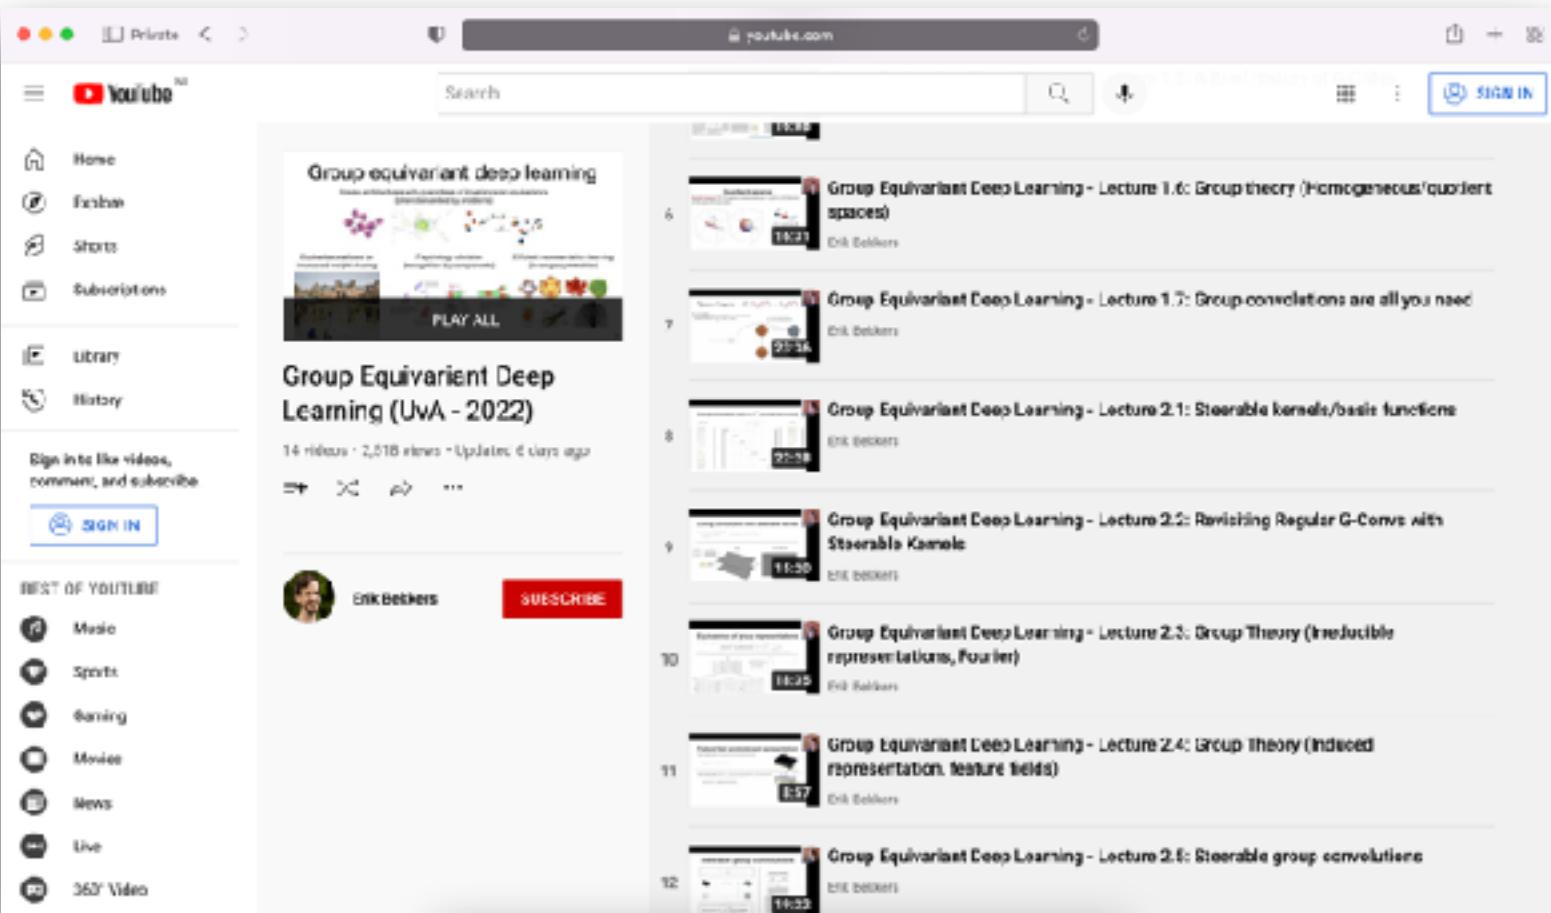
- Equivariance beyond translations
- Geometric guarantees
- Increased weight sharing

+ performance  
+ generalization

**G-CNNs are not only relevant for invariant problems but for any type of structured data!**

# UvA course on group equivariant deep learning (<https://uvagedl.github.io>)

## Youtube playlist



## Tutorial notebooks

[View on GitHub](#)

### GDL - Regular Group Convolutions

This notebook covers the implementation of regular group convolutional networks from scratch, using PyTorch. It includes sections on introduction, brief recap on CNNs, and a detailed explanation of the convolution operation as an inner product of the function  $f$  and a shifted kernel  $k$ .

[View on GitHub](#)

### GDL - Steerable CNNs

This notebook introduces steerable CNNs, a general framework for convolutional networks. It covers topics like working with the Uva cluster, research projects with PyTorch, and implementing a steerable CNN layer. It also discusses the Fourier transform and its application in steerable CNNs.

[View on GitHub](#)

### Introduction

The goal of this tutorial is to get you familiar with the ways in which positional information is used in graph neural networks. We will firstly look at superpixel CNNs, which groups pixels by similarity in order to create a graph. Then, we will apply steerable methods to MD17, a 3D dataset of molecular relaxation trajectories.

This tutorial was created by Petri van der Linden and Ruò Hesselink. If you would like to know more about these methods, feel free to contact us at [r.vanderlinden@uva.nl](mailto:r.vanderlinden@uva.nl) or [p.v.hesselink@uva.nl](mailto:p.v.hesselink@uva.nl).

Below you can find several papers that we recommend, if you're interested in some more reading.

- [Neural Message Passing for Quantum Chemistry](#)
- [Directional Message Passing for Molecular Graphs](#)
- [Equivariant Graph Neural Networks for Data-Efficient and Accurate Interatomic Potentials](#)

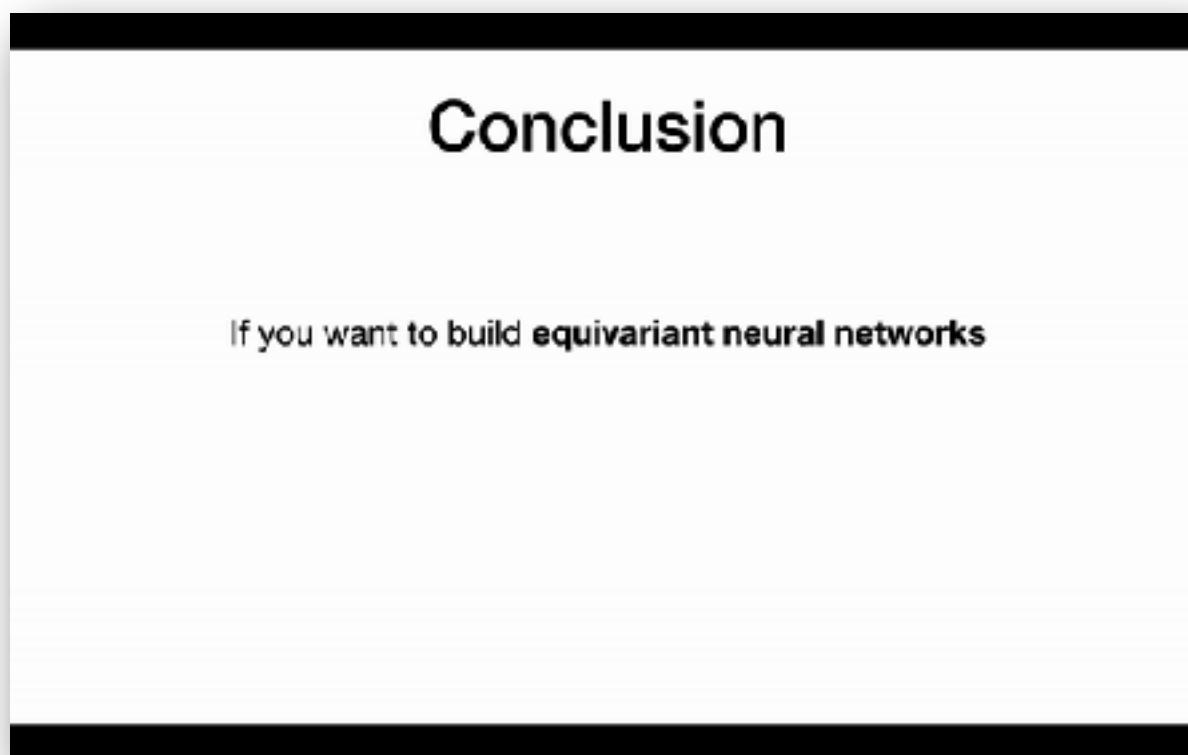
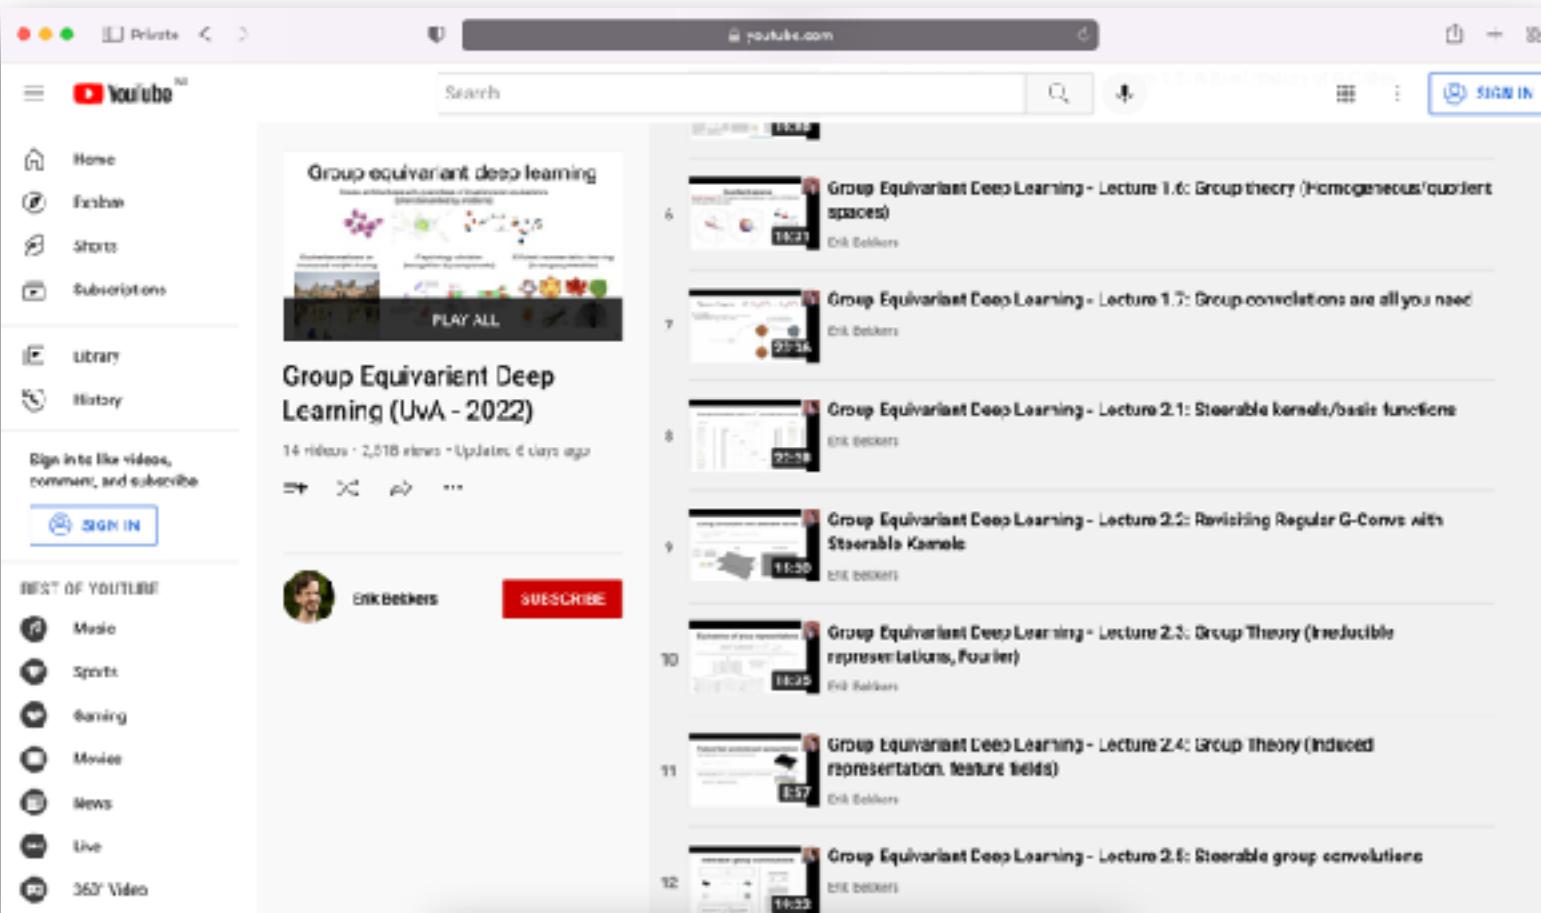
### 0. Graphs as objects embedded in Euclidean Space

In this tutorial we are going to look at graph convolution methods that can act on graphs embedded in some Euclidean space, meaning that the graph represents a some multidimensional structure, etc. we will refer to it as a **Manifold Graph**.

Let us first consider a generic graph consisting of a node set  $V$  containing nodes  $v_i$ , and a connectivity given by an edge set  $E$  consisting of

# UvA course on group equivariant deep learning (<https://uvagedl.github.io>)

## Youtube playlist



## Tutorial notebooks

[View on GitHub](#)

### GDL - Regular Group Convolutions

This notebook covers the implementation of regular group convolutional networks from scratch, using PyTorch primitives. It includes sections on introduction, brief recap on CNNs, and a detailed explanation of the convolution operation as an inner product of the function  $f$  and a shifted kernel  $k$ .

[View on GitHub](#)

### GDL - Steerable CNNs

This notebook introduces Steerable CNNs, a general framework for convolutional networks that can handle groups with infinite elements. It includes sections on introduction, prerequisite knowledge, and a detailed explanation of the Fourier transform and its application in steerable CNNs.

[View on GitHub](#)

### Introduction

The goal of this tutorial is to get you familiar with the ways in which positional information is used in graph neural networks. We will first look at superpixel CNNs, which group pixels by similarity in order to create a graph. Then, we will apply steerable methods to MD17, a 3D dataset of molecular relaxation trajectories.

This tutorial was created by Petri van der Linden and Ruò Hesselink. If you would like to know more about these methods, feel free to contact us at [r.hesselink@uva.nl](mailto:r.hesselink@uva.nl) or [p.vanderlinden@uva.nl](mailto:p.vanderlinden@uva.nl).

Below you can find several papers that we recommend, if you're interested in some more reading:

- [Neural Message Passing for Quantum Chemistry](#)
- [Directional Message Passing for Molecular Graphs](#)
- [Equivariant Graph Neural Networks for Data-Efficient and Accurate Interatomic Potentials](#)

### 0. Graphs as objects embedded in Euclidean Space

In this tutorial we are going to look at graph convolution methods that can act on graphs embedded in some Euclidean space, meaning that the graph represents a some multidimensional structure, etc. we will refer to it as a **Manifold Graph**.

Let us first consider a generic graph consisting of a node set  $V$  containing nodes  $v_i$ , and a connectivity given by an edge set  $E$  consisting of

**Thank you for listening!**